

(https://github.com/HazyResearch/deepdive)



- 🔗 Quick Start (<http://deepdive.stanford.edu/quickstart>)
- 📖 Documentation (<http://deepdive.stanford.edu/#documentation>)
- 🖼️ Showcase (<http://deepdive.stanford.edu/showcase/apps>)
- 📄 Papers (<http://deepdive.stanford.edu/papers>)
- 🗄️ Data (<http://deepdive.stanford.edu/opendata>)
- 💬 Chat (<https://gitter.im/HazyResearch/deepdive>)
- 👥 Forum (<https://groups.google.com/d/forum/deepdive-users>)



undefined
undefined
undefined
undefined
undefined
undefined
undefined

Papers

DeepDive System Overview

- *Best Reference:* DeepDive: A Data Management System for Automatic Knowledge Base Construction. (<http://pages.cs.wisc.edu/%7Ec Zhang/zhang.thesis.pdf>) Ce Zhang.Ph.D. Dissertation, University of Wisconsin-Madison, 2015.

Applications

- *Best Reference:* Machine-compiled Macroevo lutionary History of Phanerozoic Life. (<http://arxiv.org/abs/1406.2963>) Shanan E. Peters, Ce Zhang, Miron Livny, and C. Ré. PloS ONE.
- The Mobilize Center: an NIH big data to knowledge center to advance human movement research and improve mobility. (<http://jamia.oxfordjournals.org/content/early/2015/08/13/jamia.ocv071.abstract>) Joy P Ku, Jennifer L Hicks, Trevor Hastie, Jure Leskovec, C. Ré, and Scott L Delp. AMIA, 2015.
- Large-scale extraction of gene interactions from full text literature using DeepDive. Emily Mallory, Ce Zhang, C. Ré, and Russ Altman. Bioinformatics 2015.
- A Demonstration of Data Labeling in Knowledge Base Construction. (<http://cs.stanford.edu/people/chrismre/papers/p2148-shin-vldb-demo.pdf>) Jaeho Shin, Mike Cafarella, and C. Ré. VLDB Demo, 2015.
- Elementary: Large-scale Knowledge-base Construction via Machine Learning and Statistical Inference. (http://www.cs.stanford.edu/people/chrismre/papers/elementary_journal.pdf) F. Niu, Ce Zhang, C. Ré, and J. Shavlik. IJSWIS, Special Issue on Knowledge Extraction from the Web, 2012.
- Stanford’s 2014 Slot Filling Systems. (<http://i.stanford.edu/hazy/papers/2014kbp-systemdescription.pdf>) Gabor Angeli, Sonal Gupta, Melvin Jose, Christopher D. Manning, Christopher Ré, Julie Tibshirani, Jean Y. Wu, Sen Wu, Ce Zhang.
- Using Social Media to Measure Labor Market Flow. (<http://www-personal.umich.edu/%7Eshapiro/papers/LaborFlowsSocialMedia.pdf>) D. Antenucci, M. Cafarella, M. Levenstein, C. Ré, and M. Shapiro. NBER. Selected for NBER Digest (<http://www.nber.org/digest/aug14/aug14.pdf>).
- Feature Selection in Enterprise Analytics: A Demonstration using an R-based Data Analytics System (<http://cs.stanford.edu/people/chrismre/papers/ColumbusDemo-VLDB.pdf>) P. Konda, A. Kumar, C. Ré, and V. Sashikanth. VLDB Demo 2013
- GeoDeepDive: Statistical Inference using Familiar Data-Processing Languages. (http://www.cs.stanford.edu/people/chrismre/papers/gdd_demo.pdf) Ce Zhang, V. Govindaraju, J. Borchardt, T. Foltz, C. Ré, and S. Peters. SIGMOD 13 (demo).
- Hazy: Making it Easier to Build and Maintain Big-data Analytics. (<http://queue.acm.org/detail.cfm?id=2431055>) Arun Kumar, Feng Niu, and C. Ré ACM Queue, 2013. Invited to CACM March 2013
- Building an Entity-Centric Stream Filtering Test Collection for TREC 2102 (<http://trec.nist.gov/pubs/trec21/papers/KBA.OVERVIEW.pdf>). J.R. Frank, M.Kleiman-Weiner, D. A. Roberts, F.Niu, Ce Zhang, C. Ré, and I. Soboroff. TREC 2013

Feature Engineering

- *Best Reference:* Feature Engineering for Knowledge Base Construction. (<http://arxiv.org/abs/1407.6439>) DeepDive Group. Data Engineering Bulletin.
- Brainwash: A Data System for Feature Engineering. (http://www.cs.stanford.edu/people/chrismre/papers/mythical_man.pdf) M. Anderson et al. CIDR Conference 2013 (Vision Track)

- Materialization Optimizations for Feature Selection. (<http://cs.stanford.edu/people/chrismre/papers/mod539-zhang.pdf>) Ce Zhang, Arun Kumar, and C. Ré. SIGMOD 2014. Best Paper Award.
- Parallel Feature Selection Inspired by Group Testing. (<http://papers.nips.cc/paper/5296-parallel-feature-selection-inspired-by-group-testing.pdf>) Y. Zhou et al. NIPS2014.

The Underlying Engine

- *Best Overall Description*: Incremental Knowledge Base Construction Using DeepDive (<http://i.stanford.edu/hazy/papers/inc.pdf>) Sen Wu, Ce Zhang, Christopher De Sa, Jaeho Shin, Feiran Wang, and C. Ré. VLDB. 2015.
- *Best Reference for Sampler*: DimmWitted: A Study of Main-Memory Statistical Analytics. (<http://arxiv.org/abs/1403.7550>) Ce Zhang and C. Ré. VLDB 2014.
- *Best Reference for Grounding*: Tuffy: Scaling up Statistical Inference in Markov Logic Networks using an RDBMS. (<http://www.cs.stanford.edu/people/chrismre/papers/tuffy-vldb11.pdf>) F. Niu, C. Ré, A. Doan, and J.W. Shavlik. PVLDB 11.
- Rapidly Mixing Gibbs Sampling for a Class of Factor Graphs Using Hierarchy Width. Christopher De Sa, Ce Zhang, Kunle Olukotun, and C. Ré. NIPS 2015 (Spotlight).
- Taming the Wild: A Unified Analysis of Hogwild!-Style Algorithms. (<http://arxiv.org/abs/1506.06438>) Christopher De Sa, Ce Zhang, Kunle Olukotun, and C. Ré. NIPS 2015.
- Towards High-Throughput Gibbs Sampling at Scale: A Study across Storage Managers. (http://www.cs.stanford.edu/people/chrismre/papers/elementary_sigmod.pdf) Ce Zhang and C. Ré. SIGMOD 2013
- An Asynchronous Parallel Stochastic Coordinate Descent Algorithm. (<http://arxiv.org/abs/1311.1873>) J. Liu, S. Wright, C. Ré, V. Bittorf, S. Sridhar. ICML 2014.
- Hogwild!: A Lock-Free Approach to Parallelizing Stochastic Gradient Descent. (<http://i.stanford.edu/hazy/papers/hogwild-nips.pdf>) F. Niu, B. Recht, C. Ré, and S. J. Wright. NIPS, 2011
- Felix: Scaling Inference for Markov Logic with an Operator-based Approach. (<http://arxiv.org/pdf/1108.0294v1>) Feng Niu, Ce Zhang, C. Ré, and Jude Shavlik.

Some Related Techniques

- Understanding Tables in Context Using Standard NLP Toolkits. (<http://cs.stanford.edu/people/chrismre/papers/jointable-acl.pdf>) Vidhya Govindaraju, Ce Zhang, and C. Ré. ACL 2013
- Probabilistic Management of OCR using an RDBMS. (http://www.cs.stanford.edu/people/chrismre/papers/HazyOCR_VLDB2012.pdf) Arun Kumar and C. Ré. PVLDB 2012.
- Optimizing Statistical Information Extraction Programs Over Evolving Text. (<http://www.cs.stanford.edu/people/chrismre/papers/crflex-icde12.pdf>) Fei Chen, Xixuan Feng, C. Ré, and Min Wang. ICDE
- Incrementally maintaining classification using an RDBMS. (<http://www.cs.stanford.edu/people/chrismre/papers/hazy-classification-vldb11.pdf>) Mehmet Levent Koc and C. Ré. PVLDB Volume 4, 2011, p. 302-313

Studies

- Big Data versus the Crowd: Looking for Relationships in All the Right Places. (<http://www.cs.stanford.edu/people/chrismre/papers/acl2012.pdf>) Ce Zhang, Feng Niu, C. Ré, and Jude Shavlik. ACL, 2012.

Formal Foundations

- Probabilistic Databases. (<http://www.morganclaypool.com/doi/pdfplus/10.2200/S00362ED1V01Y201105DTM016>) Dan Suciu, Dan Olteanu, C. Ré, and Christoph Koch. Morgan Claypool's Synthesis Lectures, 2011
- Transducing Markov Sequences. (<http://dl.acm.org/citation.cfm?doid=2630065>) Benny Kimelfeld and C. Ré. JACM 2014.
- Queries and materialized views on probabilistic databases. (<http://dx.doi.org/10.1016/j.jcss.2010.04.006>) Nilesh N. Dalvi, C. Re, and Dan Suciu. JCSS 2011.

Hazy Research Group, InfoLab (<http://infolab.stanford.edu>), Computer Science Department, Stanford University (<http://cs.stanford.edu>).

Questions? Chat with us (<https://gitter.im/HazyResearch/deepdive>) or ask deepdive-users (<https://groups.google.com/d/forum/deepdive-users>)!
Fork DeepDive on Github (<https://github.com/HazyResearch/deepdive>) and join deepdive-dev (<https://groups.google.com/d/forum/deepdive-dev>).