

Heaton Research

Overview of Keras/TensorFlow Basic Operations

📅 2017-07-22 (/2017/07/22/keras-getting-started.html)

📁 ai (/categories/ai/)

I in the process of updating my deep learning course (<https://sites.wustl.edu/jeffheaton/t81-558/>) and books (<http://www.aifh.org>) to make use of Keras (<https://keras.io/>). This posting contains some of the basic examples that I put together. This post is not meant to be an introduction to neural networks in general. For such an introduction, refer to either my books (/book/) or this article (<http://www.heatonresearch.com/content/non-mathematical-introduction-using-neural-networks>).

I will expand on these examples greatly for both the book and course. The basic neural network operations that I needed were:

- Simple Regression
- Regression Early Stopping
- Simple Classification

About

Jeff Heaton, Ph.D. is a computer scientist, data scientist, and indie publisher. Heaton Research is the homepage for his projects and research.

Categories

ai (/categories/ai/) (4)
aifh (/categories/aifh/) (1)
datascience
(/categories/datascience/)
(10)
encog (/categories/encog/)
(1)
gpu (/categories/gpu/) (2)
kaggle
(/categories/kaggle/) (1)
learning
(/categories/learning/) (2)
phd (/categories/phd/) (7)
python
(/categories/python/) (1)
r (/categories/r/) (4)
tensorflow
(/categories/tensorflow/)

- Classification Early Stopping
- Deep Neural Networks w/Dropout and Other Regularization
- Convolutional Neural Networks
- LSTM Neural Networks
- Loading/Saving Neural Networks

These are some of the most basic operations that I need to perform when working with a new neural network package. This provides me with a sort of Rosetta Stone for a new neural network package. Once I have these operations, I can more easily create additional examples that are more complex.

The first thing to check is what versions you have of the required packages:

```
1 import keras
2 import tensorflow as tf
3 import sys
4 import sklearn as sk
5 import pandas as pd
6
7 print("Tensor Flow Version: {}".format(tf.__version__))
8 print("Keras Version: {}".format(keras.__version__))
9 print()
10 print("Python {}".format(sys.version))
11 print('Pandas {}'.format(pd.__version__))
12 print('Scikit-Learn {}'.format(sk.__version__))
```

(2)

Archives

November 2017

(/archives/2017/11/) (3)

September 2017

(/archives/2017/09/) (1)

August 2017

(/archives/2017/08/) (1)

July 2017

(/archives/2017/07/) (4)

June 2017

(/archives/2017/06/) (1)

May 2017

(/archives/2017/05/) (1)

March 2017

(/archives/2017/03/) (1)

February 2017

(/archives/2017/02/) (1)

January 2017

(/archives/2017/01/) (1)

September 2016

(/archives/2016/09/) (1)

March 2016

(/archives/2016/03/) (1)

February 2016

(/archives/2016/02/) (1)

September 2015

(/archives/2015/09/) (1)

August 2015

(/archives/2015/08/) (1)

May 2015

(/archives/2015/05/) (1)

```

1  Tensor Flow Version: 1.0.0
2  Keras Version: 2.0.6
3
4  Python 3.5.2 |Continuum Analytics, Inc.| (default, Jul
5  Pandas 0.19.2
6  Scikit-Learn 0.18.1

```

The following functions are from my set of helpful functions (https://github.com/jeffheaton/t81_558_deep_learning/blob/master) that I created for my class and use in many of my books (<http://www.aifh.org>):

```

1  import pandas as pd
2  from sklearn import preprocessing
3
4  # Encode text values to dummy variables(i.e. [1,0,0],[0,1,0],[0,0,1])
5  def encode_text_dummy(df, name):
6      dummies = pd.get_dummies(df[name])
7      for x in dummies.columns:
8          dummy_name = "{}-{}".format(name, x)
9          df[dummy_name] = dummies[x]
10     df.drop(name, axis=1, inplace=True)
11
12  # Encode text values to indexes(i.e. [1],[2],[3] for red, blue, green)
13  def encode_text_index(df, name):
14      le = preprocessing.LabelEncoder()
15      df[name] = le.fit_transform(df[name])
16      return le.classes_
17

```

March 2015
(/archives/2015/03/) (4)
December 2014
(/archives/2014/12/) (1)
September 2014
(/archives/2014/09/) (1)
May 2014
(/archives/2014/05/) (2)
February 2014
(/archives/2014/02/) (1)
July 2013
(/archives/2013/07/) (2)
June 2013
(/archives/2013/06/) (2)
April 2013
(/archives/2013/04/) (1)
March 2013
(/archives/2013/03/) (1)

Recents

AWS EC2 Data Science:
My Jupyter Workspaces for
GPU (/2017/11/30/aws-
data-science-gpu.html)
AWS EC2 Data Science:
My Jupyter Workspaces for
Docker CPU Only
(/2017/11/25/aws-data-
science-cpu.html)
AWS EC2 Data Science:
My Jupyter Workspaces
(/2017/11/15/aws-data-
science.html)

```
18 # Convert all missing values in the specified column to
19 def missing_median(df, name):
20     med = df[name].median()
21     df[name] = df[name].fillna(med)
22
23 # Convert a Pandas dataframe to the x,y inputs that TensorFlow
24 def to_xy(df, target):
25     result = []
26     for x in df.columns:
27         if x != target:
28             result.append(x)
29
30     # find out the type of the target column. Is it regression or classification?
31     target_type = df[target].dtypes
32     target_type = target_type[0] if hasattr(target_type, 'item') else target_type
33
34     # Encode to int for classification, float otherwise
35     if target_type in (np.int64, np.int32):
36         # Classification
37         dummies = pd.get_dummies(df[target])
38         return df.as_matrix(result).astype(np.float32), dummies.as_matrix().astype(np.float32)
39     else:
40         # Regression
41         return df.as_matrix(result).astype(np.float32), df[target].as_matrix().astype(np.float32)
42
```

Installing TensorFlow
(/2017/09/14/install_tf.html)
Data Science Rosetta
Stone: Classification in
Python, R, MATLAB, SAS,
& Julia
(/2017/08/17/ds_rosetta_stone.html)

Simple Regression

Regression is where a neural network accepts several values (predictors) and produces a prediction that is numeric. In this simple example we attempt to predict the miles per gallon (MPG) of several cars based on characteristics of those cars. Several parameters are used below and described here.

- Losses Supported by Keras (<https://keras.io/losses/>)
 - Typically use **mean_squared_error** for regression (the square root of mean square error is root mean square error(RMSE)).
 - and for classification use: **binary_crossentropy** for 2 classes, **categorical_crossentropy** for more than 2 classes.
- kernel_initializer supported by Keras (<https://keras.io/initializers/>) - Specifies how the weights of are randomized.
- activation (<https://keras.io/activations/>) - Usually relu or softmax will be used.

```
import Sequential
from import Dense, Activation

metrics

thubusercontent.com/jeffheaton/t81_558_deep_learning/master/c
tringIO(requests.get(url).content.decode('utf-8')),na_values=

place=True)
'horsepower')
")

)
input_dim=x.shape[1], activation='relu'))
kernel_initializer='normal'))
'mean_squared_error', optimizer='adam')

se=2, epochs=100)
```



```
1 Epoch 1/100
2 0s - loss: 2240.8034
3 Epoch 2/100
4 0s - loss: 1469.8520
5 Epoch 3/100
6 0s - loss: 1038.2052
7 Epoch 4/100
8 0s - loss: 820.4976
9 Epoch 5/100
10
11 ...
12
13 0s - loss: 560.3524
14 Epoch 99/100
15 0s - loss: 559.7951
16 Epoch 100/100
17 0s - loss: 559.2341
18 <keras.callbacks.History at 0x2263d8cc518>
```

Now that the neural network is trained, we will test how good it is and perform some sample predictions.

```
1 pred = model.predict(x)
2
3 # Measure RMSE error.  RMSE is common for regression.
4 score = np.sqrt(metrics.mean_squared_error(pred,y))
5 print("Final score (RMSE): {}".format(score))
6
7 # Sample predictions
8 for i in range(10):
9     print("{} Car name: {}, MPG: {}, predicted MPG: {}'
```

```
1 Final score (RMSE): 3.494112253189087
2 1. Car name: chevrolet chevelle malibu, MPG: [ 18.], pr
3 2. Car name: buick skylark 320, MPG: [ 15.], predicted
4 3. Car name: plymouth satellite, MPG: [ 18.], predicted
5 4. Car name: amc rebel sst, MPG: [ 16.], predicted MPG
6 5. Car name: ford torino, MPG: [ 17.], predicted MPG:
7 6. Car name: ford galaxie 500, MPG: [ 15.], predicted
8 7. Car name: chevrolet impala, MPG: [ 14.], predicted
9 8. Car name: plymouth fury iii, MPG: [ 14.], predicted
10 9. Car name: pontiac catalina, MPG: [ 14.], predicted
11 10. Car name: amc ambassador dpl, MPG: [ 15.], predicte
```

Regression (Early Stop)

Early stopping sets aside a part of the data to be used to validate the neural

network. The neural network is trained with the training data and

validated
with the validation data. Once the error no longer improves on
the validation
set, the training stops. This prevents the neural network from
overfitting (<https://en.wikipedia.org/wiki/Overfitting>).

```
1 import pandas as pd
2 import io
3 import requests
4 import numpy as np
5 from sklearn import metrics
6 from keras.models import Sequential
7 from keras.layers.core import Dense, Activation
8 from keras.callbacks import EarlyStopping
9
10 url="https://raw.githubusercontent.com/jeffheaton/t81_!
11 df=pd.read_csv(io.StringIO(requests.get(url).content.de
12
13 cars = df['name']
14 df.drop('name',1,inplace=True)
15 missing_median(df, 'horsepower')
16 x,y = to_xy(df,"mpg")
17
18 # Split into train/test
19 x_train, x_test, y_train, y_test = train_test_split(
20     x, y, test_size=0.25, random_state=45)
21
22 model = Sequential()
23 model.add(Dense(10, input_dim=x.shape[1], kernel_initia
24 model.add(Dense(1, kernel_initializer='normal'))
25 model.compile(loss='mean_squared_error', optimizer='ad
26
27 monitor = EarlyStopping(monitor='val_loss', min_delta=
28
29 model.fit(x,y,validation_data=(x_test,y_test),callback:
```

```
1 Train on 398 samples, validate on 100 samples
2 Epoch 1/1000
3 0s - loss: 374.7638 - val_loss: 179.2396
4 Epoch 2/1000
5 0s - loss: 199.9990 - val_loss: 169.4834
6 Epoch 3/1000
7 0s - loss: 197.9431 - val_loss: 153.8338
8 Epoch 4/1000
9 0s - loss: 187.7644 - val_loss: 152.2758
10 Epoch 5/1000
11 0s - loss: 185.5505 - val_loss: 149.9817
12
13 ...
14
15 Epoch 179/1000
16 0s - loss: 10.3191 - val_loss: 8.2763
17 Epoch 180/1000
18 0s - loss: 10.0629 - val_loss: 8.3435
19 Epoch 181/1000
20 0s - loss: 10.7124 - val_loss: 8.4712
21 Epoch 182/1000
22 0s - loss: 10.6406 - val_loss: 8.4272
23 <keras.callbacks.History at 0x222a8ecd1d0>
```

Classification Model (Early Stop)

Early stopping can also be used with classification. Early stopping sets aside a part of the data to be used to validate the neural network. The neural network is trained with the training

data and validated with the validation data. Once the error no longer improves on the validation set, the training stops. This prevents the neural network from overfitting (<https://en.wikipedia.org/wiki/Overfitting>).

```
1 import pandas as pd
2 import io
3 import requests
4 import numpy as np
5 from sklearn import metrics
6 from keras.models import Sequential
7 from keras.layers.core import Dense, Activation
8 from keras.callbacks import EarlyStopping
9
10 url="https://raw.githubusercontent.com/jeffheaton/t81_!
11 df=pd.read_csv(io.StringIO(requests.get(url).content.de
12
13 species = encode_text_index(df,"species")
14 x,y = to_xy(df,"species")
15
16 # Split into train/test
17 x_train, x_test, y_train, y_test = train_test_split(
18     x, y, test_size=0.25, random_state=45)
19
20 model = Sequential()
21 model.add(Dense(10, input_dim=x.shape[1], kernel_initia
22 model.add(Dense(1, kernel_initializer='normal'))
23 model.add(Dense(y.shape[1],activation='softmax'))
24 model.compile(loss='categorical_crossentropy', optimiz
25
26 monitor = EarlyStopping(monitor='val_loss', min_delta=
27
28 model.fit(x,y,validation_data=(x_test,y_test),callback:
```

```
1  Train on 150 samples, validate on 38 samples
2  Epoch 1/1000
3  0s - loss: 1.1095 - val_loss: 1.1143
4  Epoch 2/1000
5  0s - loss: 1.1065 - val_loss: 1.1096
6  Epoch 3/1000
7  0s - loss: 1.1041 - val_loss: 1.1057
8  Epoch 4/1000
9  0s - loss: 1.1020 - val_loss: 1.1038
10 Epoch 5/1000
11 0s - loss: 1.1011 - val_loss: 1.1017
12
13 ...
14
15 Epoch 325/1000
16 0s - loss: 0.1758 - val_loss: 0.1320
17 Epoch 326/1000
18 0s - loss: 0.1755 - val_loss: 0.1332
19 Epoch 00325: early stopping
20 <keras.callbacks.History at 0x222a9242d68>
```

Show the predictions (raw, probability of each class.)

```
1  # Print out the raw predictions. Because there are 3 species
2  # the probability that the flower is that type of iris.
3
4  np.set_printoptions(suppress=True)
5  pred = model.predict(x_test)
6  print(pred[0:10])
```

```
1  [[ 0.97540218  0.0245978  0.          ]
2   [ 0.94149685  0.05850318  0.          ]
3   [ 0.02133332  0.27963796  0.69902873]
4   [ 0.94382465  0.05617536  0.          ]
5   [ 0.95254719  0.04745276  0.          ]
6   [ 0.95966363  0.04033642  0.          ]
7   [ 0.94291645  0.05708356  0.          ]
8   [ 0.00293462  0.06093681  0.93612856]
9   [ 0.00873046  0.14257514  0.84869444]
10  [ 0.00293431  0.0609317   0.93613404]]
```

```
1  # The to_xy function represented the input in the same way
2  # of iris. This is the training data, we KNOW what type
3  # is 1.0 (hot)
4
5  print(y_test[0:10])
```

```
1  [[ 1.  0.  0.]
2   [ 1.  0.  0.]
3   [ 0.  0.  1.]
4   [ 1.  0.  0.]
5   [ 1.  0.  0.]
6   [ 1.  0.  0.]
7   [ 1.  0.  0.]
8   [ 0.  0.  1.]
9   [ 0.  0.  1.]
10  [ 0.  0.  1.]]
```



```
1 # Of course it is very easy to turn these indexes back :
2
3 print(species[predict_classes[1:10]])
```

```
1 ['Iris-setosa' 'Iris-virginica' 'Iris-setosa' 'Iris-seto
2 'Iris-setosa' 'Iris-virginica' 'Iris-virginica' 'Iris-\\
```

```
1 from sklearn.metrics import accuracy_score
2
3 # Accuracy might be a more easily understood error metr
4 # what percent were correct? The downside is it does n
5
6 correct = accuracy_score(expected_classes,predict_class
7 print("Accuracy: {}".format(correct))
```

```
1 Accuracy: 1.0
```

Deeper Networks

Keras makes it easy to add additional layers as shown here:

```
1 import pandas as pd
2 import io
3 import requests
4 import numpy as np
5 from sklearn import metrics
```

```
6 from keras.callbacks import EarlyStopping
7 from keras.layers import Dense, Dropout
8 from keras import regularizers
9
10 url="https://raw.githubusercontent.com/jeffheaton/t81_!
11 df=pd.read_csv(io.StringIO(requests.get(url).content.d
12
13 cars = df['name']
14 df.drop('name',1,inplace=True)
15 missing_median(df, 'horsepower')
16 x,y = to_xy(df,"mpg")
17
18 # Split into train/test
19 x_train, x_test, y_train, y_test = train_test_split(
20     x, y, test_size=0.25, random_state=45)
21
22
23 model = Sequential()
24 model.add(Dense(50, input_dim=x.shape[1], kernel_initia
25 model.add(Dropout(0.2))
26 model.add(Dense(25, input_dim=x.shape[1], kernel_initia
27 model.add(Dense(10, input_dim=64,
28             kernel_regularizer=regularizers.l2(0.0:
29             activity_regularizer=regularizers.l1(0
30 model.add(Dense(1, kernel_initializer='normal'))
31 model.compile(loss='mean_squared_error', optimizer='ad
32
33 monitor = EarlyStopping(monitor='val_loss', min_delta=
34
35 model.fit(x,y,validation_data=(x_test,y_test),callback:
36 pred = model.predict(x_test)
37
```

```
38 # Measure RMSE error. RMSE is common for regression.
39 score = np.sqrt(metrics.mean_squared_error(pred,y_test)
40 print("Final score (RMSE): {}".format(score))
```

```
1 Epoch 00064: early stopping
2 Final score (RMSE): 4.421816825866699
```

The Classic MNIST Dataset


The next examples will use the MNIST digits dataset (<http://yann.lecun.com/exdb/mnist/>). The previous examples used CSV files to load training data. Most neural network frameworks, such as Keras, have common training sets built in. This makes it easy to run the example, but hard to abstract the example to your own data. Your own data are not likely built into Keras. However, the MNIST data is complex enough that it is beyond the scope of this article to discuss how to load it. We will use the MNIST data built into Keras.

```
1 from keras.datasets import mnist
2 (x_train, y_train), (x_test, y_test) = mnist.load_data()
3
4 print("Shape of x_train: {}".format(x_train.shape))
5 print("Shape of y_train: {}".format(y_train.shape))
6 print()
7 print("Shape of x_test: {}".format(x_test.shape))
8 print("Shape of y_test: {}".format(y_test.shape))
```

```
1 Shape of x_train: (60000, 28, 28)
2 Shape of y_train: (60000,)
3
4 Shape of x_test: (10000, 28, 28)
5 Shape of y_test: (10000,)
```

```
1 # Display as image
2 %matplotlib inline
3 import matplotlib.pyplot as plt
4 import numpy as np
5
6 digit = 101 # Change to choose new digit
7
8 a = x_train[digit]
9 plt.imshow(a, cmap='gray', interpolation='nearest')
10 print("Image ({}) : Which is digit '{}'".format(digit, y_train[digit]))
```

```
1 Image (#101): Which is digit '7'
```




Convolutional Neural Networks

Convolutional Neural Networks are specifically for images. They have been applied to other cases; however, use beyond images is somewhat rarer than with images.

```
1 import keras
2 from keras.datasets import mnist
3 from keras.models import Sequential
4 from keras.layers import Dense, Dropout, Flatten
5 from keras.layers import Conv2D, MaxPooling2D
6 from keras import backend as K
7
8 batch_size = 128
9 num_classes = 10
10 epochs = 12
11
12 # input image dimensions
```

```
13 img_rows, img_cols = 28, 28
14
15 if K.image_data_format() == 'channels_first':
16     x_train = x_train.reshape(x_train.shape[0], 1, img_
17     x_test = x_test.reshape(x_test.shape[0], 1, img_row
18     input_shape = (1, img_rows, img_cols)
19 else:
20     x_train = x_train.reshape(x_train.shape[0], img_row
21     x_test = x_test.reshape(x_test.shape[0], img_rows,
22     input_shape = (img_rows, img_cols, 1)
23
24 x_train = x_train.astype('float32')
25 x_test = x_test.astype('float32')
26 x_train /= 255
27 x_test /= 255
28 print('x_train shape:', x_train.shape)
29 print("Training samples: {}".format(x_train.shape[0]))
30 print("Test samples: {}".format(x_test.shape[0]))
31
32 # convert class vectors to binary class matrices
33 y_train = keras.utils.to_categorical(y_train, num_clas:
34 y_test = keras.utils.to_categorical(y_test, num_classe:
35
36 model = Sequential()
37 model.add(Conv2D(32, kernel_size=(3, 3),
38                 activation='relu',
39                 input_shape=input_shape))
40 model.add(Conv2D(64, (3, 3), activation='relu'))
41 model.add(MaxPooling2D(pool_size=(2, 2)))
42 model.add(Dropout(0.25))
43 model.add(Flatten())
44 model.add(Dense(128, activation='relu'))
```

```
45 model.add(Dropout(0.5))
46 model.add(Dense(num_classes, activation='softmax'))
47
48 model.compile(loss=keras.losses.categorical_crossentropy,
49               optimizer=keras.optimizers.Adadelta(),
50               metrics=['accuracy'])
51
52 model.fit(x_train, y_train,
53         batch_size=batch_size,
54         epochs=epochs,
55         verbose=2,
56         validation_data=(x_test, y_test))
57 score = model.evaluate(x_test, y_test, verbose=0)
58 print('Test loss: {}'.format(score[0]))
59 print('Test accuracy: {}'.format(score[1]))
```



```
1 x_train shape: (60000, 28, 28, 1)
2 Training samples: 60000
3 Test samples: 10000
4 Train on 60000 samples, validate on 10000 samples
5 Epoch 1/12
6 271s - loss: 0.3435 - acc: 0.8950 - val_loss: 0.0817 -
7 Epoch 2/12
8 269s - loss: 0.1171 - acc: 0.9660 - val_loss: 0.0581 -
9 Epoch 3/12
10 458s - loss: 0.0885 - acc: 0.9742 - val_loss: 0.0453 -
11 Epoch 4/12
12 554s - loss: 0.0743 - acc: 0.9778 - val_loss: 0.0382 -
13 Epoch 5/12
14 261s - loss: 0.0642 - acc: 0.9810 - val_loss: 0.0346 -
15 Epoch 6/12
16 321s - loss: 0.0594 - acc: 0.9826 - val_loss: 0.0337 -
17 Epoch 7/12
18 309s - loss: 0.0515 - acc: 0.9846 - val_loss: 0.0335 -
19 Epoch 8/12
20 317s - loss: 0.0477 - acc: 0.9857 - val_loss: 0.0337 -
21 Epoch 9/12
22 308s - loss: 0.0448 - acc: 0.9870 - val_loss: 0.0330 -
23 Epoch 10/12
24 322s - loss: 0.0416 - acc: 0.9873 - val_loss: 0.0307 -
25 Epoch 11/12
26 326s - loss: 0.0394 - acc: 0.9879 - val_loss: 0.0300 -
27 Epoch 12/12
28 313s - loss: 0.0367 - acc: 0.9887 - val_loss: 0.0313 -
29 Test loss: 0.03131893762472173
30 Test accuracy: 0.9902
```


Long Short Term Memory (LSTM)

Long Short Term Memory is typically used for either time series or natural language processing (which can be thought of as a special case of natural language processing).

```
1  from keras.preprocessing import sequence
2  from keras.models import Sequential
3  from keras.layers import Dense, Embedding
4  from keras.layers import LSTM
5  from keras.datasets import imdb
6  import numpy as np
7
8  max_features = 4 # 0,1,2,3 (total of 4)
9  x = [
10     [[0],[1],[1],[0],[0],[0]],
11     [[0],[0],[0],[2],[2],[0]],
12     [[0],[0],[0],[0],[3],[3]],
13     [[0],[2],[2],[0],[0],[0]],
14     [[0],[0],[3],[3],[0],[0]],
15     [[0],[0],[0],[0],[1],[1]]
16 ]
17 x = np.array(x,dtype=np.float32)
18 y = np.array([1,2,3,2,3,1],dtype=np.int32)
19
20 # Convert y2 to dummy variables
21 y2 = np.zeros((y.shape[0], max_features),dtype=np.float32)
22 y2[np.arange(y.shape[0]), y] = 1.0
23 print(y2)
24
25 print('Build model...')
```

```
26 model = Sequential()
27 model.add(LSTM(128, dropout=0.2, recurrent_dropout=0.2,
28 model.add(Dense(4, activation='sigmoid'))
29
30 # try using different optimizers and different optimizer
31 model.compile(loss='binary_crossentropy',
32               optimizer='adam',
33               metrics=['accuracy'])
34
35 print('Train...')
36 model.fit(x,y2,epochs=200)
37 pred = model.predict(x)
38 predict_classes = np.argmax(pred,axis=1)
39 print("Predicted classes: {}".format(predict_classes))
40 print("Expected classes: {}".format(predict_classes))
```

```
1  [[ 0.  1.  0.  0.]
2   [ 0.  0.  1.  0.]
3   [ 0.  0.  0.  1.]
4   [ 0.  0.  1.  0.]
5   [ 0.  0.  0.  1.]
6   [ 0.  1.  0.  0.]]
7  Build model...
8
9
10 c:\users\jeffh\anaconda3\envs\tf-latest\lib\site-packages\keras\
11 c:\users\jeffh\anaconda3\envs\tf-latest\lib\site-packages\keras\
12
13
14 Train...
```

```
15 Epoch 1/200
16 6/6 [=====] - 2s - loss: 0.70
17 Epoch 2/200
18 6/6 [=====] - 0s - loss: 0.70
19 Epoch 3/200
20 6/6 [=====] - 0s - loss: 0.68
21 Epoch 4/200
22 6/6 [=====] - 0s - loss: 0.68
23 Epoch 5/200
24 6/6 [=====] - 0s - loss: 0.67
25
26 ...
27
28 Epoch 198/200
29 6/6 [=====] - 0s - loss: 0.22
30 Epoch 199/200
31 6/6 [=====] - 0s - loss: 0.29
32 Epoch 200/200
33 6/6 [=====] - 0s - loss: 0.19
34 Predicted classes: {} [1 2 3 2 3 1]
35 Expected classes: {} [1 2 3 2 3 1]
```

Load/Save a Neural Network

It is very important to be able to load and save neural networks. This allows your neural network to be used each time without retraining.

```
1 import pandas as pd
2 import io
3 import requests
4 import numpy as np
5 from sklearn import metrics
6 from keras.models import Sequential
7 from keras.layers.core import Dense, Activation
8 from keras.models import load_model
9
10 url="https://raw.githubusercontent.com/jeffheaton/t81_!
11 df=pd.read_csv(io.StringIO(requests.get(url).content.de
12
13 cars = df['name']
14 df.drop('name',1,inplace=True)
15 missing_median(df, 'horsepower')
16 x,y = to_xy(df,"mpg")
17
18 model = Sequential()
19 model.add(Dense(10, input_dim=x.shape[1], kernel_initi:
20 model.add(Dense(1, kernel_initializer='normal'))
21 model.compile(loss='mean_squared_error', optimizer='ad:
22
23 model.fit(x,y,verbose=2,epochs=100)
```

```
1 Epoch 1/100
2 0s - loss: 188.3123
3 Epoch 2/100
4 0s - loss: 180.3333
5 Epoch 3/100
6 0s - loss: 177.1118
7 Epoch 4/100
8 0s - loss: 173.3682
9 Epoch 5/100
10 0s - loss: 167.0144
11
12 ...
13
14 Epoch 98/100
15 0s - loss: 11.2297
16 Epoch 99/100
17 0s - loss: 11.0280
18 Epoch 100/100
19 0s - loss: 10.9314
```

```
1 pred = model.predict(x)
2
3 # Measure RMSE error. RMSE is common for regression.
4 score = np.sqrt(metrics.mean_squared_error(pred,y))
5 print("Before save score (RMSE): {}".format(score))
6
7 # save neural network structure to JSON (no weights)
8 model_json = model.to_json()
9 with open("network.json", "w") as json_file:
10     json_file.write(model_json)
11
12 # save neural network structure to YAML (no weights)
13 model_yaml = model.to_yaml()
14 with open("network.yaml", "w") as yaml_file:
15     yaml_file.write(model_yaml)
16
17 # save entire network to HDF5 (save everything, suggested)
18 model.save("network.h5")
```

```
1 Before save score (RMSE): 3.276093006134033
```

```
1 from keras.models import load_model
2
3 model2 = load_model('network.h5')
4
5 # Measure RMSE error. RMSE is common for regression.
6 score = np.sqrt(metrics.mean_squared_error(pred,y))
7 print("After load score (RMSE): {}".format(score))
```