

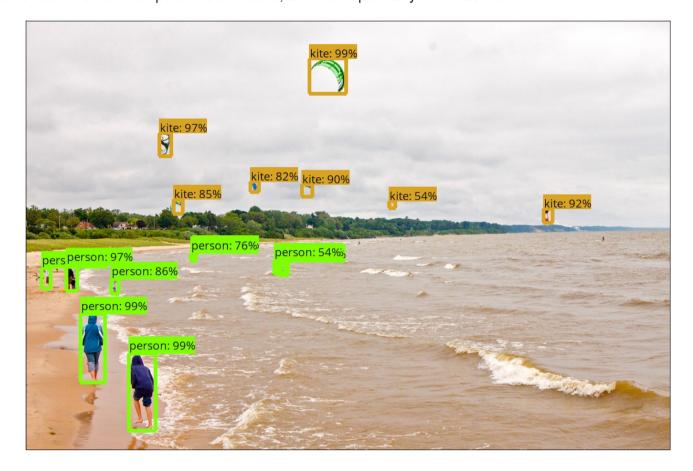
README.md

Tensorflow Object Detection API

Creating accurate machine learning models capable of localizing and identifying multiple objects in a single image remains a

1 of 3 2017年07月10日 14:48

core challenge in computer vision. The TensorFlow Object Detection API is an open source framework built on top of TensorFlow that makes it easy to construct, train and deploy object detection models. At Google we've certainly found this codebase to be useful for our computer vision needs, and we hope that you will as well.



Contributions to the codebase are welcome and we would love to hear back from you if you find this API useful. Finally if you use the Tensorflow Object Detection API for a research publication, please consider citing:

"Speed/accuracy trade-offs for modern convolutional object detectors."
Huang J, Rathod V, Sun C, Zhu M, Korattikara A, Fathi A, Fischer I, Wojna Z,
Song Y, Guadarrama S, Murphy K, CVPR 2017

[link][bibtex]

Maintainers

• Jonathan Huang, github: jch1

Vivek Rathod, github: tombstone

• Derek Chow, github: derekjchow

• Chen Sun, github: jesu9

• Menglong Zhu, github: dreamdragon

Table of contents

Before You Start:

Installation

Quick Start:

Quick Start: Jupyter notebook for off-the-shelf inference

Quick Start: Training a pet detector

Setup:

- Configuring an object detection pipeline
- Preparing inputs

Running:

- Running locally
- Running on the cloud

Extras:

• Tensorflow detection model zoo

- Exporting a trained model for inference
- Defining your own model architecture
- Bringing in your own dataset

Getting Help

Please report bugs to the tensorflow/models/ Github issue tracker, prefixing the issue name with "object_detection". To get help with issues you may encounter using the Tensorflow Object Detection API, create a new question on StackOverflow with the tags "tensorflow" and "object-detection".

Release information

June 15, 2017

In addition to our base Tensorflow detection model definitions, this release includes:

- A selection of trainable detection models, including:
 - Single Shot Multibox Detector (SSD) with MobileNet,
 - SSD with Inception V2,
 - o Region-Based Fully Convolutional Networks (R-FCN) with Resnet 101,
 - o Faster RCNN with Resnet 101,
 - o Faster RCNN with Inception Resnet v2
- Frozen weights (trained on the COCO dataset) for each of the above models to be used for out-of-the-box inference purposes.
- A Jupyter notebook for performing out-of-the-box inference with one of our released models
- Convenient local training scripts as well as distributed training and evaluation pipelines via Google Cloud.

Thanks to contributors: Jonathan Huang, Vivek Rathod, Derek Chow, Chen Sun, Menglong Zhu, Matthew Tang, Anoop Korattikara, Alireza Fathi, Ian Fischer, Zbigniew Wojna, Yang Song, Sergio Guadarrama, Jasper Uijlings, Viacheslav Kovalevskyi, Kevin Murphy

© 2017 GitHub, Inc. Terms Privacy Security Status Help

Contact GitHub API Training Shop Blog About