

luuuyi的博客


计算机视觉小硕

目录视图

摘要视图

RSS 订阅

个人资料



luuuyi



访问：45443次

积分：699

等级：BLOG > 3

排名：千里之外

原创：22篇 转载：0篇

译文：0篇 评论：12条

文章搜索

异步赠书：9月重磅新书升级，本本经典

程序员9月书讯

每周荐书：ES6、虚拟现实、物联网（评论送书）

Ubuntu上用caffe的SSD方法训练umdfaces数据集

标签：caffe 深度学习 ssd ubuntu

2017-03-13 22:42

1866人阅读

评论

分类：caffe (4) SSD (2)

版权声明：本文为博主原创文章，未经博主允许不得转载。

目录(?)

[+]

实验目的

继前一段时间用SSD训练过VOC数据集以后，这一次使用SSD+K80服务器来训练自己的人脸识别应用，选择的数据集还是之前下载的umdfaces，总共36w张人脸图像。

实验环境

训练平台：NVIDIA K80

预测平台：NVIDIA TX1

语言：C++，Python

框架：caffe

方法：SSD

数据集：umdfaces

实验准备

关闭

其实训练的步骤在上一篇博客中写得已经很清楚了，这一次主要关注的是数据集的处理，最方便的方法就是将自己的数据集也做成VOC数据集格式，也即是三个文件夹路径的结构：

名称	修改日期	类型	大小
Annotations	2007/11/7 1:46	文件夹	
ImageSets	2007/11/7 1:40	文件夹	
JPEGImages	2007/11/7 1:45	文件夹	

关于数据集的处理，在这边博客中写得很清楚：

使用faster rcnn训练umdfaces数据集

数据集处理好之后，和上一篇博客不同的主要是训练脚本（Python）的写法，这里贴出我的文件，在后面有地方都是需要自己去配置的地方：

[python]

```
01.  #!/usr/bin/python
02.
03.  from __future__ import print_function
04.  import caffe
05.  from caffe.model_libs import *
06.  from google.protobuf import text_format
07.
08.  import math
09.  import os
10.  import shutil
11.  import stat
12.  import subprocess
13.  import sys
14.
15.  # Add extra layers on top of a "base" network (e.g. VGGNet or Inception).
16.  def AddExtraLayers(net, use_batchnorm=True, lr_mult=1):
17.      use_relu = True
18.
19.      # Add additional convolutional layers.
20.      # 19 x 19
21.      from_layer = net.keys()[-1]
22.
```

文章分类

caffe (5)
Zbar (1)
感悟 (1)
Python (2)
SSD (3)
C++ (9)
机器学习 (3)
练习 (9)

文章存档

2017年07月 (2)
2017年06月 (4)
2017年05月 (3)
2017年04月 (1)
2017年03月 (4)

展开

阅读排行

Ubuntu上配置caffe+SSD (7371)
Zbar算法流程介绍 (5545)
使用caffe框架利用faster- (4530)
Ubuntu上用caffe的SSD (3704)
Windows7下Python安装 (3229)
用C++实现定积分运算 (2733)
Ubuntu上用caffe的SSD (1864)

关闭

使用faster rcnn训练umdf (1783)

LeetCode周练Contest-3 (1702)

机器学习算法逻辑回归 (1399)

评论排行

Zbar算法流程介绍 (5)

Ubuntu上配置caffe+SSD (3)

Ubuntu上用caffe的SSD (2)

Windows7下Python安装 (1)

使用faster rcnn训练umdf (1)

用C++实现定积分运算 (0)

Ubuntu上用caffe的SSD (0)

使用Python分析移动语音 (0)

2016年总结 (0)

LeetCode周练Contest-3 (0)

推荐文章

* CSDN新版博客feed流内测用户征集令

* Android检查更新下载安装

* 动手打造史上最简单的Recycleview 侧滑菜单

* TCP网络通讯如何解决分包粘包问题

* SDCC 2017之大数据技术实战线上峰会

* 快速集成一个视频直播功能

最新评论

```

23. # TODO(weiliu89): Construct the name using the last layer to avoid duplication.
24. # 10 x 10
25. out_layer = "conv6_1"
26. ConvBNLayer(net, from_layer, out_layer, use_batchnorm, use_relu, 256, 1, 0, 1,
27.               lr_mult=lr_mult)
28.
29. from_layer = out_layer
30. out_layer = "conv6_2"
31. ConvBNLayer(net, from_layer, out_layer, use_batchnorm, use_relu, 512, 3, 1, 2,
32.               lr_mult=lr_mult)
33.
34. # 5 x 5
35. from_layer = out_layer
36. out_layer = "conv7_1"
37. ConvBNLayer(net, from_layer, out_layer, use_batchnorm, use_relu, 128, 1, 0, 1,
38.               lr_mult=lr_mult)
39.
40. from_layer = out_layer
41. out_layer = "conv7_2"
42. ConvBNLayer(net, from_layer, out_layer, use_batchnorm, use_relu, 256, 3, 1, 2,
43.               lr_mult=lr_mult)
44.
45. # 3 x 3
46. from_layer = out_layer
47. out_layer = "conv8_1"
48. ConvBNLayer(net, from_layer, out_layer, use_batchnorm, use_relu, 256, 3, 1, 2,
49.               lr_mult=lr_mult)
50.
51. from_layer = out_layer
52. out_layer = "conv8_2"
53. ConvBNLayer(net, from_layer, out_layer, use_batchnorm, use_relu, 256, 3, 0, 1,
54.               lr_mult=lr_mult)
55.
56. # 1 x 1
57. from_layer = out_layer
58. out_layer = "conv9_1"
59. ConvBNLayer(net, from_layer, out_layer, use_batchnorm, use_relu, 128, 1, 0, 1,
60.               lr_mult=lr_mult)
61.

```

关闭

Zbar算法流程介绍

XXDK141: 大神有源码吗, 我的
是zbarimg.c 是基于imageMagick
版本的, opencv版本的哪里可...

Windows7下Python安装最新xgb

echorma: 楼主太棒了! 奇怪的是,
楼主分享的东西, 只有预编译的
xgboost能下载成功, 其他都下不了...

Ubuntu上用caffe的SSD方法训练

yx2017: @erfenxing:你好, 我也
是在用ssd训练widerface, 可否
交流下? 1343102499

Ubuntu上用caffe的SSD方法训练

erfenxing: 您好, 我用SSD训练
widerface数据集, 但是loss值从
20左右下降到大约5的时候就不再
下降了...

Ubuntu上配置caffe+SSD及detc

luuuyi: @Reezea:435977170

Ubuntu上配置caffe+SSD及detc

Reezea: 能否加个好友请教一
下, 最近一直没有安装成功
caffe, 总是各种问题, 希望能够
指点一二, 万分感激了。系...

Zbar算法流程介绍

windwolf2004: 大好人, 绝对原创

Zbar算法流程介绍

ATUXLI: 楼主想请教一下, 在提
取明暗宽度流之前会有density,
其定义是: int density = CFG...

Zbar算法流程介绍

luuuyi: @qq_34344335:Linux上
的, 不过clone下来自己改一改,
可以移植到Windows上

Zbar算法流程介绍

qq_34344335: 请问这个是什么
环境下的源码

```

62.     from_layer = out_layer
63.     out_layer = "conv9_2"
64.     ConvBNLayer(net, from_layer, out_layer, use_batchnorm, use_relu, 256, 3, 0, 1,
65.         lr_mult=lr_mult)
66.
67.     return net
68.
69.
70.     ### Modify the following parameters accordingly ###
71.     # The directory which contains the caffe code.
72.     # We assume you are running the script at the CAFFE_ROOT.
73.     caffe_root = os.getcwd()
74.
75.     # Set true if you want to start training right after generating all files.
76.     run_soon = True
77.     # Set true if you want to load from most recently saved snapshot.
78.     # Otherwise, we will load from the pretrain_model defined below.
79.     resume_training = True
80.     # If true, Remove old model files.
81.     remove_old_models = False
82.
83.     # The database file for training data. Created by data/VOC0712/create_data.sh
84.     train_data = "/media/scs4450/hard/VOCdevkit/FACE_LUYI/lmdb/FACE_LUYI_trainval_lmdb"
85.     # The database file for testing data. Created by data/VOC0712/create_data.sh
86.     test_data = "/media/scs4450/hard/VOCdevkit/FACE_LUYI/lmdb/FACE_LUYI_test_lmdb"
87.     # Specify the batch sampler.
88.     resize_width = 300
89.     resize_height = 300
90.     resize = "{}x{}".format(resize_width, resize_height)
91.     batch_sampler = [
92.         {
93.             'sampler': {
94.                 },
95.             'max_trials': 1,
96.             'max_sample': 1,
97.         },
98.         {
99.             'sampler': {
100.                 'min_scale': 0.3,

```

关闭

```
101.         'max_scale': 1.0,
102.         'min_aspect_ratio': 0.5,
103.         'max_aspect_ratio': 2.0,
104.     },
105.     'sample_constraint': {
106.         'min_jaccard_overlap': 0.1,
107.     },
108.     'max_trials': 50,
109.     'max_sample': 1,
110. },
111. {
112.     'sampler': {
113.         'min_scale': 0.3,
114.         'max_scale': 1.0,
115.         'min_aspect_ratio': 0.5,
116.         'max_aspect_ratio': 2.0,
117.     },
118.     'sample_constraint': {
119.         'min_jaccard_overlap': 0.3,
120.     },
121.     'max_trials': 50,
122.     'max_sample': 1,
123. },
124. {
125.     'sampler': {
126.         'min_scale': 0.3,
127.         'max_scale': 1.0,
128.         'min_aspect_ratio': 0.5,
129.         'max_aspect_ratio': 2.0,
130.     },
131.     'sample_constraint': {
132.         'min_jaccard_overlap': 0.5,
133.     },
134.     'max_trials': 50,
135.     'max_sample': 1,
136. },
137. {
138.     'sampler': {
139.         'min_scale': 0.3,
```

关闭

```
140.         'max_scale': 1.0,
141.         'min_aspect_ratio': 0.5,
142.         'max_aspect_ratio': 2.0,
143.     },
144.     'sample_constraint': {
145.         'min_jaccard_overlap': 0.7,
146.     },
147.     'max_trials': 50,
148.     'max_sample': 1,
149. },
150. {
151.     'sampler': {
152.         'min_scale': 0.3,
153.         'max_scale': 1.0,
154.         'min_aspect_ratio': 0.5,
155.         'max_aspect_ratio': 2.0,
156.     },
157.     'sample_constraint': {
158.         'min_jaccard_overlap': 0.9,
159.     },
160.     'max_trials': 50,
161.     'max_sample': 1,
162. },
163. {
164.     'sampler': {
165.         'min_scale': 0.3,
166.         'max_scale': 1.0,
167.         'min_aspect_ratio': 0.5,
168.         'max_aspect_ratio': 2.0,
169.     },
170.     'sample_constraint': {
171.         'max_jaccard_overlap': 1.0,
172.     },
173.     'max_trials': 50,
174.     'max_sample': 1,
175. },
176. ]
177. train_transform_param = {
178.     'mirror': True,
```

关闭

```
179.         'mean_value': [104, 117, 123],
180.         'resize_param': {
181.             'prob': 1,
182.             'resize_mode': P.Resize.WARP,
183.             'height': resize_height,
184.             'width': resize_width,
185.             'interp_mode': [
186.                 P.Resize.LINEAR,
187.                 P.Resize.AREA,
188.                 P.Resize.NEAREST,
189.                 P.Resize.CUBIC,
190.                 P.Resize.LANCZOS4,
191.             ],
192.         },
193.         'distort_param': {
194.             'brightness_prob': 0.5,
195.             'brightness_delta': 32,
196.             'contrast_prob': 0.5,
197.             'contrast_lower': 0.5,
198.             'contrast_upper': 1.5,
199.             'hue_prob': 0.5,
200.             'hue_delta': 18,
201.             'saturation_prob': 0.5,
202.             'saturation_lower': 0.5,
203.             'saturation_upper': 1.5,
204.             'random_order_prob': 0.0,
205.         },
206.         'expand_param': {
207.             'prob': 0.5,
208.             'max_expand_ratio': 4.0,
209.         },
210.         'emit_constraint': {
211.             'emit_type': caffe_pb2.EmitConstraint.CENTER,
212.         }
213.     }
214. test_transform_param = {
215.     'mean_value': [104, 117, 123],
216.     'resize_param': {
217.         'prob': 1,
```

关闭

```
218.         'resize_mode': P.Resize.WARP,
219.         'height': resize_height,
220.         'width': resize_width,
221.         'interp_mode': [P.Resize.LINEAR],
222.     },
223. }
224.
225. # If true, use batch norm for all newly added layers.
226. # Currently only the non batch norm version has been tested.
227. use_batchnorm = False
228. lr_mult = 1
229. # Use different initial learning rate.
230. if use_batchnorm:
231.     base_lr = 0.0004
232. else:
233.     # A learning rate for batch_size = 1, num_gpus = 1.
234.     base_lr = 0.00004
235.
236. # Modify the job name if you want.
237. job_name = "SSD_{}".format(resize)
238. # The name of the model. Modify it if you want.
239. model_name = "VGG_FACE_LUYI_{}".format(job_name)
240.
241. # Directory which stores the model .prototxt file.
242. save_dir = "/home/ly/ssd/caffe/models/VGGNet/FACE_LUYI/{}".format(job_name)
243. # Directory which stores the snapshot of models.
244. snapshot_dir = "/home/ly/ssd/caffe/models/VGGNet/FACE_LUYI/{}".format(job_name) #luyi
245. # Directory which stores the job script and log file.
246. job_dir = "/home/ly/ssd/caffe/jobs/VGGNet/FACE_LUYI/{}".format(job_name) #luyi
247. # Directory which stores the detection results.
248. output_result_dir = "
    {}/data/VOCdevkit/results/FACE_LUYI/{}/Main".format(os.environ['CAFFE_ROOT'], job_name) #luyi
249.
250. # model definition files.
251. train_net_file = "{}train.prototxt".format(save_dir)
252. test_net_file = "{}test.prototxt".format(save_dir)
253. deploy_net_file = "{}deploy.prototxt".format(save_dir)
254. solver_file = "{}solver.prototxt".format(save_dir)
255. # snapshot prefix.
```

关闭


```
256. snapshot_prefix = "{}/{}/".format(snapshot_dir, model_name)
257. # job script path.
258. job_file = "{}/{}.sh".format(job_dir, model_name)
259.
260. # Stores the test image names and sizes. Created by data/VOC0712/create_list.sh
261. name_size_file = "/home/ly/ssd/caffe/data/FACE_LUYI/test_name_size.txt"
262. # The pretrained model. We use the Fully convolutional reduced (atrous) VGGNet.
263. pretrain_model = "/home/ly/ssd/caffe/models/VGGNet/VGG_ILSVRC_16_layers_fc_reduced.caffemodel"
264. # Stores LabelMapItem.
265. label_map_file = "/home/ly/ssd/caffe/data/FACE_LUYI/labelmap_voc.prototxt"
266.
267. # MultiBoxLoss parameters.
268. num_classes = 2                                #luyi
269. share_location = True
270. background_label_id=0
271. train_on_diff_gt = True
272. normalization_mode = P.Loss.VALID
273. code_type = P.PriorBox.CENTER_SIZE
274. ignore_cross_boundary_bbox = False
275. mining_type = P.MultiBoxLoss.MAX_NEGATIVE
276. neg_pos_ratio = 3.
277. loc_weight = (neg_pos_ratio + 1.) / 4.
278. multibox_loss_param = {
279.     'loc_loss_type': P.MultiBoxLoss.SMOOTH_L1,
280.     'conf_loss_type': P.MultiBoxLoss.SOFTMAX,
281.     'loc_weight': loc_weight,
282.     'num_classes': num_classes,
283.     'share_location': share_location,
284.     'match_type': P.MultiBoxLoss.PER_PREDICTION,
285.     'overlap_threshold': 0.5,
286.     'use_prior_for_matching': True,
287.     'background_label_id': background_label_id,
288.     'use_difficult_gt': train_on_diff_gt,
289.     'mining_type': mining_type,
290.     'neg_pos_ratio': neg_pos_ratio,
291.     'neg_overlap': 0.5,
292.     'code_type': code_type,
293.     'ignore_cross_boundary_bbox': ignore_cross_boundary_bbox,
294. }
```

关闭

```
295. loss_param = {
296.     'normalization': normalization_mode,
297. }
298.
299. # parameters for generating priors.
300. # minimum dimension of input image
301. min_dim = 300
302. # conv4_3 ==> 38 x 38
303. # fc7 ==> 19 x 19
304. # conv6_2 ==> 10 x 10
305. # conv7_2 ==> 5 x 5
306. # conv8_2 ==> 3 x 3
307. # conv9_2 ==> 1 x 1
308. mbox_source_layers = ['conv4_3', 'fc7', 'conv6_2', 'conv7_2', 'conv8_2', 'conv9_2']
309. # in percent %
310. min_ratio = 20
311. max_ratio = 90
312. step = int(math.floor((max_ratio - min_ratio) / (len(mbox_source_layers) - 2)))
313. min_sizes = []
314. max_sizes = []
315. for ratio in xrange(min_ratio, max_ratio + 1, step):
316.     min_sizes.append(min_dim * ratio / 100.)
317.     max_sizes.append(min_dim * (ratio + step) / 100.)
318. min_sizes = [min_dim * 10 / 100.] + min_sizes
319. max_sizes = [min_dim * 20 / 100.] + max_sizes
320. steps = [8, 16, 32, 64, 100, 300]
321. aspect_ratios = [[2], [2, 3], [2, 3], [2, 3], [2], [2]]
322. # L2 normalize conv4_3.
323. normalizations = [20, -1, -1, -1, -1, -1]
324. # variance used to encode/decode prior bboxes.
325. if code_type == P.PriorBox.CENTER_SIZE:
326.     prior_variance = [0.1, 0.1, 0.2, 0.2]
327. else:
328.     prior_variance = [0.1]
329. flip = True
330. clip = False
331.
332. # Solver parameters.
333. # Defining which GPUs to use.
```

关闭

```
334. gpus = "1" #luyi
335. gpulist = gpus.split(",")
336. num_gpus = len(gpulist)
337.
338. # Divide the mini-batch to different GPUs.
339. batch_size = 32
340. accum_batch_size = 32
341. iter_size = accum_batch_size / batch_size
342. solver_mode = P.Solver.CPU
343. device_id = 0
344. batch_size_per_device = batch_size
345. if num_gpus > 0:
346.     batch_size_per_device = int(math.ceil(float(batch_size) / num_gpus))
347.     iter_size = int(math.ceil(float(accum_batch_size) / (batch_size_per_device * num_
348. solver_mode = P.Solver.GPU
349.     device_id = int(gpulist[0])
350.
351. if normalization_mode == P.Loss.NONE:
352.     base_lr /= batch_size_per_device
353. elif normalization_mode == P.Loss.VALID:
354.     base_lr *= 25. / loc_weight
355. elif normalization_mode == P.Loss.FULL:
356.     # Roughly there are 2000 prior bboxes per image.
357.     # TODO(weiliu89): Estimate the exact # of priors.
358.     base_lr *= 2000.
359.
360. # Evaluate on whole test set.
361. num_test_image = 4952 #luyi
362. test_batch_size = 8
363. # Ideally test_batch_size should be divisible by num_test_image,
364. # otherwise mAP will be slightly off the true value.
365. test_iter = int(math.ceil(float(num_test_image) / test_batch_size))
366.
367. solver_param = {
368.     # Train parameters
369.     'base_lr': base_lr,
370.     'weight_decay': 0.0005,
371.     'lr_policy': "multistep",
372.     'stepvalue': [80000, 100000, 120000],
```

关闭

```
373.     'gamma': 0.1,
374.     'momentum': 0.9,
375.     'iter_size': iter_size,
376.     'max_iter': 120000,
377.     'snapshot': 80000,
378.     'display': 10,
379.     'average_loss': 10,
380.     'type': "SGD",
381.     'solver_mode': solver_mode,
382.     'device_id': device_id,
383.     'debug_info': False,
384.     'snapshot_after_train': True,
385.     # Test parameters
386.     'test_iter': [test_iter],
387.     'test_interval': 10000,
388.     'eval_type': "detection",
389.     'ap_version': "11point",
390.     'test_initialization': False,
391. }
392.
393. # parameters for generating detection output.
394. det_out_param = {
395.     'num_classes': num_classes,
396.     'share_location': share_location,
397.     'background_label_id': background_label_id,
398.     'nms_param': {'nms_threshold': 0.45, 'top_k': 400},
399.     'save_output_param': {
400.         'output_directory': output_result_dir,
401.         'output_name_prefix': "comp4_det_test_",
402.         'output_format': "VOC",
403.         'label_map_file': label_map_file,
404.         'name_size_file': name_size_file,
405.         'num_test_image': num_test_image,
406.     },
407.     'keep_top_k': 200,
408.     'confidence_threshold': 0.01,
409.     'code_type': code_type,
410. }
411.
```

关闭

```
412. # parameters for evaluating detection results.
413. det_eval_param = {
414.     'num_classes': num_classes,
415.     'background_label_id': background_label_id,
416.     'overlap_threshold': 0.5,
417.     'evaluate_difficult_gt': False,
418.     'name_size_file': name_size_file,
419. }
420.
421. ### Hopefully you don't need to change the following ###
422. # Check file.
423. check_if_exist(train_data)
424. check_if_exist(test_data)
425. check_if_exist(label_map_file)
426. check_if_exist(pretrain_model)
427. make_if_not_exist(save_dir)
428. make_if_not_exist(job_dir)
429. make_if_not_exist(snapshot_dir)
430.
431. # Create train net.
432. net = caffe.NetSpec()
433. net.data, net.label = CreateAnnotatedDataLayer(train_data, batch_size=batch_size_per
434.     train=True, output_label=True, label_map_file=label_map_file,
435.     transform_param=train_transform_param, batch_sampler=batch_sampler)
436.
437. VGGNetBody(net, from_layer='data', fully_conv=True, reduced=True,
438.     dropout=False)
439.
440. AddExtraLayers(net, use_batchnorm, lr_mult=lr_mult)
441.
442. mbox_layers = CreateMultiBoxHead(net, data_layer='data', from_layers=mbox_source_layers,
443.     use_batchnorm=use_batchnorm, min_sizes=min_sizes, max_sizes=max_sizes,
444.     aspect_ratios=aspect_ratios, steps=steps, normalizations=normalizations,
445.     num_classes=num_classes, share_location=share_location, flip=flip, clip=clip,
446.     prior_variance=prior_variance, kernel_size=3, pad=1, lr_mult=lr_mult)
447.
448. # Create the MultiBoxLossLayer.
449. name = "mbox_loss"
450. mbox_layers.append(net.label)
```

关闭

```

451. net[name] = L.MultiBoxLoss(*mbox_layers, multibox_loss_param=multibox_loss_param,
452.                             loss_param=loss_param, include=dict(phase=caffe_pb2.Phase.Value('TRAIN')),
453.                             propagate_down=[True, True, False, False])
454.
455. with open(train_net_file, 'w') as f:
456.     print('name: "{}_train"'.format(model_name), file=f)
457.     print(net.to_proto(), file=f)
458. shutil.copy(train_net_file, job_dir)
459.
460. # Create test net.
461. net = caffe.NetSpec()
462. net.data, net.label = CreateAnnotatedDataLayer(test_data, batch_size=test_batch_size,
463.                                                 train=False, output_label=True, label_map_file=label_map_file,
464.                                                 transform_param=test_transform_param)
465.
466. VGGNetBody(net, from_layer='data', fully_conv=True, reduced=True, dilated=True,
467.             dropout=False)
468.
469. AddExtraLayers(net, use_batchnorm, lr_mult=lr_mult)
470.
471. mbox_layers = CreateMultiBoxHead(net, data_layer='data', from_layers=mbox_source_layers,
472.                                   use_batchnorm=use_batchnorm, min_sizes=min_sizes, max_sizes=max_sizes,
473.                                   aspect_ratios=aspect_ratios, steps=steps, normalizations=normalizations,
474.                                   num_classes=num_classes, share_location=share_location, flip=flip, clip=clip,
475.                                   prior_variance=prior_variance, kernel_size=3, pad=1,
476.                                   )
477. conf_name = "mbox_conf"
478. if multibox_loss_param["conf_loss_type"] == P.MultiBoxLoss.SOFTMAX:
479.     reshape_name = "{}_reshape".format(conf_name)
480.     net[reshape_name] = L.Reshape(net[conf_name], shape=dict(dim=[0, -1, num_classes]))
481.     softmax_name = "{}_softmax".format(conf_name)
482.     net[softmax_name] = L.Softmax(net[reshape_name], axis=2)
483.     flatten_name = "{}_flatten".format(conf_name)
484.     net[flatten_name] = L.Flatten(net[softmax_name], axis=1)
485.     mbox_layers[1] = net[flatten_name]
486. elif multibox_loss_param["conf_loss_type"] == P.MultiBoxLoss.LOGISTIC:
487.     sigmoid_name = "{}_sigmoid".format(conf_name)
488.     net[sigmoid_name] = L.Sigmoid(net[conf_name])
489.     mbox_layers[1] = net[sigmoid_name]

```

关闭

```

490.
491. net.detection_out = L.DetectionOutput(*mbox_layers,
492.     detection_output_param=det_out_param,
493.     include=dict(phase=caffe_pb2.Phase.Value('TEST')))
494. net.detection_eval = L.DetectionEvaluate(net.detection_out, net.label,
495.     detection_evaluate_param=det_eval_param,
496.     include=dict(phase=caffe_pb2.Phase.Value('TEST')))
497.
498. with open(test_net_file, 'w') as f:
499.     print('name: "{}_test"'.format(model_name), file=f)
500.     print(net.to_proto(), file=f)
501. shutil.copy(test_net_file, job_dir)
502.
503. # Create deploy net.
504. # Remove the first and last layer from test net.
505. deploy_net = net
506. with open(deploy_net_file, 'w') as f:
507.     net_param = deploy_net.to_proto()
508.     # Remove the first (AnnotatedData) and last (DetectionEvaluate) layer from test net.
509.     del net_param.layer[0]
510.     del net_param.layer[-1]
511.     net_param.name = '{}_deploy'.format(model_name)
512.     net_param.input.extend(['data'])
513.     net_param.input_shape.extend([
514.         caffe_pb2.BlobShape(dim=[1, 3, resize_height, resize_width])
515.     ])
516.     print(net_param, file=f)
517. shutil.copy(deploy_net_file, job_dir)
518.
519. # Create solver.
520. solver = caffe_pb2.SolverParameter(
521.     train_net=train_net_file,
522.     test_net=[test_net_file],
523.     snapshot_prefix=snapshot_prefix,
524.     **solver_param)
525.
526. with open(solver_file, 'w') as f:
527.     print(solver, file=f)
528. shutil.copy(solver_file, job_dir)

```

关闭

```

529. max_iter = 0
530. # Find most recent snapshot.
531. for file in os.listdir(snapshot_dir):
532.     if file.endswith(".solverstate"):
533.         basename = os.path.splitext(file)[0]
534.         iter = int(basename.split("{}_iter_".format(model_name))[1])
535.         if iter > max_iter:
536.             max_iter = iter
537.
538. train_src_param = '--weights "{}" \\n'.format(pretrain_model)
539. if resume_training:
540.     if max_iter > 0:
541.         train_src_param = '--snapshot="
542.         "{}_iter_{}.solverstate" \\n'.format(snapshot_prefix, max_iter)
543.
544. if remove_old_models:
545.     # Remove any snapshots smaller than max_iter.
546.     for file in os.listdir(snapshot_dir):
547.         if file.endswith(".solverstate"):
548.             basename = os.path.splitext(file)[0]
549.             iter = int(basename.split("{}_iter_".format(model_name))[1])
550.             if max_iter > iter:
551.                 os.remove("{}_iter_{}".format(snapshot_dir, file))
552.         if file.endswith(".caffemodel"):
553.             basename = os.path.splitext(file)[0]
554.             iter = int(basename.split("{}_iter_".format(model_name))[1])
555.             if max_iter > iter:
556.                 os.remove("{}_iter_{}".format(snapshot_dir, file))
557.
558. # Create job file.
559. with open(job_file, 'w') as f:
560.     f.write('cd {}\\n'.format(caffe_root))
561.     f.write('/home/ly/ssd/caffe/build/tools/caffe train \\n') #luyi
562.     f.write('--solver "{}" \\n'.format(solver_file))
563.     f.write(train_src_param)
564.     if solver_param['solver_mode'] == P.Solver.GPU:
565.         f.write('--gpu {} 2>&1 | tee {}/{}.log\\n'.format(gpus, job_dir, model_name))
566.     else:
567.         f.write('2>&1 | tee {}/{}.log\\n'.format(job_dir, model_name))

```

关闭


```

567.
568. # Copy the python script to job_dir.
569. py_file = os.path.abspath(__file__)
570. shutil.copy(py_file, job_dir)
571.
572. # Run the job.
573. os.chmod(job_file, stat.S_IRWXU)
574. if run_soon:
575.     subprocess.call(job_file, shell=True)

```

实验结果

计划训练的迭代次数是12w次，但是在K80上只开了一个核来进行计算，差不多一天可以迭代1w+次吧，跑了12w次，打断来测试，在K80上，检测单张人脸图片，分辨率在300X300左右，速度为40ms左右，也就是说帧率25fps，速度还是很快的。至于准确度，在log文件里面，每1w次迭代之后会计算一个mAP，第6w次的时候mAP（作者自定义过，跟mAP比较像）为0.965：

```

23416 I0306 14:48:02.205540 29719 solver.cpp:259] Train net output #0: mbox_loss = 0.698651 (* 1 = 0.6
23417 I0306 14:48:02.205556 29719 sgd_solver.cpp:138] Iteration 59990, lr = 0.001
23418 I0306 14:50:59.707754 29719 solver.cpp:433] Iteration 60000, Testing net (#0)
23419 I0306 14:50:59.707929 29719 net.cpp:693] Ignoring source layer mbox_loss
23420 I0306 14:55:33.095422 29719 solver.cpp:546] Test net output #0: detection_eval = 0.965216
23421 I0306 14:55:36.818568 29719 solver.cpp:243] Iteration 60010, loss = 0.640824
23422 I0306 14:55:36.818614 29719 solver.cpp:259] Train net output #0: mbox_loss = 0.639989 (* 1 = 0.6
23423 I0306 14:55:36.818624 29719 sgd_solver.cpp:138] Iteration 60000, lr = 0.0
23424 I0306 14:56:34.483124 29719 solver.cpp:243] Iteration 60010, loss = 0.616
23425 I0306 14:56:34.483324 29719 solver.cpp:259] Train net output #0: mbox_loss = 0.843735 (* 1 = 0.843735 loss)

```

总体来说，效果不错，不过目标是在嵌入式平台上做到实时，还需要再继续努力。

欢迎各位来交流：435977170（Q&Q）

顶 踩

0

0

关闭

[上一篇](#) [用C++实现定积分运算](#)[下一篇](#) [机器学习KNN方法介绍及实现并实践（约会满意度统计）](#)

相关文章推荐

- Ubuntu上配置caffe+SSD及demo演示（附带问题...
- Presto的服务治理与架构在京东的实践与应用--王...
- 【用Python学习Caffe】2. 使用Caffe完成图像目标...
- 深入掌握Kubernetes应用实践--王渊命
- Ubuntu上用caffe的SSD方法训练Pascal VOC数据集
- Python基础知识汇总
- RCNN系列实验的PASCAL VOC数据集格式设置
- Android核心技术详解
- SSD框架训练自己的数据集
- Retrofit 从入门封装到源码解析
- 论文笔记 SSD: Single Shot MultiBox Detector
- 自然语言处理工具Word2Vec
- SSD框架训练自己的数据集
- SSD源码解读1~~~~~ssd_pascal.py
- SSD 安装、训练、测试（ubuntu14.04+cuda7.5+
- 使用faster rcnn训练umdfaces数据集

查看评论

1楼 [erfenxing](#) 2017-06-05 17:13发表



您好，我用SSD训练widerface数据集，但是loss值从20左右下降到大约5的时候就不再下降了，请问会是数据集本身的问题吗？请问用create_list.sh 和create_data.sh生成lmdb的时候，有对文件做修改吗？谢谢！

Re: [yx2017](#) 2017-07-12 11:20发表



回复erfenxing：你好，我也是在用ssd训练widerface，可否交流下？1343102499

您还没有登录,请[\[登录\]](#)或[\[注册\]](#)

* 以上用户言论只代表其个人观点，不代表CSDN网站的观点或立场

[关闭](#)

[公司简介](#) | [招贤纳士](#) | [广告服务](#) | [联系方式](#) | [版权声明](#) | [法律顾问](#) | [问题报告](#) | [合作伙伴](#) | [论坛反馈](#)

网站客服

杂志客服

微博客服

webmaster@csdn.net

400-660-0108

| 北京创新乐知信息技术有限公司 版权所有 | 江苏知之为计算机有限公司 | 江苏乐知网络技术有限公司

京 ICP 证 09002463 号 | Copyright © 1999-2017, CSDN.NET, All Rights Reserved



关闭