

This repository

Search

Pull requests

Issues

Gist

Shark-ML / Shark

Watch26

Star160

Fork62

Code

Issues19

Pull requests1

Projects0

Pulse

Graphs

The Shark Machine Learning Library. See more: <http://shark-ml.org/>

3,782 commits

6 branches

6 releases

15 contributors

GPL-3.0

Branch: master

New pull request

Create new file

Upload files

Find file

Clone or download

Ulfgard committed on GitHub Merge pull request #178 from Shark-ML/VSfix			Latest commit 14d47f6 2 days ago
Test	fixed error in a test		2 days ago
doc	Added unified initialization with starting points for optimizers		7 days ago
examples	renamed ArgMaxConverter to Classifier		7 days ago
include	Merge pull request #178 from Shark-ML/VSfix		2 days ago
src	Merge pull request #178 from Shark-ML/VSfix		2 days ago
.gitignore	removed Rng component and replaced by single header Core/Random.h		23 days ago
.travis.yml	Test with gcc-6		24 days ago
CMakeLists.txt	added SIMD to dthe default gemm and implemented benchmarks		3 months ago
COPYING	changed license		3 years ago
COPYING.LESSER	fixed legal thingy that prevented CPack from building		3 years ago
CTestConfig.cmake	oops		2 years ago
README.txt	Update README.txt		21 days ago
SharkConfig.cmake.in	oops		2 years ago
SharkConfigVersion.cmake.in	updated shark packaging for cmake		4 years ago
UseShark.cmake	fixed a few smaller issues with NDEBUG		2 years ago
appveyor.yml	Attempt to fix MSVC/DLL linker error		22 days ago
cBlasCheck.cpp	added support for generic cblas libraries		2 years ago
cmake_uninstall.cmake.in	simple "make uninstall" added		4 years ago
update_remora.sh	added remora update script		3 months ago

README.txt

Shark is a fast, modular, general open-source C++ machine learning library.

Shark is licensed under the GNU Lesser General Public License, please see the files COPYING and COPYING.LESSER, or visit <http://www.gnu.org/licenses> .

Any application of the SHARK code toward military research and use is expressly against the wishes of the SHARK development team.

INSTALLATION / DOCUMENTATION

-----

The entry point to the Shark library documentation is located at doc/index.html . For installation instructions, please click on "Getting started" on that page. Short version of installation guide: issue "ccmake ." in the main directory to select your build options, and afterwards issue "make" in the main directory -- you should be done (assuming Boost and CMake were installed). See the documentation for detailed instructions.

BUILDING THE DOCUMENTATION: To build the documentation yourself (e.g., if you need to read it locally in order to install it, i.e., because you don't have internet), see doc/README.txt

FILE STRUCTURE  
-----

README.txt	This file (residing in the root directory of the Shark library).
CMakeLists.txt	Definitions for the CMake build system.
include/	This directory and its sub-directories hold all include files of the library. Note that some functionality is implemented in lower-level Impl/ folders and inline .inl files.
lib/	The Shark library is placed in this directory. In the source code distribution this directory is initially empty, and the library is placed into the directory as the results of compilation. Binary distributions already contain the library, pre-built in release mode.
doc/	All documentation files are found in this sub-directory. In packaged versions of Shark the html documentation is pre-built; the repository provides the corresponding sources. The documentation contains technical reference documents for all classes and functions as well as a collection of introductory and advanced tutorials.
doc/index.html	Entry point to the Shark documentation.
examples/	The examples directory contains example use-cases of the most important algorithms implemented in Shark. Besides exemplifying powerful learning algorithms, these programs are intended as starting points for experimentation with the library. The executables corresponding to the C++ example programs are found in examples/bin/.
Test/	Shark comes with a large collection of unit tests, all of which reside inside the Test directory.
bin/	The binaries of the Shark unit tests are placed here. Once the CMake build system is set up (with the "ccmake" command or equivalent) the whole test suite can be executed with the command "make test", issued in the Shark root directory.
src/	Source files of the Shark library. Note that from Shark version 3 onwards large parts of the library are templated and therefore header-only.
contrib/	The contrib directory contains (non-standard) tools by third parties. Typically, there is no need for users of Shark to deal with these tools directly.
gpl-3.0.txt	GNU general public license, version 3.

Note:  
Depending of the type of Shark distribution (binary or source package, or current repository snapshot) not all of theses files and directories are present.

PACKAGE STRUCTURE  
-----

>> Note for users of Shark 2: <<  
The internal structure of the Shark library has changed in the

transition to version 3. The old infrastructure packages Array, Rng, and FileUtil, as well as parts of LinAlg, have been replaced with more modern solutions provided by Boost. The machine learning related components EALib, MOO-EALib, Mixture, ReClam, and TimeSeries have been unified and organized into competely new interfaces. Therefore there is no one-to-one correspondance between files or even concepts in version 3 and in older versions of Shark. In fact, the lion's share of the library has been rewritten from scratch, and this is also reflected in a completely new structure. In particular, many of the rather independent sub-modules (such as Mixture and MOO-EALib) have been unified. They now share the same top-level interfaces and thus form a coherent learning architecture.

The organization of the include/ directory reflects the structure of the Shark library. It consists of the following modules:

GENERAL INFRASTRUCTURE:

LinAlg	Data structures and algorithms for typical linear algebra computations. For (dense and sparse) vector and matrix classes Shark relies on Boost uBLAS. Many higher level algorithms (such as singular value decomposition) are still implemented by the library itself.
Statistics	This component is new in Shark 3. It wraps the capabilities of Boost accumulators, and it provides tools that appear regularly in machine learning, such as the Mann-Whitney U-test (also known as the Wilcoxon rank-sum test).

LEARNING INFRASTRUCTURE:

Core	The core module is the central place for all top-level interfaces. In addition it holds a few infrastructure classes, such as exceptions.
Data	The data module hosts data containers that have been specifically designed for the needs of machine learning code. Also, data can be imported and exported from and to different standard machine learning data file formats.

MACHINE LEARNING:

Models	Models are adaptive systems, the architectures on top of which (machine) learning happens. Shark features a rich set of models, from simple linear maps to (feed-forward and recurrent) neural networks, support vector machines, and different types of trees. Models can also be concatenated with data format converters and other models.
ObjectiveFunctions	This module collects different types of cost, fitness, or objective functions for learning. The bandwidth includes data-dependent error functions based on simple loss functions, cross-validation, area under the ROC curve, and different objectives used for model selection.
Algorithms	All actual learning algorithms reside in this module. There are two main groups of learning algorithms, namely iterative optimizers and more specialized model trainers. General optimizers are organized into direct search and gradient-based optimization. Specialized algorithms for linear programming (a part of GLPK, the GNU linear programming kit) and quadratic programming for training of non-linear support vector machines are included. Shark also ships with algorithms for efficient nearest neighbor search.
Fuzzy	The fuzzy module provides classes for the representation of linguistic terms, variables, operators and rules, as well as fuzzy logic interference engines and controllers.

Unsupervised

This module contains the Shark implementation of restricted Boltzmann machines (RBMs), a recent experimental feature of Shark.