

[Start Here](#)[Blog](#)[Books](#)[About](#)[Contact](#)

Need help with Deep Learning? [Take the FREE Mini-Course](#)

Exploratory Configuration of a Multilayer Perceptron Network for Time Series Forecasting

by **Jason Brownlee** on April 26, 2017 in **Deep Learning**



It can be difficult when starting out on a new predictive modeling project with neural networks.

There is so much to configure, and no clear idea where to start.

It is important to be systematic. You can break bad assumptions and quickly hone in on configurations that work and areas for further investigation likely to payoff.

[Get Your Start in Machine Learning](#)

In this tutorial, you will discover how to use exploratory configuration of multilayer perceptron (MLP) neural networks to find good first-cut models for time series forecasting.

After completing this tutorial, you will know:

- How to design a robust experimental test harness to evaluate MLP models for time series forecasting.
- Systematic experimental designs for varying epochs, neurons, and lag configurations.
- How to interpret results and use diagnostics to learn more about well-performing models.

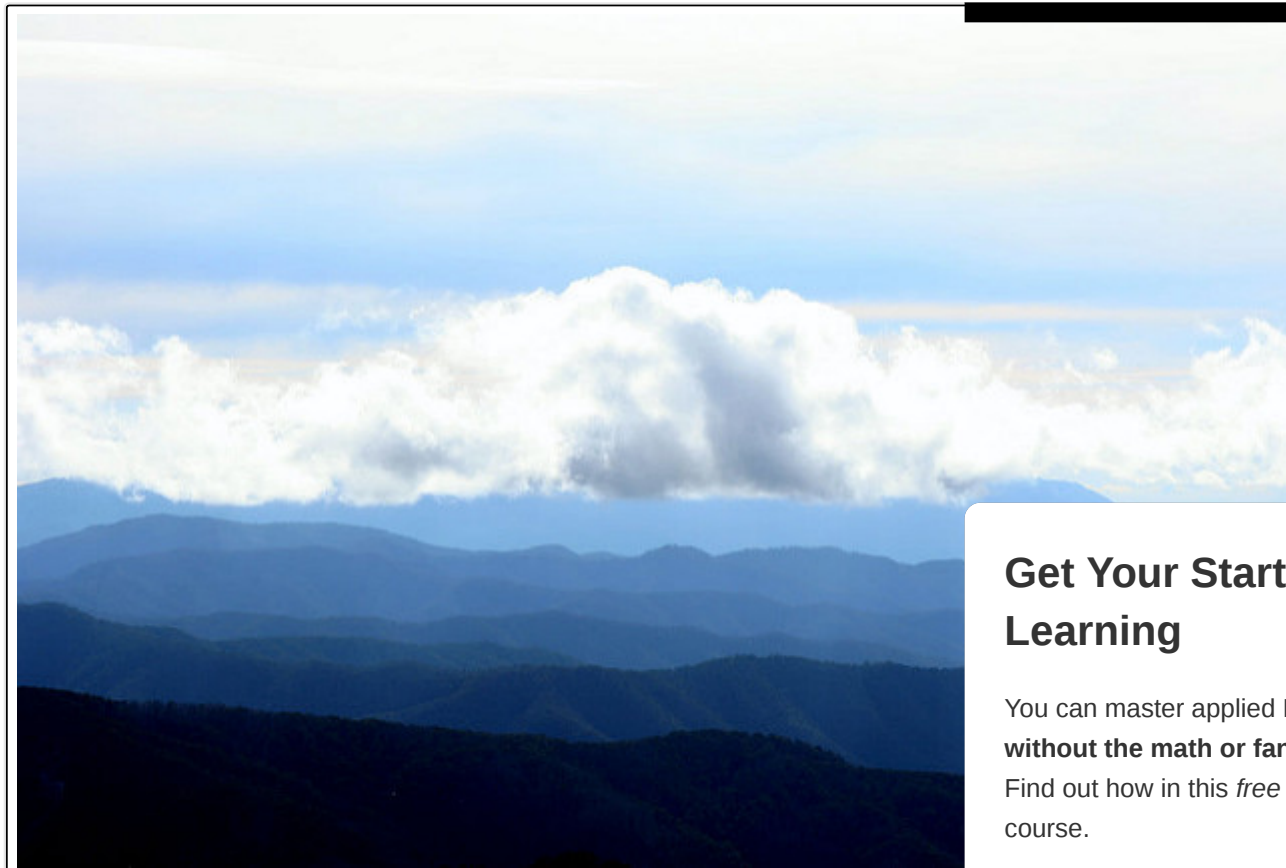
Let's get started.

- **Update July/2017:** Changed function for creating models to be more descriptive.

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**
Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Exploratory Configuration of a Multilayer Perceptron Network for Time Series Forecasting
Photo by [Lachlan Donald](#), some rights reserved.

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Tutorial Overview

This tutorial is broken down into 6 parts. They are:

1. Shampoo Sales Dataset
2. Experimental Test Harness
3. Vary Training Epochs
4. Vary Hidden Layer Neurons
5. Vary Hidden Layer Neurons with Lag
6. Review of Results

Get Your Start in Machine Learning

Environment

This tutorial assumes you have a Python SciPy environment installed. You can use either Python 2 or 3 with this example.

This tutorial assumes you have Keras v2.0 or higher installed with either the TensorFlow or Theano backend.

This tutorial also assumes you have scikit-learn, Pandas, NumPy, and Matplotlib installed.

Next, let's take a look at a standard time series forecasting problem that we can use as context for this experiment.

If you need help setting up your Python environment, see this post:

- [How to Setup a Python Environment for Machine Learning and Deep Learning with Anaconda](#)

Shampoo Sales Dataset

This dataset describes the monthly number of sales of shampoo over a 3-year period.

The units are a sales count and there are 36 observations. The original dataset is credited to Makrid

You can download and learn more about the dataset [here](#).

The example below loads and creates a plot of the loaded dataset.

```
1 # load and plot dataset
2 from pandas import read_csv
3 from pandas import datetime
4 from matplotlib import pyplot
5 # load dataset
6 def parser(x):
7     return datetime.strptime('190'+x, '%Y-%m')
8 series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True, date_parser=parser)
9 # summarize first few rows
10 print(series.head())
11 # line plot
12 series.plot()
13 pyplot.show()
```

Running the example loads the dataset as a Pandas Series and prints the first 5 rows.

Get Your Start in Machine Learning ×

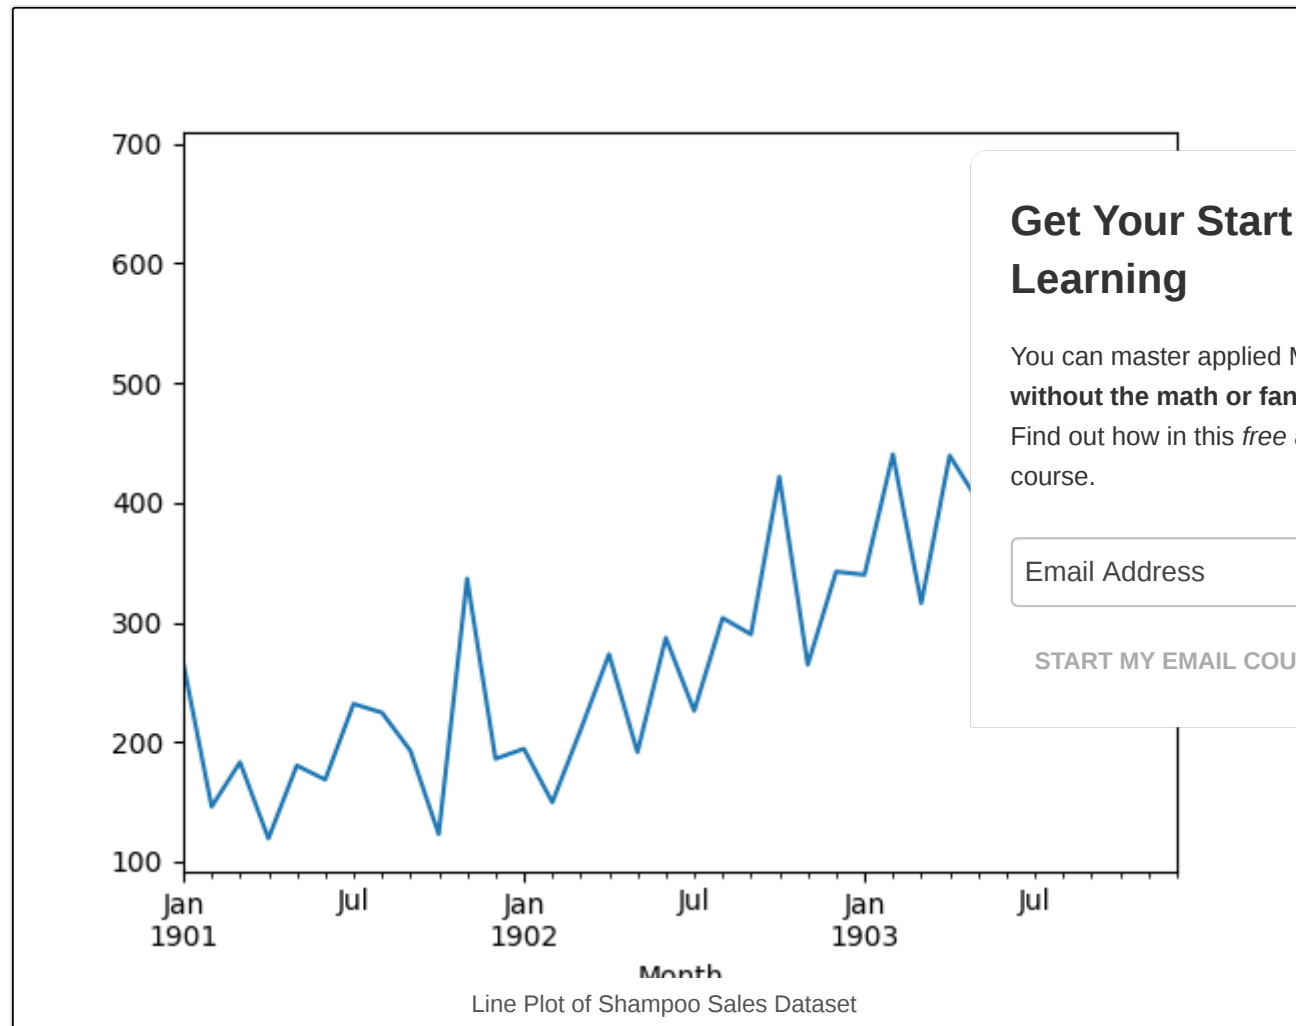
You can master applied Machine Learning **without the math or fancy degree**. Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
1 Month
2 1901-01-01 266.0
3 1901-02-01 145.9
4 1901-03-01 183.1
5 1901-04-01 119.3
6 1901-05-01 180.3
7 Name: Sales, dtype: float64
```

A line plot of the series is then created showing a clear increasing trend.



Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Next, we will take a look at the model configuration and test harness used in the experiment.

Experimental Test Harness

This section describes the test harness used in this tutorial.

Data Split

We will split the Shampoo Sales dataset into two parts: a training and a test set.

The first two years of data will be taken for the training dataset and the remaining one year of data will be used for the test set.

Models will be developed using the training dataset and will make predictions on the test dataset.

The persistence forecast (naive forecast) on the test dataset achieves an error of 136.761 monthly sales, which is the baseline bound of performance on the test set.

Model Evaluation

A rolling-forecast scenario will be used, also called walk-forward model validation.

Each time step of the test dataset will be walked one at a time. A model will be used to make a forecast on the test set will be taken and made available to the model for the forecast on the next time step.

This mimics a real-world scenario where new Shampoo Sales observations would be available each month.

This will be simulated by the structure of the train and test datasets.

All forecasts on the test dataset will be collected and an error score calculated to summarize the skill of the model. The root mean squared error (RMSE) will be used as it punishes large errors and results in a score that is in the same units as the forecast data, namely monthly shampoo sales.

Data Preparation

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Before we can fit an MLP model to the dataset, we must transform the data.

The following three data transforms are performed on the dataset prior to fitting a model and making a forecast.

1. **Transform the time series data so that it is stationary.** Specifically, a lag=1 differencing to remove the increasing trend in the data.
2. **Transform the time series into a supervised learning problem.** Specifically, the organization of data into input and output patterns where the observation at the previous time step is used as an input to forecast the observation at the current timestep
3. **Transform the observations to have a specific scale.** Specifically, to rescale the data to values between -1 and 1.

These transforms are inverted on forecasts to return them into their original scale before calculating and error score.

MLP Model

We will use a base MLP model with 1 neuron hidden layer, a rectified linear activation function on hidden neurons.

A batch size of 4 is used where possible, with the training data truncated to ensure the number of parameters is less than 2 is used.

Normally, the training dataset is shuffled after each batch or each epoch, which can aid in fitting the model to new problems. Shuffling was turned off for all experiments as it seemed to result in better performance. No shuffling was used for time series forecasting.

The model will be fit using the efficient ADAM optimization algorithm and the mean squared error loss function.

Experimental Runs

Each experimental scenario will be run 30 times and the RMSE score on the test set will be recorded from the end each run.

Let's dive into the experiments.

Vary Training Epochs

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

In this first experiment, we will investigate varying the number of training epochs for a simple MLP with one hidden layer and one neuron in the hidden layer.

We will use a batch size of 4 and evaluate training epochs 50, 100, 500, 1000, and 2000.

The complete code listing is provided below.

This code listing will be used as the basis for all following experiments, with only the changes to this code provided in subsequent sections.

```
1 from pandas import DataFrame
2 from pandas import Series
3 from pandas import concat
4 from pandas import read_csv
5 from pandas import datetime
6 from sklearn.metrics import mean_squared_error
7 from sklearn.preprocessing import MinMaxScaler
8 from keras.models import Sequential
9 from keras.layers import Dense
10 from math import sqrt
11 import matplotlib
12 # be able to save images on server
13 matplotlib.use('Agg')
14 from matplotlib import pyplot
15 import numpy
16
17 # date-time parsing function for loading the dataset
18 def parser(x):
19     return datetime.strptime('190'+x, '%Y-%m')
20
21 # frame a sequence as a supervised learning problem
22 def timeseries_to_supervised(data, lag=1):
23     df = DataFrame(data)
24     columns = [df.shift(i) for i in range(1, lag+1)]
25     columns.append(df)
26     df = concat(columns, axis=1)
27     return df
28
29 # create a differenced series
30 def difference(dataset, interval=1):
31     diff = list()
32     for i in range(interval, len(dataset)):
33         value = dataset[i] - dataset[i - interval]
34         diff.append(value)
35     return Series(diff)
```

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning


```

36
37 # invert differenced value
38 def inverse_difference(history, yhat, interval=1):
39     return yhat + history[-interval]
40
41 # scale train and test data to [-1, 1]
42 def scale(train, test):
43     # fit scaler
44     scaler = MinMaxScaler(feature_range=(-1, 1))
45     scaler = scaler.fit(train)
46     # transform train
47     train = train.reshape(train.shape[0], train.shape[1])
48     train_scaled = scaler.transform(train)
49     # transform test
50     test = test.reshape(test.shape[0], test.shape[1])
51     test_scaled = scaler.transform(test)
52     return scaler, train_scaled, test_scaled
53
54 # inverse scaling for a forecasted value
55 def invert_scale(scaler, X, yhat):
56     new_row = [x for x in X] + [yhat]
57     array = numpy.array(new_row)
58     array = array.reshape(1, len(array))
59     inverted = scaler.inverse_transform(array)
60     return inverted[0, -1]
61
62 # fit an MLP network to training data
63 def fit_model(train, batch_size, nb_epoch, neurons):
64     X, y = train[:, 0:-1], train[:, -1]
65     model = Sequential()
66     model.add(Dense(neurons, activation='relu', input_dim=X.shape[1]))
67     model.add(Dense(1))
68     model.compile(loss='mean_squared_error', optimizer='adam')
69     model.fit(X, y, epochs=nb_epoch, batch_size=batch_size, verbose=0, shuffle=False)
70     return model
71
72 # run a repeated experiment
73 def experiment(repeats, series, epochs, lag, neurons):
74     # transform data to be stationary
75     raw_values = series.values
76     diff_values = difference(raw_values, 1)
77     # transform data to be supervised learning
78     supervised = timeseries_to_supervised(diff_values, lag)
79     supervised_values = supervised.values[lag:, :]
80     # split data into train and test-sets
81     train, test = supervised_values[0:-12], supervised_values[-12:]
82     # transform the scale of the data

```

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```

83 scaler, train_scaled, test_scaled = scale(train, test)
84 # run experiment
85 error_scores = list()
86 for r in range(repeats):
87     # fit the model
88     batch_size = 4
89     train_trimmed = train_scaled[2:, :]
90     model = fit_model(train_trimmed, batch_size, epochs, neurons)
91     # forecast test dataset
92     test_reshaped = test_scaled[:, 0:-1]
93     output = model.predict(test_reshaped, batch_size=batch_size)
94     predictions = list()
95     for i in range(len(output)):
96         yhat = output[i, 0]
97         X = test_scaled[i, 0:-1]
98         # invert scaling
99         yhat = invert_scale(scaler, X, yhat)
100        # invert differencing
101        yhat = inverse_difference(raw_values, yhat, len(test_scaled)+1-i)
102        # store forecast
103        predictions.append(yhat)
104    # report performance
105    rmse = sqrt(mean_squared_error(raw_values[-12:], predictions))
106    print('%d) Test RMSE: %.3f' % (r+1, rmse))
107    error_scores.append(rmse)
108 return error_scores
109
110 # load dataset
111 series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=
112 # experiment
113 repeats = 30
114 results = DataFrame()
115 lag = 1
116 neurons = 1
117 # vary training epochs
118 epochs = [50, 100, 500, 1000, 2000]
119 for e in epochs:
120     results[str(e)] = experiment(repeats, series, e, lag, neurons)
121 # summarize results
122 print(results.describe())
123 # save boxplot
124 results.boxplot()
125 pyplot.savefig('boxplot_epochs.png')

```

Running the experiment prints the test set RMSE at the end of each experimental run.

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

At the end of all runs, a table of summary statistics is provided, one row for each statistic and one configuration for each column.

The summary statistics suggest that on average 1000 training epochs resulted in the better performance with a general decreasing trend in error with the increase of training epochs.

		50	100	500	1000	2000
1						
2	count	30.000000	30.000000	30.000000	30.000000	30.000000
3	mean	129.660167	129.388944	111.444027	103.821703	107.500301
4	std	30.926344	28.499592	23.181317	22.138705	24.780781
5	min	94.598957	94.184903	89.506815	86.511801	86.452041
6	25%	105.198414	105.722736	90.679930	90.058655	86.457260
7	50%	129.705407	127.449491	93.508245	90.118331	90.074494
8	75%	141.420145	149.625816	136.157299	135.510850	135.741340
9	max	198.716220	198.704352	141.226816	139.994388	142.097747

A box and whisker plot of the distribution of test RMSE scores for each configuration was also created.

The plot highlights that each configuration shows the same general spread in test RMSE scores (box and whiskers) with the increase of training epochs.

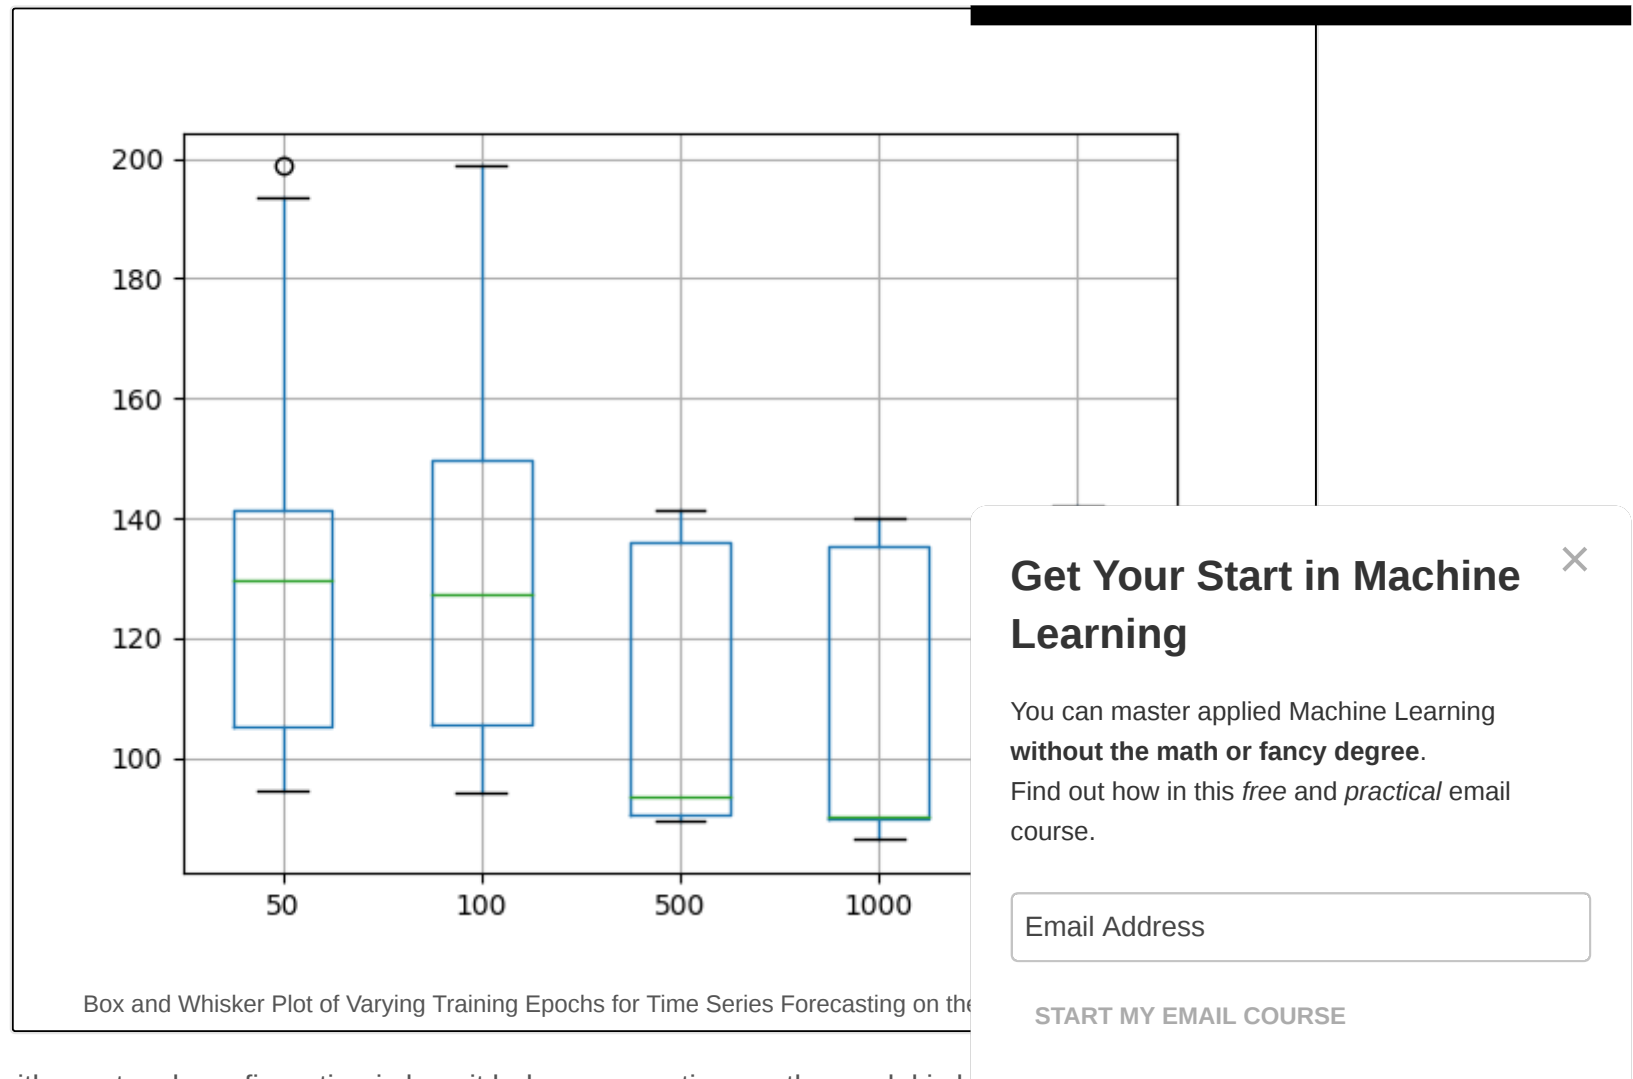
The results confirm that the configured MLP trained for 1000 is a good starting point on this problem.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning



Another angle to consider with a network configuration is how it behaves over time as the model is being trained.

We can evaluate the model on the training and test datasets after each training epoch to get an idea as to if the configuration is overfitting or underfitting the problem.

We will use this diagnostic approach on the top result from each set of experiments. A total of 10 repeats of the configuration will be run and the train and test RMSE scores after each training epoch plotted on a line plot.

In this case, we will use this diagnostic on the MLP fit for 1000 epochs.

[Get Your Start in Machine Learning](#)

The complete diagnostic code listing is provided below.

As with the previous code listing, the code listing below will be used as the basis for all diagnostics in this tutorial and only the changes to this listing will be provided in subsequent sections.

```
1 from pandas import DataFrame
2 from pandas import Series
3 from pandas import concat
4 from pandas import read_csv
5 from pandas import datetime
6 from sklearn.metrics import mean_squared_error
7 from sklearn.preprocessing import MinMaxScaler
8 from keras.models import Sequential
9 from keras.layers import Dense
10 from math import sqrt
11 import matplotlib
12 # be able to save images on server
13 matplotlib.use('Agg')
14 from matplotlib import pyplot
15 import numpy
16
17 # date-time parsing function for loading the dataset
18 def parser(x):
19     return datetime.strptime('190'+x, '%Y-%m')
20
21 # frame a sequence as a supervised learning problem
22 def timeseries_to_supervised(data, lag=1):
23     df = DataFrame(data)
24     columns = [df.shift(i) for i in range(1, lag+1)]
25     columns.append(df)
26     df = concat(columns, axis=1)
27     df = df.drop(0)
28     return df
29
30 # create a differenced series
31 def difference(dataset, interval=1):
32     diff = list()
33     for i in range(interval, len(dataset)):
34         value = dataset[i] - dataset[i - interval]
35         diff.append(value)
36     return Series(diff)
37
38 # scale train and test data to [-1, 1]
39 def scale(train, test):
40     # fit scaler
```

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```

41 scaler = MinMaxScaler(feature_range=(-1, 1))
42 scaler = scaler.fit(train)
43 # transform train
44 train = train.reshape(train.shape[0], train.shape[1])
45 train_scaled = scaler.transform(train)
46 # transform test
47 test = test.reshape(test.shape[0], test.shape[1])
48 test_scaled = scaler.transform(test)
49 return scaler, train_scaled, test_scaled
50
51 # inverse scaling for a forecasted value
52 def invert_scale(scaler, X, yhat):
53     new_row = [x for x in X] + [yhat]
54     array = numpy.array(new_row)
55     array = array.reshape(1, len(array))
56     inverted = scaler.inverse_transform(array)
57     return inverted[0, -1]
58
59 # evaluate the model on a dataset, returns RMSE in transformed units
60 def evaluate(model, raw_data, scaled_dataset, scaler, offset, batch_size):
61     # separate
62     X, y = scaled_dataset[:,0:-1], scaled_dataset[:, -1]
63     # forecast dataset
64     output = model.predict(X, batch_size=batch_size)
65     # invert data transforms on forecast
66     predictions = list()
67     for i in range(len(output)):
68         yhat = output[i,0]
69         # invert scaling
70         yhat = invert_scale(scaler, X[i], yhat)
71         # invert differencing
72         yhat = yhat + raw_data[i]
73         # store forecast
74         predictions.append(yhat)
75     # report performance
76     rmse = sqrt(mean_squared_error(raw_data[1:], predictions))
77     return rmse
78
79 # fit an MLP network to training data
80 def fit(train, test, raw, scaler, batch_size, nb_epoch, neurons):
81     X, y = train[:, 0:-1], train[:, -1]
82     # prepare model
83     model = Sequential()
84     model.add(Dense(neurons, activation='relu', input_dim=X.shape[1]))
85     model.add(Dense(1))
86     model.compile(loss='mean_squared_error', optimizer='adam')
87     # fit model

```

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```

88 train_rmse, test_rmse = list(), list()
89 for i in range(nb_epoch):
90     model.fit(X, y, epochs=1, batch_size=batch_size, verbose=0, shuffle=False)
91     # evaluate model on train data
92     raw_train = raw[-(len(train)+len(test)+1):-len(test)]
93     train_rmse.append(evaluate(model, raw_train, train, scaler, 0, batch_size))
94     # evaluate model on test data
95     raw_test = raw[-(len(test)+1):]
96     test_rmse.append(evaluate(model, raw_test, test, scaler, 0, batch_size))
97 history = DataFrame()
98 history['train'], history['test'] = train_rmse, test_rmse
99 return history
100
101 # run diagnostic experiments
102 def run():
103     # config
104     repeats = 10
105     n_batch = 4
106     n_epochs = 1000
107     n_neurons = 1
108     n_lag = 1
109     # load dataset
110     series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True)
111     # transform data to be stationary
112     raw_values = series.values
113     diff_values = difference(raw_values, 1)
114     # transform data to be supervised learning
115     supervised = timeseries_to_supervised(diff_values, n_lag)
116     supervised_values = supervised.values[n_lag:,:]
117     # split data into train and test-sets
118     train, test = supervised_values[0:-12], supervised_values[-12:]
119     # transform the scale of the data
120     scaler, train_scaled, test_scaled = scale(train, test)
121     # fit and evaluate model
122     train_trimmed = train_scaled[2:, :]
123     # run diagnostic tests
124     for i in range(repeats):
125         history = fit(train_trimmed, test_scaled, raw_values, scaler, n_batch, n_epochs, n_neurons)
126         pyplot.plot(history['train'], color='blue')
127         pyplot.plot(history['test'], color='orange')
128         print('%d) TrainRMSE=%f, TestRMSE=%f' % (i, history['train'].iloc[-1], history['test'].iloc[-1]))
129     pyplot.savefig('diagnostic_epochs.png')
130
131 # entry point
132 run()

```

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE


Running the diagnostic prints the final train and test RMSE for each run. More interesting is the final

Get Your Start in Machine Learning

The line plot shows the train RMSE (blue) and test RMSE (orange) after each training epoch.

In this case, the diagnostic plot shows little difference in train and test RMSE after about 400 training epochs. Both train and test performance level out on a near flat line.

This rapid leveling out suggests the model is reaching capacity and may benefit from more information in terms of lag observations or additional neurons.



Vary Hidden Layer Neurons

Get Your Start in Machine Learning

In this section, we will look at varying the number of neurons in the single hidden layer.

Increasing the number of neurons can increase the learning capacity of the network at the risk of overfitting the training data.

We will explore increasing the number of neurons from 1 to 5 and fit the network for 1000 epochs.

The differences in the experiment script are listed below.

```
1 # load dataset
2 series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True, date_parser=parser)
3 # experiment
4 repeats = 30
5 results = DataFrame()
6 lag = 1
7 epochs = 1000
8 # vary neurons
9 neurons = [1, 2, 3, 4, 5]
10 for n in neurons:
11     results[str(n)] = experiment(repeats, series, epochs, lag, n)
12 # summarize results
13 print(results.describe())
14 # save boxplot
15 results.boxplot()
16 pyplot.savefig('boxplot_neurons.png')
```

Running the experiment prints summary statistics for each configuration.

Looking at the average performance, it suggests a decrease of test RMSE with an increase in the number of neurons.

The best results appear to be with 3 neurons.

	1	2	3	4	5
count	30.000000	30.000000	30.000000	30.000000	30.000000
mean	105.107026	102.836520	92.675912	94.889952	96.577617
std	23.130824	20.102353	10.266732	9.751318	6.421356
min	86.565630	84.199871	83.388967	84.385293	87.208454
25%	88.035396	89.386670	87.643954	89.154866	89.961809
50%	90.084895	91.488484	90.670565	91.204303	96.717739
75%	136.145248	104.416518	93.117926	100.228730	101.969331
max	143.428154	140.923087	136.883946	135.891663	106.797563

A box and whisker plot is also created to summarize and compare the distributions of results.

Get Your Start in Machine Learning

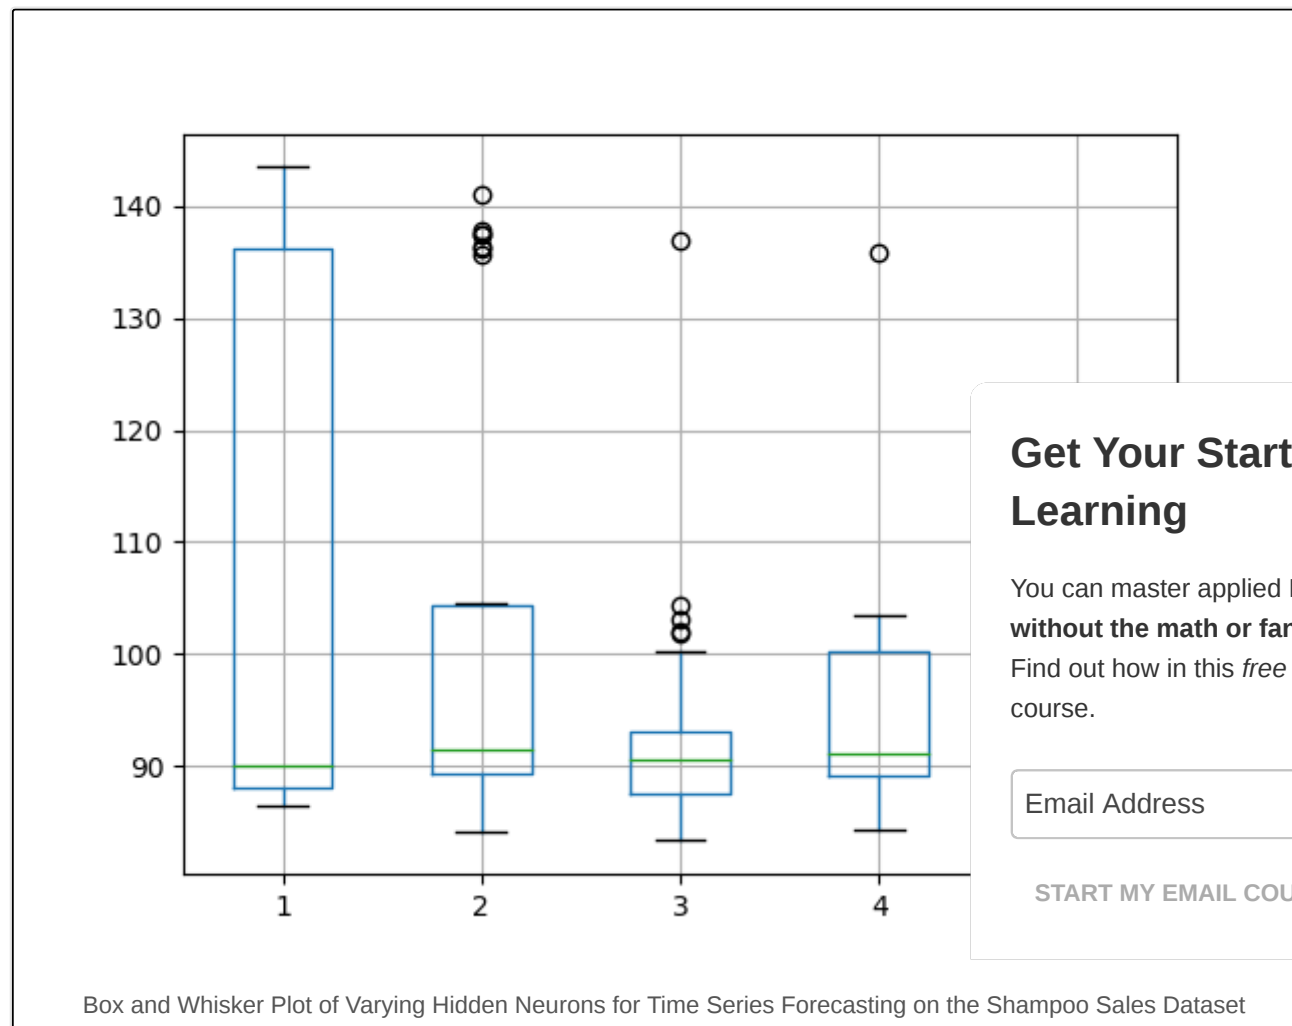
You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

The plot confirms the suggestion of 3 neurons performing well compared to the other configurations and suggests in addition that the spread of results is also smaller. This may indicate a more stable configuration.



Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree**. Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Again, we can dive a little deeper by reviewing diagnostics of the chosen configuration of 3 neurons fit for 1000 epochs.

The changes to the diagnostic script are limited to the `run()` function and listed below.

```
1 # run diagnostic experiments
2 def run():
```

Get Your Start in Machine Learning

```

3  # config
4  repeats = 10
5  n_batch = 4
6  n_epochs = 1000
7  n_neurons = 3
8  n_lag = 1
9  # load dataset
10 series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True, date_parser=parser)
11 # transform data to be stationary
12 raw_values = series.values
13 diff_values = difference(raw_values, 1)
14 # transform data to be supervised learning
15 supervised = timeseries_to_supervised(diff_values, n_lag)
16 supervised_values = supervised.values[n_lag:,:]
17 # split data into train and test-sets
18 train, test = supervised_values[0:-12], supervised_values[-12:]
19 # transform the scale of the data
20 scaler, train_scaled, test_scaled = scale(train, test)
21 # fit and evaluate model
22 train_trimmed = train_scaled[2:, :]
23 # run diagnostic tests
24 for i in range(repeats):
25     history = fit(train_trimmed, test_scaled, raw_values, scaler, n_batch, n_epochs,
26                 pyplot.plot(history['train'], color='blue')
27                 pyplot.plot(history['test'], color='orange')
28                 print('%d) TrainRMSE=%f, TestRMSE=%f' % (i, history['train'].iloc[-1], history['
29                 pyplot.savefig('diagnostic_neurons.png')

```

Running the diagnostic script provides a line plot of train and test RMSE for each training epoch.

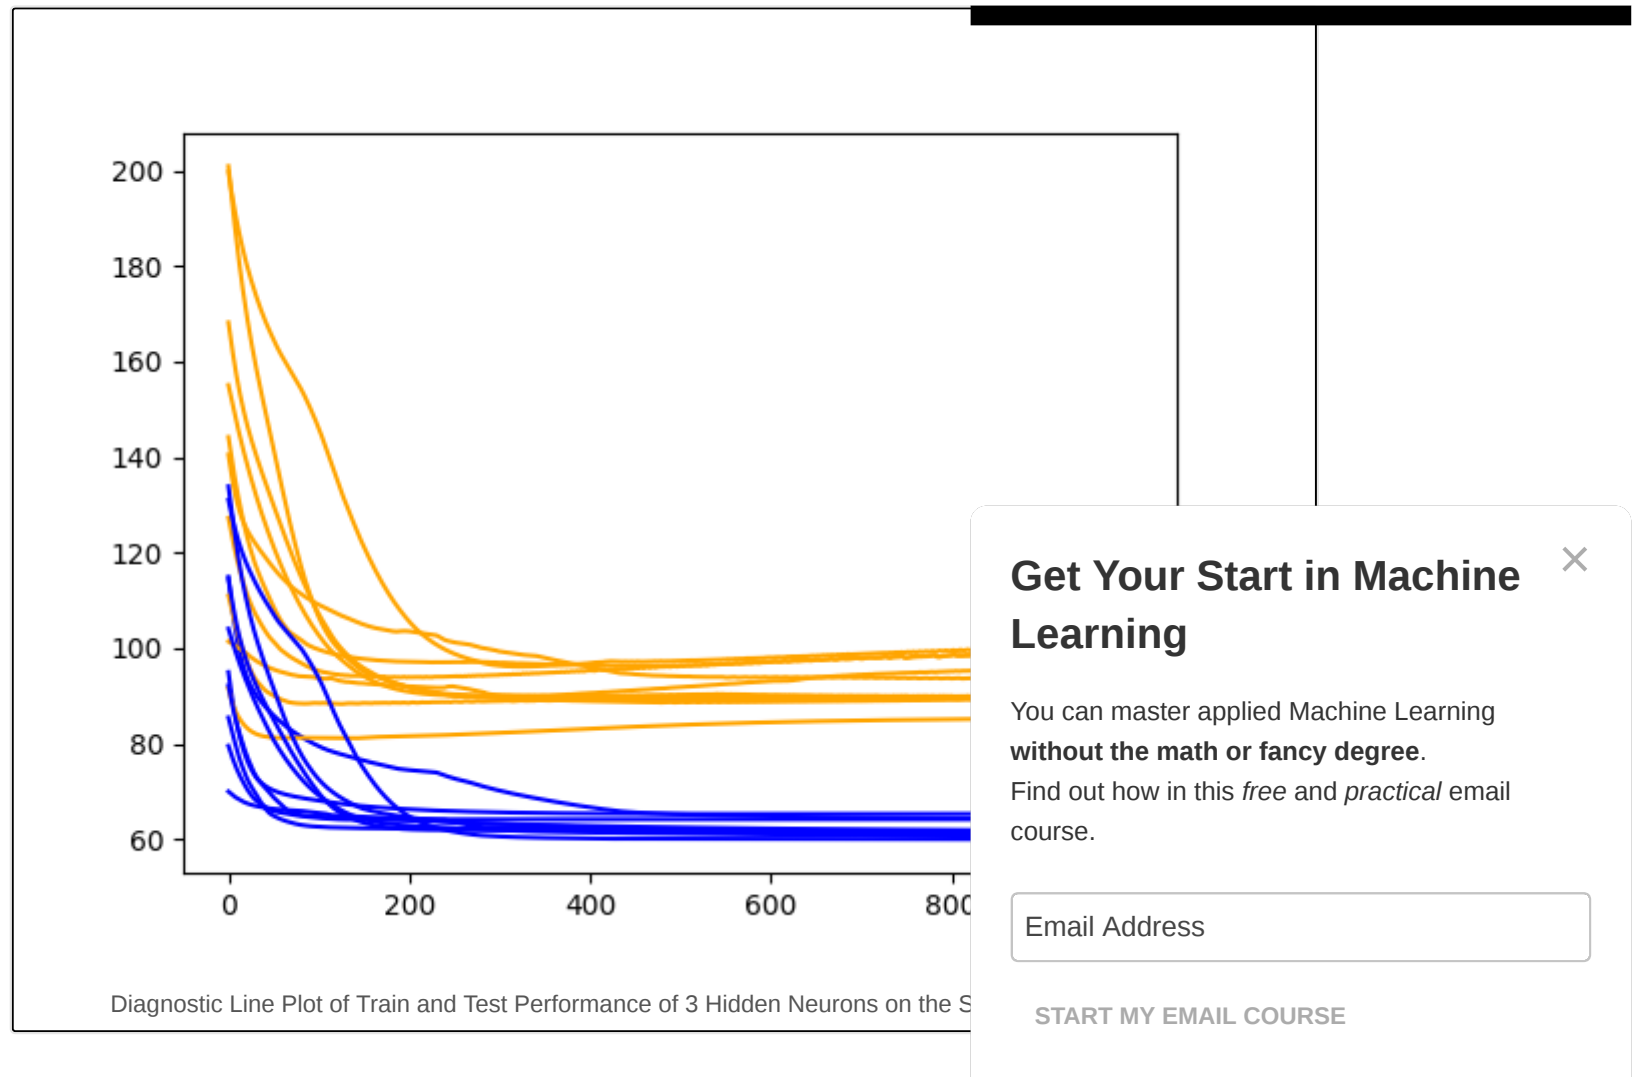
The diagnostics suggest a flattening out of model skill, perhaps around 400 epochs. The plot also suggests a slight increase in test RMSE over the last 500 training epochs, but not a strong increase in training RMSE.

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning



Vary Hidden Layer Neurons with Lag

In this section, we will look at increasing the lag observations as input, whilst at the same time increasing the capacity of the network.

Increased lag observations will automatically scale the number of input neurons. For example, 3 lag observations as input will result in 3 input neurons.

The added input will require additional capacity in the network. As such, we will also scale the number of neurons in the one hidden layer with the number of lag observations used as input.

Get Your Start in Machine Learning

We will use odd numbers of lag observations as input from 1, 3, 5, and 7 and use the same number of neurons respectively.

The change to the number of inputs affects the total number of training patterns during the conversion of the time series data to a supervised learning problem. As such, the batch size was reduced from 4 to 2 for all experiments in this section.

A total of 1000 training epochs are used in each experimental run.

The changes from the base experiment script are limited to the `experiment()` function and the running of the experiment, listed below.

```

1 # run a repeated experiment
2 def experiment(repeats, series, epochs, lag, neurons):
3     # transform data to be stationary
4     raw_values = series.values
5     diff_values = difference(raw_values, 1)
6     # transform data to be supervised learning
7     supervised = timeseries_to_supervised(diff_values, lag)
8     supervised_values = supervised.values[lag,:]
9     # split data into train and test-sets
10    train, test = supervised_values[0:-12], supervised_values[-12:]
11    # transform the scale of the data
12    scaler, train_scaled, test_scaled = scale(train, test)
13    # run experiment
14    error_scores = list()
15    for r in range(repeats):
16        # fit the model
17        batch_size = 2
18        model = fit_model(train_scaled, batch_size, epochs, neurons)
19        # forecast test dataset
20        test_resaped = test_scaled[:,0:-1]
21        output = model.predict(test_resaped, batch_size=batch_size)
22        predictions = list()
23        for i in range(len(output)):
24            yhat = output[i,0]
25            X = test_scaled[i, 0:-1]
26            # invert scaling
27            yhat = invert_scale(scaler, X, yhat)
28            # invert differencing
29            yhat = inverse_difference(raw_values, yhat, len(test_scaled)+1-i)
30            # store forecast
31            predictions.append(yhat)
32        # report performance
33        rmse = sqrt(mean_squared_error(raw_values[-12:], predictions))
34        print('%d) Test RMSE: %.3f' % (r+1, rmse))
35        error_scores.append(rmse)

```

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```

36     return error_scores
37
38 # load dataset
39 series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True, date_parser=parser)
40 # experiment
41 repeats = 30
42 results = DataFrame()
43 epochs = 1000
44 # vary neurons
45 neurons = [1, 3, 5, 7]
46 for n in neurons:
47     results[str(n)] = experiment(repeats, series, epochs, n, n)
48 # summarize results
49 print(results.describe())
50 # save boxplot
51 results.boxplot()
52 pyplot.savefig('boxplot_neurons_lag.png')

```

Running the experiment summarizes the results using descriptive statistics for each configuration.

The results suggest that all increases in lag input variables with increases with hidden neurons decrease

Of note is the 1 neuron and 1 input configuration, which compared to the results from the previous section shows a smaller standard deviation.

It is possible that the decrease in performance is related to the smaller batch size and that the results may be biased. I will tease this out.

	1	3	5	7
count	30.000000	30.000000	30.000000	30.000000
mean	105.465038	109.447044	158.894730	147.024776
std	20.827644	15.312300	43.177520	22.717514
min	89.909627	77.426294	88.515319	95.801699
25%	92.187690	102.233491	125.008917	132.335683
50%	92.587411	109.506480	166.438582	145.078842
75%	135.386125	118.635143	189.457325	166.329000
max	139.941789	144.700754	232.962778	186.185471

A box and whisker plot of the distribution of results was also created allowing configurations to be compared.

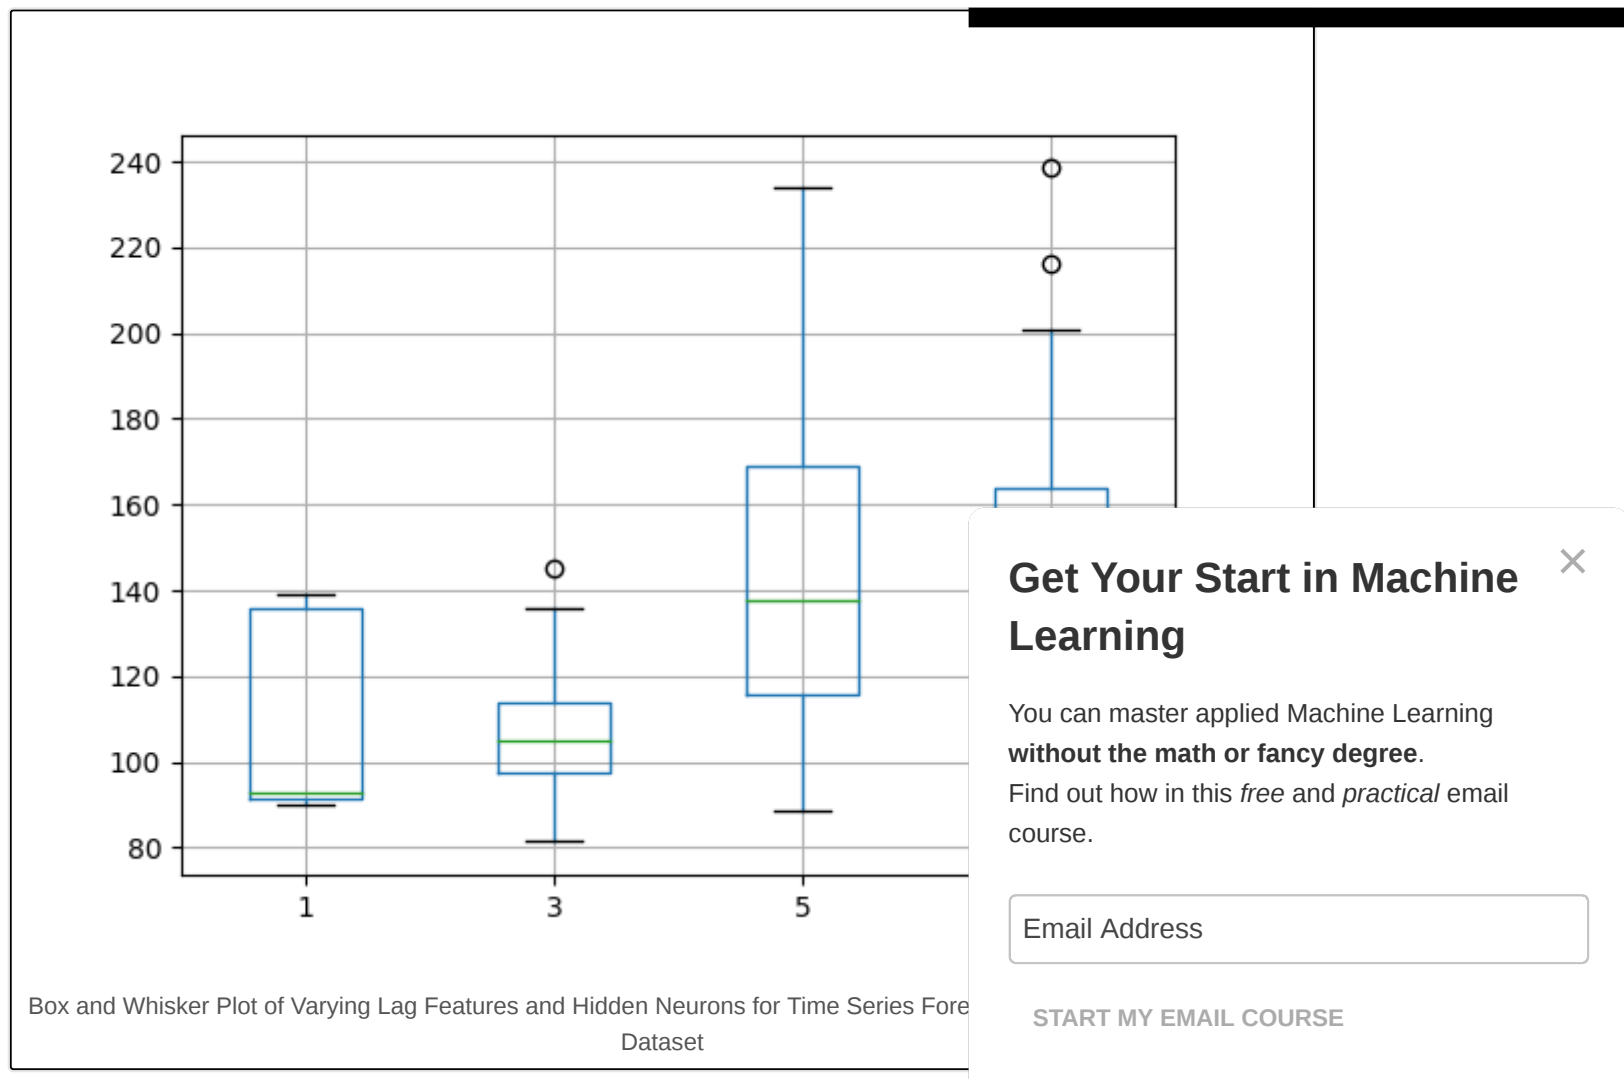
Interestingly, the use of 3 neurons and 3 input variables shows a tighter spread compared to the other configurations. This is similar to the observation from 3 neurons and 1 input variable seen in the previous section.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree**. Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

We can also use diagnostics to tease out how the dynamics of the model might have changed while fitting the model.

The results for 3-lags/3-neurons are interesting and we will investigate them further.

The changes to the diagnostic script are confined to the `run()` function.

```
1 # run diagnostic experiments
2 def run():
3     # config
```

Get Your Start in Machine Learning

```
4 repeats = 10
5 n_batch = 2
6 n_epochs = 1000
7 n_neurons = 3
8 n_lag = 3
9 # load dataset
10 series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True, date_parser=parser)
11 # transform data to be stationary
12 raw_values = series.values
13 diff_values = difference(raw_values, 1)
14 # transform data to be supervised learning
15 supervised = timeseries_to_supervised(diff_values, n_lag)
16 supervised_values = supervised.values[n_lag:,:]
17 # split data into train and test-sets
18 train, test = supervised_values[0:-12:], supervised_values[-12:]
19 # transform the scale of the data
20 scaler, train_scaled, test_scaled = scale(train, test)
21 # fit and evaluate model
22 train_trimmed = train_scaled[2:, :]
23 # run diagnostic tests
24 for i in range(repeats):
25     history = fit(train_trimmed, test_scaled, raw_values, scaler, n_batch, n_epochs,
26     pyplot.plot(history['train'], color='blue')
27     pyplot.plot(history['test'], color='orange')
28     print('%d) TrainRMSE=%f, TestRMSE=%f' % (i, history['train'].iloc[-1], history['
29     pyplot.savefig('diagnostic_neurons_lag.png')
```

Running the diagnostics script creates a line plot showing the train and test RMSE after each training

The results suggest good learning during the first 500 epochs and perhaps overfitting in the remaining trend and the train RMSE showing a decreasing trend.


Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning



Review of Results

We have covered a lot of ground in this tutorial. Let's review.

- **Epochs.** We looked at how model skill varied with the number of training epochs and found that 1000 might be a good starting point.
- **Neurons.** We looked at varying the number of neurons in the hidden layer and found that 3 neurons might be a good configuration.

Get Your Start in Machine Learning

- **Lag Inputs.** We looked at varying the number of lag observations as inputs whilst at the same time increasing the number of neurons in the hidden layer and found that results generally got worse, but again, 3 neurons in the hidden layer shows interest. Poor results may have been related to the change of batch size from 4 to 2 compared to other experiments.

The results suggest using a 1 lag input, 3 neurons in the hidden layer, and fit for 1000 epochs as a first-cut model configuration.

This can be improved upon in many ways; the next section lists some ideas.

Extensions

This section lists extensions and follow-up experiments you might like to explore.

- **Shuffle vs No Shuffle.** No shuffling was used, which is abnormal. Develop an experiment to compare shuffling to no shuffling of the training set when fitting the model for time series forecasting.
- **Normalization Method.** Data was rescaled to -1 to 1, typical for a tanh activation function, not u rescaling, such as 0-1 normalization and standardization and the impact on model performance.
- **Multiple Layers.** Explore the use of multiple hidden layers to add network capacity to learn more complex patterns.
- **Feature Engineering.** Explore the use of additional features, such as an error time series and e

Also, check out the post:

- [How To Improve Deep Learning Performance](#)

Did you try any of these extensions?

Post your results in the comments below.

Summary

In this tutorial, you discovered how to use systematic experiments to explore the configuration of a multilayer perceptron for time series forecasting and develop a first-cut model.

Specifically, you learned:

- How to develop a robust test harness for evaluating MLP models for time series forecasting.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE


Get Your Start in Machine Learning

- How to systematically evaluate training epochs, hidden layer neurons, and lag inputs.
- How to use diagnostics to help interpret results and suggest follow-up experiments.

Do you have any questions about this tutorial?

Ask your questions in the comments below and I will do my best to answer.

Frustrated With Your Progress In Deep Learning?



What If You Could Develop A Network in Minutes

...with just a few lines of code

Discover how in my new Ebook: [Deep Learning with Python](#)

It covers **self-study tutorials** and **end-to-end projects** including *Multilayer Perceptrons, Convolutional Nets and Recurrent Neural Networks*

Finally Bring Deep Learning To Your Own Projects

Skip the Academics. Just learn how to build a network in minutes.

[Click to learn more](#)

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



About Jason Brownlee

[Get Your Start in Machine Learning](#)



Dr. Jason Brownlee is a husband, proud father, academic researcher, author, professional developer and a machine learning practitioner. He is dedicated to helping developers get started and get good at applied machine learning. [Learn more.](#)

[View all posts by Jason Brownlee](#) →

< Instability of Online Learning for Stateful LSTM for Time Series Forecasting

How to Use Dropout with LSTM Networks for Time Series Forecasting >

20 Responses to *Exploratory Configuration of a Multilayer Perceptron Network for Time Series Forecasting*



elina April 26, 2017 at 11:00 pm #

what is mean by transforming observations to have specific scale? can u more elaborate it with



Jason Brownlee April 27, 2017 at 8:40 am #

Yes, normalize each column to the range 0-1.



Kunpeng Zhang April 28, 2017 at 3:12 am #

Hi Jason,

I have a question regarding the prediction result evaluation.

In my code, I get MSE MAE MAPE calculated by keras as follows:

```
model.compile(loss='mse', optimizer=optimizer, metrics=['mae', 'mape'])
```

...

Epoch 150/150

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

REPLY ↩

Get Your Start in Machine Learning

0s – loss: 0.0038 – mean_absolute_error: 0.0455 – mean_absolute_percentage_error: 164095.5176
mse=0.003538, mae=0.043654, mape=58238.251235

But when I compute these values by the code:

```
mse = mean_squared_error(test_Y, predicted_output)
rmse = math.sqrt(mse)
mae = mean_absolute_error(test_Y, predicted_output)
mape = np.mean(np.abs(np.divide(np.subtract(test_Y, predicted_output), test_Y))) * 100
RMSE: 14.992
MAE: 10.462
MAPE: 2.208
```

The two results are very different from each other. What happened?



Jason Brownlee April 28, 2017 at 7:53 am #

That is interesting. I'm not sure what is going on here.

I would trust the manual results. I have not had any issues like this myself, my epoch scores always

Consider preparing a small self-contained example and posting it as a bug to the Keras project:

<https://github.com/fchollet/keras>



Kunpeng Zhang April 29, 2017 at 11:46 pm #

I give it a try on your example.

<http://machinelearningmastery.com/time-series-prediction-lstm-recurrent-neural-networks-python-keras/>

Change

```
model.compile(loss='mean_squared_error', optimizer='adam')
```

to

```
model.compile(loss='mean_squared_error', optimizer='adam', metrics=['mae', 'mape'])
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Got the same result.

0s – loss: 0.0019 – mean_absolute_error: 0.0343 – mean_absolute_percentage_error: 620693.8651

Epoch 96/100

0s – loss: 0.0020 – mean_absolute_error: 0.0351 – mean_absolute_percentage_error: 326543.9207

Epoch 97/100

0s – loss: 0.0019 – mean_absolute_error: 0.0345 – mean_absolute_percentage_error: 488762.9108

Epoch 98/100

0s – loss: 0.0019 – mean_absolute_error: 0.0345 – mean_absolute_percentage_error: 514091.1566

Epoch 99/100

0s – loss: 0.0019 – mean_absolute_error: 0.0345 – mean_absolute_percentage_error: 531419.0410

Epoch 100/100

0s – loss: 0.0019 – mean_absolute_error: 0.0341 – mean_absolute_percentage_error: 454424.3737

Train Score: 22.34 RMSE

Test Score: 45.66 RMSE

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Hans April 28, 2017 at 6:12 pm #

What is the most robust sign for overfitting?



Jason Brownlee April 29, 2017 at 7:22 am #

Skill on the test set is worse than the training set.



Hans April 30, 2017 at 3:50 pm #

Is there a list of overfitting signs available?

Alternative:

If I want to prepare such a list, how should it look like?

REPLY ↩

Get Your Start in Machine Learning

1. Skill on the test set is worse than the training set.
2. ?
3. ?
4. ?
5. ?



Hans April 30, 2017 at 3:51 pm #

REPLY ↩

Is there a function available for our code, which 'alerts' overfitting?



Hans April 30, 2017 at 4:00 pm #

I monitor every result and multiple parameters of your examples in a Sqlite database.
Therefore I could easelly tagging overfitting.
For example if a test result is worse then the training set.
Could there be more relations of parameters and results to monitor for overfitting in general?



Jason Brownlee May 1, 2017 at 5:55 am #

It is a trend you are seeking for overfitting, not necessarily one result being worse.
<https://en.wikipedia.org/wiki/Overfitting>



Jason Brownlee May 1, 2017 at 5:54 am #

No.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning



Jason Brownlee May 1, 2017 at 5:54 am #

REPLY ↩

Nope. The first is all you need.

Instead, you develop a list of 100s of ways to address overfitting/early convergence.



Hans May 1, 2017 at 9:21 am #

A)

Does “Skill on the test set is worse than the training set”
mean trainRmse is less than testRmse?

In my baseline-test my trainRmse is always zero and the testRmse is always higher? Is

B)

Could we say that a higher variance of test values is an indicator for overfitting?



Jason Brownlee May 2, 2017 at 5:56 am #

This post might make things clearer for you Hans:

<http://machinelearningmastery.com/overfitting-and-underfitting-with-machine-learning-al>

Get Your Start in Machine Learning



You can master applied Machine Learning
without the math or fancy degree.

Find out how in this *free* and *practical* email
course.

START MY EMAIL COURSE



Hans April 30, 2017 at 4:52 pm #

REPLY ↩

Could we assume that the best performing epochs, neurons and lag inputs found with this MLP setup are also the best performing ones in a LSTM model or any other model?

Get Your Start in Machine Learning



Jason Brownlee May 1, 2017 at 5:56 am #

REPLY ↩

No, generally findings are not transferable to other problems or other algorithms.



Hans May 1, 2017 at 8:59 pm #

REPLY ↩

Is it possible to save a trained model on HD, to predict unseen data later, in a shorter time?



Magnus May 12, 2017 at 11:22 pm #

REPLY ↩

Would it be possible to achieve the same study using GridSearchCV, like you did in another post?



Jason Brownlee May 13, 2017 at 6:15 am #

No, we must use walk-forward validation to evaluate time series models correctly:

<http://machinelearningmastery.com/backtest-machine-learning-models-time-series-forecasting/>

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Leave a Reply

Get Your Start in Machine Learning

Name (required)

Email (will not be published) (required)

Website

Welcome to Machine Learning Mastery



Hi, I'm Dr. Jason Brownlee.

My goal is to make practitioners like YOU awesome at applied machine learning.

[Read More](#)

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

[START MY EMAIL COURSE](#)


Finally Get Started With Deep Learning

Sick of the fancy math and need for super computers?
Looking for step-by-step tutorials?

[Get Your Start in Machine Learning](#)

Want end-to-end projects?

Get Started With Deep Learning in Python Today!



POPULAR



Time Series Prediction with LSTM Recurrent Neural Networks in Python with Keras

JULY 21, 2016



Your First Machine Learning Project in Python Step-By-Step

JUNE 10, 2016



Develop Your First Neural Network in Python With Keras Step-By-Step

MAY 24, 2016



Sequence Classification with LSTM Recurrent Neural Networks in Python with Keras

JULY 26, 2016

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning



How to Setup a Python Environment for Machine Learning and Deep Learning with Anaconda

MARCH 13, 2017



Time Series Forecasting with the Long Short-Term Memory Network in Python

APRIL 7, 2017



Multi-Class Classification Tutorial with the Keras Deep Learning Library

JUNE 2, 2016



Regression Tutorial with the Keras Deep Learning Library in Python

JUNE 9, 2016



Multivariate Time Series Forecasting with LSTMs in Keras

AUGUST 14, 2017



How to Implement the Backpropagation Algorithm From Scratch In Python

NOVEMBER 7, 2016

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

© 2017 Machine Learning Mastery. All Rights Reserved.

[Privacy](#) | [Contact](#) | [About](#)

Get Your Start in Machine Learning