

[Start Here](#)[Blog](#)[Books](#)[About](#)[Contact](#)

Search...



How to Convert a Time Series to a Supervised Learning Problem in Python

by **Jason Brownlee** on May 8, 2017 in **Time Series**



Machine learning methods like deep learning can be used for time series forecasting.

Before machine learning can be used, time series forecasting problems must be re-framed as supervised learning problems with input and output sequences.

In this tutorial, you will discover how to transform univariate and multivariate time series forecasting problems into supervised learning problems with machine learning algorithms.

After completing this tutorial, you will know:

- How to develop a function to transform a time series dataset into a supervised learning dataset.
- How to transform univariate time series data for machine learning.

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

- How to transform multivariate time series data for machine learning.

Let's get started.



How to Convert a Time Series to a Supervised Learning Problem in Python
Photo by [Quim Gil](#), some rights reserved.

Time Series vs Supervised Learning

Before we get started, let's take a moment to better understand the form of time series and supervised learning.

A time series is a sequence of numbers that are ordered by a time index. This can be thought of as a sequence of observations over time.

For example:

1	0
2	1

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree**. Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

```
3 2
4 3
5 4
6 5
7 6
8 7
9 8
10 9
```

Get Your Start in Machine Learning

A supervised learning problem is comprised of input patterns (X) and output patterns (y), such that an algorithm can learn how to predict the output patterns from the input patterns.

For example:

```
1 X, y
2 1, 2
3 2, 3
4 3, 4
5 4, 5
6 5, 6
7 6, 7
8 7, 8
9 8, 9
```

For more on this topic, see the post:

- [Time Series Forecasting as Supervised Learning](#)

Pandas `shift()` Function

A key function to help transform time series data into a supervised learning problem is the Pandas `shift()` function.

Given a `DataFrame`, the `shift()` function can be used to create copies of columns that are pushed forward or pulled back (rows of NaN values added to the end).

This is the behavior required to create columns of lag observations as well as columns of forecast observations in supervised learning format.

Let's look at some examples of the `shift` function in action.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

We can define a mock time series dataset as a sequence of 10 numbers, in this case a single column in a DataFrame as follows:

```
1 from pandas import DataFrame
2 df = DataFrame()
3 df['t'] = [x for x in range(10)]
4 print(df)
```

Get Your Start in Machine Learning

Running the example prints the time series data with the row indices for each observation.

```
1      t
2  0  0
3  1  1
4  2  2
5  3  3
6  4  4
7  5  5
8  6  6
9  7  7
10 8  8
11 9  9
```

We can shift all the observations down by one time step by inserting one new row at the top. Because the new row has no data, we can use NaN to represent “no data”.

The shift function can do this for us and we can insert this shifted column next to our original series.

```
1 from pandas import DataFrame
2 df = DataFrame()
3 df['t'] = [x for x in range(10)]
4 df['t-1'] = df['t'].shift(1)
5 print(df)
```

Running the example gives us two columns in the dataset. The first with the original observations and the second with the shifted observations.

We can see that shifting the series forward one time step gives us a primitive supervised learning problem. We can ignore the column of row labels. The first row would have to be discarded because of the NaN value in the second column (input or X) and the value of 1 in the first column (output or y).

```
1      t  t-1
2  0  0 NaN
3  1  1 0.0
4  2  2 1.0
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

```

5 3 3 2.0
6 4 4 3.0
7 5 5 4.0
8 6 6 5.0
9 7 7 6.0
10 8 8 7.0
11 9 9 8.0

```

Get Your Start in Machine Learning

We can see that if we can repeat this process with shifts of 2, 3, and more, how we could create long input sequences (X) that can be used to forecast an output value (y).

The shift operator can also accept a negative integer value. This has the effect of pulling the observations up by inserting new rows at the end. Below is an example:

```

1 from pandas import DataFrame
2 df = DataFrame()
3 df['t'] = [x for x in range(10)]
4 df['t+1'] = df['t'].shift(-1)
5 print(df)

```

Running the example shows a new column with a NaN value as the last value.

We can see that the forecast column can be taken as an input (X) and the second as an output value (y). That is the input value of 0 can be used to forecast the output value of 1.

```

1      t  t+1
2  0  0  1.0
3  1  1  2.0
4  2  2  3.0
5  3  3  4.0
6  4  4  5.0
7  5  5  6.0
8  6  6  7.0
9  7  7  8.0
10 8  8  9.0
11 9  9  NaN

```

Technically, in time series forecasting terminology the current time (t) and future times ($t+1$, $t+n$) are used to make forecasts.

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree**. Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

We can see how positive and negative shifts can be used to create a new DataFrame from a time series with sequences of input and output patterns for a supervised learning problem.

This permits not only classical $X \rightarrow y$ prediction, but also $X \rightarrow Y$ where both input and output can be

Get Your Start in Machine Learning

Further, the shift function also works on so-called multivariate time series problems. That is where instead of having one set of observations for a time series, we have multiple (e.g. temperature and pressure). All variates in the time series can be shifted forward or backward to create multivariate input and output sequences. We will explore this more later in the tutorial.

The `series_to_supervised()` Function

We can use the `shift()` function in Pandas to automatically create new framings of time series problems given the desired length of input and output sequences.

This would be a useful tool as it would allow us to explore different framings of a time series problem with machine learning algorithms to see which might result in better performing models.

In this section, we will define a new Python function named `series_to_supervised()` that takes a univariate or multivariate time series and frames it as a supervised learning dataset.

The function takes four arguments:

- **data**: Sequence of observations as a list or 2D NumPy array. Required.
- **n_in**: Number of lag observations as input (X). Values may be between $[1..\text{len}(\text{data})]$ Optional. Defaults to 1.
- **n_out**: Number of observations as output (y). Values may be between $[0..\text{len}(\text{data})-1]$. Optional. Defaults to 1.
- **dropnan**: Boolean whether or not to drop rows with NaN values. Optional. Defaults to True.

The function returns a single value:

- **return**: Pandas DataFrame of series framed for supervised learning.

The new dataset is constructed as a DataFrame, with each column suitably named both by variable and time step. This allows for a variety of different time step sequence type forecasting problems from a given univariate or multivariate time series.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Once the DataFrame is returned, you can decide how to split the rows of the returned DataFrame into X and y components for supervised learning any way you wish.

The function is defined with default parameters so that if you call it with just your data, it will construct

Get Your Start in Machine Learning

The function is confirmed to be compatible with Python 2 and Python 3.

The complete function is listed below, including function comments.

```

1 from pandas import DataFrame
2 from pandas import concat
3
4 def series_to_supervised(data, n_in=1, n_out=1, dropnan=True):
5     """
6     Frame a time series as a supervised learning dataset.
7     Arguments:
8         data: Sequence of observations as a list or NumPy array.
9         n_in: Number of lag observations as input (X).
10        n_out: Number of observations as output (y).
11        dropnan: Boolean whether or not to drop rows with NaN values.
12    Returns:
13        Pandas DataFrame of series framed for supervised learning.
14    """
15    n_vars = 1 if type(data) is list else data.shape[1]
16    df = DataFrame(data)
17    cols, names = list(), list()
18    # input sequence (t-n, ... t-1)
19    for i in range(n_in, 0, -1):
20        cols.append(df.shift(i))
21        names += [('var%d(t-%d)' % (j+1, i)) for j in range(n_vars)]
22    # forecast sequence (t, t+1, ... t+n)
23    for i in range(0, n_out):
24        cols.append(df.shift(-i))
25        if i == 0:
26            names += [('var%d(t)' % (j+1)) for j in range(n_vars)]
27        else:
28            names += [('var%d(t+%d)' % (j+1, i)) for j in range(n_vars)]
29    # put it all together
30    agg = concat(cols, axis=1)
31    agg.columns = names
32    # drop rows with NaN values
33    if dropnan:
34        agg.dropna(inplace=True)
35    return agg

```

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Can you see obvious ways to make the function more robust or more readable?
Please let me know in the comments below.

Now that we have the whole function, we can explore how it may be used.

Get Your Start in Machine Learning

One-Step Univariate Forecasting

It is standard practice in time series forecasting to use lagged observations (e.g. $t-1$) as input variables to forecast the current time step (t).

This is called one-step forecasting.

The example below demonstrates a one lag time step ($t-1$) to predict the current time step (t).

```

1 from pandas import DataFrame
2 from pandas import concat
3
4 def series_to_supervised(data, n_in=1, n_out=1, dropnan=True):
5     """
6     Frame a time series as a supervised learning dataset.
7     Arguments:
8         data: Sequence of observations as a list or NumPy array.
9         n_in: Number of lag observations as input (X).
10        n_out: Number of observations as output (y).
11        dropnan: Boolean whether or not to drop rows with NaN values.
12    Returns:
13        Pandas DataFrame of series framed for supervised learning.
14    """
15    n_vars = 1 if type(data) is list else data.shape[1]
16    df = DataFrame(data)
17    cols, names = list(), list()
18    # input sequence (t-n, ... t-1)
19    for i in range(n_in, 0, -1):
20        cols.append(df.shift(i))
21        names += [('var%d(t-%d)' % (j+1, i)) for j in range(n_vars)]
22    # forecast sequence (t, t+1, ... t+n)
23    for i in range(0, n_out):
24        cols.append(df.shift(-i))
25        if i == 0:
26            names += [('var%d(t)' % (j+1)) for j in range(n_vars)]
27        else:
28            names += [('var%d(t+%d)' % (j+1, i)) for j in range(n_vars)]
29    # put it all together

```

Get Your Start in Machine Learning ×

You can master applied Machine Learning
without the math or fancy degree.

Find out how in this *free* and *practical* email
course.

START MY EMAIL COURSE


```

30     agg = concat(cols, axis=1)
31     agg.columns = names
32     # drop rows with NaN values
33     if dropnan:
34         agg.dropna(inplace=True)
35     return agg
36
37
38 values = [x for x in range(10)]
39 data = series_to_supervised(values)
40 print(data)

```

Get Your Start in Machine Learning

Running the example prints the output of the reframed time series.

	var1(t-1)	var1(t)
1		
2	1	0.0
3	2	1.0
4	3	2.0
5	4	3.0
6	5	4.0
7	6	5.0
8	7	6.0
9	8	7.0
10	9	8.0

We can see that the observations are named “var1” and that the input observation is suitably named (t-1) and the output time step is named (t).

We can also see that rows with NaN values have been automatically removed from the DataFrame.

We can repeat this example with an arbitrary number length input sequence, such as 3. This can be as an argument; for example:

```

1 data = series_to_supervised(values, 3)

```

The complete example is listed below.

```

1 from pandas import DataFrame
2 from pandas import concat
3
4 def series_to_supervised(data, n_in=1, n_out=1, dropnan=True):
5     """
6     Frame a time series as a supervised learning dataset.
7     Arguments:
8         data: Sequence of observations as a list or NumPy array.

```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree**. Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

```

9     n_in: Number of lag observations as input (X).
10    n_out: Number of observations as output (y).
11    dropnan: Boolean whether or not to drop rows with NaN values.
12    Returns:
13        Pandas DataFrame of series framed for supervised learning.
14    """
15    n_vars = 1 if type(data) is list else data.shape[1]
16    df = DataFrame(data)
17    cols, names = list(), list()
18    # input sequence (t-n, ... t-1)
19    for i in range(n_in, 0, -1):
20        cols.append(df.shift(i))
21        names += [('var%d(t-%d)' % (j+1, i)) for j in range(n_vars)]
22    # forecast sequence (t, t+1, ... t+n)
23    for i in range(0, n_out):
24        cols.append(df.shift(-i))
25        if i == 0:
26            names += [('var%d(t)' % (j+1)) for j in range(n_vars)]
27        else:
28            names += [('var%d(t+%d)' % (j+1, i)) for j in range(n_vars)]
29    # put it all together
30    agg = concat(cols, axis=1)
31    agg.columns = names
32    # drop rows with NaN values
33    if dropnan:
34        agg.dropna(inplace=True)
35    return agg
36
37
38 values = [x for x in range(10)]
39 data = series_to_supervised(values, 3)
40 print(data)

```

Again, running the example prints the reframed series. We can see that the input sequence is in the be predicted on the far right.

	var1(t-3)	var1(t-2)	var1(t-1)	var1(t)
1 3	0.0	1.0	2.0	3
3 4	1.0	2.0	3.0	4
4 5	2.0	3.0	4.0	5
5 6	3.0	4.0	5.0	6
6 7	4.0	5.0	6.0	7
7 8	5.0	6.0	7.0	8
8 9	6.0	7.0	8.0	9

Get Your Start in Machine Learning

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Multi-Step or Sequence Forecasting

A different type of forecasting problem is using past observations to forecast a sequence of future observations.

This may be called sequence forecasting or multi-step forecasting.

We can frame a time series for sequence forecasting by specifying another argument. For example, we could frame a forecast problem with an input sequence of 2 past observations to forecast 2 future observations as follows:

```
1 data = series_to_supervised(values, 2, 2)
```

The complete example is listed below:

```
1 from pandas import DataFrame
2 from pandas import concat
3
4 def series_to_supervised(data, n_in=1, n_out=1, dropnan=True):
5     """
6     Frame a time series as a supervised learning dataset.
7     Arguments:
8         data: Sequence of observations as a list or NumPy array.
9         n_in: Number of lag observations as input (X).
10        n_out: Number of observations as output (y).
11        dropnan: Boolean whether or not to drop rows with NaN values.
12    Returns:
13        Pandas DataFrame of series framed for supervised learning.
14    """
15    n_vars = 1 if type(data) is list else data.shape[1]
16    df = DataFrame(data)
17    cols, names = list(), list()
18    # input sequence (t-n, ... t-1)
19    for i in range(n_in, 0, -1):
20        cols.append(df.shift(i))
21        names += [('var%d(t-%d)' % (j+1, i)) for j in range(n_vars)]
22    # forecast sequence (t, t+1, ... t+n)
23    for i in range(0, n_out):
24        cols.append(df.shift(-i))
25        if i == 0:
26            names += [('var%d(t)' % (j+1)) for j in range(n_vars)]
27        else:
28            names += [('var%d(t+%d)' % (j+1, i)) for j in range(n_vars)]
29    # put it all together
30    agg = concat(cols, axis=1)
```

Get Your Start in Machine Learning

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

```

31     agg.columns = names
32     # drop rows with NaN values
33     if dropnan:
34         agg.dropna(inplace=True)
35     return agg
36
37
38 values = [x for x in range(10)]
39 data = series_to_supervised(values, 2, 2)
40 print(data)

```

Get Your Start in Machine Learning

Running the example shows the differentiation of input (t-n) and output (t+n) variables with the current observation (t) considered an output.

	var1(t-2)	var1(t-1)	var1(t)	var1(t+1)
1				
2	2	0.0	1.0	2
3	3	1.0	2.0	3
4	4	2.0	3.0	4
5	5	3.0	4.0	5
6	6	4.0	5.0	6
7	7	5.0	6.0	7
8	8	6.0	7.0	8

Multivariate Forecasting

Another important type of time series is called multivariate time series.

This is where we may have observations of multiple different measures and an interest in forecasting

For example, we may have two sets of time series observations obs1 and obs2 and we wish to forecast

We can call `series_to_supervised()` in exactly the same way.

For example:

```

1 from pandas import DataFrame
2 from pandas import concat
3
4 def series_to_supervised(data, n_in=1, n_out=1, dropnan=True):
5     """
6     Frame a time series as a supervised learning dataset.
7     Arguments:
8         data: Sequence of observations as a list or NumPy array.

```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

```

9         n_in: Number of lag observations as input (X).
10        n_out: Number of observations as output (y).
11        dropnan: Boolean whether or not to drop rows with NaN values.
12    Returns:
13        Pandas DataFrame of series framed for supervised learning.
14    """
15    n_vars = 1 if type(data) is list else data.shape[1]
16    df = DataFrame(data)
17    cols, names = list(), list()
18    # input sequence (t-n, ... t-1)
19    for i in range(n_in, 0, -1):
20        cols.append(df.shift(i))
21        names += [('var%d(t-%d)' % (j+1, i)) for j in range(n_vars)]
22    # forecast sequence (t, t+1, ... t+n)
23    for i in range(0, n_out):
24        cols.append(df.shift(-i))
25        if i == 0:
26            names += [('var%d(t)' % (j+1)) for j in range(n_vars)]
27        else:
28            names += [('var%d(t+%d)' % (j+1, i)) for j in range(n_vars)]
29    # put it all together
30    agg = concat(cols, axis=1)
31    agg.columns = names
32    # drop rows with NaN values
33    if dropnan:
34        agg.dropna(inplace=True)
35    return agg
36
37
38 raw = DataFrame()
39 raw['ob1'] = [x for x in range(10)]
40 raw['ob2'] = [x for x in range(50, 60)]
41 values = raw.values
42 data = series_to_supervised(values)
43 print(data)

```

Running the example prints the new framing of the data, showing an input pattern with one time step for both variables.

Again, depending on the specifics of the problem, the division of columns into X and Y components of observation of *var1* was also provided as input and only *var2* was to be predicted.

		var1(t-1)	var2(t-1)	var1(t)	var2(t)
1					
2	1	0.0	50.0	1	51
3	2	1.0	51.0	2	52

Get Your Start in Machine Learning

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

4	3	2.0	52.0	3	53
5	4	3.0	53.0	4	54
6	5	4.0	54.0	5	55
7	6	5.0	55.0	6	56
8	7	6.0	56.0	7	57
9	8	7.0	57.0	8	58
10	9	8.0	58.0	9	59

Get Your Start in Machine Learning

You can see how this may be easily used for sequence forecasting with multivariate time series by specifying the length of the input and output sequences as above.

For example, below is an example of a reframing with 1 time step as input and 2 time steps as forecast sequence.

```

1 from pandas import DataFrame
2 from pandas import concat
3
4 def series_to_supervised(data, n_in=1, n_out=1, dropnan=True):
5     """
6     Frame a time series as a supervised learning dataset.
7     Arguments:
8         data: Sequence of observations as a list or NumPy array.
9         n_in: Number of lag observations as input (X).
10        n_out: Number of observations as output (y).
11        dropnan: Boolean whether or not to drop rows with NaN values.
12    Returns:
13        Pandas DataFrame of series framed for supervised learning.
14    """
15    n_vars = 1 if type(data) is list else data.shape[1]
16    df = DataFrame(data)
17    cols, names = list(), list()
18    # input sequence (t-n, ... t-1)
19    for i in range(n_in, 0, -1):
20        cols.append(df.shift(i))
21        names += [('var%d(t-%d)' % (j+1, i)) for j in range(n_vars)]
22    # forecast sequence (t, t+1, ... t+n)
23    for i in range(0, n_out):
24        cols.append(df.shift(-i))
25        if i == 0:
26            names += [('var%d(t)' % (j+1)) for j in range(n_vars)]
27        else:
28            names += [('var%d(t+%d)' % (j+1, i)) for j in range(n_vars)]
29    # put it all together
30    agg = concat(cols, axis=1)
31    agg.columns = names
32    # drop rows with NaN values
33    if dropnan:

```

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

```

34     agg.dropna(inplace=True)
35     return agg
36
37
38 raw = DataFrame()
39 raw['ob1'] = [x for x in range(10)]
40 raw['ob2'] = [x for x in range(50, 60)]
41 values = raw.values
42 data = series_to_supervised(values, 1, 2)
43 print(data)

```

Get Your Start in Machine Learning

Running the example shows the large reframed DataFrame.

	var1(t-1)	var2(t-1)	var1(t)	var2(t)	var1(t+1)	var2(t+1)
1						
2	1	0.0	50.0	1	51	2.0
3	2	1.0	51.0	2	52	3.0
4	3	2.0	52.0	3	53	4.0
5	4	3.0	53.0	4	54	5.0
6	5	4.0	54.0	5	55	6.0
7	6	5.0	55.0	6	56	7.0
8	7	6.0	56.0	7	57	8.0
9	8	7.0	57.0	8	58	9.0

Experiment with your own dataset and try multiple different framings to see what works best.

Summary

In this tutorial, you discovered how to reframe time series datasets as supervised learning problems

Specifically, you learned:

- About the Pandas *shift()* function and how it can be used to automatically define supervised learning problems
- How to reframe a univariate time series into one-step and multi-step supervised learning problems
- How to reframe multivariate time series into one-step and multi-step supervised learning problems

Do you have any questions?

Ask your questions in the comments below and I will do my best to answer.


Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Want to Develop Time Series Forecasts with Python?



Develop Your Own Forecasts with Python

...with just a few lines of python code

Discover how in my new Ebook:

[Introduction to Time Series Forecasting With Python](#)

It covers **self-study tutorials** and **end-to-end projects** on topics like:
Loading data, visualization, modeling, algorithm tuning, and much more...

Finally Bring Time Series Forecasting to Your Own Projects

Skip the Academics. Just Results.

[Click to learn more.](#)



About Jason Brownlee

Dr. Jason Brownlee is a husband, proud father, academic researcher, author, professional developer, and data scientist. He is passionate about helping developers get started and get good at applied machine learning. [Learn more.](#)

[View all posts by Jason Brownlee](#) →

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

[START MY EMAIL COURSE](#)

[← How to Use Weight Regularization with LSTM Networks for Time Series Forecasting](#)

[Multi-step Time Series Forecasting](#)

[Time Series Forecasting](#) >

40 Responses to *How to Convert a Time Series to a Supervised Learning Problem in Python*

Get Your Start in Machine Learning

**Mikkel** May 8, 2017 at 7:07 pm #

REPLY ↩

Hi Jason, thanks for your highly relevant article 😊

I am having a hard time following the structure of the dataset. I understand the basics of $t-n$, $t-1$, t , $t+1$, $t+n$ and so forth. Although, what exactly are we describing here in the t and $t-1$ column? Is it the change over time for a specific explanatory variable? In that case, wouldn't it make more sense to transpose the data, so that the time were described in the rows rather than columns?

Also, how would you then characterise following data:

Customer_ID Month Balance

1 01 1,500
1 02 1,600
1 03 1,700
1 04 1,900
2 01 1,000
2 02 900
2 03 700
2 04 500
3 01 3,500
3 02 1,500
3 03 2,500
3 04 4,500

Let's say, that we wanna forecast their balance using supervised learning, or classify the customers as "s...

**Jason Brownlee** May 9, 2017 at 7:40 am #

Yes, it is transposing each variable, but allowing control over the length of each row back in

Get Your Start in Machine Learning

✕

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Daniel May 9, 2017 at 4:56 pm #

Hey Jason,

this is an awesome article! I was looking for that the whole time.

The only thing is I am general programming in R, so I only found something similar like your code, but I am not sure if it is the same. I have got this from <https://www.r-bloggers.com/generating-a-laglead-variables/> and it deals with lagged and leaded values. Also the output includes NA values.

```
shift1)
return(sapply(shift_by,shift, x=x))

out 0 )
out<-c(tail(x,-abs_shift_by),rep(NA,abs_shift_by))
else if (shift_by < 0 )
out<-c(rep(NA,abs_shift_by), head(x,-abs_shift_by))
else
out<-x
out
}
```

Output:

```
x df_lead2 df_lag2
1 1 3 NA
2 2 4 NA
3 3 5 1
4 4 6 2
5 5 7 3
6 6 8 4
7 7 9 5
8 8 10 6
9 9 NA 7
10 10 NA 8
```

I also tried to recompile your code in R, but it failed.

Get Your Start in Machine Learning

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree**. Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee May 10, 2017 at 8:44 am #

I would recommend contacting the authors of the R code you reference.

Get Your Start in Machine Learning



chris May 12, 2017 at 3:28 am #

REPLY ↩

Can you answer this in Go, Java, C# and COBOL as well????? Thanks, I really don't want to do anything



Jason Brownlee May 12, 2017 at 7:45 am #

REPLY ↩

I do my best to help, some need more help than others.



Lee May 9, 2017 at 11:40 pm #

Hi Jason, good article, but could be much better if you illustrated everything with some actual time series code 5 times 😊 Gripes aside, this was very timely as I'm just about to get into some time series forecasting

REPLY ↩



Jason Brownlee May 10, 2017 at 8:48 am #

Thanks for the suggestion Lee.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Christopher May 12, 2017 at 9:16 pm #



Hi Jason,

thank you for the good article! I really like the shifting approach for reframing the training data!

But my question about this topic is: What do you think is the next step for a one-step univariate most suitable for that?

Obviously a regressor is the best choice but how can I determine the size of the sliding window for the training?

Thanks a lot for your help and work

~ Christopher

Get Your Start in Machine Learning



Jason Brownlee May 13, 2017 at 6:14 am #

REPLY ↩

My advice is to trial a suite of different window sizes and see what works best.

You can use an ACF/PACF plots as a heuristic:

<http://machinelearningmastery.com/gentle-introduction-autocorrelation-partial-autocorrelation/>



tom June 8, 2017 at 4:06 pm #

REPLY ↩

hi Jason :

In this post, you create new framings of time series ,such as $t-1$, t , $t+1$. But, what's the use of these time series? What's the good effect on model? Maybe

my question is too simple ,because I am a newer ,please understand! thank you !

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

Email Address

START MY EMAIL COURSE



Jason Brownlee June 9, 2017 at 6:19 am #

I am providing a technique to help you convert a series into a supervised learning problem.

This is valuable because you can then transform your time series problems into supervised learning and regression techniques in order to make forecasts.



tom June 9, 2017 at 11:40 am #

REPLY ↩

Wow, your answer always makes me learn a lot. Thank you Jason!

Get Your Start in Machine Learning



Jason Brownlee June 10, 2017 at 8:12 am #

REPLY ↩

You're welcome.



Brad Suzon June 23, 2017 at 11:32 pm #

REPLY ↩

If there are multiple variables varXi to train and only one variable varY to predict will the same technique be used in the below way:

$\text{varX1}(t-1)$ $\text{varX2}(t-1)$ $\text{varX1}(t)$ $\text{varX2}(t)$... $\text{varY}(t-1)$ $\text{varY}(t)$

..

and then use linear regression and as Response= $\text{varY}(t)$?

Thanks in advance



Jason Brownlee June 24, 2017 at 8:03 am #

Not sure I follow your question Brad, perhaps you can restate it?



Brad June 25, 2017 at 4:47 pm #

In case there are multiple measures and then make the transformation in order to forecast

$\text{var1}(t-1)$ $\text{var2}(t-1)$ $\text{var1}(t)$ $\text{var2}(t)$... $\text{varN}(t-1)$ $\text{varN}(t)$

linear regression should use as the response variable the $\text{varN}(t)$?

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee June 26, 2017 at 6:06 am #

Sure.

Get Your Start in Machine Learning



Geoff June 24, 2017 at 8:10 am #

REPLY ↩

Hi Jason,

I've found your articles very useful during my capstone at a bootcamp I'm attending. I have two questions that I hope you could advise where to find better info about.

First, I've run into an issue with running PCA on the newly supervised version only the data. Does PCA recognize that the lagged series are actually the same data? If one was to do PCA do they need to perform it before supervising the data?

Secondly, what do you propose as the best learning algorithms and proper ways to perform train test splits on the data?

Thanks again,



Jason Brownlee June 25, 2017 at 5:58 am #

Sorry, I have not used PCA on time series data. I expect careful handling of the framing of t

I have tips for backtesting on time series data here:

<http://machinelearningmastery.com/backtest-machine-learning-models-time-series-forecasting/>

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Kushal July 1, 2017 at 1:31 pm #

Hi Jason

Great post.

Just one question. If the some of the input variables are continuous and some are categorical with one binary, predicting two output variables.

How does the shift work then?

Thanks

Kushal

Get Your Start in Machine Learning



Jason Brownlee July 2, 2017 at 6:26 am #

REPLY ↩

The same, but consider encoding your categorical variables first (e.g. number encoding or one hot encoding).



Kushal July 15, 2017 at 5:22 pm #

REPLY ↩

Thanks

Should I then use the lagged versions of the predictors?

Kushal



Jason Brownlee July 16, 2017 at 7:57 am #

Perhaps, I do not follow your question, perhaps you can restate it with more information.



Viorel Emilian Teodorescu July 8, 2017 at 9:45 am #

great article, Jason!

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee July 9, 2017 at 10:50 am #

REPLY ↩

Thanks, I hope it helps.

Get Your Start in Machine Learning



Chinesh August 10, 2017 at 5:15 pm #

REPLY ↩

very helpful article !!

I am working on developing an algorithm which will predict the future traffic for the restaurant. The features I am using are: Day, whether there was festival, temperature, climatic condition, current rating, whether there was holiday, service rating, number of reviews etc. Can I solve this problem using time series analysis along with these features, If yes how.

Please guide me



Jason Brownlee August 11, 2017 at 6:36 am #

REPLY ↩

This process will guide you:

<http://machinelearningmastery.com/start-here/#process>

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

Email Address

START MY EMAIL COURSE



Hossein August 23, 2017 at 1:16 am #

Great article Jason. Just a naive question: How does this method different from moving average

Thanks



Jason Brownlee August 23, 2017 at 6:56 am #

This post is just about the framing of the problem.

Moving average is something to do to the data once it is framed.



pkl520 August 26, 2017 at 10:29 pm #

Hi , Jason! Good article as always~

I have a question.

“Running the example shows the differentiation of input (t-n) and output (t+n) variables with the current observation (t) considered an output.”

```
values = [x for x in range(10)]
data = series_to_supervised(values, 2, 2)
print(data)
```

```
var1(t-2) var1(t-1) var1(t) var1(t+1)
2 0.0 1.0 2 3.0
3 1.0 2.0 3 4.0
4 2.0 3.0 4 5.0
5 3.0 4.0 5 6.0
6 4.0 5.0 6 7.0
7 5.0 6.0 7 8.0
8 6.0 7.0 8 9.0
```

So above example, var1(t-2) var1(t-1) are input , var1(t) var1(t+1) are output, am I right?

Then,below example.

```
raw = DataFrame()
raw['ob1'] = [x for x in range(10)]
raw['ob2'] = [x for x in range(50, 60)]
values = raw.values
data = series_to_supervised(values, 1, 2)
print(data)
```

Running the example shows the large reframed DataFrame.

Get Your Start in Machine Learning

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

var1(t-1) var2(t-1) var1(t) var2(t) var1(t+1) var2(t+1)

1 0.0 50.0 1 51 2.0 52.0

2 1.0 51.0 2 52 3.0 53.0

3 2.0 52.0 3 53 4.0 54.0

4 3.0 53.0 4 54 5.0 55.0

5 4.0 54.0 5 55 6.0 56.0

6 5.0 55.0 6 56 7.0 57.0

7 6.0 56.0 7 57 8.0 58.0

8 7.0 57.0 8 58 9.0 59.0

var1(t-1) var2(t-1) are input, var1(t) var2(t) var1(t+1) var2(t+1) are output.

can u answer my question? I will be very appreciate!



Jason Brownlee August 27, 2017 at 5:48 am #

REPLY ↩

Yes, or you can interpret and use the columns any way you wish.



Thabet August 30, 2017 at 7:36 am #

Thank you Jason!!

You are the best teacher ever



Jason Brownlee August 30, 2017 at 4:17 pm #

Thanks Thabet.

Charles September 29, 2017 at 12:24 am #

Get Your Start in Machine Learning

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason,

I love your articles! Keep it up! I have a generalization question. In this data set:

var1(t-1) var2(t-1) var1(t) var2(t)

1 0.0 50.0 1 51

2 1.0 51.0 2 52

3 2.0 52.0 3 53

4 3.0 53.0 4 54

5 4.0 54.0 5 55

6 5.0 55.0 6 56

7 6.0 56.0 7 57

8 7.0 57.0 8 58

9 8.0 58.0 9 59

If I was trying to predict var2(t) from the other 3 data, would the input data X shape would be (9,1,3) and the target data Y would be (9,1)? To generalize, what if this was just one instance of multiple time series that I wanted to use. Say I have 1000 instances of time series. Would my data X have the shape (1000,9,3)? And the input target set Y would have shape (1000,9)?

Is my reasoning off? Am I framing my problem the wrong way?

Thanks!

Charles



Jason Brownlee September 29, 2017 at 5:06 am #

This post provides help regarding how to reshape data for LSTMs:

<https://machinelearningmastery.com/reshape-input-data-long-short-term-memory-networks-keras/>



Sean Maloney October 1, 2017 at 5:24 pm #

Hi Jason!

Get Your Start in Machine Learning

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

I'm really struggling to make a new prediction once the model has been build. Could you give an example? I've been trying to write a method that takes the past time data and returns the yhat for the next time.

Thanks you.

Get Your Start in Machine Learning



Jason Brownlee October 2, 2017 at 9:37 am #

REPLY ↩

Yes, see this post:

<https://machinelearningmastery.com/make-sample-forecasts-arma-python/>



Sean Maloney October 1, 2017 at 5:28 pm #

REPLY ↩

P.S. I'm the most stuck at how to scale the new input values.



Jason Brownlee October 2, 2017 at 9:38 am #

REPLY ↩

Any data transforms performed on training data must be performed on new data for which y

Leave a Reply

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

SUBMIT COMMENT

Welcome to Machine Learning Mastery



Hi, I'm Dr. Jason Brownlee.
My goal is to make practitioners like YOU awesome at applied machine learning.

[Read More](#)

Get Good at Time Series Forecasting

Need visualizations and forecast models?
Looking for step-by-step tutorials?


Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree**.
Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Want end-to-end projects?

Get Started with Time Series Forecasting in Python!



Get Your Start in Machine Learning

POPULAR



Time Series Prediction with LSTM Recurrent Neural Networks in Python with Keras

JULY 21, 2016



Your First Machine Learning Project in Python Step-By-Step

JUNE 10, 2016



Develop Your First Neural Network in Python With Keras Step-By-Step

MAY 24, 2016



Sequence Classification with LSTM Recurrent Neural Networks in Python with Keras

JULY 26, 2016

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



How to Setup a Python Environment for Machine Learning and Deep Learning with Anaconda

MARCH 13, 2017



Time Series Forecasting with the Long Short-Term Memory Network in Python

APRIL 7, 2017




Multi-Class Classification Tutorial with the Keras Deep Learning Library

JUNE 2, 2016




Regression Tutorial with the Keras Deep Learning Library in Python

JUNE 9, 2016



Multivariate Time Series Forecasting with LSTMs in Keras

AUGUST 14, 2017



How to Implement the Backpropagation Algorithm From Scratch In Python

NOVEMBER 7, 2016

Get Your Start in Machine Learning

© 2017 Machine Learning Mastery. All Rights Reserved.

[Privacy](#) | [Contact](#) | [About](#)

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE