

[Start Here](#)[Blog](#)[Books](#)[About](#)[Contact](#)

Need help with Python Machine Learning? [Take the FREE Mini-Course](#)

How to Handle Missing Data with Python

by **Jason Brownlee** on March 20, 2017 in **Python Machine Learning**



Real-world data often has missing values.

Data can have missing values for a number of reasons such as observations that were not recorded and data corruption.

Handling missing data is important as many machine learning algorithms do not support data with missing values.

In this tutorial, you will discover how to handle missing data for machine learning with Python.

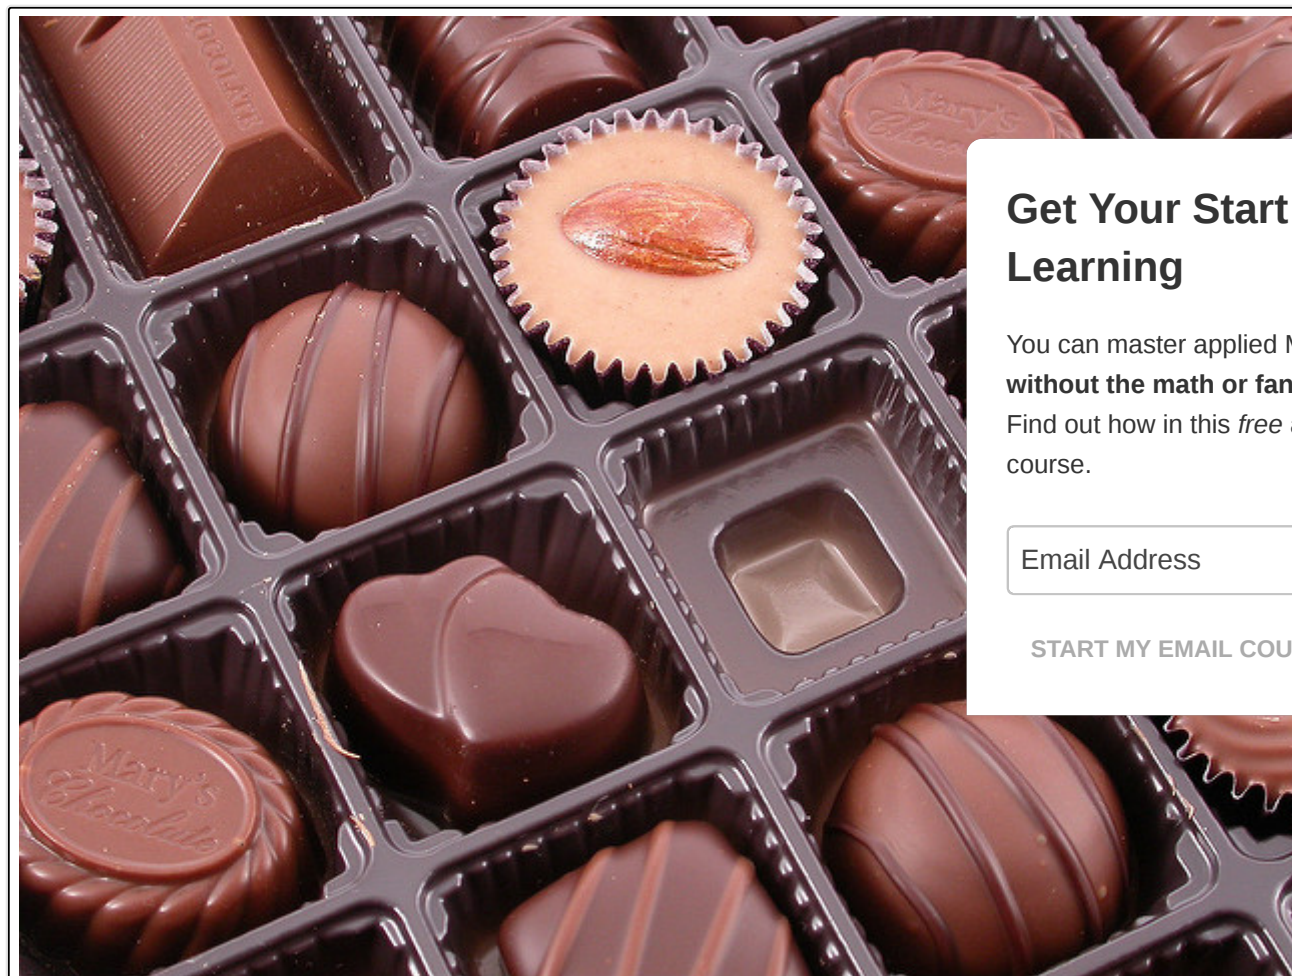
Specifically, after completing this tutorial you will know:

[Get Your Start in Machine Learning](#)

- How to marking invalid or corrupt values as missing in your dataset.
- How to remove rows with missing data from your dataset.
- How to impute missing values with mean values in your dataset.

Let's get started.

Note: The examples in this post assume that you have Python 2 or 3 with Pandas, NumPy and Scikit-Learn installed, specifically scikit-learn version 0.18 or higher.



How to Handle Missing Values with Python
Photo by CoCreatr, some rights reserved.

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Overview

This tutorial is divided into 6 parts:

1. **Pima Indians Diabetes Dataset:** where we look at a dataset that has known missing values.
2. **Mark Missing Values:** where we learn how to mark missing values in a dataset.
3. **Missing Values Causes Problems:** where we see how a machine learning algorithm can fail when it contains missing values.
4. **Remove Rows With Missing Values:** where we see how to remove rows that contain missing values.
5. **Impute Missing Values:** where we replace missing values with sensible values.
6. **Algorithms that Support Missing Values:** where we learn about algorithms that support missing values.

First, let's take a look at our sample dataset with missing values.

1. Pima Indians Diabetes Dataset

The [Pima Indians Diabetes Dataset](#) involves predicting the onset of diabetes within 5 years in Pima

It is a binary (2-class) classification problem. The number of observations for each class is not balanced. There are 8 input variables and 1 output variable. The variable names are as follows:

- 0. Number of times pregnant.
- 1. Plasma glucose concentration a 2 hours in an oral glucose tolerance test.
- 2. Diastolic blood pressure (mm Hg).
- 3. Triceps skinfold thickness (mm).
- 4. 2-Hour serum insulin (mu U/ml).
- 5. Body mass index (weight in kg/(height in m)^2).
- 6. Diabetes pedigree function.
- 7. Age (years).
- 8. Class variable (0 or 1).

The baseline performance of predicting the most prevalent class is a classification accuracy of approximately 65%. Top results achieve a classification accuracy of approximately 77%.

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

A sample of the first 5 rows is listed below.

```
1 6,148,72,35,0,33.6,0.627,50,1
2 1,85,66,29,0,26.6,0.351,31,0
3 8,183,64,0,0,23.3,0.672,32,1
4 1,89,66,23,94,28.1,0.167,21,0
5 0,137,40,35,168,43.1,2.288,33,1
```

This dataset is known to have missing values.

Specifically, there are missing observations for some columns that are marked as a zero value.

We can corroborate this by the definition of those columns and the domain knowledge that a zero value is invalid for those measures, e.g. a zero for body mass index or blood pressure is invalid.

Download the dataset from [here](#) and save it to your current working directory with the file name *pima*

2. Mark Missing Values

In this section, we will look at how we can identify and mark values as missing.

We can use plots and summary statistics to help identify missing or corrupt data.

We can load the dataset as a Pandas DataFrame and print summary statistics on each attribute.

```
1 from pandas import read_csv
2 dataset = read_csv('pima-indians-diabetes.csv', header=None)
3 print(dataset.describe())
```

Running this example produces the following output:

	0	1	2	3	4	5	\
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```

10
11          6          7          8
12 count  768.000000  768.000000  768.000000
13 mean    0.471876   33.240885   0.348958
14 std     0.331329   11.760232   0.476951
15 min     0.078000   21.000000   0.000000
16 25%     0.243750   24.000000   0.000000
17 50%     0.372500   29.000000   0.000000
18 75%     0.626250   41.000000   1.000000
19 max     2.420000   81.000000   1.000000

```

This is useful.

We can see that there are columns that have a minimum value of zero (0). On some columns, a value of zero does not make sense and indicates an invalid or missing value.

Specifically, the following columns have an invalid zero minimum value:

- 1: Plasma glucose concentration
- 2: Diastolic blood pressure
- 3: Triceps skinfold thickness
- 4: 2-Hour serum insulin
- 5: Body mass index

Let's confirm this by looking at the raw data, the example prints the first 20 rows of data.

```

1 from pandas import read_csv
2 import numpy
3 dataset = read_csv('pima-indians-diabetes.csv', header=None)
4 # print the first 20 rows of data
5 print(dataset.head(20))

```

Running the example, we can clearly see 0 values in the columns 2, 3, 4, and 5.

```

1      0      1      2      3      4      5      6      7      8
2  0      6    148    72    35      0    33.6    0.627    50    1
3  1      1     85    66    29      0    26.6    0.351    31    0
4  2      8    183    64     0      0    23.3    0.672    32    1
5  3      1     89    66    23     94    28.1    0.167    21    0
6  4      0    137    40    35    168    43.1    2.288    33    1
7  5      5    116    74     0      0    25.6    0.201    30    0

```

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

8	6	3	78	50	32	88	31.0	0.248	26	1
9	7	10	115	0	0	0	35.3	0.134	29	0
10	8	2	197	70	45	543	30.5	0.158	53	1
11	9	8	125	96	0	0	0.0	0.232	54	1
12	10	4	110	92	0	0	37.6	0.191	30	0
13	11	10	168	74	0	0	38.0	0.537	34	1
14	12	10	139	80	0	0	27.1	1.441	57	0
15	13	1	189	60	23	846	30.1	0.398	59	1
16	14	5	166	72	19	175	25.8	0.587	51	1
17	15	7	100	0	0	0	30.0	0.484	32	1
18	16	0	118	84	47	230	45.8	0.551	31	1
19	17	7	107	74	0	0	29.6	0.254	31	1
20	18	1	103	30	38	83	43.3	0.183	33	0
21	19	1	115	70	30	96	34.6	0.529	32	1

We can get a count of the number of missing values on each of these columns. We can do this by marking all of the values in the subset of the DataFrame we are interested in that have zero values as True. We can then count the number of true values in each column.

We can do this by marking all of the values in the subset of the DataFrame we are interested in that have zero values as True. We can then count the number of true values in each column.

```
1 from pandas import read_csv
2 dataset = read_csv('pima-indians-diabetes.csv', header=None)
3 print((dataset[[1,2,3,4,5]] == 0).sum())
```

Running the example prints the following output:

```
1 1 5
2 2 35
3 3 227
4 4 374
5 5 11
```

We can see that columns 1,2 and 5 have just a few zero values, whereas columns 3 and 4 show a lot more, nearly half of the rows.

This highlights that different “missing value” strategies may be needed for different columns, e.g. to ensure that there are still a sufficient number of records left to train a predictive model.

In Python, specifically Pandas, NumPy and Scikit-Learn, we mark missing values as NaN.

Values with a NaN value are ignored from operations like sum, count, etc.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

We can mark values as NaN easily with the Pandas DataFrame by using the `replace()` function on a subset of the columns we are interested in.

After we have marked the missing values, we can use the `isnull()` function to mark all of the NaN values in the dataset as True and get a count of the missing values for each column.

```
1 from pandas import read_csv
2 import numpy
3 dataset = read_csv('pima-indians-diabetes.csv', header=None)
4 # mark zero values as missing or NaN
5 dataset[[1,2,3,4,5]] = dataset[[1,2,3,4,5]].replace(0, numpy.NaN)
6 # count the number of NaN values in each column
7 print(dataset.isnull().sum())
```

Running the example prints the number of missing values in each column. We can see that the columns 1:5 have the same number of missing values as zero values identified above. This is a sign that we have marked the identified missing values correctly.

We can see that the columns 1 to 5 have the same number of missing values as zero values identified above. This is a sign that we have marked the identified missing values correctly.

```
1 0      0
2 1      5
3 2     35
4 3    227
5 4    374
6 5     11
7 6      0
8 7      0
9 8      0
```

This is a useful summary. I always like to look at the actual data though, to confirm that I have not forced any values to be NaN.

Below is the same example, except we print the first 20 rows of data.

```
1 from pandas import read_csv
2 import numpy
3 dataset = read_csv('pima-indians-diabetes.csv', header=None)
4 # mark zero values as missing or NaN
5 dataset[[1,2,3,4,5]] = dataset[[1,2,3,4,5]].replace(0, numpy.NaN)
6 # print the first 20 rows of data
7 print(dataset.head(20))
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Running the example, we can clearly see NaN values in the columns 2, 3, 4 and 5. There are only 5 missing values in column 1, so it is not surprising we did not see an example in the first 20 rows.

It is clear from the raw data that marking the missing values had the intended effect.

	0	1	2	3	4	5	6	7	8
1	0	6	148.0	72.0	35.0	NaN	33.6	0.627	50
2	1	1	85.0	66.0	29.0	NaN	26.6	0.351	31
3	2	8	183.0	64.0	NaN	NaN	23.3	0.672	32
4	3	1	89.0	66.0	23.0	94.0	28.1	0.167	21
5	4	0	137.0	40.0	35.0	168.0	43.1	2.288	33
6	5	5	116.0	74.0	NaN	NaN	25.6	0.201	30
7	6	3	78.0	50.0	32.0	88.0	31.0	0.248	26
8	7	10	115.0	NaN	NaN	NaN	35.3	0.134	29
9	8	2	197.0	70.0	45.0	543.0	30.5	0.158	53
10	9	8	125.0	96.0	NaN	NaN	NaN	0.232	54
11	10	4	110.0	92.0	NaN	NaN	37.6	0.191	30
12	11	10	168.0	74.0	NaN	NaN	38.0	0.537	34
13	12	10	139.0	80.0	NaN	NaN	27.1	1.441	57
14	13	1	189.0	60.0	23.0	846.0	30.1	0.398	59
15	14	5	166.0	72.0	19.0	175.0	25.8	0.587	51
16	15	7	100.0	NaN	NaN	NaN	30.0	0.484	32
17	16	0	118.0	84.0	47.0	230.0	45.8	0.551	31
18	17	7	107.0	74.0	NaN	NaN	29.6	0.254	31
19	18	1	103.0	30.0	38.0	83.0	43.3	0.183	33
20	19	1	115.0	70.0	30.0	96.0	34.6	0.529	32

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Before we look at handling missing values, let's first demonstrate that having missing values in a dataset can cause errors with some machine learning algorithms.

3. Missing Values Causes Problems

Having missing values in a dataset can cause errors with some machine learning algorithms.

In this section, we will try to evaluate a the Linear Discriminant Analysis (LDA) algorithm on the dataset with missing values.

This is an algorithm that does not work when there are missing values in the dataset.

The below example marks the missing values in the dataset, as we did in the previous section, then attempts to evaluate LDA using 3-fold cross validation and print the mean accuracy.

Get Your Start in Machine Learning


```
1 from pandas import read_csv
2 import numpy
3 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
4 from sklearn.model_selection import KFold
5 from sklearn.model_selection import cross_val_score
6 dataset = read_csv('pima-indians-diabetes.csv', header=None)
7 # mark zero values as missing or NaN
8 dataset[[1,2,3,4,5]] = dataset[[1,2,3,4,5]].replace(0, numpy.NaN)
9 # split dataset into inputs and outputs
10 values = dataset.values
11 X = values[:,0:8]
12 y = values[:,8]
13 # evaluate an LDA model on the dataset using k-fold cross validation
14 model = LinearDiscriminantAnalysis()
15 kfold = KFold(n_splits=3, random_state=7)
16 result = cross_val_score(model, X, y, cv=kfold, scoring='accuracy')
17 print(result.mean())
```

Running the example results in an error, as follows:

```
1 ValueError: Input contains NaN, infinity or a value too large for dtype('float64').
```

This is as we expect.

We are prevented from evaluating an LDA algorithm (and other algorithms) on the dataset with missing values.

Now, we can look at methods to handle the missing values.

4. Remove Rows With Missing Values

The simplest strategy for handling missing data is to remove records that contain a missing value.

We can do this by creating a new Pandas DataFrame with the rows containing missing values removed.

Pandas provides the `dropna()` function that can be used to drop either columns or rows with missing data. We can use `dropna()` to remove all rows with missing data, as follows:

```
1 from pandas import read_csv
2 import numpy
3 dataset = read_csv('pima-indians-diabetes.csv', header=None)
4 # mark zero values as missing or NaN
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
5 dataset[[1,2,3,4,5]] = dataset[[1,2,3,4,5]].replace(0, numpy.NaN)
6 # drop rows with missing values
7 dataset.dropna(inplace=True)
8 # summarize the number of rows and columns in the dataset
9 print(dataset.shape)
```

Running this example, we can see that the number of rows has been aggressively cut from 768 in the original dataset to 392 with all rows containing a NaN removed.

```
1 (392, 9)
```

We now have a dataset that we could use to evaluate an algorithm sensitive to missing values like LDA.

```
1 from pandas import read_csv
2 import numpy
3 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
4 from sklearn.model_selection import KFold
5 from sklearn.model_selection import cross_val_score
6 dataset = read_csv('pima-indians-diabetes.csv', header=None)
7 # mark zero values as missing or NaN
8 dataset[[1,2,3,4,5]] = dataset[[1,2,3,4,5]].replace(0, numpy.NaN)
9 # drop rows with missing values
10 dataset.dropna(inplace=True)
11 # split dataset into inputs and outputs
12 values = dataset.values
13 X = values[:,0:8]
14 y = values[:,8]
15 # evaluate an LDA model on the dataset using k-fold cross validation
16 model = LinearDiscriminantAnalysis()
17 kfold = KFold(n_splits=3, random_state=7)
18 result = cross_val_score(model, X, y, cv=kfold, scoring='accuracy')
19 print(result.mean())
```

The example runs successfully and prints the accuracy of the model.

```
1 0.78582892934
```

Removing rows with missing values can be too limiting on some predictive modeling problems, an alternative is to impute missing values.

5. Impute Missing Values

Imputing refers to using a model to replace missing values.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

There are many options we could consider when replacing a missing value, for example:

- A constant value that has meaning within the domain, such as 0, distinct from all other values.
- A value from another randomly selected record.
- A mean, median or mode value for the column.
- A value estimated by another predictive model.

Any imputing performed on the training dataset will have to be performed on new data in the future when predictions are needed from the finalized model. This needs to be taken into consideration when choosing how to impute the missing values.

For example, if you choose to impute with mean column values, these mean column values will need to be stored to file for later use on new data that has missing values.

Pandas provides the `fillna()` function for replacing missing values with a specific value.

For example, we can use `fillna()` to replace missing values with the mean value for each column, as

```
1 from pandas import read_csv
2 import numpy
3 dataset = read_csv('pima-indians-diabetes.csv', header=None)
4 # mark zero values as missing or NaN
5 dataset[[1,2,3,4,5]] = dataset[[1,2,3,4,5]].replace(0, numpy.NaN)
6 # fill missing values with mean column values
7 dataset.fillna(dataset.mean(), inplace=True)
8 # count the number of NaN values in each column
9 print(dataset.isnull().sum())
```

Running the example provides a count of the number of missing values in each column, showing zero

```
1 0    0
2 1    0
3 2    0
4 3    0
5 4    0
6 5    0
7 6    0
8 7    0
9 8    0
```

The scikit-learn library provides the `Imputer()` pre-processing class that can be used to replace missing

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

It is a flexible class that allows you to specify the value to replace (it can be something other than NaN) and the technique used to replace it (such as mean, median, or mode). The Imputer class operates directly on the NumPy array instead of the DataFrame.

The example below uses the Imputer class to replace missing values with the mean of each column then prints the number of NaN values in the transformed matrix.

```
1 from pandas import read_csv
2 from sklearn.preprocessing import Imputer
3 import numpy
4 dataset = read_csv('pima-indians-diabetes.csv', header=None)
5 # mark zero values as missing or NaN
6 dataset[[1,2,3,4,5]] = dataset[[1,2,3,4,5]].replace(0, numpy.NaN)
7 # fill missing values with mean column values
8 values = dataset.values
9 imputer = Imputer()
10 transformed_values = imputer.fit_transform(values)
11 # count the number of NaN values in each column
12 print(numpy.isnan(transformed_values).sum())
```

Running the example shows that all NaN values were imputed successfully.

```
1
```

In either case, we can train algorithms sensitive to NaN values in the transformed dataset, such as LDA.

The example below shows the LDA algorithm trained in the Imputer transformed dataset.

```
1 from pandas import read_csv
2 import numpy
3 from sklearn.preprocessing import Imputer
4 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
5 from sklearn.model_selection import KFold
6 from sklearn.model_selection import cross_val_score
7 dataset = read_csv('pima-indians-diabetes.csv', header=None)
8 # mark zero values as missing or NaN
9 dataset[[1,2,3,4,5]] = dataset[[1,2,3,4,5]].replace(0, numpy.NaN)
10 # split dataset into inputs and outputs
11 values = dataset.values
12 X = values[:,0:8]
13 y = values[:,8]
14 # fill missing values with mean column values
15 imputer = Imputer()
16 transformed_X = imputer.fit_transform(X)
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
17 # evaluate an LDA model on the dataset using k-fold cross validation
18 model = LinearDiscriminantAnalysis()
19 kfold = KFold(n_splits=3, random_state=7)
20 result = cross_val_score(model, transformed_X, y, cv=kfold, scoring='accuracy')
21 print(result.mean())
```

Running the example prints the accuracy of LDA on the transformed dataset.

```
1 0.766927083333
```

Try replacing the missing values with other values and see if you can lift the performance of the model.

Maybe missing values have meaning in the data.

Next we will look at using algorithms that treat missing values as just another value when modeling.

6. Algorithms that Support Missing Values

Not all algorithms fail when there is missing data.

There are algorithms that can be made robust to missing data, such as k-Nearest Neighbors that can handle missing values. If a value is missing.

There are also algorithms that can use the missing value as a unique and different value when building regression trees.

Sadly, the scikit-learn implementations of decision trees and k-Nearest Neighbors are not robust to missing data.

Nevertheless, this remains as an option if you consider using another algorithm implementation (such as [xgboost](#)) or developing your own implementation.

Further Reading

- [Working with missing data, in Pandas](#)
- [Imputation of missing values, in scikit-learn](#)

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Summary

In this tutorial, you discovered how to handle machine learning data that contains missing values.

Specifically, you learned:

- How to mark missing values in a dataset as `numpy.nan`.
- How to remove rows from the dataset that contain missing values.
- How to replace missing values with sensible values.

Do you have any questions about handling missing values?

Ask your questions in the comments and I will do my best to answer.

Frustrated With Python Machine Learning?

Develop Your Own Models

...with just a few lines of sci

Discover how in my new
[Machine Learning Mastery](#)

Covers **self-study tutorials** and **end-to-end projects**.
Loading data, visualization, modeling, tuning, and much more...

Finally Bring Machine Learning To Your Own Projects


Skip the Academics. Just Results.


[Click to learn more.](#)

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**
Find out how in this *free* and *practical* email course.

[START MY EMAIL COURSE](#)





About Jason Brownlee

Dr. Jason Brownlee is a husband, proud father, academic researcher, author, professional developer and a machine learning practitioner. He is dedicated to helping developers get started and get good at applied machine learning. [Learn more.](#)

[View all posts by Jason Brownlee](#) →

< How to Train a Final Machine Learning Model

17 Responses to *How to Handle Missing Data with Python*



Mike March 20, 2017 at 3:16 pm #

Fancy impute is a library i've turned too for imputation:

<https://github.com/hammerlab/fancyimpute>

Also missingno is great for visualizations!

<https://github.com/ResidentMario/missingno>

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee March 21, 2017 at 8:37 am #

REPLY ↩

Get Your Start in Machine Learning

Thanks for the tip Mike.



bakyalakshmi September 27, 2017 at 2:56 pm #

REPLY ↩

please tell me about how to impute median using one dataset



Jozo Kovac April 1, 2017 at 8:06 am #

REPLY ↩

Thanks for pointing on interesting problem. I would love another one about how to deal with categorical attributes in Python

And dear reader, please never ever remove rows with missing values. It changes the distribution of your data.
Learn from mistakes of others and don't repeat them 😊



Jason Brownlee April 2, 2017 at 6:22 am #

Thanks Jozo.

This post will help with categorical input data:

<http://machinelearningmastery.com/data-preparation-gradient-boosting-xgboost-python/>

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Tommy Carstensen April 4, 2017 at 3:56 am #

REPLY ↩

Super duper! Thanks for writing! Would it have been worth mentioning interpolate of Pandas? <http://pandas.pydata.org/pandas-docs/stable/generated/pandas.Series.interpolate.html>

Jason Brownlee April 4, 2017 at 9:18 am #

Get Your Start in Machine Learning



Thanks Tommy.



Aswathy April 14, 2017 at 12:10 pm #

REPLY ↩

Hi Jason,

I was just wondering if there is a way to use a different imputation strategy for each column. Say, for a categorical feature you want to impute using the mode but for a continuous attribute, you want to impute using mean.



Jason Brownlee April 15, 2017 at 9:30 am #

Yes, try lots of techniques, go with whatever results in the most accurate models.



Salu Khadka June 11, 2017 at 12:29 am #

thanks for your tutorial sir.

I would also seek help from you for multi label classification of a textual data , if possible.

For example, categorizing a twitter post as related to sports, business , tech , or others.

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

REPLY ↩



Jason Brownlee June 11, 2017 at 8:26 am #

Sure, see this post:

<http://machinelearningmastery.com/sequence-classification-lstm-recurrent-neural-networks-python-keras/>

Ali Gabriel Lara June 13, 2017 at 4:51 am #

Get Your Start in Machine Learning



Hello Mr. Brownlee. Thank you so much for your post.

Do you know any approach to recognize the pattern of missing data? I mean, I am interested in discovering the pattern of missing data on a time series data.

The database is historical data of a chemical process. I think I should apply some pattern recognition approach columnwise because each column represents a process variable and the value coming from a transmissor.

My goal is to predict if the missing data is for a mechanical fault or a deviation in registration process or for any other causes. Then I should apply a kind of filling methods if it is required.

Have you any advice? Thanks in advance



Jason Brownlee June 13, 2017 at 8:25 am #

I would invert the problem and model the series of missing data and mark all data you do have as "1".

Great problem!

Let me know how you go.



Patricia Villa October 5, 2017 at 3:45 pm #

You helped me keep my sanity. THANK YOU!!



Jason Brownlee October 5, 2017 at 5:23 pm #

I'm really glad to hear that Patricia!

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

REPLY ↩

Get Your Start in Machine Learning



Sachin Raj October 6, 2017 at 7:58 pm #

REPLY ↩

How to know whether to apply mean or to replace it with mode?



Jason Brownlee October 7, 2017 at 5:54 am #

REPLY ↩

Try both and see what results in the most skillful models.

Leave a Reply

Name (required)

Email (will not be published) (required)

Website

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

[SUBMIT COMMENT](#)

Welcome to Machine Learning Mastery



Hi, I'm Dr. Jason Brownlee.

My goal is to make practitioners like YOU awesome at applied machine learning.

[Read More](#)

Develop Predictive Models With Python

Want to develop your own models in scikit-learn?

Want step-by-step tutorials?

Looking for sample code and templates?


Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

[START MY EMAIL COURSE](#)

[Get Your Start in Machine Learning](#)



POPULAR

**Time Series Prediction with LSTM Recurrent Neural Networks in Python with Keras**

JULY 21, 2016

**Your First Machine Learning Project in Python Step-By-Step**

JUNE 10, 2016

**Develop Your First Neural Network in Python With Keras Step-By-Step**

MAY 24, 2016

**Sequence Classification with LSTM Recurrent Neural Networks in Python with Keras**

JULY 26, 2016

How to Setup a Python Environment for Machine Learning and Deep Learning with Anaconda

MARCH 13, 2017

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

[START MY EMAIL COURSE](#)[Get Your Start in Machine Learning](#)



Time Series Forecasting with the Long Short-Term Memory Network in Python

APRIL 7, 2017




Multi-Class Classification Tutorial with the Keras Deep Learning Library

JUNE 2, 2016




Regression Tutorial with the Keras Deep Learning Library in Python

JUNE 9, 2016



Multivariate Time Series Forecasting with LSTMs in Keras

AUGUST 14, 2017



How to Implement the Backpropagation Algorithm From Scratch In Python

NOVEMBER 7, 2016

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

© 2017 Machine Learning Mastery. All Rights Reserved.

[Privacy](#) | [Contact](#) | [About](#)

Get Your Start in Machine Learning