

[Start Here](#)[Blog](#)[Books](#)[About](#)[Contact](#)

Need help with machine learning? [Take the FREE Crash-Course.](#)

How to Prepare Text Data for Deep Learning with Keras

by **Jason Brownlee** on October 2, 2017 in **Natural Language Processing**



You cannot feed raw text directly into deep learning models.

Text data must be encoded as numbers to be used as input or output for machine learning and deep learning models.

The Keras deep learning library provides some basic tools to help you prepare your text data.

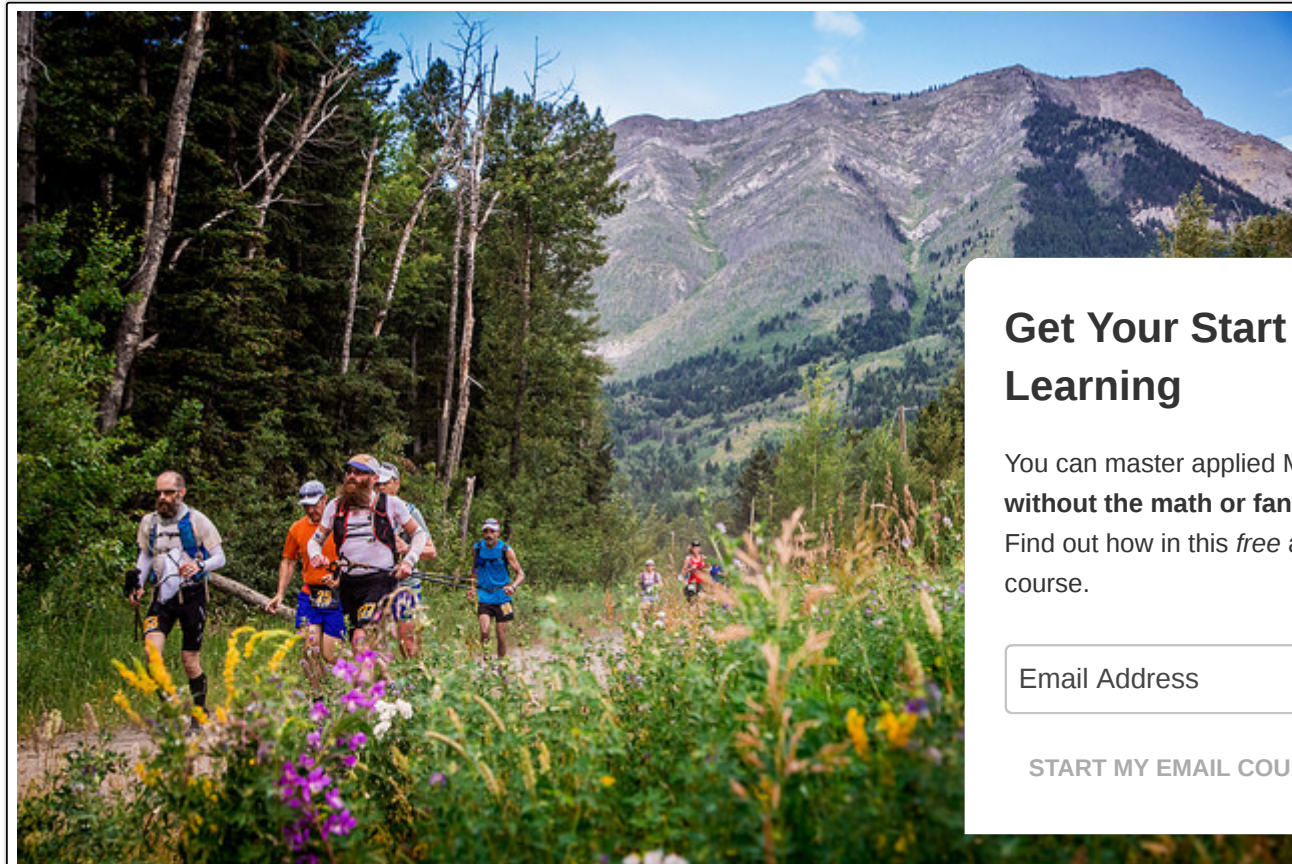
In this tutorial, you will discover how you can use Keras to prepare your text data.

After completing this tutorial, you will know:

[Get Your Start in Machine Learning](#)

- About the convenience methods that you can use to quickly prepare text data.
- The Tokenizer API that can be fit on training data and used to encode training, validation, and test documents.
- The range of 4 different document encoding schemes offered by the Tokenizer API.

Let's get started.



How to Prepare Text Data for Deep Learning with Keras
Photo by [ActiveSteve](#), some rights reserved.

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Tutorial Overview

This tutorial is divided into 4 parts; they are:

Get Your Start in Machine Learning

1. Split words with `text_to_word_sequence`.
2. Encoding with `one_hot`.
3. Hash Encoding with `hashing_trick`.
4. Tokenizer API

Split Words with `text_to_word_sequence`

A good first step when working with text is to split it into words.

Words are called tokens and the process of splitting text into tokens is called tokenization.

Keras provides the `text_to_word_sequence()` function that you can use to split text into a list of words.

By default, this function automatically does 3 things:

- Splits words by space (`split=" "`).
- Filters out punctuation (`filters='!"#$%&()*+,-./:;<=>?@[\\]^_`{|}~\t\n'`).
- Converts text to lowercase (`lower=True`).

You can change any of these defaults by passing arguments to the function.

Below is an example of using the `text_to_word_sequence()` function to split a document (in this case

```
1 from keras.preprocessing.text import text_to_word_sequence
2 # define the document
3 text = 'The quick brown fox jumped over the lazy dog.'
4 # tokenize the document
5 result = text_to_word_sequence(text)
6 print(result)
```

Running the example creates an array containing all of the words in the document. The list of words is printed for review.

```
1 ['the', 'quick', 'brown', 'fox', 'jumped', 'over', 'the', 'lazy', 'dog']
```

This is a good first step, but further pre-processing is required before you can work with the text.

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree**. Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Encoding with one_hot

It is popular to represent a document as a sequence of integer values, where each word in the document is represented as a unique integer.

Keras provides the `one_hot()` function that you can use to tokenize and integer encode a text document in one step. The name suggests that it will create a one-hot encoding of the document, which is not the case. Instead, the function is a wrapper for the `hashing_trick()` function described in the next section. The function returns an integer encoded version of the document. The use of a hash function means that there may be collisions and not all words will be assigned unique integer values.

Instead, the function is a wrapper for the `hashing_trick()` function described in the next section. The function returns an integer encoded version of the document. The use of a hash function means that there may be collisions and not all words will be assigned unique integer values.

As with the `text_to_word_sequence()` function in the previous section, the `one_hot()` function will make words based on white space.

In addition to the text, the vocabulary size (total words) must be specified. This could be the total number of words in the text, or the number of words in the vocabulary. The size of the vocabulary defines the number of words to encode additional documents that contains additional words. The size of the vocabulary defines the number of words to encode. Ideally, this should be larger than the vocabulary by some percentage (perhaps 25%) to minimize the number of collisions. is used, although as we will see in the next section, alternate hash functions can be specified when creating the encoder.

We can use the `text_to_word_sequence()` function from the previous section to split the document into unique words in the document. The size of this set can be used to estimate the size of the vocabulary.

For example:

```
1 from keras.preprocessing.text import text_to_word_sequence
2 # define the document
3 text = 'The quick brown fox jumped over the lazy dog.'
4 # estimate the size of the vocabulary
5 words = set(text_to_word_sequence(text))
6 vocab_size = len(words)
7 print(vocab_size)
```

We can put this together with the `one_hot()` function and one hot encode the words in the document. The complete example is listed below.

The vocabulary size is increased by one-third to minimize collisions when hashing words.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

```
1 from keras.preprocessing.text import one_hot
2 from keras.preprocessing.text import text_to_word_sequence
3 # define the document
4 text = 'The quick brown fox jumped over the lazy dog.'
5 # estimate the size of the vocabulary
6 words = set(text_to_word_sequence(text))
7 vocab_size = len(words)
8 print(vocab_size)
9 # integer encode the document
10 result = one_hot(text, round(vocab_size*1.3))
11 print(result)
```

Running the example first prints the size of the vocabulary as 8. The encoded document is then printed as an array of integer encoded words.

```
1 8
2 [5, 9, 8, 7, 9, 1, 5, 3, 8]
```

Hash Encoding with hashing_trick

A limitation of integer and count base encodings is that they must maintain a vocabulary of words and

An alternative to this approach is to use a one-way hash function to convert words to integers. This approach is faster and requires less memory.

Keras provides the `hashing_trick()` function that tokenizes and then integer encodes the document, j flexibility, allowing you to specify the hash function as either 'hash' (the default) or other hash function.

Below is an example of integer encoding a document using the md5 hash function.

```
1 from keras.preprocessing.text import hashing_trick
2 from keras.preprocessing.text import text_to_word_sequence
3 # define the document
4 text = 'The quick brown fox jumped over the lazy dog.'
5 # estimate the size of the vocabulary
6 words = set(text_to_word_sequence(text))
7 vocab_size = len(words)
8 print(vocab_size)
9 # integer encode the document
10 result = hashing_trick(text, round(vocab_size*1.3), hash_function='md5')
11 print(result)
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Running the example prints the size of the vocabulary and the integer encoded document.

We can see that the use of a different hash function results in consistent, but different integers for words as the `one_hot()` function in the previous section.

```
1 8
2 [6, 4, 1, 2, 7, 5, 6, 2, 6]
```

Tokenizer API

So far we have looked at one-off convenience methods for preparing text with Keras.

Keras provides a more sophisticated API for preparing text that can be fit and reused to prepare multiple text documents. This may be the preferred approach for large projects.

Keras provides the `Tokenizer` class for preparing text documents for deep learning. The `Tokenizer` model can be used to fit on a set of documents or integer encoded text documents.

For example:

```
1 from keras.preprocessing.text import Tokenizer
2 # define 5 documents
3 docs = ['Well done!',
4         'Good work',
5         'Great effort',
6         'nice work',
7         'Excellent!']
8 # create the tokenizer
9 t = Tokenizer()
10 # fit the tokenizer on the documents
11 t.fit_on_texts(docs)
```

Once fit, the `Tokenizer` provides 4 attributes that you can use to query what has been learned about your documents:

- **word_counts**: A dictionary of words and their counts.
- **word_docs**: An integer count of the total number of documents that were used to fit the `Tokenizer`.
- **word_index**: A dictionary of words and their uniquely assigned integers.
- **document_count**: A dictionary of words and how many documents each appeared in.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

For example:

```
1 # summarize what was learned
2 print(t.word_counts)
3 print(t.document_count)
4 print(t.word_index)
5 print(t.word_docs)
```

Once the Tokenizer has been fit on training data, it can be used to encode documents in the train or test datasets.

The `texts_to_matrix()` function on the Tokenizer can be used to create one vector per document provided per input. The length of the vectors is the total size of the vocabulary.

This function provides a suite of standard bag-of-words model text encoding schemes that can be provided via a mode argument to the function.

The modes available include:

- `'binary'`: Whether or not each word is present in the document. This is the default.
- `'count'`: The count of each word in the document.
- `'tfidf'`: The Text Frequency-Inverse DocumentFrequency (TF-IDF) scoring for each word in the document.
- `'freq'`: The frequency of each word as a ratio of words within each document.

We can put all of this together with a worked example.

```
1 from keras.preprocessing.text import Tokenizer
2 # define 5 documents
3 docs = ['Well done!',
4         'Good work',
5         'Great effort',
6         'nice work',
7         'Excellent!']
8 # create the tokenizer
9 t = Tokenizer()
10 # fit the tokenizer on the documents
11 t.fit_on_texts(docs)
12 # summarize what was learned
13 print(t.word_counts)
14 print(t.document_count)
15 print(t.word_index)
16 print(t.word_docs)
17 # integer encode documents
```

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
18 encoded_docs = t.texts_to_matrix(docs, mode='count')
19 print(encoded_docs)
```

Running the example fits the Tokenizer with 5 small documents. The details of the fit Tokenizer are printed. Then the 5 documents are encoded using a word count.

Each document is encoded as a 9-element vector with one position for each word and the chosen encoding scheme value for each word position. In this case, a simple word count mode is used.

```
1 OrderedDict([('well', 1), ('done', 1), ('good', 1), ('work', 2), ('great', 1), ('effort', 1), ('nice', 1), ('excellent', 1)])
2 5
3 {'work': 1, 'effort': 6, 'done': 3, 'great': 5, 'good': 4, 'excellent': 8, 'well': 2, 'nice': 7}
4 {'work': 2, 'effort': 1, 'done': 1, 'well': 1, 'good': 1, 'great': 1, 'excellent': 1, 'nice': 1}
5 [[ 0.  0.  1.  1.  0.  0.  0.  0.  0.]
6  [ 0.  1.  0.  0.  1.  0.  0.  0.  0.]
7  [ 0.  0.  0.  0.  0.  1.  1.  0.  0.]
8  [ 0.  1.  0.  0.  0.  0.  0.  1.  0.]
9  [ 0.  0.  0.  0.  0.  0.  0.  0.  1.]]
```

Further Reading

This section provides more resources on the topic if you are looking go deeper.

- [Text Preprocessing Keras API](#)
- [text_to_word_sequence Keras API](#)
- [one_hot Keras API](#)
- [hashing_trick Keras API](#)
- [Tokenizer Keras API](#)

Summary

In this tutorial, you discovered how you can use the Keras API to prepare your text data for deep learning.

Specifically, you learned:

- About the convenience methods that you can use to quickly prepare text data.
- The Tokenizer API that can be fit on training data and used to encode training, validation, and test data.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

- The range of 4 different document encoding schemes offered by the Tokenizer API.

Do you have any questions?

Ask your questions in the comments below and I will do my best to answer.



About Jason Brownlee

Dr. Jason Brownlee is a husband, proud father, academic researcher, author, professional developer and a machine learning practitioner. He is dedicated to helping developers get started and get good at applied machine learning. [Learn more.](#)

[View all posts by Jason Brownlee](#) →

< How to Prepare Text Data for Machine Learning with scikit-learn

How to Prepare Text Data for Machine Learning with Keras >

6 Responses to *How to Prepare Text Data for Deep Learning with Keras*



Chiedu October 2, 2017 at 6:40 am #

Hi Jason,

Do you have any plans to cover word embeddings using either word2vec or GloVe and how they work with Keras?



Jason Brownlee October 2, 2017 at 9:40 am #

Yes! I have many posts on word embeddings scheduled for the coming days/weeks.

REPLY ↩

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Lalit Parihar October 6, 2017 at 6:59 pm #

REPLY ↩

Hello Jason,

It seems the attributes mentioned for Tokenizer have been typed incorrectly, document_count and word_docs have been inter-changed.

Thanks,
Lalit



Jason Brownlee October 7, 2017 at 5:52 am #

Thanks Lalit, in which part of the tutorial exactly?



Gopika Bhardwaj October 16, 2017 at 3:07 am #

How do we further apply a neural network on this data?



Jason Brownlee October 16, 2017 at 5:45 am #

Great question, I will have many tutorials about how to do that coming out on the blog in co

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Leave a Reply

Get Your Start in Machine Learning

Name (required)

Email (will not be published) (required)

Website

Welcome to Machine Learning Mastery



Hi, I'm Dr. Jason Brownlee.
My goal is to make practitioners like YOU awesome at applied machine learning.

[Read More](#)

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**
Find out how in this *free* and *practical* email course.

You're a Professional (and you need results)!

The field moves quickly...
How Long Can You Afford To Wait?

Get Your Start in Machine Learning

GET THE TRAINING YOU NEED

POPULAR

**Time Series Prediction with LSTM Recurrent Neural Networks in Python with Keras**

JULY 21, 2016

**Your First Machine Learning Project in Python Step-By-Step**

JUNE 10, 2016

**Develop Your First Neural Network in Python With Keras Step-By-Step**

MAY 24, 2016

**Sequence Classification with LSTM Recurrent Neural Networks in Python with Keras**

JULY 26, 2016

**How to Setup a Python Environment for Machine Learning and Deep Learning with Anaconda**

MARCH 13, 2017

**Time Series Forecasting with the Long Short-Term Memory Network in Python**

APRIL 7, 2017

**Multi-Class Classification Tutorial with the Keras Deep Learning Library**

JUNE 2, 2016

**Regression Tutorial with the Keras Deep Learning Library in Python**


JUNE 9, 2016

Get Your Start in Machine Learning




You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

[START MY EMAIL COURSE](#)[Get Your Start in Machine Learning](#)



Multivariate Time Series Forecasting with LSTMs in Keras

AUGUST 14, 2017



How to Implement the Backpropagation Algorithm From Scratch In Python

NOVEMBER 7, 2016

© 2017 Machine Learning Mastery. All Rights Reserved.

[Privacy](#) | [Contact](#) | [About](#)

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE