

[Start Here](#)[Blog](#)[Books](#)[About](#)[Contact](#)

Need help with LSTMs in Python? [Take the FREE Mini-Course.](#)

# How to Use Features in LSTM Networks for Time Series Forecasting

by **Jason Brownlee** on April 19, 2017 in **Long Short-Term Memory Networks**



The Long Short-Term Memory (LSTM) network in Keras supports multiple input features.

This raises the question as to whether lag observations for a univariate time series can be used as features for an LSTM and whether or not this improves forecast performance.

In this tutorial, we will investigate the use of lag observations as features in LSTM models in Python.

After completing this tutorial, you will know:

- How to develop a test harness to systematically evaluate LSTM features for time series forecasts

[Get Your Start in Machine Learning](#)

- The impact of using a varied number of lagged observations as input features for LSTM models.
- The impact of using a varied number of lagged observations and matching numbers of neurons for LSTM models.

Let's get started.



How to Use Features in LSTM Networks for Time Series Forecasting  
Photo by [Tom Hodgkinson](#), some rights reserved.

## Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

## Tutorial Overview

This tutorial is divided into 4 parts. They are:

1. Shampoo Sales Dataset

Get Your Start in Machine Learning

2. Experimental Test Harness
3. Experiments with Timesteps
4. Experiments with Timesteps and Neurons

## Environment

This tutorial assumes you have a Python SciPy environment installed. You can use either Python 2 or 3 with this example.

This tutorial assumes you have Keras v2.0 or higher installed with either the TensorFlow or Theano backend.

This tutorial also assumes you have scikit-learn, Pandas, NumPy, and Matplotlib installed.

If you need help setting up your Python environment, see this post:

- [How to Setup a Python Environment for Machine Learning and Deep Learning with Anaconda](#)

### Need help with LSTMs for Sequence Prediction?

Take my free 7-day email course and discover 6 different LSTM architectures.

Click to sign-up and also get a free PDF Ebook version of the course.

**Start Your FREE Mini-Course Now!**

### Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

## Shampoo Sales Dataset

This dataset describes the monthly number of sales of shampoo over a 3-year period.

The units are a sales count and there are 36 observations. The original dataset is credited to Makridakis, Wheelwright, and McGee (1983).

**Get Your Start in Machine Learning**

You can download and learn more about the dataset [here](#).

The example below loads and creates a plot of the loaded dataset.

```
1 # load and plot dataset
2 from pandas import read_csv
3 from pandas import datetime
4 from matplotlib import pyplot
5 # load dataset
6 def parser(x):
7     return datetime.strptime('190'+x, '%Y-%m')
8 series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True, date_parser=parser)
9 # summarize first few rows
10 print(series.head())
11 # line plot
12 series.plot()
13 pyplot.show()
```

Running the example loads the dataset as a Pandas Series and prints the first 5 rows.

```
1 Month
2 1901-01-01 266.0
3 1901-02-01 145.9
4 1901-03-01 183.1
5 1901-04-01 119.3
6 1901-05-01 180.3
7 Name: Sales, dtype: float64
```

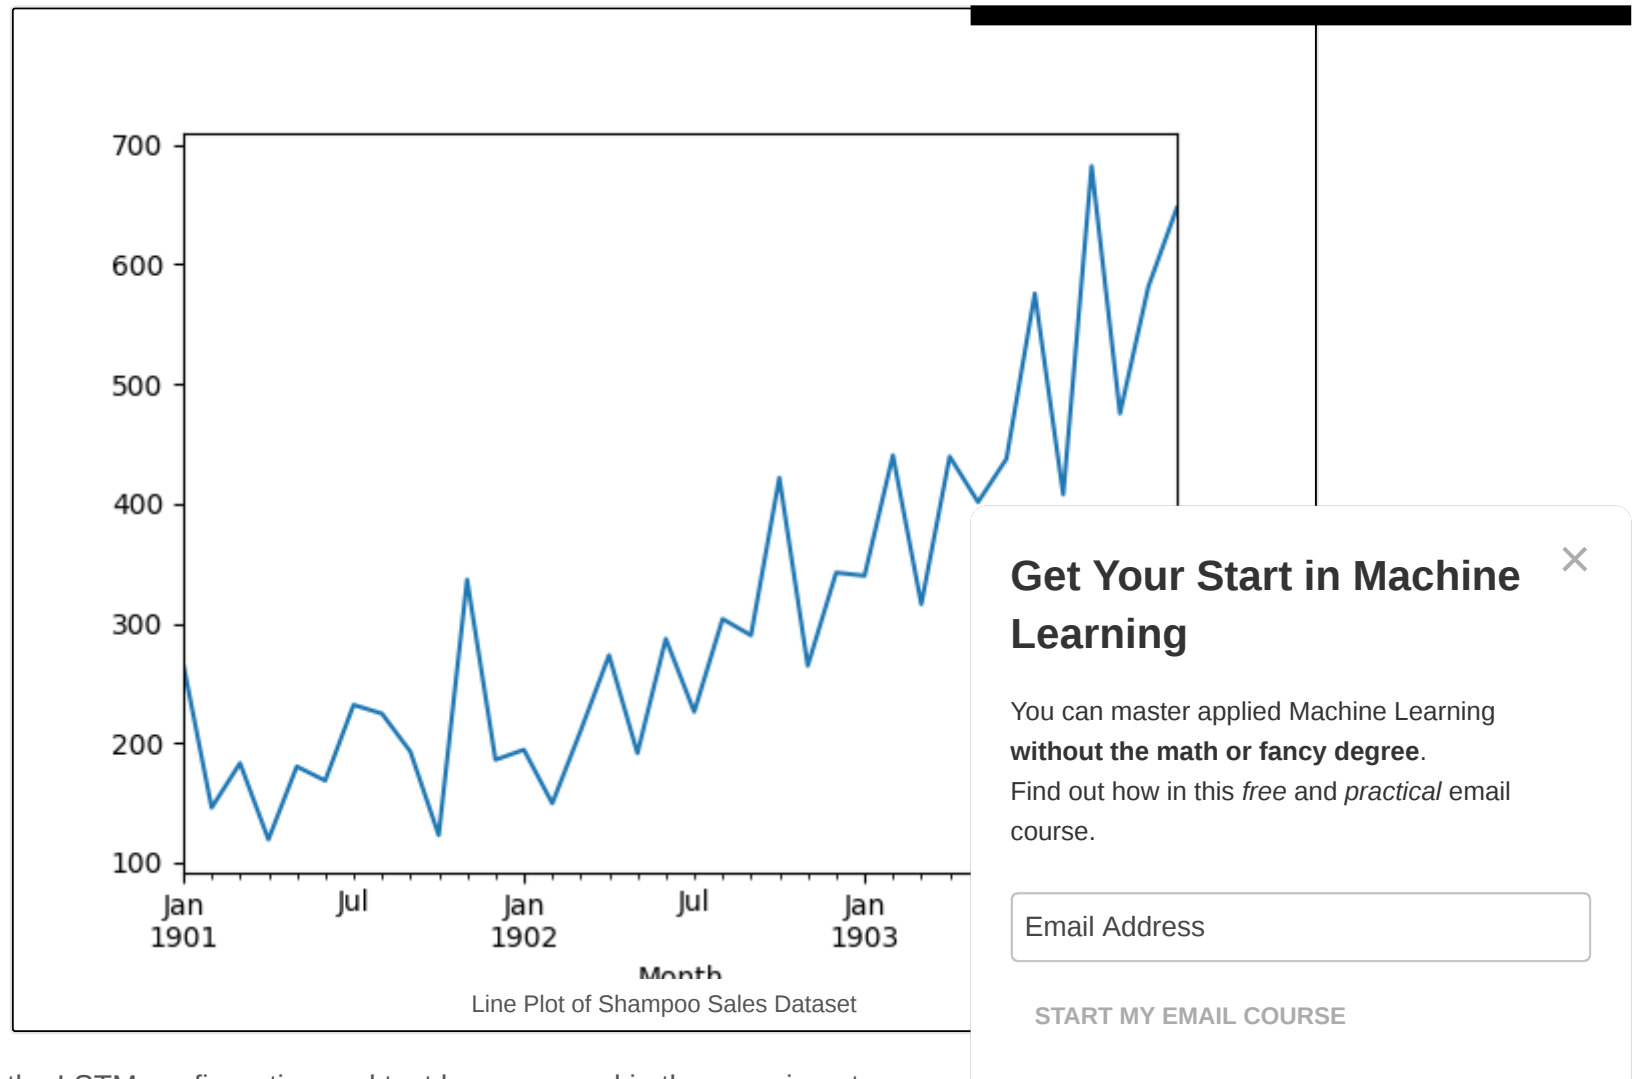
A line plot of the series is then created showing a clear increasing trend.

## Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning



Next, we will take a look at the LSTM configuration and test harness used in the experiment.

## Experimental Test Harness

This section describes the test harness used in this tutorial.

## Data Split

Get Your Start in Machine Learning

We will split the Shampoo Sales dataset into two parts: a training and a test set.

The first two years of data will be taken for the training dataset and the remaining one year of data will be used for the test set.

Models will be developed using the training dataset and will make predictions on the test dataset.

The persistence forecast (naive forecast) on the test dataset achieves an error of 136.761 monthly shampoo sales. This provides a lower acceptable bound of performance on the test set.

## Model Evaluation

A rolling-forecast scenario will be used, also called walk-forward model validation.

Each time step of the test dataset will be walked one at a time. A model will be used to make a forecast from the test set will be taken and made available to the model for the forecast on the next time step.

This mimics a real-world scenario where new Shampoo Sales observations would be available each month.

This will be simulated by the structure of the train and test datasets.

All forecasts on the test dataset will be collected and an error score calculated to summarize the skill will be used as it punishes large errors and results in a score that is in the same units as the forecast.

## Data Preparation

Before we can fit an LSTM model to the dataset, we must transform the data.

The following three data transforms are performed on the dataset prior to fitting a model and making a forecast.

1. **Transform the time series data so that it is stationary.** Specifically, a lag=1 differencing to remove the increasing trend in the data.
2. **Transform the time series into a supervised learning problem.** Specifically, the organization of data into input and output patterns where the observation at the previous time step is used as an input to forecast the observation at the current time step

### Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

3. **Transform the observations to have a specific scale.** Specifically, to rescale the data to values between -1 and 1 to meet the default hyperbolic tangent activation function of the LSTM model.

These transforms are inverted on forecasts to return them into their original scale before calculating an error score.

## LSTM Model

We will use a base stateful LSTM model with 1 neuron fit for 500 epochs.

A batch size of 1 is required as we will be using walk-forward validation and making one-step forecasts for each of the final 12 months of test data.

A batch size of 1 means that the model will be fit using online training (as opposed to batch training or mini-batch training). As a result, it is expected that the model fit will have some variance.

Ideally, more training epochs would be used (such as 1000 or 1500), but this was truncated to 500 to speed up the experiment.

The model will be fit using the efficient ADAM optimization algorithm and the mean squared error loss.

## Experimental Runs

Each experimental scenario will be run 10 times.

The reason for this is that the random initial conditions for an LSTM network can result in very different results.

Let's dive into the experiments.

## Experiments with Features

We will perform 5 experiments; each will use a different number of lag observations as features from 1 to 5.

A representation with a 1 input feature would be the default representation when using a stateful LSTM. Using 2 to 5 features is contrived. The hope would be that the additional context from the lagged observations may improve performance of the predictive model.

The univariate time series is converted to a supervised learning problem before training the model. The input variables ( $X$ ) used to predict the next observation ( $y$ ). As such, for each feature used in the model,

### Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

from the beginning of the dataset. This is because there are no prior observations to use as features for the first values in the dataset.

The complete code listing for testing 1 input feature is provided below.

The features parameter in the `run()` function is varied from 1 to 5 for each of the 5 experiments. In addition, the results are saved to file at the end of the experiment and this filename must also be changed for each different experimental run, e.g. `experiment_features_1.csv`, `experiment_features_2.csv`, etc.

```
1 from pandas import DataFrame
2 from pandas import Series
3 from pandas import concat
4 from pandas import read_csv
5 from pandas import datetime
6 from sklearn.metrics import mean_squared_error
7 from sklearn.preprocessing import MinMaxScaler
8 from keras.models import Sequential
9 from keras.layers import Dense
10 from keras.layers import LSTM
11 from math import sqrt
12 import matplotlib
13 import numpy
14 from numpy import concatenate
15
16 # date-time parsing function for loading the dataset
17 def parser(x):
18     return datetime.strptime('190'+x, '%Y-%m')
19
20 # frame a sequence as a supervised learning problem
21 def timeseries_to_supervised(data, lag=1):
22     df = DataFrame(data)
23     columns = [df.shift(i) for i in range(1, lag+1)]
24     columns.append(df)
25     df = concat(columns, axis=1)
26     return df
27
28 # create a differenced series
29 def difference(dataset, interval=1):
30     diff = list()
31     for i in range(interval, len(dataset)):
32         value = dataset[i] - dataset[i - interval]
33         diff.append(value)
34     return Series(diff)
35
36 # invert differenced value
37 def inverse_difference(history, yhat, interval=1):
38     return yhat + history[-interval]
```

## Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning



```

39
40 # scale train and test data to [-1, 1]
41 def scale(train, test):
42     # fit scaler
43     scaler = MinMaxScaler(feature_range=(-1, 1))
44     scaler = scaler.fit(train)
45     # transform train
46     train = train.reshape(train.shape[0], train.shape[1])
47     train_scaled = scaler.transform(train)
48     # transform test
49     test = test.reshape(test.shape[0], test.shape[1])
50     test_scaled = scaler.transform(test)
51     return scaler, train_scaled, test_scaled
52
53 # inverse scaling for a forecasted value
54 def invert_scale(scaler, X, yhat):
55     new_row = [x for x in X] + [yhat]
56     array = numpy.array(new_row)
57     array = array.reshape(1, len(array))
58     inverted = scaler.inverse_transform(array)
59     return inverted[0, -1]
60
61 # fit an LSTM network to training data
62 def fit_lstm(train, batch_size, nb_epoch, neurons):
63     X, y = train[:, 0:-1], train[:, -1]
64     X = X.reshape(X.shape[0], 1, X.shape[1])
65     model = Sequential()
66     model.add(LSTM(neurons, batch_input_shape=(batch_size, X.shape[1], X.shape[2]), stateful=True))
67     model.add(Dense(1))
68     model.compile(loss='mean_squared_error', optimizer='adam')
69     for i in range(nb_epoch):
70         model.fit(X, y, epochs=1, batch_size=batch_size, verbose=0, shuffle=False)
71         model.reset_states()
72     return model
73
74 # make a one-step forecast
75 def forecast_lstm(model, batch_size, X):
76     X = X.reshape(1, 1, len(X))
77     yhat = model.predict(X, batch_size=batch_size)
78     return yhat[0,0]
79
80 # run a repeated experiment
81 def experiment(repeats, series, features):
82     # transform data to be stationary
83     raw_values = series.values
84     diff_values = difference(raw_values, 1)
85     # transform data to be supervised learning

```

## Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```

86 supervised = timeseries_to_supervised(diff_values, features)
87 supervised_values = supervised.values[features,:]
88 # split data into train and test-sets
89 train, test = supervised_values[0:-12, :], supervised_values[-12:, :]
90 # transform the scale of the data
91 scaler, train_scaled, test_scaled = scale(train, test)
92 # run experiment
93 error_scores = list()
94 for r in range(repeats):
95     # fit the base model
96     lstm_model = fit_lstm(train_scaled, 1, 500, 1)
97     # forecast test dataset
98     predictions = list()
99     for i in range(len(test_scaled)):
100         # predict
101         X, y = test_scaled[i, 0:-1], test_scaled[i, -1]
102         yhat = forecast_lstm(lstm_model, 1, X)
103         # invert scaling
104         yhat = invert_scale(scaler, X, yhat)
105         # invert differencing
106         yhat = inverse_difference(raw_values, yhat, len(test_scaled)+1-i)
107         # store forecast
108         predictions.append(yhat)
109     # report performance
110     rmse = sqrt(mean_squared_error(raw_values[-12:], predictions))
111     print('%d) Test RMSE: %.3f' % (r+1, rmse))
112     error_scores.append(rmse)
113 return error_scores
114
115 # execute the experiment
116 def run():
117     # load dataset
118     series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True)
119     # experiment
120     repeats = 10
121     results = DataFrame()
122     # run experiment
123     features = 1
124     results['results'] = experiment(repeats, series, features)
125     # summarize results
126     print(results.describe())
127     # save results
128     results.to_csv('experiment_features_1.csv', index=False)
129
130 # entry point
131 run()

```

## Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Run the 5 different experiments for the 5 different numbers of features.

You can run them in parallel if you have sufficient memory and CPU resources. GPU resources are not required for these experiments and runs should be complete in minutes to tens of minutes.

After running the experiments, you should have 5 files containing the results, as follows:

- `experiment_features_1.csv`
- `experiment_features_2.csv`
- `experiment_features_3.csv`
- `experiment_features_4.csv`
- `experiment_features_5.csv`

We can write some code to load and summarize these results.

Specifically, it is useful to review both descriptive statistics from each run and compare the results for

Code to summarize the results is listed below.

```
1 from pandas import DataFrame
2 from pandas import read_csv
3 from matplotlib import pyplot
4 # load results into a dataframe
5 filenames = ['experiment_features_1.csv', 'experiment_features_2.csv',
6             'experiment_features_3.csv', 'experiment_features_4.csv', 'experiment_features_5.csv']
7 results = DataFrame()
8 for name in filenames:
9     results[name[11:-4]] = read_csv(name, header=0)
10 # describe all results
11 print(results.describe())
12 # box and whisker plot
13 results.boxplot()
14 pyplot.show()
```

Running the code first prints descriptive statistics for each set of results.

We can see from the average performance alone that the default of using a single feature resulted in the best performance. This is also shown when reviewing the median test RMSE (50th percentile).

## Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

	features_1	features_2	features_3	features_4	features_5
1					
2	count	10.000000	10.000000	10.000000	10.000000
3	mean	104.588249	126.597800	118.268251	107.694178
4	std	10.205840	18.639757	14.359983	8.683271
5	min	89.046814	93.857991	103.900339	93.702085
6	25%	97.850827	120.296634	107.664087	102.992045
7	50%	103.713285	133.582095	116.123790	106.116922
8	75%	111.441655	134.362198	121.794533	111.498255
9	max	122.341580	149.807155	152.412861	123.006088

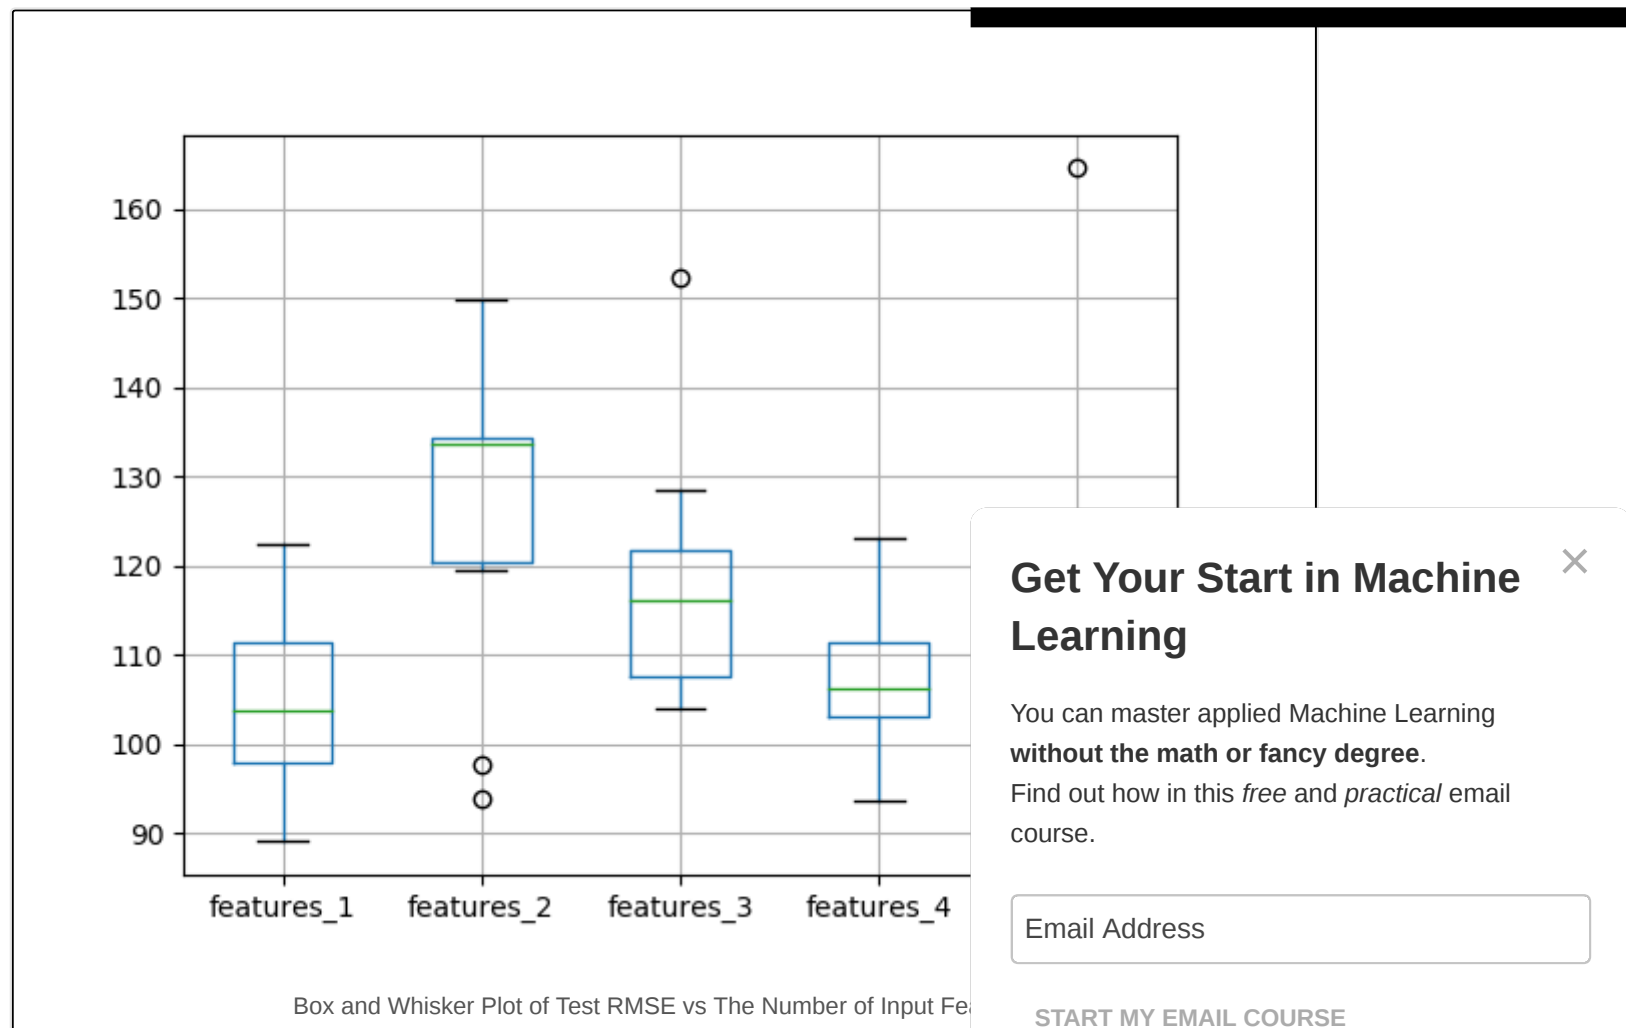
A box and whisker plot comparing the distributions of results is also created.

The plot tells the same story as the descriptive statistics. The test RMSE seems to leap up with 2 features and trend upward as the number of features is increased.

## Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



## Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree**. Find out how in this *free* and *practical* email course.

**START MY EMAIL COURSE**

The expectation of decreased error with the increase of features was not observed, at least with the dataset and LSTM configuration used.

This raises the question as to whether the capacity of the network is a limiting factor. We will look at this in the next section.

## Experiments with Features and Neurons

The number of neurons (also called units) in the LSTM network defines its learning capacity.

Get Your Start in Machine Learning

It is possible that in the previous experiments the use of one neuron limited the learning capacity of the network such that it was not capable of making effective use of the lagged observations as features.

We can repeat the above experiments and increase the number of neurons in the LSTM with the increase in features and see if it results in an increase in performance.

This can be achieved by changing the line in the experiment function from:

```
1 lstm_model = fit_lstm(train_scaled, 1, 500, 1, features)
```

to

```
1 lstm_model = fit_lstm(train_scaled, 1, 500, features, features)
```

In addition, we can keep the results written to file separate from the results from the first experiment example, changing:

```
1 results.to_csv('experiment_features_1.csv', index=False)
```

to

```
1 results.to_csv('experiment_features_1_neurons.csv', index=False)
```

Repeat the same 5 experiments with these changes.

After running these experiments, you should have 5 result files.

- `experiment_features_1_neurons.csv`
- `experiment_features_2_neurons.csv`
- `experiment_features_3_neurons.csv`
- `experiment_features_4_neurons.csv`
- `experiment_features_5_neurons.csv`

As in the previous experiment, we can load the results, calculate descriptive statistics, and create a box and whisker plot. The complete code listing is below.

```
1 from pandas import DataFrame
```

## Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```

2 from pandas import read_csv
3 from matplotlib import pyplot
4 # load results into a dataframe
5 filenames = ['experiment_features_1_neurons.csv', 'experiment_features_2_neurons.csv',
6             'experiment_features_3_neurons.csv', 'experiment_features_4_neurons.csv', 'experiment_features_5_neurons.csv']
7 results = DataFrame()
8 for name in filenames:
9     results[name[11:-12]] = read_csv(name, header=0)
10 # describe all results
11 print(results.describe())
12 # box and whisker plot
13 results.boxplot()
14 pyplot.show()

```

Running the code first prints descriptive statistics from each of the 5 experiments.

The results tell a different story to the first set of experiments with a one neuron LSTM. The average neurons and the number of features is set to one, then error increases as neurons and features are

	features_1	features_2	features_3	features_4	features_5
1 count	10.000000	10.000000	10.000000	10.000000	10.000000
2 mean	106.219189	138.411111	127.687128	154.281694	175.951500
3 std	16.100488	29.700981	21.411766	30.526294	44.839217
4 min	91.073598	92.641030	103.503546	94.063639	117.017109
5 25%	97.263723	125.748973	108.972440	134.805621	142.146601
6 50%	99.036766	133.639168	128.627349	162.295657	182.406707
7 75%	110.625302	146.896608	134.012859	176.969980	197.913894
8 max	146.638148	206.760081	170.899267	188.911768	250.685187

A box and whisker plot is created to compare the distributions.

The trend in spread and median performance almost shows a linear increase in test RMSE as the n

The linear trend may suggest that the increase network capacity is not given sufficient time to fit the data. Perhaps an increase in the number of epochs would be required as well.

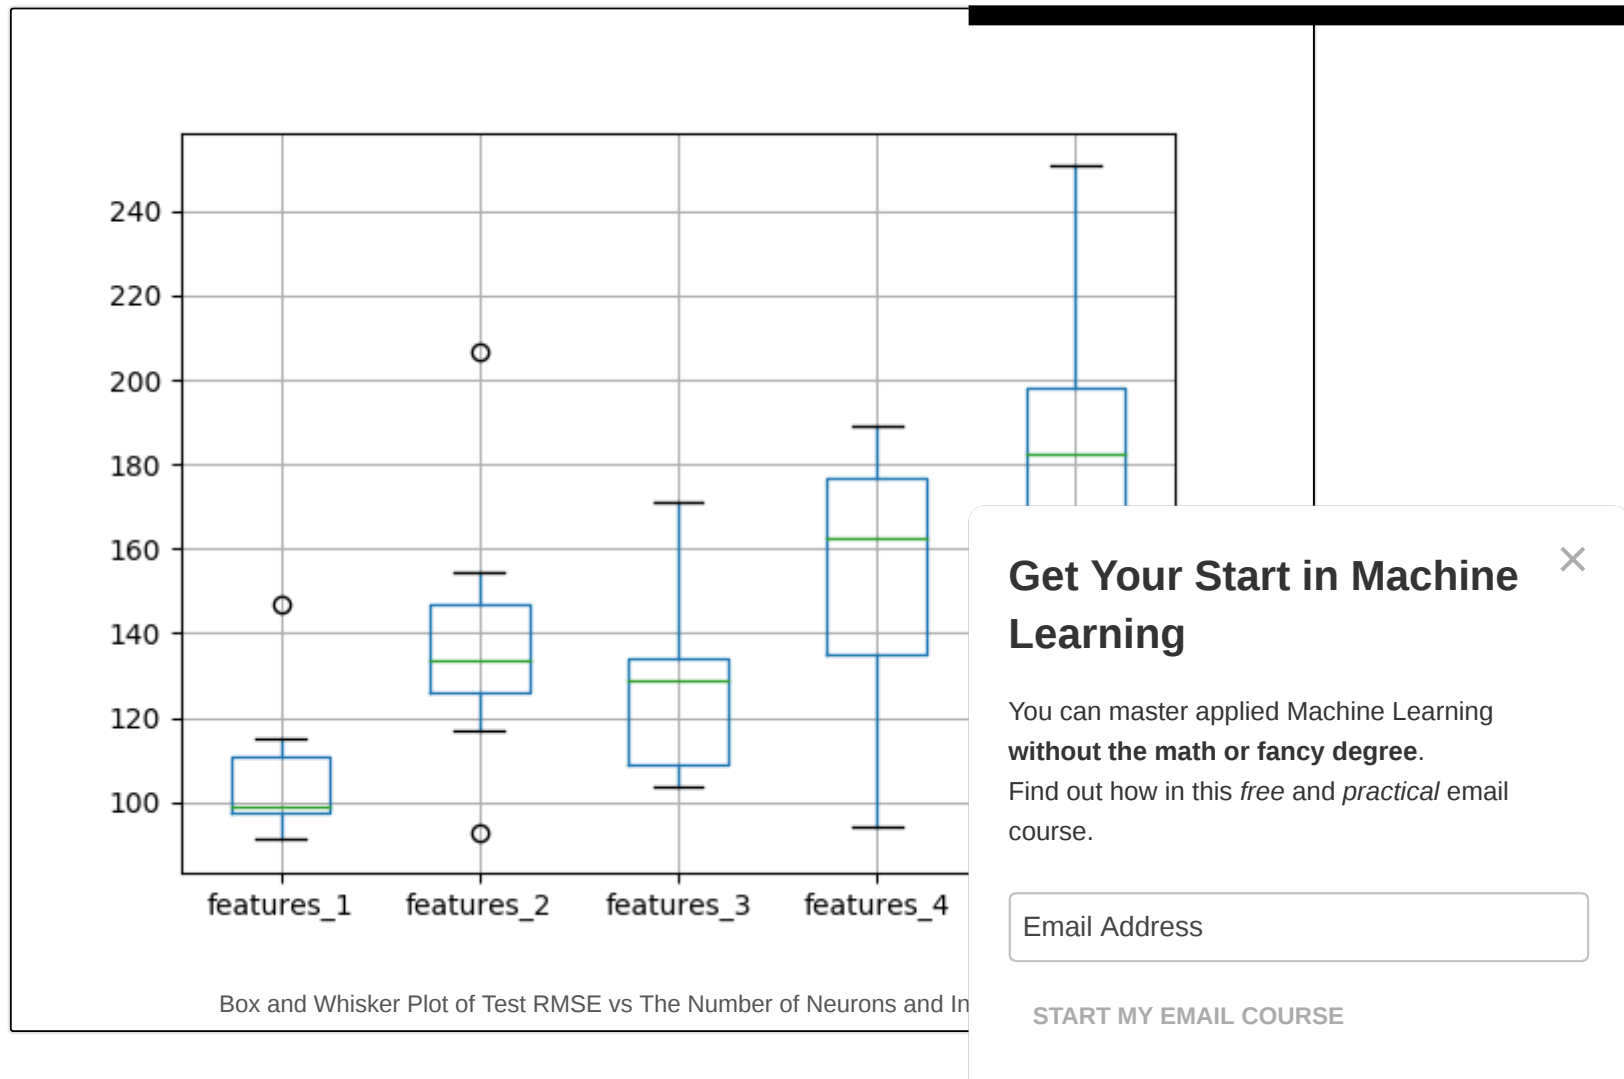
## Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning



## Experiments with Features and Neurons More Epochs

In this section, we repeat the above experiment to increase the number of neurons with the number of features but double the number of training epochs from 500 to 1000.

This can be achieved by changing the line in the experiment function from:

```
1 lstm_model = fit_lstm(train_scaled, 1, 500, features, features)
```

Get Your Start in Machine Learning



to

```
1 lstm_model = fit_lstm(train_scaled, 1, 1000, features, features)
```

In addition, we can keep the results written to file separate from the results from the previous experiment by adding a “1000” suffix to the filenames, for example, changing:

```
1 results.to_csv('experiment_features_1_neurons.csv', index=False)
```

to

```
1 results.to_csv('experiment_features_1_neurons1000.csv', index=False)
```

Repeat the same 5 experiments with these changes.

After running these experiments, you should have 5 result files.

- *experiment\_features\_1\_neurons1000.csv*
- *experiment\_features\_2\_neurons1000.csv*
- *experiment\_features\_3\_neurons1000.csv*
- *experiment\_features\_4\_neurons1000.csv*
- *experiment\_features\_5\_neurons1000.csv*

As in the previous experiment, we can load the results, calculate descriptive statistics, and create a plot below.

```
1 from pandas import DataFrame
2 from pandas import read_csv
3 from matplotlib import pyplot
4 # load results into a dataframe
5 filenames = ['experiment_features_1_neurons1000.csv', 'experiment_features_2_neurons1000.csv',
6             'experiment_features_3_neurons1000.csv', 'experiment_features_4_neurons1000.csv', 'experiment_features_5_neurons1000.csv']
7 results = DataFrame()
8 for name in filenames:
9     results[name[11:-16]] = read_csv(name, header=0)
10 # describe all results
11 print(results.describe())
12 # box and whisker plot
13 results.boxplot()
14 pyplot.show()
```

## Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Running the code first prints descriptive statistics from each of the 5 experiments.

The results tell a very similar story to the previous experiment with half the number of training epochs. On average, a model with 1 input feature and 1 neuron outperformed the other configurations.

1		features_1	features_2	features_3	features_4	features_5
2	count	10.000000	10.000000	10.000000	10.000000	10.000000
3	mean	109.262674	158.295172	120.340623	149.741882	201.992209
4	std	13.850525	32.288109	45.219564	53.121113	82.986691
5	min	95.927393	111.936394	83.983325	111.017837	78.040385
6	25%	98.754253	130.875314	95.198556	122.287208	148.840499
7	50%	103.990988	167.915523	110.256517	129.552084	188.498836
8	75%	116.055435	180.679252	122.158321	154.283676	234.519359
9	max	133.270446	204.260072	242.186747	288.907803	335.595974

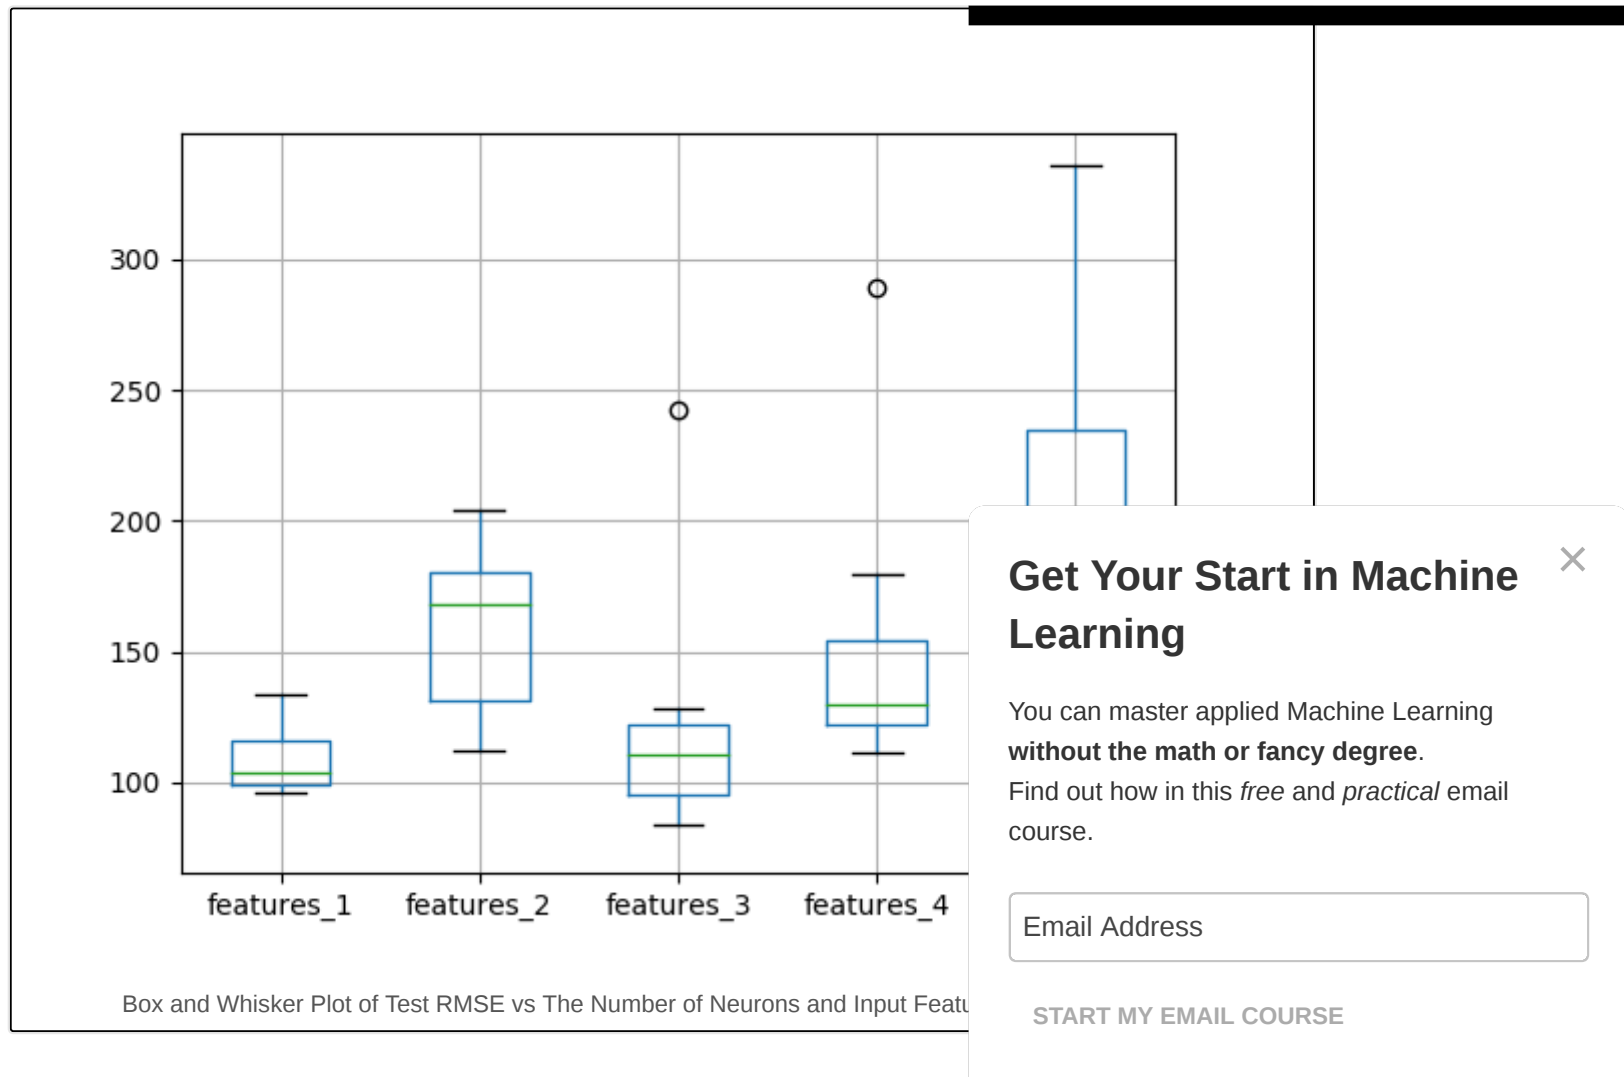
A box and whisker plot was also created to compare the distributions. In the plot, we see the same trend.

At least on this problem and with the chosen LSTM configuration, we do not see any clear benefit in

## Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



## Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree**. Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

## Extensions

This section lists some areas for further investigation that you may consider exploring.

- **Diagnostic Run Plots.** It may be helpful to review plots of train and test RMSE over epochs for multiple runs for a given experiment. This might help tease out whether overfitting or underfitting is taking place, and in turn, methods to address it.
- **Increase Repeats.** Using 10 repeats results in a relatively small population of test RMSE results. It is possible that increasing repeats to 30 or 100 (or even higher) may result in a more stable outcome.

Get Your Start in Machine Learning

Did you explore any of these extensions?

Share your findings in the comments below; I'd love to hear what you found.

## Summary

In this tutorial, you discovered how to investigate using lagged observations as input features in an LSTM network.

Specifically, you learned:

- How to develop a robust test harness for experimenting with input representation with LSTMs.
- How to use lagged observations as input features for time series forecasting with LSTMs.
- How to increase the learning capacity of the network with the increase of input features.

You discovered that the expectation that *“the use of lagged observations as input features improves chosen problem and LSTM configuration.”*

Do you have any questions?

Ask your questions in the comments below and I will do my best to answer them.

## Develop LSTMs for Sequence Prediction

### Develop Your Own LSTM models in Minutes

...with just a few lines of python code

Discover how in my new Ebook:

[Long Short-Term Memory Networks with Python](#)

It provides **self-study tutorials** on topics like:

*CNN LSTMs, Encoder-Decoder LSTMs, generative models, data preparation, making predictions* and much more...

**Finally Bring LSTM Recurrent Neural Networks to  
Your Sequence Predictions Projects**

## Get Your Start in Machine Learning


You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Skip the Academics. Just Results.

[Click to learn more.](#)



### About Jason Brownlee

Dr. Jason Brownlee is a husband, proud father, academic researcher, author, professional developer, and entrepreneur. He is passionate about helping developers get started and get good at applied machine learning. [Learn more.](#)

[View all posts by Jason Brownlee](#) →

## Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

[START MY EMAIL COURSE](#)

[◀ How to Use Timesteps in LSTM Networks for Time Series Forecasting](#)

[Stateful and Stateless LSTM for Time Series Forecasting with Python ▶](#)

[Get Your Start in Machine Learning](#)

## 14 Responses to *How to Use Features in LSTM Networks for Time Series Forecasting*



**Zach** May 5, 2017 at 7:40 am #

REPLY ↩

When using multiple features, is a stateful model the same as a stateless one with timesteps? Can you use timesteps in a stateful model with multiple features?



**Jason Brownlee** May 5, 2017 at 11:23 am #

REPLY ↩

It really depends on how you structure your data.

Stateless with all data in a batch or all time steps in a sample within a batch is the same as stateful.

I hope that makes things clearer. Also see this post:

<http://machinelearningmastery.com/stateful-stateless-lstm-time-series-forecasting-python/>

### Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



**Jason Ho** May 5, 2017 at 1:30 pm #

Hi, Dr. Jason Brownlee. I am confused that difference between time\_steps and features in a LSTM



**Jason Brownlee** May 6, 2017 at 7:32 am #

REPLY ↩

Great question.

Time steps are the past observations that the network will learn from (e.g. backpropagation through time). I have a few posts on BPTT coming out on the blog soon.

Features are different measurements taken at a given time step (e.g. pressure and temperature).

Get Your Start in Machine Learning

Does that help?



**Jason Ho** May 11, 2017 at 12:34 pm #

REPLY ↩

Thanks for Dr.Jason Brownlee's reply.

So we can take time steps and features just 2 type of LSTM's input?

And for multi-dimension input , just one neuron in input layer , does it really work? And does the number of input layer's neuron of LSTM need to be same as the input's dimension(time steps' dimension +features' dimension ) like the traditional BP Neural Network?



**Jason Brownlee** May 12, 2017 at 7:33 am #

No, the number of memory cells in the input layer does not have to match the numk



**Jason Ho** May 11, 2017 at 12:52 pm #

Compared to traditional method ,like AR, MA and ARIMA , What does the LSTM's advantages i



**Jason Brownlee** May 12, 2017 at 7:35 am #

Great question.

The promise of LSTMs (I have a scheduled post on this) is that they can learn the temporal dependence. That you just feed in sequences and do not have to specify the required window/lag (e.g. outcome from an ACF/PACF analysis).

## Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



**Jason Ho** May 12, 2017 at 1:59 pm #

REPLY ↩

And except that the LTSM do not have to specify the required window/lag ,is the LSTM's prediction accuracy better than ARIMA or the LSTM can do better than other tradition method thanks to its special model struction and deep layers?



**Jason Brownlee** May 13, 2017 at 6:11 am #

REPLY ↩

It may be, that is problem specific (e.g. no free lunch).



**Siddharth** May 31, 2017 at 10:51 pm #

Thanks Jason, yet again a very well structured blog post. I am learning a lot from this set of pos  
I have two questions for you:

- 1) What affect does adding more layers (going deeper) to the network have? More specifically, how do y  
when to make the network deeper (more layers)?
- 2) This might be a broad question, but specifically for time series forecasting, what "size" of data do you  
for a problem, is it possible to transfer learn from another LSTM?

Looking forward to your response!

## Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



**Jason Brownlee** June 2, 2017 at 12:50 pm #

REPLY ↩

More layers mean more hierarchical representational capacity, but slower to train and perhaps more likely to overfit. Try different configurations on your problem.

As much data as you can get.

You may be able to do transfer learning, I have not read anything about it for time series with LSTMs

Get Your Start in Machine Learning





**Ukesh** June 14, 2017 at 6:59 am #

REPLY ↩

Hi Jason,

Thank you for your wonderful post.

I am new to deep learning and LSTM. I have a very simple question. I have taken a sample of demands for 50 time steps and I am trying to forecast the demand value for the next 10 time steps using the sample to train the model.

But unfortunately, the closest I came is splitting the sample demands into 67 training % and 33 testing % and my forecast is only forecasting for the 33%. Can anybody help me with this issue?

I have borrowed your code and put it in the in github explaining this issue.

Thank you in advance.

<https://github.com/ukeshchawal/hello-world/blob/master/trial.ipynb>



**Jason Brownlee** June 14, 2017 at 8:53 am #

See this post for ways of evaluating time series forecasting methods:

<http://machinelearningmastery.com/backtest-machine-learning-models-time-series-forecasting/>

## Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Leave a Reply

Get Your Start in Machine Learning

Name (required)

Email (will not be published) (required)

Website

## Welcome to Machine Learning Mastery



Hi, I'm Dr. Jason Brownlee.  
My goal is to make practitioners like YOU awesome at applied machine learning.

[Read More](#)

## Get Your Start in Machine Learning ×


You can master applied Machine Learning **without the math or fancy degree.**  
Find out how in this *free* and *practical* email course.

## Deep Learning for Sequence Prediction

Cut through the math and research papers.  
Discover 4 Models, 6 Architectures, and 14 Tutorials.

[Get Your Start in Machine Learning](#)

Get Started With LSTMs in Python Today!



## POPULAR

**Time Series Prediction with LSTM Recurrent Neural Networks in Python with Keras**

JULY 21, 2016

**Your First Machine Learning Project in Python Step-By-Step**

JUNE 10, 2016

**Develop Your First Neural Network in Python With Keras Step-By-Step**

MAY 24, 2016

**Sequence Classification with LSTM Recurrent Neural Networks in Python with Keras**

JULY 26, 2016

**How to Setup a Python Environment for Machine Learning and Deep Learning with Anaconda**

MARCH 13, 2017

## Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

[START MY EMAIL COURSE](#)[Get Your Start in Machine Learning](#)



### Time Series Forecasting with the Long Short-Term Memory Network in Python

APRIL 7, 2017




### Multi-Class Classification Tutorial with the Keras Deep Learning Library

JUNE 2, 2016




### Regression Tutorial with the Keras Deep Learning Library in Python

JUNE 9, 2016



### Multivariate Time Series Forecasting with LSTMs in Keras

AUGUST 14, 2017



### How to Implement the Backpropagation Algorithm From Scratch In Python

NOVEMBER 7, 2016

## Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

© 2017 Machine Learning Mastery. All Rights Reserved.

[Privacy](#) | [Contact](#) | [About](#)

Get Your Start in Machine Learning