

[Start Here](#)[Blog](#)[Books](#)[About](#)[Contact](#)

Need help with machine learning? [Take the FREE Crash-Course.](#)

How to Use Word Embedding Layers for Deep Learning with Keras

by **Jason Brownlee** on October 4, 2017 in **Natural Language Processing**



Word embeddings provide a dense representation of words and their relative meanings.

They are an improvement over sparse representations used in simpler bag of word model representations.

Word embeddings can be learned from text data and reused among projects. They can also be learned as part of fitting a neural network on text data.

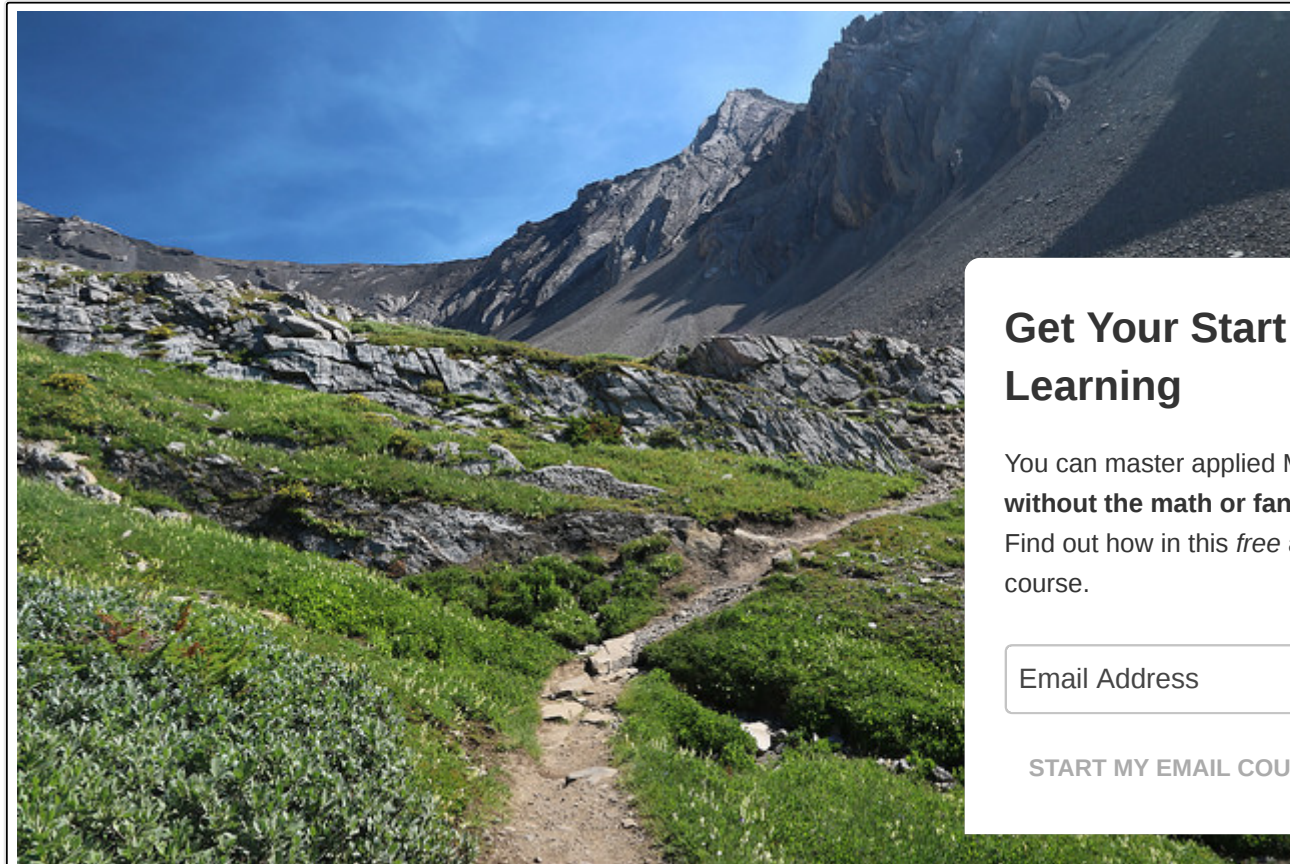
In this tutorial, you will discover how to use word embeddings for deep learning in Python with Keras.

After completing this tutorial, you will know:

[Get Your Start in Machine Learning](#)

- About word embeddings and that Keras supports word embeddings via the Embedding layer.
- How to learn a word embedding while fitting a neural network.
- How to use a pre-trained word embedding in a neural network.

Let's get started.



How to Use Word Embedding Layers for Deep Learning with Keras
Photo by [thisguy](#), some rights reserved.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Tutorial Overview

This tutorial is divided into 3 parts; they are:

Get Your Start in Machine Learning

1. Word Embedding
2. Keras Embedding Layer
3. Example of Learning an Embedding
4. Example of Using Pre-Trained GloVe Embedding

1. Word Embedding

A word embedding is a class of approaches for representing words and documents using a dense vector representation.

It is an improvement over more the traditional bag-of-word model encoding schemes where large sparse vectors were used to represent each word or to score each word within a vector to represent an entire vocabulary. These representations were sparse because the vocabularies were vast and a given word or document would be represented by a large vector comprised mostly of zero values.

Instead, in an embedding, words are represented by dense vectors where a vector represents the position of a word within the vector space.

The position of a word within the vector space is learned from text and is based on the words that surround it.

The position of a word in the learned vector space is referred to as its embedding.

Two popular examples of methods of learning word embeddings from text include:

- Word2Vec.
- GloVe.

In addition to these carefully designed methods, a word embedding can be learned as part of a deep learning model that is trained on a specific task, which tailors the model to a specific training dataset.

2. Keras Embedding Layer

Keras offers an [Embedding](#) layer that can be used for neural networks on text data.

It requires that the input data be integer encoded, so that each word is represented by a unique integer. This data preparation step can be performed using the [Tokenizer API](#) also provided with Keras.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

The Embedding layer is initialized with random weights and will learn an embedding for all of the words in the training dataset.

It is a flexible layer that can be used in a variety of ways, such as:

- It can be used alone to learn a word embedding that can be saved and used in another model later.
- It can be used as part of a deep learning model where the embedding is learned along with the model itself.
- It can be used to load a pre-trained word embedding model, a type of transfer learning.

The Embedding layer is defined as the first hidden layer of a network. It must specify 3 arguments:

It must specify 3 arguments:

- **input_dim**: This is the size of the vocabulary in the text data. For example, if your data is integer encoded with values between 0-10, then the size of the vocabulary would be 11 words.
- **output_dim**: This is the size of the vector space in which words will be embedded. It defines the dimensions in which words will be embedded. For example, it could be 32 or 100 or even larger. Test different values for your problem.
- **input_length**: This is the length of input sequences, as you would define for any input layer of a neural network. If your documents are comprised of 1000 words, this would be 1000.

For example, below we define an Embedding layer with a vocabulary of 200 (e.g. integer encoded words), 32 dimensions in which words will be embedded, and input documents that have 50 words each.

```
1 e = Embedding(200, 32, input_length=50)
```

The Embedding layer has weights that are learned. If you save your model to file, this will include weights for the Embedding layer.

The output of the *Embedding* layer is a 2D vector with one embedding for each word in the input sequence.

If you wish to connect a *Dense* layer directly to an Embedding layer, you must first flatten the 2D output matrix to a 1D vector using the *Flatten* layer.

Now, let's see how we can use an Embedding layer in practice.

3. Example of Learning an Embedding

In this section, we will look at how we can learn a word embedding while fitting a neural network on a dataset.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

We will define a small problem where we have 10 text documents, each with a comment about a piece of work a student submitted. Each text document is classified as positive “1” or negative “0”. This is a simple sentiment analysis problem.

First, we will define the documents and their class labels.

```
1 # define documents
2 docs = ['Well done!',
3         'Good work',
4         'Great effort',
5         'nice work',
6         'Excellent!',
7         'Weak',
8         'Poor effort!',
9         'not good',
10        'poor work',
11        'Could have done better.']
12 # define class labels
13 labels = [1,1,1,1,1,0,0,0,0,0]
```

Next, we can integer encode each document. This means that as input the Embedding layer will have other more sophisticated bag of word model encoding like counts or TF-IDF.

Keras provides the `one_hot()` function that creates a hash of each word as an efficient integer encoding. The vocabulary size is much larger than needed to reduce the probability of collisions from the hash function.

```
1 # integer encode the documents
2 vocab_size = 50
3 encoded_docs = [one_hot(d, vocab_size) for d in docs]
4 print(encoded_docs)
```

The sequences have different lengths and Keras prefers inputs to be vectorized and all inputs to have the length of 4. Again, we can do this with a built in Keras function, in this case the `pad_sequences()` function.

```
1 # pad documents to a max length of 4 words
2 max_length = 4
3 padded_docs = pad_sequences(encoded_docs, maxlen=max_length, padding='post')
4 print(padded_docs)
```

We are now ready to define our *Embedding* layer as part of our neural network model.

The *Embedding* has a vocabulary of 50 and an input length of 4. We will choose a small embedding

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

The model is a simple binary classification model. Importantly, the output from the *Embedding* layer will be 4 vectors of 8 dimensions each, one for each word. We flatten this to a one 32-element vector to pass on to the *Dense* output layer.

```
1 # define the model
2 model = Sequential()
3 model.add(Embedding(vocab_size, 8, input_length=max_length))
4 model.add(Flatten())
5 model.add(Dense(1, activation='sigmoid'))
6 # compile the model
7 model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['acc'])
8 # summarize the model
9 print(model.summary())
```

Finally, we can fit and evaluate the classification model.

```
1 # fit the model
2 model.fit(padded_docs, labels, epochs=50, verbose=0)
3 # evaluate the model
4 loss, accuracy = model.evaluate(padded_docs, labels, verbose=0)
5 print('Accuracy: %f' % (accuracy*100))
```

The complete code listing is provided below.

```
1 from keras.preprocessing.text import one_hot
2 from keras.preprocessing.sequence import pad_sequences
3 from keras.models import Sequential
4 from keras.layers import Dense
5 from keras.layers import Flatten
6 from keras.layers.embeddings import Embedding
7 # define documents
8 docs = ['Well done!',
9         'Good work',
10        'Great effort',
11        'nice work',
12        'Excellent!',
13        'Weak',
14        'Poor effort!',
15        'not good',
16        'poor work',
17        'Could have done better.']
18 # define class labels
19 labels = [1,1,1,1,1,0,0,0,0,0]
20 # integer encode the documents
21 vocab_size = 50
22 encoded_docs = [one_hot(d, vocab_size) for d in docs]
```

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```

23 print(encoded_docs)
24 # pad documents to a max length of 4 words
25 max_length = 4
26 padded_docs = pad_sequences(encoded_docs, maxlen=max_length, padding='post')
27 print(padded_docs)
28 # define the model
29 model = Sequential()
30 model.add(Embedding(vocab_size, 8, input_length=max_length))
31 model.add(Flatten())
32 model.add(Dense(1, activation='sigmoid'))
33 # compile the model
34 model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['acc'])
35 # summarize the model
36 print(model.summary())
37 # fit the model
38 model.fit(padded_docs, labels, epochs=50, verbose=0)
39 # evaluate the model
40 loss, accuracy = model.evaluate(padded_docs, labels, verbose=0)
41 print('Accuracy: %f' % (accuracy*100))

```

Running the example first prints the integer encoded documents.

```
1 [[6, 16], [42, 24], [2, 17], [42, 24], [18], [17], [22, 17], [27, 42], [22, 24], [49, 46],
```

Then the padded versions of each document are printed, making them all uniform length.

```

1 [[ 6 16  0  0]
2  [42 24  0  0]
3  [ 2 17  0  0]
4  [42 24  0  0]
5  [18  0  0  0]
6  [17  0  0  0]
7  [22 17  0  0]
8  [27 42  0  0]
9  [22 24  0  0]
10 [49 46 16 34]]

```

After the network is defined, a summary of the layers is printed. We can see that as expected, the output of the Embedding layer is a 4×8 matrix and this is squashed to a 32-element vector by the Flatten layer.

```

1 -----
2 Layer (type)                Output Shape          Param #
3 -----
4 embedding_1 (Embedding)      (None, 4, 8)          400
5 -----

```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning


```

6 flatten_1 (Flatten)          (None, 32)          0
7 -----
8 dense_1 (Dense)              (None, 1)           33
9 =====
10 Total params: 433
11 Trainable params: 433
12 Non-trainable params: 0
13 -----

```

Finally, the accuracy of the trained model is printed, showing that it learned the training dataset perfectly (which is not surprising).

```

1 Accuracy: 100.000000

```

You could save the learned weights from the Embedding layer to file for later use in other models.

You could also use this model generally to classify other documents that have the same kind vocabulary.

Next, let's look at loading a pre-trained word embedding in Keras.

4. Example of Using Pre-Trained GloVe Embedding

The Keras Embedding layer can also use a word embedding learned elsewhere.

It is common in the field of Natural Language Processing to learn, save, and make freely available word embeddings.

For example, the researchers behind GloVe method provide a suite of pre-trained word embeddings under a permissive license. See:

- [GloVe: Global Vectors for Word Representation](#)

The smallest package of embeddings is 822Mb, called “*glove.6B.zip*”. It was trained on a dataset of one billion tokens (words) with a vocabulary of 400 thousand words. There are a few different embedding vector sizes, including 50, 100, 200 and 300 dimensions.

You can download this collection of embeddings and we can seed the Keras *Embedding* layer with weights from the pre-trained embedding for the words in your training dataset.

This example is inspired by an example in the Keras project: [pretrained_word_embeddings.py](#).

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree**. Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

After downloading and unzipping, you will see a few files, one of which is “*glove.6B.100d.txt*”, which contains a 100-dimensional version of the embedding.

If you peek inside the file, you will see a token (word) followed by the weights (100 numbers) on each line. For example, below are the first line of the embedding ASCII text file showing the embedding for “*the*”.

```
1 the -0.038194 -0.24487 0.72812 -0.39961 0.083172 0.043953 -0.39141 0.3344 -0.57545 0.087459 0.28787 -0.06731 0.30906 -0.26384 -0.13231 -0.
```

As in the previous section, the first step is to define the examples, encode them as integers, then pad the sequences to be the same length.

In this case, we need to be able to map words to integers as well as integers to words.

Keras provides a `Tokenizer` class that can be fit on the training data, can convert text to sequences consistently by calling the `texts_to_sequences()` method on the `Tokenizer` class, and provides access to the dictionary mapping of words to integers i

```
1 # define documents
2 docs = ['Well done!',
3         'Good work',
4         'Great effort',
5         'nice work',
6         'Excellent!',
7         'Weak',
8         'Poor effort!',
9         'not good',
10        'poor work',
11        'Could have done better.']
12 # define class labels
13 labels = [1,1,1,1,1,0,0,0,0,0]
14 # prepare tokenizer
15 t = Tokenizer()
16 t.fit_on_texts(docs)
17 vocab_size = len(t.word_index) + 1
18 # integer encode the documents
19 encoded_docs = t.texts_to_sequences(docs)
20 print(encoded_docs)
21 # pad documents to a max length of 4 words
22 max_length = 4
23 padded_docs = pad_sequences(encoded_docs, maxlen=max_length, padding='post')
24 print(padded_docs)
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Next, we need to load the entire GloVe word embedding file into memory as a dictionary of word to e

Get Your Start in Machine Learning

```

1 # load the whole embedding into memory
2 embeddings_index = dict()
3 f = open('glove.6B.100d.txt')
4 for line in f:
5     values = line.split()
6     word = values[0]
7     coefs = asarray(values[1:], dtype='float32')
8     embeddings_index[word] = coefs
9 f.close()
10 print('Loaded %s word vectors.' % len(embeddings_index))

```

This is pretty slow. It might be better to filter the embedding for the unique words in your training data.

Next, we need to create a matrix of one embedding for each word in the training dataset. We can do that by enumerating all unique words in the `Tokenizer.word_index` and locating the embedding weight vector from the loaded GloVe embedding.

The result is a matrix of weights only for words we will see during training.

```

1 # create a weight matrix for words in training docs
2 embedding_matrix = zeros((vocab_size, 100))
3 for word, i in t.word_index.items():
4     embedding_vector = embeddings_index.get(word)
5     if embedding_vector is not None:
6         embedding_matrix[i] = embedding_vector

```

Now we can define our model, fit, and evaluate it as before.

The key difference is that the embedding layer can be seeded with the GloVe word embedding weights. The Embedding layer must be defined with `output_dim` set to 100. Finally, we do not want to update the weights, so we will set the `trainable` attribute for the model to be `False`.

```

1 e = Embedding(vocab_size, 100, weights=[embedding_matrix], input_length=4, trainable=False)

```

The complete worked example is listed below.

```

1 from numpy import asarray
2 from numpy import zeros
3 from keras.preprocessing.text import Tokenizer
4 from keras.preprocessing.sequence import pad_sequences
5 from keras.models import Sequential
6 from keras.layers import Dense
7 from keras.layers import Flatten

```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

```

8 from keras.layers import Embedding
9 # define documents
10 docs = ['Well done!',
11         'Good work',
12         'Great effort',
13         'nice work',
14         'Excellent!',
15         'Weak',
16         'Poor effort!',
17         'not good',
18         'poor work',
19         'Could have done better.']
20 # define class labels
21 labels = [1,1,1,1,1,0,0,0,0,0]
22 # prepare tokenizer
23 t = Tokenizer()
24 t.fit_on_texts(docs)
25 vocab_size = len(t.word_index) + 1
26 # integer encode the documents
27 encoded_docs = t.texts_to_sequences(docs)
28 print(encoded_docs)
29 # pad documents to a max length of 4 words
30 max_length = 4
31 padded_docs = pad_sequences(encoded_docs, maxlen=max_length, padding='post')
32 print(padded_docs)
33 # load the whole embedding into memory
34 embeddings_index = dict()
35 f = open('../glove_data/glove.6B/glove.6B.100d.txt')
36 for line in f:
37     values = line.split()
38     word = values[0]
39     coefs = asarray(values[1:], dtype='float32')
40     embeddings_index[word] = coefs
41 f.close()
42 print('Loaded %s word vectors.' % len(embeddings_index))
43 # create a weight matrix for words in training docs
44 embedding_matrix = zeros((vocab_size, 100))
45 for word, i in t.word_index.items():
46     embedding_vector = embeddings_index.get(word)
47     if embedding_vector is not None:
48         embedding_matrix[i] = embedding_vector
49 # define model
50 model = Sequential()
51 e = Embedding(vocab_size, 100, weights=[embedding_matrix], input_length=4, trainable=False)
52 model.add(e)
53 model.add(Flatten())
54 model.add(Dense(1, activation='sigmoid'))

```

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```

55 # compile the model
56 model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['acc'])
57 # summarize the model
58 print(model.summary())
59 # fit the model
60 model.fit(padded_docs, labels, epochs=50, verbose=0)
61 # evaluate the model
62 loss, accuracy = model.evaluate(padded_docs, labels, verbose=0)
63 print('Accuracy: %f' % (accuracy*100))

```

Running the example may take a bit longer, but then demonstrates that it is just as capable of fitting this simple problem.

```

1  [[6, 2], [3, 1], [7, 4], [8, 1], [9], [10], [5, 4], [11, 3], [5, 1], [12, 13, 2, 14]]
2
3  [[ 6  2  0  0]
4   [ 3  1  0  0]
5   [ 7  4  0  0]
6   [ 8  1  0  0]
7   [ 9  0  0  0]
8   [10  0  0  0]
9   [ 5  4  0  0]
10  [11  3  0  0]
11  [ 5  1  0  0]
12  [12 13  2 14]]
13
14 Loaded 400000 word vectors.
15
16 -----
17 Layer (type)                Output Shape          Param #
18 -----
19 embedding_1 (Embedding)      (None, 4, 100)        1500
20 -----
21 flatten_1 (Flatten)         (None, 400)           0
22 -----
23 dense_1 (Dense)             (None, 1)             401
24 =====
25 Total params: 1,901
26 Trainable params: 401
27 Non-trainable params: 1,500
28 -----
29
30
31 Accuracy: 100.000000

```

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

In practice, I would encourage you to experiment with learning a word embedding using a pre-trainer learning on top of a pre-trained embedding.

Get Your Start in Machine Learning

See what works best for your specific problem.

Further Reading

This section provides more resources on the topic if you are looking go deeper.

- [Word Embedding on Wikipedia](#)
- [Keras Embedding Layer API](#)
- [Using pre-trained word embeddings in a Keras model, 2016](#)
- [Example of using a pre-trained GloVe Embedding in Keras](#)
- [GloVe Embedding](#)
- [An overview of word embeddings and their connection to distributional semantic models, 2016](#)
- [Deep Learning, NLP, and Representations, 2014](#)

Summary

In this tutorial, you discovered how to use word embeddings for deep learning in Python with Keras.

Specifically, you learned:

- About word embeddings and that Keras supports word embeddings via the Embedding layer.
- How to learn a word embedding while fitting a neural network.
- How to use a pre-trained word embedding in a neural network.

Do you have any questions?

Ask your questions in the comments below and I will do my best to answer.



Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

About Jason Brownlee

Get Your Start in Machine Learning



Dr. Jason Brownlee is a husband, proud father, academic researcher, author, professional developer and a machine learning practitioner. He is dedicated to helping developers get started and get good at applied machine learning. [Learn more.](#)

[View all posts by Jason Brownlee](#) →

< [How to Prepare Text Data for Deep Learning with Keras](#)

[How to Develop Word Embeddings in Python with Gensim](#) >

24 Responses to *How to Use Word Embedding Layers for Deep Learning with Keras*



Mohammad October 4, 2017 at 7:58 am #

Thank you Jason,
I am excited to read more NLP posts.



Jason Brownlee October 4, 2017 at 8:03 am #

Thanks.



shiv October 5, 2017 at 10:07 am #

I split my data into 80-20 test-train and I'm still getting 100% accuracy. Any idea why? It is ~99% on epoch 1 and the rest its 100%.

Jason Brownlee October 5, 2017 at 5:22 pm #

[Get Your Start in Machine Learning](#)

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Consider using the procedure in this post to evaluate your model:
<https://machinelearningmastery.com/evaluate-skill-deep-learning-models/>



trulia October 6, 2017 at 12:47 pm #

REPLY ↩

Use drop-out 20%, your model is overfit!!



Sandy October 6, 2017 at 2:44 pm #

REPLY ↩

Thank you Jason. I always find things easier when reading your post.

I have a question about the vector of each word after training. For example, the word “done” in sentence “I could have done better!” from that word in sentence “Could have done better!”. Is that right? I mean the presentation of each word

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee October 7, 2017 at 5:48 am #

No, each word in the dictionary is represented differently, but the same word in different contexts

It is the word in its different contexts that is used to define the representation of the word.

Does that help?



Sandy October 7, 2017 at 5:37 pm #

REPLY ↩

Yes, thank you. But I still have a question. We will train each context separately, then after training the first context, in this case is “Well done!”, we will have a vector representation of the word “done”. After training the second context, “Could have done better”, we have another vector representation of the word “done”. So, which vector will we choose to be the representation of the word “done”?

I might misunderstand the procedure of training. Thank you for clarifying it for me.

Get Your Start in Machine Learning



Jason Brownlee October 8, 2017 at 8:32 am #

REPLY ↩

No. All examples where a word is used are used as part of the training of the representation of the word. There is only one representation for each word during and after training.



Sandy October 8, 2017 at 2:46 pm #

I got it. Thank you, Jason.



Chiedu October 7, 2017 at 5:36 pm #

Hi Jason,
any ideas on how to “filter the embedding for the unique words in your training data” as mentioned in the



Jason Brownlee October 8, 2017 at 8:32 am #

The mapping of word to vector dictionary is built into Gensim, you can access it directly to r
model.wv.vocab



Abbey October 8, 2017 at 2:19 am #

Hi, Jason

Good day.

REPLY ↩

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**
Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

I just need your suggestion and example. I have two different dataset, where one is structured and the other is unstructured. The goal is to use the structured to construct a representation for the unstructured, so apply use word embedding on the two input data but how can I find the average of the two embedding and flatten it to one before feeding the layer into CNN and LSTM.

Looking forward to your response.

Regards

Abbey



Jason Brownlee October 8, 2017 at 8:40 am #

REPLY ↩

Sorry, what was your question?

If your question was if this is a good approach, my advice is to try it and see.



Abiodun Modupe October 9, 2017 at 7:46 pm #

Hi, Jason

How can I find the average of the word embedding from the two input?

Regards

Abbey



Jason Brownlee October 10, 2017 at 7:43 am #

Perhaps you could retrieve the vectors for each word and take their average?

Perhaps you can use the Gensim API to achieve this result?

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Vinu October 9, 2017 at 5:54 pm #

Get Your Start in Machine Learning



Hi Jason...Could you also help us with R codes for using Pre-Trained GloVe Embedding



Jason Brownlee October 10, 2017 at 7:43 am #

REPLY ↩

Sorry, I don't have R code for word embeddings.



Hao October 12, 2017 at 5:49 pm #

REPLY ↩

Hi Jason, really appreciate that you answered all the replies! I am planning to try both CNN and RNN (maybe LSTM & GRU) on text classification. Most of my documents are less than 100 words long, but about 5 % are longer than 500 words. How do I set it to be 1000, will it degrade the learning result? Should I just use 100? Will it be different in the case? Thank you!

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee October 13, 2017 at 5:45 am #

I would recommend experimenting with different configurations and see how the impact mo



Michael October 13, 2017 at 10:22 am #

I'd like to thank you for this post. I've been struggling to understand this precise way of using keras for a week now and this is the only post I've found that actually explains what each step in the process is doing – and provides code that self-documents what the data looks like as the model is constructed and trained. This makes it so much easier to adapt to my particular requirements.



Jason Brownlee October 13, 2017 at 2:55 pm #

REPLY ↩

Get Your Start in Machine Learning

Thanks, I'm glad it helped.



Azim October 17, 2017 at 5:47 pm #

REPLY ↩

In the above Keras example, how can we predict a list of context words given a word? Lets say i have a word named 'sudoku' and want to predict the sourrounding words. how can we use word2vec from keras to do that?



Jason Brownlee October 18, 2017 at 5:32 am #

REPLY ↩

It sounds like you are describing a language model. We can use LSTMs to learn these rela

Here is an example:

<https://machinelearningmastery.com/text-generation-lstm-recurrent-neural-networks-python-keras/>

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Leave a Reply

 Name (required)

Email (will not be published) (required)

Website

[SUBMIT COMMENT](#)

Welcome to Machine Learning Mastery



Hi, I'm Dr. Jason Brownlee.

My goal is to make practitioners like YOU awesome at applied machine learning.

[Read More](#)

You're a Professional (and you need results)!

The field moves quickly...

How Long Can You Afford To Wait?

Take Action Now!

GET THE TRAINING YOU NEED

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

POPULAR



Time Series Prediction with LSTM Recurrent Neural Networks in Python with Keras

JULY 21, 2016

[Get Your Start in Machine Learning](#)

**Your First Machine Learning Project in Python Step-By-Step**

JUNE 10, 2016

**Develop Your First Neural Network in Python With Keras Step-By-Step**

MAY 24, 2016

**Sequence Classification with LSTM Recurrent Neural Networks in Python with Keras**

JULY 26, 2016

**How to Setup a Python Environment for Machine Learning and Deep Learning with Anaconda**

MARCH 13, 2017

**Time Series Forecasting with the Long Short-Term Memory Network in Python**

APRIL 7, 2017

**Multi-Class Classification Tutorial with the Keras Deep Learning Library**

JUNE 2, 2016

**Regression Tutorial with the Keras Deep Learning Library in Python**

JUNE 9, 2016

**Multivariate Time Series Forecasting with LSTMs in Keras**

AUGUST 14, 2017

**How to Implement the Backpropagation Algorithm From Scratch In Python**

NOVEMBER 7, 2016

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

[START MY EMAIL COURSE](#)

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning