

[Start Here](#)[Blog](#)[Books](#)[About](#)[Contact](#)

Need help with LSTMs in Python? [Take the FREE Mini-Course.](#)

# Instability of Online Learning for Stateful LSTM for Time Series Forecasting

by **Jason Brownlee** on April 24, 2017 in **Long Short-Term Memory Networks**



Some neural network configurations can result in an unstable model.

This can make them hard to characterize and compare to other model configurations on the same problem using descriptive statistics.

One good example of a seemingly unstable model is the use of online learning (a batch size of 1) for a stateful Long Short-Term Memory (LSTM) model.

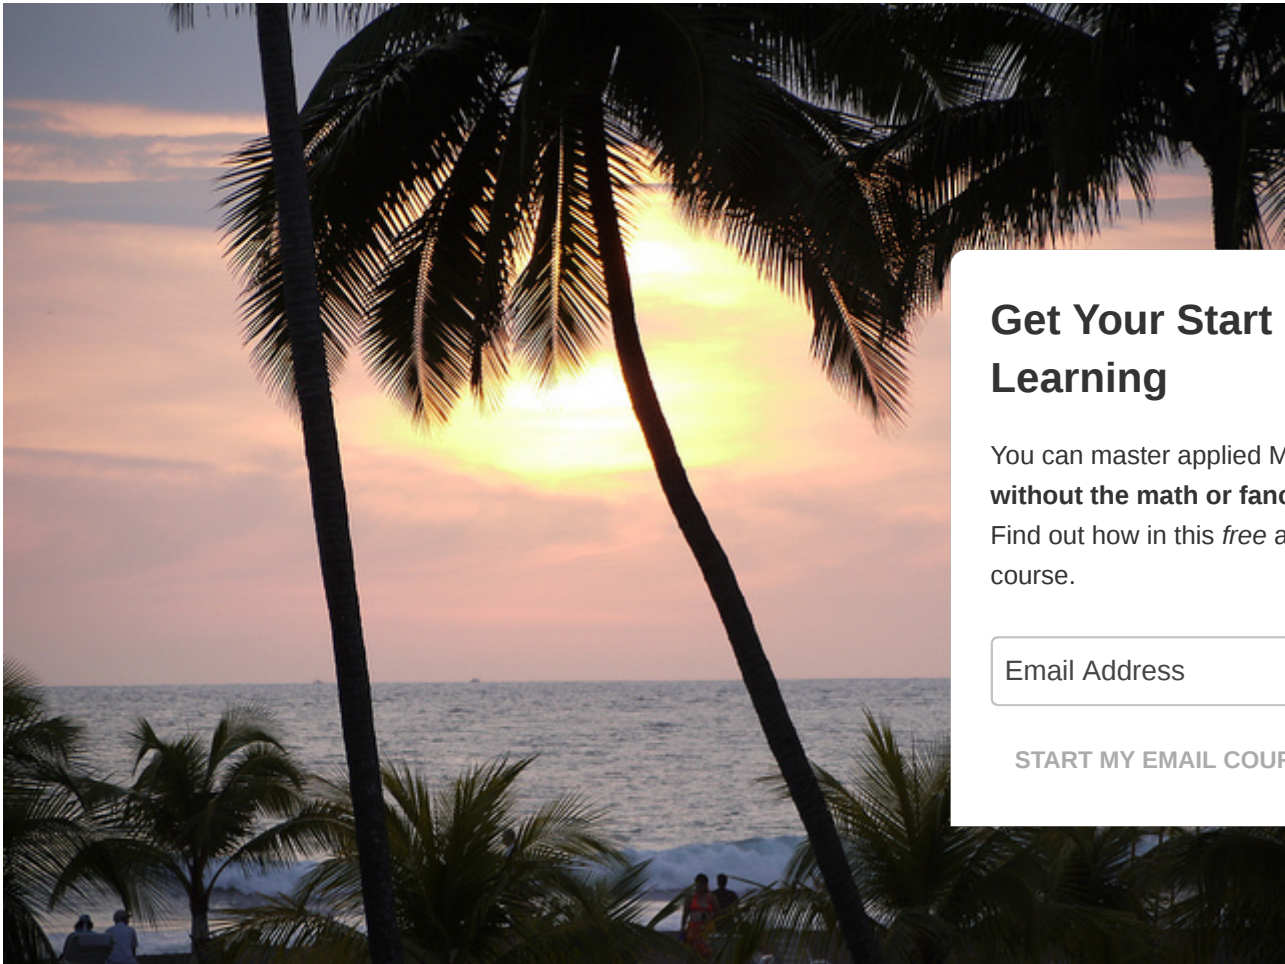
In this tutorial, you will discover how to explore the results of a stateful LSTM fit using online learning on a standard time series forecasting problem.

After completing this tutorial, you will know:

[Get Your Start in Machine Learning](#)

- How to design a robust test harness for evaluating LSTM models on time series forecasting problems.
- How to analyze a population of results, including summary statistics, spread, and distribution of results.
- How to analyze the impact of increasing the number of repeats for an experiment.

Let's get started.



## Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

[START MY EMAIL COURSE](#)

Instability of Online Learning for Stateful LSTM for Time Series Forecasting  
Photo by [Magnus Brath](#), some rights reserved.

## Model Instability

Get Your Start in Machine Learning

When you train the same network on the same data more than once, you may get very different results.

This is because neural networks are initialized randomly and the optimization nature of how they are fit to the training data can result in different final weights within the network. These different networks can in turn result in varied predictions given the same input data.

As a result, it is important to repeat any experiment on neural networks multiple times to find an averaged expected performance.

For more on the stochastic nature of machine learning algorithms like neural networks, see the post:

- [Embrace Randomness in Machine Learning](#)

The batch size in a neural network defines how often the weights within the network are updated given exposure to a training dataset.

A batch size of 1 means that the network weights are updated after each single row of training data. that can learn quickly, but a configuration that can be quite unstable.

In this tutorial, we will explore the instability of online learning for a stateful LSTM configuration for time series forecasting.

We will explore this by looking at the average performance of an LSTM configuration on a standard time series forecasting task. We will look at the number of repeats of the experiment.

That is, we will re-train the same model configuration on the same data many times and look at the performance. We will then review how unstable the model can be.

## Tutorial Overview

This tutorial is broken down into 6 parts. They are:

1. Shampoo Sales Dataset
2. Experimental Test Harness
3. Code and Collect Results
4. Basic Statistics on Results
5. Repeats vs Test RMSE
6. Review of Results

### Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

## Environment

This tutorial assumes you have a Python SciPy environment installed. You can use either Python 2 or 3 with this example.

This tutorial assumes you have Keras v2.0 or higher installed with either the TensorFlow or Theano backend.

This tutorial also assumes you have scikit-learn, Pandas, NumPy, and Matplotlib installed.

Next, let's take a look at a standard time series forecasting problem that we can use as context for this experiment.

If you need help setting up your Python environment, see this post:

- [How to Setup a Python Environment for Machine Learning and Deep Learning with Anaconda](#)

### Need help with LSTMs for Sequence Prediction?

Take my free 7-day email course and discover 6 different LSTM architectures.

Click to sign-up and also get a free PDF Ebook version of the course.

[Start Your FREE Mini-Course Now!](#)

### Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

[START MY EMAIL COURSE](#)

## Shampoo Sales Dataset

This dataset describes the monthly number of sales of shampoo over a 3-year period.

The units are a sales count and there are 36 observations. The original dataset is credited to Makridakis, Wheelwright, and Hyndman (1998).

You can download and learn more about the dataset [here](#).

[Get Your Start in Machine Learning](#)

The example below loads and creates a plot of the loaded dataset.

```
1 # load and plot dataset
2 from pandas import read_csv
3 from pandas import datetime
4 from matplotlib import pyplot
5 # load dataset
6 def parser(x):
7     return datetime.strptime('190'+x, '%Y-%m')
8 series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True, date_parser=parser)
9 # summarize first few rows
10 print(series.head())
11 # line plot
12 series.plot()
13 pyplot.show()
```

Running the example loads the dataset as a Pandas Series and prints the first 5 rows.


```
1 Month
2 1901-01-01 266.0
3 1901-02-01 145.9
4 1901-03-01 183.1
5 1901-04-01 119.3
6 1901-05-01 180.3
7 Name: Sales, dtype: float64
```

A line plot of the series is then created showing a clear increasing trend.

## Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Next, we will take a look at the LSTM configuration and test harness used in the experiment.

## Experimental Test Harness

This section describes the test harness used in this tutorial.

## Data Split

Get Your Start in Machine Learning

We will split the Shampoo Sales dataset into two parts: a training and a test set.

The first two years of data will be taken for the training dataset and the remaining one year of data will be used for the test set.

Models will be developed using the training dataset and will make predictions on the test dataset.

The persistence forecast (naive forecast) on the test dataset achieves an error of 136.761 monthly shampoo sales. This provides a lower acceptable bound of performance on the test set.

## Model Evaluation

A rolling-forecast scenario will be used, also called walk-forward model validation.

Each time step of the test dataset will be walked one at a time. A model will be used to make a forecast from the test set will be taken and made available to the model for the forecast on the next time step.

This mimics a real-world scenario where new Shampoo Sales observations would be available each month.

This will be simulated by the structure of the train and test datasets.

All forecasts on the test dataset will be collected and an error score calculated to summarize the skill will be used as it punishes large errors and results in a score that is in the same units as the forecast.

## Data Preparation

Before we can fit an LSTM model to the dataset, we must transform the data.

The following three data transforms are performed on the dataset prior to fitting a model and making a forecast.

1. **Transform the time series data so that it is stationary.** Specifically a lag=1 differencing to remove the increasing trend in the data.
2. **Transform the time series into a supervised learning problem.** Specifically the organization of data into input and output patterns where the observation at the previous time step is used as an input to forecast the observation at the current time step

### Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

3. **Transform the observations to have a specific scale.** Specifically, to rescale the data to values between -1 and 1 to meet the default hyperbolic tangent activation function of the LSTM model.

These transforms are inverted on forecasts to return them into their original scale before calculating an error score.

## LSTM Model

We will use a base stateful LSTM model with 1 neuron fit for 1000 epochs.

A batch size of 1 is required as we will be using walk-forward validation and making one-step forecasts for each of the final 12 months of test data.

A batch size of 1 means that the model will be fit using online training (as opposed to batch training or mini-batch training). As a result, it is expected that the model fit will have some variance.

Ideally, more training epochs would be used (such as 1500), but this was truncated to 1000 to keep the training time reasonable.

The model will be fit using the efficient ADAM optimization algorithm and the mean squared error loss function.

## Experimental Runs

Each experimental scenario will be run 100 times and the RMSE score on the test set will be recorded.

All test RMSE scores are written to file for later analysis.

Let's dive into the experiments.

## Code and Collect Results

The complete code listing is provided below.

It may take a few hours to run on modern hardware.

```
1 from pandas import DataFrame
2 from pandas import Series
3 from pandas import concat
4 from pandas import read_csv
```

### Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning



```
5 from pandas import datetime
6 from sklearn.metrics import mean_squared_error
7 from sklearn.preprocessing import MinMaxScaler
8 from keras.models import Sequential
9 from keras.layers import Dense
10 from keras.layers import LSTM
11 from math import sqrt
12 import matplotlib
13 import numpy
14 from numpy import concatenate
15
16 # date-time parsing function for loading the dataset
17 def parser(x):
18     return datetime.strptime('190'+x, '%Y-%m')
19
20 # frame a sequence as a supervised learning problem
21 def timeseries_to_supervised(data, lag=1):
22     df = DataFrame(data)
23     columns = [df.shift(i) for i in range(1, lag+1)]
24     columns.append(df)
25     df = concat(columns, axis=1)
26     return df
27
28 # create a differenced series
29 def difference(dataset, interval=1):
30     diff = list()
31     for i in range(interval, len(dataset)):
32         value = dataset[i] - dataset[i - interval]
33         diff.append(value)
34     return Series(diff)
35
36 # invert differenced value
37 def inverse_difference(history, yhat, interval=1):
38     return yhat + history[-interval]
39
40 # scale train and test data to [-1, 1]
41 def scale(train, test):
42     # fit scaler
43     scaler = MinMaxScaler(feature_range=(-1, 1))
44     scaler = scaler.fit(train)
45     # transform train
46     train = train.reshape(train.shape[0], train.shape[1])
47     train_scaled = scaler.transform(train)
48     # transform test
49     test = test.reshape(test.shape[0], test.shape[1])
50     test_scaled = scaler.transform(test)
51     return scaler, train_scaled, test_scaled
```

## Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```

52
53 # inverse scaling for a forecasted value
54 def invert_scale(scaler, X, yhat):
55     new_row = [x for x in X] + [yhat]
56     array = numpy.array(new_row)
57     array = array.reshape(1, len(array))
58     inverted = scaler.inverse_transform(array)
59     return inverted[0, -1]
60
61 # fit an LSTM network to training data
62 def fit_lstm(train, batch_size, nb_epoch, neurons):
63     X, y = train[:, 0:-1], train[:, -1]
64     X = X.reshape(X.shape[0], 1, X.shape[1])
65     model = Sequential()
66     model.add(LSTM(neurons, batch_input_shape=(batch_size, X.shape[1], X.shape[2]), stateful=True))
67     model.add(Dense(1))
68     model.compile(loss='mean_squared_error', optimizer='adam')
69     for i in range(nb_epoch):
70         model.fit(X, y, epochs=1, batch_size=batch_size, verbose=0, shuffle=False)
71         model.reset_states()
72     return model
73
74 # make a one-step forecast
75 def forecast_lstm(model, batch_size, X):
76     X = X.reshape(1, 1, len(X))
77     yhat = model.predict(X, batch_size=batch_size)
78     return yhat[0,0]
79
80 # run a repeated experiment
81 def experiment(repeats, series):
82     # transform data to be stationary
83     raw_values = series.values
84     diff_values = difference(raw_values, 1)
85     # transform data to be supervised learning
86     supervised = timeseries_to_supervised(diff_values, 1)
87     supervised_values = supervised.values[1:,:]
88     # split data into train and test-sets
89     train, test = supervised_values[0:-12, :], supervised_values[-12:, :]
90     # transform the scale of the data
91     scaler, train_scaled, test_scaled = scale(train, test)
92     # run experiment
93     error_scores = list()
94     for r in range(repeats):
95         # fit the base model
96         lstm_model = fit_lstm(train_scaled, 1, 1000, 1)
97         # forecast test dataset
98         predictions = list()

```

## Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```

99     for i in range(len(test_scaled)):
100         # predict
101         X, y = test_scaled[i, 0:-1], test_scaled[i, -1]
102         yhat = forecast_lstm(lstm_model, 1, X)
103         # invert scaling
104         yhat = invert_scale(scaler, X, yhat)
105         # invert differencing
106         yhat = inverse_difference(raw_values, yhat, len(test_scaled)+1-i)
107         # store forecast
108         predictions.append(yhat)
109     # report performance
110     rmse = sqrt(mean_squared_error(raw_values[-12:], predictions))
111     print('%d) Test RMSE: %.3f' % (r+1, rmse))
112     error_scores.append(rmse)
113     return error_scores
114
115 # execute the experiment
116 def run():
117     # load dataset
118     series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True)
119     # experiment
120     repeats = 100
121     results = DataFrame()
122     # run experiment
123     results['results'] = experiment(repeats, series)
124     # summarize results
125     print(results.describe())
126     # save results
127     results.to_csv('experiment_stateful.csv', index=False)
128
129 # entry point
130 run()

```

Running the experiment saves the RMSE scores of the fit model on the test dataset.

Results are saved to the file “*experiment\_stateful.csv*”.

A truncated listing of the results is provided below.

Re-running the experiment yourself may give a different population of results because we did not seed the random number generator.

```

1  ...
2  116.39769471284067
3  105.0459745537738
4  93.57827109861229

```

## Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```

5 128.973001927212
6 97.02915084460737
7 198.56877142225886
8 113.09568645243242
9 97.84127724751188
10 124.60413895331735
11 111.62139008607713

```

## Basic Statistics on Results

We can start off by calculating some basic statistics on the entire population of 100 test RMSE scores.

Generally, we expect machine learning results to have a Gaussian distribution. This allows us to report the mean and standard deviation of a model and indicate a confidence interval for the model when making predictions on unseen data.

The snippet below loads the result file and calculates some descriptive statistics.

```

1 from pandas import DataFrame
2 from pandas import read_csv
3 from numpy import mean
4 from numpy import std
5 from matplotlib import pyplot
6 # load results file
7 results = read_csv('experiment_stateful.csv', header=0)
8 # descriptive stats
9 print(results.describe())
10 # box and whisker plot
11 results.boxplot()
12 pyplot.show()

```

Running the example prints descriptive statistics from the results.

We can see that on average, the configuration achieved an RMSE of about 107 monthly shampoo sales with a standard deviation of about 17.

We can also see that the best test RMSE observed was about 90 sales, whereas the worse was just under 200, which is quite a spread of scores.

```

1          results
2 count  100.000000
3 mean   107.051146
4 std     17.694512
5 min     90.407323

```

### Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

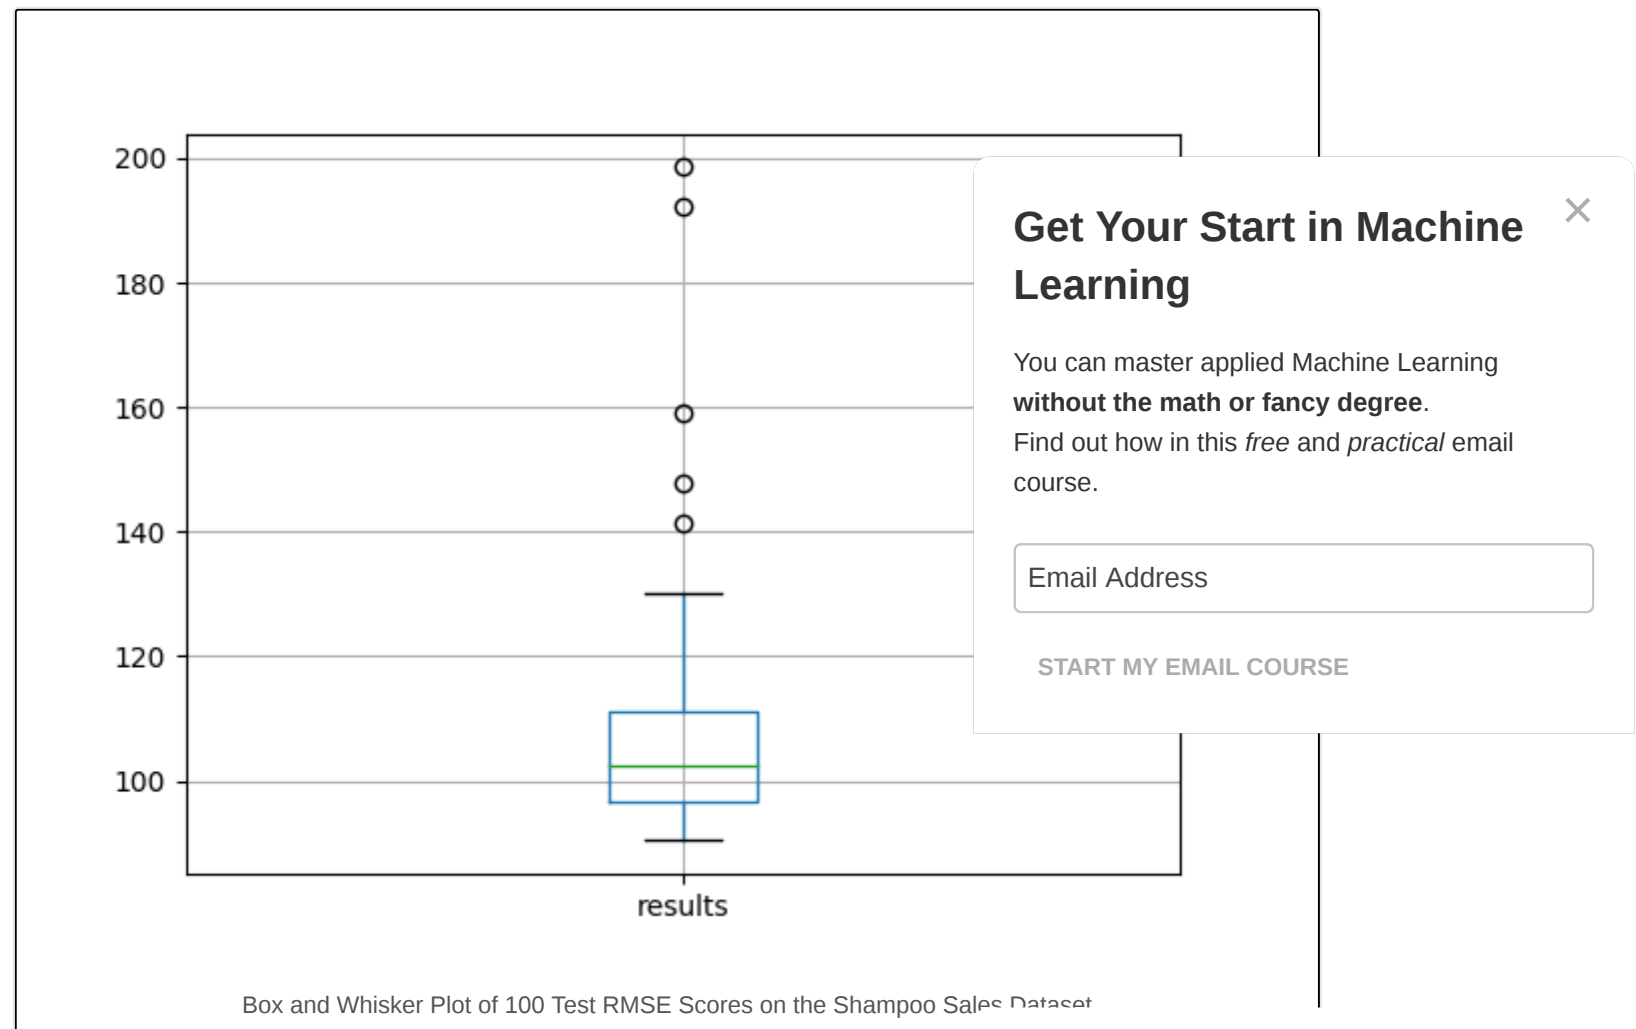
START MY EMAIL COURSE

Get Your Start in Machine Learning

6	25%	96.630800
7	50%	102.603908
8	75%	111.199574
9	max	198.568771

To get a better idea of the spread of the data, a box and whisker plot is also created.

The plot shows the median (green line), middle 50% of the data (box), and outliers (dots). We can see quite a spread to the data towards poor RMSE scores.

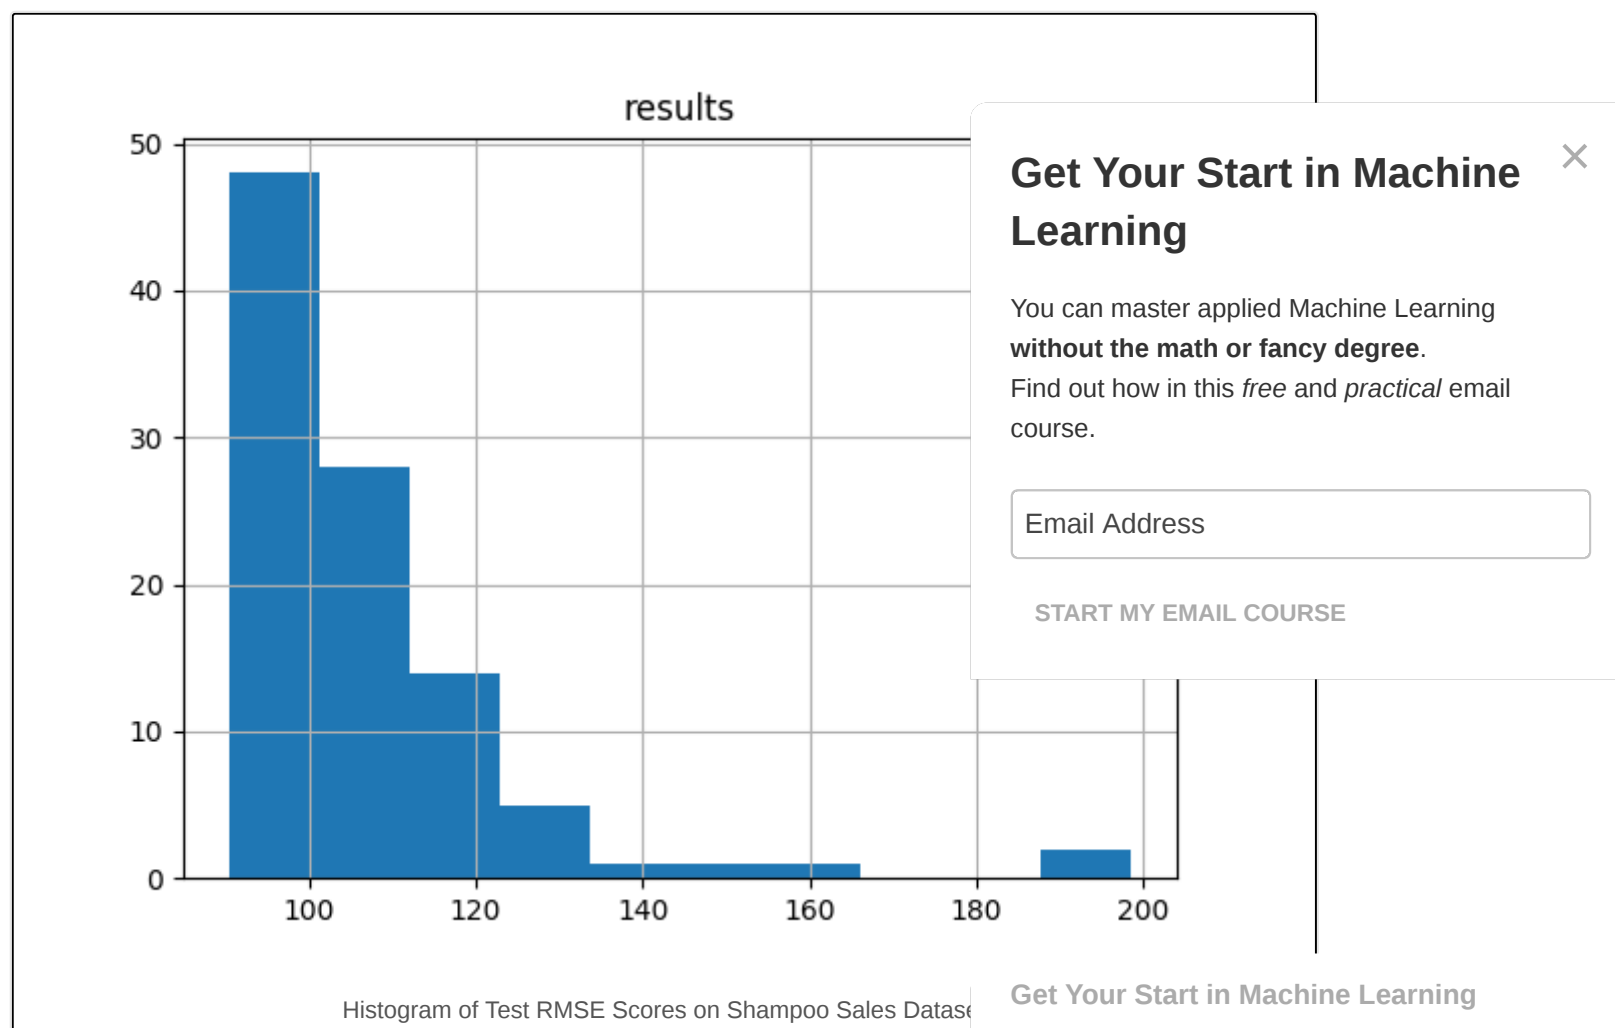


A histogram of the raw result values is also created.

The plot suggests a skewed or even an exponential distribution with a mass around an RMSE of 100 and a long tail leading out towards an RMSE of 200.

The distribution of the results are clearly not Gaussian. This is unfortunate, as the mean and standard deviation cannot be used directly to estimate a confidence interval for the model (e.g. [95% confidence as 2x the standard deviation around the mean](#)).

The skewed distribution also highlights that the median (50th percentile) would be a better central tendency to use instead of the mean for these results. The median should be more robust to outlier results than the mean.



## Repeats vs Test RMSE

We can start to look at how the summary statistics for the experiment change as the number of repeats is increased from 1 to 100.

We can accumulate the test RMSE scores and calculate descriptive statistics. For example, the score from one repeat, the scores from the first and second repeats, the scores from the first 3 repeats, and so on to 100 repeats.

We can review how the central tendency changes as the number of repeats is increased as a line plot. We'll look at both the mean and median.

Generally, we would expect that as the number of repeats of the experiment is increased, the distribution would increasingly better match the underlying distribution, including the central tendency, such as the mean.

The complete code listing is provided below.

```
1 from pandas import DataFrame
2 from pandas import read_csv
3 from numpy import median
4 from numpy import mean
5 from matplotlib import pyplot
6 import numpy
7 # load results file
8 results = read_csv('experiment_stateful.csv', header=0)
9 values = results.values
10 # collect cumulative stats
11 medians, means = list(), list()
12 for i in range(1, len(values)+1):
13     data = values[0:i, 0]
14     mean_rmse, median_rmse = mean(data), median(data)
15     means.append(mean_rmse)
16     medians.append(median_rmse)
17     print(i, mean_rmse, median_rmse)
18 # line plot of cumulative values
19 line1, = pyplot.plot(medians, label='Median RMSE')
20 line2, = pyplot.plot(means, label='Mean RMSE')
21 pyplot.legend(handles=[line1, line2])
22 pyplot.show()
```

The cumulative size of the distribution, mean, and median is printed as the number of repeats is increased. A truncated output is listed below.

1 ...

### Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

2	90	105.759546832	101.477640071
3	91	105.876449555	102.384620485
4	92	105.867422653	102.458057114
5	93	105.735281239	102.384620485
6	94	105.982491033	102.458057114
7	95	105.888245347	102.384620485
8	96	106.853667494	102.458057114
9	97	106.918018205	102.531493742
10	98	106.825398399	102.458057114
11	99	107.004981637	102.531493742
12	100	107.051145721	102.603907965

A line plot is also created showing how the mean and median change as the number of repeats is increased.

The results show that the mean is more influenced by outlier results than the median, as expected.

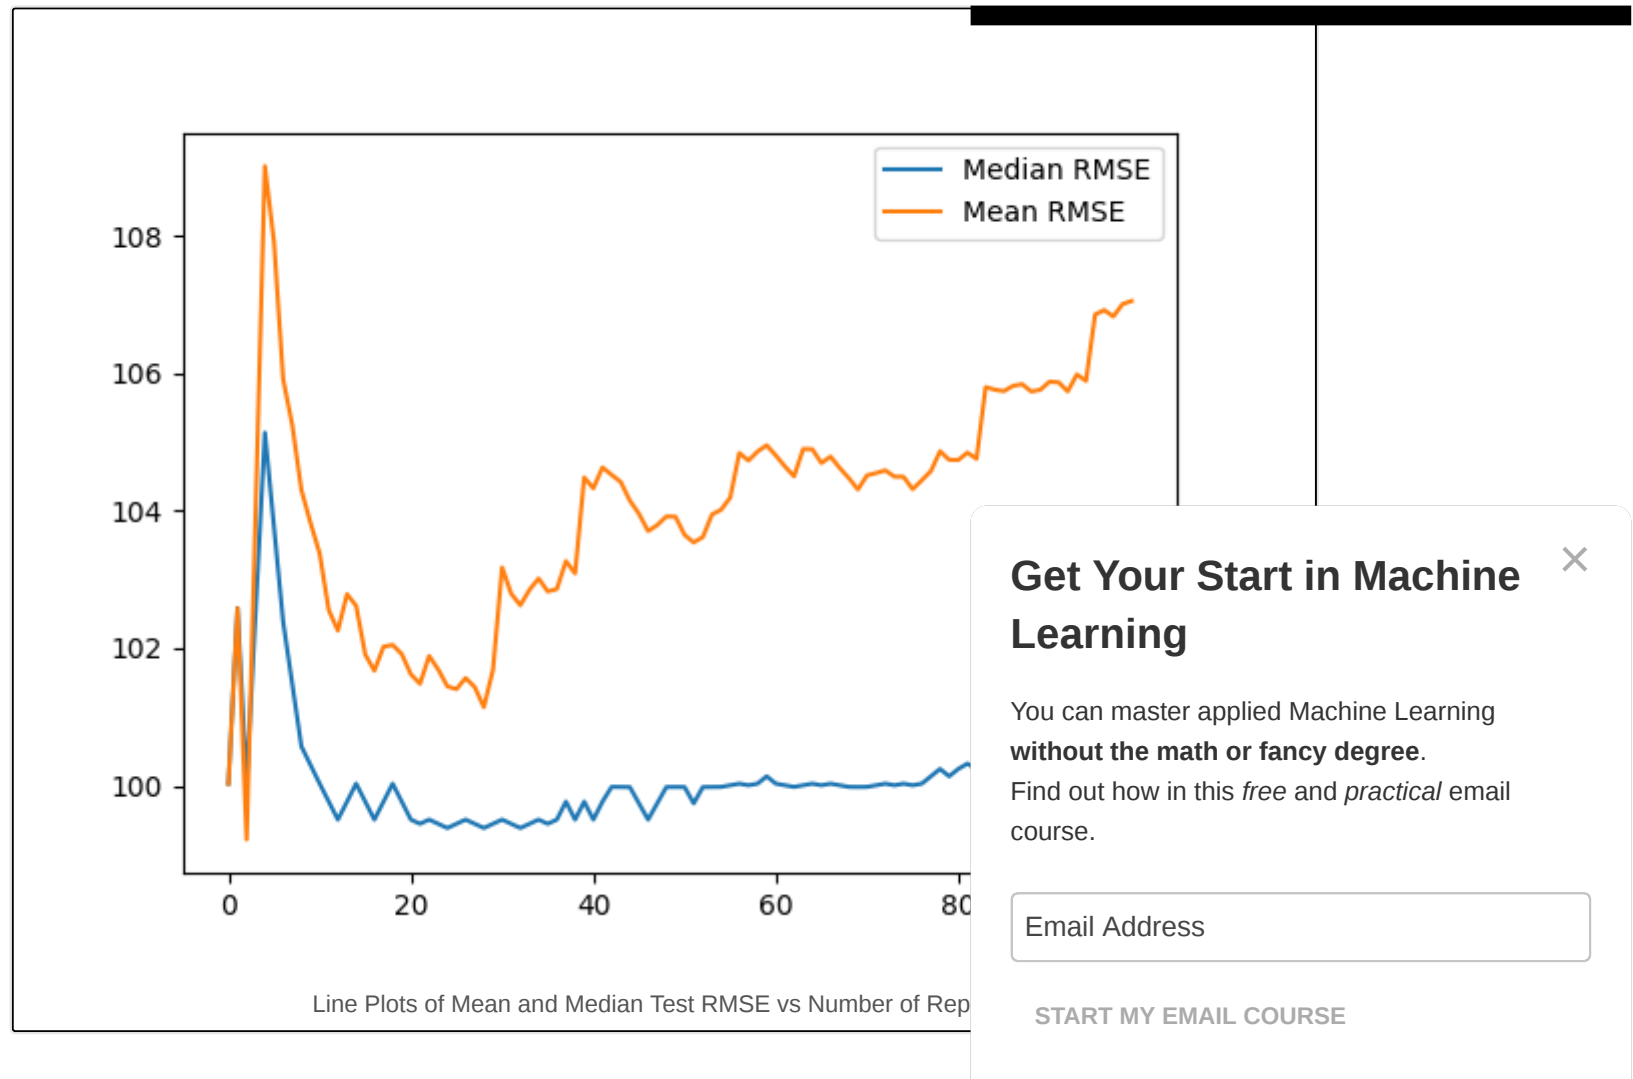
We can see that the median appears quite stable down around 99-100. This jumps to 102 towards the scores at later repeats.

## Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE





## Review of Results

We made some useful observations from 100 repeats of a stateful LSTM on a standard time series forecasting problem.

Specifically:

- We observed that the distribution of results is not Gaussian. It may be a skewed Gaussian or an exponential distribution with a long tail and outliers.
- We observed that the distribution of results did not stabilize with the increase of repeats from 1

Get Your Start in Machine Learning

The observations suggest a few important properties:

- The choice of online learning for the LSTM and problem results in a relatively unstable model.
- The chosen number of repeats (100) may not be sufficient to characterize the behavior of the model.

This is a useful finding as it would be a mistake to make strong conclusions about the model from 100 or fewer repeats of the experiment.

This is an important caution to consider when describing your own machine learning results.

This suggests some extensions to this experiment, such as:

- Explore the impact of the number of repeats on a more stable model, such as one using batch or mini-batch learning.
- Increase the number of repeats to thousands or more in an attempt to account for the general instability of the model with online learning.

## Summary

In this tutorial, you discovered how to analyze experimental results from LSTM models fit using online learning.

You learned:

- How to design a robust test harness for evaluating LSTM models on time series forecast problems.
- How to analyze experimental results, including summary statistics.
- How to analyze the impact of increasing the number of experiment repeats and how to identify a stable model.

Do you have any questions?

Ask your questions in the comments below and I will do my best to answer.

### Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree**. Find out how in this *free* and *practical* email course.


START MY EMAIL COURSE

## Develop LSTMs for Sequence Prediction Today!

Develop Your Own LSTM models in Minutes

...with just a few lines of python code

Get Your Start in Machine Learning



Discover how in my new EBOOK.

### Long Short-Term Memory Networks with Python

It provides **self-study tutorials** on topics like:

*CNN LSTMs, Encoder-Decoder LSTMs, generative models, data preparation, making predictions* and much more...

## Finally Bring LSTM Recurrent Neural Networks to Your Sequence Predictions Projects

Skip the Academics. Just Results.

[Click to learn more.](#)

### Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

**START MY EMAIL COURSE**



#### About Jason Brownlee

Dr. Jason Brownlee is a husband, proud father, academic researcher, author, professional developer, and helps developers get started and get good at applied machine learning. [Learn more.](#)

[View all posts by Jason Brownlee](#) →

[◀ Stateful and Stateless LSTM for Time Series Forecasting with Python](#)

[Exploratory Configuration of a Multilayer Perceptron Network for Time Series Forecasting ▶](#)

[Get Your Start in Machine Learning](#)

## 2 Responses to *Instability of Online Learning for Stateful LSTM for Time Series Forecasting*



**Cao** May 31, 2017 at 8:48 pm #

REPLY ↩

Dear Jason,

always thank you for your posts, which are really helpful.

i am a beginner of machine learning, and i stuck in some questions for a longtime.

1.what is the online learning? batch size=1 is online and dynamic?

2.how to set a lstm structure, which can dynamic add new data every minutes, and update the weight and bias for next minute data.

thank you again, jason.



**Jason Brownlee** June 2, 2017 at 12:47 pm #

Online learning means the model is updated after each training pattern.

The structure does not change. You need to find a structure that achieves good results on your problem.

### Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Leave a Reply

Get Your Start in Machine Learning

Name (required)

Email (will not be published) (required)

Website

## Welcome to Machine Learning Mastery



Hi, I'm Dr. Jason Brownlee.  
My goal is to make practitioners like YOU awesome at applied machine learning.

[Read More](#)

## Get Your Start in Machine Learning ×

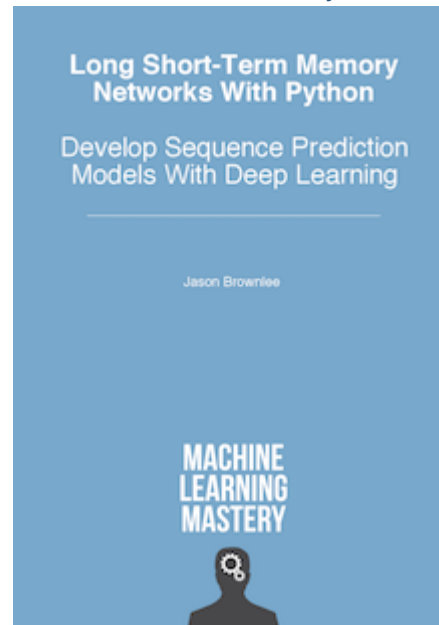
You can master applied Machine Learning **without the math or fancy degree.**  
Find out how in this *free* and *practical* email course.

## Deep Learning for Sequence Prediction

Cut through the math and research papers.  
Discover 4 Models, 6 Architectures, and 14 Tutorials.

[Get Your Start in Machine Learning](#)

Get Started With LSTMs in Python Today!



## POPULAR

**Time Series Prediction with LSTM Recurrent Neural Networks in Python with Keras**

JULY 21, 2016

**Your First Machine Learning Project in Python Step-By-Step**

JUNE 10, 2016

**Develop Your First Neural Network in Python With Keras Step-By-Step**

MAY 24, 2016

**Sequence Classification with LSTM Recurrent Neural Networks in Python with Keras**

JULY 26, 2016

**How to Setup a Python Environment for Machine Learning and Deep Learning with Anaconda**

MARCH 13, 2017

## Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

[START MY EMAIL COURSE](#)[Get Your Start in Machine Learning](#)



Time Series Forecasting with the Long Short-Term Memory Network in Python

APRIL 7, 2017




Multi-Class Classification Tutorial with the Keras Deep Learning Library

JUNE 2, 2016




Regression Tutorial with the Keras Deep Learning Library in Python

JUNE 9, 2016



Multivariate Time Series Forecasting with LSTMs in Keras

AUGUST 14, 2017



How to Implement the Backpropagation Algorithm From Scratch In Python

NOVEMBER 7, 2016

## Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

© 2017 Machine Learning Mastery. All Rights Reserved.

[Privacy](#) | [Contact](#) | [About](#)

Get Your Start in Machine Learning