

Home  
About  
Research  
FAST  
Blog posts



# Erik Smistad

FOLLOW:



OPENCL

149

MORE

## RECENT COMMENTS



TIRTHESH SAYS:  
correct the spelling of SquareRoot in tests.cpp code.



CAT&ANT SAYS:  
Hello Erik, I don't find the gmock after...



CAT&ANT SAYS:  
Great. Thanks a lot for this helpful info!!



ERIK SMISTAD SAYS:  
The image sampler in OpenCL handles the edges (the variable sampler...



ROYI SAYS:  
Hi, Great post. I would like to...

## POPULAR POSTS



# OpenCL

OPENCL

Getting started with OpenCL and GPU Computing  
21 JUN, 2010



## Getting started with OpenCL and GPU Computing

BY [ERIK SMISTAD](#) · PUBLISHED JUNE 21, 2010 ·  
UPDATED MAY 28, 2016

OpenCL (Open Computing Language) is a new framework for writing programs that execute in parallel on different compute devices (such as CPUs and GPUs) from different vendors (AMD, Intel, ATI, Nvidia etc.). The framework defines a language to write "kernels" in. These kernels are the functions which are to run on the different compute devices. In this post I explain how to get started with OpenCL and how to make a small OpenCL program that will compute the sum of two lists in parallel.



# OpenCL

## Installing and setting up OpenCL on your computer

First of all you need to download the newest drivers to your graphics card. This is



Erik Smistad Retweeted



**NTNUmedisin og helse**  
@NTNUhelse

Møt blant annet Svein Erik Måsøy fra #CIUS som skal fortelle om det norske ultralydeventyret [forskningdagene.no/artikler/verdi...](https://forskningdagene.no/artikler/verdi...) @Forskningdag @NTNU

16h

Erik Smistad Retweeted



**Dan Hulme**  
@nasaldemons

Replying to @caffe2ai @bobpoekert  
Why are you competing with NNEF rather than supporting it?


Sep 8, 2017

Erik Smistad Retweeted



**P Cantillon-Murphy**  
@cantillonmurphy

Thank you Thomas and team @SINTEF Trondheim for hosting & sharing your work in airway navigation. Looking forward to future collaboration!



## GITHUB PROJECTS

- 3D-Gradient-Vector-Flow-for-Matlab
- caffe
- cpd
- CustusX
- Detect-Screen-GPU-With-OpenCL




IMAGE PROCESSING & GFX / OPENCL / OPENG

Marching Cubes implementation using OpenCL and OpenGL

7 OCT, 2011



GENERAL

Getting started with Google Test (GTest) on Ubuntu

5 JUL, 2012



IMAGE PROCESSING & GFX / OPENCL

GPU-based Gradient Vector Flow using OpenCL

23 JUL, 2012

important because OpenCL will not work if you don't have drivers that support OpenCL.

To install OpenCL you need to download an implementation of OpenCL. The major graphic vendors Nvidia and AMD/ATI have both released implementations of OpenCL for their GPUs. These implementation come in a so called software development kits and often include some useful tools such as a visual profiler. The next step is to download and install the SDK for the GPU you have on your computer. Note that not all graphic cards are supported. A list of which graphic cards are supported can be found on the vendors websites.

For AMD/ATI GPUs download the [AMD APP SDK \(formerly known as AMD Stream SDK\)](#) For Nvidia GPUs download the [CUDA Toolkit](#)

The installation steps differ for each SDK and the OS you are running. Follow the installation manual of the SDK carefully. Personally I use Ubuntu Linux and have an AMD 7970 graphics card. Below are some installation steps for this specific setup.

## Installing OpenCL on Ubuntu Linux with AMD graphics card

To install the latest AMD drivers on Ubuntu 12.04 open additional drivers and install/active the one called "ATI/AMD proprietary FGLRX graphic driver (post-release updates)".

After that is done, restart and download and extract the [AMD APP SDK](#).

AMD APP SDK 2.8 includes an installer. Run this with the command:

```
sudo sh Install-AMD-APP.sh
```

Next, install the OpenCL headers files

```
sudo apt-get install opencl-headers
```

And your done! Note that the AMD APP SDK and its samples is located at /opt/AMDAPP.

- [django-tinymce-gallery](#)
- [FAST](#)
- [FAST-example-project](#)
- [FLTK-OpenGL-OpenCL-Interoperability](#)
- [freenect2-test](#)
- [GPU-Marching-Cubes](#)
- [GPU-Multigrid-Gradient-Vector-Flow](#)
- [Gradient-Vector-Flow](#)
- [GTest](#)
- [keras](#)
- [Level-Set-Segmentation](#)
- [libfreenect2](#)
- [Memory-mapped-file](#)
- [midas-journal-111](#)
- [Multiple-Windows-Single-OpenGL-Context](#)
- [Nerve-Segmentation](#)
- [node-webchat](#)
- [OpenCL-Gaussian-Blur](#)
- [OpenCL-Getting-Started](#)
- [OpenCL-GVF](#)
- [OpenCL-Level-Set-Segmentation](#)
- [OpenCLUtilities](#)
- [OpenCLUtilityLibrary](#)
- [opengles-book-samples](#)
- [Qt-Context-Sharing](#)

## Installing OpenCL on Ubuntu Linux with NVIDIA graphics card

Download the CUDA toolkit for Ubuntu from [NVIDIA's CUDA site](#). Open a terminal and run the installation file with the command:

```
sudo sh cuda-toolkit_3.1_linux_64_ubuntu12.04.deb
```

Download the Developer Drivers for Linux at the same website and install it by first stopping X, running the file and start X again. To stop X use:

```
sudo /etc/init.d/gdm stop
```

Then get a terminal up by pressing CTRL+ALT+F5, login and navigate to where you downloaded the devdriver then type:

```
sudo sh dev-driver_3.1_linux_64_256.40.deb
```

After the driver has been installed start x again by typing

```
startx
```

Before compiling an OpenCL application you need to add the path to the lib folder of CUDA to LD\_LIBRARY\_PATH like so:

```
export LD_LIBRARY_PATH=/usr/local/cuda/lib64
```

## Your first OpenCL program – Vector addition

To demonstrate OpenCL I explain how to perform the simple task of vector addition. Suppose we have two lists of numbers, A and B, of equal size. The task of vector addition is to add the elements of A with the elements of B and put the result in the element of a new list called C of the same size. The figure below explains the operation.

A	3	6	2	0	-2	...
+						
B	2	3	1	1	2	...
=						
C	5	9	3	1	0	...

Two lists A and B and the result list C of vector addition on A and B

The naive way of performing this operation is to simply loop through the list and perform the operation on one element at a time like the C++ code below:

```
for(int i = 0; i < LIST_SIZE; i++) {
    C[i] = A[i] + B[i];
}
```

This algorithm is simple but has a linear time complexity,  $O(n)$  where  $n$  is the size of the list. But since each iteration of this loop is independent on the other iterations this operation is data parallel, meaning that each iteration can be computed simultaneously. So if we have  $n$  cores on a processor this operation can be performed in constant time  $O(1)$ .

To make OpenCL perform this operation in parallel we need to make the kernel. The kernel is the function which will run on the compute device.

## The kernel

The kernel is written in the OpenCL language which is a subset of C and has a lot of math and vector functions included. The kernel to perform the vector addition operation is defined below.

```
__kernel void vector_add(__global const
    // Get the index of the current element
    int i = get_global_id(0);

    // Do the operation
    C[i] = A[i] + B[i];
}
```

## The host program

The host program controls the execution of kernels on the compute devices. The host program is written in C, but bindings for other languages like C++ and Python exists. The OpenCL API is defined in the `cl.h` (or `opencl.h` for apple) header file. Below is the code for the host program that executes the kernel above on compute device. I will not go into details on each step as this is supposed to be an introductory article although I can recommend the book [“The OpenCL Programming Book”](#) if you want to dive into the details. The main steps of a host program is as follows:

- Get information about the platform and the devices available on the computer (line 42)
- Select devices to use in execution (line 43)
- Create an OpenCL context (line 47)
- Create a command queue (line 50)
- Create memory buffer objects (line 53-58)
- Transfer data (list A and B) to memory buffers on the device (line 61-64)
- Create program object (line 67)
- Load the kernel source code (line 24-35) and compile it (line 71) (online execution) or load the precompiled binary OpenCL program (offline execution)
- Create kernel object (line 74)
- Set kernel arguments (line 77-79)
- Execute the kernel (line 84)
- Read memory objects. In this case we read the list C from the compute device (line 88-90)

```
1 #include <stdio.h>
2 #include <stdlib.h>
3
4 #ifdef __APPLE__
5 #include <OpenCL/opencl.h>
6 #else
7 #include <CL/cl.h>
8 #endif
9
10 #define MAX_SOURCE_SIZE (0x100000)
11
12 int main(void) {
13     // Create the two input vectors
14     int i;
15     const int LIST_SIZE = 1024;
```

```

16  int *A = (int*)malloc(sizeof(int)*LIST_SIZE);
17  int *B = (int*)malloc(sizeof(int)*LIST_SIZE);
18  for(i = 0; i < LIST_SIZE; i++)
19      A[i] = i;
20      B[i] = LIST_SIZE - i;
21  }
22
23  // Load the kernel source code
24  FILE *fp;
25  char *source_str;
26  size_t source_size;
27
28  fp = fopen("vector_add_kernel.cl", "r");
29  if (!fp) {
30      fprintf(stderr, "Failed to open kernel source file\n");
31      exit(1);
32  }
33  source_str = (char*)malloc(MAX_SOURCE_SIZE);
34  source_size = fread(source_str, sizeof(char), MAX_SOURCE_SIZE, fp);
35  fclose(fp);
36
37  // Get platform and device info
38  cl_platform_id platform_id = NULL;
39  cl_device_id device_id = NULL;
40  cl_uint ret_num_devices;
41  cl_uint ret_num_platforms;
42  cl_int ret = clGetPlatformIDs(1, &platform_id, &ret_num_platforms);
43  ret = clGetDeviceIDs(platform_id, CL_DEVICE_TYPE_DEFAULT, 1, &device_id, &ret_num_devices);
44
45  // Create an OpenCL context
46  cl_context context = clCreateContext(NULL, 1, &device_id, NULL, NULL, NULL);
47
48  // Create a command queue
49  cl_command_queue command_queue = clCreateCommandQueue(context, device_id, 0, NULL);
50
51  // Create memory buffers on the device
52  cl_mem a_mem_obj = clCreateBuffer(context, CL_MEM_READ_WRITE, LIST_SIZE * sizeof(int), NULL, NULL);
53  cl_mem b_mem_obj = clCreateBuffer(context, CL_MEM_READ_WRITE, LIST_SIZE * sizeof(int), NULL, NULL);
54  cl_mem c_mem_obj = clCreateBuffer(context, CL_MEM_READ_WRITE, LIST_SIZE * sizeof(int), NULL, NULL);
55
56  // Copy the lists A and B to the device
57  ret = clEnqueueWriteBuffer(command_queue, a_mem_obj, CL_TRUE, 0, LIST_SIZE * sizeof(int), A, 0, 0, 0);
58  ret = clEnqueueWriteBuffer(command_queue, b_mem_obj, CL_TRUE, 0, LIST_SIZE * sizeof(int), B, 0, 0, 0);
59
60  // Create a program from the kernel source
61  cl_program program = clCreateProgramWithSource(context, 1, (const char **)&source_str, &source_size, NULL, NULL);
62
63  // Build the program
64  ret = clBuildProgram(program, 1, &device_id, NULL, NULL, NULL);
65
66  // Create the OpenCL kernel
67  cl_kernel kernel = clCreateKernel(program, "vector_add", &ret);
68
69  // Set the arguments of the kernel
70  ret = clSetKernelArg(kernel, 0, sizeof(cl_mem), &a_mem_obj);
71  ret = clSetKernelArg(kernel, 1, sizeof(cl_mem), &b_mem_obj);
72  ret = clSetKernelArg(kernel, 2, sizeof(cl_mem), &c_mem_obj);
73
74  // Execute the OpenCL kernel on the device
75  size_t global_item_size = LIST_SIZE;
76  size_t local_item_size = 64;
77  ret = clEnqueueNDRangeKernel(command_queue, 1, &global_item_size, &local_item_size, NULL, NULL, 0, 0, 0);
78
79  // Read the memory buffer C on the host
80  int *C = (int*)malloc(sizeof(int)*LIST_SIZE);
81  ret = clEnqueueReadBuffer(command_queue, c_mem_obj, CL_TRUE, 0, LIST_SIZE * sizeof(int), C, 0, 0, 0);
82
83  // Display the result to the screen
84  for(i = 0; i < LIST_SIZE; i++)
85      printf("%d + %d = %d\n", A[i], B[i], C[i]);
86
87  // Clean up
88  ret = clFlush(command_queue);
89  ret = clFinish(command_queue);
90  ret = clReleaseKernel(kernel);
91  ret = clReleaseCommandQueue(command_queue);
92  ret = clReleaseContext(context);
93

```

```
100     ret = clReleaseProgram(program);
101     ret = clReleaseMemObject(a_mem);
102     ret = clReleaseMemObject(b_mem);
103     ret = clReleaseMemObject(c_mem);
104     ret = clReleaseCommandQueue(command_queue);
105     ret = clReleaseContext(context);
106     free(A);
107     free(B);
108     free(C);
109     return 0;
110 }
```

To make OpenCL run the kernel on the GPU you can change the constant `CL_DEVICE_TYPE_DEFAULT` to `CL_DEVICE_TYPE_GPU` in line 43. To run on CPU you can set it to `CL_DEVICE_TYPE_CPU`. This shows how easy OpenCL makes it to run different programs on different compute devices.

[The source code for this example can be downloaded here.](#)

## Compiling an OpenCL program

If the OpenCL header and library files are located in their proper folders (`/usr/include` and `/usr/lib`) the following command will compile the `vectorAddition` program.

```
gcc main.c -o vectorAddition -l OpenCL
```


## How to learn more


To learn more about OpenCL I recommend the book from Fixstars called [The OpenCL programming book](#). Below are some links to useful sites with information on OpenCL:

- [Nvidia's page on OpenCL](#)
- [AMD/ATI's page on OpenCL](#)
- [Official website of OpenCL](#)

Tags: GPGPU GPU computing OpenCL

### 👍 YOU MAY ALSO LIKE...





Level set  
segmentati  
on on GPUs  
using  
OpenCL

Real-time  
medical  
surface  
extraction  
with FAST  
DECEMBER 5,  
2015

## OpenCL

OpenCL C++  
Utilities  
SEPTEMBER  
16, 2011

JULY 11, 2013

### 149 RESPONSES

Comments 139 Pingbacks 10



**Gabriel** August 20, 2016 at 00:51

Thanks guy xD

Reply



**shilpa** March 12, 2015 at 13:08

Hi Erik,

Your program is pretty good to understand . I ran this program on a octacore machine but could not see the expected parallel processing of instructions.

The time taken for this code was more than a sequential code for vector addition. I doubt , this program is not running on multiple cores parallely on my system. Could you plz help ?

Regards,

Shilpa

Reply



**Erik Smistad**

March 23, 2015 at 17:18

This is probably because the vectors are so small (only 1024 items). Try to increase the size of the vector to lets say 1024\*1024\*1024, and you will probably see a speedup.

Reply



**celio**

August 31, 2016 at 10:01

i changed item\_size to 1024\*1024\*32 but i did not defferene between GPU &CPU

Reply



**KSSR** November 28, 2014 at 16:51

hello all,

I need instruction about setting up openc1 environemnt for Multicore system i.e on GPU.

Reply



**kevinkit** July 21, 2014 at 15:28

Sry I was wrong with the cl\_mem stuff but nevertheless the "+1" is missing. (THIS I NOT C++!)



Reply



**kevinkit** · July 20, 2014 at 20:01

Doesn't it has to be

```
int *A = (int*)malloc(sizeof(int)*(LIST_SIZE
+ 1));
```

and this in every other memory allocation  
furthermore when you allocate the memory  
objects it should be

```
cl_mem a_mem_obj =
clCreateBuffer(context,
CL_MEM_READ_ONLY,
(LIST_SIZE+1) * sizeof(cl_int), NULL, &ret);
```

instead of cl\_int.

Reply



**Fabio** · July 17, 2014 at 11:11

Hi, I have a question.

When you call `clEnqueueReadBuffer`, you  
don't have to pass as parameter also the list  
of events to wait before read the buffer? I  
mean, you can get the event identifier from  
`clEnqueueNDRangeKernel` and pass it to  
`clEnqueueReadBuffer`, otherwise the program  
may read the results before the sum is  
completed.

If it's not needed, why?

Reply



**Erik Smistad** · July 17, 2014 at 13:31



The third argument of  
`clEnqueueReadBuffer` with the value  
`CL_TRUE` ensures that this call is  
blocking. That means that the  
function will not return until it has  
read the buffer and thus explicit  
synchronization is not needed.  
However, if you set this argument to  
`CL_FALSE` you have to do explicit  
synchronization using events as you  
suggest.

Reply



**Fabio** · July 17, 2014 at 14:27

I know that the third  
argument make the call  
blocking, but I think that  
means that (as you say) "the  
function will not return until it  
has read the buffer". However  
the documentation doesn't  
say anything about the  
previous enqueued  
operations for this argument.  
Maybe is not so clear.  
The documentation says also  
that you must ensure that  
"All commands that use this  
buffer object have finished  
execution before the read  
command begins execution"

Reply

**Erik Smistad**

🕒 July 17, 2014 at 15:18

Ah yes, that is a good point. The clue here is that your command queue is created with in-order execution (this is default and most devices doesn't support out-of-order execution). In-order execution guarantees that all of the queued commands will be executed in the order in which they were added to the queue. Thus, for `clEnqueueReadBuffer` to finish, all other queued operations have to finish first.

See

<http://www.khronos.org/registry/cl/sdk/1.0/docs/man/xhtml/clCreateCommandQueue.html>  
for more info on this

Reply

**Fabio**

🕒 July 17, 2014 at 15:33

Aaahh sorry 😊

Your queue is not out-of-order.

I'm working with an aout-of-order queue and I have some problems so I'm trying to understand...

Thanks

Reply

**Anonymous** 🕒 April 9, 2014 at 10:15

Have you found out how to fix the failure?

*Name (required):*

I added a `printf()` to your code after line 43:  
`if (ret != CL_SUCCESS) {printf("Error: Failed to query platforms! (%d)\n", ret);return EXIT_FAILURE;}`  
after compiling, running it gives me this error:  
"Failed to query platforms (-1001)"

Reply

**meenu** • February 13, 2014 at 10:03

I am currently using ubuntu13.04 and have a  
VGA compatible controller: NVIDIA  
Corporation GK107 [GeForce GT 630 OEM]  
(rev a1) ... My open CL samples are running  
fine for CL\_DEVICE\_TYPE\_DEFAULT and  
CL\_DEVICE\_TYPE\_CPU... But they are not  
able to find OPENCL devices for  
CL\_DEVICE\_TYPE\_GPU...

Reply

**Erik Smistad**

★ • February 13, 2014 at 12:28

The NVIDIA OpenCL platform only  
support GPUs. So most likely you  
have more than one platform  
installed (or the NVIDIA platform is  
not installed). Try to select the correct  
platform. You can do this by  
increasing the number of entries in  
the clGetPlatformIDs function, see  
[https://www.khronos.org/registry  
/cl/sdk/1.1/docs/man/xhtml  
/clGetPlatformIDs.html](https://www.khronos.org/registry/cl/sdk/1.1/docs/man/xhtml/clGetPlatformIDs.html)

Reply

**Anonymous**

• February 18, 2014 at 06:17

Okie i will tell u what exactly  
my configuration is and what  
i have observed then  
probably u might help me  
out... My CPU is intel core  
processor -i7 and i have an  
inbuilt gpu of nvidia. Now i  
have installed both intel sdk  
for opencl and nvidia  
proprietary drivers for  
ubuntu13.04. How can i  
create a run time so that it  
identifies both CPU and GPU.  
Currently i feel that only intel  
platform is getting  
recognised ... hence opencl is  
working fine for option CPU  
and not GPU. Is there a way  
around where both my  
devices are identified and  
probably i can transfer data  
between my CPU and GPU.  
Also in my vendor directory i  
can observe both  
intelocl64.icd and nvidia.icd.

Reply

**Erik Smistad**★ • February 18, 2014  
at 13:47

An OpenCL context  
can only be  
associated with ONE

platform. You have  
TWO platforms  
installed and the  
code above only  
selects ONE  
platform, whichever  
one comes first, in  
your case the intel  
platform. To select  
the NVIDIA platform  
you need to increase  
the number of  
entries:

```
cl_platform_id  
platform_ids[2];  
clGetPlatformIDs(2,  
platform_ids,  
&ret_num_platforms);
```

and then select the  
other platform like  
this:

```
clGetDeviceIDs(  
platform_ids[1],  
CL_DEVICE_TYPE_DEFAULT,  
1, &device_id,  
&ret_num_devices);
```

Reply



**meenu**

🕒 February  
19, 2014 at  
09:46

Thanks Eric  
!!! This was  
the  
solution...  
Thanks  
again...

Reply



**Gustavo Rozolin da Silva**

🕒 October 12, 2013 at 21:25

Excellent post Erik,

Erik what I need to change in this code for to  
pass \_\_local arg to kernel.

Thanks you.

Reply



**Erik Smistad**



🕒 October 14, 2013 at 16:23

Local memory, or shared memory as  
it is also called, is not accessible from  
the host. So if you want to use it you  
have to transfer data from global to  
local explicitly in a kernel.

Reply



**Laxator2** · October 12, 2013 at 19:19

Here is how it worked on my machine, using a (rather old) Nvidia card under PCLinuxOS :

```
gcc main.c -I/usr/local/cuda-5.0/include/  
-L/usr/lib/nvidia-current -lOpenCL -o  
vectorAddition
```

Great example, short and to the point.

Reply



**Kareti** · July 31, 2013 at 06:37

Hi Erik, The blog was very helpful. Thanks for that.

I have no GPU on my laptop, so is there a way to practise the opencl programs by emulating a GPU! I am using Fedora 18, 64 bit !

Thank you.

Reply



**Erik Smistad** · July 31, 2013 at 11:11

★ Yes, you can use your CPU instead.

That's the nice thing about OpenCL:

The code can run on different types of processors. To do so simply install the Intel or AMD OpenCL runtime, depending on which type of processor you have. Afterwards execute the example above and it should run on the CPU.

Reply



**Victor** · July 18, 2013 at 11:58

Hi Erik. Could u tell me what does this mean?

```
gcc -c -I /usr/include/CL main.c -o main.o  
gcc -o local16 main.o -L /usr/lib -lOpenCL  
/usr/lib64/gcc/x86_64-suse-linux/4.7/..  
/usr/lib64/gcc/x86_64-suse-linux/bin/ld: skipping  
incompatible /usr/lib/libOpenCL.so when  
searching for -lOpenCL  
/usr/lib64/gcc/x86_64-suse-linux/4.7/..  
/usr/lib64/gcc/x86_64-suse-linux/bin/ld: skipping  
incompatible /usr/lib/libc.so when searching  
for -lc
```

I dont know if I successfully link to the lib.  
What's weird is that the result keeps the same even if I comment the whole kernel code.

Reply



**Erik Smistad** · July 31, 2013 at 11:13

★ This is a problem with the linking, not the source code. Not sure what the problem is, never seen that error message before

Reply



**Xocoatzin** · June 6, 2013 at 20:42

Wow, it just worked. Awesome!

Reply



**Anonymous** · May 12, 2013 at 18:04

Very nice article, thanks!

I'm wondering whether the statements:

```
ret = clFlush(command_queue);
```

```
ret = clFinish(command_queue);
```

are actually needed. If I'm getting it right, since the command queue is in-order, when `clEnqueueReadBuffer` (which is blocking thanks to the `CL_TRUE` parameter) returns, the command queue should be empty.

Another point that would be worth explaining is why there is no `clFinish` between `clEnqueueNDRangeKernel` and `clEnqueueReadBuffer`.

Reply



**Erik Smistad** · May 13, 2013 at 10:44



Since the program terminates right after all the clean up statements, none of them are actually needed.

In either case, you are correct, the flush and finish statements are not necessary.

When the command queue is in-order, OpenCL will make sure that all the enqueue commands are performed in the order in which they are called in your C program.

Reply



**Rohit Vashistha** · April 25, 2013 at 18:51

Hi

All those getting 'zero' result change the option "GPU" to "CPU" or vice versa

Regards,

Rohit

Reply



**Jai** · April 22, 2013 at 15:03

I have a desktop pc with configuration as:  
Intel(R) Core(TM) 2 Duo CPU [E700@2.93GHz](#),  
2GB RAM 32 bit OS

and I am willing to purchase a graphic card with config as:

Sapphire AMD/ATI RADEON HD 6450 2GB  
RAM

Can you please tell me is it compatible for my pc...?

Thanks in advance.

Reply



**Jack** · April 8, 2013 at 13:39

Can i run openCL on a CPU? (I do not have a GPU but want to experiment with openCL before buying one).

I have an intel i3 processor with a dual boot for Windows 7 and Ubuntu 12.04

Reply





**Erik Smistad** · April 8, 2013 at 13:56  
Sure. Just install the Intel OpenCL SDK  
from <http://software.intel.com/en-us/vcsource/tools/opencv-sdk>

Reply



**Anonymous** · February 22, 2013 at 02:41  
The sample code all worked fine.  
But when I changed  
CL\_DEVICE\_TYPE\_DEFAULT into  
CL\_DEVICE\_TYPE\_GPU, it runs, but give me:  
0 + 1024 = 763461144  
1 + 1023 = 32716  
2 + 1022 = 763461144  
3 + 1021 = 32716  
4 + 1020 = 15489024  
5 + 1019 = 0  
6 + 1018 = 15489024  
7 + 1017 = 0  
8 + 1016 = 0  
9 + 1015 = 0  
10 + 1014 = 0  
11 + 1013 = 0  
12 + 1012 = 0  
13 + 1011 = 0  
14 + 1010 = 0  
15 + 1009 = 0  
16 + 1008 = 0  
17 + 1007 = 0  
18 + 1006 = 0  
19 + 1005 = 0  
20 + 1004 = 0  
21 + 1003 = 0  
22 + 1002 = 0  
23 + 1001 = 0  
24 + 1000 = 0  
25 + 999 = 0  
26 + 998 = 124817  
27 + 997 = 0  
28 + 996 = 1801415779  
29 + 995 = 1717531240  
30 + 994 = 540292720  
31 + 993 = 1633643619  
32 + 992 = 1717527661  
33 + 991 = 540292720  
34 + 990 = 1801415779  
35 + 989 = 1734308456  
36 + 988 = 1633841004  
37 + 987 = 1852399468  
38 + 986 = 1597125492  
39 + 985 = 1702060386  
40 + 984 = 1869898079  
41 + 983 = 1935894893  
42 + 982 = 1600938784  
43 + 981 = 1601333355  
44 + 980 = 1651469415  
45 + 979 = 1767861345  
46 + 978 = 842232942  
47 + 977 = 1954047327  
48 + 976 = 1701080677  
49 + 975 = 1952538468

50 + 974 = 1667853679

....

i tried CL\_DEVICE\_TYPE\_CPU, and it worked fine.

why is GPU not working?

Reply



**safwan** · December 29, 2012 at 21:32

Thank you for this briefly example and its work with me but I have a problem when I change the value of LIST\_SIZE to another value, the program execute but she doesn't execute the kernel fonction an finally the result i have for C[i]=0 :

0+1024=0

1+1023=0

...

..

how can I resolve this problem?

Thanks

Reply



**Erik Smistad**

★

· January 2, 2013 at 14:53

If you change the LIST\_SIZE variable to another variable that is not dividable with 64 it will not run because local size is set to 64 (see line 83). This means that 64 work-items are grouped together.

If you only get 0, you probably haven't installed OpenCL correctly.

Reply



**Yaknan** · December 28, 2012 at 06:30

Hi Erik,

I am enjoying your tutorials on OpenCL.

Thanks for the good work. Please, am testing a hypothesis here for my thesis work on enhancing graphic rendering of the ray tracing algorithm. Am wondering if it is possible to integrate cilk++ with openCL. The idea is to see if cilk++ will take full and better advantage of CPUs while OpnCL takes advantage of the GPUs more efficiently.

Thanks!

Reply



**Erik Smistad**

★

· January 2, 2013 at 14:50

As far as I know cilk++ can run regular C/C++ code as well. And OpenCL is written in C, so I think it should work..

Reply



**prince** · November 14, 2012 at 12:35

i am using gpu of type nvidia,i am using opencl but when i run the program using ctrl+F5(start without debugging )then i get result in which gpu takes more time than cpu but when i run the program cpu takes more time than gpu and result is also i am giving



```

start without debugging -> cpu time=6127
ms gpu time= 6240 ms
start with debug-> cpu time= 18354 ms gpu
time= 9125 ms

wt is the reason in this difference.....
visual studio 2010 i am using
the code is here. wt is going wrong?..thanks

// Hello.cpp : Defines the entry point for the
console application.
//

#include
#include
#include
#include
#include
#include "CL/cl.h"
#define DATA_SIZE 100000
const char *KernelSource =
"kernel void hello(global float *input , global
float *output)\n"
"{\n"
" size_t id =get_global_id(0);\n"
"output[id] =input[id]*input[id];\n"
"}"
"\n"
"\n";
//float start_time,end_time;

int main(void)
{
double start_time,end_time;
start_time=clock();
cl_context context;
cl_context_properties properties[3];
cl_kernel kernel;
cl_command_queue command_queue;
cl_program program;
cl_int err;
cl_uint num_of_platforms=0;
cl_platform_id platform_id;
cl_device_id device_id;
cl_uint num_of_devices=0;
cl_mem input,output;
size_t global;
float inputData[100000];
for(int j=0;j<100000;j++)
{
inputData[j]=(float)j;
}

float results[DATA_SIZE];//={0};

// int i;

//retrieve a list of platform variable
if(clGetPlatformIDs(1,&platform_id,&num_of_platforms)!=CL_SUCCESS)
{
printf("Unable to get platform_id\n");
return 1;
}

//try to get supported GPU Device
if(clGetDeviceIDs(platform_id,CL_DEVICE_TYPE_CPU,1,&device_id,

```

```
&num_of_devices)!=CL_SUCCESS)
{
    printf("unable to get device_id\n");
    return 1;
}

//context properties list -must be terminated
with 0
properties[0]=CL_CONTEXT_PLATFORM;
properties[1]=(cl_context_properties)
platform_id;
properties[2]=0;

//create a context with the GPU device
context=clCreateContext(properties,1,&
device_id,NULL,NULL,&err);

//create command queue using the context
and device
command_queue=clCreateCommandQueue(context,device_id,0,&
err);

//create a program from the kernel source
code
program=clCreateProgramWithSource(context,1,
(const char**)
&KernelSource,NULL,&err);

//compile the program
err=clBuildProgram(program,0,NULL,NULL,NULL);
if((err!=CL_SUCCESS))
{
    printf("build error \n",err);
    size_t len;
    char buffer[4096];
    //get the build log
    clGetProgramBuildInfo(program,device_id,CL_PROGRAM_BUILD_LOG,sizeof(buffer),buffer,&len);
    printf("——build Log——\n%s\n",buffer);
    exit(1);
}

// return 1;
}

//specify which kernel from the program to
execute
kernel=clCreateKernel(program,"hello",&err);

//create buffers for the input and output
input=clCreateBuffer(context,CL_MEM_READ_ONLY,sizeof(float)*DATA_SIZE,NULL,NULL);
output=clCreateBuffer(context,CL_MEM_WRITE_ONLY,sizeof(float)*DATA_SIZE,NULL,NULL);

//load data into the input buffer
clEnqueueWriteBuffer(command_queue,input,CL_TRUE,0,
sizeof(float)*DATA_SIZE,inputData,0,NULL,NULL);

//set the argument list for the kernel
command
clSetKernelArg(kernel,0,sizeof(cl_mem),&input);
clSetKernelArg(kernel,1,sizeof(cl_mem),&output);
global=DATA_SIZE;

//enqueue the kernel command for execution
clEnqueueNDRangeKernel(command_queue,kernel,1,NULL,&global,NULL,0,NULL,NULL);
clFinish(command_queue);

//copy the results from out of the buffer
clEnqueueReadBuffer(command_queue,output,CL_TRUE,0,sizeof(float)*DATA_SIZE,results,0,
NULL,NULL);
```

```
//print the results
printf("output:");
for(int i=0;i<DATA_SIZE;i++)
{
    printf("%f\n",results[i]);
    //printf("no. of times loop run %d\n",count);
}

//cleanup-release OpenCL resources

clReleaseMemObject(input);
clReleaseMemObject(output);
clReleaseProgram(program);
clReleaseKernel(kernel);
clReleaseCommandQueue(command_queue);
clReleaseContext(context);
end_time=clock();
printf("execution time is%f",end_time-
start_time);
_getch();
return 0;

}
```

Reply



**Erik Smistad**



🕒 November 15, 2012 at 16:05

It is normal for the execution time to increase when debugging an application. It doesn't mean that anything is wrong with the program

Reply



**swap** 🕒 August 7, 2012 at 22:26

How will i execute same program on windows 7 with intel HD 4000 GPU.

I have installed Intel openc1 SDK

Reply



**Erik Smistad**



🕒 August 10, 2012 at 14:23

Just remember to link to the lib files included in the Intel OpenCL SDK and add the include folder. If you are using Visual Studio you can add these things in the project settings menu.

Reply



**swap** 🕒 August 7, 2012 at 22:24

how will i execute same program on windows 7 with intel HD 4000.

I have downloaded Intel Openc1 SDK.

Reply



**Lennart** 🕒 July 10, 2012 at 01:03

Thanks!

You officially got me started with OpenCL

(Hilsen fra Bergen)

Reply



**Ricardas** 🕒 July 4, 2012 at 17:56

Great Tutorial. Thanks ! 😊

Reply



**Shinchuan** · June 3, 2012 at 05:26

Hi,

I followed your steps and was able to get  
main.o.

But when i did

```
gcc main.o -o host -L /home/mydir  
/Downloads/ati-stream-sdk-v2.1-lnx64/lib  
/x86_64 -l OpenCL
```

I got

```
/usr/bin/ld: cannot find -lopenCL
```

```
collect2: ld returned 1 exit status
```

I have no idea what this means. Please help!

Reply



**Erik Smistad** · June 4, 2012 at 11:35



It means that the linker can't find the  
libOpenCL file which should lie in your  
/home/mydir/Downloads/ati-  
stream-sdk-v2.1-lnx64/lib/x86\_64  
folder. Make sure you are using a  
large O in "-lOpenCL" and not  
"-lopenCL" as it says in your error  
message: "/usr/bin/ld: cannot find  
-lopenCL"

Reply



**spandei** · April 4, 2012 at 04:05

Thank you!

You write very understandable code.)

To tell the truth your article helped me to  
understand the OpenCL Mechanism better  
than the whole AMD SDK. Keep it up!

Reply



**Bishoy Mikhael** · March 7, 2012 at 23:35

i failed to compile the example, can you please  
review my configuration for the IDE, i've tried  
MS Visual Studio 2010, NetBeans 7.1 and  
Eclipse Indigo using both AMD and NVIDIA  
SDKs on Windows 7 x64 with Nvidia GeForce  
330M graphics card.

i've declared the environment variables for  
CUDA SDK as follows \$(CUDA\_LIB\_PATH),  
\$(CUDA\_BIN\_PATH) and \$(CUDA\_INC\_PATH)  
for (CUDA installation path\lib\x64), (CUDA  
installation path\bin), (CUDA installation  
path\include) respectively. in NetBeans  
TOOLS->Options in C/C++ Code Assistance  
tab i've added the include directory for CUDA  
SDK, then in the project properties in "C  
Compiler" tab i've added the include directory  
path in "Include Directory", and in the "Linker"  
tab i've added the library path in "Additional  
Library Directories" then "openccl.lib" in  
"Additional Dependencies", i don't know what  
to add in the "Compilation Line" or if there is  
another settings i'm missing.

when i build the project i get an error:

```
"/bin/sh: -c: line 0: syntax error near  
unexpected token `{'
```

```
/bin/sh: -c: line 0: `gcc.exe -m64 -c
-(CUDA_INC_PATH) -MM -MP -MF
build/Release/MinGW-Windows
/_ext/141055651/vector_add.o.d -o
build/Release/MinGW-Windows
/_ext/141055651/vector_add.o /C/Users
/Arch/Documents/NetBeansProjects
/vector_add/vector_add.c'
make[2]: *** [build/Release/MinGW-
Windows/_ext/141055651/vector_add.o]
Error 2"
```

Reply



**Erik Smistad**

★ ⌚ March 12, 2012 at 10:44

Always a nightmare to compile on windows... but it should work on Visual Studio 2010 and setting include and library settings should be enough. Don't know what's wrong in your case

Reply



**Bishoy Mikhael**

⌚ March 13, 2012 at 11:20

i've uninstalled VS 2010, netbeans, MinGW, cygwin and the newly installed Microsoft Visual C++ 2008 Redistributables and .NET frameworks, then installed Code::Blocks then i copied the CL directory from the include directory of NVIDIA Toolkit and set the additional libraries in Code::Blocks to point at the OpenCL.lib, guess what, it worked fine without any errors

Reply



**Hao Wang** ⌚ March 1, 2012 at 18:39

Hi,

I'm trying to run OpenCL program on a simulator (gem5).

The simulator supports Full-System mode, that is, first boot Linux from a disk-image and then run the program.

I borrowed the ICD record and libOpenCL.so from AMD SDK, and put them into the proper place in the image file.

But the simulation trace shows that, it fails to find a platform, and then crashes when trying to create a context.

Do you have any suggestions on my situation>

Thank you.

Reply



**Erik Smistad** ⌚ March 5, 2012 at 10:14



Hi

Make sure the ICD files can be read by the program and that the AMD APP SDK and display drivers are properly installed.

Reply



**rupam** ○ February 29, 2012 at 08:00

hi, I am pretty new in GPU. Currently I am working on AMD RADEON 6500 series.... I have written some program on MATLAB 2011b...I want to run the codes on GPU...can u plz instruct me how to run .m codes on GPU...thanx in advance...

Reply



**vegihat** ○ February 27, 2012 at 13:15

Hello Erik,

i try to understand what is the physical meaning of below  
clEnqueueNDRangeKernel's arguments

```
const size_t *global_work_size
```

```
const size_t *local_work_size
```

you used values

```
size_t global_item_size = LIST_SIZE;
```

```
size_t local_item_size = 64
```

which means that we have LIST\_SIZE/64 work-groups, right?

what's the difference between  
local\_item\_size=1 and local\_item\_size = 64?

i want to understand which is the perfect value of the local\_item\_size .

Reply



**Erik Smistad**



○ February 27, 2012 at 14:46

The "perfect" value of local\_item\_size is very system and application dependent. You can omit this parameter and let OpenCL decide by itself. Note that for NVIDIA GPUs the local\_item\_size should be a multiple of 32 (one warp) or else some work-items will be idle. The same applies to AMD GPUs, but with a multiple of 64 instead (one wavefront as they call it). This is because NVIDIA and AMD schedules 32 and 64 work-items atomically on each compute unit.

Reply



**The Ham** ○ February 25, 2012 at 18:53

Hi, to all people that get wrong output (garbage or all 0)

I used this code under gForce M8400 GS and get garbage and error -52 in clEnqueueNDRangeKernel which is wrong kernel arguments.

that is what i changed in this code:  
add 4-th argument for the kernel

```
ret = clSetKernelArg(kernel, 3,  
sizeof(int), &LIST_SIZE);
```

remember that in this code LIST\_SIZE must  
be 64 \* m (m is integer)

This code works ok on my AMD HD 6670  
without any changes,  
dont know why (just started with OpenCL)  
(cant rly add comment on this site!!)

Reply



**The Ham** • February 25, 2012 at 18:52

Hi, to all people that get wrong output  
(garbage or all 0)

I used this code under gForce M8400 GS and  
get garbage and error -52 in  
clEnqueueNDRangeKernel which is wrong  
kernel arguments.

that is what i changed in this code:  
add 4-th argument for the kernel

```
ret = clSetKernelArg(kernel, 3,  
sizeof(int), &LIST_SIZE);
```

remember that in this code LIST\_SIZE must  
be 64 \* m (m is integer)

This code works ok on my AMD HD 6670  
without any changes,  
dont know why (just started with OpenCL)

Reply



**The Ham** • February 25, 2012 at 18:59

i just can get it right..

```
another thing in the kernel  
change output buffer type  
__global int *C to __global  
float *C  
or you get all = 0
```

Reply



**LaKing** • February 22, 2012 at 16:26

Hello. ..

I compiled some sample applications, and get  
this error when running any OpenCL  
application. ...

OpenCL SW Info:

Error -1001 in clGetPlatformIDs Call !!!

I was googling for several hours, got some  
useful info's but could not solve the problem  
yet. Any ideas? ... Thanks.

Reply



**Erik Smistad**

🕒 February 23, 2012 at 13:46

Hi

I think it means that it can't find any OpenCL platforms. Check to see if the .icd files are properly installed. They should exist under /etc/OpenCL/vendors. If no .icd files exists there you have to find them in the downloaded SDK and extract them manually to this location.

Reply



**Otto** 🕒 February 15, 2012 at 19:53

The OpenCL library seems to slow initial execution down a bit. The example above does not really work the GPU. I created the same thing in pure C and it's almost 3 times faster. 😊

time ./cpu >dump)

real 0m0.320s

user 0m0.292s

sys 0m0.024s

time ./gpu >dump)

real 0m0.825s

user 0m0.736s

sys 0m0.100s

Reply



**Erik Smistad**

★ 🕒 February 16, 2012 at 15:48

There is always some overhead with using the GPU such as transferring data back and forth and setup time compared to the CPU. But if you try the same example with a VERY large list/vector you will definitely see a speedup. So if your data parallel problem is very small, using the GPU is usually not faster, but large problems such as volume processing where you have several million elements to be processed you get an enormous speedup.

Reply



**Otto** 🕒 February 15, 2012 at 19:25

Create a Makefile

#—cut—

GCC=gcc

CFLAGS= -c -I/opt/AMDAPP/include

LDFLAGS= -L/opt/AMDAPP/lib/x86\_64

-L/usr/lib -lOpenCL

all: host

host:

\$(GCC) \$(CFLAGS) main.c -o main.o

\$(GCC) main.o -o host \$(LDFLAGS)

clean:

rm -rf host main.o

#—cut—



I'm using Ubuntu 11.10 64bit + AMD-APP-SDK 2.6 + 11.12 drivers and got it working with the Makefile. In Erik's example above, the linking params are in the wrong order, therefore it fails every time.

Reply



**Erik Smistad**



February 16, 2012 at 15:45

I can swear that my compilation command worked before, but when I tried it myself just now on the same system it failed. Maybe it's a new version of gcc or something... anyway thanks for letting me know, I've updated the post.

Reply



**Ajay** February 7, 2012 at 10:48

Hey Erik,

Is there a way to write functions inside the opencl kernel?

Reply



**Erik Smistad**



February 7, 2012 at 13:15

I don't now why you would want to write a function inside a kernel... but you can create a function in OpenCL that is inside the same file as the kernels and then call that function in the kernel (see below). If this is not what you want you may have to use macros or something (like "#define MAX(a,b) a > b ? a : b").

```
int myFunction() {  
    ...  
}
```

```
__kernel void myKernel() {  
    myFunction();  
}
```

Reply



**José** February 4, 2012 at 00:07

*José:*

I tried to install and run your example but the answer is that it can't find the file cl.h  
What should I look for?

I tried to reinstall but it simply don't want to run.

The icd files and the .h are where they are supposed to be.

The installation of the openCL is on the /opt/AMDAPP/

Also, running on a Phenom II 1090T + HD6850

Forgot this information:

The comand lines that I used:

```
main.c @ ~/openCL-teste/
```

comand line used:

```
gcc -I /opt/AMDAPP/include/CL/ main.c -o  
main.o
```

Reply



**Erik Smistad**



February 7, 2012 at 13:11

It is because the include line is  
"#include ", but your include path is  
inside the CL directory. Try this  
instead: "gcc -I /opt/AMDAPP  
/include/ main.c -o main.o"

Reply



**José** February 4, 2012 at 00:03

I tried to install and run your example but the  
answer is that it can't find the file cl.h

What should I look for?

I tried to reinstall but it simply don't want to  
run.

The icd files and the .h are where they are  
supposed to be.

The instalation of the openCL is on the  
/opt/AMDAPP/

Also, running on a Phenom II 1090T +  
HD6850

Reply



**Jesse** January 31, 2012 at 20:25

This is a good guide, so much easier on linux  
than windows

Reply



**Anjil** January 16, 2012 at 17:43

I tried with CPP vesrion of your example and it  
simply displays: clGetPlatformIDs(-1001)

Does this mean the OpenCL is not installed  
properly? Should I re install the NVIDIA driver?

Reply



**Erik Smistad**



January 16, 2012 at 21:32

Yes, it looks like OpenCL is not  
properly installed. Try reinstalling and  
check that the .icd files are present in  
the /etc/OpenCL/vendors folder.  
Download the development driver  
and CUDA toolkit from here:  
[http://developer.nvidia.com/cuda-  
toolkit-40](http://developer.nvidia.com/cuda-toolkit-40) and follow the install  
instructions for linux

Reply



**Anjil** January 13, 2012 at 22:00

The program compiles fine but the results are  
not as expected:

989 + 35 = 0  
990 + 34 = 0  
991 + 33 = 0  
992 + 32 = 0  
993 + 31 = 0  
994 + 30 = 0  
995 + 29 = 0  
996 + 28 = 0  
997 + 27 = 0  
998 + 26 = 0  
999 + 25 = 0  
1000 + 24 = 0  
1001 + 23 = 0  
1002 + 22 = 0  
1003 + 21 = 0  
1004 + 20 = 0  
1005 + 19 = 0  
1006 + 18 = 0  
1007 + 17 = 0  
1008 + 16 = 0  
1009 + 15 = 0  
1010 + 14 = 0  
1011 + 13 = 0  
1012 + 12 = 0  
1013 + 11 = 0  
1014 + 10 = 0  
1015 + 9 = 0  
1016 + 8 = 0  
1017 + 7 = 0  
1018 + 6 = 0  
1019 + 5 = 0  
1020 + 4 = 0  
1021 + 3 = 0  
1022 + 2 = 0  
1023 + 1 = 0

Please let me know whats the issue here?

Reply



**klm123** ⌚ June 4, 2012 at 17:58  
Why nobody answer? No one can explain it???!!

Reply



**klm123**  
⌚ June 4, 2012 at 17:59  
I've got same result and it does not depend on the content of vector\_add\_kernel.cl file.

Reply



**Erik Smistad** ★  
⌚ June 10, 2012 at 09:15  
Some of the function calls must be failing. Use the error argument to check which and check which error type it is. Most likely OpenCL is not properly

installed.

Reply



**Pratik Anand**

⌚ July 8,  
2016 at 03:37

The platform  
selection call  
can be  
modified as  
per this URL  
[http://stackoverflow.com  
/questions  
/30079550  
/opengl-  
clgetdeviceids-  
returns-  
1-when-  
gpu-and-0-  
when-cpu](http://stackoverflow.com/questions/30079550/opengl-getdeviceids-returns-1-when-gpu-and-0-when-cpu)

In that case,  
the OpenCL  
starts  
detecting the  
GPU.

Reply



**Paraita Wohler**

⌚ July 11, 2012 at 13:57

You should double check if your  
NVidia card actually support OpenCL,  
depending on your card, some  
operations might not be  
implemented. Also check what  
clGetPlatformIDs gives you back  
(error -1001 ?), I had alot of trouble  
with this too, at the end it was the  
card I was using...

Reply



**Paraita Wohler**

⌚ July 11, 2012 at 14:00

in fact, even if your hardware  
doesn't support OpenCL, you  
can still run OpenCL code, but  
it won't work (returning error  
1001 which is a "catch-all"  
error for almost everything  
related to hardware)

Reply



**blogger** ⌚ December 16, 2011 at 05:20

Great tutorial! Thanks a lot, it helped me get  
started with opengl

Reply



**Guilherme** ⌚ August 21, 2011 at 10:55

Great tutorial! Congrats dude.

I'm having just a little problem. I'm trying to  
use M\$ Visual C++ 2010 Express: I created a  
new empty project, included the path to the  
necessary files (AMD SDK – wich are located

on C:\Program Files (x86)\AMD APP\include)  
on the "include settings" of the project,  
created two files on this project with the  
names test.cpp and vector\_add\_kernel.cl  
with the exact same codes on this tutorial. I  
get a lot of errors:  
(first line in the output)- 1>classe.obj : error  
LNK2019: unresolved external symbol  
\_clReleaseContext@4 referenced in function  
\_main

Am I doing something wrong? Should I use  
another IDE? I'm really newbie on OpenCL  
(and a little bit in C++). And sorry for my  
english 'cause I'm brazilian.

Reply



**Erik Smistad**

★ ⌚ August 21, 2011 at 13:23

You have forgot to add additional  
library paths in Visual Studio. You  
need to go to project settings, c++  
linker, and then add C:\Program Files  
(x86)\AMD APP\lib\x86 to additional  
libraries, and on input write  
OpenCL.lib...

I think I will add some notes on how  
to set up OpenCL in visual studio in  
this post in a few days

Reply



**Guilherme**

⌚ August 22, 2011 at 03:02

It worked! Thanks. =]

Reply



**Ben** ⌚ August 13, 2011 at 23:19

Yours is probably the clearest and most  
concise "getting started" guide I've seen, so  
many thanks for this.

Incidentally, both your C and C++ example  
programs also work perfectly on Windows  
using AMD's OpenCL implementation with  
MinGW-w64.

Reply



**rwi** ⌚ July 13, 2011 at 10:55

Hi Erik,  
your guide helped me getting started. I'm  
using the nVIDIA CUDA Toolkit 4 on Ubuntu  
10.10. I had trouble compiling with gcc and  
ended up with:

```
gcc -lOpenCL -I /usr/lib64/ -I /usr/local  
/cuda/include main.c -o host
```

where again lOpenCL starts with a lowercase  
L and the two locations are prefixed with  
-(capital i) (just in case). This worked for me.

Thanks a lot!

Reply



**Jeffin** ⌚ March 21, 2012 at 15:26

Hi rwi,  
Thanks a lot for your post. The  
command options that you provided  
was very helpful for my build issues  
while trying with Netbeans IDE.  
Thanks again  
Jf

Reply



**none** · June 30, 2011 at 08:03

I feel with you and did not read your  
desperate endeavours (trying to help the ppl  
here) till the end. But one thing i had to fight  
with a year ago was that the nvidia drivers  
that came with ubuntu (i mean the proprietary  
ones) did not place two links right. One of  
them was the opencl... (as above). One  
solution was to install the driver directly from  
nvidia. The other one is to search the net and  
the answer will come. It's just two softlinks.

Reply



**Erik Smistad** · July 2, 2011 at 16:00

★ Thanks for the tip

Reply



**Dan** · June 9, 2011 at 01:07

Can we run OpenCL without installing one of  
the SDKs? if so, what is the best way to do it?

Reply



**Erik Smistad** · June 9, 2011 at 10:26

★ You need an OpenCL implementation  
installed to use it. If you have an intel  
CPU and want to run OpenCL on it  
check out their SDK at  
[http://software.intel.com/en-  
us/articles/opencl-sdk/](http://software.intel.com/en-us/articles/opencl-sdk/)

Reply



**michael** · May 27, 2011 at 05:21

Here is my result, it's not correct. why?

0 + 1000 = 0

1 + 999 = 0

2 + 998 = 0

3 + 997 = 0

4 + 996 = 0

5 + 995 = 0

6 + 994 = 0

7 + 993 = 0

8 + 992 = 0

9 + 991 = 0

10 + 990 = 0

11 + 989 = 0

12 + 988 = 0

Reply



**Anjil** · January 13, 2012 at 21:58

I am getting the same response, all  
results as zero's any idea how to fix  
this?

Reply



Harshit · May 26, 2011 at 16:33

Hi Erik,

I'm a noob..

I get this error when i Compile the program using

```
Documents/cuda/Assignment/Ex. 2$ gcc -c -Wall -I /usr/local/cuda/include cpu.c -o main
```

```
cpu.c:8: warning: return type defaults to 'int'
```

```
cpu.c: In function 'main':
```

```
cpu.c:52: warning: control reaches end of non-void function
```

But I'm getting the right result.. 😊

Is something wrong??

Reply



Erik Smistad · May 26, 2011 at 18:03



They are only warnings... and it seems you only need to add a "return 0;" in the end of your main method to get rid of it.

Reply



Ajay · May 26, 2011 at 10:51

I changed it. But I want to know that whether the program is running on the GPU or not..

Will I get error message if the program is not running on GPU or will it run on CPU if it doesn't find the GPU after changing CL\_DEVICE\_TYPE\_DEFAULT to CL\_DEVICE\_TYPE\_GPU.

I'm asking this coz I wanted to measure performance of an OpenCL program by using clGetEventProfilingInfo. I have done everything right but I'm getting timings only if i set CL\_DEVICE\_TYPE\_CPU. If i set it to CL\_DEVICE\_TYPE\_DEFAULT or CL\_DEVICE\_TYPE\_GPU, then I'm getting answer as 0.00.

Why am I getting timings only on CPU?? Is the program running on CPU even after setting CL\_DEVICE\_TYPE\_GPU?? Thats my doubt..

Reply



Erik Smistad · May 26, 2011 at 18:02



There is no error checking in the example above, because I wanted the example to be as simple as possible. So it can happen that it doesn't run at all. You can print out ret\_num\_devices and see if it is larger than 0.

I recommend the free OpenCL/OpenGL profiler tool gDEBugger – <http://www.gremedy.com/>

With this program you should be able to see exactly where your code is run and how fast.

Reply



Ajay · May 25, 2011 at 18:30

thanks Erik,

I have another doubt..

How can i find whether the code is actually running on GPU or not??

Reply



Erik Smistad · May 25, 2011 at 21:40

★ Change `CL_DEVICE_TYPE_DEFAULT` to `CL_DEVICE_TYPE_GPU` on line 43 and it should only run on GPUs.

Reply



Ajay · May 18, 2011 at 10:11

I am new to OpenCL programming. I successfully compiled the sample program. then I tried './host' to run it. but I'm getting an error..

./host: error while loading shared libraries: libOpenCL.so.1: cannot open shared object file: No such file or directory

Reply



Erik Smistad · May 18, 2011 at 10:35

★ That error means that you have forgot to set the `LD_LIBRARY_PATH`, see the description in the article above.

Reply



Kurt · May 16, 2011 at 14:40

Hi Erik – Thanks for your blog. Has given me good help.

Now on a slightly different topic: ATI plus NVidia plus CPU.

The reason I want to use OpenCL is so my software can work on any of these parallel platforms – I might even create a driver for high-end DSPs after I learn a bit more.

However, it seems difficult to find platform-neutral info on how to get set up. Is the platform (NV vs ATI) entirely based on what `LD_LIBRARY_PATH` points to? If so, is it impossible to use –device cpu in a machine equipped with NVidia GPU?

IBM has what I think is intended to be a solution called OpenCLCommonRuntime (aplhaworks), but I can't seem to get it to work with AMD –device cpu. Do you have any experience with this? IBM is, at least, a neutral source.

Reply



Erik Smistad · May 18, 2011 at 10:34

★ Hi Kurt

My personal experience with NVidia's and ATI/AMD's OpenCL implementation is that NVidia still only has a buggy version of OpenCL 1.0. Their implementation has some serious issues when using images.



ATI/AMD's implementation on the other hand has worked great and is updated to 1.1.

The thing about OpenCL is that it is to some degree portable, (it is portable if you don't take into account all of NVidia's bugs), but it is not necessary performance portable. Some OpenCL code might be fast on an ATI card while it may not be on an NVidia card. Because of their different architectures. This might and hopefully will be less true when their OpenCL compilers have matured.

It is possible to use the CPU with OpenCL on a machine with a NVidia GPU.

I haven't tried OpenCLCommonRuntime, but it looks quite promising

Reply



**Krish** · April 6, 2011 at 08:04

hi erik

it worked .....

thanx

Reply



**Krish** · March 30, 2011 at 07:04

Hi erik

(for the post no:28)

even after specifying cuda installation path

```
[root@ghost OpenCL]# gcc -I /usr/local
```

```
/cuda/include/CL/ -lOpenCL hello.c
```

In file included from hello.c:12:0:

```
cl.h:32:28: fatal error: CL/cl_platform.h: No
```

```
such file or directory compilation terminated.
```

i dont understand what is happening..... how to correct the error.....

Reply



**Erik Smistad**



· March 30, 2011 at 23:39

It is supposed to be -l (capitol l) and

not -l in front of the include path

Reply



**paolo** · March 29, 2011 at 18:43

Hi Erik,

thanks for posting these info

I tried to compile your example on a Mac OSX 10.6 with Envidia GEFORCE 9400M.

I have the drivers installed correctly (I believe), and so it the opencl toolkit.

When I call the linking part of it, I get the error on "host"

it says: "host: no such file or directory"

what is host exactly?

thank you

Reply



**Erik Smistad**

★ ○ March 31, 2011 at 14:31

Host is just the name of the executable you are creating when compiling. The name of the executable is not important. There must be something wrong with your compile command. Paste your compile command here

Reply



**Krish** ○ March 28, 2011 at 12:17

hi

i am just started learning open cl

i have installed nvidia sdk etc... as given on Nvidia website...

i am using Fedora 14 – 64 bit edition when i am compiling

as

```
$ gcc hello.c -lopenCL
```

it gives me

```
cl.h:32:28: fatal error: CL/cl_platform.h: No such file or directory; compilation terminated.
```

pls need help from u to rectify it ....

Reply



**Erik Smistad**

★ ○ March 28, 2011 at 14:43

Have you remembered to specify the include path -I to the include folder of your CUDA installation?

Reply



**Freddan** ○ October 23, 2010 at 18:50

```
main.c:(.text+0x129): undefined reference to `clGetPlatformIDs'
```

```
main.c:(.text+0x150): undefined reference to `clGetDeviceIDs'
```

```
main.c:(.text+0x17b): undefined reference to `clCreateContext'
```

```
main.c:(.text+0x19e): undefined reference to `clCreateCommandQueue'
```

```
main.c:(.text+0x1cc): undefined reference to `clCreateBuffer'
```

```
main.c:(.text+0x1fa): undefined reference to `clCreateBuffer'
```

```
main.c:(.text+0x228): undefined reference to `clCreateBuffer'
```

```
main.c:(.text+0x27c): undefined reference to `clEnqueueWriteBuffer'
```

```
main.c:(.text+0x2cc): undefined reference to `clEnqueueWriteBuffer'
```

```
main.c:(.text+0x2ef): undefined reference to `clCreateProgramWithSource'
```

```
main.c:(.text+0x31f): undefined reference to `clBuildProgram'
```

```
main.c:(.text+0x33a): undefined reference to `clCreateKernel'
```

```
main.c:(.text+0x361): undefined reference to `clSetKernelArg'
```

```
main.c:(text+0x384): undefined reference to
`clSetKernelArg'
main.c:(text+0x3aa): undefined reference to
`clSetKernelArg'
main.c:(text+0x411): undefined reference to
`clEnqueueNDRangeKernel'
main.c:(text+0x47f): undefined reference to
`clEnqueueReadBuffer'
main.c:(text+0x4e7): undefined reference to
`clFlush'
main.c:(text+0x4f6): undefined reference to
`clFinish'
main.c:(text+0x508): undefined reference to
`clReleaseKernel'
main.c:(text+0x51a): undefined reference to
`clReleaseProgram'
main.c:(text+0x529): undefined reference to
`clReleaseMemObject'
main.c:(text+0x538): undefined reference to
`clReleaseMemObject'
main.c:(text+0x54a): undefined reference to
`clReleaseMemObject'
main.c:(text+0x559): undefined reference to
`clReleaseCommandQueue'
main.c:(text+0x568): undefined reference to
`clReleaseContext'
```

Reply



**Erik Smistad**



🕒 October 27, 2010 at 16:57

You need to link the OpenCL library when you compile: (example: gcc -L /path-to-the-lib-folder-with-OpenCL-libfile/ -l OpenCL main.o -o host)

Reply



**Rohan**

🕒 February 23, 2011 at 09:17

I too got the same errors and it actually happens during the first command not the second. I even created the symlinks as stated above by a user, but nothing works 😞

Reply



**Erik Smistad**



🕒 February 23, 2011 at 14:12

What OS and OpenCL implementation are you using? And what is the exact commands you are running? Because this works perfectly on my system...

Reply



**snk**

🕒 February 23, 2011

at 23:26

Hello,

I got these errors, but the solution was simple. Try the `-c` option for the first command. This option will prevent the gcc from linking, and then the linking can be done with the second command. It worked for me. I hope it will work for you too.

Reply



**Erik Smistad**

★  
🕒 February 24, 2011 at 17:26

Thanks for clearing that up snk. It seems that the `-c` flag was missing in the first example.

Reply



**Lucas Campos**

🕒 May 30, 2011 at 14:15

I had these same problems on linking. I'm using nVidia's SDK, on Mint 10 x64 and this solution did not work for me. But, when I compiled using just

```
gcc -I path-to-opencl-common-inc -lopenCL main.c -o test, worked like a charm. Using SDK 4.0 was
```

```
gcc -lOpenCL -I ~/NVIDIA_GPU_Computing_SDK
```

```
/OpenCL
/common/inc/
main.c -o
test
```

Important to  
notice that  
its lOpenCL  
(lower case  
l) and -l  
(upper case  
-l)

Reply



tom · September 4, 2010 at 02:25

```
0 + 1000 = 1561096336
1 + 999 = 32714
2 + 998 = 1561096336
3 + 997 = 32714
4 + 996 = 6311776
5 + 995 = 0
6 + 994 = 6311776
7 + 993 = 0
8 + 992 = 2147483647
9 + 991 = 0
10 + 990 = 12370275
11 + 989 = 0
12 + 988 = 788225073
```

Reply



tom · September 4, 2010 at 02:24

once everything compiles, running the  
program, it actually returns garbage !?  
the vectors are not added but garbage is  
displayed.

Reply



Erik Smistad

★ · September 4, 2010 at 13:06

I have experienced this myself when  
trying to run this on a NVidia graphics  
card that doesn't support OpenCL.

You can check to see if your GPU  
supports OpenCL here:

[http://www.nvidia.com/object/cuda\\_gpus.html](http://www.nvidia.com/object/cuda_gpus.html)

Also make sure that you have the  
proper dev drivers. If you are using  
linux, check out "Developer Drivers  
for Linux (256.40)" at

[http://developer.nvidia.com/object/cuda\\_3\\_1\\_downloads.html](http://developer.nvidia.com/object/cuda_3_1_downloads.html)

Reply



tom

· September 5, 2010 at 08:55

my nvidia quadro fx3700 is  
cuda/opencl compatible

Reply



tom

· September 5, 2010 at 08:56

and I installed the newest

drivers, 256.53. this is on  
opensuse 11.1 (64bit)

Reply



**Erik Smistad**

★ @ September 5, 2010

at 22:58

Hm.. should have  
worked...

You could try the C++  
version at

<http://www.thebigblob.com/using-the-cpp-bindings-for-opencl/>

which use exceptions  
and see if it catches  
some error. Maybe  
that could help you  
to find out what's  
failing

Reply



**tom** @ September 4, 2010 at 01:25

then I get: "undefined reference to  
`clGetPlatformIDs'" and similar errors all  
relating to the functions defined in  
cl\_platform.h (which gcc finds, I checked). any  
hints appreciated.

Reply



**tom** @ September 4, 2010 at 02:23

symlinks need to be created.  
libOpenCL.so.1.0.0 comes from  
nvidia. libOpenCL.so and  
libOpenCL.so.1 are needed symlinks  
that need to be created (command  
'ln')

Reply



**Rohan**

@ February 23, 2011 at 08:59

I am getting the same  
problem. Could please write  
the exact command that  
need to be run to create the  
symlinks?

Reply



**tom** @ September 3, 2010 at 23:57

also, the link above to download the source  
code leads to a 404 file not found.

Reply



**Erik Smistad**

★ @ September 4, 2010 at 12:59

Sorry about that. The link has been  
fixed now

Reply



**tom**

@ September 6, 2010 at 10:11  
using your files exactly, still  
getting garbage output of the

vector addition.

Reply



**Erik Smistad**

🕒 September 6, 2010  
at 13:22

Then I'm out of ideas.. it works fine on all of my ATI cards. I will try to run the code on some machines with nvidia cards at the university and see if I get the same problem there.

Reply



**Name**  
(required)

🕒 September 7, 2010 at 02:29

```
I added a
printf() to
your code
after line 43:
if (ret !=
CL_SUCCESS)
{
printf("Error:
Failed to
query
platforms!
(%d)\n", ret);
return
EXIT_FAILURE;
}
```

after  
compiling,  
running it  
gives me this  
error: "Failed  
to query  
platforms  
(-1001)"

Reply



**tom**

🕒 September 10, 2010 at 11:47

have you had  
a chance to  
try it out at  
your  
university?

it's still not  
working on  
my side 😞

Reply



Erik  
Smistad

○

September  
10, 2010  
at 12:27

Yes, I  
tested

it

just

now

on a

geforce

8400

at

the

university.

The

installation

was

no

problem.

First

I

downloaded

the

CUDA

toolkit

and

installed

that.

Then

I

installed

the

CUDA

devdriver,

set

LD\_LIBRARY\_PATH

to

/usr/local

/cuda/lib64

and

ran

the

example,

no

problem...

I will

add

the

procedure

for

nvidia

to

the

post

soon.

I

don't

understand



why  
it  
doesn't  
work  
for  
you..  
you  
have  
a  
nvidia.icd  
file  
in  
your  
/etc/OpenCL  
/vendors  
folder,  
right?



**Anjil**

⊙

January

14,  
2012

at  
20:48

Hi  
Erik,

Is  
the  
garbage  
value  
issue  
posted  
by  
few  
developers  
above  
fixed?

Is  
there  
any  
solution  
available.  
Please

let  
me  
know,  
I am  
facing  
the  
same  
problem,  
Every  
time  
I  
execute  
the  
sample,  
it  
simply  
returns

either  
zero  
some  
garbage  
value  
which  
is  
not  
relevant.  
I am  
working  
with  
Ubuntu  
11.04  
with  
NVIDIA  
graphics  
driver  
installed.



**Erik  
Smistad**

January  
15,  
2012  
at  
15:01  
There  
is no  
error  
checking  
in  
the  
example  
above,  
because  
I  
wanted  
the  
example  
to be  
as  
simple  
as  
possible.  
So it  
can  
happen  
that  
the  
kernel  
doesn't  
run  
at  
all,  
maybe  
because  
OpenCL  
is  
not  
properly

installed.  
Try  
using  
the  
C++  
bindings  
in  
this  
post  
<http://www.thebigblob.com/using-the-cpp-bindings-for-opengl/>  
as  
this  
example  
has  
error  
checking.  
See  
what  
error  
message  
you  
get,  
maybe  
that  
can  
help  
you  
find  
out  
what  
the  
problem  
is.  
Post  
it  
here  
afterwards.  
I got  
this  
error  
myself  
once  
and  
found  
that  
it  
was  
because  
OpenCL  
wasn't  
properly  
installed,  
was  
lacking  
icd

files  
or  
something  
like  
that.  
Hope  
that  
helps!



**tom** · September 3, 2010 at 23:52  
gcc question: gcc can't find the cl.h, even though it's in the same directory as the source file, also it can't find it even after explicitly specifying the path!!!!???

Reply



**tom** · September 4, 2010 at 01:23  
ok, the -l is not a lower case "l" it is an uppercase "I": I

Reply



**Achtkraut** · August 28, 2010 at 09:18  
SSE2 support was only added in ATI Stream SDK v2.2

Reply



**Jillian Marohnic** · July 21, 2010 at 19:47  
Is it the case that the CPU must be sse3? I followed your blog but the OpenCL samples don't run successfully (have sse2). Also, when you say to decompress the registration in the root folder, I assumed you meant the sdk root folder ... might want to clarify (for dummies like me) that it is the file system root folder.  
Thanks for the article ... nicely written.

Reply



**Erik Smistad** · July 21, 2010 at 20:24  
★ I don't believe that the CPU must support SSE3. I don't see why it shouldn't support CPUs with SSE2. Can you compile the example and does it run without any errors? What exactly happens when you try to run the example?

I don't mean that you should decompress the SDK to the root folder of your system. What I mean is that you should extract the .icd files of the icd-registration.tgz file to the folder /etc/OpenCL/vendors. This is ofcourse only if you are running on linux, I haven't tried setting OpenCL on windows just yet. You have to do this to make it work.

Reply

LEAVE A REPLY

Comment

Name

Email

Website

Post Comment

- Home
- About
- Research
- FAST
- Blog posts

