

Pre-compile the OpenCL Kernel Program - Part 2

Sat 22 November 2014

Tags [opengl](#)

Posted by [Logan](#)

In the [part 1](#) of this article, we have mentioned how to pre-compile the OpenCL kernel program and load the pre-compiled binaries with the OpenCL API.

However, I was using the `ioc64` command from the [Intel OpenCL SDK](#) to pre-compile the kernel program. This command might be unavailable in the other OpenCL implementations. How could we get the compiled binaries in those implementations?

After checking the [manual](#), it shows that we can retrieve the compiled binaries with `clGetProgramInfo()` after the execution of `clBuildProgram()`. In detail, we need four steps to get the binaries:

1. Allocate the an array of `size_t` to save the size of each binaries.
2. Get the **size** of each binaries with `clGetProgramInfo(program, CL_PROGRAM_BINARY_SIZES, ...)`.
3. Allocate the buffers for the binaries. The size of each buffers should be greater than or equal to the size returned in the second step.
4. Get the **binaries** with `clGetProgramInfo(program, CL_PROGRAM_BINARIES, ...)`.

Here's the code listing:

```
cl_int write_binaries(cl_program program, unsigned num_devices,
                    cl_uint platform_idx) {
    unsigned i;
```

```
cl_int err = CL_SUCCESS;
size_t *binaries_size = NULL;
unsigned char **binaries_ptr = NULL;

// Read the binaries size
size_t binaries_size_alloc_size = sizeof(size_t) * num_devices;
binaries_size = (size_t *)malloc(binaries_size_alloc_size);
if (!binaries_size) {
    err = CL_OUT_OF_HOST_MEMORY;
    goto cleanup;
}

err = clGetProgramInfo(program, CL_PROGRAM_BINARY_SIZES,
                        binaries_size_alloc_size, binaries_size, N
if (err != CL_SUCCESS) {
    goto cleanup;
}

// Read the binaries
size_t binaries_ptr_alloc_size = sizeof(unsigned char *) * num_de
binaries_ptr = (unsigned char **)malloc(binaries_ptr_alloc_size);
if (!binaries_ptr) {
    err = CL_OUT_OF_HOST_MEMORY;
    goto cleanup;
}
memset(binaries_ptr, 0, binaries_ptr_alloc_size);
for (i = 0; i < num_devices; ++i) {
    binaries_ptr[i] = (unsigned char *)malloc(binaries_size[i]);
    if (!binaries_ptr[i]) {
        err = CL_OUT_OF_HOST_MEMORY;
        goto cleanup;
    }
}

err = clGetProgramInfo(program, CL_PROGRAM_BINARIES,
                        binaries_ptr_alloc_size,
                        binaries_ptr, NULL);
if (err != CL_SUCCESS) {
    goto cleanup;
}
```

```
}

// Write the binaries to file
for (i = 0; i < num_devices; ++i) {
    // Create output file name
    char filename[128];
    snprintf(filename, sizeof(filename), "cl-out_%u-%u.bin",
              (unsigned)platform_idx, (unsigned)i);

    // Write the binary to the output file
    write_file(filename, binaries_ptr[i], binaries_size[i]);
}

cleanup:
// Free the return value buffer
if (binaries_ptr) {
    for (i = 0; i < num_devices; ++i) {
        free(binaries_ptr[i]);
    }
    free(binaries_ptr);
}
free(binaries_size);

return err;
}
```

Based on these OpenCL APIs, I have written a simple OpenCL kernel program compiler to translate `.cl` files into pre-compiled binaries. Please refer to [cl-compile.c](#) for the source code.

This completes our discussion on the compilation of OpenCL kernel binaries. In the next post, I would like to give an introduction to OpenCL SPIR, the official intermediate representation for OpenCL kernel programs.