

This repository

Search

Pull requests

Issues

Gist

clMathLibraries / cIBLAS

Watch

84

Star

518

Fork

159

<> Code

Issues

63

Pull requests

0

Projects

0

Wiki

Pulse

Graphs

a software library containing BLAS functions written in OpenCL

371 commits

4 branches

7 releases

31 contributors

Apache-2.0

Branch: master

New pull request

Create new file

Upload files

Find file

Clone or download

| | | |
|-----------------|--|---------------------------------|
| tingxingdong | committed on GitHub Merge pull request #312 from clMathLibraries/revert-311-device-all | Latest commit cf91139 on 24 Mar |
| doc | add perf data for fiji | a year ago |
| src | Revert "Use CL_DEVICE_TYPE_ALL for all samples" | 2 months ago |
| .gitattributes | Initial check-in of open source cIBLAS code | 4 years ago |
| .gitignore | Removing the pedantic flag from gcc compiles | a year ago |
| .travis.yml | Update .travis.yml and appveyor.yml | a year ago |
| CHANGELOG | Cleanup of txt not in the spirit of the clMath Apache license | 4 years ago |
| CONTRIBUTING.md | Point the CONTRIBUTING wiki links to the correct repository | 9 months ago |
| LICENSE | delete appendix in license file | 2 years ago |
| NOTICE | Initial check-in of open source cIBLAS code | 4 years ago |
| README.md | Update README with release notes | 4 months ago |
| appveyor.yml | Update .travis.yml and appveyor.yml | a year ago |

README.md

Build Status

| Build branch | master | develop |
|-------------------|---------------|---------------|
| GCC/Clang x64 | build passing | build passing |
| Visual Studio x64 | build passing | build passing |

cIBLAS

This repository houses the code for the OpenCL™ BLAS portion of clMath. The complete set of BLAS level 1, 2 & 3 routines is implemented. Please see Netlib BLAS for the list of supported routines. In addition to GPU devices, the library also supports running on CPU devices to facilitate debugging and multicore programming. APPML 1.12 is the most current generally available pre-packaged binary version of the library available for download for both Linux and Windows platforms.

The primary goal of cIBLAS is to make it easier for developers to utilize the inherent performance and power efficiency benefits of heterogeneous computing. cIBLAS interfaces do not hide nor wrap OpenCL interfaces, but rather leaves OpenCL state management to the control of the user to allow for maximum performance and flexibility. The cIBLAS library does generate and enqueue optimized OpenCL kernels, relieving the user from the task of writing, optimizing and maintaining kernel code themselves.

cIBLAS update notes 01/2017

- v2.12 is a bugfix release as a rollup of all fixes in /develop branch
 - Thanks to @pavanky, @iotamudelta, @shahsan10, @psyhtest, @haahh, @hughperkins, @tfauck @abhiShandy, @IvanVergiliev, @zouglob, @mgates3 for contributions to cIBLAS v2.12
- Summary of fixes available to read on the releases tab

clBLAS library user documentation

[Library and API documentation](#) for developers is available online as a GitHub Pages website

Google Groups

Two mailing lists have been created for the clMath projects:

- clmath@googlegroups.com - group whose focus is to answer questions on using the library or reporting issues
- clmath-developers@googlegroups.com - group whose focus is for developers interested in contributing to the library code itself

clBLAS Wiki

The [project wiki](#) contains helpful documentation, including a [build primer](#)

Contributing code

Please refer to and read the [Contributing](#) document for guidelines on how to contribute code to this open source project. The code in the /master branch is considered to be stable, and all pull-requests should be made against the /develop branch.

License

The source for clBLAS is licensed under the [Apache License, Version 2.0](#)

Example

The simple example below shows how to use clBLAS to compute an OpenCL accelerated SGEMM

```
#include <sys/types.h>
#include <stdio.h>

/* Include the clBLAS header. It includes the appropriate OpenCL headers */
#include <clBLAS.h>

/* This example uses predefined matrices and their characteristics for
 * simplicity purpose.
 */

#define M  4
#define N  3
#define K  5

static const cl_float alpha = 10;

static const cl_float A[M*K] = {
11, 12, 13, 14, 15,
21, 22, 23, 24, 25,
31, 32, 33, 34, 35,
41, 42, 43, 44, 45,
};
static const size_t lda = K;          /* i.e. lda = K */

static const cl_float B[K*N] = {
11, 12, 13,
21, 22, 23,
31, 32, 33,
41, 42, 43,
51, 52, 53,
};
static const size_t ldb = N;          /* i.e. ldb = N */

static const cl_float beta = 20;

static cl_float C[M*N] = {
    11, 12, 13,
```

```

    21, 22, 23,
    31, 32, 33,
    41, 42, 43,
};
static const size_t ldc = N;          /* i.e. ldc = N */

static cl_float result[M*N];

int main( void )
{
    cl_int err;
    cl_platform_id platform = 0;
    cl_device_id device = 0;
    cl_context_properties props[3] = { CL_CONTEXT_PLATFORM, 0, 0 };
    cl_context ctx = 0;
    cl_command_queue queue = 0;
    cl_mem bufA, bufB, bufC;
    cl_event event = NULL;
    int ret = 0;

    /* Setup OpenCL environment. */
    err = clGetPlatformIDs( 1, &platform, NULL );
    err = clGetDeviceIDs( platform, CL_DEVICE_TYPE_GPU, 1, &device, NULL );

    props[1] = (cl_context_properties)platform;
    ctx = clCreateContext( props, 1, &device, NULL, NULL, &err );
    queue = clCreateCommandQueue( ctx, device, 0, &err );

    /* Setup clBLAS */
    err = clblasSetup( );

    /* Prepare OpenCL memory objects and place matrices inside them. */
    bufA = clCreateBuffer( ctx, CL_MEM_READ_ONLY, M * K * sizeof(*A),
                          NULL, &err );
    bufB = clCreateBuffer( ctx, CL_MEM_READ_ONLY, K * N * sizeof(*B),
                          NULL, &err );
    bufC = clCreateBuffer( ctx, CL_MEM_READ_WRITE, M * N * sizeof(*C),
                          NULL, &err );

    err = clEnqueueWriteBuffer( queue, bufA, CL_TRUE, 0,
                               M * K * sizeof( *A ), A, 0, NULL, NULL );
    err = clEnqueueWriteBuffer( queue, bufB, CL_TRUE, 0,
                               K * N * sizeof( *B ), B, 0, NULL, NULL );
    err = clEnqueueWriteBuffer( queue, bufC, CL_TRUE, 0,
                               M * N * sizeof( *C ), C, 0, NULL, NULL );

    /* Call clBLAS extended function. Perform gemm for the lower right sub-matrices */
    err = clblasSgemm( clblasRowMajor, clblasNoTrans, clblasNoTrans,
                      M, N, K,
                      alpha, bufA, 0, lda,
                      bufB, 0, ldb, beta,
                      bufC, 0, ldc,
                      1, &queue, 0, NULL, &event );

    /* Wait for calculations to be finished. */
    err = clWaitForEvents( 1, &event );

    /* Fetch results of calculations from GPU memory. */
    err = clEnqueueReadBuffer( queue, bufC, CL_TRUE, 0,
                              M * N * sizeof(*result),
                              result, 0, NULL, NULL );

    /* Release OpenCL memory objects. */
    clReleaseMemObject( bufC );
    clReleaseMemObject( bufB );
    clReleaseMemObject( bufA );

    /* Finalize work with clBLAS */
    clblasTeardown( );

    /* Release OpenCL working objects. */
    clReleaseCommandQueue( queue );
    clReleaseContext( ctx );

    return ret;

```

```
}
```

Build dependencies

Library for Windows

- Windows® 7/8
- Visual Studio 2010 SP1, 2012
- An OpenCL SDK, such as APP SDK 2.8
- Latest CMake

Library for Linux

- GCC 4.6 and onwards
- An OpenCL SDK, such as APP SDK 2.9
- Latest CMake

Library for Mac OSX

- Recommended to generate Unix makefiles with cmake

Test infrastructure

- Googletest v1.6
- Latest Boost
- CPU BLAS
- Netlib CBLAS (recommended) Ubuntu: install by "apt-get install libblas-dev" Windows: download & install lapack-3.6.0 which comes with CBLAS
- or ACML on windows/linux; Accelerate on Mac OSX

Performance infrastructure

- Python