

Chuck Anderson

Professor, Computer Science, Colorado State University

[About](#) [Research](#) [Publications](#) [Teaching](#) [CV](#)

Numpy versus Pytorch

August 26, 2017 by anderson

Here we compare the accuracy and computation time of the training of simple fully-connected neural networks using numpy and pytorch implementations and applied to the MNIST data set. The Adam optimization algorithm in numpy and pytorch are compared, as well as the Scaled Conjugate Gradient optimization algorithm in numpy.

The original notebook is [available here](#).

Additional comments and explanations will be added shortly. If you have suggestions or corrections, please write to chuck.anderson@colostate.edu.

In []:

```
!nvidia-smi
```

Recent Posts

[Comparing Numpy, Pytorch, and autograd on CPU and GPU](#)
October 13, 2017
[Numpy versus Pytorch](#)
August 26, 2017

Mon Jul 17 15:16:36 2017

```

+-----+
| NVIDIA-SMI 375.66                Driver Version: 375.66
+-----+-----+-----+
| GPU Name      Persistence-M| Bus-Id  Disp |
| Fan  Temp  Perf  Pwr:Usage/Cap|  Memory- |
|=====+=====|
|  0  NVS 315     Off | 0000:02:00.0  N/A |
| 30%  42C   P8   N/A /  N/A |  35MiB /  964MiB |
+-----+-----+-----+
|  1  GeForce GTX 980  Off | 0000:03:00.0  Of |
| 26%  36C   P8   13W / 180W |  281MiB /  4038MiB |
+-----+-----+-----+

+-----+
| Processes:                         GPU Memory Usage
| GPU      PID  Type  Process name                        Memory |
|=====+=====|
|  0              Not Supported
|  1    8001   C   .../anderson/anaconda3/envs/i
+-----+

```

In []:

```

import numpy as np
import pickle
import gzip
import json
import time
import sys
import subprocess

# For my numpy neural network implementation
import neuralnetworks as nn

```

[Fast Reinforcement Learning After Pretraining](#)
 June 30, 2017
[Javascript for Zooming and Panning in a Canvas](#)
 September 1, 2012
[Record Your Own Lectures](#)
 September 1, 2011

Tags

[C \(1\)](#) [climate \(1\)](#)

[deep learning](#)

[\(2\)](#) [GPU \(1\)](#)

[graphics \(1\)](#)

[image analysis \(1\)](#)

```
# for pytorch
import os
import torch
import torch.nn as tnn
import torchvision.datasets as dsets
import torchvision.transforms as transforms
from torch.autograd import Variable
```

```
# for reading and plotting results
import matplotlib.pyplot as plt
%matplotlib inline
```

In []:

```
#####;
## Read mnist data into dataTrain
```

```
with gzip.open('mnist.pkl.gz', 'rb') as f:
    train_set, valid_set, test_set = pickle.load(f, encoding='latin1')
    Xtrain = np.array(train_set[0])
    Xval = np.array(valid_set[0])
    Xtest = np.array(test_set[0])
    Ttrain = np.array(train_set[1]).reshape((-1, 1))
    Tval = np.array(valid_set[1]).reshape((-1, 1))
    Ttest = np.array(test_set[1]).reshape((-1, 1))
```

```
# to match with main-gpu.py in res/pytorch...01...
Xtrain = np.vstack((Xtrain, Xval))
Ttrain = np.vstack((Ttrain, Tval))
dataTrain = np.hstack((Xtrain, Ttrain)) # so we can
dataTrain = dataTrain.astype(np.float32)
nSamples = dataTrain.shape[0]
```

In []:

[javascript \(1\)](#)

[matlab](#)

[b \(3\)](#)

[neural network](#)

[s \(2\)](#)

[python \(2\)](#)

[pytorch \(2\)](#)

[reinforcement learning](#)

[\(2\)](#)

[Recent Comments](#)

[Categories](#)

```
#####;
## Write results
```

```
def writeResults(filename, timeAcc, label):
    if filename == 'stdout':
        f = sys.stdout
    else:
        f = open(filename, 'a')
        f.write(label+'\n')
        f.write(str(len(timeAcc)) + '\n')
        for ta in timeAcc:
            f.write('{:.2f} {:.3f}\n'.format(ta[0], ta[1]))
        if filename != 'stdout':
            f.close()
```

In []:

```
def runnumpy(batchSize=None, numEpochs=10,

    label = 'Numpy '
    label += 'Adam' if useAdam else 'SCG'
    label += ' batch {} epochs {:d} hids {} nIter {:d}'.f
    label += ' ReLU ' if useRelu else ' Tanh '
    label += time.strftime('%m/%d/17-%H:%M')

    Xtrain = dataTrain[:, :-1]
    Ttrain = dataTrain[:, -1:]

    nnet = nn.NeuralNetworkClassifier([Xtrain.shape[1],
                                        np.arange(10), useRelu=use
    # NOT STANDARDIZING THE INPUTS!!!
    nnet.setStandardize(False)

    secsAcc = []
```

[atmospheric science](#)
(1)

[deep learning](#)
(2)

[javascript](#)
(1)

[lecture recording](#)
(1)

[neural networks](#)
(2)

[reinforcement learning](#)
(1)

[reinforcement learning](#)
(1)

[Uncategorized](#) (1)

[Archives](#)

[October 2017](#)

```
# numEpochs determines number of breaks du
if batchSize is None:
```

```
    startTime = time.time()
    for i in range(numEpochs):
        Xtrain = dataTrain[:, :-1]
        Ttrain = dataTrain[:, -1:]
        nnet.train(Xtrain, Ttrain, nIterations=nItera
        ptest = nnet.use(Xtest)
        secsAcc.append([time.time() - startTime, nI
```

```
else: # numpyg on batches
```

```
    nSamples = dataTrain.shape[0]
    if nSamples % batchSize != 0:
        print('WARNING: nSamples {} is not divisib
              nSampmles, batchSize))
    nBatches = nSamples // batchSize
    startTime = time.time()
    for epoch in range(numEpochs):
        np.random.shuffle(dataTrain)
        for traini in range(0, nSamples, batchSize):
            Xtrain = dataTrain[traini:traini+batchSize
            Ttrain = dataTrain[traini:traini+batchSize
            nnet.train(Xtrain, Ttrain, restart=True, nI
            ptest = nnet.use(Xtest)
            secsAcc.append([time.time() - startTime, nI
```

```
writeResults(resultsFilename, secsAcc, label)
if numEpochs <= 10:
    writeResults('stdout', secsAcc, label)
```

In []:

```
def runpytorch(batchSize=100, numEpochs=10, h
```

[August 2017](#)

[June 2017](#)

[September 2012](#)

[September 2011](#)

[January 2007](#)

[October 2001](#)

[May 2001](#)

[June 1997](#)

[August 1996](#)

[Meta](#)

[Log in](#)

[Entries RSS](#)

[Comments](#)

[RSS](#)

[WordPress .org](#)

learningRate=0.001, nIterations=100, use

if useGPU:

```
os.environ['CUDA_VISIBLE_DEVICES'] = '0'
print('torch.cuda.is_available() is', torch.cuda
```

if not torch.cuda.is_available():

```
print('GPU is not available. Not running sgc
return
```

label = 'Pytorch '

if useGPU:

```
label += 'GPU '
```

```
label += 'Adam batch {} epochs {:d} lr {:.6f} hids
```

```
label += ' ReLU ' if useRelu else ' Tanh '
```

```
label += time.strftime('%m/%d/17-%H:%M')
```

Neural Network Model (1 hidden layer)

class Net(tnn.Module):

```
def __init__(self, input_size, hidden_size, num
```

```
self.hidden_size = hidden_size
```

```
super(Net, self).__init__()
```

```
self.fc1 = tnn.Linear(input_size, hidden_size)
```

```
self.relu = tnn.ReLU() if useRelu else tnn.Ta
```

```
if len(hidden_size) > 1:
```

```
self.fc2 = tnn.Linear(hidden_size[0], hidd
```

```
self.relu2 = tnn.ReLU() if useRelu else tnn
```

```
self.fc3 = tnn.Linear(hidden_size[1], num
```

```
else:
```

```
self.fc3 = tnn.Linear(hidden_size[0], num
```

```
def forward(self, x):
```

```
out = self.fc1(x)
```

```
out = self.relu(out)
```

```
if len(self.hidden_size) > 1:
```

```
out = self.fc2(out)
```

```
out = self.relu2(out)
```

```
        out = self.fc3(out)
    return out

train_dataset = datasets.MNIST(root='./data',
                                train=True,
                                transform=transforms.ToTensor(),
                                download=True)

test_dataset = datasets.MNIST(root='./data',
                               train=False,
                               transform=transforms.ToTensor())

# Data Loader (Input Pipeline)
train_loader = torch.utils.data.DataLoader(data_loader_args=train_args,
                                             batch_size=batchSize,
                                             shuffle=True)

test_loader = torch.utils.data.DataLoader(data_loader_args=test_args,
                                           batch_size=batchSize,
                                           shuffle=False)

num_classes = 10
net = Net(784, hidden, num_classes)
if useGPU:
    net.cuda()

# Loss and Optimizer
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(net.parameters())

global dataTrain, nSamples

if nSamples % batchSize != 0:
    print('WARNING: nSamples {} is not divisible
          nSamples, batchSize))
nBatches = nSamples // batchSize
```

```
secsAcc = []
startTime = time.time()
for epoch in range(numEpochs):
    np.random.shuffle(dataTrain)
    for i, (images, labels) in enumerate(train_loader):
        # Forward + Backward + Optimize
        if useGPU:
            images = Variable(images.view(-1, 28*28))
            labels = Variable(labels).cuda()
        else:
            images = Variable(images.view(-1, 28*28))
            labels = Variable(labels)
        # Forward + Backward + Optimize
        for iter in range(nIterations):
            optimizer.zero_grad() # zero the gradier
            outputs = net(images)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

correct = 0
total = 0
for images, labels in test_loader:
    if useGPU:
        images = Variable(images.view(-1, 28*28))
    else:
        images = Variable(images.view(-1, 28*28))
    outputs = net(images)
    _, predicted = torch.max(outputs.data, 1)
    total += labels.size(0)
    if useGPU:
        correct += (predicted.cpu() == labels).sum()
    else:
        correct += (predicted == labels).sum()
secsAcc.append([time.time() - startTime, (total - correct) / total])
```



```
writeResults(resultsFilename, secsAcc, label)
if numEpochs <= 10:
    writeResults('stdout', secsAcc, label)
```

In []:

```
resultsFilename = 'test.results'
subprocess.call(['rm', resultsFilename])
hidden = [500, 500] # can contain one or two ints
batchSize = 100
numEpochs = 50
```

In []:

```
runpytorch(batchSize=batchSize, numEpochs=numEpochs,
            learningRate=0.001, nIterations=1, useRelu=True)
```

```
torch.cuda.is_available() is True
```

In []:

```
runpytorch(batchSize=batchSize, numEpochs=numEpochs,
            learningRate=0.001, nIterations=1, useRelu=True)
```

In []:

```
runpytorch(batchSize=batchSize, numEpochs=numEpochs,
            learningRate=0.001, nIterations=1, useRelu=True)
```

In []:

```
runpytorch(batchSize=batchSize, numEpochs=numEpochs,
            learningRate=0.001, nIterations=1, useRelu=useRelu)
```

In []:

```
def plotFromFile(filename='test.results'):
    results = {}
    with open(filename, 'r') as f:
        while True:
            label = f.readline()
            if label is None or label == "":
                break;
            n = int(f.readline())
            secsAcc = []
            for i in range(n):
                secsAcc.append([float(s) for s in f.readline().split()])
            results[label] = secsAcc

    markers = ['s', '8', '>', '^', '<', 'v', 'o', 'X', 'P', 'd', 'h', '*', 'p', 'x']
    mi = 0
    print(sorted(results))
    for key in sorted(results):
        value = results[key]
        value = np.array(value)
        plt.plot(value[:, 0], value[:, 1], '-',
                 marker=markers[mi], label=key, lw=4,
                 markersize=15)
        mi = (mi + 1) % len(markers)
    plt.xlabel('Seconds')
    plt.ylabel('Fraction of test samples incorrectly classified')
    plt.legend();
```

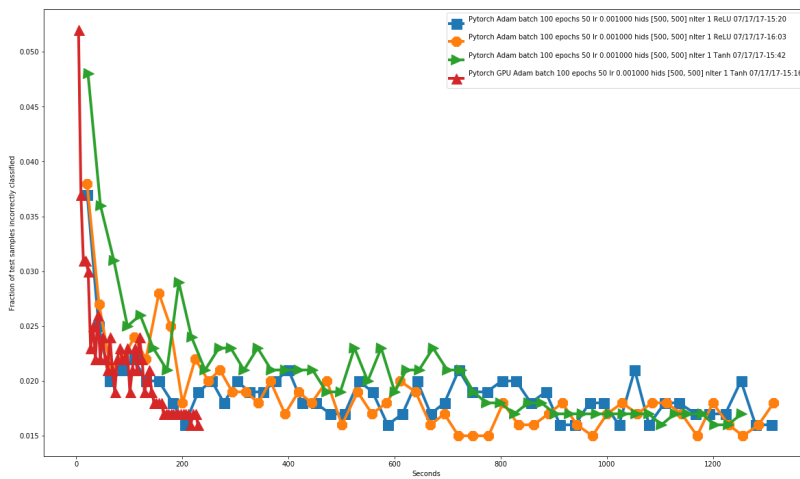
In []:

```
# cat test.results
```

```
In [ ]:
```

```
plt.figure(figsize=(20, 12))
plotFromFile('test.results')
```

```
['Pytorch Adam batch 100 epochs 50 lr 0.001000 |
```



```
In [ ]:
```

```
runnumpy(batchSize=batchSize, numEpochs=nur
          useRelu=False, useAdam=True)
runnumpy(batchSize=batchSize, numEpochs=nur
          useRelu=False, useAdam=False)

runnumpy(batchSize=batchSize, numEpochs=nur
          useRelu=True, useAdam=True)
runnumpy(batchSize=batchSize, numEpochs=nur
          useRelu=True, useAdam=False)
```

In []:

```

runnumpy(batchSize=None, numEpochs=50, hids=[500, 500],
          useRelu=False, useAdam=True)
runnumpy(batchSize=None, numEpochs=50, hids=[500, 500],
          useRelu=False, useAdam=False)

runnumpy(batchSize=None, numEpochs=50, hids=[500, 500],
          useRelu=True, useAdam=True)
runnumpy(batchSize=None, numEpochs=50, hids=[500, 500],
          useRelu=True, useAdam=False)

```

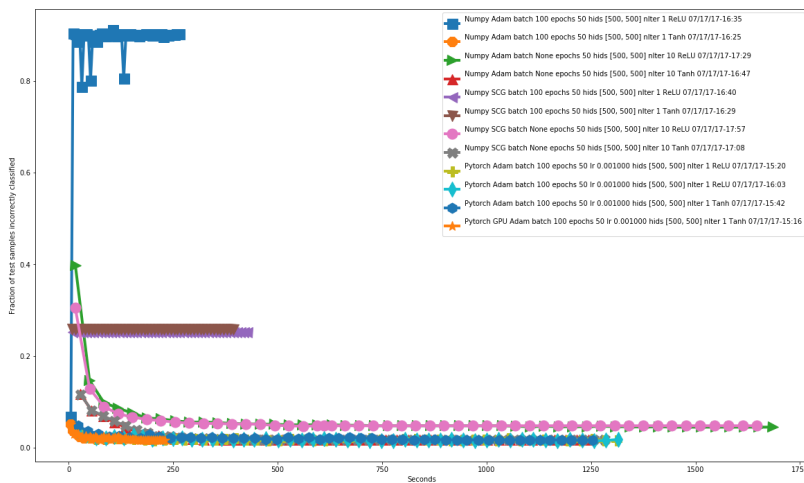
In []:

```

plt.figure(figsize=(20, 12))
plotFromFile('test.results')

```

['Numpy Adam batch 100 epochs 50 hids [500, 500] niter 1 ReLU 07/17/17-16:35



In []:

```

plt.figure(figsize=(20, 12))

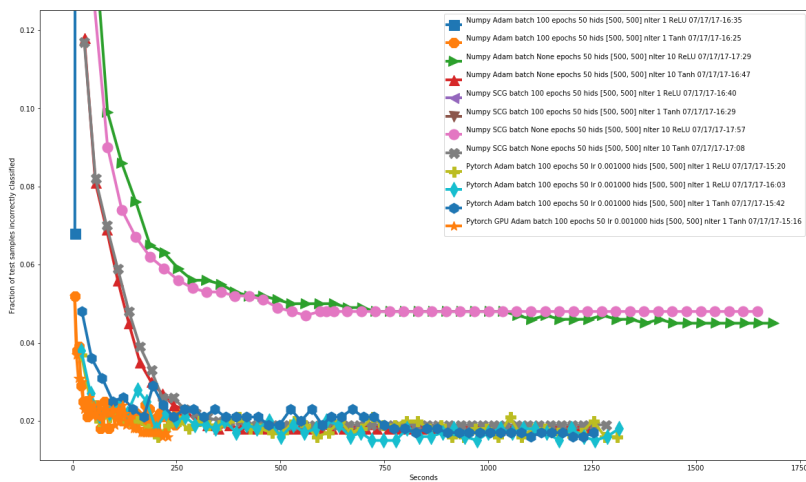
```

```
plotFromFile('test.results')
plt.ylim(0.01,0.125)
# plt.xlim(0,40)
```

```
['Numpy Adam batch 100 epochs 50 hids [500, 50
```

Out[]:

```
(0.01, 0.125)
```



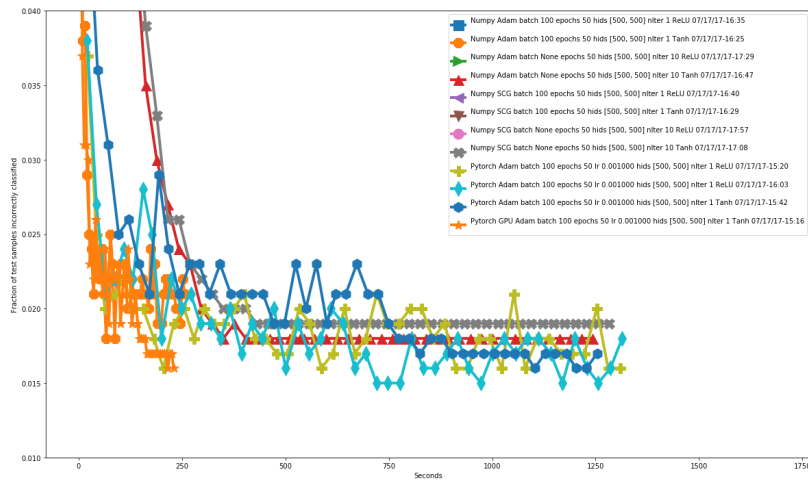
In []:

```
plt.figure(figsize=(20, 12))
plotFromFile('test.results')
plt.ylim(0.01,0.04)
# plt.xlim(0,40)
```

```
['Numpy Adam batch 100 epochs 50 hids [500, 50
```

Out[]:

(0.01, 0.04)



In []:

- deep learning
- 🔍 deep learning, python, pytorch
- < Fast Reinforcement Learning After Pretraining
- > Comparing Numpy, Pytorch, and autograd on CPU and GPU

Leave a Comment

Name *

Email *

Website

Post Comment

© 2018 • GeneratePress