

Cross-entropy method

From Wikipedia, the free encyclopedia

The cross-entropy (CE) method attributed to Reuven Rubinstein is a general Monte Carlo approach to combinatorial and continuous multi-extremal optimization and importance sampling. The method originated from the field of *rare event simulation*, where very small probabilities need to be accurately estimated, for example in network reliability analysis, queueing models, or performance analysis of telecommunication systems. The CE method can be applied to static and noisy combinatorial optimization problems such as the traveling salesman problem, the quadratic assignment problem, DNA sequence alignment, the max-cut problem and the buffer allocation problem, as well as continuous global optimization problems with many local extrema.

In a nutshell the CE method consists of two phases:

1. Generate a random data sample (trajectories, vectors, etc.) according to a specified mechanism.
2. Update the parameters of the random mechanism based on the data to produce a "better" sample in the next iteration. This step involves minimizing the *cross-entropy* or Kullback–Leibler divergence.

Contents

- 1 Estimation via importance sampling
- 2 Generic CE algorithm
- 3 Continuous optimization—example
 - 3.1 Pseudo-code
- 4 Related methods
- 5 See also
- 6 References
- 7 External links

Estimation via importance sampling

Consider the general problem of estimating the quantity

$$\ell = \mathbb{E}_{\mathbf{u}}[H(\mathbf{X})] = \int H(\mathbf{x}) f(\mathbf{x}; \mathbf{u}) \, \mathrm{d}\mathbf{x},$$

where *H* is some *performance function* and *f*(**x**; **u**) is a member of some parametric family of distributions. Using importance sampling this quantity can be estimated as

$$\hat{\ell} = \frac{1}{N} \sum_{i=1}^N H(\mathbf{X}_i) \frac{f(\mathbf{X}_i; \mathbf{u})}{g(\mathbf{X}_i)},$$

where **X**₁, . . . , **X**_N is a random sample from *g* . For positive *H*, the theoretically *optimal* importance sampling density (pdf) is given by

$$g^*(\mathbf{x}) = H(\mathbf{x})f(\mathbf{x}; \mathbf{u})/\ell.$$

This, however, depends on the unknown *ℓ*. The CE method aims to approximate the optimal PDF by adaptively selecting members of the parametric family that are closest (in the Kullback–Leibler sense) to the optimal PDF *g*^{*}.

Generic CE algorithm

1. Choose initial parameter vector **v**⁽⁰⁾; set t = 1.
2. Generate a random sample **X**₁, . . . , **X**_N from *f*(·; **v**^(t−1))
3. Solve for **v**^(t), where

$$\mathbf{v}^{(t)} = \operatorname{argmax}_{\mathbf{u}} \frac{1}{N} \sum_{i=1}^N H(\mathbf{X}_i) \frac{f(\mathbf{X}_i; \mathbf{u})}{f(\mathbf{X}_i; \mathbf{v}^{(t-1)})} \log f(\mathbf{X}_i; \mathbf{v}^{(t-1)})$$

4. If convergence is reached then stop; otherwise, increase t by 1 and reiterate from step 2.

In several cases, the solution to step 3 can be found *analytically*. Situations in which this occurs are

- When *f* belongs to the natural exponential family
- When *f* is discrete with finite support
- When *H*(**X**) = I_{**x** ∈ *A*} and *f*(**X**_{*i*}; **u**) = *f*(**X**_{*i*}; **v**^(t−1)), then **v**^(t) corresponds to the maximum likelihood estimator based on

those $\mathbf{X}_k \in A$.

Continuous optimization—example

The same CE algorithm can be used for optimization, rather than estimation. Suppose the problem is to maximize some function $S(\boldsymbol{x})$, for example, $S(\boldsymbol{x}) = \mathrm{e}^{-(\boldsymbol{x}-2)^2} + 0.8\,\mathrm{e}^{-(\boldsymbol{x}+2)^2}$. To apply CE, one considers first the *associated stochastic problem* of estimating $\mathbb{P}_{\boldsymbol{\theta}}(S(\boldsymbol{X}) \geq \gamma)$ for a given *level* γ , and parametric family $\{f(\cdot; \boldsymbol{\theta})\}$, for example the 1-dimensional Gaussian distribution, parameterized by its mean μ_t and variance σ_t^2 (so $\boldsymbol{\theta} = (\mu, \sigma^2)$ here). Hence, for a given γ , the goal is to find $\boldsymbol{\theta}$ so that $D_{\text{KL}}(\mathbf{I}_{\{S(\boldsymbol{x}) \geq \gamma\}} \| f_{\boldsymbol{\theta}})$ is minimized. This is done by solving the sample version (stochastic counterpart) of the KL divergence minimization problem, as in step 3 above. It turns out that parameters that minimize the stochastic counterpart for this choice of target distribution and parametric family are the sample mean and sample variance corresponding to the *elite samples*, which are those samples that have objective function value $\geq \gamma$. The worst of the elite samples is then used as the level parameter for the next iteration. This yields the following randomized algorithm that happens to coincide with the so-called Estimation of Multivariate Normal Algorithm (EMNA), an estimation of distribution algorithm.

Pseudo-code

```
1. mu:=-6; sigma2:=100; t:=0; maxits=100; // Initialize parameters
2. N:=100; Ne:=10; //
3. while t < maxits and sigma2 > epsilon // While maxits not exceeded and not converged
4. X = SampleGaussian(mu,sigma2,N); // Obtain N samples from current sampling distribution
5. S = exp(-(X-2)^2) + 0.8 exp(-(X+2)^2); // Evaluate objective function at sampled points
6. X = sort(X,S); // Sort X by objective function values (in descending order)
7. mu = mean(X(1:Ne)); sigma2=var(X(1:Ne)); // Update parameters of sampling distribution
8. t = t+1; // Increment iteration counter
9. return mu // Return mean of final sampling distribution as solution
```

Related methods

- Simulated annealing
- Genetic algorithms
- Harmony search
- Estimation of distribution algorithm
- Tabu search

See also

- Cross entropy
- Kullback–Leibler divergence
- Randomized algorithm
- Importance sampling

References

- De Boer, P-T., Kroese, D.P, Mannor, S. and Rubinstein, R.Y. (2005). A Tutorial on the Cross-Entropy Method. *Annals of Operations Research*, 134 (1), 19–67.[1] (<http://www.maths.uq.edu.au/~kroese/ps/aortut.pdf>)
- Rubinstein, R.Y. (1997). Optimization of Computer simulation Models with Rare Events, *European Journal of Operations Research*, 99, 89–112.
- Rubinstein, R.Y., Kroese, D.P. (2004). *The Cross-Entropy Method: A Unified Approach to Combinatorial Optimization, Monte-Carlo Simulation, and Machine Learning*, Springer-Verlag, New York.

External links

- CEoptim R package (<https://cran.r-project.org/web/packages/CEoptim/index.html>)

Retrieved from "https://en.wikipedia.org/w/index.php?title=Cross-entropy_method&oldid=786353572"

- This page was last edited on 18 June 2017, at 22:52.
- Text is available under the Creative Commons Attribution-ShareAlike License; additional terms may apply. By using this site, you agree to the Terms of Use and Privacy Policy. Wikipedia® is a registered trademark of the Wikimedia Foundation, Inc., a non-profit organization.