

 openai / gym

## Join GitHub today

[Dismiss](#)








GitHub is home to over 20 million developers working together to host and review code, manage projects, and build software together.











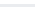
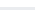


[Sign up](#)

A toolkit for developing and comparing reinforcement learning algorithms. <https://gym.openai.com/read-only.html>

 **751** commits **38** branches **2** releases **85** contributorsBranch: **master** ▼[New pull request](#)[Find file](#)[Clone or download](#) ▼**jhseu** committed with **wojzaremba** Fix Dict space for list arguments (#713)

Latest commit afe0c42 9 days ago

 <a href="#">bin</a>	Docker - Minor build fix (#286)	a year ago
 <a href="#">docs</a>	Fix a typo	6 months ago
 <a href="#">examples</a>	Better error when called on non-discrete action spaces	7 months ago
 <a href="#">gym</a>	Fix Dict space for list arguments (#713)	9 days ago
 <a href="#">misc</a>	TimeLimit refactor with Monitor Simplification (#482)	8 months ago
 <a href="#">scripts</a>	minor fixes to rollout generation	4 months ago
 <a href="#">vendor</a>	Switch to Docker for tests (#285)	a year ago

 <a href="#">.dockerignore</a>	Switch to Docker for tests (#285)	a year ago
 <a href="#">.gitignore</a>	Ignore .cache directory created by py-test	a month ago
 <a href="#">.travis.yml</a>	comment out docker build line (#710)	13 days ago
 <a href="#">CODE_OF_CONDUCT.rst</a>	Initial release. Hello world :).	a year ago
 <a href="#">Dockerfile</a>	Allow double-creating the gym directory	a year ago
 <a href="#">LICENSE.md</a>	Update LICENSE.md	5 months ago
 <a href="#">Makefile</a>	Switch to quay.io/openai/gym:base for tests	a year ago
 <a href="#">README.rst</a>	Make sure env.spec always exists and is valid. (#621)	3 months ago
 <a href="#">requirements.txt</a>	Pin scipy	a year ago
 <a href="#">requirements_dev.txt</a>	switch to pytest (#495)	7 months ago
 <a href="#">setup.py</a>	Pin mujoco-py to <1.0.0	3 months ago
 <a href="#">test.dockerfile</a>	added ffmpeg to dockerfile and fixed test (#325)	a year ago
 <a href="#">tox.ini</a>	Update tox.ini	3 months ago
 <a href="#">unittest.cfg</a>	Capture logs in tests	a year ago

## README.rst

# OpenAI Gym

**OpenAI Gym is a toolkit for developing and comparing reinforcement learning algorithms.** This is the `gym` open-source library, which gives you access to an ever-growing variety of environments.

build passing

[See What's New section below](#)

`gym` makes no assumptions about the structure of your agent, and is compatible with any numerical computation library, such as TensorFlow or Theano. You can use it from Python code, and soon from other languages.

If you're not sure where to start, we recommend beginning with the [docs](#) on our site. See also the [FAQ](#).

A whitepaper for OpenAI Gym is available at <http://arxiv.org/abs/1606.01540>, and here's a BibTeX entry that you can use to cite it in a publication:

```
@misc{1606.01540,  
  Author = {Greg Brockman and Vicki Cheung and Ludwig Pettersson and Jonas Schneider and John Schulma  
  Title = {OpenAI Gym},  
  Year = {2016},  
  Eprint = {arXiv:1606.01540},  
}
```

## Contents of this document

- [Basics](#)
- [Installation](#)
  - [Installing everything](#)
  - [Supported systems](#)
  - [Pip version](#)
  - [Rendering on a server](#)
  - [Installing dependencies for specific environments](#)
- [Environments](#)
  - [Algorithmic](#)
  - [Atari](#)

- [Board games](#)
- [Box2d](#)
- [Classic control](#)
- [MuJoCo](#)
- [Toy text](#)
- [Examples](#)
- [Testing](#)
- [What's new](#)

## Basics

---

There are two basic concepts in reinforcement learning: the environment (namely, the outside world) and the agent (namely, the algorithm you are writing). The agent sends actions to the environment, and the environment replies with observations and rewards (that is, a score).

The core gym interface is [Env](#), which is the unified environment interface. There is no interface for agents; that part is left to you. The following are the `Env` methods you should know:

- `reset(self)`: Reset the environment's state. Returns observation.
- `step(self, action)`: Step the environment by one timestep. Returns observation, reward, done, info.
- `render(self, mode='human', close=False)`: Render one frame of the environment. The default mode will do something human friendly, such as pop up a window. Passing the close flag signals the renderer to close any such windows.

## Installation

---

You can perform a minimal install of `gym` with:

```
git clone https://github.com/openai/gym.git
cd gym
pip install -e .
```

If you prefer, you can do a minimal install of the packaged version directly from PyPI:

```
pip install gym
```

You'll be able to run a few environments right away:

- [algorithmic](#)
- [toy\\_text](#)
- [classic\\_control](#) (you'll need `pyglet` to render though)

We recommend playing with those environments at first, and then later installing the dependencies for the remaining environments.

## Installing everything

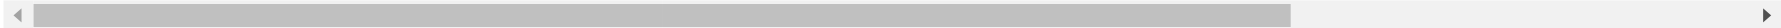
To install the full set of environments, you'll need to have some system packages installed. We'll build out the list here over time; please let us know what you end up installing on your platform.

On OSX:

```
brew install cmake boost boost-python sdl2 swig wget
```

On Ubuntu 14.04:

```
apt-get install -y python-numpy python-dev cmake zlib1g-dev libjpeg-dev xvfb libav-tools xorg-dev python-op
```



MuJoCo has a proprietary dependency we can't set up for you. Follow the [instructions](#) in the `mujoco-py` package for help.

Once you're ready to install everything, run `pip install -e '[all]'` (or `pip install 'gym[all]'`).

## Supported systems

We currently support Linux and OS X running Python 2.7 or 3.5. Some users on OSX + Python3 may need to run

```
brew install boost-python --with-python3
```

If you want to access Gym from languages other than python, we have limited support for non-python frameworks, such as lua/Torch, using the OpenAI Gym [HTTP API](#).

## Pip version

To run `pip install -e '[all]'`, you'll need a semi-recent pip. Please make sure your pip is at least at version `1.5.0`. You can upgrade using the following: `pip install --ignore-installed pip`. Alternatively, you can open [setup.py](#) and install the dependencies by hand.

## Rendering on a server

If you're trying to render video on a server, you'll need to connect a fake display. The easiest way to do this is by running under `xvfb-run` (on Ubuntu, install the `xvfb` package):

```
xvfb-run -s "-screen 0 1400x900x24" bash
```

## Installing dependencies for specific environments

If you'd like to install the dependencies for only specific environments, see [setup.py](#). We maintain the lists of dependencies on a per-environment group basis.

## Environments

---

The code for each environment group is housed in its own subdirectory [gym/envs](#). The specification of each task is in [gym/envs/\\_\\_init\\_\\_.py](#). It's worth browsing through both.

### Algorithmic

These are a variety of algorithmic tasks, such as learning to copy a sequence.

```
import gym
env = gym.make('Copy-v0')
env.reset()
env.render()
```

### Atari

The Atari environments are a variety of Atari video games. If you didn't do the full install, you can install dependencies via `pip install -e '.[atari]'` (you'll need `cmake` installed) and then get started as follow:

```
import gym
env = gym.make('SpaceInvaders-v0')
env.reset()
env.render()
```

This will install `atari-py`, which automatically compiles the [Arcade Learning Environment](#). This can take quite a while (a few minutes on a decent laptop), so just be prepared.

## Board games

The board game environments are a variety of board games. If you didn't do the full install, you can install dependencies via `pip install -e '[board_game]'` (you'll need `cmake` installed) and then get started as follow:

```
import gym
env = gym.make('Go9x9-v0')
env.reset()
env.render()
```

## Box2d

Box2d is a 2D physics engine. You can install it via `pip install -e '[box2d]'` and then get started as follow:

```
import gym
env = gym.make('LunarLander-v2')
env.reset()
env.render()
```

## Classic control

These are a variety of classic control tasks, which would appear in a typical reinforcement learning textbook. If you didn't do the full install, you will need to run `pip install -e '[classic_control]'` to enable rendering. You can get started with them via:

```
import gym
env = gym.make('CartPole-v0')
env.reset()
env.render()
```



## MuJoCo

[MuJoCo](#) is a physics engine which can do very detailed efficient simulations with contacts. It's not open-source, so you'll have to follow the instructions in [mujoco-py](#) to set it up. You'll have to also run `pip install -e '[mujoco]'` if you didn't do the full install.

```
import gym
env = gym.make('Humanoid-v1')
env.reset()
env.render()
```

## Toy text

Toy environments which are text-based. There's no extra dependency to install, so to get started, you can just do:

```
import gym
env = gym.make('FrozenLake-v0')
env.reset()
env.render()
```

## Examples

---

See the `examples` directory.

- Run [examples/agents/random\\_agent.py](#) to run an simple random agent and upload the results to the scoreboard.
- Run [examples/agents/cem.py](#) to run an actual learning agent (using the cross-entropy method) and upload the results to the scoreboard.
- Run [examples/scripts/list\\_envs](#) to generate a list of all environments. (You see also just [browse](#) the list on our site.

- Run [examples/scripts/upload](#) to upload the recorded output from `random_agent.py` or `cem.py`. Make sure to obtain an [API key](#).

## Testing

---

We are using [pytest](#) for tests. You can run them via:

```
pytest
```

## What's new

---

- 2017-06-16: Make `env.spec` into a property to fix a bug that occurs when you try to print out an unregistered Env.
- 2017-05-13: BACKWARDS INCOMPATIBILITY: The Atari environments are now at *v4*. To keep using the old *v3* environments, keep `gym <= 0.8.2` and `atari-py <= 0.0.21`. Note that the *v4* environments will not give identical results to existing *v3* results, although differences are minor. The *v4* environments incorporate the latest Arcade Learning Environment (ALE), including several ROM fixes, and now handle loading and saving of the emulator state. While seeds still ensure determinism, the effect of any given seed is not preserved across this upgrade because the random number generator in ALE has changed. The *\*NoFrameSkip-v4* environments should be considered the canonical Atari environments from now on.
- 2017-03-05: BACKWARDS INCOMPATIBILITY: The `configure` method has been removed from Env. `configure` was not used by gym, but was used by some dependent libraries including `universe`. These libraries will migrate away from the `configure` method by using wrappers instead. This change is on master and will be released with 0.8.0.
- 2016-12-27: BACKWARDS INCOMPATIBILITY: The gym monitor is now a wrapper. Rather than starting monitoring as `env.monitor.start(directory)`, envs are now wrapped as follows: `env = wrappers.Monitor(env, directory)`. This change is on master and will be released with 0.7.0.
- 2016-11-1: Several experimental changes to how a running monitor interacts with environments. The monitor will now raise an error if `reset()` is called when the env has not returned `done=True`. The monitor will only record complete

episodes where `done=True`. Finally, the monitor no longer calls `seed()` on the underlying env, nor does it record or upload seed information.

- 2016-10-31: We're experimentally expanding the environment ID format to include an optional username.
- 2016-09-21: Switch the Gym automated logger setup to configure the root logger rather than just the 'gym' logger.
- 2016-08-17: Calling `close` on an env will also close the monitor and any rendering windows.
- 2016-08-17: The monitor will no longer write manifest files in real-time, unless `write_upon_reset=True` is passed.
- 2016-05-28: For controlled reproducibility, envs now support seeding (cf #91 and #135). The monitor records which seeds are used. We will soon add seed information to the display on the scoreboard.