

 [toddhester](#) / [rl-texplore-ros-pkg](#)

Join GitHub today

[Dismiss](#)





GitHub is home to over 20 million developers working together to host and review code, manage projects, and build software together.

[Sign up](#)

Automatically exported from code.google.com/p/rl-texplore-ros-pkg

 **72** commits **2** branches **0** releases **2** contributorsBranch: **master** ▼[New pull request](#)[Find file](#)[Clone or download](#) ▼**toddhester** Update package.xml

Latest commit 9c927f8 on 6 Sep 2015

 src	Update package.xml	2 years ago
 .catkin_workspace	catkinized version of rl-texplore-ros-pkg	2 years ago
 .gitignore	catkinized version of rl-texplore-ros-pkg	2 years ago
 README	Update Readme	2 years ago

 [README](#)

rl-texplore-ros-pkg

Reinforcement Learning Framework and Repository

Todd Hester
4 May 2011

July 2015 Update

The indigo-devel branch of this repository contains a catkinized version of the original codebase. If you are using ROS Groovy or newer, you should check out this version.

The stacks directory contains a ROS stack for reinforcement learning (RL). The code in this repository provides agents, environments, and multiple ways for them to communicate (through ROS messages, or by including the agent and environment libraries). There are 5 packages in the repository:

rl_common:	Some files that are common to both agents and environments.
rl_msgs:	Definitions of ROS messages for agents and envs to communicate.
rl_agent:	A library of some RL agents including Q-Learning and TEXPLORE.
rl_env:	A library of some RL environments such as Taxi and Fuel World.
rl_experiment:	Code to run some RL experiments without ROS message passing.

Documentation for these packages is here:

http://www.ros.org/wiki/reinforcement_learning

And there is a more thorough tutorial here:

http://www.ros.org/wiki/reinforcement_learning/Tutorials/Reinforcement%20Learning%20Tutorial

Working with the ROS build system is difficult at first. Expect a steep learning curve. Working through some of the ROS tutorials is definitely worthwhile when starting out. Documentation:

<http://www.ros.org/wiki/ROS>.

The interfaces themselves are quite good. So are the visualization and debugging tools. And, there is an impressive amount of useful robotics software already available in the ROS package system.

Directions:

- (1) Check out a copy of the RL code base with ROS interfaces

Our packages have their own Github project. Check out the trunk of rl-texplore-ros-pkg:

```
$ git clone https://github.com/toddhester/rl-texplore-ros-pkg.git
```

This creates a copy of the source tree without commit access.

To commit changes yourself, create a github account and e-mail the project owners requesting commit access.

- (2) Install ROS as in <http://www.ros.org/wiki/ROS/Installation>

This is time-consuming, but not especially difficult. Be careful

to follow the directions exactly. It works fine on Ubuntu and can be made to work for Mac OSX with some effort. Windows will probably not work, but there are future plans to support it.

Add the ROS environment setup to the end of your `.bashrc`:

```
$ echo "source ~/ros/setup.sh" >> ~/.bashrc
```

Append our packages as well as the ROS ones (this example does it from the shell, or you can use any editor):

```
$ cat <<'EOF' >> ~/.bashrc
export ROS_PACKAGE_PATH=${ROS_PACKAGE_PATH}:~/svn/rl-texplore-ros-pkg
EOF
```

(3) Log out and log back in.

(4) Try to run some ROS commands:

```
$ rospack find rl_agent
$ roscd rl_agent
$ rosmake rl_agent
```

This should compile cleanly.

(5) Try to run an RL experiment. First compile the code:

```
$ rosmake rl_agent
$ rosmake rl_env
```

Now we want to run the code to run our experiment. We need to start both an agent and an environment, and they will interact by passing `rl_msg` messages back and forth. Open three tabs in your terminal. Run each of these commands in a different tab:

```
$ roscore
$ rosrun rl_agent agent --agent qlearner
$ rosrun rl_env env --env taxi
```

The second line starts an Q-Learner agent. There are other options to start different agents. The third line starts a taxi environment. There are other options here as well to start different environments. Options for both of these can be seen with the `-h` option. In the agent window, you'll see it print out the sum of rewards at the end of each episode.

- (6) ROS has some nice tools for plotting and recording data. To look at the messages being passed during the experiment, type:

```
$ rostopic list
```

This prints a list of the current messages to the screen. Let's view a particular message:

```
$ rostopic echo /rl_env/rl_state_reward
```

Here we can see the contents of the state reward message being sent by the environment, which contains the current state vector, the reward, and a boolean of whether this is a terminal state or not.

Another thing we can do using ros is make live plots. To make a live plot

of reward per episode, type:

```
$ rxplot /rl_agent/rl_experiment_info/episode_reward
```

One more thing that can be done is to record all the messages being passed in order to view, replay, or plot them later. To record all the messages sent by the agent, type:

```
$ rosbag record /rl_agent/rl_action /rl_agent/rl_experiment_info
```

Now we can play them back in place of an agent, and just have the saved messages control the agent in the environment. Kill the agent, and let's play back the recorded actions:

```
$ rosbag play <bagfile>
```

- (7) Finally, we can also run experiments without passing ROS messages using the `rl_experiment` package. It includes both the agent and env libraries directly, instead of doing message passing. First, compile this package:

```
$ rosmake rl_experiment
```

To run an experiment, we pass an agent and an environment on the command line, and we can also pass other options. Once it starts running, it will print out the sum of rewards per episode for episodic tasks, or the per step rewards for non-episodic tasks. Let's try running R-Max with $M = 1$ on deterministic Fuel World:

```
$ rosrn rl_experiment experiment --agent rmax --m 1 --env fuelworld --deterministic
```

