

Branch: master ▾

agents / agents / ppo / algorithm.py

Find file

Copy path

 **vincentvanhoucke** Fix obesrv typo

104a68b 7 days ago

2 contributors



559 lines (493 sloc) 23.1 KB

```
1  # Copyright 2017 The TensorFlow Agents Authors.
2  #
3  # Licensed under the Apache License, Version 2.0 (the "License");
4  # you may not use this file except in compliance with the License.
5  # You may obtain a copy of the License at
6  #
7  #     http://www.apache.org/licenses/LICENSE-2.0
8  #
9  # Unless required by applicable law or agreed to in writing, software
10 # distributed under the License is distributed on an "AS IS" BASIS,
11 # WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
12 # See the License for the specific language governing permissions and
13 # limitations under the License.
14
15 """Proximal Policy Optimization algorithm.
16
17 Based on John Schulman's implementation in Python and Theano:
18 https://github.com/joschu/modular_rl/blob/master/modular_rl/ppo.py
19 """
20
21 from __future__ import absolute_import
22 from __future__ import division
```

```
23 from __future__ import print_function
24
25 import collections
26
27 import tensorflow as tf
28
29 from agents.ppo import memory
30 from agents.ppo import normalize
31 from agents.ppo import utility
32
33
34 _NetworkOutput = collections.namedtuple(
35     'NetworkOutput', 'policy, mean, logstd, value, state')
36
37
38 class PPOAlgorithm(object):
39     """A vectorized implementation of the PPO algorithm by John Schulman."""
40
41     def __init__(self, batch_env, step, is_training, should_log, config):
42         """Create an instance of the PPO algorithm.
43
44         Args:
45             batch_env: In-graph batch environment.
46             step: Integer tensor holding the current training step.
47             is_training: Boolean tensor for whether the algorithm should train.
48             should_log: Boolean tensor for whether summaries should be returned.
49             config: Object containing the agent configuration as attributes.
50         """
51         self._batch_env = batch_env
52         self._step = step
53         self._is_training = is_training
54         self._should_log = should_log
55         self._config = config
56         self._observ_filter = normalize.StreamingNormalize(
57             self._batch_env.observ[0], center=True, scale=True, clip=5,
```

```
58         name='normalize_observ')
59     self._reward_filter = normalize.StreamingNormalize(
60         self._batch_env.reward[0], center=False, scale=True, clip=10,
61         name='normalize_reward')
62     # Memory stores tuple of observ, action, mean, logstd, reward.
63     template = (
64         self._batch_env.observ[0], self._batch_env.action[0],
65         self._batch_env.action[0], self._batch_env.action[0],
66         self._batch_env.reward[0])
67     self._memory = memory.EpisodeMemory(
68         template, config.update_every, config.max_length, 'memory')
69     self._memory_index = tf.Variable(0, False)
70     use_gpu = self._config.use_gpu and utility.available_gpus()
71     with tf.device('/gpu:0' if use_gpu else '/cpu:0'):
72         # Create network variables for later calls to reuse.
73         self._network(
74             tf.zeros_like(self._batch_env.observ)[: , None],
75             tf.ones(len(self._batch_env)), reuse=None)
76         cell = self._config.network(self._batch_env.action.shape[1].value)
77         with tf.variable_scope('ppo_temporary'):
78             self._episodes = memory.EpisodeMemory(
79                 template, len(batch_env), config.max_length, 'episodes')
80             self._last_state = utility.create_nested_vars(
81                 cell.zero_state(len(batch_env), tf.float32))
82             self._last_action = tf.Variable(
83                 tf.zeros_like(self._batch_env.action), False, name='last_action')
84             self._last_mean = tf.Variable(
85                 tf.zeros_like(self._batch_env.action), False, name='last_mean')
86             self._last_logstd = tf.Variable(
87                 tf.zeros_like(self._batch_env.action), False, name='last_logstd')
88     self._penalty = tf.Variable(
89         self._config.kl_init_penalty, False, dtype=tf.float32)
90     self._policy_optimizer = self._config.policy_optimizer(
91         self._config.policy_lr, name='policy_optimizer')
92     self._value_optimizer = self._config.value_optimizer(
```

```
93         self._config.value_lr, name='value_optimizer')
94
95     def begin_episode(self, agent_indices):
96         """Reset the recurrent states and stored episode.
97
98         Args:
99             agent_indices: 1D tensor of batch indices for agents starting an episode.
100
101         Returns:
102             Summary tensor.
103         """
104         with tf.name_scope('begin_episode/'):
105             reset_state = utility.reinit_nested_vars(self._last_state, agent_indices)
106             reset_buffer = self._episodes.clear(agent_indices)
107             with tf.control_dependencies([reset_state, reset_buffer]):
108                 return tf.constant('')
109
110     def perform(self, observ):
111         """Compute batch of actions and a summary for a batch of observation.
112
113         Args:
114             observ: Tensor of a batch of observations for all agents.
115
116         Returns:
117             Tuple of action batch tensor and summary tensor.
118         """
119         with tf.name_scope('perform/'):
120             observ = self._observ_filter.transform(observ)
121             network = self._network(
122                 observ[:, None], tf.ones(observ.shape[0]), self._last_state)
123             action = tf.cond(
124                 self._is_training, network.policy.sample, lambda: network.mean)
125             logprob = network.policy.log_prob(action)[:, 0]
126             # pylint: disable=g-long-lambda
127             summary = tf.cond(self._should_log, lambda: tf.summary.merge([
```

```
128         tf.summary.histogram('mean', network.mean[:, 0]),
129         tf.summary.histogram('std', tf.exp(network.logstd[:, 0])),
130         tf.summary.histogram('action', action[:, 0]),
131         tf.summary.histogram('logprob', logprob)], str)
132     # Remember current policy to append to memory in the experience callback.
133     with tf.control_dependencies([
134         utility.assign_nested_vars(self._last_state, network.state),
135         self._last_action.assign(action[:, 0]),
136         self._last_mean.assign(network.mean[:, 0]),
137         self._last_logstd.assign(network.logstd[:, 0])]):
138         return tf.check_numerics(action[:, 0], 'action'), tf.identity(summary)
139
140     def experience(self, observ, action, reward, unused_done, unused_nextob):
141         """Process the transition tuple of the current step.
142
143         When training, add the current transition tuple to the memory and update
144         the streaming statistics for observations and rewards. A summary string is
145         returned if requested at this step.
146
147         Args:
148             observ: Batch tensor of observations.
149             action: Batch tensor of actions.
150             reward: Batch tensor of rewards.
151             unused_done: Batch tensor of done flags.
152             unused_nextob: Batch tensor of successor observations.
153
154         Returns:
155             Summary tensor.
156         """
157         with tf.name_scope('experience/'):
158             return tf.cond(
159                 self._is_training,
160                 lambda: self._define_experience(observ, action, reward), str)
161
162     def _define_experience(self, observ, action, reward):
```

```
163 """Implement the branch of experience() entered during training."""
164 update_filters = tf.summary.merge([
165     self._observ_filter.update(observ),
166     self._reward_filter.update(reward)])
167 with tf.control_dependencies([update_filters]):
168     if self._config.train_on_agent_action:
169         # NOTE: Doesn't seem to change much.
170         action = self._last_action
171         batch = observ, action, self._last_mean, self._last_logstd, reward
172         append = self._episodes.append(batch, tf.range(len(self._batch_env)))
173 with tf.control_dependencies([append]):
174     norm_observ = self._observ_filter.transform(observ)
175     norm_reward = tf.reduce_mean(self._reward_filter.transform(reward))
176     # pylint: disable=g-long-lambda
177     summary = tf.cond(self._should_log, lambda: tf.summary.merge([
178         update_filters,
179         self._observ_filter.summary(),
180         self._reward_filter.summary(),
181         tf.summary.scalar('memory_size', self._memory_index),
182         tf.summary.histogram('normalized_observ', norm_observ),
183         tf.summary.histogram('action', self._last_action),
184         tf.summary.scalar('normalized_reward', norm_reward)]), str)
185     return summary
186
187 def end_episode(self, agent_indices):
188     """Add episodes to the memory and perform update steps if memory is full.
189
190     During training, add the collected episodes of the batch indices that
191     finished their episode to the memory. If the memory is full, train on it,
192     and then clear the memory. A summary string is returned if requested at
193     this step.
194
195     Args:
196         agent_indices: 1D tensor of batch indices for agents starting an episode.
```

```
198     Returns:
199         Summary tensor.
200     """
201     with tf.name_scope('end_episode/'):
202         return tf.cond(
203             self._is_training,
204             lambda: self._define_end_episode(agent_indices), str)
205
206     def _define_end_episode(self, agent_indices):
207         """Implement the branch of end_episode() entered during training."""
208         episodes, length = self._episodes.data(agent_indices)
209         space_left = self._config.update_every - self._memory_index
210         use_episodes = tf.range(tf.minimum(
211             tf.shape(agent_indices)[0], space_left))
212         episodes = [tf.gather(elem, use_episodes) for elem in episodes]
213         append = self._memory.replace(
214             episodes, tf.gather(length, use_episodes),
215             use_episodes + self._memory_index)
216         with tf.control_dependencies([append]):
217             inc_index = self._memory_index.assign_add(tf.shape(use_episodes)[0])
218         with tf.control_dependencies([inc_index]):
219             memory_full = self._memory_index >= self._config.update_every
220             return tf.cond(memory_full, self._training, str)
221
222     def _training(self):
223         """Perform multiple training iterations of both policy and value baseline.
224
225         Training on the episodes collected in the memory. Reset the memory
226         afterwards. Always returns a summary string.
227
228         Returns:
229             Summary tensor.
230         """
231         with tf.name_scope('training'):
232             assert_full = tf.assert_equal(
```

```
233         self._memory_index, self._config.update_every)
234     with tf.control_dependencies([assert_full]):
235         data = self._memory.data()
236         (observ, action, old_mean, old_logstd, reward), length = data
237     with tf.control_dependencies([tf.assert_greater(length, 0)]):
238         length = tf.identity(length)
239         observ = self._observ_filter.transform(observ)
240         reward = self._reward_filter.transform(reward)
241         policy_summary = self._update_policy(
242             observ, action, old_mean, old_logstd, reward, length)
243     with tf.control_dependencies([policy_summary]):
244         value_summary = self._update_value(observ, reward, length)
245     with tf.control_dependencies([value_summary]):
246         penalty_summary = self._adjust_penalty(
247             observ, old_mean, old_logstd, length)
248     with tf.control_dependencies([penalty_summary]):
249         clear_memory = tf.group(
250             self._memory.clear(), self._memory_index.assign(0))
251     with tf.control_dependencies([clear_memory]):
252         weight_summary = utility.variable_summaries(
253             tf.trainable_variables(), self._config.weight_summaries)
254     return tf.summary.merge([
255         policy_summary, value_summary, penalty_summary, weight_summary])
256
257 def _update_value(self, observ, reward, length):
258     """Perform multiple update steps of the value baseline.
259
260     We need to decide for the summary of one iteration, and thus choose the one
261     after half of the iterations.
262
263     Args:
264         observ: Sequences of observations.
265         reward: Sequences of reward.
266         length: Batch of sequence lengths.
267
```



```
268     Returns:
269         Summary tensor.
270     """
271     with tf.name_scope('update_value'):
272         loss, summary = tf.scan(
273             lambda _1, _2: self._update_value_step(observ, reward, length),
274             tf.range(self._config.update_epochs_value),
275             [0., ''], parallel_iterations=1)
276         print_loss = tf.Print(0, [tf.reduce_mean(loss)], 'value loss: ')
277         with tf.control_dependencies([loss, print_loss]):
278             return summary[self._config.update_epochs_value // 2]
279
280     def _update_value_step(self, observ, reward, length):
281         """Compute the current value loss and perform a gradient update step.
282
283         Args:
284             observ: Sequences of observations.
285             reward: Sequences of reward.
286             length: Batch of sequence lengths.
287
288         Returns:
289             Tuple of loss tensor and summary tensor.
290         """
291         loss, summary = self._value_loss(observ, reward, length)
292         gradients, variables = (
293             zip(*self._value_optimizer.compute_gradients(loss)))
294         optimize = self._value_optimizer.apply_gradients(
295             zip(gradients, variables))
296         summary = tf.summary.merge([
297             summary,
298             tf.summary.scalar('gradient_norm', tf.global_norm(gradients)),
299             utility.gradient_summaries(
300                 zip(gradients, variables), dict(value=r'*. *'))])
301         with tf.control_dependencies([optimize]):
302             return [tf.identity(loss), tf.identity(summary)]
```

```
303
304 def _value_loss(self, observ, reward, length):
305     """Compute the loss function for the value baseline.
306
307     The value loss is the difference between empirical and approximated returns
308     over the collected episodes. Returns the loss tensor and a summary strin.
309
310     Args:
311         observ: Sequences of observations.
312         reward: Sequences of reward.
313         length: Batch of sequence lengths.
314
315     Returns:
316         Tuple of loss tensor and summary tensor.
317     """
318     with tf.name_scope('value_loss'):
319         value = self._network(observ, length).value
320         return_ = utility.discounted_return(
321             reward, length, self._config.discount)
322         advantage = return_ - value
323         value_loss = 0.5 * self._mask(advantage ** 2, length)
324         summary = tf.summary.merge([
325             tf.summary.histogram('value_loss', value_loss),
326             tf.summary.scalar('avg_value_loss', tf.reduce_mean(value_loss))])
327         value_loss = tf.reduce_mean(value_loss)
328         return tf.check_numerics(value_loss, 'value_loss'), summary
329
330 def _update_policy(
331     self, observ, action, old_mean, old_logstd, reward, length):
332     """Perform multiple update steps of the policy.
333
334     The advantage is computed once at the beginning and shared across
335     iterations. We need to decide for the summary of one iteration, and thus
336     choose the one after half of the iterations.
337
```

```
338     Args:
339         observ: Sequences of observations.
340         action: Sequences of actions.
341         old_mean: Sequences of action means of the behavioral policy.
342         old_logstd: Sequences of action log std devs of the behavioral policy.
343         reward: Sequences of rewards.
344         length: Batch of sequence lengths.
345
346     Returns:
347         Summary tensor.
348     """
349     with tf.name_scope('update_policy'):
350         return_ = utility.discounted_return(
351             reward, length, self._config.discount)
352         value = self._network(observ, length).value
353         if self._config.gae_lambda:
354             advantage = utility.lambda_return(
355                 reward, value, length, self._config.discount,
356                 self._config.gae_lambda)
357         else:
358             advantage = return_ - value
359         mean, variance = tf.nn.moments(advantage, axes=[0, 1], keep_dims=True)
360         advantage = (advantage - mean) / (tf.sqrt(variance) + 1e-8)
361         advantage = tf.Print(
362             advantage, [tf.reduce_mean(return_), tf.reduce_mean(value)],
363             'return and value: ')
364         advantage = tf.Print(
365             advantage, [tf.reduce_mean(advantage)],
366             'normalized advantage: ')
367         # pylint: disable=g-long-lambda
368         loss, summary = tf.scan(
369             lambda _1, _2: self._update_policy_step(
370                 observ, action, old_mean, old_logstd, advantage, length),
371             tf.range(self._config.update_epochs_policy),
372             [0., ''], parallel_iterations=1)
```

```
373     print_loss = tf.Print(0, [tf.reduce_mean(loss)], 'policy loss: ')
374     with tf.control_dependencies([loss, print_loss]):
375         return summary[self._config.update_epochs_policy // 2]
376
377 def _update_policy_step(
378     self, observ, action, old_mean, old_logstd, advantage, length):
379     """Compute the current policy loss and perform a gradient update step.
380
381     Args:
382         observ: Sequences of observations.
383         action: Sequences of actions.
384         old_mean: Sequences of action means of the behavioral policy.
385         old_logstd: Sequences of action log std devs of the behavioral policy.
386         advantage: Sequences of advantages.
387         length: Batch of sequence lengths.
388
389     Returns:
390         Tuple of loss tensor and summary tensor.
391     """
392     network = self._network(observ, length)
393     loss, summary = self._policy_loss(
394         network.mean, network.logstd, old_mean, old_logstd, action,
395         advantage, length)
396     gradients, variables = (
397         zip(*self._policy_optimizer.compute_gradients(loss)))
398     optimize = self._policy_optimizer.apply_gradients(
399         zip(gradients, variables))
400     summary = tf.summary.merge([
401         summary,
402         tf.summary.scalar('gradient_norm', tf.global_norm(gradients)),
403         utility.gradient_summaries(
404             zip(gradients, variables), dict(policy=r'*. *'))])
405     with tf.control_dependencies([optimize]):
406         return [tf.identity(loss), tf.identity(summary)]
407
```

```
408 def _policy_loss(
409     self, mean, logstd, old_mean, old_logstd, action, advantage, length):
410     """Compute the policy loss composed of multiple components.
411
412     1. The policy gradient loss is importance sampled from the data-collecting
413         policy at the beginning of training.
414     2. The second term is a KL penalty between the policy at the beginning of
415         training and the current policy.
416     3. Additionally, if this KL already changed more than twice the target
417         amount, we activate a strong penalty discouraging further divergence.
418
419     Args:
420         mean: Sequences of action means of the current policy.
421         logstd: Sequences of action log stddevs of the current policy.
422         old_mean: Sequences of action means of the behavioral policy.
423         old_logstd: Sequences of action log stddevs of the behavioral policy.
424         action: Sequences of actions.
425         advantage: Sequences of advantages.
426         length: Batch of sequence lengths.
427
428     Returns:
429         Tuple of loss tensor and summary tensor.
430     """
431     with tf.name_scope('policy_loss'):
432         entropy = utility.diag_normal_entropy(mean, logstd)
433         kl = tf.reduce_mean(self._mask(utility.diag_normal_kl(
434             old_mean, old_logstd, mean, logstd), length), 1)
435         policy_gradient = tf.exp(
436             utility.diag_normal_logpdf(mean, logstd, action) -
437             utility.diag_normal_logpdf(old_mean, old_logstd, action))
438         surrogate_loss = -tf.reduce_mean(self._mask(
439             policy_gradient * tf.stop_gradient(advantage), length), 1)
440         kl_penalty = self._penalty * kl
441         cutoff_threshold = self._config.kl_target * self._config.kl_cutoff_factor
442         cutoff_count = tf.reduce_sum(
```

```
443         tf.cast(kl > cutoff_threshold, tf.int32))
444     with tf.control_dependencies([tf.cond(
445         cutoff_count > 0,
446         lambda: tf.Print(0, [cutoff_count], 'kl cutoff! '), int)]):
447         kl_cutoff = (
448             self._config.kl_cutoff_coef *
449             tf.cast(kl > cutoff_threshold, tf.float32) *
450             (kl - cutoff_threshold) ** 2)
451     policy_loss = surrogate_loss + kl_penalty + kl_cutoff
452     summary = tf.summary.merge([
453         tf.summary.histogram('entropy', entropy),
454         tf.summary.histogram('kl', kl),
455         tf.summary.histogram('surrogate_loss', surrogate_loss),
456         tf.summary.histogram('kl_penalty', kl_penalty),
457         tf.summary.histogram('kl_cutoff', kl_cutoff),
458         tf.summary.histogram('kl_penalty_combined', kl_penalty + kl_cutoff),
459         tf.summary.histogram('policy_loss', policy_loss),
460         tf.summary.scalar('avg_surr_loss', tf.reduce_mean(surrogate_loss)),
461         tf.summary.scalar('avg_kl_penalty', tf.reduce_mean(kl_penalty)),
462         tf.summary.scalar('avg_policy_loss', tf.reduce_mean(policy_loss))]
463     )
464     policy_loss = tf.reduce_mean(policy_loss, 0)
465     return tf.check_numerics(policy_loss, 'policy_loss'), summary
466
467 def _adjust_penalty(self, observ, old_mean, old_logstd, length):
468     """Adjust the KL policy between the behavioral and current policy.
469
470     Compute how much the policy actually changed during the multiple
471     update steps. Adjust the penalty strength for the next training phase if we
472     overshot or undershot the target divergence too much.
473
474     Args:
475         observ: Sequences of observations.
476         old_mean: Sequences of action means of the behavioral policy.
477         old_logstd: Sequences of action log stddevs of the behavioral policy.
478         length: Batch of sequence lengths.
```

```
478
479 Returns:
480     Summary tensor.
481 """
482 with tf.name_scope('adjust_penalty'):
483     network = self._network(observ, length)
484     assert_change = tf.assert_equal(
485         tf.reduce_all(tf.equal(network.mean, old_mean)), False,
486         message='policy should change')
487     print_penalty = tf.Print(0, [self._penalty], 'current penalty: ')
488     with tf.control_dependencies([assert_change, print_penalty]):
489         kl_change = tf.reduce_mean(self._mask(utility.diag_normal_kl(
490             old_mean, old_logstd, network.mean, network.logstd), length))
491         kl_change = tf.Print(kl_change, [kl_change], 'kl change: ')
492         maybe_increase = tf.cond(
493             kl_change > 1.3 * self._config.kl_target,
494             # pylint: disable=g-long-lambda
495             lambda: tf.Print(self._penalty.assign(
496                 self._penalty * 1.5), [0], 'increase penalty '),
497             float)
498         maybe_decrease = tf.cond(
499             kl_change < 0.7 * self._config.kl_target,
500             # pylint: disable=g-long-lambda
501             lambda: tf.Print(self._penalty.assign(
502                 self._penalty / 1.5), [0], 'decrease penalty '),
503             float)
504     with tf.control_dependencies([maybe_increase, maybe_decrease]):
505         return tf.summary.merge([
506             tf.summary.scalar('kl_change', kl_change),
507             tf.summary.scalar('penalty', self._penalty)])
508
509 def _mask(self, tensor, length):
510     """Set padding elements of a batch of sequences to zero.
511
512     Useful to then safely sum along the time dimension.
```

```
513
514     Args:
515         tensor: Tensor of sequences.
516         length: Batch of sequence lengths.
517
518     Returns:
519         Masked sequences.
520     """
521     with tf.name_scope('mask'):
522         range_ = tf.range(tensor.shape[1].value)
523         mask = tf.cast(range_[None, :] < length[:, None], tf.float32)
524         masked = tensor * mask
525         return tf.check_numerics(masked, 'masked')
526
527 def _network(self, observ, length=None, state=None, reuse=True):
528     """Compute the network output for a batched sequence of observations.
529
530     Optionally, the initial state can be specified. The weights should be
531     reused for all calls, except for the first one. Output is a named tuple
532     containing the policy as a TensorFlow distribution, the policy mean and log
533     standard deviation, the approximated state value, and the new recurrent
534     state.
535
536     Args:
537         observ: Sequences of observations.
538         length: Batch of sequence lengths.
539         state: Batch of initial recurrent states.
540         reuse: Python boolean whether to reuse previous variables.
541
542     Returns:
543         NetworkOutput tuple.
544     """
545     with tf.variable_scope('network', reuse=reuse):
546         observ = tf.convert_to_tensor(observ)
547         use_gpu = self._config.use_gpu and utility.available_gpus()
```



```
548     with tf.device('/gpu:0' if use_gpu else '/cpu:0'):
549         observ = tf.check_numerics(observ, 'observ')
550         cell = self._config.network(self._batch_env.action.shape[1].value)
551         (mean, logstd, value), state = tf.nn.dynamic_rnn(
552             cell, observ, length, state, tf.float32, swap_memory=True)
553         mean = tf.check_numerics(mean, 'mean')
554         logstd = tf.check_numerics(logstd, 'logstd')
555         value = tf.check_numerics(value, 'value')
556         policy = tf.contrib.distributions.MultivariateNormalDiag(
557             mean, tf.exp(logstd))
558     return _NetworkOutput(policy, mean, logstd, value, state)
```