

## Join GitHub today


[Dismiss](#)

GitHub is home to over 20 million developers working together to host and review code, manage projects, and build software together.

[Sign up](#)

### Efficient Batched Reinforcement Learning in TensorFlow

[#reinforcement-learning](#) [#tensorflow](#) [#multi-processing](#) [#artificial-intelligence](#) [#python](#) [#vectorized-computation](#) [#control](#)

 5 commits 1 branch 1 release 3 contributors Apache-2.0

Branch: master ▾

[New pull request](#)[Find file](#)[Clone or download ▾](#)**danijar** Update link to the paper.

Latest commit 19e5e55 4 days ago

 <a href="#">agents</a>	Fix typos in network comments. (#2)	5 days ago
 <a href="#">.gitignore</a>	Initial release.	8 days ago
 <a href="#">AUTHORS</a>	Initial release.	8 days ago
 <a href="#">CONTRIBUTING.md</a>	Initial release.	8 days ago
 <a href="#">LICENSE</a>	Initial release.	8 days ago
 <a href="#">README.md</a>	Update link to the paper.	4 days ago

# TensorFlow Agents

This project provides optimized infrastructure for reinforcement learning. It extends the [OpenAI gym interface](#) to multiple parallel environments and allows agents to be implemented in TensorFlow and perform batched computation. As a starting point, we provide BatchPPO, an optimized implementation of [Proximal Policy Optimization](#).



Please cite the [TensorFlow Agents paper](#) if you use code from this project in your research:

```
@article{hafner2017agents,  
  title={TensorFlow Agents: Efficient Batched Reinforcement Learning in TensorFlow},  
  author={Hafner, Danijar and Davidson, James and Vanhoucke, Vincent},  
  journal={arXiv preprint arXiv:1709.02878},  
  year={2017}  
}
```

Dependencies: Python 2/3, TensorFlow 1.3+, Gym, rumamel.yaml

## Instructions

Clone the repository and run the PPO algorithm by typing:

```
python3 -m agents.scripts.train --logdir=/path/to/logdir --config=pendulum
```

The algorithm to use is defined in the configuration and `pendulum` started here uses the included PPO implementation. Check out more pre-defined configurations in `agents/scripts/configs.py`.

If you want to resume a previously started run, add the `--timestamp=<time>` flag to the last command and provide the timestamp in the directory name of your run.

To visualize metrics start TensorBoard from another terminal, then point your browser to `http://localhost:2222`:

```
tensorboard --logdir=/path/to/logdir --port=2222
```

To render videos and gather OpenAI Gym statistics to upload to the scoreboard, type:

```
python3 -m agents.scripts.visualize --logdir=/path/to/logdir/<time>-<config> --outdir=/path/to/outdir/
```

## Modifications

We release this project as a starting point that makes it easy to implement new reinforcement learning ideas. These files are good places to start when modifying the code:

File	Content
<code>scripts/configs.py</code>	Experiment configurations specifying the tasks and algorithms.
<code>scripts/networks.py</code>	Neural network models defined as <a href="#">TensorFlow RNNCells</a> .
<code>scripts/train.py</code>	The executable file containing the training setup.
<code>ppo/algorithm.py</code>	The TensorFlow graph for the PPO algorithm.

To run all unit tests, type:

```
python3 -m unittest discover -p "*_test.py"
```

For further questions, please open an issue on Github.

## Implementation

---

We include a batched interface for OpenAI Gym environments that fully integrates with TensorFlow for efficient algorithm implementations. This is achieved through these core components:

- **`agents.tools.wrappers.ExternalProcess`** is an environment wrapper that constructs an OpenAI Gym environment inside of an external process. Calls to `step()` and `reset()`, as well as attribute access, are forwarded to the process and wait for the result. This allows to run multiple environments in parallel without being restricted by Python's global interpreter lock.
- **`agents.tools.BatchEnv`** extends the OpenAI Gym interface to batches of environments. It combines multiple OpenAI Gym environments, with `step()` accepting a batch of actions and returning a batch of observations, rewards, done flags, and info objects. If the individual environments live in external processes, they will be stepped in parallel.
- **`agents.tools.InGraphBatchEnv`** integrates a batch environment into the TensorFlow graph and makes its `step()` and `reset()` functions accessible as operations. The current batch of observations, last actions, rewards, and done flags is stored in variables and made available as tensors.
- **`agents.tools.simulate()`** fuses the step of an in-graph batch environment and a reinforcement learning algorithm together into a single operation to be called inside the training loop. This reduces the number of session calls and provides a simple way to train future algorithms.

To understand all the code, please make yourself familiar with TensorFlow's control flow operations, especially `tf.cond()`, `tf.scan()`, and `tf.control_dependencies()`.

## Disclaimer

---

This is not an official Google product.