

Please note that GitHub no longer supports old versions of Firefox.
We recommend upgrading to the latest [Safari](#), [Google Chrome](#), or [Firefox](#).

Learn more

Ignore

This repository

Search

Pull requests

Issues

Marketplace

Gist

dennybritz / **reinforcement-learning**

Watch 301

Star 3,734

Fork 1,125

<> Code

Issues 31

Pull requests 2

Projects 0

Wiki

Insights

Branch: master

reinforcement-learning / **PolicyGradient** / **a3c** /

Create new file

Upload files

Find file

History

andytwigg committed on GitHub Update train.py	Latest commit f6945fb Feb 22, 2017
..	
README.md	Fixed the typo Nov 18, 2016
estimator_test.py	Scope summaries by worker Nov 5, 2016
estimators.py	Fix tensorflow deprecations Jan 26, 2017
policy_monitor.py	env.monitor is deprecated Jan 25, 2017
policy_monitor_test.py	Refactoring Nov 4, 2016
train.py	Update train.py Feb 22, 2017
worker.py	fix-value-prediction Dec 3, 2016
worker_test.py	Scope summaries by worker Nov 5, 2016

README.md

Implementation of A3C (Asynchronous Advantage Actor-Critic)

Running

```
./train.py --model_dir /tmp/a3c --env Breakout-v0 --t_max 5 --eval_every 300 --parallelism 8
```

See `./train.py --help` for a full list of options. Then, monitor training progress in Tensorboard:

```
tensorboard --logdir=/tmp/a3c
```

Components

- `train.py` contains the main method to start training.
- `estimators.py` contains the Tensorflow graph definitions for the Policy and Value networks.
- `worker.py` contains code that runs in each worker threads.
- `policy_monitor.py` contains code that evaluates the policy network by running an episode and saving rewards to Tensorboard.