

Please note that GitHub no longer supports old versions of Firefox.
We recommend upgrading to the latest [Safari](#), [Google Chrome](#), or [Firefox](#).

Learn more

Ignore

This repository

Search

Pull requests

Issues

Marketplace

Gist



yukezhu / tensorflow-reinforce

Watch 15

Star 224

Fork 77

Code

Issues 3

Pull requests 1

Projects 0

Wiki

Insights

Implementations of Reinforcement Learning Models in Tensorflow

reinforcement-learning

tensorflow

policy-gradient

deep-reinforcement-learning

deep-q-network

actor-critic

39 commits

1 branch

0 releases

4 contributors

MIT

Branch: master

New pull request

Create new file

Upload files

Find file

Clone or download

yukezhu committed on GitHub Update README.md		Latest commit f3f2321 Apr 14, 2017
rl	update APIs	Apr 14, 2017
.gitignore	Initial commit	May 23, 2016
LICENSE	Initial commit	May 23, 2016
README.md	Update README.md	Apr 14, 2017
run_actor_critic_acrobot.py	upgrade to tensorflow v1.0 APIs	Apr 14, 2017
run_cem_cartpole.py	update APIs	Apr 14, 2017
run_ddpg_mujoco.py	merge API upgrade	Apr 14, 2017
run_dqn_cartpole.py	upgrade to tensorflow v1.0 APIs	Apr 14, 2017
run_ql_cartpole.py	upgrade to tensorflow v1.0 APIs	Apr 14, 2017
run_reinforce_cartpole.py	upgrade to tensorflow v1.0 APIs	Apr 14, 2017

README.md

Tensorflow-Reinforce

A collection of [Tensorflow](#) implementations of reinforcement learning models. Models are evaluated in [OpenAI Gym](#) environments. Any contribution/feedback is more than welcome. **Disclaimer:** These implementations are used for educational purposes only (i.e., to learn deep RL myself). There is no guarantee that the exact models will work on any of your particular RL problems without changes.

Environments

This codebase works in both Python 2.7 and 3.5. The models are implemented in Tensorflow 1.0.

Models

Model	Code	References
Cross-Entropy Method	run_cem_cartpole	Cross-entropy method
Tabular Q Learning	rl/tabular_q_learner	Sutton and Barto, Chapter 8
Deep Q Network	rl/neural_q_learner	Mnih et al.
Double Deep Q Network	rl/neural_q_learner	van Hasselt et al.
REINFORCE Policy Gradient	rl/pg_reinforce	Sutton et al.

Model	Code	References
Actor-critic Policy Gradient	rl/pg_actor_critic	Minh et al.
Deep Deterministic Policy Gradient	rl/pg_ddpg	Lillicrap et al.

License

MIT