

Decision stump

From Wikipedia, the free encyclopedia

A decision stump is a machine learning model consisting of a one-level decision tree.^[1] That is, it is a decision tree with one internal node (the root) which is immediately connected to the terminal nodes (its leaves). A decision stump makes a prediction based on the value of just a single input feature. Sometimes they are also called 1-rules.^[2]

Depending on the type of the input feature, several variations are possible. For nominal features, one may build a stump which contains a leaf for each possible feature value^{[3][4]} or a stump with the two leaves, one of which corresponds to some chosen category, and the other leaf to all the other categories.^[5] For binary features these two schemes are identical. A missing value may be treated as a yet another category.^[5]

For continuous features, usually, some threshold feature value is selected, and the stump contains two leaves — for values below and above the threshold. However, rarely, multiple thresholds may be chosen and the stump therefore contains three or more leaves.

Decision stumps are often^[6] used as components (called "weak learners" or "base learners") in machine learning ensemble techniques such as bagging and boosting. For example, a state-of-the-art Viola–Jones face detection algorithm employs AdaBoost with decision stumps as weak learners.^[7]

The term "decision stump" was coined in a 1992 ICML paper by Wayne Iba and Pat Langley.^{[1][8]}

References

- Iba, Wayne; and Langley, Pat (1992); *Induction of One-Level Decision Trees* (<http://lyonesse.stanford.edu/~langley/papers/stump.ml92.pdf>), in *ML92: Proceedings of the Ninth International Conference on Machine Learning, Aberdeen, Scotland, 1–3 July 1992*, San Francisco, CA: Morgan Kaufmann, pp. 233–240
- Holte, Robert C. (1993). "Very Simple Classification Rules Perform Well on Most Commonly Used Datasets". CiteSeerX 10.1.1.67.2711 (<https://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.67.2711>) .
- Loper, Edward L.; Bird,, Steven; Klein, Ewan (2009). *Natural language processing with Python* (<http://nltk.googlecode.com/svn/trunk/doc/book/ch06.html>). Sebastopol, CA: O'Reilly. ISBN 0-596-51649-5.
- This classifier is implemented in Weka under the name OneR (for "1-rule").
- This is what has been implemented in Weka's DecisionStump classifier.
- Reyzin, Lev; and Schapire, Robert E. (2006); *How Boosting the Margin Can Also Boost Classifier Complexity* (http://www.cs.princeton.edu/~schapire/papers/boost_complexity.pdf), in *ICML’06: Proceedings of the 23rd international conference on Machine Learning*, pp. 753-760
- Viola, Paul; and Jones, Michael J. (2004); *Robust Real-Time Face Detection* (<http://www.face-rec.org/algorithms/Boosting-Ensemble/16981346.pdf>), *International Journal of Computer Vision*, 57(2), 137–154
- Oliver, Jonathan J.; and Hand, David (1994); *Averaging Over Decision Stumps*, in *Machine Learning: ECML-94, European Conference on Machine Learning, Catania, Italy, April 6–8, 1994, Proceedings*, Lecture Notes in Computer Science (LNCS) 784, Springer, pp. 231–241 ISBN 3-540-57868-4 doi:10.1007/3-540-57868-4_61 (https://dx.doi.org/10.1007%2F3-540-57868-4_61)
Quote: "These simple rules are in effect severely pruned decision trees and have been termed *decision stumps* [cites Iba and Langley]".

Retrieved from "https://en.wikipedia.org/w/index.php?title=Decision_stump&oldid=784065698"

Categories: Decision trees

- This page was last edited on 6 June 2017, at 08:14.
- Text is available under the Creative Commons Attribution-ShareAlike License; additional terms may apply. By using this site, you agree to the Terms of Use and Privacy Policy. Wikipedia® is a registered trademark of the Wikimedia Foundation, Inc., a non-profit organization.

