Gentlest Intro to Tensorflow (Part 3)

Khor Soon Hin, @neth_6, re:Culture

In collaboration with Sam & Edmund

Overview

- Multi-feature Linear Regression
- Logistic Regression
  - Multi-class prediction
  - Cross-entropy
  - Softmax
- Tensorflow Cheatsheet #1

Review: Predict from Single Feature with Linear Regression

Quick Review: Predict from Single Feature (House Size)

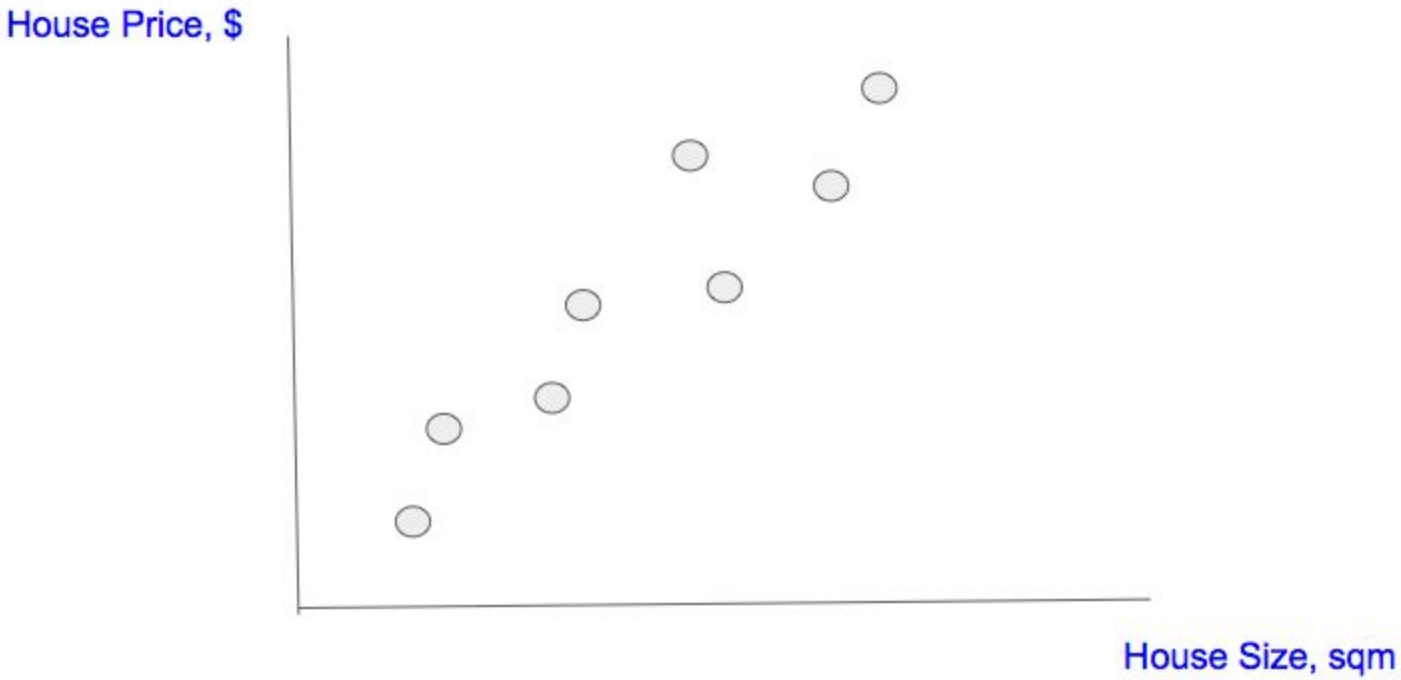

Quick Review: Use Linear Regression

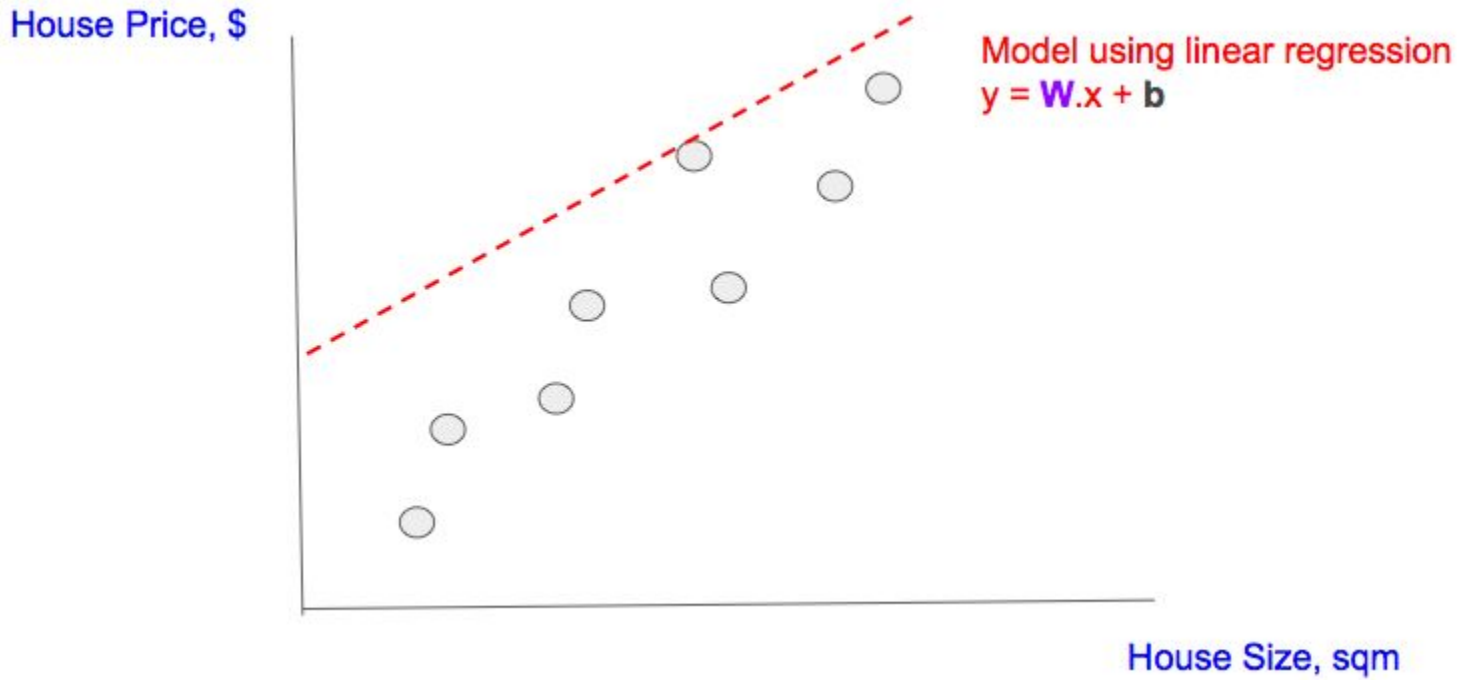

Quick Review: Predict using Linear Regression

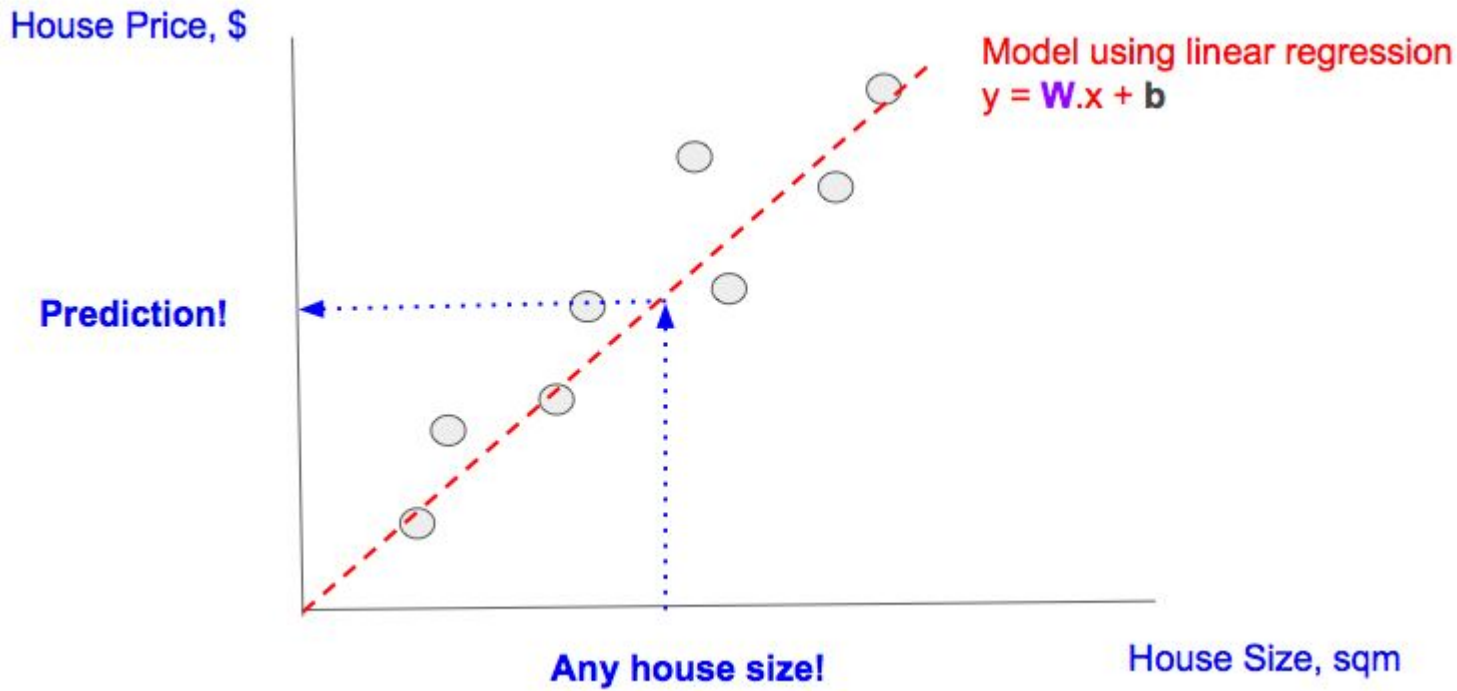

Linear Regression: Predict from Two (or More) Features

Two Features: House Size, Rooms

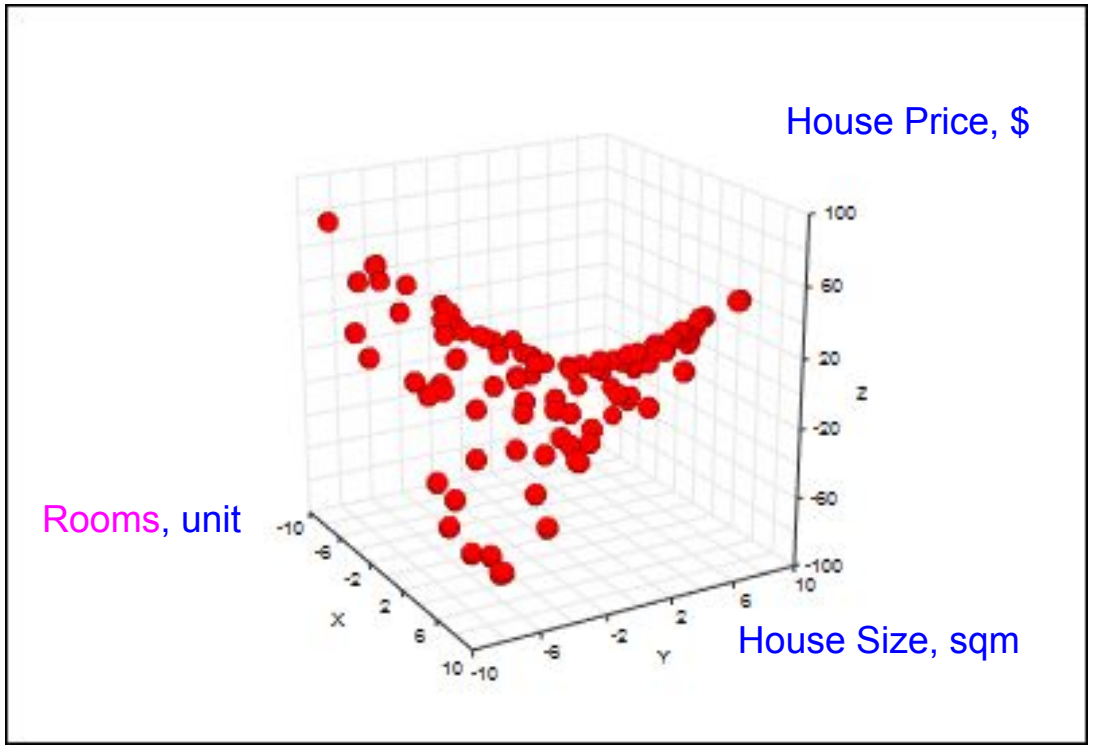

Source: teraplot.com

Same Issue: Predict for Values without Datapoint

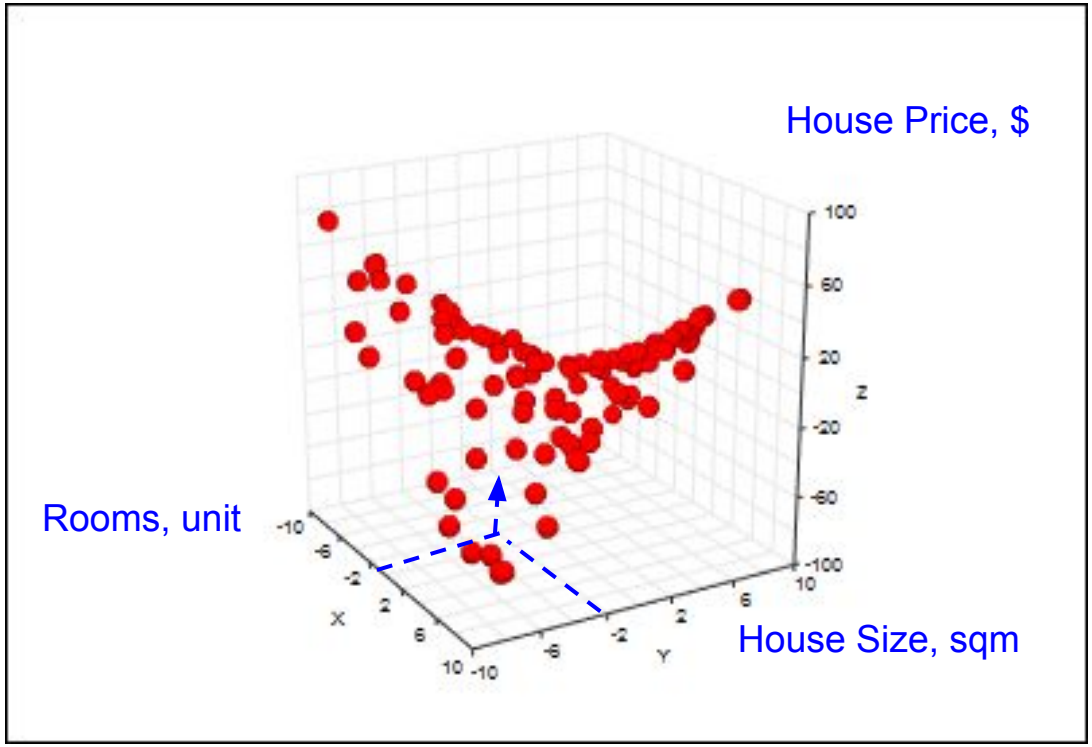

Source: teraplot.com

Same Solution: Find Best-Fit

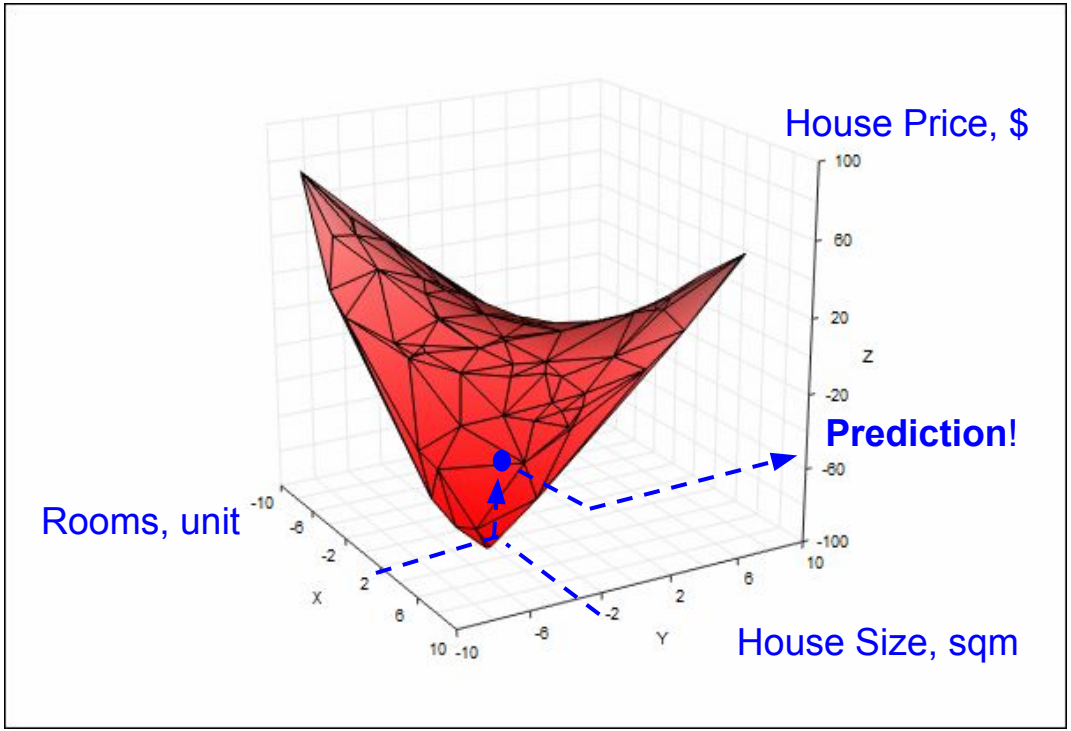

Source: teraplot.com

Review: Tensorflow Code

Tensorflow Code

```
# Model linear regression y = Wx + b
x = tf.placeholder(tf.float32, [None, 1])
W = tf.Variable(tf.zeros([1,1]))
b = tf.Variable(tf.zeros([1]))
product = tf.matmul(x,W)
y = product + b
y_ = tf.placeholder(tf.float32, [None, 1])

# Cost function 1/n * sum((y_-y)**2)
cost = tf.reduce_mean(tf.square(y_-y))

# Training using Gradient Descent to minimize cost
train_step = tf.train.GradientDescentOptimizer(0.0000001).minimize(cost)
```

Multi-feature: Change in Model & Cost Function

Model

1 Feature

$y = W.x + b$

y: House price prediction

x: House size

Goal: Find scalars **W,b**

Model

1 Feature

y = W.x + b

y: House price prediction

x: House size

Goal: Find scalars **W,b**

2 Features

y = W.x + W2.x2 + b

y: House price prediction

x: House size

x2: Rooms

Goal: Find scalars **W, W2, b**

Tensorflow Graph

1 Feature

```
y = tf.matmul(x, W) + b

W = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1])
```

Tensorflow Graph

1 Feature

```
y = tf.matmul(x, W) + b

W = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1])
```

2 Features

```
y = matmul(x, W) + matmul(x2, W2) + b

W = tf.Variable(tf.zeros[1,1])
W2 = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 1])
x2 = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1])
```

Tensorflow Graph: Train

1 Feature

```
y = tf.matmul(x, W) + b

W = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1])

Train: feed = { x: … , y_: … }
```

2 Features

```
y = matmul(x, W) + matmul(x2, W2) + b

W = tf.Variable(tf.zeros[1,1])
W2 = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 1])
x2 = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1])

Train: feed = { x: … , x2: ... y_: … }
```

Tensorflow Graph: Scalability Issue

Model gets messy!

1 Feature

```
y = tf.matmul(x, W) + b

W = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])
x = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1])
```

Train: feed = { x: … , y_: … }

2 Features

```
y = tf.matmul(x, W) + tf.matmul(x2, W2) + b

W = tf.Variable(tf.zeros[1,1])
W2 = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])
x = tf.placeholder(tf.float, [None, 1])
x2 = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1]
```

Train: feed = { x: … , x2: ... y_: … }

3 Features

```
y = tf.matmul(x, W) + tf.matmul(x2, W2) + tf.matmul(x3, W3) +b

W = tf.Variable(tf.zeros[1,1])
W2 = tf.Variable(tf.zeros[1,1])
W3 = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])
x = tf.placeholder(tf.float, [None, 1])
x2 = tf.placeholder(tf.float, [None, 1])
x3 = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1]
```

Train: feed = { x: … , x2: ... , x3: … , y_: … }

Data Representation

House #1	Size_1	Rooms_1	Price_1
House #2	Size_2	Rooms_2	Price_2
…	…	…	...
House #m	Size_m	Rooms_m	Price_m

Lots of Data Manipulation

House #1	Size_1	Rooms_1	Price_1
House #2	Size_2	Rooms_2	Price_2
…	…	…	...
House #m	Size_m	Rooms_m	Price_m

2 Features

```
y = tf.matmul(x, W) + tf.matmul(x2, W2) + b

W = tf.Variable(tf.zeros[1,1])
W2 = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 1])
x2 = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1]
```

Train: feed = { x: … , x2: ... y_: … }

Lots of Data Manipulation 2

House #1	Size_1	Rooms_1	Price_1
House #2	Size_2	Rooms_2	Price_2
…	…	…	...
House #m	Size_m	Rooms_m	Price_m

2 Features

```
y = tf.matmul(x, W) + tf.matmul(x2, W2) + b

W = tf.Variable(tf.zeros[1,1])
W2 = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 1])
x2 = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1]
```

Train: feed = { x: … , x2: ... y_: … }

Lots of Data Manipulation 3

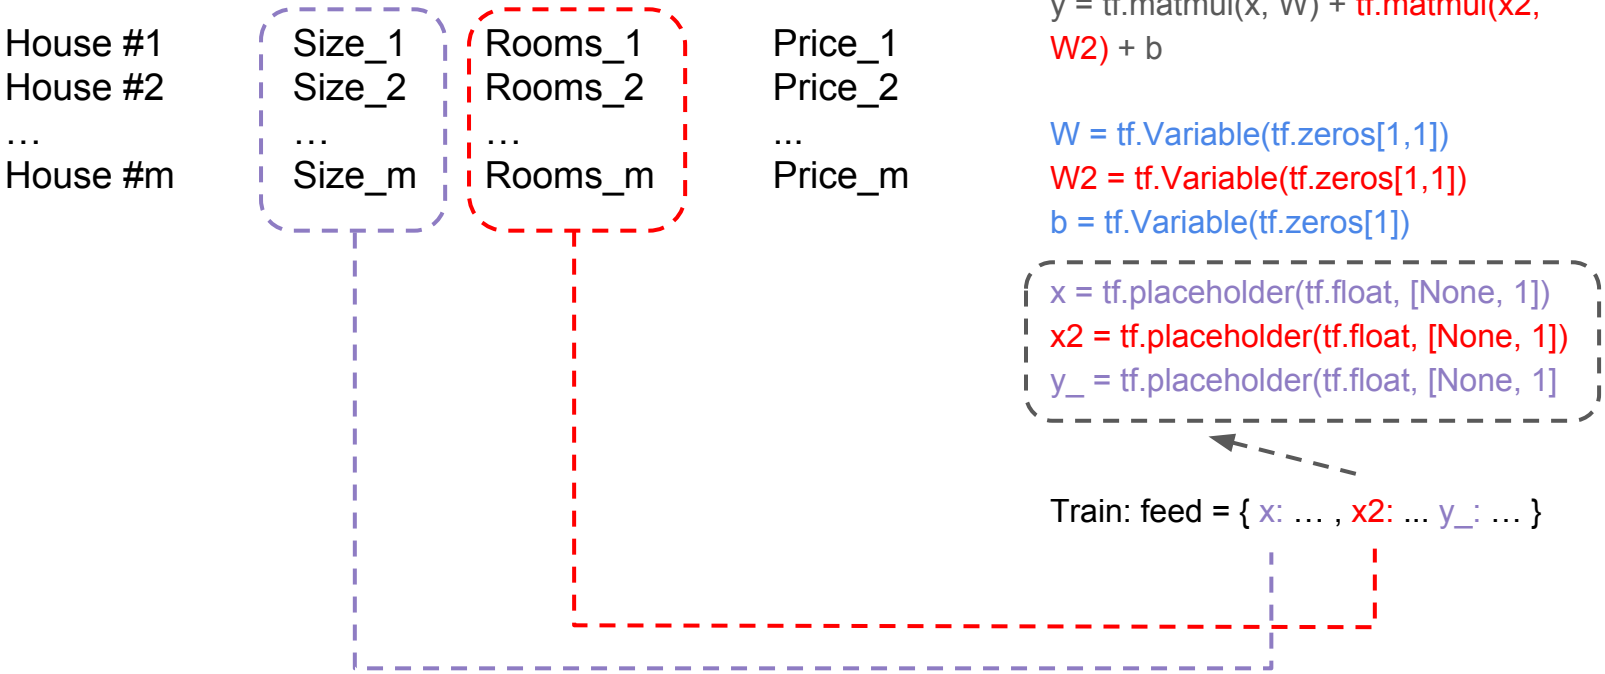

Lots of Data Manipulation 4

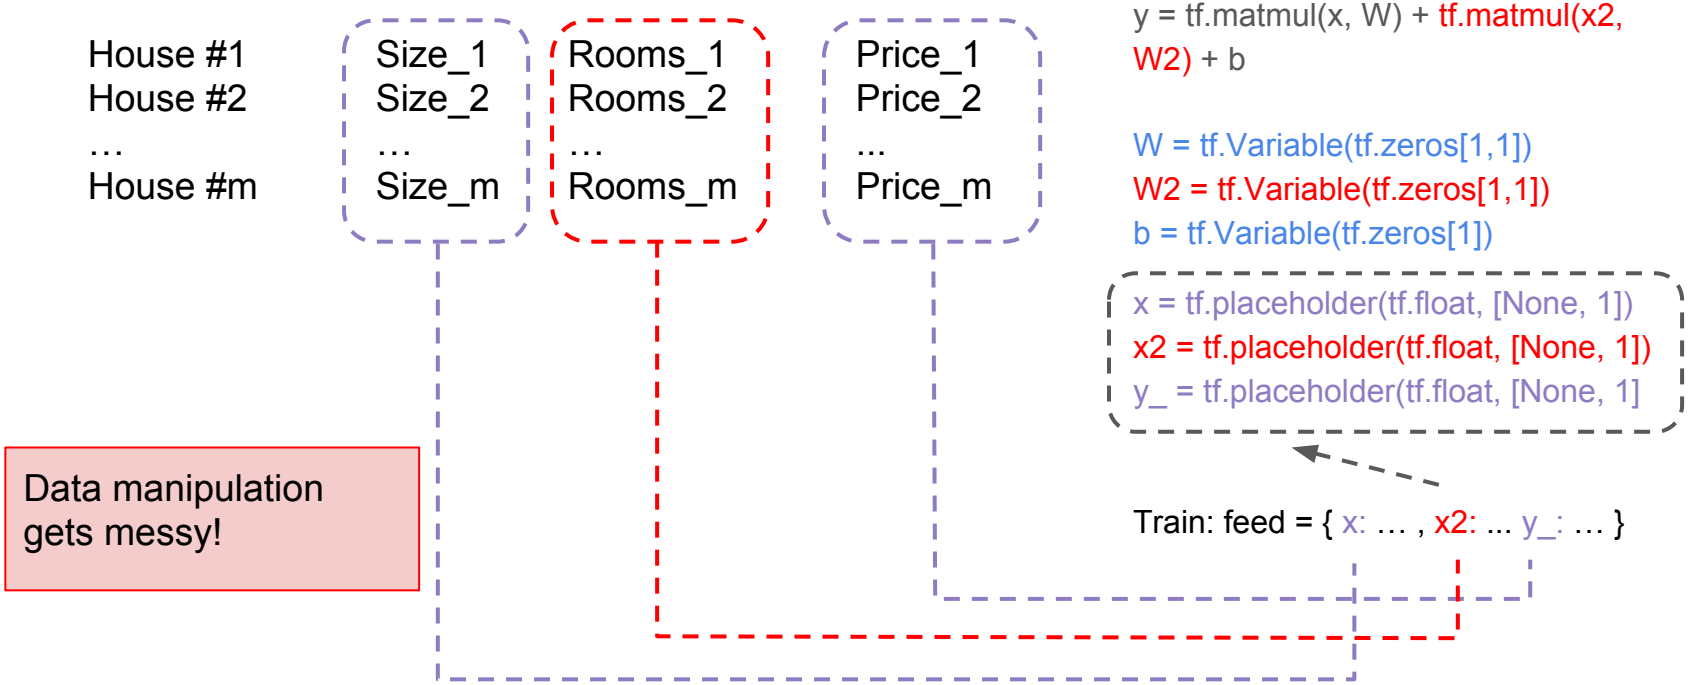

Lots of Data Manipulation 4

Find better way

$y = W.x + W2.x2 + b$

2 Features

```
y = tf.matmul(x, W) + tf.matmul(x2,
W2) + b
```

```
W = tf.Variable(tf.zeros[1,1])
W2 = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])
```

```
x = tf.placeholder(tf.float, [None, 1])
x2 = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1]
```

Train: feed = { x: … , x2: ... y_: … }

	Size_1	Rooms_1	Price_1
House #1	Size_2	Rooms_2	Price_2
…	…	…	...
House #m	Size_m	Rooms_m	Price_m

Matrix: Cleaning Up Representations

Matrix Representation

House #0	Size_0	Rooms_0	Price_0
House #1	Size_1	Rooms_1	Price_1
…	…	…	...
House #m	Size_m	Rooms_m	Price_m

Matrix Representation

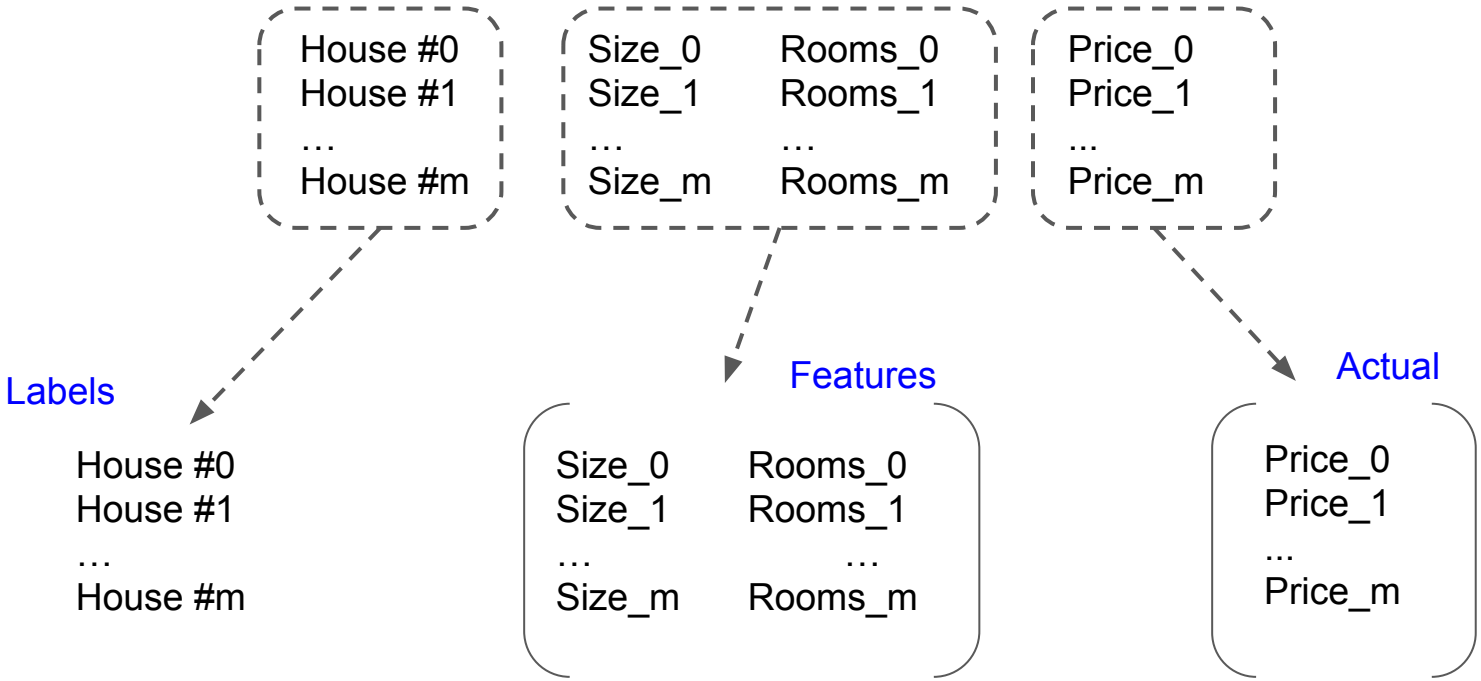

Matrix Representation

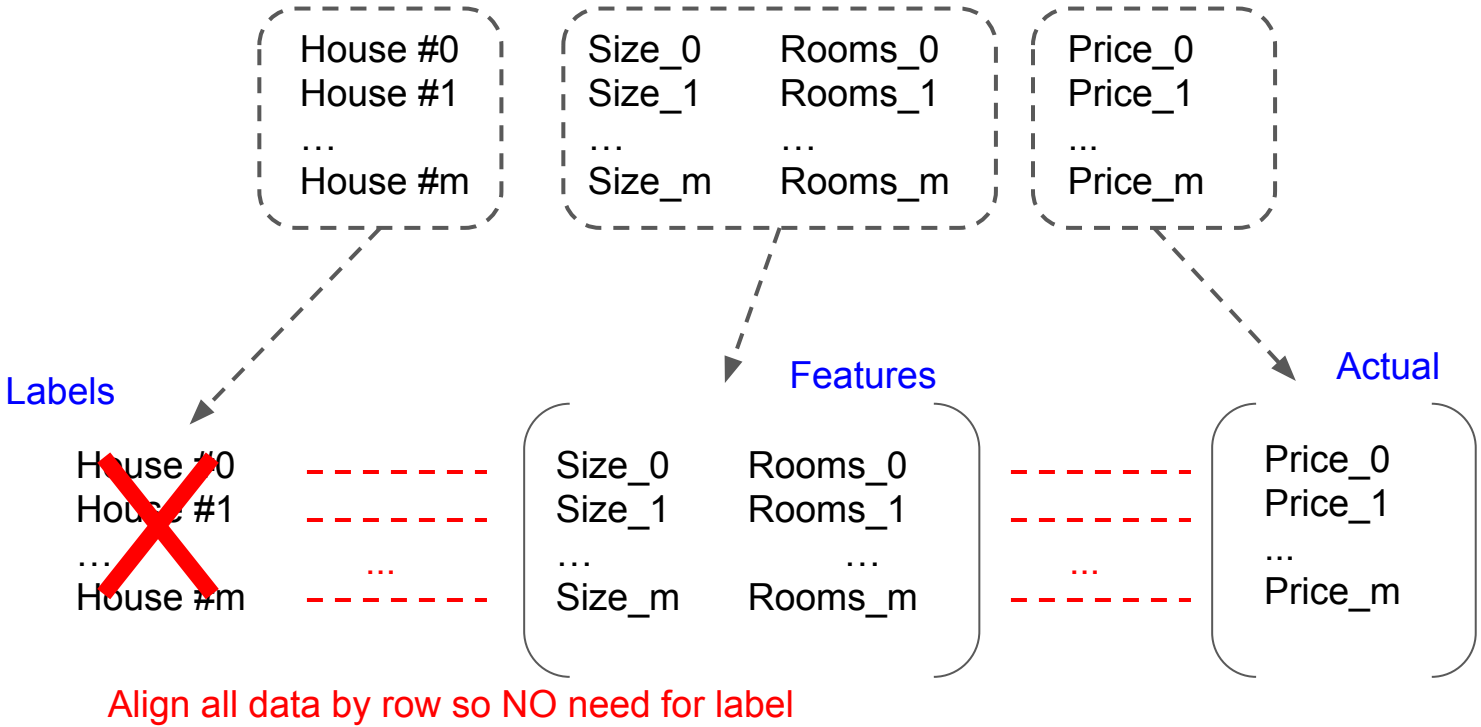

Matrix Representation

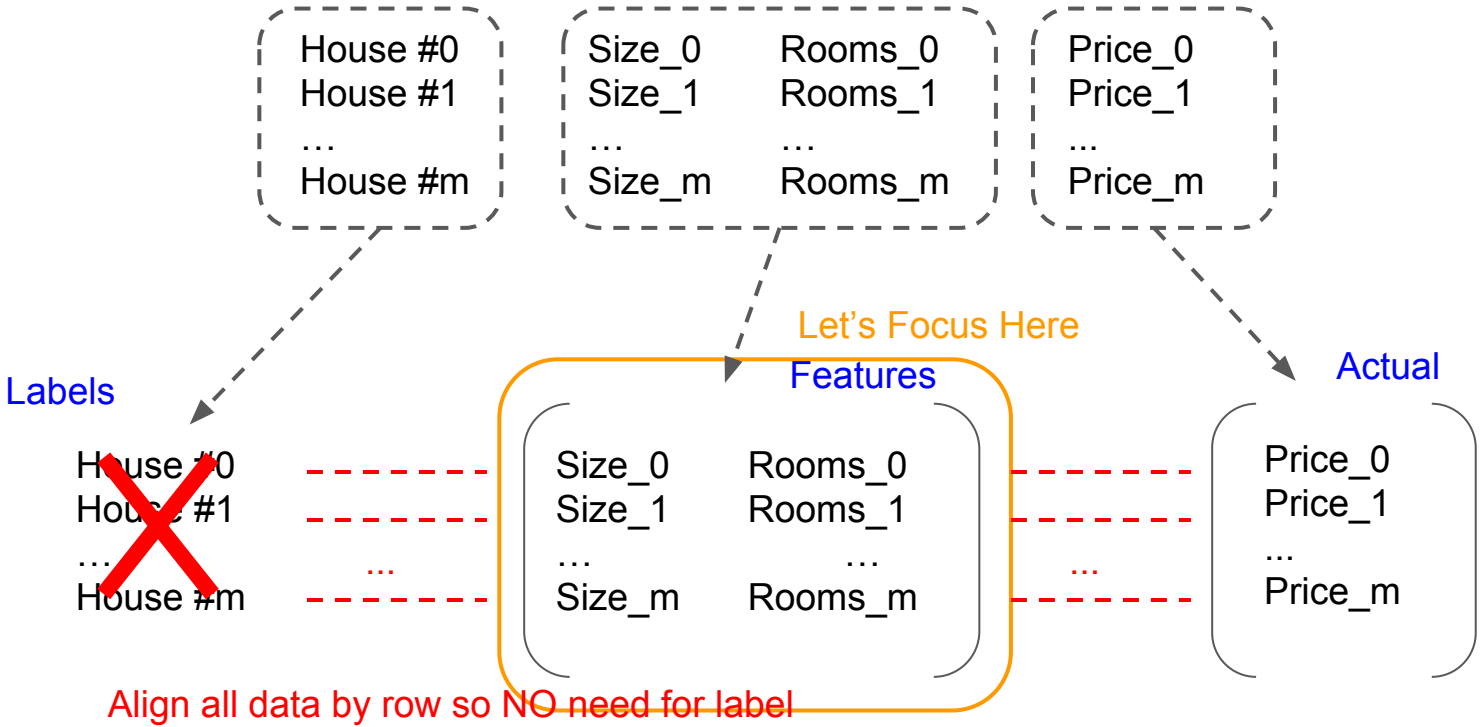

Matrix: Cleaning Up Models

Better Model Equation

Find better way

$$y = W.x + W2.x2 + b$$

2 Features

```
y = tf.matmul(x, W) + tf.matmul(x2,
W2) + b
```

```
W = tf.Variable(tf.zeros[1,1])
W2 = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])
```

```
x = tf.placeholder(tf.float, [None, 1])
x2 = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1]
```

Train: feed = { x: … , x2: … , y_: …
}

Better Model Equation

Find better way

$$y = W.x + W2.x2 + b$$

2 Features

```
y = tf.matmul(x, W) + tf.matmul(x2,
W2) + b
```

```
W = tf.Variable(tf.zeros[1,1])
W2 = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])
```

```
x = tf.placeholder(tf.float, [None, 1])
x2 = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1]
```

Train: feed = { x: … , x2: … , y_: … }

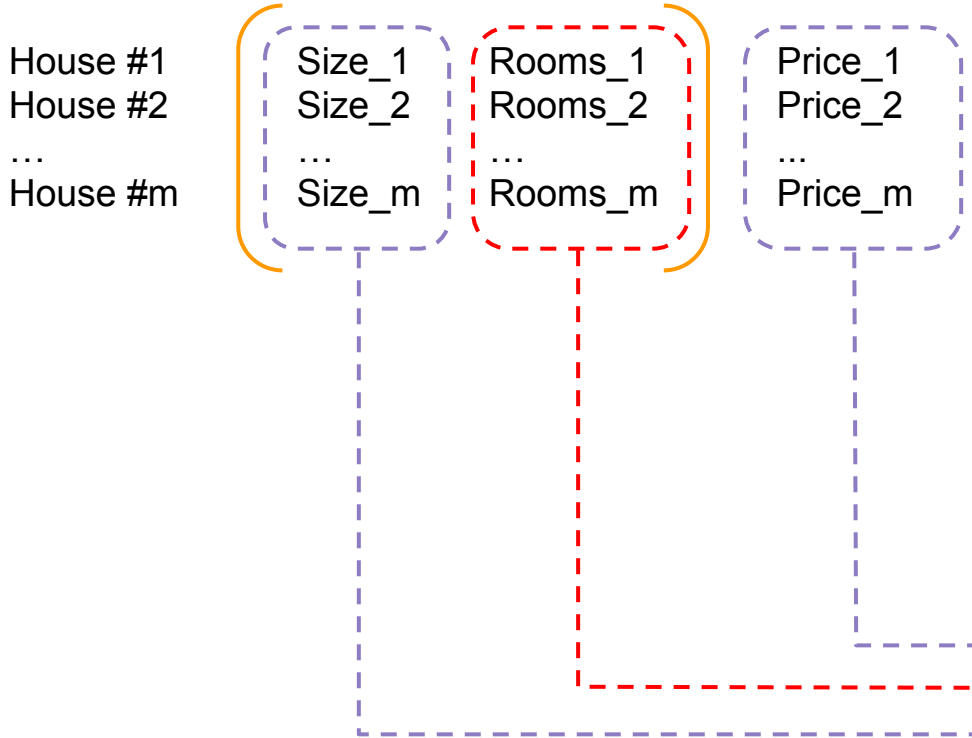

Better Model Equation

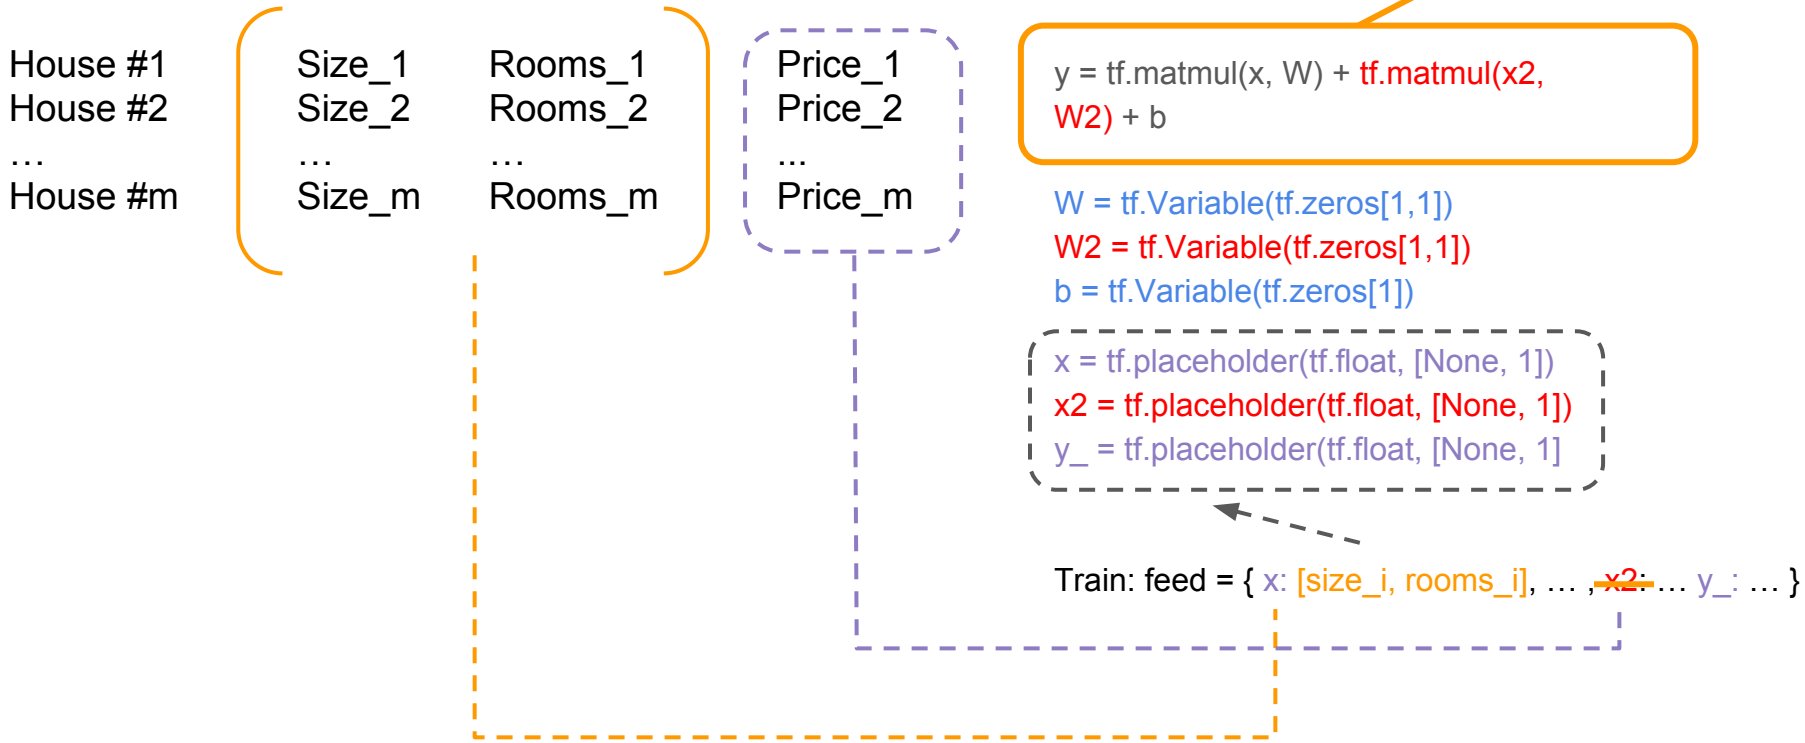

Better Model Equation

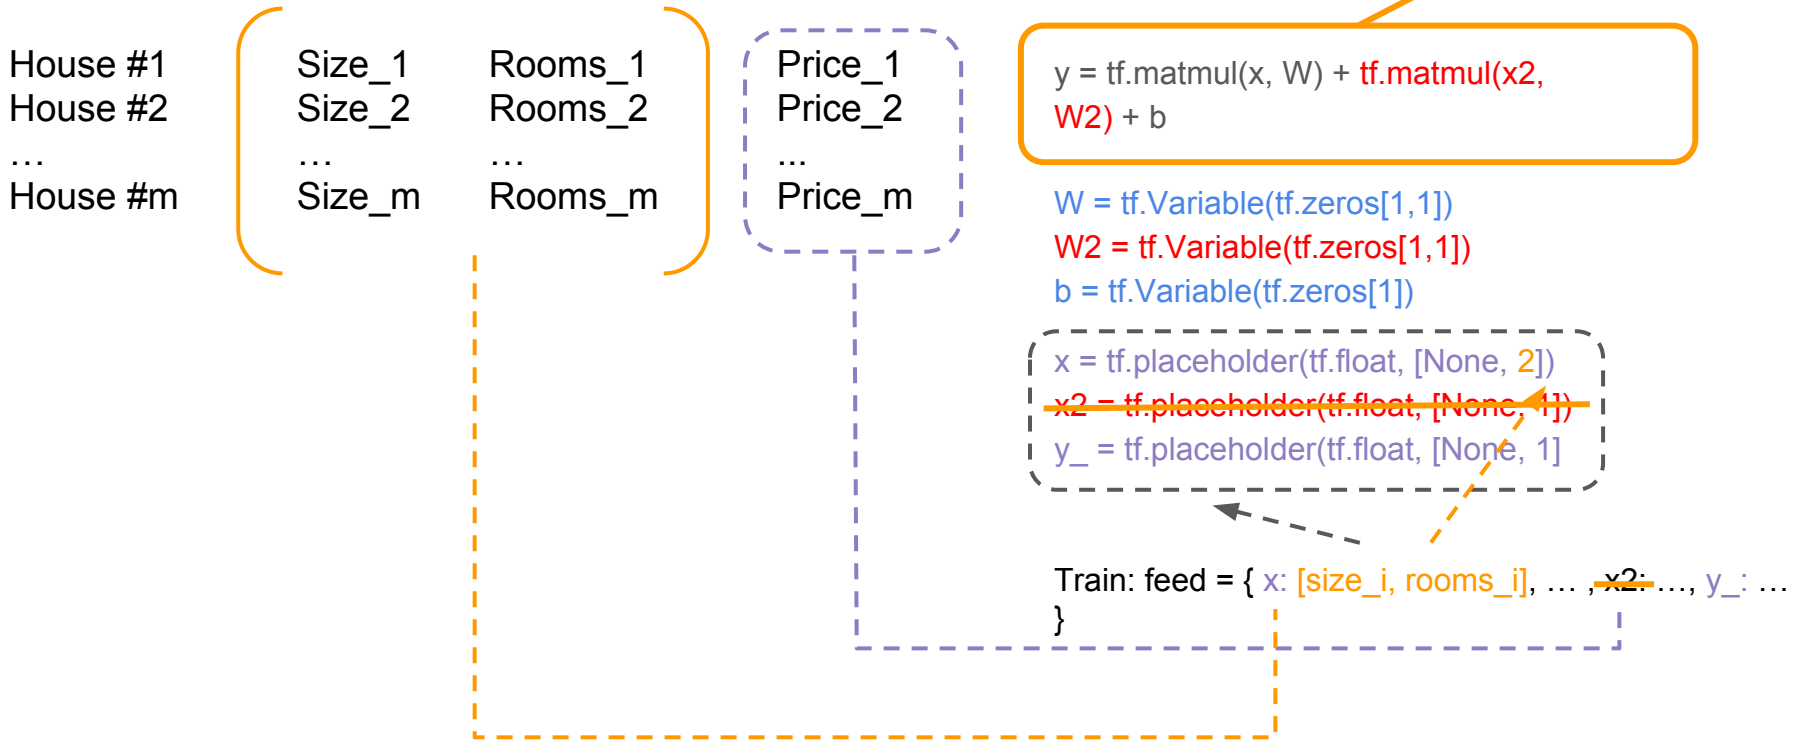

Better Model Equation

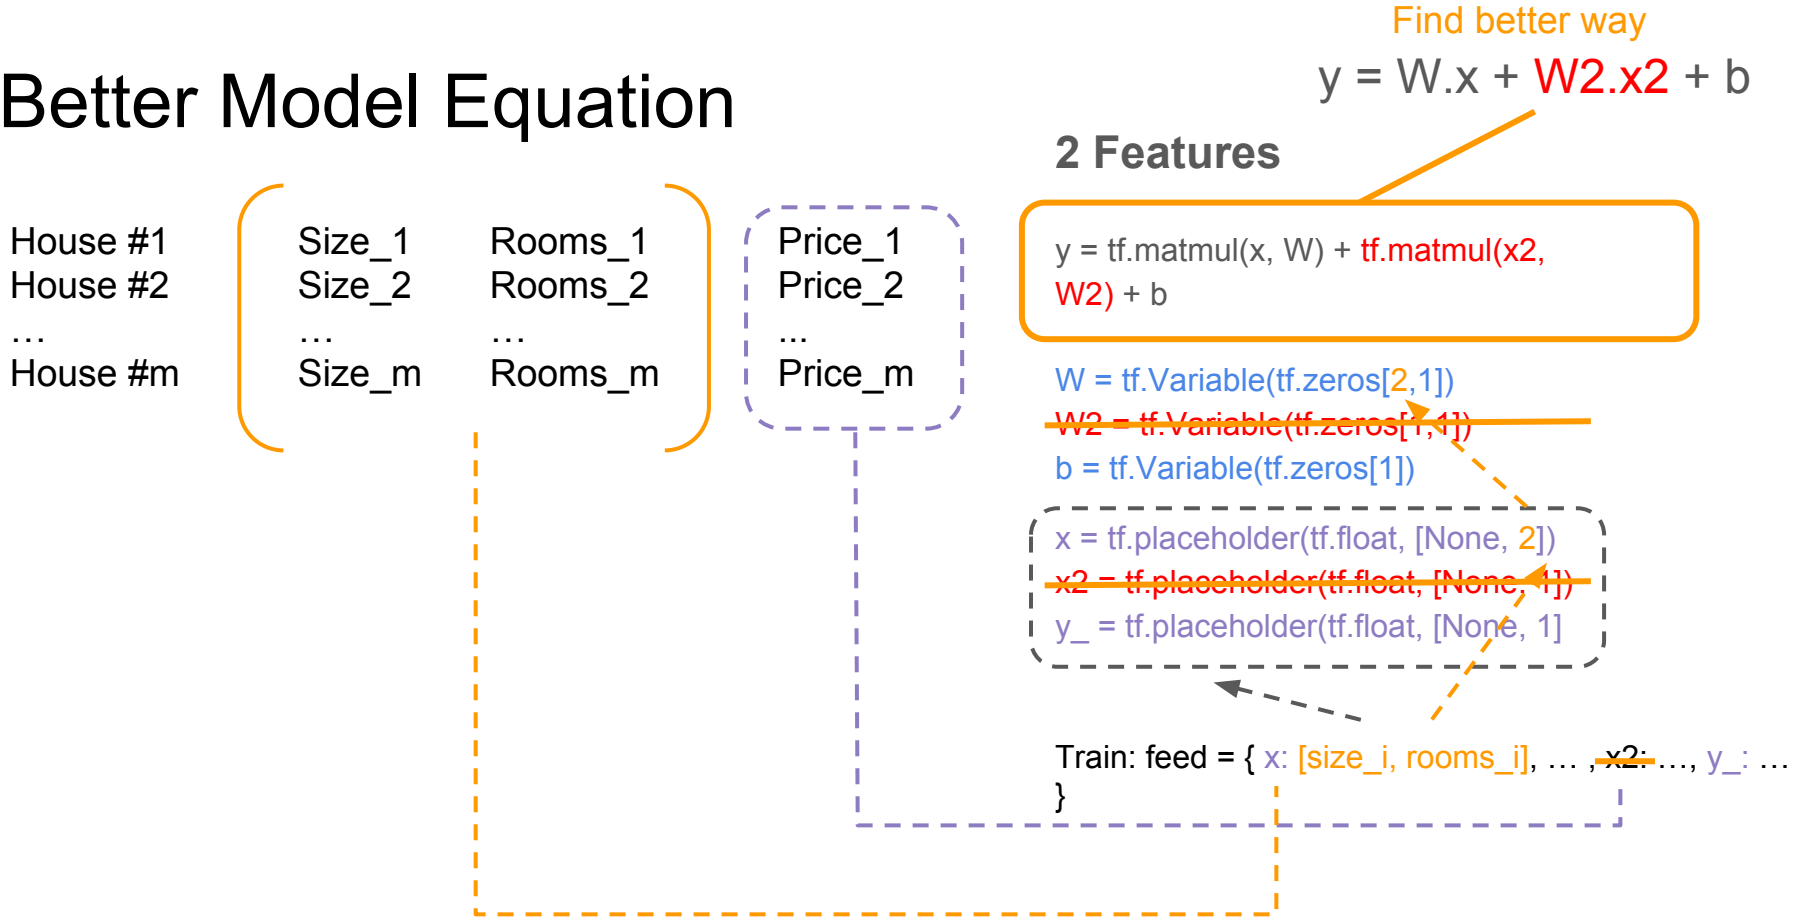

Better Model Equation

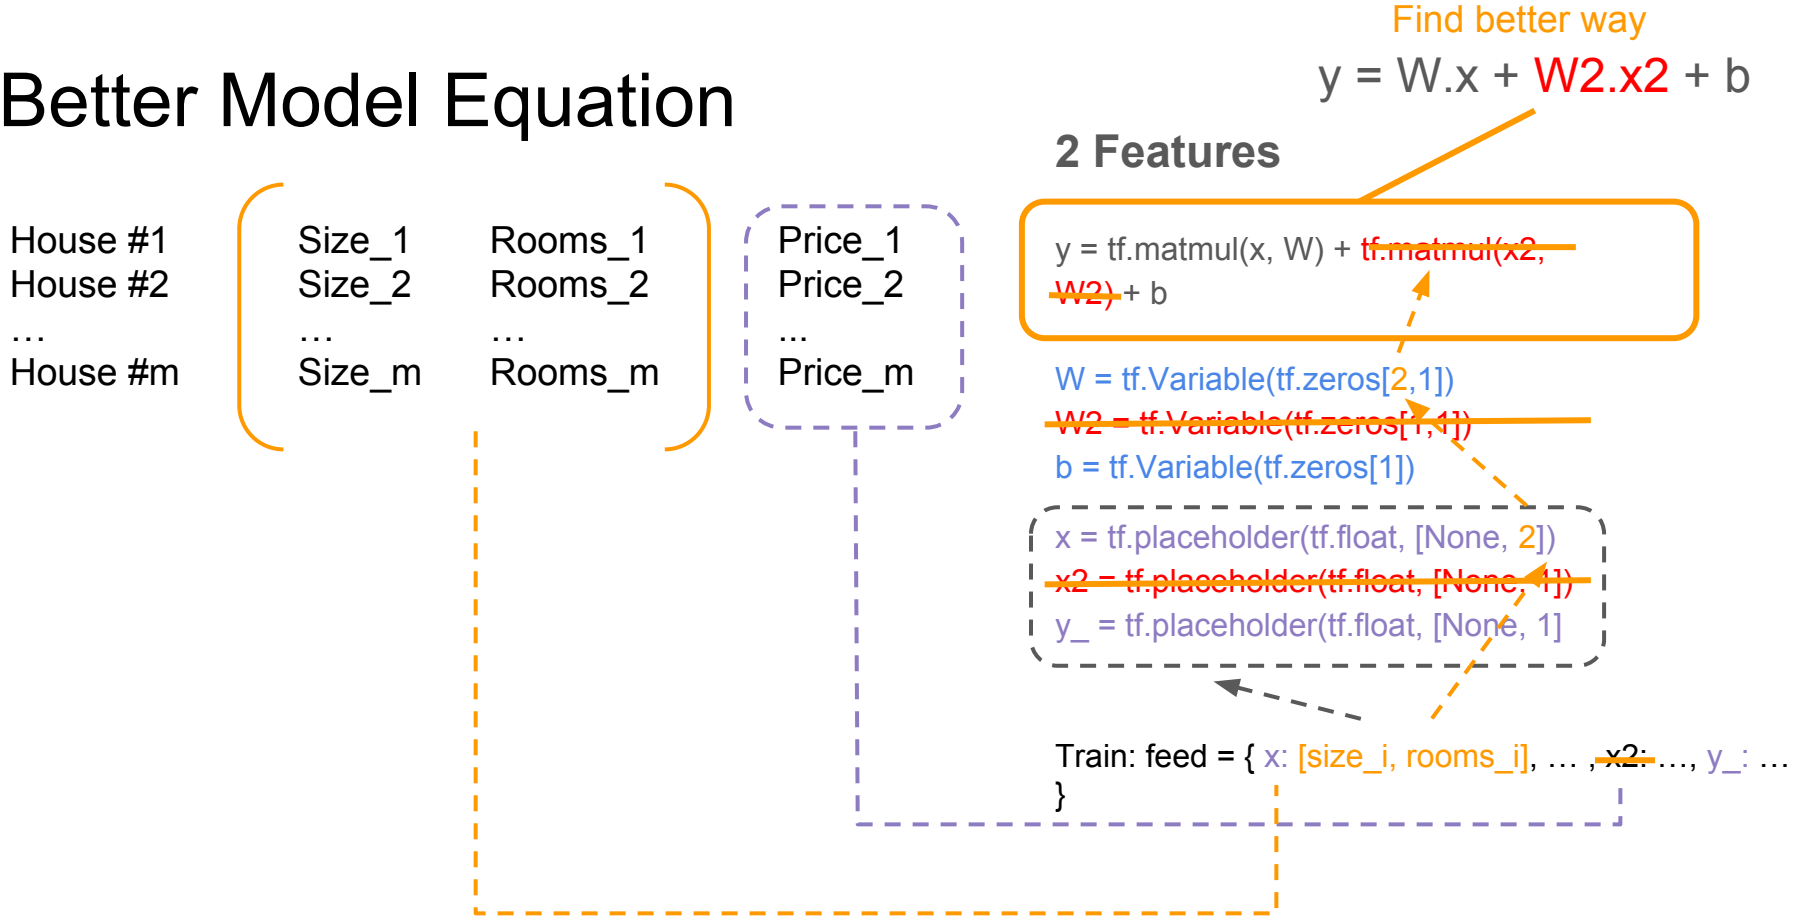

Better Model Equation

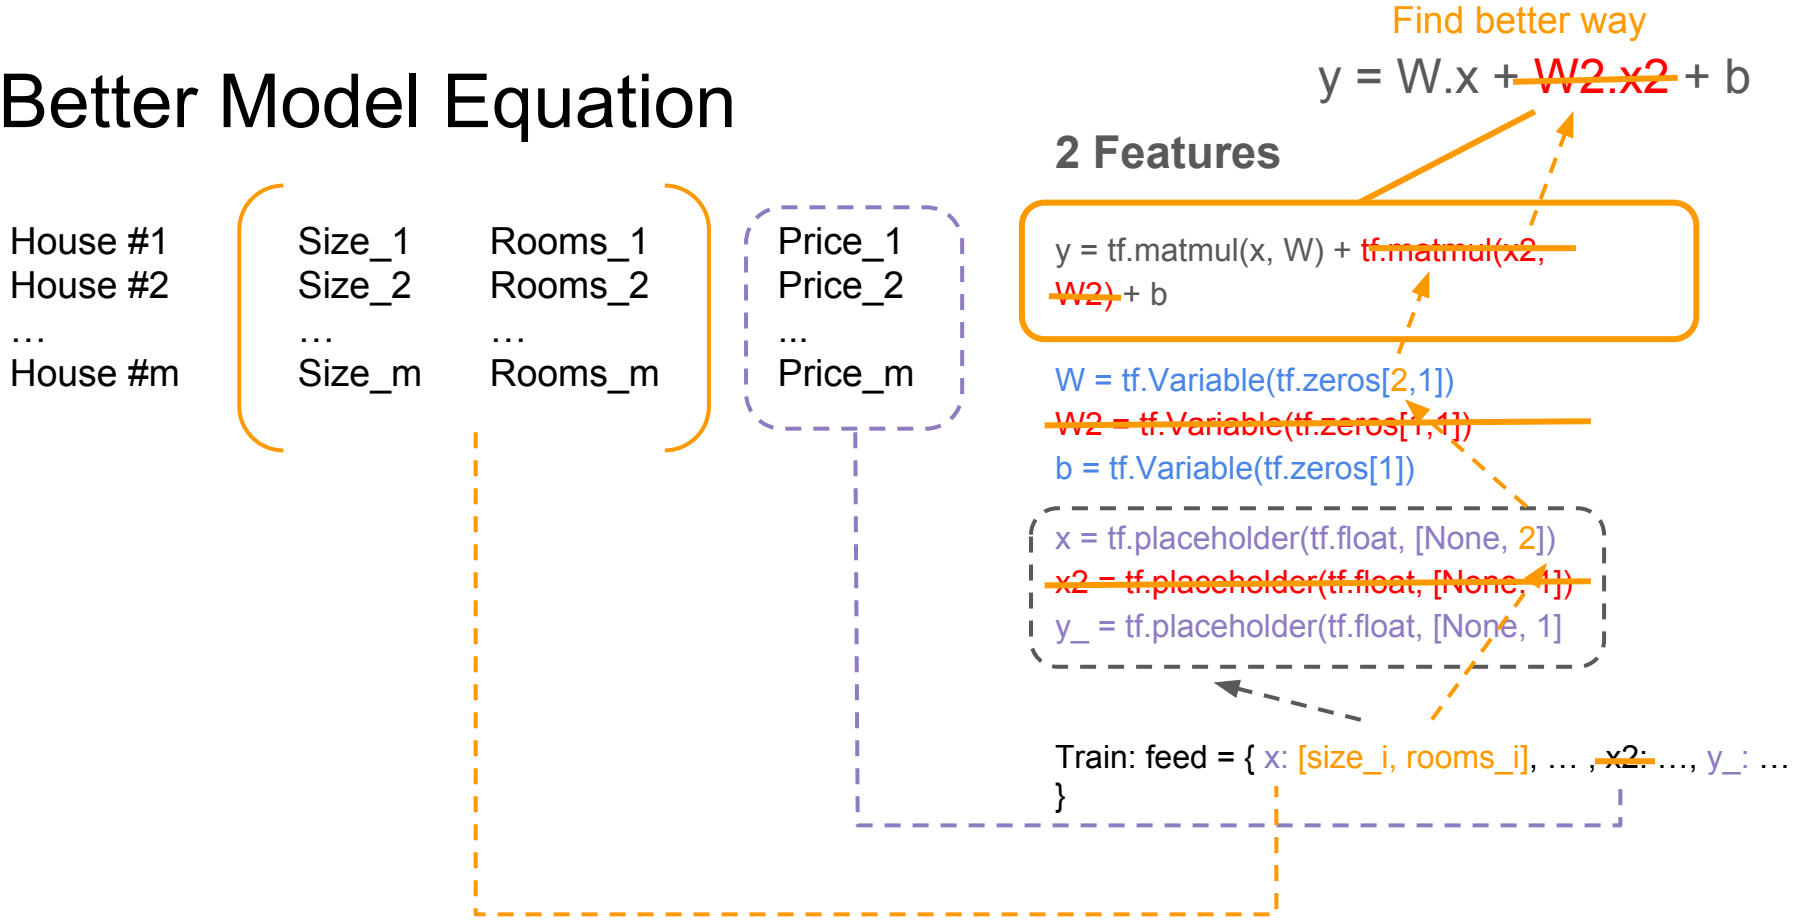

Tensorflow Graph (Messy)

1 Feature

```
y = tf.matmul(x, W) + b

W = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1])
```

2 Features

```
y = matmul(x, W) + matmul(x2, W2) + b

W = tf.Variable(tf.zeros[1,1])
W2 = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 1])
x2 = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1])
```

Tensorflow Graph (Clean)

1 Feature

```
y = tf.matmul(x, W) + b

W = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1])
```

2 Features

```
y = matmul(x, W) + matmul(x2, W2) + b

W = tf.Variable(tf.zeros[2,1])
W2 = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 2])
x2 = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1])
```

Tensorflow Graph (Clean and Formatted)

1 Feature

```
y = tf.matmul(x, W) + b

W = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1])
```

2 Features

```
y = matmul(x, W) + b

W = tf.Variable(tf.zeros[2,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 2])
y_ = tf.placeholder(tf.float, [None, 1])
```

Tensorflow Graph (Illustration)

1 Feature

y = tf.matmul(x, W) + b

W = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1])

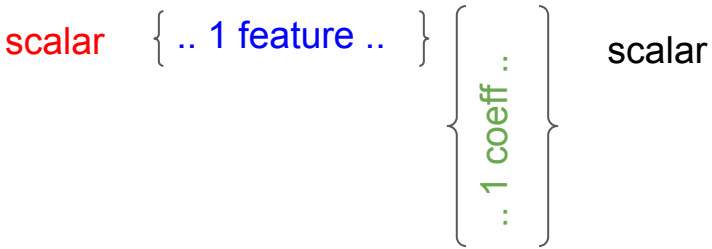

2 Features

y = matmul(x, W) + b

W = tf.Variable(tf.zeros[2,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 2])
y_ = tf.placeholder(tf.float, [None, 1])

Logistic Regression

Linear vs. Logistic Regression

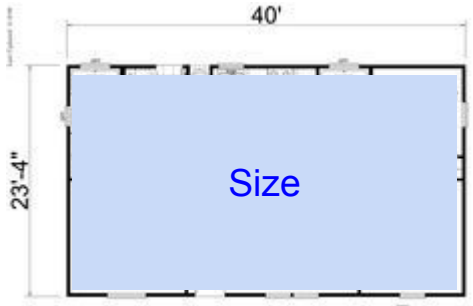

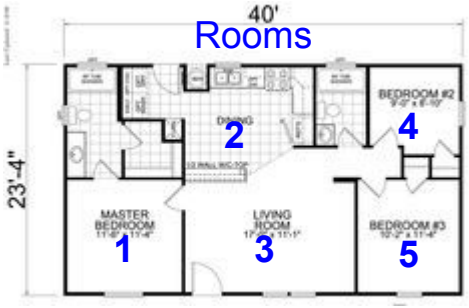

ML

Linear Regression

price (scalar)

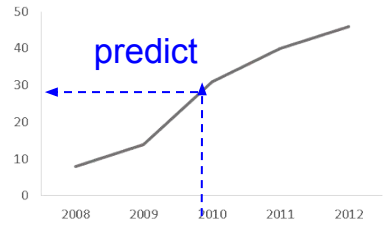

Linear vs. Logistic Regression

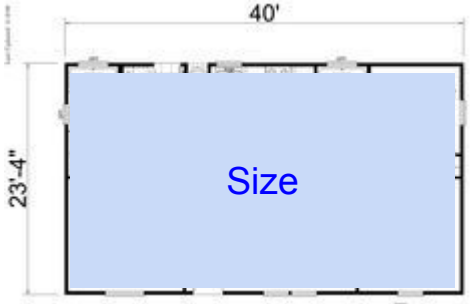

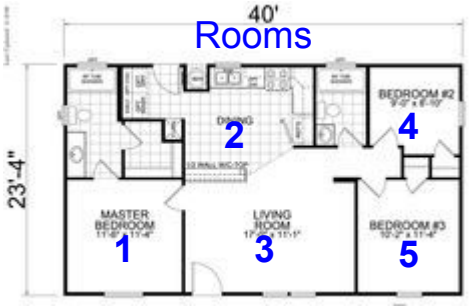

ML

Linear Regression

price (scalar)

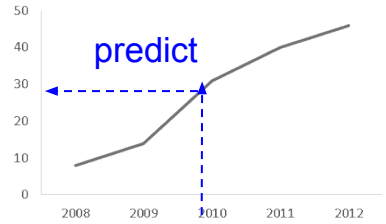

Image

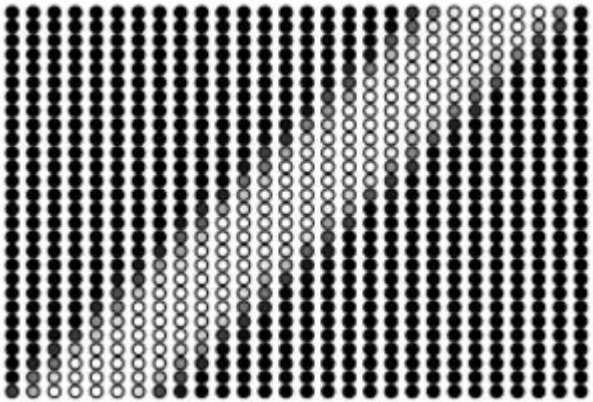

ML

Logistic Regression

number (discrete classes)

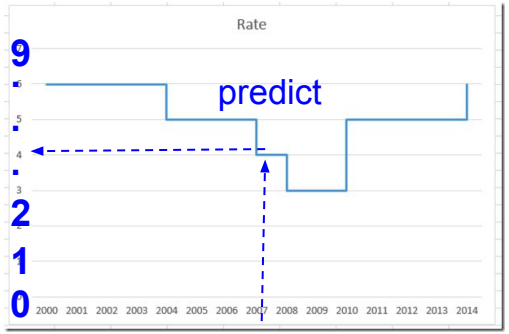

Image Features

Image Features 1

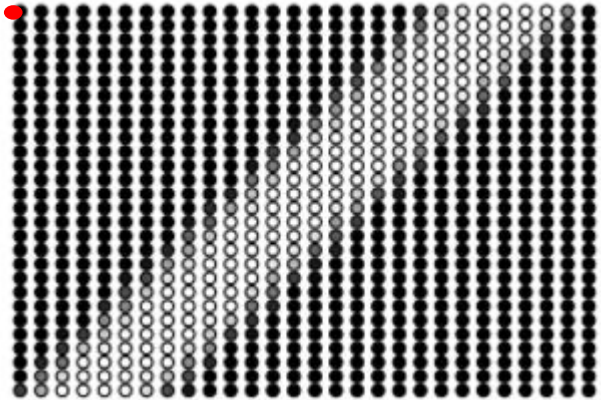

Grayscale value of each pixel

23	53	…	33
53	20	…	88
…	…	…	....
62	2	…	193

Image Features 2

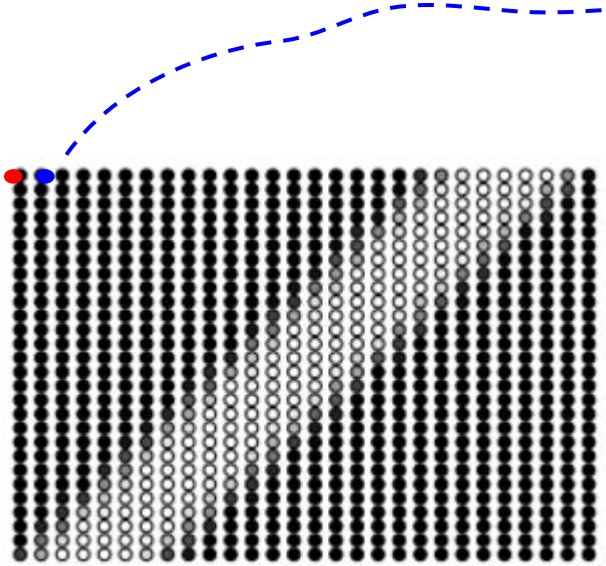

Grayscale value of each pixel

23	53	…	33
53	20	…	88
…	…	…	....
62	2	…	193

Image Features 3

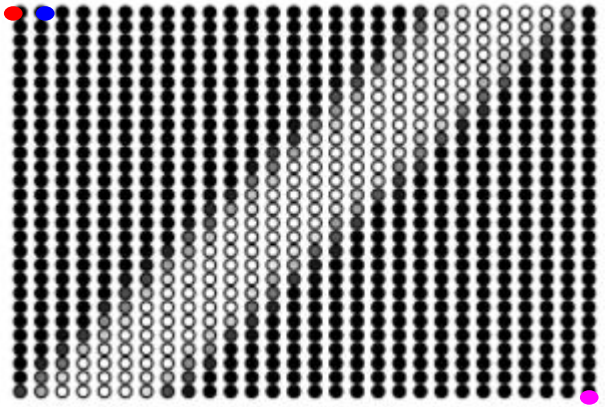

Grayscale value of each pixel

23	53	…	33
53	20	…	88
…	…	…	....
62	2	…	193

Logistic Regression: Change in Models

Model

Linear Regression

y = W.x + b

y: House price (scalar) prediction

x: [House size, Rooms]

Goal: Find scalars **W,b**

Logistic Regression

y = W.x + b

y: Discrete class [0,1,...9] prediction

x: [2-Dim pixel grayscale colors]

Goal: Find scalars **W, b**

Data Representation Comparison

Feature values x | Actual outcome y_

House #1	Size_1	Rooms_1	Price_1	
House #2	Size_2	Rooms_2	Price_2	
…	…	…	...	
House #m	Size_m	Rooms_m	Price_m	

Feature values x | Actual outcome y_

Image #1	23	53	…	33	5
	53	20	…	88	
	…	…	…	....	
	62	2	…	193	
Image #2	250	10	…	33	2
	103	5	…	88	
	…	…	…	...	
	5	114	…	193	
Image #m			...		3

Data Representation Comparison

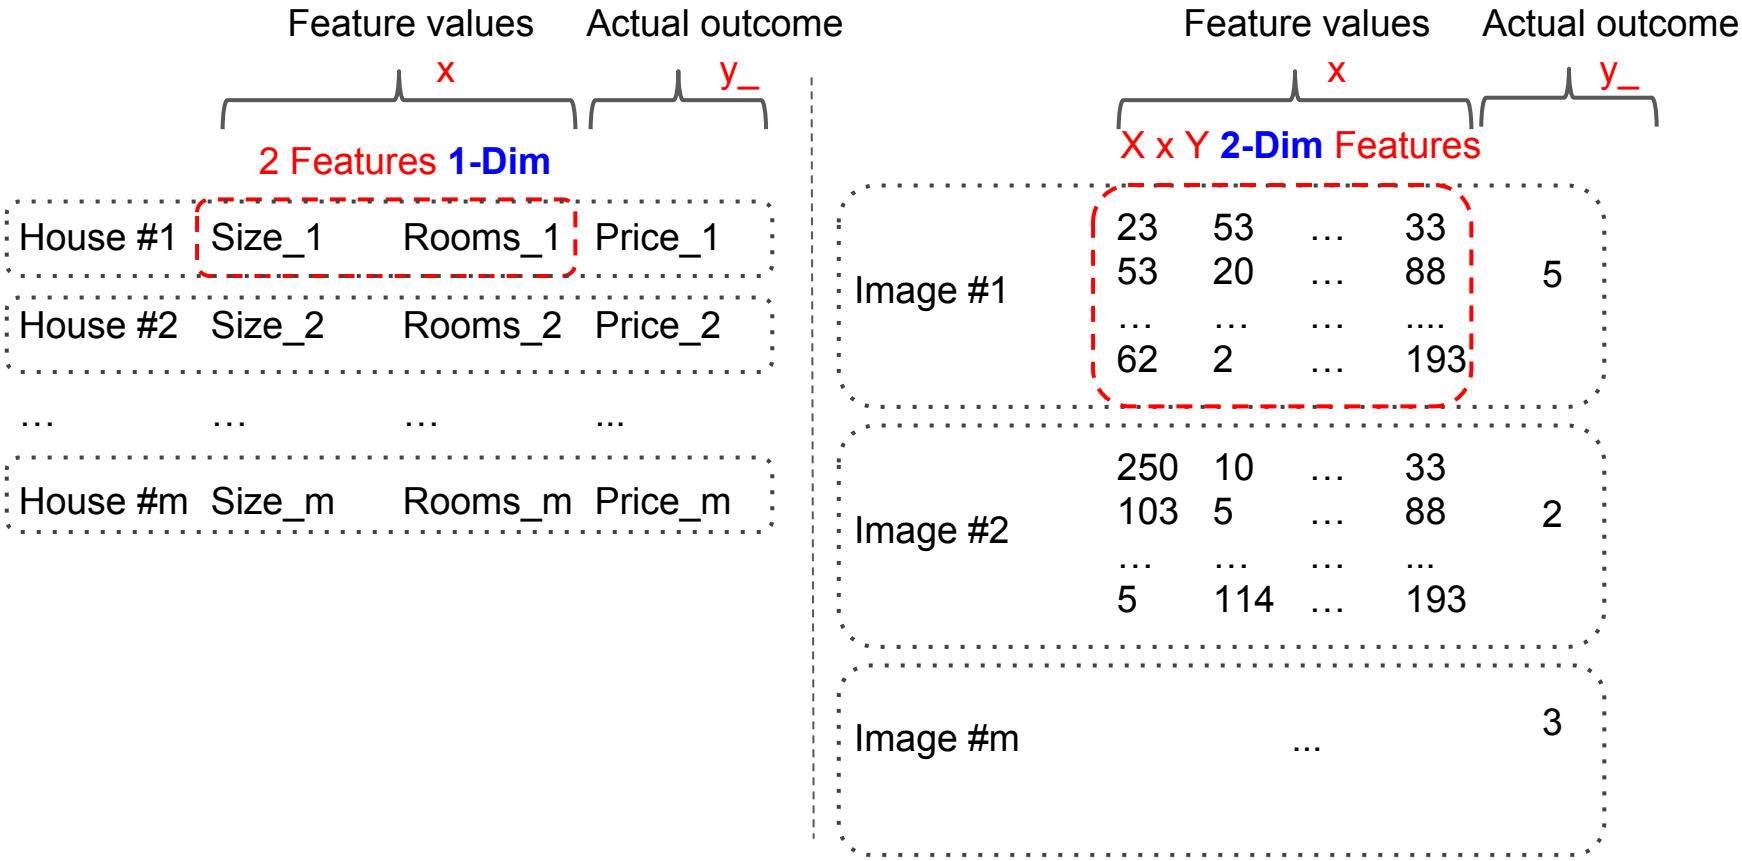

Data Representation Comparison

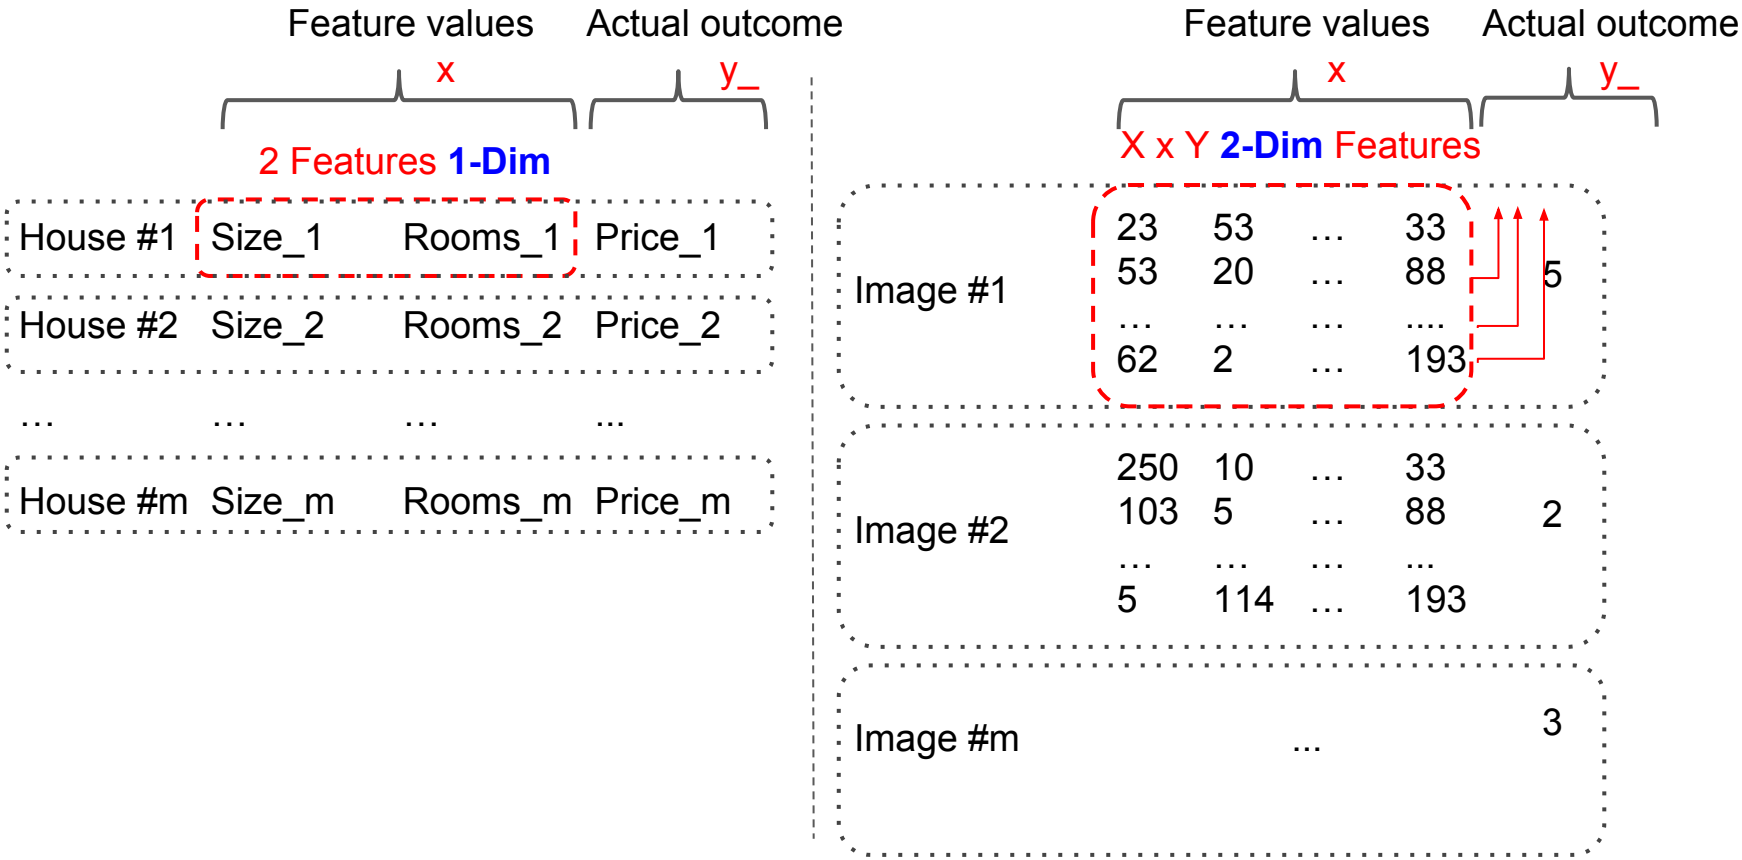

Data Representation Comparison

Feature values x | Actual outcome y_

2 Features **1-Dim**

House #1	Size_1	Rooms_1	Price_1
House #2	Size_2	Rooms_2	Price_2
…	…	…	…
House #m	Size_m	Rooms_m	Price_m

Feature values x | Actual outcome y_

Generalize into a multi-feature problem

X.Y ~~2~~ **1-Dim** Features

Image #1	23 53 … 33 53 20 … 88 … 62 2 … 193	5
Image #2	250 10 … 33 103 5 … 88 … 5 114 … 193	2
Image #m	…	3

Model

Linear Regression

y = W.x + b

y: House price (scalar) prediction

x: [House size, Rooms]

Goal: Find scalars **W,b**

Logistic Regression

y = W.x + b

y: Discrete class [0,1,...9] prediction

~~x: [2-Dim pixel grayscale colors]~~

x: [Pixel 1, Pixel 2, …, Pixel X.Y]

Goal: Find scalars **W, b**

Model

Linear Regression		Logistic Regression
y = W.x + b	This needs change as well!	y = W.x + b
y: House price (scalar) prediction		y: Discrete class [0,1,...9] prediction
x: [House size, Rooms]		~~x: [2-Dim pixel grayscale colors]~~
		x: [Pixel 1, Pixel 2, …, Pixel X.Y]
Goal: Find scalars **W,b**		Goal: Find scalars **W, b**

Why Can’t ‘y’ be left as scalar of 0-9?

HINT: Beside the model, when doing ML we need this function!

Logistic Regression: Change in Cost Function

Linear Regression: Cost

Scalar

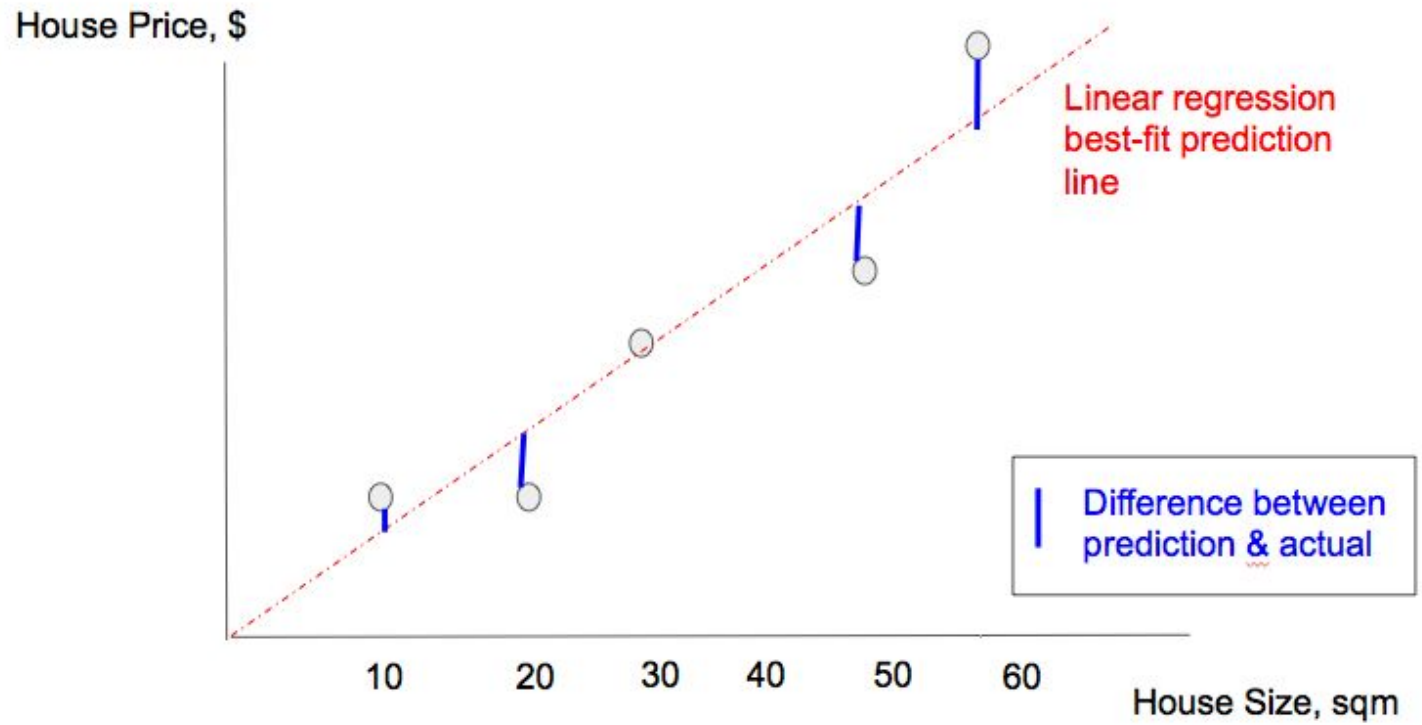

Linear Regression: Cost, NOT

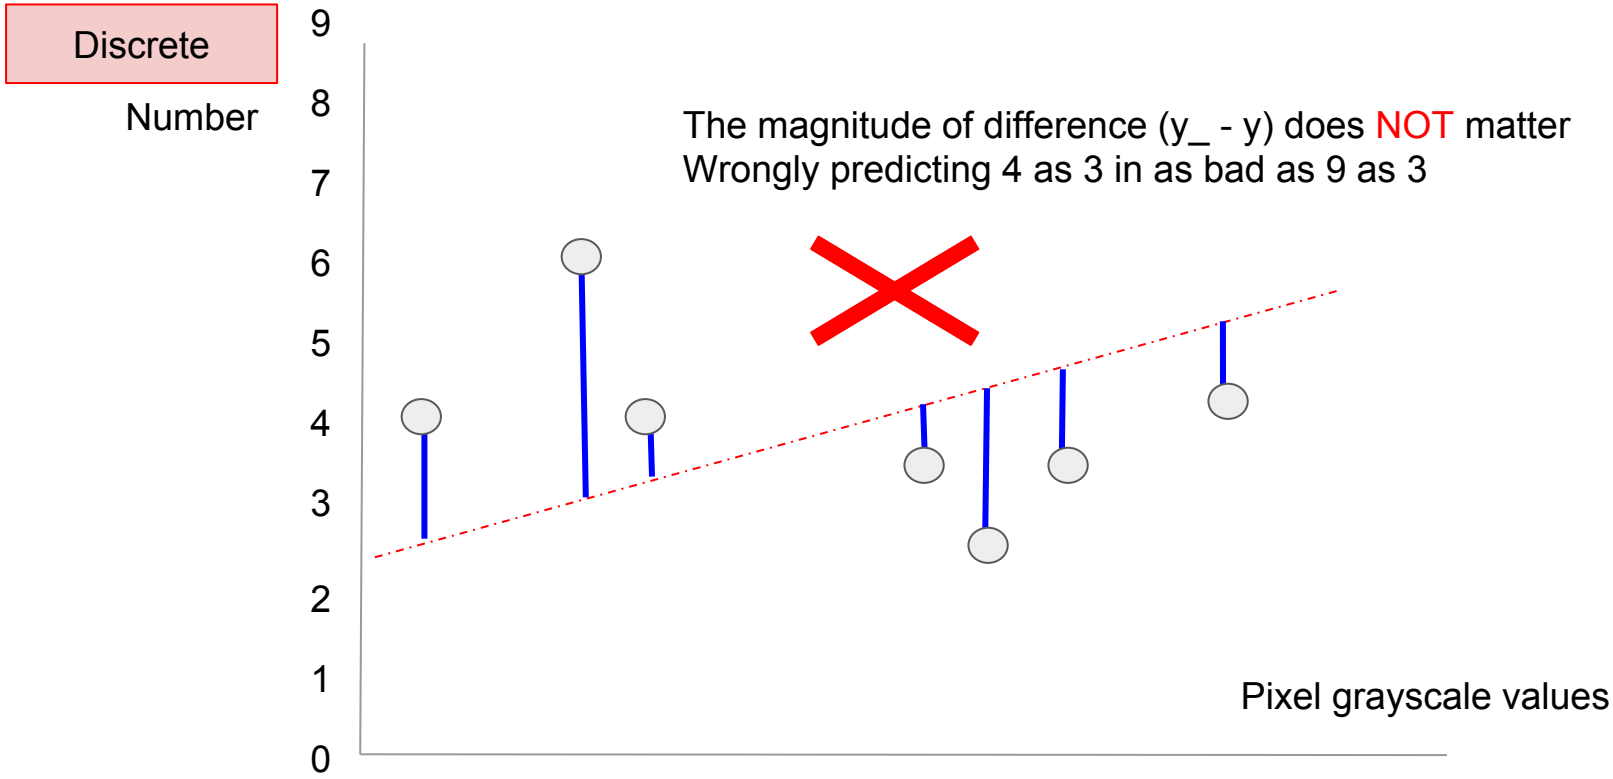

Logistic Regression: Prediction

Image

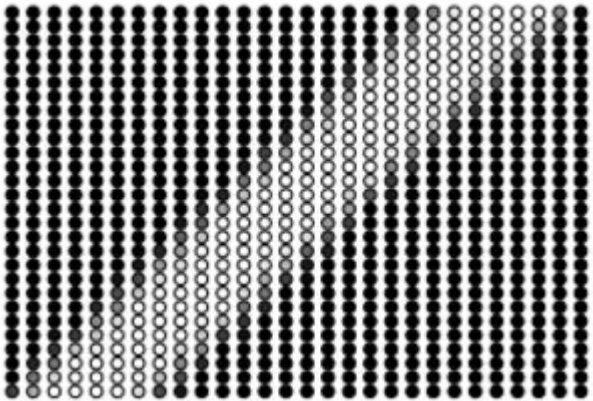

ML

Logistic Regression

1

Logistic Regression: Prediction

Image

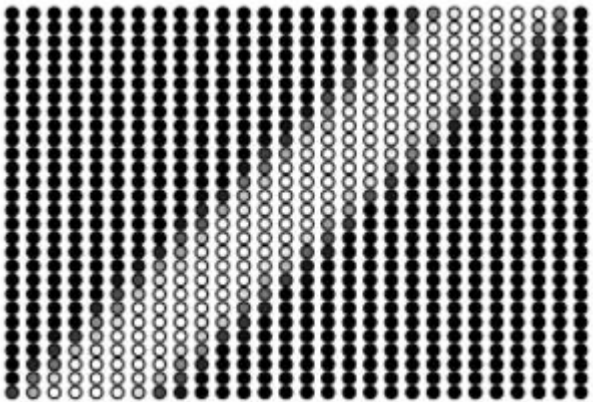

ML

Logistic Regression

Probability

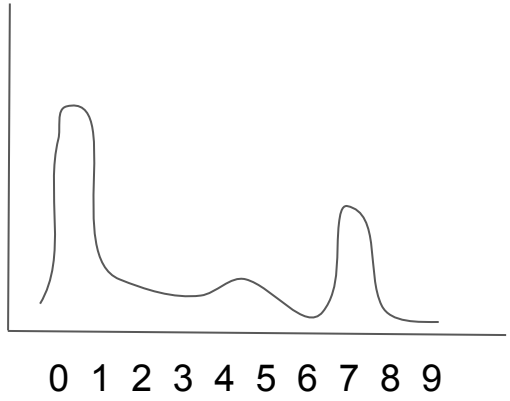

Logistic Regression: Prediction

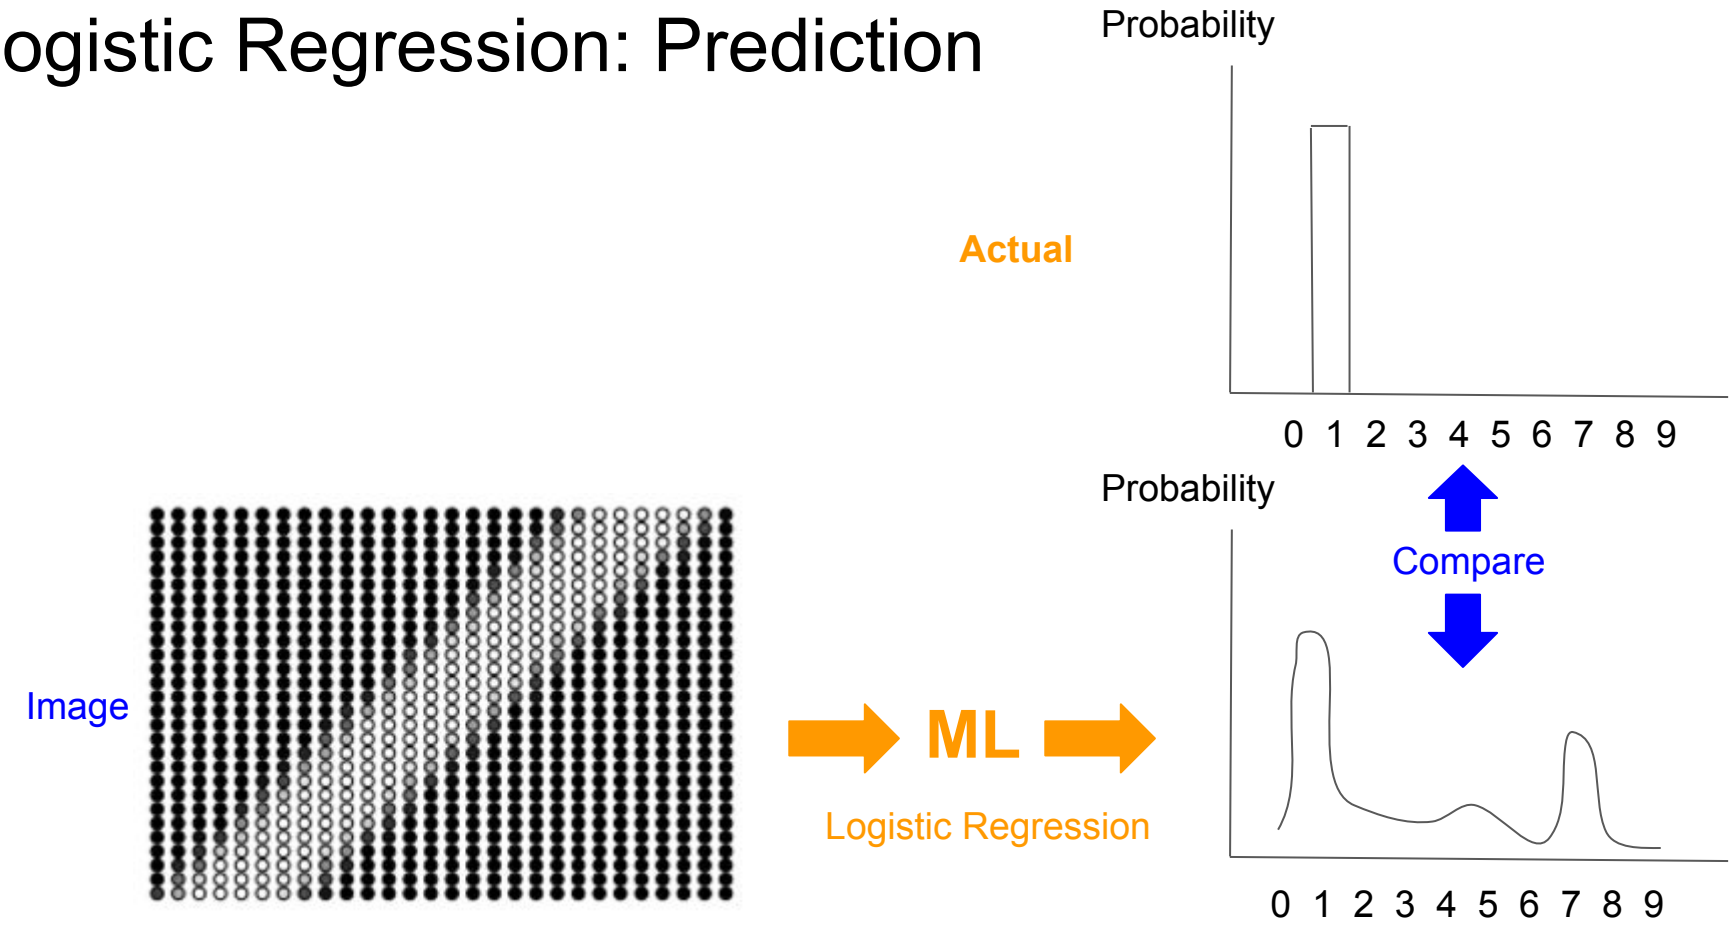

Cross-Entropy: Cost Function 1

$$H_{y'}(y) = -\sum_i y'_i \log(y_i)$$

cross_entropy = -tf.reduce_sum(y_*tf.log(y))

Cross-Entropy: Cost Function 2

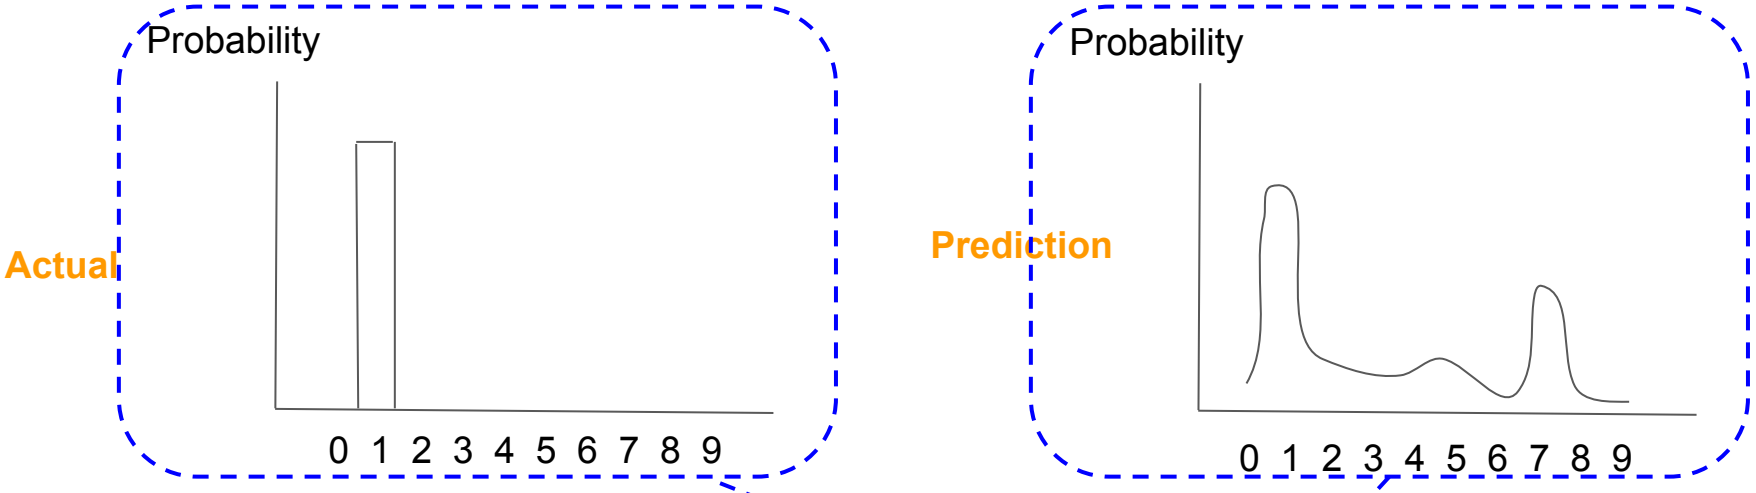

$$H_{y'}(y) = -\sum_i y'_i \log(y_i)$$

```
cross_entropy = -tf.reduce_sum(y_*tf.log(y))
```

Cross-Entropy: Cost Function 3

Almost inverse because Probability < 1

Probability

Actual

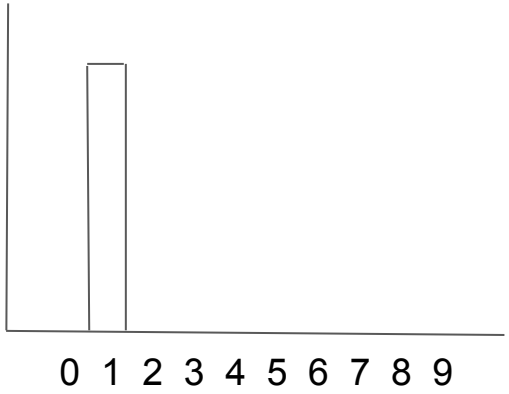

-log (Probability**)**

Prediction

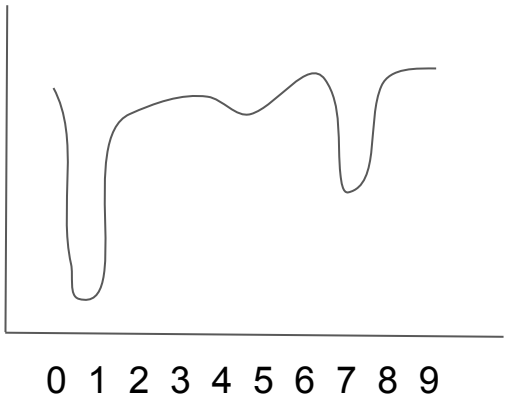

$$H_{y'}(y) = -\sum_i y'_i \log(y_i)$$

cross_entropy = -tf.reduce_sum(y_*tf.log(y))

Graph: -log(probability)

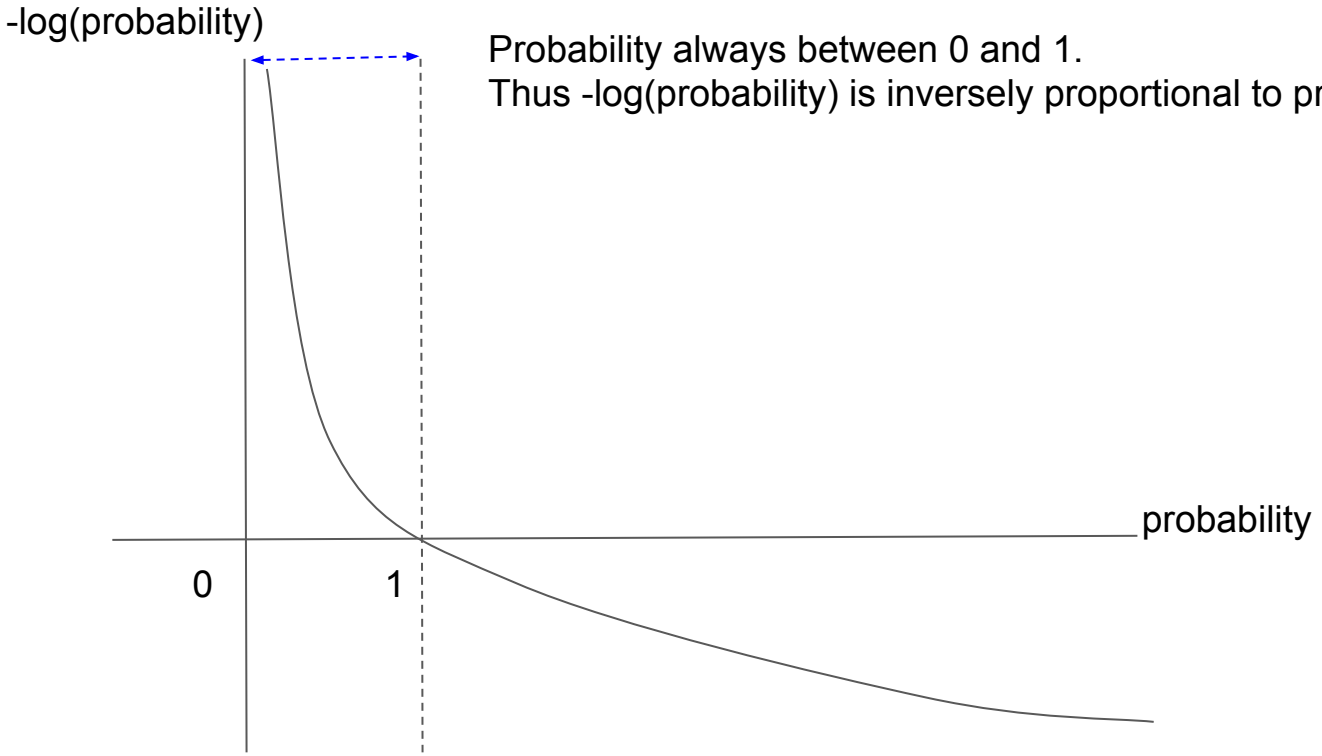

Probability always between 0 and 1.
Thus -log(probability) is inversely proportional to probability

Cross-Entropy: Cost Function x x xx

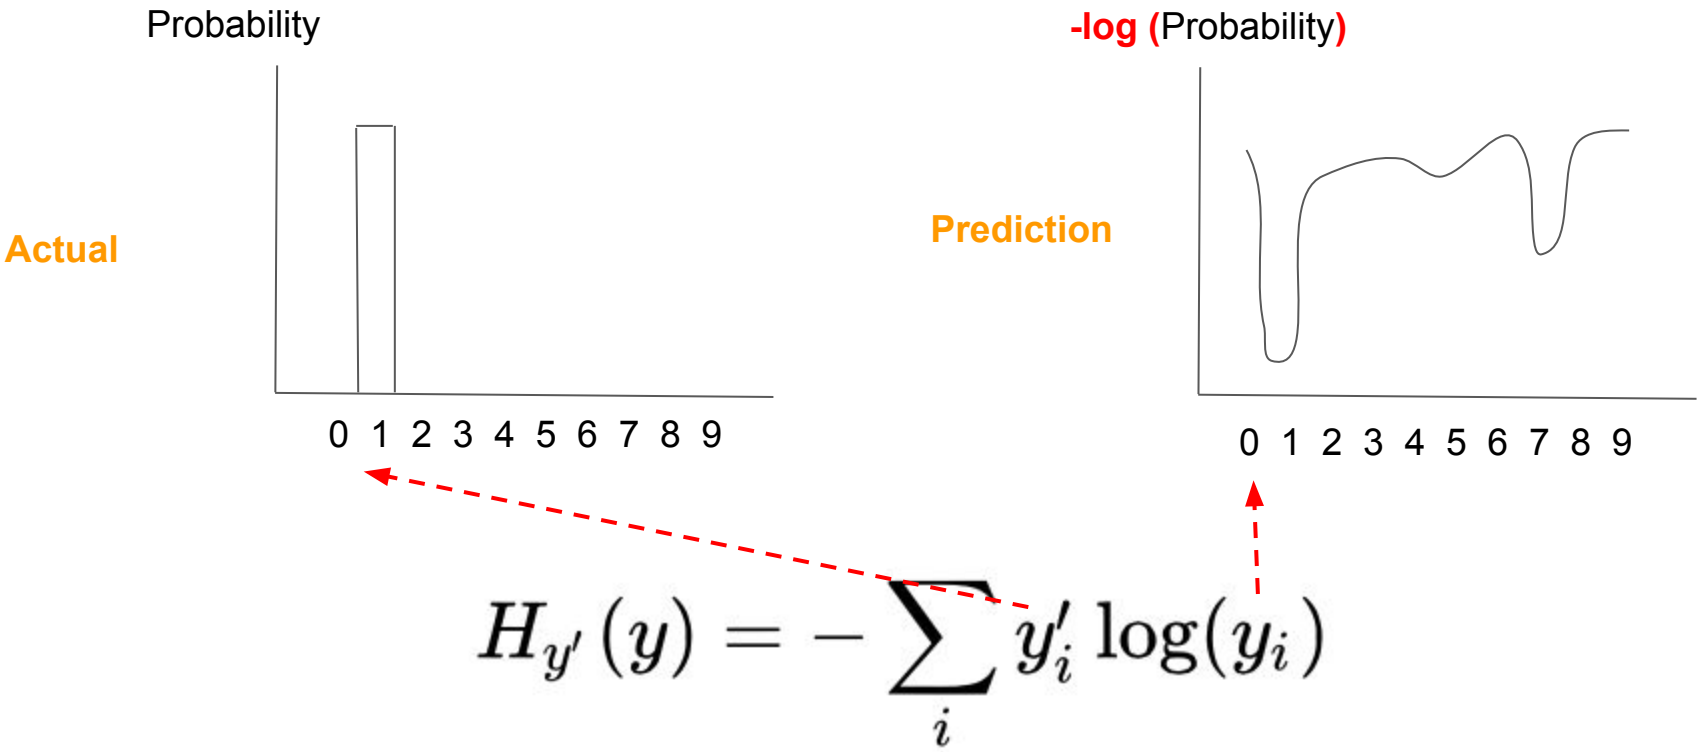

$$H_{y'}(y) = -\sum_i y'_i \log(y_i)$$

cross_entropy = -tf.reduce_sum(y_*tf.log(y))

Cross-Entropy: Cost Function

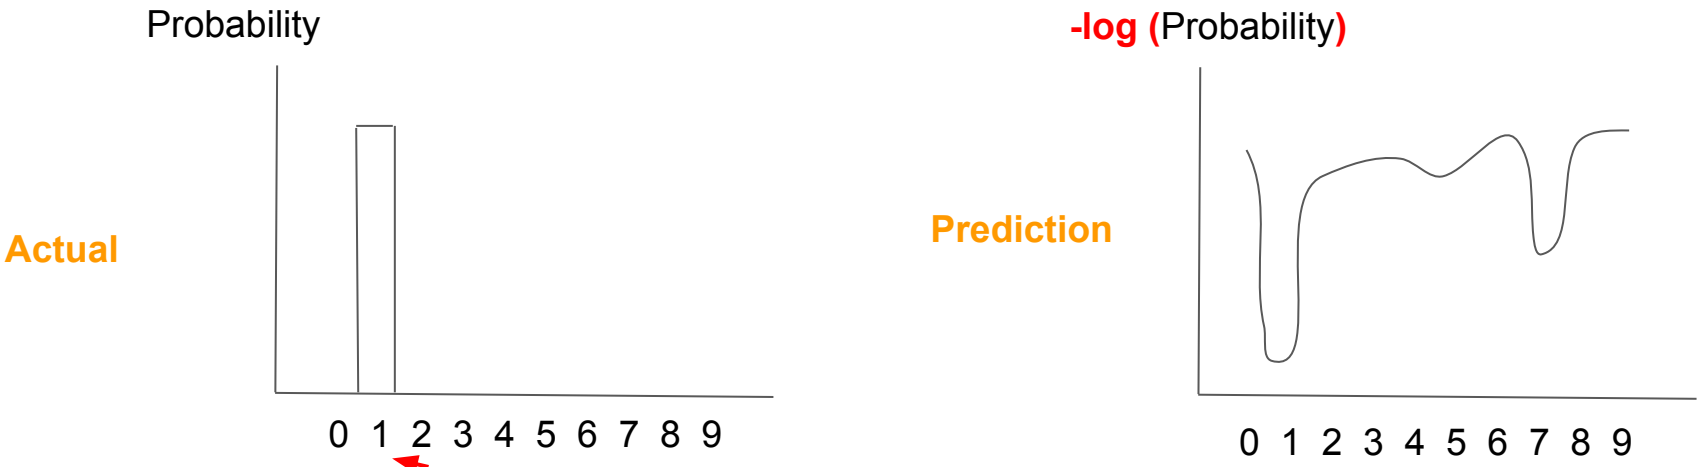

$$H_{y'}(y) = -\sum_i y'_i \log(y_i)$$

cross_entropy = -tf.reduce_sum(y_*tf.log(y))

Cross-Entropy: Cost Function

Probability

-log (Probability**)**

Actual

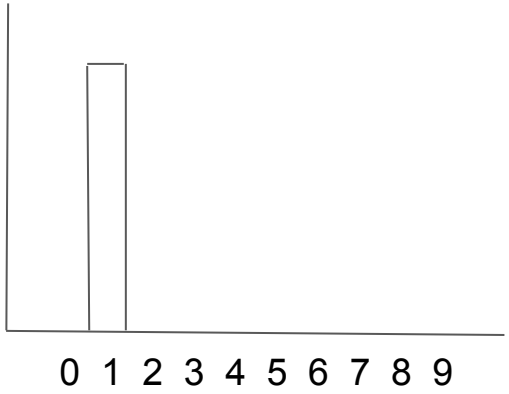

Prediction

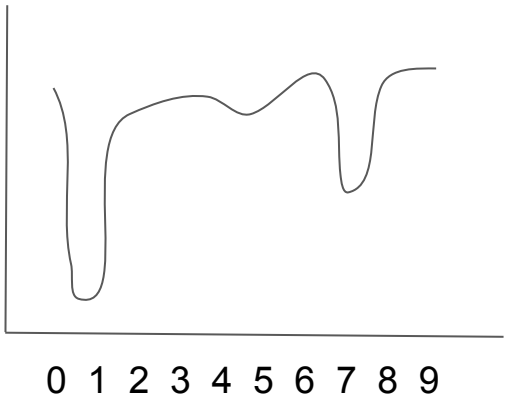

$$H_{y'}(y) = -\sum_i y'_i \log(y_i)$$

cross_entropy = -tf.reduce_sum(y_*tf.log(y))

Cross-Entropy: Cost Function

Probability

- Only y'i = 1 and log(yi) matters
- The smaller the log(yi) the lower cost
- The higher the yi the lower the cost

-log (Probability**)**

Actual

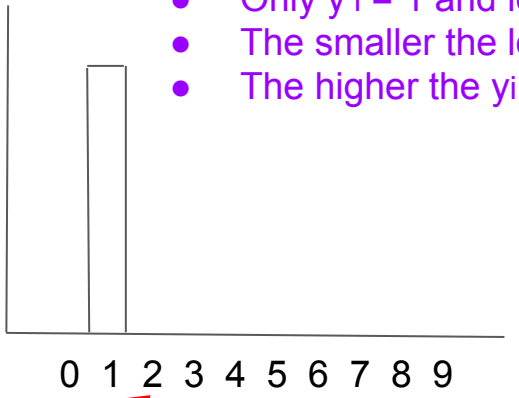

Prediction

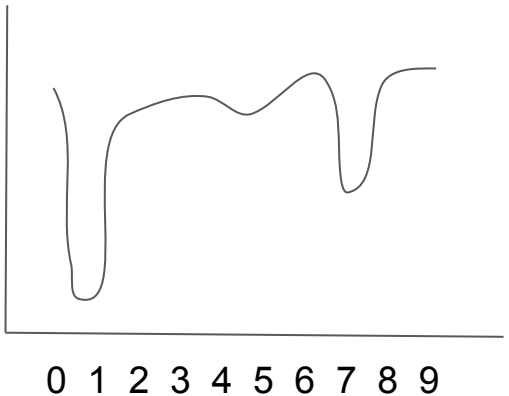

$$H_{y'}(y) = -\sum_i y'_i \log(y_i)$$

cross_entropy = -tf.reduce_sum(y_*tf.log(y))

Logistic Regression: Prediction

Image

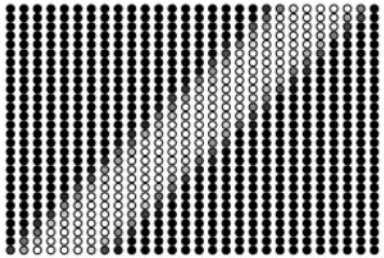

ML

Linear Regression

1

Image

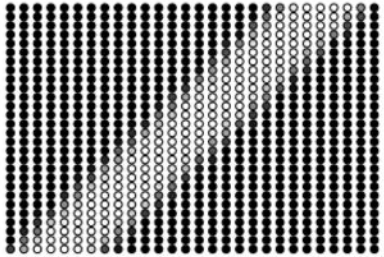

ML

Logistic Regression

Probability

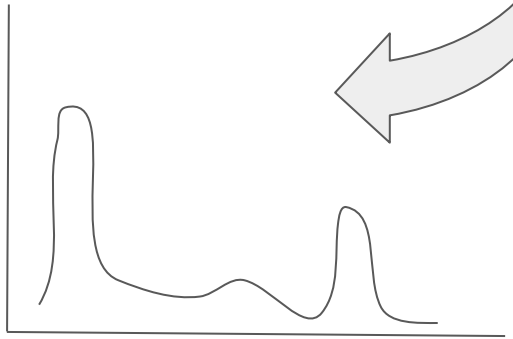

0 1 2 3 4 5 6 7 8 9

Tensorflow Graph (Review)

1 Feature

```
y = tf.matmul(x, W) + b

W = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1])
```

2 Features

```
y = matmul(x, W) + b

W = tf.Variable(tf.zeros[2,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 2])
y_ = tf.placeholder(tf.float, [None, 1])
```

Tensorflow Graph (Review 2)

1 Feature

y = tf.matmul(x, W) + b

W = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1])

scalar { .. 1 feature .. } { .. 1 coeff .. } scalar

2 Features

y = matmul(x, W) + b

W = tf.Variable(tf.zeros[2,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 2])
y_ = tf.placeholder(tf.float, [None, 1])

scalar { .. 2 features .. } { .. 2 coeffs .. } scalar

Tensorflow Graph (Review 2)

1 Feature

y = tf.matmul(x, W) + b

W = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1])

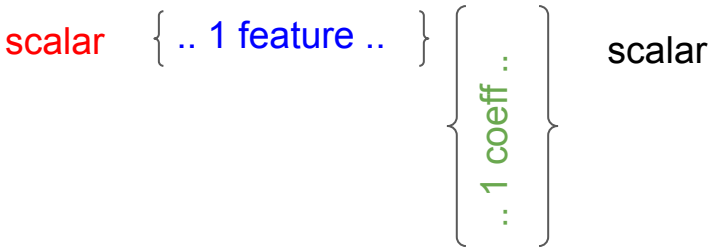

2 Features

y = matmul(x, W) + b

W = tf.Variable(tf.zeros[2,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 2])
y_ = tf.placeholder(tf.float, [None, 1])

Apply multi-feature linear regression

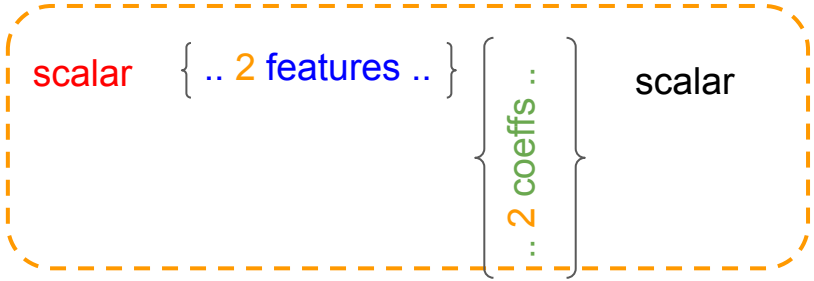

Logistic Regression: Prediction

y = tf.matmul(x,W) + b

scalar { .. n features .. } { .. n coeffs .. } scalar

Image

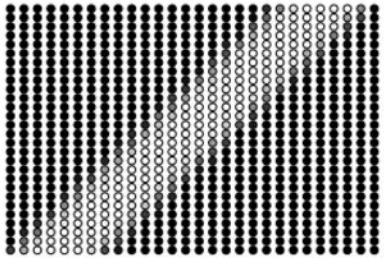

ML

Linear Regression

1

Image

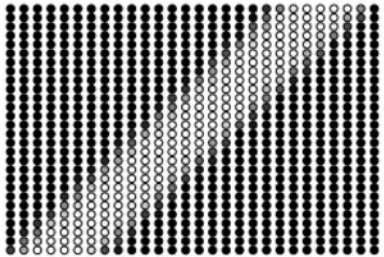

ML

Logistic Regression

Probability

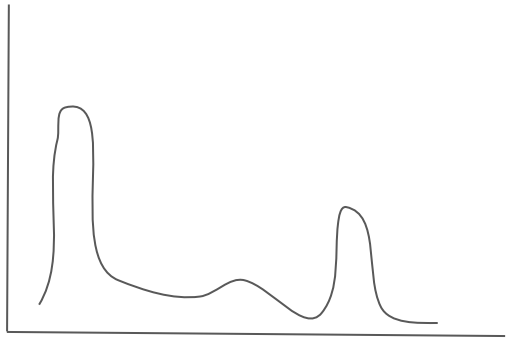

0 1 2 3 4 5 6 7 8 9

Logistic Regression: Prediction

y = tf.matmul(x,W) + b

Image

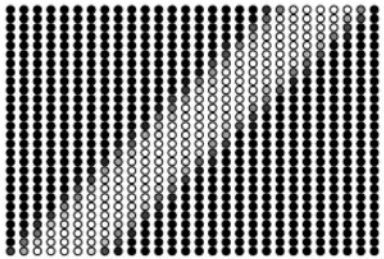

ML

Linear Regression

scalar

1

{ .. n features .. }

{ .. n coeffs .. }

scalar

Probability

{.. k class ..}

Image

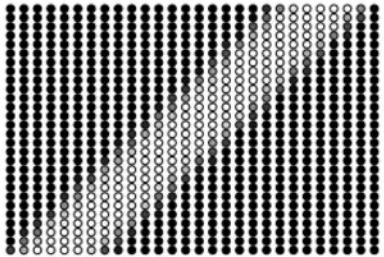

ML

Logistic Regression

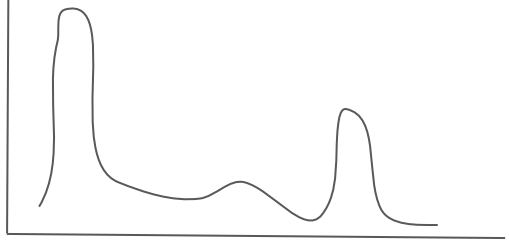

0 1 2 3 4 5 6 7 8 9

Logistic Regression: Prediction

y = tf.matmul(x,W) + b

Image

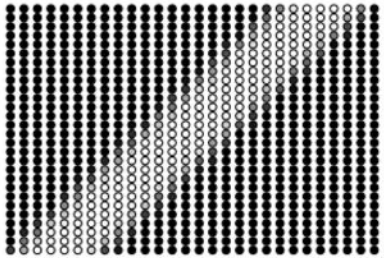

ML

Linear Regression

1

scalar .. n features .. .. n coeffs .. scalar

y = tf.matmul(x,W) + b

Probability

Image

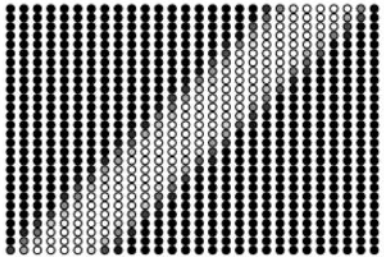

ML

Logistic Regression

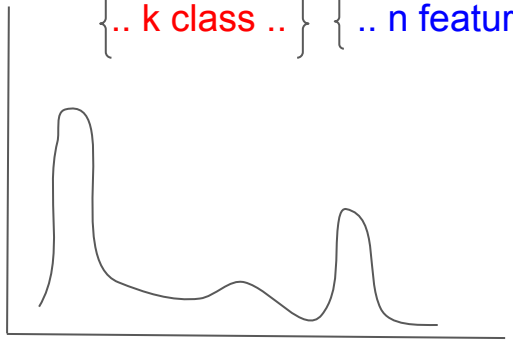

.. k class .. .. n features .. .. k class .. .. k class .. .. n coeffs ..

0 1 2 3 4 5 6 7 8 9

Tensorflow Graph: Basic, Multi-feature, Multi-class

1 Feature

```
y = tf.matmul(x, W) + b

W = tf.Variable(tf.zeros[1,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 1])
y_ = tf.placeholder(tf.float, [None, 1])
```

2 Features

```
y = matmul(x, W) + b

W = tf.Variable(tf.zeros[2,1])
b = tf.Variable(tf.zeros[1])

x = tf.placeholder(tf.float, [None, 2])
y_ = tf.placeholder(tf.float, [None, 1])
```

2 Features, 10 Classes

```
y = matmul(x, W) + b

W = tf.Variable(tf.zeros[2,10])
b = tf.Variable(tf.zeros[10])

x = tf.placeholder(tf.float, [None, 2])
y_ = tf.placeholder(tf.float, [None, 10])
```

Logistic Regression: Prediction

Great but….sum of all prediction ‘probability’ NOT 1

Actual

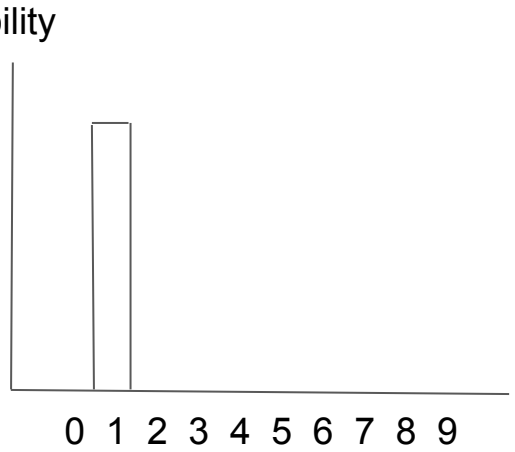

Image

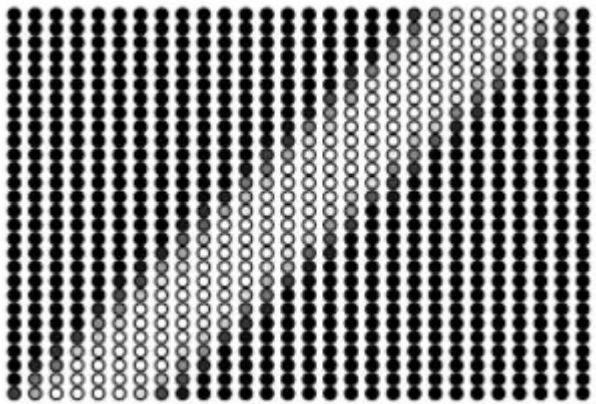

ML

Logistic Regression

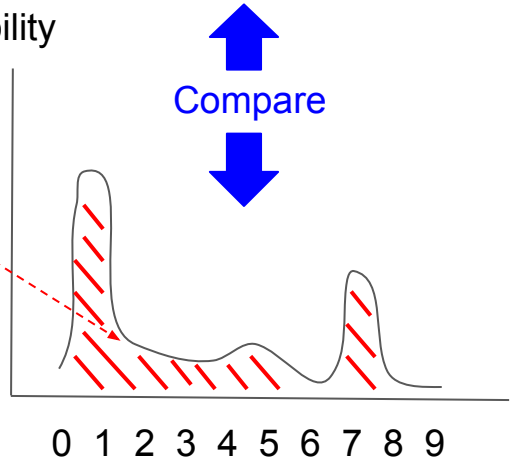

$$\text{softmax}(x)_i = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$$

Image

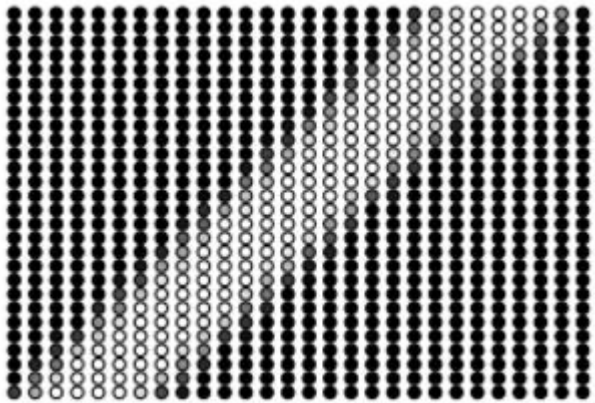

ML

Logistic Regression

Value

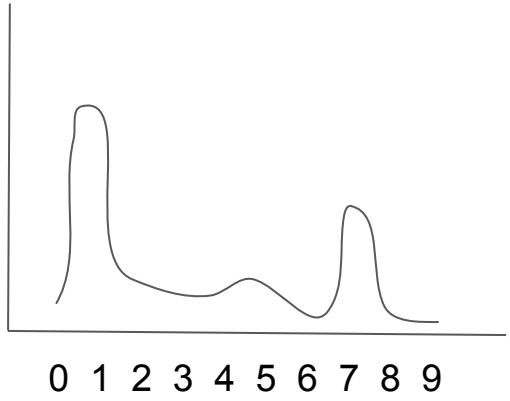

$$\text{softmax}(x)_i = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$$

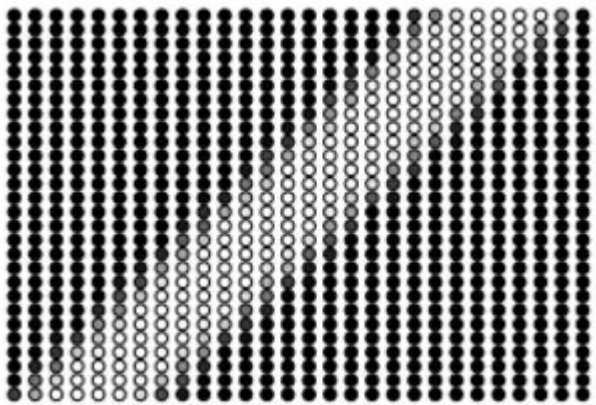

ML

Logistic Regression

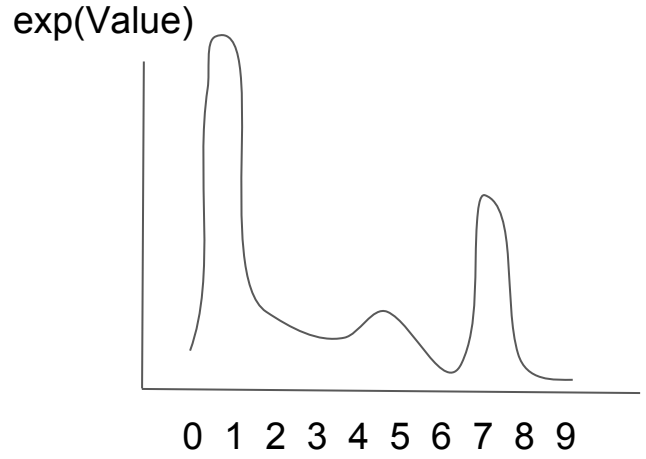

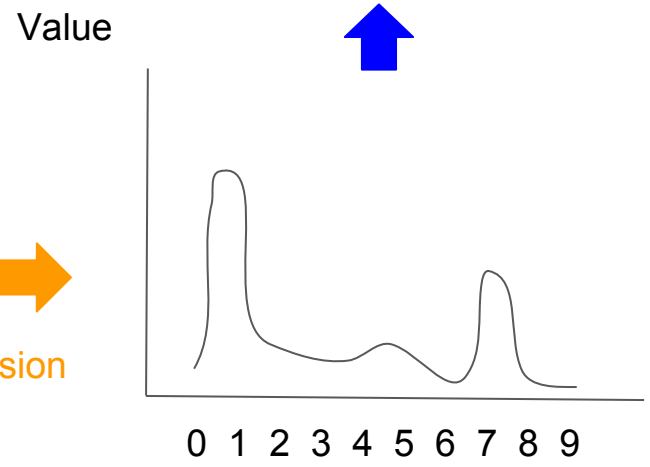

$$\mathrm{softmax}(x)_i = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$$

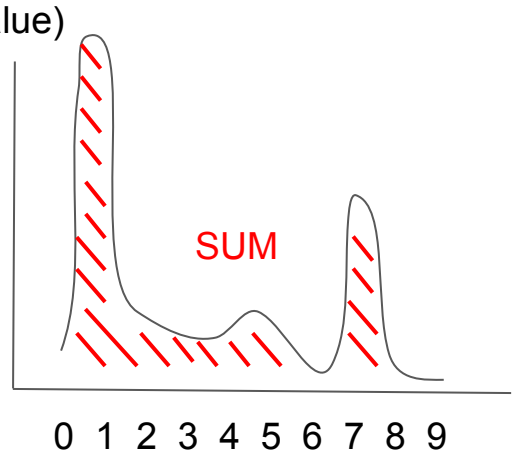

Image

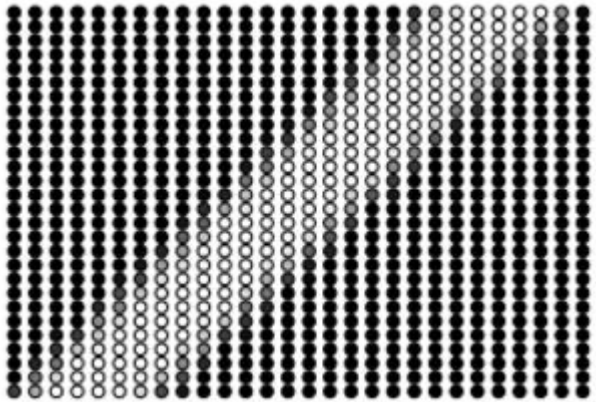

ML

Logistic Regression

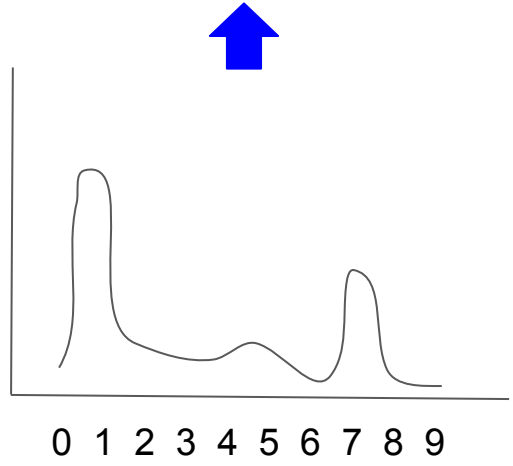

$$\text{softmax}(x)_i = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$$

exp(Value)

SUM

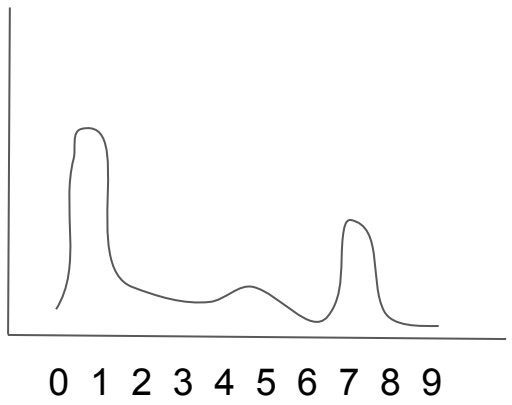

exp(Value)

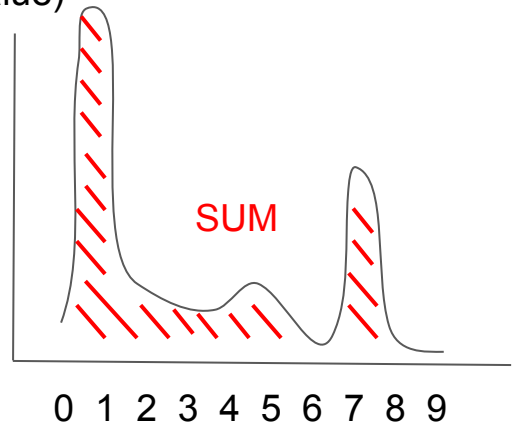

Image

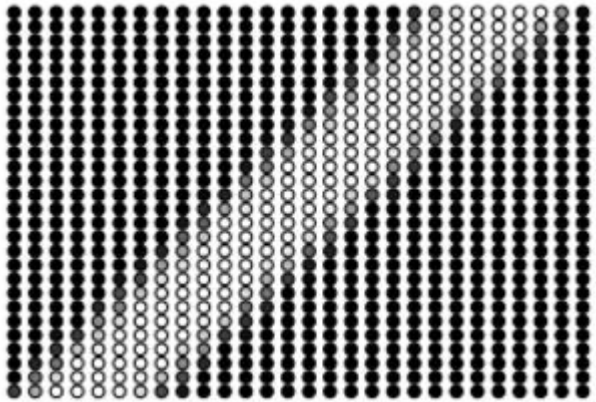

ML

Logistic Regression

Value

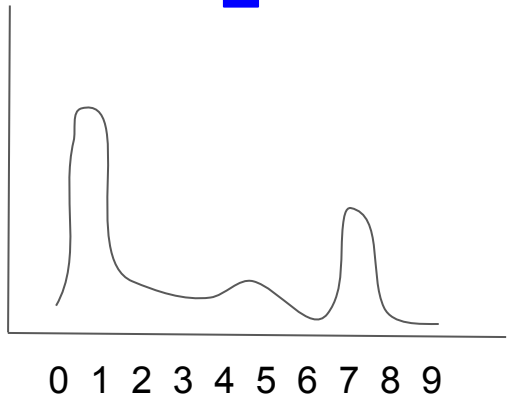

exp(Value)

SUM

$$\text{softmax}(x)_i = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$$

sum of all prediction 'probability' is 1

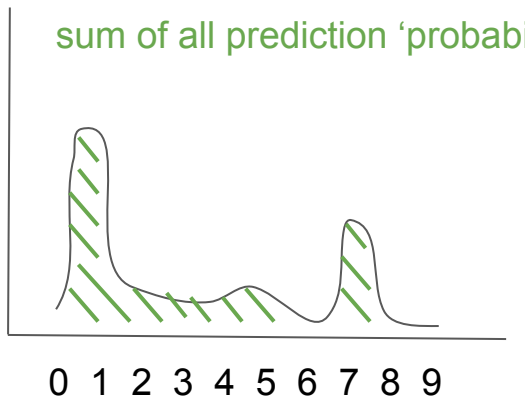

exp(Value)

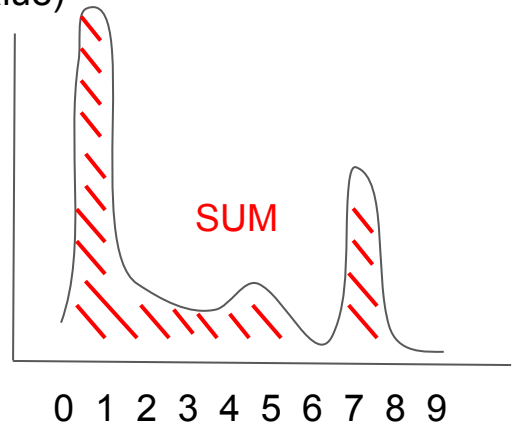

Image

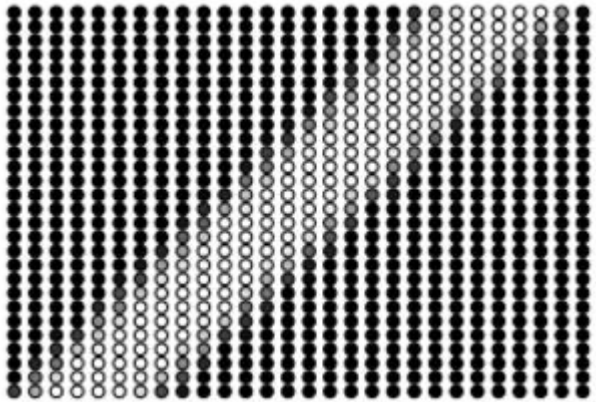

ML

Logistic Regression

Value

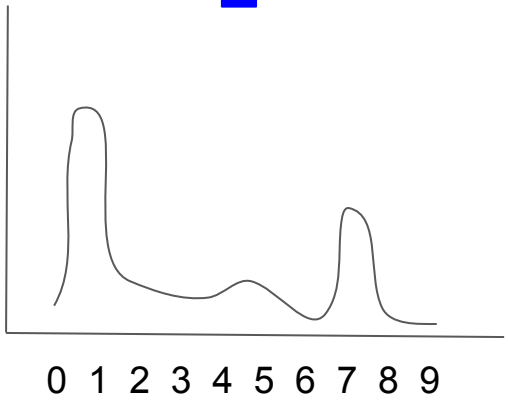

Before softmax

$$y = tf.matmul(x, W) + b$$

With softmax

~~y = tf.matmul(x, W) + b~~

$$\text{softmax}(x)_i = \frac{\exp(x_i)}{\sum_j \exp(x_j)}$$

y = tf.nn.softmax(tf.matmul(x, W) + b)

Summary

- Cheat sheet: Single feature, Multi-feature, Multi-class
- Logistic regression:
  - Multi-class prediction: Ensure that prediction is one of a discrete set of values
  - Cross-entropy: Measure difference between multi-class prediction and actual
  - Softmax: Ensure the multi-class prediction probability is a valid distribution (sum = 1)

Congrats!

You can now understand Google’s Tensorflow Beginner’s Tutorial

(https://www.tensorflow.org/versions/r0.7/tutorials/mnist/beginners/index.html)

References

- Perform ML with TF using multi-feature linear regression (the wrong way)
  - https://github.com/nethsix/gentle_tensorflow/blob/master/code/linear_regression_multi_feature_using_mini_batch_with_tensorboard.py
- Perform ML with TF using multi-feature linear regression
  - https://github.com/nethsix/gentle_tensorflow/blob/master/code/linear_regression_multi_feature_using_mini_batch_without_matrix_with_tensorboard.py
- Tensorflow official tutorial for character recognition
  - https://www.tensorflow.org/versions/r0.7/tutorials/mnist/beginners/index.html
- Colah's excellent explanation of cross-entropy
  - http://colah.github.io/posts/2015-09-Visual-Information/