

[Start Here](#)[Blog](#)[Books](#)[About](#)[Contact](#)

Need help with Deep Learning? [Take the FREE Mini-Course](#)

Develop Your First Neural Network in Python With Keras Step-By-Step

by **Jason Brownlee** on May 24, 2016 in **Deep Learning**



Keras is a powerful easy-to-use Python library for developing and evaluating [deep learning](#) models.

It wraps the efficient numerical computation libraries Theano and TensorFlow and allows you to define and train neural network models in a few short lines of code.

In this post, you will discover how to create your first neural network model in Python using Keras.

Let's get started.

- **Update Feb/2017:** Updated prediction example so rounding works in Python 2 and Python 3.

[Get Your Start in Machine Learning](#)

- **Update Mar/2017:** Updated example for Keras 2.0.2, TensorFlow 1.0.1 and Theano 0.9.0.



Develop Your First Neural Network in Python With Keras Step-By-Step
Photo by Phil Whitehouse, some rights reserved.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Tutorial Overview

There is not a lot of code required, but we are going to step over it slowly so that you will know how to create your own models in the future.

The steps you are going to cover in this tutorial are as follows:

Get Your Start in Machine Learning

1. Load Data.
2. Define Model.
3. Compile Model.
4. Fit Model.
5. Evaluate Model.
6. Tie It All Together.

This tutorial has a few requirements:

1. You have Python 2 or 3 installed and configured.
2. You have SciPy (including NumPy) installed and configured.
3. You have Keras and a backend (Theano or TensorFlow) installed and configured.

If you need help with your environment, see the tutorial:

- [How to Setup a Python Environment for Machine Learning and Deep Learning with Anaconda](#)

Create a new file called **keras_first_network.py** and type or copy-and-paste the code into the file and

Need help with Deep Learning in Python?

Take my free 2-week email course and discover MLPs, CNNs and LSTMs

Click to sign-up now and also get a free PDF Ebook version of the course!

Start Your FREE Mini-Course Now!

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

1. Load Data

Get Your Start in Machine Learning

Whenever we work with machine learning algorithms that use a stochastic process (e.g. random numbers), it is a good idea to set the random number seed.

This is so that you can run the same code again and again and get the same result. This is useful if you need to demonstrate a result, compare algorithms using the same source of randomness or to debug a part of your code.

You can initialize the random number generator with any seed you like, for example:

```
1 from keras.models import Sequential
2 from keras.layers import Dense
3 import numpy
4 # fix random seed for reproducibility
5 numpy.random.seed(7)
```

Now we can load our data.

In this tutorial, we are going to use the [Pima Indians onset of diabetes dataset](#). This is a standard machine learning repository. It describes patient medical record data for Pima Indians and whether they had

As such, it is a binary classification problem (onset of diabetes as 1 or not as 0). All of the input variables makes it easy to use directly with neural networks that expect numerical input and output values, and

Download the [Pima Indian dataset from the UCI Machine Learning repository](#) and place it in your local directory. Save it with the file name:

```
1 pima-indians-diabetes.csv
```

You can now load the file directly using the NumPy function `loadtxt()`. There are eight input variables. After we have loaded we can split the dataset into input variables (X) and the output class variable (Y).

```
1 # load pima indians dataset
2 dataset = numpy.loadtxt("pima-indians-diabetes.csv", delimiter=",")
3 # split into input (X) and output (Y) variables
4 X = dataset[:,0:8]
5 Y = dataset[:,8]
```

We have initialized our random number generator to ensure our results are reproducible and loaded our data. We are now ready to define our neural network model.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree**.

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

2. Define Model

Models in Keras are defined as a sequence of layers.

We create a Sequential model and add layers one at a time until we are happy with our network topology.

The first thing to get right is to ensure the input layer has the right number of inputs. This can be specified when creating the first layer with the **input_dim** argument and setting it to 8 for the 8 input variables.

How do we know the number of layers and their types?

This is a very hard question. There are heuristics that we can use and often the best network structure is found through a process of trial and error experimentation. Generally, you need a network large enough to capture the structure of the problem.

In this example, we will use a fully-connected network structure with three layers.

Fully connected layers are defined using the Dense class. We can specify the number of neurons in the layer as the second argument as **units** and specify the activation function using the **activation** argument.

In this case, we initialize the network weights to a small random number generated from a uniform distribution between 0 and 0.05 because that is the default uniform weight initialization in Keras. Another traditional alternative is to initialize weights generated from a Gaussian distribution.

We will use the **rectifier** ('relu') activation function on the first two layers and the sigmoid function on the output layer. In the past, sigmoid and tanh activation functions were preferred for all layers. These days, better performance is achieved using the rectifier and sigmoid on the output layer to ensure our network output is between 0 and 1 and easy to map to either class of classification of either class with a default threshold of 0.5.

We can piece it all together by adding each layer. The first layer has 12 neurons and expects 8 input variables. The second hidden layer has 8 neurons and finally, the output layer has 1 neuron to predict the class (onset of diabetes or not).

```
1 # create model
2 model = Sequential()
3 model.add(Dense(12, input_dim=8, activation='relu'))
4 model.add(Dense(8, activation='relu'))
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
5 model.add(Dense(1, activation='sigmoid'))
```

3. Compile Model

Now that the model is defined, we can compile it.

Compiling the model uses the efficient numerical libraries under the covers (the so-called backend) such as Theano or TensorFlow. The backend automatically chooses the best way to represent the network for training and making predictions to run on your hardware, such as CPU or GPU or even distributed.

When compiling, we must specify some additional properties required when training the network. Remember training a network means finding the best set of weights to make predictions for this problem.

We must specify the loss function to use to evaluate a set of weights, the optimizer used to search the space of weights, and optional metrics we would like to collect and report during training.

In this case, we will use logarithmic loss, which for a binary classification problem is defined in Keras as the efficient gradient descent algorithm “**adam**” for no other reason that it is an efficient default. Learn more about it in the paper “[Adam: A Method for Stochastic Optimization](#)”.

Finally, because it is a classification problem, we will collect and report the classification accuracy as

```
1 # Compile model
2 model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

4. Fit Model

We have defined our model and compiled it ready for efficient computation.

Now it is time to execute the model on some data.

We can train or fit our model on our loaded data by calling the **fit()** function on the model.

The training process will run for a fixed number of iterations through the dataset called epochs, that we can also set the number of instances that are evaluated before a weight update in the network is performed.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

batch_size argument.

For this problem, we will run for a small number of iterations (150) and use a relatively small batch size of 10. Again, these can be chosen experimentally by trial and error.

```
1 # Fit the model
2 model.fit(X, Y, epochs=150, batch_size=10)
```

This is where the work happens on your CPU or GPU.

5. Evaluate Model

We have trained our neural network on the entire dataset and we can evaluate the performance of the network on the same dataset.

This will only give us an idea of how well we have modeled the dataset (e.g. train accuracy), but no idea of how well the model performs on new data. We have done this for simplicity, but ideally, you could separate your data into train and test data.

You can evaluate your model on your training dataset using the **evaluate()** function on your model and the **model.evaluate()** function.

This will generate a prediction for each input and output pair and collect scores, including the average accuracy.

```
1 # evaluate the model
2 scores = model.evaluate(X, Y)
3 print("\n%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))
```

6. Tie It All Together

You have just seen how you can easily create your first neural network model in Keras.

Let's tie it all together into a complete code example.

```
1 # Create your first MLP in Keras
2 from keras.models import Sequential
3 from keras.layers import Dense
4 import numpy
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning


```

5 # fix random seed for reproducibility
6 numpy.random.seed(7)
7 # load pima indians dataset
8 dataset = numpy.loadtxt("pima-indians-diabetes.csv", delimiter=",")
9 # split into input (X) and output (Y) variables
10 X = dataset[:,0:8]
11 Y = dataset[:,8]
12 # create model
13 model = Sequential()
14 model.add(Dense(12, input_dim=8, activation='relu'))
15 model.add(Dense(8, activation='relu'))
16 model.add(Dense(1, activation='sigmoid'))
17 # Compile model
18 model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
19 # Fit the model
20 model.fit(X, Y, epochs=150, batch_size=10)
21 # evaluate the model
22 scores = model.evaluate(X, Y)
23 print("\n%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))

```

Running this example, you should see a message for each of the 150 epochs printing the loss and accuracy of the trained model on the training dataset.

It takes about 10 seconds to execute on my workstation running on the CPU with a Theano backend.

```

1 ...
2 Epoch 145/150
3 768/768 [=====] - 0s - loss: 0.5105 - acc: 0.7396
4 Epoch 146/150
5 768/768 [=====] - 0s - loss: 0.4900 - acc: 0.7591
6 Epoch 147/150
7 768/768 [=====] - 0s - loss: 0.4939 - acc: 0.7565
8 Epoch 148/150
9 768/768 [=====] - 0s - loss: 0.4766 - acc: 0.7773
10 Epoch 149/150
11 768/768 [=====] - 0s - loss: 0.4883 - acc: 0.7591
12 Epoch 150/150
13 768/768 [=====] - 0s - loss: 0.4827 - acc: 0.7656
14 32/768 [>.....] - ETA: 0s
15 acc: 78.26%

```

Note: If you try running this example in an IPython or Jupyter notebook you may get an error. The reason is the output progress bars during training. You can easily turn these off by setting **verbose=0** in the call to **model.fit()**.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

7. Bonus: Make Predictions

The number one question I get asked is:

“After I train my model, how can I use it to make predictions on new data?”

Great question.

We can adapt the above example and use it to generate predictions on the training dataset, pretending it is a new dataset we have not seen before.

Making predictions is as easy as calling **model.predict()**. We are using a sigmoid activation function on the output layer, so the predictions will be in the range between 0 and 1. We can easily convert them into a crisp binary prediction for this classification.

The complete example that makes predictions for each record in the training data is listed below.

```
1 # Create first network with Keras
2 from keras.models import Sequential
3 from keras.layers import Dense
4 import numpy
5 # fix random seed for reproducibility
6 seed = 7
7 numpy.random.seed(seed)
8 # load pima indians dataset
9 dataset = numpy.loadtxt("pima-indians-diabetes.csv", delimiter=",")
10 # split into input (X) and output (Y) variables
11 X = dataset[:,0:8]
12 Y = dataset[:,8]
13 # create model
14 model = Sequential()
15 model.add(Dense(12, input_dim=8, init='uniform', activation='relu'))
16 model.add(Dense(8, init='uniform', activation='relu'))
17 model.add(Dense(1, init='uniform', activation='sigmoid'))
18 # Compile model
19 model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
20 # Fit the model
21 model.fit(X, Y, epochs=150, batch_size=10, verbose=2)
22 # calculate predictions
23 predictions = model.predict(X)
24 # round predictions
25 rounded = [round(x[0]) for x in predictions]
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
26 print(rounded)
```

Running this modified example now prints the predictions for each input pattern. We could use these predictions directly in our application if needed.

```
1 [1.0, 0.0, 1.0, 0.0, 1.0, 0.0, 0.0, 1.0, 1.0, 0.0, 0.0, 1.0, 1.0, 1.0, 1.0, 0.0, 1.0, 0.0, 0.0, 0.0, 1.0, 0.0, 1.0, 0.0, 1.0, 1.0, 1.0, 0.
```

Summary

In this post, you discovered how to create your first neural network model using the powerful Keras Python library for deep learning.

Specifically, you learned the five key steps in using Keras to create a neural network or deep learning model, step-by-step including:

1. How to load data.
2. How to define neural network in Keras.
3. How to compile a Keras model using the efficient numerical backend.
4. How to train a model on data.
5. How to evaluate a model on data.

Do you have any questions about Keras or about this tutorial?

Ask your question in the comments and I will do my best to answer.

Related Tutorials

Are you looking for some more Deep Learning tutorials with Python and Keras?

Take a look at some of these:

- [5 Step Life-Cycle for Neural Network Models in Keras](#)
- [How to Grid Search Hyperparameters for Deep Learning Models in Python With Keras](#)
- [Time Series Prediction With Deep Learning in Keras](#)
- [Multi-Class Classification Tutorial with the Keras Deep Learning Library](#)
- [Regression Tutorial with the Keras Deep Learning Library in Python](#)

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Frustrated With Your Progress In Deep Learning?

What If You Could Develop A Network in Minutes

...with just a few lines of Python


Discover how in my new Ebook: [Deep Learning With Python](#)

It covers **self-study tutorials** and **end-to-end projects** on topics like:
Multilayer Perceptrons, Convolutional Nets and Recurrent Neural Nets, and more...

Finally Bring Deep Learning To Your Own Projects

Skip the Academics. Just

[Click to learn more](#)



About Jason Brownlee

Dr. Jason Brownlee is a husband, proud father, academic researcher, author, professional developer and a machine learning practitioner. He is dedicated to helping developers get started and get good at applied machine learning.

[Learn more.](#)

[View all posts by Jason Brownlee](#) →

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

431 Responses to *Develop Your First Neural Network in Python With Keras Step-By-Step*

Saurav May 27, 2016 at 11:08 pm #

REPLY ↩

The input layer doesn't have any activation function, but still activation="relu" is mentioned in the first layer of the model. Why?



Jason Brownlee May 28, 2016 at 6:32 am #

Hi Saurav,

The first layer in the network here is technically a hidden layer, hence it has an activation function.

sam Johnson December 21, 2016 at 2:44 am #

Why have you made it a hidden layer though? the input layer is not usually represented



Jason Brownlee December 21, 2016 at 8:41 am #

Hi sam,

Note this line:

```
1 model.add(Dense(12, input_dim=8, init='uniform', activation='relu'))
```

It does a few things.

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

REPLY ↩

Get Your Start in Machine Learning

It defines the input layer as having 8 inputs.

It defines a hidden layer with 12 neurons, connected to the input layer that use relu activation function.

It initializes all weights using a sample of uniform random numbers.

Does that help?

Pavidevi May 17, 2017 at 2:31 am #

Hi Jason,

U have used two different activation functions so how can we know which activation function fit the model?



Jason Brownlee May 17, 2017 at 8:38 am #

Sorry, I don't understand the question.

Marco Cheung August 23, 2017 at 12:51 am #

Hi Jason,

I am interested in deep learning and machine learning. You mentioned "It defines a hidden layer with 12 neurons, connected to the input layer that use relu activation function." I wonder how can we determine the number of neurons in the hidden layer for a given model?

Thanks a lot!!!



Jason Brownlee August 23, 2017 at 6:55 am #

Use trial and error. We cannot specify the "best" number of neurons analytically.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Ramzan Shahid November 10, 2017 at 4:32 am #

Sir, thanks for your tutorial. Would you like to make tutorial on stock Data Prediction through Neural Network Model and training this on any stock data. If you have on this so please share the link. Thanks



Jason Brownlee November 10, 2017 at 10:39 am #

I am reticent to post tutorials on stock market prediction given the random walk hypothesis of security prices:
<https://machinelearningmastery.com/gentle-introduction-random-walk-times-series-forec>

Geoff May 29, 2016 at 6:18 am #

Can you explain how to implement weight regularization into the layers?



Jason Brownlee June 15, 2016 at 5:50 am #

Yep, see here:
<http://keras.io/regularizers/>

KWC June 14, 2016 at 12:08 pm #

Import statements if others need them:

```
from keras.models import Sequential
from keras.layers import Dense, Activation
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

REPLY ↩

Get Your Start in Machine Learning



Jason Brownlee June 15, 2016 at 5:49 am #

REPLY ↩

Thanks.

I had them in Part 6, but I have also added them to Part 1.

Aakash Nain June 29, 2016 at 6:00 pm #

REPLY ↩

If there are 8 inputs for the first layer then why we have taken them as '12' in the following line :

```
model.add(Dense(12, input_dim=8, init='uniform', activation='relu'))
```



Jason Brownlee June 30, 2016 at 6:47 am #

Hi Aakash.

The input layer is defined by the input_dim parameter, here set to 8.

The first hidden layer has 12 neurons.

Joshua July 2, 2016 at 12:04 am #

I ran your program and i have an error:

ValueError: could not convert string to float:

what could be the reason for this, and how may I solve it.

thanks.

great post by the way.

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning



Jason Brownlee July 2, 2016 at 6:20 am #

REPLY ↩

It might be a copy-paste error. Perhaps try to copy and run the whole example listed in section 6?

cheikh brahim July 5, 2016 at 7:40 pm #

REPLY ↩

thank you for your simple and useful example.



Jason Brownlee July 6, 2016 at 6:22 am #

You're welcome cheikh.

Nikhil Thakur July 6, 2016 at 6:39 pm #

Hello Sir, I am trying to use Keras for NLP , specifically sentence classification. I have given the code to execute. I am using Pycharm IDE.

```
batch_size = 32
nb_filter = 250
filter_length = 3
nb_epoch = 2
pool_length = 2
output_dim = 5
hidden_dims = 250

# Build the model

model1 = Sequential()
```

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
model1.add(Convolution1D(nb_filter, filter_length, activation='relu', border_mode='valid',
input_shape=(len(embb_weights), dim), weights=[embb_weights]))

model1.add(Dense(hidden_dims))
model1.add(Dropout(0.2))
model1.add(Activation('relu'))

model1.add(MaxPooling1D(pool_length=pool_length))

model1.add(Dense(output_dim, activation='sigmoid'))

sgd = SGD(lr=0.1, decay=1e-6, momentum=0.9, nesterov=True)

model1.compile(loss='mean_squared_error',
optimizer=sgd,
metrics=['accuracy'])
```



Jason Brownlee July 7, 2016 at 7:31 am #

You may want a larger network. You may also want to use a standard repeating structure like

See this post on using a CNN:

<http://machinelearningmastery.com/handwritten-digit-recognition-using-convolutional-neural-network/>

Later, you may also want to try some stacked LSTMs.

Andre Norman July 15, 2016 at 10:40 am #

Hi Jason, thanks for the awesome example. Given that the accuracy of this model is 79.56%. From here on, what steps would you take to improve the accuracy?

Given my nascent understanding of Machine Learning, my initial approach would have been:

Implement forward propagation, then compute the cost function, then implement back propagation, use gradient checking to evaluate my network (disable after use), then use gradient descent.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

REPLY ↩

Get Your Start in Machine Learning

However, this approach seems arduous compared to using Keras. Thanks for your response.



Jason Brownlee July 15, 2016 at 10:52 am #

REPLY ↩

Hi Andre, indeed Keras makes working with neural nets so much easier. Fun even!

We may be maxing out on this problem, but here is some general advice for lifting performance.

- data prep – try lots of different views of the problem and see which is best at exposing the structure of the problem to the learning algorithm (data transforms, feature engineering, etc.)
- algorithm selection – try lots of algorithms and see which one or few are best on the problem (try on all views)
- algorithm tuning – tune well performing algorithms to get the most out of them (grid search or random search hyperparameter tuning)
- ensembles – combine predictions from multiple algorithms (stacking, boosting, bagging, etc.)

For neural nets, there are a lot of things to tune, I think there are big gains in trying different network in concert with training epochs and learning rate (bigger nets need more training).

I hope that helps as a start.

Andre Norman July 18, 2016 at 7:19 am #

Awesome! Thanks Jason =)



Jason Brownlee July 18, 2016 at 8:03 am #

You're welcome Andre.

REPLY ↩

quentin August 7, 2017 at 8:41 pm #

REPLY ↩

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Some interesting stuff here
<https://youtu.be/vq2nnJ4g6N0>



Jason Brownlee August 8, 2017 at 7:49 am #

REPLY ↩

Thanks for sharing. What did you like about it?

Romilly Cocking July 21, 2016 at 12:31 am #

REPLY ↩

Hi Jason, it's a great example but if anyone runs it in an IPython/Jupyter notebook they are likely to get an error. This is due to a known bug in IPython.

The solution is to set verbose=0 like this

```
# Fit the model
model.fit(X, Y, nb_epoch=40, batch_size=10, verbose=0)
```



Jason Brownlee July 21, 2016 at 5:36 am #

Great, thanks for sharing Romilly.

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Anirban July 23, 2016 at 10:20 pm #

REPLY ↩

Great example. Have a query though. How do I now give a input and get the output (0 or 1). Can you pls give the cmd for that.
Thanks

Get Your Start in Machine Learning



Jason Brownlee July 24, 2016 at 6:53 am #

REPLY ↩

You can call `model.predict()` to get predictions and round on each value to snap to a binary value.

For example, below is a complete example showing you how to round the predictions and print them to console.

```
1 # Create first network with Keras
2 from keras.models import Sequential
3 from keras.layers import Dense
4 import numpy
5 # fix random seed for reproducibility
6 seed = 7
7 numpy.random.seed(seed)
8 # load pima indians dataset
9 dataset = numpy.loadtxt("pima-indians-diabetes.csv", delimiter=",")
10 # split into input (X) and output (Y) variables
11 X = dataset[:,0:8]
12 Y = dataset[:,8]
13 # create model
14 model = Sequential()
15 model.add(Dense(12, input_dim=8, init='uniform', activation='relu'))
16 model.add(Dense(8, init='uniform', activation='relu'))
17 model.add(Dense(1, init='uniform', activation='sigmoid'))
18 # Compile model
19 model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
20 # Fit the model
21 model.fit(X, Y, nb_epoch=150, batch_size=10, verbose=2)
22 # calculate predictions
23 predictions = model.predict(X)
24 # round predictions
25 rounded = [round(x) for x in predictions]
26 print(rounded)
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Debanjan March 27, 2017 at 12:04 pm #

REPLY ↩

Hi, Why you are not using any test set? You are predicting from the training set , I think.

Get Your Start in Machine Learning



Jason Brownlee March 28, 2017 at 8:19 am #

REPLY ↩

Correct, it is just an example to get you started with Keras.

David June 26, 2017 at 12:24 am #

REPLY ↩

Jason, I'm not quite understanding how the predicted values ([1.0, 0.0, 1.0, 0.0, 1.0,...]) map to the real world problem. For instance, what does that first "1.0" in the results indicate?

I get that it's a prediction of 'true' for diabetes...but to which patient is it predicting that—the first in the list? So then the second result, "0.0," is the prediction for the second patient/row in the dataset?



Jason Brownlee June 26, 2017 at 6:08 am #

Remember the original file has 0 and 1 values in the final class column where 0 is r

We are predicting new values in this column.

We are making predictions for special rows, we pass in their medical info and predict the on
number of rows at a time.

Rachel June 28, 2017 at 8:28 pm #

REPLY ↩

Hi Jason,

Can I ask why you use the same data X you fit the model to do the prediction?

Fit the model

```
model.fit(X, Y, epochs = 150, batch_size = 10, verbose = 2)
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
# calculate predictions  
predictions = model.predict(X)
```

Rachel



Jason Brownlee June 29, 2017 at 6:34 am #

REPLY ↩

It is all I have at hand. X means data matrix.

Replace X in predict() with Xprime or whatever you like.

Anirban July 23, 2016 at 10:52 pm #

I am not able to get to the last epoch. Getting error before that:

Epoch 11/150

390/768 [=====>.....]Traceback (most recent call last):.6921

ValueError: I/O operation on closed file

I could resolve this by varying the epoch and batch size.

Now to predict a unknown value, i loaded a new dataset and used predict cmd as below :

dataset_test = numpy.loadtxt("pima-indians-diabetes_test.csv",delimiter=",") –has only one row

```
X = dataset_test[:,0:8]
```

```
model.predict(X)
```

But I am getting error :

```
X = dataset_test[:,0:8]
```

IndexError: too many indices for array

Can you help pls.

Thanks

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning



Jason Brownlee July 24, 2016 at 6:55 am #

REPLY ↩

I see problems like this when you run from a notebook or from an IDE.

Consider running examples from the console to ensure they work.

Consider tuning off verbose output (verbose=0 in the call to fit()) to disable the progress bar.

David Kluszczynski July 28, 2016 at 12:42 am #

REPLY ↩

Hi Jason!

Loved the tutorial! I have a question however.

Is there a way to save the weights to a file after the model is trained for uses, such as kaggle?

Thanks,

David



Jason Brownlee July 28, 2016 at 5:47 am #

Thanks David.

You can save the network weights to file by calling `model.save_weights("model.h5")`

You can learn more in this post:

<http://machinelearningmastery.com/save-load-keras-deep-learning-models/>

Alex Hopper July 29, 2016 at 5:45 am #

REPLY ↩

Hey, Jason! Thank you for the awesome tutorial! I've use your tutorial to learn about CNN. I have Keras to classificate images and I have 3 or more classes to classify, How could my algorithm know about

Get Your Start in Machine Learning

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

a dog and a horse. Is there any way to code this? I've tried it:

```
target_names = ['class 0(Cats)', 'class 1(Dogs)', 'class 2(Horse)']  
print(classification_report(np.argmax(Y_test,axis=1), y_pred,target_names=target_names))
```

But my results are not classifying correctly.

```
precision recall f1-score support  
class 0(Cat) 0.00 0.00 0.00 17  
class 1(Dog) 0.00 0.00 0.00 14  
class 2(Horse) 0.99 1.00 0.99 2526  
  
avg / total 0.98 0.99 0.98 2557
```



Jason Brownlee July 29, 2016 at 6:41 am #

Great question Alex.

This is an example of a multi-class classification problem. You must use a one hot encoding on the class labels, build a neural network and specify the number of classes as the number of outputs on the final layer of your network.

I provide a tutorial with the famous iris dataset that has 3 output classes here:

<http://machinelearningmastery.com/multi-class-classification-tutorial-keras-deep-learning-library/>

Alex Hopper August 1, 2016 at 1:22 am #

Thank you.

I'll check it.



Jason Brownlee August 1, 2016 at 6:25 am #

No problem Alex.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

REPLY ↩

Get Your Start in Machine Learning

Anonymouse August 2, 2016 at 11:28 pm #

REPLY ↩

This was really useful, thank you

I'm using keras (with CNNs) for sentiment classification of documents and I'd like to improve the performance, but I'm completely at a loss when it comes to tuning the parameters in a non-arbitrary way. Could you maybe point me somewhere that will help me go about this in a more systematic fashion? There must be some heuristics or rules-of-thumb that could guide me.



Jason Brownlee August 3, 2016 at 8:09 am #

I have a tutorial coming out soon (next week) that provide lots of examples of tuning the hyperparameters, but it's currently limited to MLPs.

For CNNs, I would advise tuning the number of repeating layers (conv + max pool), the number of filters, and the number of dense layers at the predicting part of your network. Also consider using some fixed layers from pre-trained models and try just training some input and output layers around it for your problem.

I hope that helps as a start.

Shopon August 14, 2016 at 5:04 pm #

Hello Jason , My Accuracy is : 0.0104 , but yours is 0.7879 and my loss is : -9.5414 . Is there any problem with the dataset ? I downloaded the dataset from a different site .



Jason Brownlee August 15, 2016 at 12:36 pm #

REPLY ↩

I think there might be something wrong with your implementation or your dataset. Your num

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

mohamed August 15, 2016 at 9:30 am #

REPLY ↩

after training, how i can use the trained model on new sample



Jason Brownlee August 15, 2016 at 12:36 pm #

REPLY ↩

You can call `model.predict()`

See an above comment for a specific code example.

Omachi Okolo August 16, 2016 at 10:21 pm #

Hi Jason,

i'm a student conducting a research on how to use artificial neural network to predict the business viability. I intend to use python as a programming language. The application of ANN fascinates me but i'm new to how to go about this.

Many thanks

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee August 17, 2016 at 9:51 am #

Consider getting a good grounding in how to work through a machine learning problem end to end in python first.

Here is a good tutorial to get you started:

<http://machinelearningmastery.com/machine-learning-in-python-step-by-step/>

Agni August 17, 2016 at 6:23 am #

Get Your Start in Machine Learning

Dear Jeson, this is a great tutorial for beginners. It will satisfy the need of many students who are looking for the initial help. But I have a question. Could you please light on a few things: i) how to test the trained model using test dataset (i.e., loading of test dataset and applied the model and suppose the test file name is test.csv) ii) print the accuracy obtained on test dataset iii) the o/p has more than 2 class (suppose 4-class classification problem).

Please show the whole program to overcome any confusion.

Thanks a lot.



Jason Brownlee August 17, 2016 at 10:03 am #

REPLY ↩

I provide an example elsewhere in the comments, you can also see how to make predictions on new data in this post:

<http://machinelearningmastery.com/5-step-life-cycle-neural-network-models-keras/>

For an example of multi-class classification, you can see this tutorial:

<http://machinelearningmastery.com/multi-class-classification-tutorial-keras-deep-learning-library/>

Doron Vetlzer August 17, 2016 at 9:29 am #

I am trying to build a Neural Network with some recursive connections but not a full recursive la

Doron Vetlzer August 17, 2016 at 9:31 am #

I could print a diagram of the network but what I want Basically is that each neuron in the current time frame to know only its own previous output and not the output of all the neurons in the output layer.



Jason Brownlee August 17, 2016 at 10:04 am #

REPLY ↩

I don't know off hand Doron.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Doron Veltzer August 23, 2016 at 2:28 am #

REPLY ↩

Thanks for replying though, have a good day.

sairam August 30, 2016 at 8:49 am #

REPLY ↩

Hello Jason,

This is a great tutorial . Thanks for sharing.

I am having a dataset of 100 finger prints and i want to extract minutiae of 100 finger prints using python really confused.



Jason Brownlee August 31, 2016 at 8:43 am #

If your fingerprints are images, you may want to consider using convolutional neural network data.

See this tutorial on digit recognition for a start:

<http://machinelearningmastery.com/handwritten-digit-recognition-using-convolutional-neural-network>

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

padmashri July 6, 2017 at 10:12 pm #

REPLY ↩

Hi Jason

Thanks for this great tutorial, i am new to machine learning i went through your basic tutorial on keras and also handwritten-digit-recognition. I would like to understand how i can train a set of image data, for eg. the set of image data can be some thing like square, circle, pyramid. pl. let me know how the input data needs to fed to the program and how we need to export the model.

Get Your Start in Machine Learning



Jason Brownlee July 9, 2017 at 10:30 am #

REPLY ↩

Start by preparing a high-quality dataset.

CM September 1, 2016 at 4:23 pm #

REPLY ↩

Hi Jason,

Thanks for the great article. But I had 1 query.

Are there any inbuilt functions in keras that can give me the feature importance for the ANN model?

If not, can you suggest a technique I can use to extract variable importance from the loss function? I am which involves permuting the values of the selected variable and calculating the relative increase in loss

Regards,

CM



Jason Brownlee September 2, 2016 at 8:07 am #

I don't believe so CM.

I would suggest using a wrapper method and evaluate subsets of features to develop a feature importance

I talk a lot more about feature selection in this post:

<http://machinelearningmastery.com/an-introduction-to-feature-selection/>

I provide an example of feature selection in scikit-learn here:

<http://machinelearningmastery.com/feature-selection-machine-learning-python/>

I hope that helps as a start.

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Minesh Jethva May 15, 2017 at 7:49 pm #

REPLY ↩

have you develop any progress for this approach? I also have same problem.

Kamal September 7, 2016 at 2:09 am #

REPLY ↩

Dear Jason, I am new to Deep learning. Being a novice, I am asking you a technical question which may seem silly. My question is that- can we use features (for example length of the sentence etc.) of a sentence while classifying a sentence (suppose the o/p are +ve sentence and -ve sentence) using deep neural network?



Jason Brownlee September 7, 2016 at 10:27 am #

Great question Kamal, yes you can. I would encourage you to include all such features and

Saurabh September 11, 2016 at 12:42 pm #

Hi, How would I use this on a dataset that has multiple outputs? For example a dataset with out
or 4 ?



Jason Brownlee September 12, 2016 at 8:30 am #

REPLY ↩

You could use two neurons in the output layer and normalize the output variables to both be in the range of 0 to 1.

This tutorial on multi-class classification might give you some ideas:

<http://machinelearningmastery.com/multi-class-classification-tutorial-keras-deep-learning-library/>

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Tom_P September 17, 2016 at 1:47 pm #

REPLY ↩

Hi Jason,

The tutorial looks really good but unfortunately I keep getting an error when importing Dense from keras.layers, I get the error : AttributeError: module 'theano' has no attribute 'gof'

I have tried reinstalling Theano but it has not fixed the issue.

Best wishes

Tom



Jason Brownlee September 18, 2016 at 7:57 am #

Hi Tom, sorry to hear that. I have not seen this problem before.

Have you searched google? I can see a few posts and it might be related to your version of scipy or

Let me know how you go.

shudhan September 21, 2016 at 5:54 pm #

Hey Jason,

Can you please make a tutorial on how to add additional train data into the already trained model? This warm start is used for random forest. But not sure how to implement as algorithm. A generalised version of how to implement would be great. Thank you.



Jason Brownlee September 22, 2016 at 8:08 am #

REPLY ↩

Great question Shudhan!

Yes, you could save your weights, load them later into a new network topology and start training on

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

I'll work out an example in coming weeks, time permitting.

Joanna September 22, 2016 at 1:09 am #

REPLY ↩

Hi Jason,
first of all congratulations for this amazing work that you have done!
Here is my question:
What about if my .csv file includes also both nominal and numerical attributes?
Should I change my nominal values to numerical?

Thank you in advance



Jason Brownlee September 22, 2016 at 8:19 am #

Hi Joanna, yes.

You can use a label encoder to convert nominal to integer, and then even convert the integer to one

This post will give you code you can use:

<http://machinelearningmastery.com/data-preparation-gradient-boosting-xgboost-python/>

ATM October 2, 2016 at 5:47 am #

A small bug:-

Line 25 : rounded = [round(x) for x in predictions]

should have numpy.round instead, for the code to run!

Great tutorial, regardless. The best i've seen for intro to ANN in python. Thanks!

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**
Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee October 2, 2016 at 8:20 am #

REPLY ↩

Perhaps it's your version of Python or environment?

In Python 2.7 the round() function is built-in.

AC January 14, 2017 at 2:11 am #

REPLY ↩

If there is comment for python3, should be better.
#use unmpy.round instead, if using python3,



Jason Brownlee January 15, 2017 at 5:24 am #

Thanks for the note AC.

Ash October 9, 2016 at 1:36 am #

This is simple to grasp! Great post! How can we perform dropout in keras?



Jason Brownlee October 9, 2016 at 6:49 am #

REPLY ↩

Thanks Ash.

You can learn about drop out with Keras here:

<http://machinelearningmastery.com/dropout-regularization-deep-learning-models-keras/>

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Homagni Saha October 14, 2016 at 4:15 am #

REPLY ↩

Hello Jason,

You are using `model.predict` in the end to predict the results. Is it possible to save the model somewhere in the harddisk and transfer it to another machine(turtlebot running on ROS for my instance) and then use the model directly on turtlebot to predict the results?

Please tell me how

Thanking you

Homagni Saha



Jason Brownlee October 14, 2016 at 9:07 am #

REPLY ↩

Hi Homagni, great question.

Absolutely!

Learn exactly how in this tutorial I wrote:

<http://machinelearningmastery.com/save-load-keras-deep-learning-models/>

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree**. Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Rimi October 16, 2016 at 8:21 pm #

Hi Jason,

I implemented you code to begin with. But I am getting an accuracy of 45.18% with the same parameter

Cant figure out why.

Thanks



Jason Brownlee October 17, 2016 at 10:29 am #

REPLY ↩

There does sound like a problem there Rimi.

Confirm the code and data match exactly.

Get Your Start in Machine Learning

Ankit October 26, 2016 at 8:12 pm #

REPLY ↩

Hi Jason,

I am little confused with first layer parameters. You said that first layer has 12 neurons and expects 8 input variables.

Why there is a difference between number of neurons, input_dim for first layer.

Regards,

Ankit



Jason Brownlee October 27, 2016 at 7:45 am #

Hi Ankit,

The problem has 8 input variables and the first hidden layer has 12 neurons. Inputs are the columns are whatever we design based on whatever capacity we think we need to represent the complexity of the data. 12 neurons for the first hidden layer.

I hope that is clearer.

Tom October 27, 2016 at 3:04 am #

Hi,

I have a data , IRIS like data but with more colmunns.

I want to use MLP and DBN/CNNClassifier (or any other Deep Learning classificaiton algorithm) on my data to see how correctly it does classified into 6 groups.

Previously using DEEP LEARNING FOR J, today first time see KERAS.

does KERAS has examples (code examples) of DL Classification algorithms?

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Kindly,
Tom



Jason Brownlee October 27, 2016 at 7:48 am #

REPLY ↩

Yes Tom, the example in this post is an example of a neural network (deep learning) applied to a classification problem.

Rumesa October 30, 2016 at 1:57 am #

REPLY ↩

I have installed theano but it gives me the error of tensorflow.is it mendatory to install both pack
wndows.the only way to get it on windows is to install virtual machine



Jason Brownlee October 30, 2016 at 8:57 am #

Keras will work just fine with Theano.

Just install Theano, and configure Keras to use the Theano backend.

More information about configuring the Keras backend here:

<http://machinelearningmastery.com/introduction-python-deep-learning-library-keras/>

Rumesa October 31, 2016 at 4:36 am #

REPLY ↩

hey jason I have run your code but got the following error.Although I have already installed theano backend.help me out.I just stuck.

Using TensorFlow backend.

Traceback (most recent call last):

File "C:\Users\pc\Desktop\first.py", line 2, in

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning


```
from keras.models import Sequential
File "C:\Users\pc\Anaconda3\lib\site-packages\keras\__init__.py", line 2, in
from . import backend
File "C:\Users\pc\Anaconda3\lib\site-packages\keras\backend\__init__.py", line 64, in
from .tensorflow_backend import *
File "C:\Users\pc\Anaconda3\lib\site-packages\keras\backend\tensorflow_backend.py", line 1, in
import tensorflow as tf
ImportError: No module named 'tensorflow'
>>>
```



Jason Brownlee October 31, 2016 at 5:34 am #

Change the backend used by Keras from TensorFlow to Theano.

You can do this either by using the command line switch or changing the Keras config file.

See the link I posted in the previous post for instructions.

Maria January 6, 2017 at 1:05 pm #

Hello Rumesa!

Have you solved your problem? I have the same one. Everywhere is the same answer with keras.js you tell me what have worked for you?



Jason Brownlee January 7, 2017 at 8:20 am #

Interesting.

Maybe there is an issue with the latest version and a tight coupling to tensorflow? I have not seen this myself

Perhaps it might be worth testing prior versions of Keras, such as 1.1.0?

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

REPLY ↩

Get Your Start in Machine Learning

Try this:

```
1 pip install --upgrade --no-deps keras==1.1.0
```

Alexon November 1, 2016 at 6:54 am #

REPLY ↩

Hi Jason,

First off, thanks so much for creating these resources, I have been keeping an eye on your newsletter for a while now, and I finally have the free time to start learning more about it myself, so your work has been really appreciated.

My question is: How can I set/get the weights of each hidden node?

I am planning to create several arrays randomized weights, then use a genetic algorithm to see which weights perform best over several generations. How would be the best way to go about this, and if I use a “relu” activation function, am I right that the weights should be between 0 and 0.05?

Many thanks for your help 😊

Alexon



Jason Brownlee November 1, 2016 at 8:05 am #

Thanks Alexon,

You can get and set the weights from a network.

You can learn more about how to do this in the context of saving the weights to file here:

<http://machinelearningmastery.com/save-load-keras-deep-learning-models/>

I hope that helps as a start, I'd love to hear how you go.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Alexon November 6, 2016 at 6:36 am #

Get Your Start in Machine Learning

Thats great, thanks for pointing me in the right direction.

I'd be happy to let you know how it goes, but might take a while as this is very much a "when I can find the time" project between jobs 😊

Cheers!

Arnaldo Gunzi November 2, 2016 at 10:17 pm #

REPLY ↩

Nice introduction, thanks!



Jason Brownlee November 3, 2016 at 7:59 am #

I'm glad you found it useful Arnaldo.

Abbey November 14, 2016 at 11:05 pm #

Good day

I have a question, how can I represent a character as a vector that could be an input for the neural network LSTM

For instance, I have bf to predict boy friend or best friend and similarly I have 2mor to predict tomorrow. represented as vector, so that it can be train with RNN/LSTM to predict the output.

Thank you.

Kind Regards

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee November 15, 2016 at 7:54 am #

REPLY ↩

Get Your Start in Machine Learning

Hi Abbey, You can map characters to integers to get integer vectors.

Abbey November 15, 2016 at 6:17 pm #

REPLY ↩

Thank you Jason, if i map characters to integers value to get vectors using English Alphabets, numbers and special characters

The question is how will LSTM predict the character. Please example in more details for me.

Regards



Jason Brownlee November 16, 2016 at 9:27 am #

Hi Abbey,

If your output values are also characters, you can map them onto integers, and reverse the

Abbey November 16, 2016 at 8:39 pm #

The output value of the characters encoding will be text

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Abbey November 15, 2016 at 6:22 pm #

REPLY ↩

Thank you, Jason, if I map characters to integers value to get vectors representation of the informal text using English Alphabets, numbers and special characters

The question is how will LSTM predict the character or words that have close meaning to the input value. Please example in more details for me. I understand how RNN/LSTM work based on your tutorial example but the logic in designing processing is what I am stress with.

Regards

Get Your Start in Machine Learning

Ammar November 27, 2016 at 10:35 am #

REPLY ↩

hi Jason,

i am trying to implement CNN one dimension on my data. so, i bluit my network.

the issue is:

```
def train_model(model, X_train, y_train, X_test, y_test):
```

```
X_train = X_train.reshape(-1, 1, 41)
```

```
X_test = X_test.reshape(-1, 1, 41)
```

```
numpy.random.seed(seed)
```

```
model.fit(X_train, y_train, validation_data=(X_test, y_test), nb_epoch=100, batch_size=64)
```

```
# Final evaluation of the model
```

```
scores = model.evaluate(X_test, y_test, verbose=0)
```

```
print("Accuracy: %.2f%%" % (scores[1] * 100))
```

this method above does not work and does not give me any error message.

could you help me with this please?



Jason Brownlee November 28, 2016 at 8:40 am #

Hi Ammar, I'm surprised that there is no error message.

Perhaps run from the command line and add some print() statements to see exactly where it stops.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

KK November 28, 2016 at 6:55 pm #

REPLY ↩

Hi Jason

Great work. I have another doubt. How can we apply this to text mining. I have a csv file containing review document and label. I want to apply classify the documents based on the text available. Can U do this favor.

Get Your Start in Machine Learning



Jason Brownlee November 29, 2016 at 8:48 am #

REPLY ↩

I would recommend converting the chars to ints and then using an Embedding layer.

Alex M November 30, 2016 at 10:52 pm #

REPLY ↩

Mr Jason, this is great tutorial but I am stack with some errors.

First I can't load data set correctly, tried to correct error but can't make it. (FileNotFoundError: [Errno 2] No such file or directory: 'pima-indians-diabetes.csv').

Second: While trying to evaluate the model it says (X is not defined) May be this is because uploading failed

Thanks!



Jason Brownlee December 1, 2016 at 7:29 am #

You need to download the file and place it in your current working directory Alex.

Does that help?

Alex M December 1, 2016 at 6:45 pm #

Sir, it is now successful....

Thanks!



Jason Brownlee December 2, 2016 at 8:15 am #

REPLY ↩

Glad to hear it Alex.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Bappaditya December 2, 2016 at 7:35 pm #

REPLY ↩

Hi Jason,

First of all a special thanks to you for providing such a great tutorial. I am very new to machine learning and truly speaking i had no background in data science. The concept of ML overwhelmed me and now i have a desire to be an expert of this field. I need your advice to start from a scratch. Also i am a PhD student in Computer Engineering (computer hardware)and i want to apply it as a tool for fault detection and testing for ICs.Can you provide me some references on this field?



Jason Brownlee December 3, 2016 at 8:29 am #

Hi Bappaditya,

My best advice for getting started is here:

<http://machinelearningmastery.com/start-here/#getstarted>

I believe machine learning and deep learning are good tools for use on problems in fault detection. A

<http://scholar.google.com>

Best of luck with your project.

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Alex M December 3, 2016 at 8:00 pm #

Well as usual in our daily coding life errors happen, now I have this error how can I correct it? Thanks!

```
” _____  
NoBackendError Traceback (most recent call last)  
in ()  
16 import librosa.display  
17 audio_path = ('/Users/MA/Python Notebook/OK.mp3')  
—> 18 y, sr = librosa.load(audio_path)
```

```
C:\Users\MA\Anaconda3\lib\site-packages\librosa\core\audio.py in load(path, sr, mono, offset, duration, dtype)
```

```
107
```

```
108 y = []
```

```
→ 109 with audioread.audio_open(os.path.realpath(path)) as input_file:
```

```
110 sr_native = input_file.samplerate
```

```
111 n_channels = input_file.channels
```

```
C:\Users\MA\Anaconda3\lib\site-packages\audioread\__init__.py in audio_open(path)
```

```
112
```

```
113 # All backends failed!
```

```
→ 114 raise NoBackendError()
```

NoBackendError:

”

That is the error I am getting just when trying to load a song into librosa...

Thanks!! @Jason Brownlee



Jason Brownlee December 4, 2016 at 5:30 am #

Sorry, this looks like an issue with your librosa library, not a machine learning issue. I can't

Alex M December 4, 2016 at 10:30 pm #

Thanks I have managed to correct the error...

Happy Sunday to you all.....



Jason Brownlee December 5, 2016 at 6:49 am #

REPLY ↩

Get Your Start in Machine Learning

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Glad to hear it Alex.

Lei December 4, 2016 at 10:52 pm #

REPLY ↩

Hi, Jason, thank you for your amazing examples.

I run the same code on my laptop. But I did not get the same results. What could be the possible reasons?

I am using windows 8.1 64bit+eclipse+anaconda 4.2+theano 0.9.4+CUDA7.5

I got results like follows.

... ..

Epoch 145/150

```
10/768 [.....] - ETA: 0s - loss: 0.3634 - acc: 0.8000
80/768 [==>.....] - ETA: 0s - loss: 0.4066 - acc: 0.7750
150/768 [====>.....] - ETA: 0s - loss: 0.4059 - acc: 0.8067
220/768 [=====>.....] - ETA: 0s - loss: 0.4047 - acc: 0.8091
300/768 [=====>.....] - ETA: 0s - loss: 0.4498 - acc: 0.7867
380/768 [=====>.....] - ETA: 0s - loss: 0.4595 - acc: 0.7895
450/768 [=====>.....] - ETA: 0s - loss: 0.4568 - acc: 0.7911
510/768 [=====>.....] - ETA: 0s - loss: 0.4553 - acc: 0.7882
580/768 [=====>.....] - ETA: 0s - loss: 0.4677 - acc: 0.7776
660/768 [=====>.....] - ETA: 0s - loss: 0.4697 - acc: 0.7788
740/768 [=====>.....] - ETA: 0s - loss: 0.4611 - acc: 0.7838
768/768 [=====>.....] - 0s - loss: 0.4614 - acc: 0.7799
```

Epoch 146/150

```
10/768 [.....] - ETA: 0s - loss: 0.3846 - acc: 0.8000
90/768 [==>.....] - ETA: 0s - loss: 0.5079 - acc: 0.7444
170/768 [====>.....] - ETA: 0s - loss: 0.4500 - acc: 0.7882
250/768 [=====>.....] - ETA: 0s - loss: 0.4594 - acc: 0.7840
330/768 [=====>.....] - ETA: 0s - loss: 0.4574 - acc: 0.7818
400/768 [=====>.....] - ETA: 0s - loss: 0.4563 - acc: 0.7775
470/768 [=====>.....] - ETA: 0s - loss: 0.4654 - acc: 0.7723
```

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
540/768 [=====>.....] - ETA: 0s - loss: 0.4537 - acc: 0.7870
620/768 [=====>.....] - ETA: 0s - loss: 0.4615 - acc: 0.7806
690/768 [=====>....] - ETA: 0s - loss: 0.4631 - acc: 0.7739
750/768 [=====>.] - ETA: 0s - loss: 0.4649 - acc: 0.7733
768/768 [=====] - 0s - loss: 0.4636 - acc: 0.7734
Epoch 147/150
```

```
10/768 [.....] - ETA: 0s - loss: 0.3561 - acc: 0.9000
90/768 [==>.....] - ETA: 0s - loss: 0.4167 - acc: 0.8556
170/768 [====>.....] - ETA: 0s - loss: 0.4824 - acc: 0.8059
250/768 [=====>.....] - ETA: 0s - loss: 0.4534 - acc: 0.8080
330/768 [=====>.....] - ETA: 0s - loss: 0.4679 - acc: 0.7848
400/768 [=====>.....] - ETA: 0s - loss: 0.4590 - acc: 0.7950
460/768 [=====>.....] - ETA: 0s - loss: 0.4619 - acc: 0.7913
530/768 [=====>.....] - ETA: 0s - loss: 0.4562 - acc: 0.7868
600/768 [=====>.....] - ETA: 0s - loss: 0.4497 - acc: 0.7883
680/768 [=====>....] - ETA: 0s - loss: 0.4525 - acc: 0.7853
760/768 [=====>.] - ETA: 0s - loss: 0.4568 - acc: 0.7803
768/768 [=====] - 0s - loss: 0.4561 - acc: 0.7812
Epoch 148/150
```

```
10/768 [.....] - ETA: 0s - loss: 0.4183 - acc: 0.9000
80/768 [==>.....] - ETA: 0s - loss: 0.3674 - acc: 0.8750
160/768 [====>.....] - ETA: 0s - loss: 0.4340 - acc: 0.8250
240/768 [=====>.....] - ETA: 0s - loss: 0.4799 - acc: 0.7583
320/768 [=====>.....] - ETA: 0s - loss: 0.4648 - acc: 0.7719
400/768 [=====>.....] - ETA: 0s - loss: 0.4596 - acc: 0.7775
470/768 [=====>.....] - ETA: 0s - loss: 0.4475 - acc: 0.7809
540/768 [=====>.....] - ETA: 0s - loss: 0.4545 - acc: 0.7778
620/768 [=====>.....] - ETA: 0s - loss: 0.4590 - acc: 0.7742
690/768 [=====>....] - ETA: 0s - loss: 0.4769 - acc: 0.7652
760/768 [=====>.] - ETA: 0s - loss: 0.4748 - acc: 0.7658
768/768 [=====] - 0s - loss: 0.4734 - acc: 0.7669
Epoch 149/150
```

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```

10/768 [.....] - ETA: 0s - loss: 0.3043 - acc: 0.9000
90/768 [==>.....] - ETA: 0s - loss: 0.4913 - acc: 0.7111
170/768 [=====>.....] - ETA: 0s - loss: 0.4779 - acc: 0.7588
250/768 [=====>.....] - ETA: 0s - loss: 0.4794 - acc: 0.7640
320/768 [=====>.....] - ETA: 0s - loss: 0.4957 - acc: 0.7562
370/768 [=====>.....] - ETA: 0s - loss: 0.4891 - acc: 0.7703
450/768 [=====>.....] - ETA: 0s - loss: 0.4737 - acc: 0.7867
520/768 [=====>.....] - ETA: 0s - loss: 0.4675 - acc: 0.7865
600/768 [=====>.....] - ETA: 0s - loss: 0.4668 - acc: 0.7833
680/768 [=====>.....] - ETA: 0s - loss: 0.4677 - acc: 0.7809
760/768 [=====>.....] - ETA: 0s - loss: 0.4648 - acc: 0.7803
768/768 [=====>.....] - 0s - loss: 0.4625 - acc: 0.7826
Epoch 150/150

```

```

10/768 [.....] - ETA: 0s - loss: 0.2751 - acc: 1.0000
100/768 [==>.....] - ETA: 0s - loss: 0.4501 - acc: 0.8100
170/768 [=====>.....] - ETA: 0s - loss: 0.4588 - acc: 0.8059
250/768 [=====>.....] - ETA: 0s - loss: 0.4299 - acc: 0.8200
310/768 [=====>.....] - ETA: 0s - loss: 0.4298 - acc: 0.8129
380/768 [=====>.....] - ETA: 0s - loss: 0.4365 - acc: 0.8053
460/768 [=====>.....] - ETA: 0s - loss: 0.4469 - acc: 0.7957
540/768 [=====>.....] - ETA: 0s - loss: 0.4436 - acc: 0.8000
620/768 [=====>.....] - ETA: 0s - loss: 0.4570 - acc: 0.7871
690/768 [=====>.....] - ETA: 0s - loss: 0.4664 - acc: 0.7783
760/768 [=====>.....] - ETA: 0s - loss: 0.4617 - acc: 0.7789
768/768 [=====>.....] - 0s - loss: 0.4638 - acc: 0.7773

32/768 [>.....] - ETA: 0s
448/768 [=====>.....] - ETA: 0sacc: 79.69%

```

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee December 5, 2016 at 6:50 am #

REPLY ↩

There is randomness in the learning process that we cannot control for yet.

Get Your Start in Machine Learning

See this post:

<http://machinelearningmastery.com/randomness-in-machine-learning/>

Nanya December 10, 2016 at 2:55 pm #

REPLY ↩

Hello Jason Brownlee, Thx for sharing~

I'm new in deep learning. And I am wondering can what you discussed here: "Keras" be used to build a CNN in tensorflow and train some csv files for classification. May be this is a stupid question, but waiting for your reply. I'm working on my graduation project for Word sense disambiguation with CNN, and just can't move on. Hope for your help~ Best wishes!



Jason Brownlee December 11, 2016 at 5:22 am #

Sorry Nanya, I'm not sure I understand your question. Are you able to rephrase it?

Anon December 16, 2016 at 12:51 am #

I've just installed Anaconda with Keras and am using python 3.5.

It seems there's an error with the rounding using Py3 as opposed to Py2. I think it's because of this character

I removed the rounding and just used print(predictions) and it seemed to work outputting floats instead.

Does this look correct?

...

Epoch 150/150

0s - loss: 0.4593 - acc: 0.7839

[[0.79361773]

[0.10443526]

[0.90862554]

...,

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
[ 0.33652252]
[ 0.63745886]
[ 0.11704451]]
```



Jason Brownlee December 16, 2016 at 5:44 am #

REPLY ↩

Nice, it does look good!

Florin Claudiu Mihalache December 19, 2016 at 2:37 am #

Hi Jason Brownlee

I tried to modified your exemple for my problem (Letter Recognition ,<http://archive.ics.uci.edu/ml/dataset>
My data set look like <http://archive.ics.uci.edu/ml/machine-learning-databases/letter-recognition/letter-rec>
try to split the data in input and ouput like this :

```
X = dataset[:,1:17]
```

```
Y = dataset[:,0]
```

but a have some error (something related that strings are not recognized) .

I tried to modified each letter whit the ASCII code (A became 65 and so on).The string error disappeared

The program compiles now but the output look like this :

```
17445/20000 [=====>....] - ETA: 0s - loss: -1219.4768 - acc:0.0000e+00
17605/20000 [=====>....] - ETA: 0s - loss: -1219.4706 - acc:0.0000e+00
17730/20000 [=====>....] - ETA: 0s - loss: -1219.4566 - acc:0.0000e+00
17890/20000 [=====>....] - ETA: 0s - loss: -1219.4071 - acc:0.0000e+00
18050/20000 [=====>...] - ETA: 0s - loss: -1219.4599 - acc:0.0000e+00
18175/20000 [=====>...] - ETA: 0s - loss: -1219.3972 - acc:0.0000e+00
18335/20000 [=====>...] - ETA: 0s - loss: -1219.4642 - acc:0.0000e+00
18495/20000 [=====>...] - ETA: 0s - loss: -1219.5032 - acc:0.0000e+00
18620/20000 [=====>...] - ETA: 0s - loss: -1219.4391 - acc:0.0000e+00
18780/20000 [=====>..] - ETA: 0s - loss: -1219.5652 - acc:0.0000e+00
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
18940/20000 [=====>..] - ETA: 0s - loss: -1219.5520 - acc:0.0000e+00
19080/20000 [=====>..] - ETA: 0s - loss: -1219.5381 - acc:0.0000e+00
19225/20000 [=====>..] - ETA: 0s - loss: -1219.5182 - acc:0.0000e+00
19385/20000 [=====>..] - ETA: 0s - loss: -1219.6742 - acc:0.0000e+00
19535/20000 [=====>..] - ETA: 0s - loss: -1219.7030 - acc:0.0000e+00
19670/20000 [=====>..] - ETA: 0s - loss: -1219.7634 - acc:0.0000e+00
19830/20000 [=====>..] - ETA: 0s - loss: -1219.8336 - acc:0.0000e+00
19990/20000 [=====>..] - ETA: 0s - loss: -1219.8532 - acc:0.0000e+00
20000/20000 [=====] - 1s - loss: -1219.8594 - acc: 0.0000e+00
18880/20000 [=====>..] - ETA: 0sacc: 0.00%
```

I do not understand why. Can you please help me

Anon December 26, 2016 at 6:44 am #

What version of Python are you running?

karishma sharma December 22, 2016 at 10:03 am #

Hi Jason,

Since the epoch is set to 150 and batch size is 10, does the training algorithm pick 10 training examples 768 total in X. Or does it sample randomly after it has finished covering all.

Thanks



Jason Brownlee December 23, 2016 at 5:27 am #

Good question,

REPLY ↩

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

It iterates over the dataset 150 times and within one epoch it works through 10 rows at a time before doing an update to the weights. The patterns are shuffled before each epoch.

I hope that helps.

Kaustuv January 9, 2017 at 4:57 am #

REPLY ↩

Hi Jason

Thanks a lot for this blog. It really helps me to start learning deep learning which was in a planning state for last few months. Your simple enrich blogs are awesome. No questions from my side before completing all tutorials.

One question regarding availability of your book. How can I buy those books from India ?



Jason Brownlee January 9, 2017 at 7:53 am #

All my books and training are digital, you can purchase them from here:

<http://machinelearningmastery.com/products>

Stephen Wilson January 15, 2017 at 4:00 pm #

Hi Jason, firstly your work here is a fantastic resource and I am very thankful for the effort you put in. I am a slightly-better-than-beginner at python and an absolute novice at ML, I wonder if you could help me from.

My data is thus:

Column Names: 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, Result

Values: 4, 4, 6, 6, 3, 2, 5, 5, 0, 0, 0, 0, 0, 0, 0, 4

I want to find the percentage chance of each Column Names category being the Result based off the configuration of all the values present from 1-15. Then if need be compare the configuration of Values with another row of values to find the same, Resulting in the total needed calculation as:

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Column Names: 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, Result

Values: 4, 4, 6, 6, 3, 2, 5, 5, 0, 0, 0, 0, 0, 0, 0, 4

Values2: 7, 3, 5, 1, 4, 8, 6, 2, 9, 9, 9, 9, 9, 9, 9

I apologize if my explanation is not clear, and appreciate any help you can give me thank you.



Jason Brownlee January 16, 2017 at 10:39 am #

REPLY ↩

Hi Stephen,

This process might help you work through your problem:

<http://machinelearningmastery.com/start-here/#process>

Specifically the first step in defining your problem.

Let me know how you go.

Rohit January 16, 2017 at 10:37 pm #

Thanks Jason for such a nice and concise example.

Just wanted to ask if it is possible to save this model in a file and port it to may be an Android or iOS dev same?

Thanks

Rohit

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee January 17, 2017 at 7:38 am #

REPLY ↩

Thanks Rohit,

Get Your Start in Machine Learning

Here's an example of saving a Keras model to file:

<http://machinelearningmastery.com/save-load-keras-deep-learning-models/>

I don't know about running Keras on an Android or iOS device. Let me know how you go.

zaheer khan June 16, 2017 at 7:17 pm #

REPLY ↩

Dear Jason, Thanks for sharing this article.

I am novice to the deep learning, and my apology if my question is not clear. my question is could we call all that functions and program from any .php,.aspx, or .html webpage. i mean i load the variables and other files selection from user interface and then make them input to this functions.

will be waiting for your kind reply.

thanks in advance.

zaheer



Jason Brownlee June 17, 2017 at 7:25 am #

Perhaps, this sounds like a systems design question, not really machine learning.

I would suggest you gather requirements, assess risks like any software engineering project

Hsiang January 18, 2017 at 3:35 pm #

Hi, Jason

Thank you for your blog! It is wonderful!

I used tensorflow as backend, and implemented the procedures using Jupyter.

I did "source activate tensorflow" -> "ipython notebook".

I can successfully use Keras and import tensorflow.

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

However, it seems that such environment doesn't support pandas and sklearn.

Do you have any way to incorporate pandas, sklearn and keras?

(I wish to use sklearn to revisit the classification problem and compare the accuracy with the deep learning method. But I also wish to put the works together in the same interface.)

Thanks!



Jason Brownlee January 19, 2017 at 7:24 am #

REPLY ↩

Sorry, I do not use notebooks myself. I cannot offer you good advice.

Hsiang January 19, 2017 at 12:53 pm #

Thanks, Jason!

Actually the problem is not on notebooks. Even I used the terminal mode, i.e. doing "source activate" that mean tensorflow library is not compatible with sklearn? Thanks again!



Jason Brownlee January 20, 2017 at 10:17 am #

Sorry Hsiang, I don't have experience using sklearn and tensorflow with virtual env

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Hsiang January 21, 2017 at 12:46 am #

Thank you!

Get Your Start in Machine Learning



Jason Brownlee January 21, 2017 at 10:34 am #

You're welcome Hsiang.

keshav bansal January 24, 2017 at 12:45 am #

REPLY ↩

hello sir,

A very informative post indeed . I know my question is a very trivial one but can you please show me how to predict on a explicitly mentioned data tuple say $v=[6,148,72,35,0,33.6,0.627,50]$
thanks for the tutorial anyway



Jason Brownlee January 24, 2017 at 11:04 am #

Hi keshav,

You can make predictions by calling `model.predict()`

CATRINA WEBB January 25, 2017 at 9:06 am #

When I rerun the file (without predictions) does it reset the model and weights?

Ericson January 30, 2017 at 8:04 pm #

REPLY ↩

excuse me sir, i wanna ask you a question about this paragraph "dataset = numpy.loadtxt("pima-indians-diabetes.csv",delimiter=',')", i used the mac and downloaded the dataset,then i exchanged the text into csv file. Running the program

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
,hen i got:{Python 2.7.13 (v2.7.13:a06454b1afa1, Dec 17 2016, 12:39:47)
```

```
[GCC 4.2.1 (Apple Inc. build 5666) (dot 3)] on darwin
```

```
Type "copyright", "credits" or "license()" for more information.
```

```
>>>
```

```
===== RESTART: /Users/luowenbin/Documents/database_test.py =====
```

```
Using TensorFlow backend.
```

```
Traceback (most recent call last):
```

```
File "/Users/luowenbin/Documents/database_test.py", line 9, in
```

```
dataset = numpy.loadtxt("pima-indians-diabetes.csv",delimiter=',')
```

```
File "/Library/Frameworks/Python.framework/Versions/2.7/lib/python2.7/site-packages/numpy/lib/npio.py", line 985, in loadtxt
```

```
items = [conv(val) for (conv, val) in zip(converters, vals)]
```

```
File "/Library/Frameworks/Python.framework/Versions/2.7/lib/python2.7/site-packages/numpy/lib/npio.py", line 687, in floatconv
```

```
return float(x)
```

```
ValueError: could not convert string to float: book
```

```
>>> }
```

```
How can i solve this problem? give me a hand thank you!
```



Jason Brownlee February 1, 2017 at 10:22 am #

Hi Ericson,

Confirm that the contents of "pima-indians-diabetes.csv" meet your expectation of a list of CSV lines

Sukhpal February 7, 2017 at 9:00 pm #

REPLY ↩

excuse me sir,when i run this code for my data set ,I encounter this problem...please help me finding solution to this problem

```
runfile('C:/Users/sukhpal/.spyder/temp.py', wdir='C:/Users/sukhpal/.spyder')
```

```
Using TensorFlow backend.
```

```
Traceback (most recent call last):
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
File "", line 1, in
runfile('C:/Users/sukhpal/.spyder/temp.py', wdir='C:/Users/sukhpal/.spyder')

File "C:\Users\sukhpal\Anaconda2\lib\site-packages\spyder\utils\site\sitecustomize.py", line 866, in runfile
execfile(filename, namespace)

File "C:\Users\sukhpal\Anaconda2\lib\site-packages\spyder\utils\site\sitecustomize.py", line 87, in execfile
exec(compile(scripttext, filename, 'exec'), glob, loc)

File "C:/Users/sukhpal/.spyder/temp.py", line 1, in
from keras.models import Sequential

File "C:\Users\sukhpal\Anaconda2\lib\site-packages\keras\__init__.py", line 2, in
from . import backend

File "C:\Users\sukhpal\Anaconda2\lib\site-packages\keras\backend\__init__.py", line 67, in
from .tensorflow_backend import *

File "C:\Users\sukhpal\Anaconda2\lib\site-packages\keras\backend\tensorflow_backend.py", line 1, in
import tensorflow as tf

ImportError: No module named tensorflow
```



Jason Brownlee February 8, 2017 at 9:34 am #

This is a change with the most recent version of tensorflow, I will investigate and change the

For now, consider installing and using an older version of tensorflow.

Will February 14, 2017 at 5:33 am #

REPLY ↩

Great tutorial! Amazing amount of work you've put in and great marketing skills (I also have an email list, ebooks and sequence, etc). I ran this in Jupyter notebook... I noticed the 144th epoch (acc .7982) had more accuracy than at 150. Why is that?

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

P.S. i did this for the print: `print(numpy.round(predictions))`

It seems to avoid a list of arrays which when printing includes the dtype (messy)



Jason Brownlee February 14, 2017 at 10:07 am #

REPLY ↩

Thanks Will.

The model will fluctuate in performance while learning. You can configure triggered check points to save the model if/when conditions like a decrease in train/validation performance is detected. Here's an example:

<http://machinelearningmastery.com/check-point-deep-learning-models-keras/>

Sukhpal February 14, 2017 at 3:50 pm #

Please help me to find out this error

```
runfile('C:/Users/sukhpal/.spyder/temp.py', wdir='C:/Users/sukhpal/.spyder')ERROR: execution aborted
```



Jason Brownlee February 15, 2017 at 11:32 am #

I'm not sure Sukhpal.

Consider getting code working from the command line, I don't use IDEs myself.

Kamal February 14, 2017 at 5:15 pm #

REPLY ↩

please help me to find this error find this error

Epoch 194/195

195/195 [=====] - 0s - loss: 0.2692 - acc: 0.8667

Epoch 195/195

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

195/195 [=====] – 0s – loss: 0.2586 – acc: 0.8667

195/195 [=====] – 0s

Traceback (most recent call last):



Jason Brownlee February 15, 2017 at 11:32 am #

REPLY ↩

What was the error exactly Kamal?

Kamal February 15, 2017 at 3:24 pm #

sir when i run the code on my data set
then it doesnot show overall accuracy although it shows the accuracy and loss for the whole iterations



Jason Brownlee February 16, 2017 at 11:06 am #

I'm not sure I understand your question Kamal, please you could restate it?

Val February 15, 2017 at 9:00 pm #

Hi Jason, im just starting deep learning in python using keras and theano. I have followed the installation instructions without a hitch. Tested some examples but when i run this one line by line i get a lot of exceptions and errors once i run the “model.fit(X,Y, nb_epochs=150, batch_size=10”



Jason Brownlee February 16, 2017 at 11:06 am #

REPLY ↩

What errors are you getting?

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

CrisH February 17, 2017 at 8:12 pm #

REPLY ↩

Hi, how do I know what number to use for `random.seed()` ? I mean you use 7, is there any reason for that? Also is it enough to use it only once, in the beginning of the code?



Jason Brownlee February 18, 2017 at 8:38 am #

REPLY ↩

You can use any number CrisH. The fixed random seed makes the example reproducible.

You can learn more about randomness and random seeds in this post:
<http://machinelearningmastery.com/randomness-in-machine-learning/>

kk February 18, 2017 at 1:53 am #

am new to deep learning and found this great tutorial. keep it up and look forward!!



Jason Brownlee February 18, 2017 at 8:41 am #

Thanks!

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Iqra Ameer February 21, 2017 at 5:20 am #

REPLY ↩

Hi, I have a problem in execution the above example as it. It seems that it's not running properly and stops at Using TensorFlow backend.

Epoch 147/150

768/768 [=====] – 0s – loss: 0.4709 – acc: 0.7878

Get Your Start in Machine Learning


```
Epoch 148/150
768/768 [=====] - 0s - loss: 0.4690 - acc: 0.7812
Epoch 149/150
768/768 [=====] - 0s - loss: 0.4711 - acc: 0.7721
Epoch 150/150
768/768 [=====] - 0s - loss: 0.4731 - acc: 0.7747
32/768 [>.....] - ETA: 0sacc: 76.43%
```

I am new in this field, could you please guide me about this error.
I also executed on another data set, it stops with the same behavior.



Jason Brownlee February 21, 2017 at 9:39 am #

What is the error exactly? The example hangs?

Maybe try the Theano backend and see if that makes a difference. Also make sure all of your libraries

Iqra Ameer February 22, 2017 at 5:47 am #

Dear Jason,

Thank you so much for your valuable suggestions. I tried Theano backend and also updated all my libraries

```
768/768 [=====] - 0s - loss: 0.4656 - acc: 0.7799
Epoch 149/150
768/768 [=====] - 0s - loss: 0.4589 - acc: 0.7826
Epoch 150/150
768/768 [=====] - 0s - loss: 0.4611 - acc: 0.7773
32/768 [>.....] - ETA: 0sacc: 78.91%
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Jason Brownlee February 22, 2017 at 10:05 am #

Get Your Start in Machine Learning



I'm sorry to hear that, I have not seen this issue before.

Perhaps a RAM issue or a CPU overheating issue? Are you able to try different hardware?

frd March 8, 2017 at 2:50 am #

REPLY ↩

Hi!

Were you able to find a solution for that?

I'm having exactly the same problem

(...)

Epoch 149/150

768/768 [=====] - 0s - loss: 0.4593 - acc: 0.7773

Epoch 150/150

768/768 [=====] - 0s - loss: 0.4586 - acc: 0.7891

32/768 [>.....] - ETA: 0s acc: 76.69%

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Bhanu February 23, 2017 at 1:51 pm #

Hello sir,

i want to ask wether we can convert this code to deep learning wid increasing number of layers..

REPLY ↩



Jason Brownlee February 24, 2017 at 10:12 am #

Sure you can increase the number of layers, try it and see.

Ananya Mohapatra February 28, 2017 at 6:40 pm #

Get Your Start in Machine Learning

hello sir,

could you please tell me how do i determine the no.of neurons in each layer, because i am using a different dataset and am unable to know the no.of neurons in each layer



Jason Brownlee March 1, 2017 at 8:33 am #

REPLY ↩

Hi Ananya, great question.

Sorry, there is no good theory on how to configure a neural net.

You can configure the number of neurons in a layer by trial and error. Also consider tuning the number of epochs and batch size at the same time.

Ananya Mohapatra March 1, 2017 at 4:42 pm #

thank you so much sir. It worked ! 😊



Jason Brownlee March 2, 2017 at 8:11 am #

Glad to hear it Ananya.

Jayant Sahewal February 28, 2017 at 8:11 pm #

REPLY ↩

Hi Jason,

really helpful blog. I have a question about how much time does it take to converge?

I have a dataset with around 4000 records, 3 input columns and 1 output column. I came up with the following model

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
def create_model(dropout_rate=0.0, weight_constraint=0, learning_rate=0.001, activation='linear'):
    # create model
    model = Sequential()
    model.add(Dense(6, input_dim=3, init='uniform', activation=activation, W_constraint=maxnorm(weight_constraint)))
    model.add(Dropout(dropout_rate))
    model.add(Dense(1, init='uniform', activation='sigmoid'))
    # Optimizer
    optimizer = Adam(lr=learning_rate)
    # Compile model
    model.compile(loss='binary_crossentropy', optimizer=optimizer, metrics=['accuracy'])
    return model

# create model
model = KerasRegressor(build_fn=create_model, verbose=0)
# define the grid search parameters
batch_size = [10]
epochs = [100]
weight_constraint = [3]
dropout_rate = [0.9]
learning_rate = [0.01]
activation = ['linear']
param_grid = dict(batch_size=batch_size, nb_epoch=epochs, dropout_rate=dropout_rate, \
weight_constraint=weight_constraint, learning_rate=learning_rate, activation=activation)
grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1, cv=5)
grid_result = grid.fit(X_train, Y_train)
```

I have a 32 core machine with 64 GB RAM and it does not converge even in more than an hour. I can see the accuracy, but it's not increasing. However, if I change the input neurons to 3 then it converges in around 2 minutes.

Keras version: 1.1.1

Tensorflow version: 0.10.0rc0

theano version: 0.8.2.dev-90127534cbfe3fbb290ce85d1abf8bb9a5b203

It's using Tensorflow backend. Can you help me understand what is going on or point me in the right direction? Do you think switching to theano will help?

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Best,
Jayant



Jason Brownlee March 1, 2017 at 8:36 am #

REPLY ↩

This post might help you tune your deep learning model:

<http://machinelearningmastery.com/improve-deep-learning-performance/>

I hope that helps as a start.

Animesh Mohanty March 1, 2017 at 9:21 pm #

hello sir,

could you please tell me how can i plot the results of the code on a graph . I made a few adjustments to



Jason Brownlee March 2, 2017 at 8:16 am #

What do you want to plot exactly Animesh?

Animesh Mohanty March 2, 2017 at 4:56 pm #

Accuracy vs no.of neurons in the input layer and the no.of neurons in the hidden layer

param March 2, 2017 at 12:15 am #

REPLY ↩

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

REPLY ↩

Get Your Start in Machine Learning

sir can u plz explain

the different attributes used in this statement

```
print("%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))
```

param March 2, 2017 at 12:16 am #

REPLY ↩

precisely, what is model.metrics_names



Jason Brownlee March 2, 2017 at 8:22 am #

model.metrics_names is a list of names of the metrics collected during training.

More details here:

<https://keras.io/models/sequential/>

REPLY ↩

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee March 2, 2017 at 8:20 am #

Hi param,

It is using string formatting. %s formats a string, %.2f formats a floating point value with 2 decimal places.

You can learn more about the print function here:

<https://docs.python.org/3/library/functions.html#print>

More info on string formatting here:

<https://pyformat.info/>

REPLY ↩

Vijin K P March 2, 2017 at 4:01 am #

Get Your Start in Machine Learning

Hi Jason,

It was an awesome post. Could you please tell me how to we decide the following in a DNN

1. number of neurons in the hidden layers
2. number of hidden layers

Thanks.

Vijin



Jason Brownlee March 2, 2017 at 8:22 am #

REPLY ↩

Great question Vijin.

Generally, trial and error. There are no good theories on how to configure a neural network.

Vijin K P March 3, 2017 at 5:23 am #

We do cross validation, grid search etc to find the hyper parameters in machine algorithm. How to choose the above parameters??



Jason Brownlee March 3, 2017 at 7:46 am #

Yes, we can use grid search and tuning for neural nets.

The stochastic nature of neural nets means that each experiment (set of configs) will have to be run many times (30? 100?) so that you can take the mean performance.

More general info on tuning neural nets here:

<http://machinelearningmastery.com/improve-deep-learning-performance/>

More on randomness and stochastic algorithms here:

<http://machinelearningmastery.com/randomness-in-machine-learning/>

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Bogdan March 2, 2017 at 11:48 pm #

REPLY ↩

Jason, Please tell me about these lines in your code:

```
seed = 7
numpy.random.seed(seed)
```

What do they do? And why do they do it?

One more question is why do you call the last section Bonus:Make a prediction?

I thought this what ANN was created for. What the point if your network's output is just what you have already know?



Jason Brownlee March 3, 2017 at 7:44 am #

They seed the random number generator so that it produces the same sequence of random numbers. This way you get the same result as me.

I'm not convinced it works with Keras though.

More on randomness in machine learning here:

<http://machinelearningmastery.com/randomness-in-machine-learning/>

I was showing how to build and evaluate the model in this tutorial. The part about standalone prediction

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Sounak saho March 3, 2017 at 7:39 pm #

REPLY ↩

what exactly is the work of "seed" in the neural network code? what does it do?



Jason Brownlee March 6, 2017 at 10:44 am #

REPLY ↩

Seed refers to seeding the random number generator so that the same sequence of random numbers is generated each time the example is run.

The aim is to make the examples 100% reproducible, but this is hard with symbolic math libs like Theano and TensorFlow backends.

For more on randomness in machine learning, see this post:

<http://machinelearningmastery.com/randomness-in-machine-learning/>

Priya Sundari March 3, 2017 at 10:19 pm #

REPLY ↩

hello sir

could you plz tell me what is the role of optimizer and binary_crossentropy exactly? it is written that optimizer is used to search through the weights of the network which weights are we talking about exactly?



Jason Brownlee March 6, 2017 at 10:48 am #

Hi Priya,

You can learn more about the fundamentals of neural nets here:

<http://machinelearningmastery.com/neural-networks-crash-course/>

Bogdan March 3, 2017 at 10:23 pm #

If I am not mistaken, those lines I commented about used when we write

init = 'uniform'

?

Bogdan March 3, 2017 at 10:44 pm #

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Could you explain in more details what is the batch size?



Jason Brownlee March 6, 2017 at 10:50 am #

REPLY ↩

Hi Bogdan,

Batch size is how many patterns to show to the network before the weights are updated with the accumulated errors. The smaller the batch, the faster the learning, but also the more noisy the learning (higher variance).

Try exploring different batch sizes and see the effect on the train and test performance over each epoch.

Mohammad March 7, 2017 at 6:50 am #

Dear Jason

Firstly, thanks for your great tutorials.

I am trying to classify computer networks packets using first 500 bytes of every packet to identify its protocol. I just want to do binary classification and then tackle multilabel classification for 10 protocols. Here is my code. How can I improve the performance? should I Use RNNs?

#####

```
model=Sequential()
model.add(Convolution1D(64,10,border_mode='valid',
activation='relu',subsample_length=1, input_shape=(500, 1)))
#model.add(Convolution2D(32,5,5,border_mode='valid',input_shape=(1,28,28),))
model.add(MaxPooling1D(2))
model.add(Flatten())
model.add(Dense(200,activation='relu'))
model.add(Dense(1,activation='sigmoid'))
model.compile(loss='binary_crossentropy',
optimizer='adam',metrics=['accuracy'])
model.fit(train_set, y_train,
batch_size=250,
```

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
nb_epoch=30,
show_accuracy=True)
#x2= get_activations(model, 0,xprim )
#score = model.evaluate(t, y_test, show_accuracy = True, verbose = 0)
#print(score[0])
```



Jason Brownlee March 7, 2017 at 9:37 am #

REPLY ↩

This post lists some ideas to try an lift performance:

<http://machinelearningmastery.com/improve-deep-learning-performance/>

Damiano March 7, 2017 at 10:13 pm #

Hi Jason, thank you so much for this awesome tutorial. I have just started with python and mac
I am joking with the code doing few changes, for example i have changed..

this:

```
# create model
model = Sequential()
model.add(Dense(250, input_dim=8, init='uniform', activation='relu'))
model.add(Dense(200, init='uniform', activation='relu'))
model.add(Dense(200, init='uniform', activation='relu'))
model.add(Dense(1, init='uniform', activation='sigmoid'))
```

and this:

```
model.fit(X, Y, nb_epoch=250, batch_size=10)
```

then i would like to pass some arrays for prediction so...

```
new_input = numpy.array([[3,88,58,11,54,24.8,267,22],[6,92,92,0,0,19.9,188,28], [10,101,76,48,180,32.9,171,63], [2,122,70,27,0,36.8,0.34,27],
[5,121,72,23,112,26.2,245,30]])
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**
Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
predictions = model.predict(new_input)
print predictions # [1.0, 1.0, 1.0, 0.0, 1.0]
```

is this correct? In this example i used the same series of training (that have 0 class), but i am getting wrong results. Only one array is correctly predicted.

Thank you so much!



Jason Brownlee March 8, 2017 at 9:41 am #

REPLY ↩

Looks good. Perhaps you could try changing the configuration of your model to make it more skillful?

See this post:

<http://machinelearningmastery.com/improve-deep-learning-performance/>

ANJI March 13, 2017 at 8:48 pm #

hello sir,

could you please tell me to rectify my error below it is raised while model is training:

```
str(array.shape))
```

ValueError: Error when checking model input: expected convolution2d_input_1 to have 4 dimensions, but got array of shape (1, 1, 1, 1)

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee March 14, 2017 at 8:17 am #

It looks like you are working with CNN, not related to this tutorial.

Consider trying this tutorial to get familiar with CNNs:

<http://machinelearningmastery.com/handwritten-digit-recognition-using-convolutional-neural-networks-python-keras/>

Get Your Start in Machine Learning

Rimjhim March 14, 2017 at 8:21 pm #

REPLY ↩

I want a neural that can predict sin values. Further from a given data set i need to determine the function(for example if the data is of tan or cos, then how to determine that data is of tan only or cos only)

Thanks in advance

Sudarshan March 15, 2017 at 11:19 pm #

REPLY ↩

Keras just updated to Keras 2.0. I have an updated version of this code here: https://github.com/sudarshan85/keras-projects/tree/master/mlm/pima_indians



Jason Brownlee March 16, 2017 at 7:59 am #

Nice work.

subhasish March 16, 2017 at 5:09 pm #

hello sir,
can we use PSO (particle swarm optimisation) in this? if so can you tell how?



Jason Brownlee March 17, 2017 at 8:25 am #

Sorry, I don't have an example of PSO for fitting neural network weights.

REPLY ↩

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Ananya Mohapatra March 16, 2017 at 10:03 pm #

REPLY ↩

hello sir,

what type of neural network is used in this code? as there are 3 types of Neural network that are... feedforward, radial basis function and recurrent neurak network.



Jason Brownlee March 17, 2017 at 8:28 am #

REPLY ↩

A multilayer perceptron (MLP) neural network. A classic type from the 1980s.

Diego March 17, 2017 at 3:58 am #

got this error while compiling..

sigmoid_cross_entropy_with_logits() got an unexpected keyword argument 'labels'



Jason Brownlee March 17, 2017 at 8:30 am #

Perhaps confirm that your libraries are all up to date (Keras, Theano or TensorFlow)?

Rohan March 20, 2017 at 5:20 am #

REPLY ↩

Hi Jason!

I am trying to use two odd frames of a video to predict the even one. Thus I need to give two images as input to the network and get one image as output. Can you help me with the syntax for the first model.add()? I have X_train of dimension (190, 2, 240, 320, 3) where 190 are the number of odd nairs 2 are the two odd images, and (240,320,3) are the (height, width, depth) of each image.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

REPLY

Warnings like these:

UserWarning: Update your Dense call to the Keras 2 API: Dense(12, activation="relu", kernel_initializer="uniform", input_dim=8)

UserWarning: Update your Dense call to the Keras 2 API: Dense(8, activation="relu", kernel_initializer="

UserWarning: Update your Dense call to the Keras 2 API: Dense(1, activation="sigmoid", kernel_initializer="glorot_uniform")

UserWarning: The nb_epoch argument in fit has been renamed epochs`.

I think these are due to some package update..

But, the output of predictions was an array of zeros...

[illegible]

I am running in a Linux Machine, Fedora 24,

Python 2.7.13 (default, Jan 12 2017, 17:59:37)

[GCC 6.3.1 20161221 (Red Hat 6.3.1-1)] on linux2

START MY EMAIL COURSE

Get Your Start in Machine Learning

Why?

Thank you!



Jason Brownlee March 21, 2017 at 8:45 am #

REPLY ↩

These look like warnings related to the recent Keras 2.0 release.

They look like just warning and that you can still run the example.

I do not know why you are getting all zeros. I will investigate.

Ananya Mohapatra March 21, 2017 at 6:21 pm #

hello sir,

can you please help me build a recurrent neural network with the above given dataset. i am having a bit



Jason Brownlee March 22, 2017 at 7:56 am #

Hi Ananya ,

The Pima Indian diabetes dataset is a binary classification problem. It is not appropriate for a Recurrent Neural Network. You need more information to learn.

Ananya Mohapatra March 22, 2017 at 8:04 pm #

REPLY ↩

sir so could you tell on which type of dataset would the recurrent neural network accurately work? i have the dataset of EEG signals of epileptic patients...will recurrent network work on this?

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning



Jason Brownlee March 23, 2017 at 8:49 am #

REPLY ↩

It may if it is regular enough.

LSTMs are excellent at sequence problems that have regularity or clear signals to detect.

Shane March 22, 2017 at 5:18 am #

REPLY ↩

Hi Jason, I have a quick question related to an error I am receiving when running the code in the tutorial...

When I run

```
# Compile model
```

```
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

Python returns the following error:

```
sigmoid_cross_entropy_with_logits() got an unexpected keyword argument 'labels'
```



Jason Brownlee March 22, 2017 at 8:09 am #

Sorry, I have not seen this error Shane.

Perhaps check that your environment is up to date with the latest versions of the deep learning library.

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Tejes March 24, 2017 at 1:04 am #

REPLY ↩

Hi Jason,

Thanks for this awesome post.

Get Your Start in Machine Learning

I ran your code with tensorflow back end, just out of curiosity. The accuracy returned was different every time I ran the code. That didn't happen with theano. Can you tell me why?

Thanks in advance!



Jason Brownlee March 24, 2017 at 7:56 am #

REPLY ↩

You will get different accuracy each time you run the code because neural networks are stochastic.

This is not related to the backend (I expect).

More on randomness in machine learning here:

<http://machinelearningmastery.com/randomness-in-machine-learning/>

Saurabh Bhagvatula March 27, 2017 at 9:49 pm #

Hi Jason,

I'm new to deep learning and learning it from your tutorials, which previously helped me understand Machine Learning. In the following code, I want to know why the number of neurons differ from input_dim in first layer of Neural Network.

```
# create model
model = Sequential()
model.add(Dense(12, input_dim=8, init='uniform', activation='relu'))
model.add(Dense(8, init='uniform', activation='relu'))
model.add(Dense(1, init='uniform', activation='sigmoid'))
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee March 28, 2017 at 8:22 am #

REPLY ↩

You can specify the number of inputs via "input_dim", you can specify the number of neurons in the first hidden layer as the first parameter to Dense().

Get Your Start in Machine Learning

Saurabh Bhagvatula March 28, 2017 at 4:15 pm #

REPLY ↩

Thanx a lot.



Jason Brownlee March 29, 2017 at 9:05 am #

REPLY ↩

You're welcome.

Nalini March 29, 2017 at 2:52 am #

Hi Jason

while running this code for k fold cross validation it is not working.please give the code for k fold cross va



Jason Brownlee March 29, 2017 at 9:10 am #

Generally neural nets are too slow/large for k-fold cross validation.

Nevertheless, you can use a sklearn wrapper for a keras model and use it with any sklearn resampli

<http://machinelearningmastery.com/evaluate-performance-machine-learning-algorithms-python-using-sklearn>

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

trangtruong March 29, 2017 at 7:04 pm #

REPLY ↩

Hi Jason, why i use function evaluate to get accuracy score my model with test dataset, it return result >1, i can't understand.

Get Your Start in Machine Learning

enixon April 3, 2017 at 3:08 am #

REPLY ↩

Hey Jason, thanks for this great article! I get the following error when running the code above:

TypeError: Received unknown keyword arguments: {'epochs': 150}

Any ideas on why that might be? I can't get 'epochs', nb_epochs, etc to work...



Jason Brownlee April 4, 2017 at 9:07 am #

REPLY ↩

You need to update to Keras version 2.0 or higher.

Ananya Mohapatra April 5, 2017 at 9:30 pm #

```
def baseline_model():  
# create model  
model = Sequential()  
model.add(Dense(10, input_dim=25, init='normal', activation='softplus'))  
model.add(Dense(3, init='normal', activation='softmax'))  
# Compile model  
model.compile(loss='mean_squared_error', optimizer='adam', metrics=['accuracy'])  
return model
```

sir here mean_square_error has been used for loss calculation. Is it the same as LMS algorithm. If not, Can we use LMS, NLMS or RES to calculate the loss?

Ahmad Hijazi April 5, 2017 at 10:19 pm #

REPLY ↩

Hello Jason, thank you a lot for this example.

My question is, after I trained the model and an accuracy of 79.2% for example is obtained successfully,

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

for example if a new patient with new records appear, I want to guess the result (0 or 1) for him, how can I do that in the code?



Jason Brownlee April 9, 2017 at 2:36 pm #

REPLY ↩

You can fit your model on all available training data then make predictions on new data as follows:

```
1 yhat = model.predict(X)
```

Perick Flaus April 6, 2017 at 12:16 am #

REPLY ↩

Thanks Jason, how can we test if new patient will be diabetic or no (0 or 1) ?



Jason Brownlee April 9, 2017 at 2:36 pm #

Fit the model on all training data and call:

```
1 yhat = model.predict(X)
```

Gangadhar April 12, 2017 at 1:28 am #

Dr Jason,

In compiling the model i got below error

TypeError: compile() got an unexpected keyword argument 'metrics'

unable to resolve the below error

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning



Jason Brownlee April 12, 2017 at 7:53 am #

REPLY ↩

Ensure you have the latest version of Keras, v2.0 or higher.

Omogbehin Azeez April 13, 2017 at 1:48 am #

REPLY ↩

Hello sir,

Thank you for the post. A quick question, my dataset has 24 input and 1 binary output(170 instances, 100 epoch , hidden layer=6 and 10 batch, kernel_initializer='normal') . I adapted your code using Tensor flow and keras. I am having an accuracy of 98 to 100 percent. I am scared of over-fitting in my model. I need your candid advice. Kind regards sir



Jason Brownlee April 13, 2017 at 10:07 am #

Yes, evaluate your model using k-fold cross-validation to ensure you are not tricking yourself

Omogbehin Azeez April 14, 2017 at 1:08 am #

Thank you sir

Sethu Baktha April 13, 2017 at 5:19 am #

Hi Jason,

If I want to use the diabetes dataset (NOT Pima) <https://archive.ics.uci.edu/ml/datasets/Diabetes> to predict Blood Glucose which tutorials and e-books of yours would I need to start with.... Also, the data in its current format with time, code and value is it usable as is or do I need to convert the data in another format to be able to use it.

Thanks for your help

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

REPLY ↩

Get Your Start in Machine Learning



Jason Brownlee April 13, 2017 at 10:13 am #

REPLY ↩

This process will help you frame and work through your dataset:

<http://machinelearningmastery.com/start-here/#process>

I hope that helps as a start.

Sethu Baktha April 13, 2017 at 10:25 am #

REPLY ↩

Dr. Jason,

The data is time series(time based data) with categorical(20) with two numbers one for insulin le
series data does not have every categorical data... For example one category is blood sugar be
breakfast, before lunch and after lunch... Some times some of these category data is missing...
time series, categorical data with some category of data missing what to do in those cases.... Pl
these points?



Jason Brownlee April 14, 2017 at 8:43 am #

Hi Sethu,

I have many posts on time series that will help. Get started here:

<http://machinelearningmastery.com/start-here/#timeseries>

With categorical data, I would recommend an integer encoding perhaps followed by a one-hot encoding. You can learn more about these encodings here:

<http://machinelearningmastery.com/data-preparation-gradient-boosting-xgboost-python/>

I hope that helps.

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Omogbehin Azeez April 14, 2017 at 9:49 am #

REPLY ↩

Hello sir,

Is it compulsory to normalize the data before using ANN model. I read it somewhere I which the author insisted that each attribute be comparable on the scale of [0,1] for a meaningful model. What is your take on that sir. Kind regards.



Jason Brownlee April 15, 2017 at 9:29 am #

REPLY ↩

Yes. You must scale your data to the bounds of the activation used.

shiva April 14, 2017 at 10:38 am #

Hi Jason, You are simply awesome. I'm one of the many who got benefited from your book "mastering machine learning with keras". I have a medical image classification problem. I have two classes of medical images (each class having 1000 images). I have used convolutional neural networks. Could you guide me how to load this data to the keras dataset? Or how to preprocess it? I would be very kindly help.



Jason Brownlee April 15, 2017 at 9:30 am #

Load the data as numpy arrays and then you can use it with Keras.

Omogbehin Azeez April 18, 2017 at 12:09 am #

REPLY ↩

Hello sir,

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

I adapted your code with the cross validation pipelined with ANN (Keras) for my model. It gave me 100% sm. I got the data from UCI (Chronic Kidney Disease). It was 400 instances, 24 input attributes and 1 binary attribute. When I removed the rows with missing data I was left with 170 instances. Is my dataset too small for (24 input layer, 24 hidden layer and 1 output layer ANN, using adam and kernel initializer as uniform)?



Jason Brownlee April 18, 2017 at 8:32 am #

REPLY ↩

It is not too small.

Generally, the size of the training dataset really depends on how you intend to use the model.

Omogbehin Azeez April 18, 2017 at 11:10 pm #

Thank you sir for the response, I guess I have to contend with the over-fitting of my mo

Padmanabhan Krishnamurthy April 19, 2017 at 6:26 pm #

Hi Jason,

Great tutorial. Love the site 😊

Just a quick query : why have you used adam as an optimizer over sgd? Moreover, when do we use sgd

Thanks

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee April 20, 2017 at 9:23 am #

REPLY ↩

ADAM seems to consistently work well with little or no customization.

SGD requires configuration of at least the learning rate and momentum.

Get Your Start in Machine Learning

Try a few methods and use the one that works best for your problem.

Padmanabhan Krishnamurthy April 20, 2017 at 4:32 pm #

REPLY ↩

Thanks 😊

Omogbehin Azeez April 25, 2017 at 8:13 am #

REPLY ↩

Hello sir,

Good day sir, how can I get all the weights and biases of the keras ANN. Kind regards.



Jason Brownlee April 26, 2017 at 6:19 am #

You can save the network weights, see this post:

<http://machinelearningmastery.com/save-load-keras-deep-learning-models/>

You can also use the API to access the weights directly.

Shiva April 27, 2017 at 5:43 am #

Hi Jason,

I am currently working with the IMDB sentiment analysis problem as mentioned in your book. Am using Anaconda 3 with Python 3.5.2. In an attempt to summarize the review length as you have mentioned in your book, When i try to execute the command:

```
result = map(len, X)
print("Mean %.2f words (%f)" % (numpy.mean(result), numpy.std(result)))
```

it returns the error: unsupported operand type(s) for /: 'map' and 'int'

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

kindly help with the modified syntax. looking forward...



Jason Brownlee April 27, 2017 at 8:47 am #

REPLY ↩

I'm sorry to hear that. Perhaps comment out that line?
Or change it to remove the formatting and just print the raw mean and stdev values for you to review?

Elikplim May 1, 2017 at 1:58 am #

REPLY ↩

Hello, quite new to Python, Numpy and Keras(background in PHP, MYSQL etc). If there are 8 in
Array indexing starts from zero(from what I've gathered it's a Numpy Array, which is built on Python lists)
input variable(X) be X = dataset[:,0:7] (where we select from the 1st to 8th columns, ie. 0th to 7th indices
we the 9th column, ie. 8th index)?



Jason Brownlee May 1, 2017 at 5:59 am #

You can learn more about array indexing in numpy here:
<https://docs.scipy.org/doc/numpy/reference/arrays.indexing.html>

Jackie Lee May 1, 2017 at 12:47 pm #

REPLY ↩

I'm having troubles with the predictions part. It saves ValueError: Error when checking model input: expected dense_1_input to have shape (None, 502) but got array with shape (170464, 502)

MAKE PREDICTIONS

```
testset = numpy.loadtxt("right_stim_FD1.csv", delimiter=",")  
A = testset[:,0:502]
```

Get Your Start in Machine Learning

You can master applied Machine Learning
without the math or fancy degree.
Find out how in this *free* and *practical* email
course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
B = testset[:,502]
probabilities = model.predict(A, batch_size=10, verbose=1)
predictions = float(round(a) for a in probabilities)
accuracy = numpy.mean(predictions == B)
#round predictions
#rounded = [round(x[0]) for x in predictions]
print(predictions)
print("Prediction Accuracy: %.2f%%" % (accuracy*100))
```



Jason Brownlee May 2, 2017 at 5:55 am #

REPLY ↩

It looks like you might be giving the entire dataset as the output (y) rather than just the output

Anastasios Selalmazidis May 2, 2017 at 12:27 am #

Hi there,

I have a question regarding deep learning. In this tutorial we build a MLP with Keras. Is this Deep Learning



Jason Brownlee May 2, 2017 at 5:59 am #

Deep learning is MLP backprop these days:

<http://machinelearningmastery.com/what-is-deep-learning/>

Generally, deep learning refers to MLPs with lots of layers.

Eric T May 2, 2017 at 8:59 pm #

REPLY ↩

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Hi,

Would you mind if I use this code as an example of a simple network in a school project of mine?

Need to ask before using it, since I cannot find anywhere in this tutorial that you are OK with anyone using the code, and the ethics moment of my course requires me to ask (and of course give credit where credit is due).

Kind regards

Eric T



Jason Brownlee May 3, 2017 at 7:35 am #

REPLY ↩

Yes it's fine but I take no responsibility and you must credit the source.

I answer this question in my FAQ:

<http://machinelearningmastery.com/start-here/#faq>

BinhLN May 7, 2017 at 3:11 am #

Hi Jason

I have a problem

My Dataset have 500 record. But My teacher want my dataset have 100.000 record. I must have a new

Dp May 11, 2017 at 2:26 am #

Can you give a deep cnn code which includes 25 layers , in the first conv layer the filter size should be 39×39 with a total of 64 filters , in the 2nd conv layer , 21 ×21 with 32 filters , in the 3rd conv layer 11×11 with 64 filters , 4th Conv layer 7×7 with 32 layers . For an input size of image 256×256. I'm completely new in this Deep learning thing but if you can code that for me it would be a great help. Thanks

Jason Brownlee May 11, 2017 at 8:33 am #

Get Your Start in Machine Learning

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Consider using an off-the-shelf model like VGG:

<https://keras.io/applications/>

Maple May 13, 2017 at 12:58 pm #

REPLY ↩

I have to follow with the facebook metrics. But the result is very low. Help me.

I changed the input but did not improve

<http://archive.ics.uci.edu/ml/datasets/Facebook+metrics>



Jason Brownlee May 14, 2017 at 7:24 am #

I have a list of suggestions that may help as a start:

<http://machinelearningmastery.com/improve-deep-learning-performance/>

Alessandro May 14, 2017 at 1:01 am #

Hi Jason,

Great Tutorial and thanks for your effort.

I have a question, since I am beginner with keras and tensorflow.

I have installed both of them, keras and tensorflow, the latest version and I have run your example but I get always the same error.

Traceback (most recent call last):

File "CNN.py", line 18, in

model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

File "/Users/MacBookPro1/.virtualenvs/keras_tf/lib/python2.7/site-packages/keras/models.py", line 777, in compile

**kwargs)

File "/Users/MacBookPro1/.virtualenvs/keras_tf/lib/python2.7/site-packages/keras/engine/training.py", line 910, in compile

sample_weight, mask)

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

File "/Users/MacBookPro1/.virtualenvs/keras_tf/lib/python2.7/site-packages/keras/engine/training.py", line 430, in weighted

score_array = fn(y_true, y_pred)

File "/Users/MacBookPro1/.virtualenvs/keras_tf/lib/python2.7/site-packages/keras/losses.py", line 51, in binary_crossentropy

return K.mean(K.binary_crossentropy(y_pred, y_true), axis=-1)

File "/Users/MacBookPro1/.virtualenvs/keras_tf/lib/python2.7/site-packages/keras/backend/tensorflow_backend.py", line 2771, in binary_crossentropy
logits=output)

TypeError: sigmoid_cross_entropy_with_logits() got an unexpected keyword argument 'labels'

Could you help? Thanks

Alessandro



Jason Brownlee May 14, 2017 at 7:30 am #

Ouch, I have not seen this error before.

Some ideas:

- Consider trying the theano backend and see if that makes a difference.
- Try searching/posting on the keras user group and slack channel.
- Try searching/posting on stackoverflow or cross validated.

Let me know how you go.

Alessandro May 14, 2017 at 9:44 am #

Hi Jason,

I found the issue. The tensorflow installation was outdated; so I have updated it and everything is working nicely.

Good night,
Alessandro

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning



Jason Brownlee May 15, 2017 at 5:50 am #

REPLY ↩

I'm glad to hear it Alessandro.

Sheikh Rafiul Islam May 25, 2017 at 3:36 pm #

REPLY ↩

Thank you Mr. Brownlee for your wonderful easy to understand explanation



Jason Brownlee June 2, 2017 at 11:41 am #

Thnaks.

WAZED May 29, 2017 at 12:31 am #

Hi Jason,

Thank you very much for your wonderful tutorial. I have a question regarding the metrices.Is there default addition with the "Accurace".

Br

WAZED

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee June 2, 2017 at 12:15 pm #

REPLY ↩

Yes, see here:

<https://keras.io/metrics/>

Get Your Start in Machine Learning

chiranjib konwar May 29, 2017 at 4:30 am #

REPLY ↩

Hi Jason,

please send me a small note containing resources from where i can learn deep learning from scratch. thanks for the wonderful read you had prepared.

Thanks in advance

yes, my email id is chiranjib.konwar@gmail.com



Jason Brownlee June 2, 2017 at 12:16 pm #

REPLY ↩

Here:

<http://machinelearningmastery.com/start-here/#deeplearning>

Jeff June 1, 2017 at 11:48 am #

Why the NN have mistakes many times?



Jason Brownlee June 2, 2017 at 12:54 pm #

What do you mean exactly?

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

kevin June 2, 2017 at 5:53 pm #

REPLY ↩

Hi Jason,

I seem to be getting an error when applying the fit method:

ValueError: Error when checking input: expected dense_1_input to have shape (None, 12) but got array with shape (767, 8)

I looked this up and the most prominent suggestion seemed to be upgrade keras and theno, which I did, but that didn't resolve the problem.



Jason Brownlee June 3, 2017 at 7:24 am #

REPLY ↩

Ensure you have copied the code exactly from the post.

Hemanth Kumar K June 3, 2017 at 2:15 pm #

REPLY ↩

hi Jason,

I am stuck with an error

TypeError: sigmoid_cross_entropy_with_logits() got an unexpected keyword argument 'labels'

my tensor flow and keras virsions are

keras: 2.0.4

Tensorflow: 0.12

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee June 4, 2017 at 7:46 am #

I'm sorry to hear that, I have not seen that error before. Perhaps you could post a question

xena June 4, 2017 at 6:36 pm #

REPLY ↩

can anyone tell me which neural network is being used here? Is it MLP??

Get Your Start in Machine Learning



Jason Brownlee June 5, 2017 at 7:40 am #

REPLY ↩

Yes, it is a multilayer perceptron (MLP) feedforward neural network.

Nirmesh Shah June 9, 2017 at 11:00 pm #

REPLY ↩

Hi Jason,

I have run this code successfully on PC with CPU.

If I have to run the same code n another PC which contains GPU, What line should I add to make it sure that it runs on the GPU.



Jason Brownlee June 10, 2017 at 8:24 am #

The code would stay the same, your configuration of the Keras backend would change.

Please refer to TensorFlow or Theano documentation.

Prachi June 12, 2017 at 7:30 pm #

What if I want to train my neural which should detect whether the luggage is abandoned or not



Jason Brownlee June 13, 2017 at 8:18 am #

REPLY ↩

This process will help you work through your predictive modeling problem end to end:

<http://machinelearningmastery.com/start-here/#process>

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Ebtesam June 14, 2017 at 11:15 pm #

REPLY ↩

Hi

I was build neural machine translation model but the score i was get is 0 i am not sure why



Jason Brownlee June 15, 2017 at 8:45 am #

REPLY ↩

Here is a good list of things to try:

<http://machinelearningmastery.com/improve-deep-learning-performance/>

Sarvottam Patel June 20, 2017 at 7:31 pm #

HHey Jason , first of all thank you very much from the core of my heart to make me understand iteration.

File "keras_first_network.py", line 53, in
print("\n%s: %.2f" %(model.metrics_names[1]*100))
TypeError: not enough arguments for format string

Sarvottam Patel June 20, 2017 at 8:05 pm #

Sorry Sir my bad , actually I wrote it wrongly



Jason Brownlee June 21, 2017 at 8:12 am #

REPLY ↩

Confirm that you have copied the line exactly:

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

```
1 print("\n%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))
```

Joydeep June 30, 2017 at 4:15 pm #

REPLY ↩

Hi Dr Jason,

Thanks for the tutorial to get started using Keras.

I used the below snippet to directly load the dataset from the URL rather than downloading and saving as this makes the code more streamlined without having to navigate elsewhere.

```
# load pima indians dataset
datasource = numpy.DataSource().open("http://archive.ics.uci.edu/ml/machine-learning-databases/pima")
dataset = numpy.loadtxt(datasource, delimiter=",")
```



Jason Brownlee July 1, 2017 at 6:28 am #

Thanks for the tip.

Yvette July 7, 2017 at 9:01 pm #

Thanks for this helpful resource!



Jason Brownlee July 9, 2017 at 10:38 am #

I'm glad it helped.

REPLY ↩

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Andeep July 10, 2017 at 1:14 am #

REPLY ↩

Hi Dr Brownlee,

thank you very much for this great tutorial!

I would be grateful, if you could answer some questions:

1. What does the 7 in “numpy.random.seed(7)” means?

2. In my case I have 3 input neurons and 2 output neurons. Is the correct notation:

X = dataset[:,0:3]

Y = dataset[:,3:4] ?

3. The batch size means how many training data are used in one epoch, am I right?

I have thought we have to use the whole training data set for the training. In this case I would determine have achieved through experiments etc.. In your example, does the batch (sized 10) means that the con epoch or are the 10 training data randomly chosen among all training data before every epoch?

4. When evaluating the model what does the loss means (e.g. in loss: 0.5105 – acc: 0.7396)?

Is it the sum of values of the error function (e.g. mean_squared_error) of the output neurons?



Jason Brownlee July 11, 2017 at 10:19 am #

You can use any random seed you like, more here:

<http://machinelearningmastery.com/reproducible-results-neural-networks-keras/>

You are referring to the columns in your data. Your network will also need to be configured with the correct number of inputs and outputs (e.g. input and output layers).

Batch size is the number of samples in the dataset to work through before updating network weights. One epoch is comprised of one or more batches.

Loss is the term being optimized by the network. Here we use log loss:

https://en.wikipedia.org/wiki/Cross_entropy

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Andeep July 16, 2017 at 7:43 am #

REPLY ↩

Thank you for your response, Dr Brownlee !!



Jason Brownlee July 16, 2017 at 8:00 am #

REPLY ↩

I hope it helps.

Patrick Zawadzki July 11, 2017 at 5:35 am #

REPLY ↩

Is there anyway to see the relationship between these inputs? Essentially understand which input of inputs affect the output the most?

Maybe pairing this with unsupervised deep learning? I want to have less of a “black box” for the development content!



Jason Brownlee July 11, 2017 at 10:34 am #

Yes, try and RFE:

<http://machinelearningmastery.com/feature-selection-machine-learning-python/>

Bernt July 13, 2017 at 10:12 pm #

REPLY ↩

Hi Jason,

Thank you for sharing your skills and competence.

I want to study the change in weights and predictions between each epoch run.

Have tried to use the `model.train_on_batch` method and the `model.fit` method with `epoch=1` and `batch_size`

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

But it seems like the model doesn't save the new updated weights.
I print predictions before and after I don't see a change in the evaluation scores.

Parts of the code is printed below.

Any idea?

Thanks.

```
# Compile model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

# evaluate the model
scores = model.evaluate(X, Y)
print("\n%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))

# Run one update of the model trained run with X and compared with Y
model.train_on_batch(X, Y)

# Fit the model
model.fit(X, Y, epochs=1, batch_size=768)

scores = model.evaluate(X, Y)
print("\n%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))
```



Jason Brownlee July 14, 2017 at 8:29 am #

Sorry, I have not explored evaluating a Keras model this way.

Perhaps it is a fault, I would recommend preparing the smallest possible example that demonstrates the issue and post to the Keras GitHub issues.

iman July 18, 2017 at 11:18 pm #

Hi, I tried to apply this to the titanic data set, however the predictions were all 0.4. What do you suggest for:

```
# create model
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

REPLY ↩

Get Your Start in Machine Learning


```
model = Sequential()
model.add(Dense(12, input_dim=4, activation='relu'))
model.add(Dense(4, activation='relu'))
model.add(Dense(1, activation='sigmoid'))

model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy']) # 'sgd'

model.fit(X, Y, epochs=15, batch_size=10)
```



Jason Brownlee July 19, 2017 at 8:26 am #

REPLY ↩

This post will give you some ideas to list the skill of your model:
<http://machinelearningmastery.com/improve-deep-learning-performance/>

Camus July 19, 2017 at 2:14 am #

Hi Dr Jason,

This is probably a stupid question but I cannot find out how to do it ... and I am beginner on Neural Netw
I have relatively same number of inputs (7) and one output. This output can take numbers between -300
I want to build a neural network model in python but I don't know how to do it.
Do you have an example with outputs different from 0-1.
Tanks in advance

Camus

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**
Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee July 19, 2017 at 8:28 am #

REPLY ↩

Ensure you scale your data then use the above tutorial to get started.

Khalid Hussain July 21, 2017 at 11:28 pm #

REPLY ↩

Hi Jason Brownlee

I am using the same data "pima-indians-diabetes.csv" but all predicted values are less than 1 and are in fraction which could not distinguish any class.

If I round off then all become 0.

I am using model.predict(x) function

You are requested to kindly guide me what I am doing wrong and how can I achieve correct predicted value.

Thank you



Jason Brownlee July 22, 2017 at 8:36 am #

Consider you have copied all of the code exactly from the tutorial.

Ludo July 25, 2017 at 6:59 pm #

Hello Jason,

Thanks you for your great example. I have some comments.

- Why you have choice "12" inputs hidden layers ? and not 24 / 32 .. it's arbitrary ?
- Same question about epochs and batch_size ?

This value are very sensible !! i have try with 32 inputs first layer , epochs=500 and batch_size=1000 and the result is very different... i'am at 65% accuracy.

Thx for you help.

Regards.

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Jason Brownlee July 26, 2017 at 7:50 am #

Get Your Start in Machine Learning



Yes, it is arbitrary. Tune the parameters of the model to your problem.

Almoutasem Bellah Rajab July 25, 2017 at 7:32 pm #

REPLY ↩

Wow, you're still replying to comments more than a year later!!!... you're great,, thanks..



Jason Brownlee July 26, 2017 at 7:50 am #

REPLY ↩

Yep.

Jane July 26, 2017 at 1:23 am #

Thanks for your tutorial, I found it very useful to get me started with Keras. I've previously tried to have a question for you though. I have both Theano and TensorFlow installed, how do I know which backend to use?



Jason Brownlee July 26, 2017 at 8:02 am #

Keras will print which backend it uses every time you run your code.

You can change the backend in the Keras configuration file (`~/.keras/keras.json`) which looks like:

```
1 {  
2     "image_data_format": "channels_last",  
3     "backend": "tensorflow",  
4     "epsilon": 1e-07,  
5     "floatx": "float32"  
6 }
```

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Masood Imran July 28, 2017 at 12:00 am #

REPLY ↩

Hello Jason,

My understanding of Machine Learning or evaluating deep learning models is almost 0. But, this article gives me lot of information. It is explained in a simple and easy to understand language.

Thank you very much for this article. Would you suggest any good read to further explore Machine Learning or deep learning models please?



Jason Brownlee July 28, 2017 at 8:31 am #

REPLY ↩

Thanks.

Yes, start right here:

<http://machinelearningmastery.com/start-here/#deeplearning>

Peggy August 3, 2017 at 7:14 pm #

If I have trained prediction models or neural network function scripts. How can I use them to make predictions for end users? I want to use python but it seems I will have to redo the training in Python again. Is there a way to reuse the training and just call the function of predicting?

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

REPLY ↩



Jason Brownlee August 4, 2017 at 6:58 am #

You need to train and save the final model then load it to make predictions.

This post will make it clear:

<http://machinelearningmastery.com/train-final-machine-learning-model/>

Get Your Start in Machine Learning

Shane August 8, 2017 at 2:38 pm #

REPLY ↩

Jason, I used your tutorial to install everything needed to run this tutorial. I followed your tutorial and ran the resulting program successfully. Can you please describe what the output means? I would like to thank you for your very informative tutorials.

Shane August 8, 2017 at 2:39 pm #

REPLY ↩

```
768/768 [=====] - 0s - loss: 0.4807 - acc: 0.7826
Epoch 148/150
768/768 [=====] - 0s - loss: 0.4686 - acc: 0.7812
Epoch 149/150
768/768 [=====] - 0s - loss: 0.4718 - acc: 0.7617
Epoch 150/150
768/768 [=====] - 0s - loss: 0.4772 - acc: 0.7812
32/768 [>.....] - ETA: 0s
acc: 77.99%
```



Jason Brownlee August 8, 2017 at 5:12 pm #

It is summarizing the training of the model.

The final line evaluates the accuracy of the model's predictions – really just to demonstrate how

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee August 8, 2017 at 5:11 pm #

REPLY ↩

Well done Shane.

Which output?

Get Your Start in Machine Learning

Bene August 9, 2017 at 1:02 am #

REPLY ↩

Hello Jason, i really liked your Work and it helped me a lot with my first steps.

But i am not really familiar with the numpy stuff:

So here is my Question:

```
dataset = numpy.loadtxt("pima-indians-diabetes.csv", delimiter=",")  
# split into input (X) and output (Y) variables  
X = dataset[:,0:8]  
Y = dataset[:,8]
```

I get that the numpy.loadtxt is extracting the information from the cvs File

but what does the stuff in the Brackets mean like X = dataset[:,0:8]

why the ":" and why , 0:8

its probably pretty dumb but i can't find a good explanation online 😊

thanks really much!



Jason Brownlee August 9, 2017 at 6:37 am #

Good question Bene, it's called array slicing:

<https://docs.scipy.org/doc/numpy/reference/arrays.indexing.html>

Bene August 9, 2017 at 10:59 pm #

REPLY ↩

That helped me out tank you Jason 😊

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Chen August 12, 2017 at 5:43 pm #

REPLY ↩

Can I translate it to Chinese and put it to Internet in order to let other Chinese people can read your article?



Jason Brownlee August 13, 2017 at 9:46 am #

REPLY ↩

No, please do not.

Deep Learning August 12, 2017 at 7:36 pm #

It seems that using this line:

```
np.random.seed(5)
```

...is redundant i.e. the Keras output in a loop running the same model with the same configuration will yield the same results all the time, regardless of which number it is set to. Or am I missing something?



Jason Brownlee August 13, 2017 at 9:52 am #

Deep learning algorithms are stochastic (random within a range). That means that they will produce different results each time the same model is trained on the same data. This is a feature:

<http://machinelearningmastery.com/randomness-in-machine-learning/>

You can fix the random seed to ensure you get the same result, and it is a good idea for tutorials to help beginners out:

<http://machinelearningmastery.com/reproducible-results-neural-networks-keras/>

When evaluating the skill of a model, I would recommend repeating the experiment n times and taking skill as the average of the runs. See here for the procedure:

<http://machinelearningmastery.com/evaluate-skill-deep-learning-models/>

Does that help?

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree**. Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Deep Learning August 14, 2017 at 3:08 am #

REPLY ↩

Thanks Jason 😊

I totally get what it should do, but as I had pointed out, it does not do it. If you run the codes you have provided above in a loop for say 10 times. First 10 with random seed set and the other 10 times without that line of code all together. Then compare the result. At least the result I'm getting, is suggesting the effect is not there i.e. both sets of 10 times will have similar variation in the result.



Jason Brownlee August 14, 2017 at 6:26 am #

It may suggest that the model is overprescribed and easily addresses the training data

Deep Learning August 14, 2017 at 3:12 am #

Nice post by the way > <http://machinelearningmastery.com/evaluate-skill-deep-learning-models>

Thanks for sharing it. Been lately thinking about the aspect of accuracy a lot, it seems that at the moment we are doing it out of the box. I think a lot of non PhD / non expert crowd (most people) will at least initially be easily convinced by your post.

Thanks for all the amazing contributions you are making in this field!

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee August 14, 2017 at 6:26 am #

REPLY ↩

I'm glad it helped.

Get Your Start in Machine Learning

RATNA NITIN PATIL August 14, 2017 at 8:16 pm #

REPLY ↩

Hello Jason, Thanks for a wonderful tutorial.

Can I use Genetic Algorithm for feature selection??

If yes, Could you please provide the link for it???

Thanks in advance.



Jason Brownlee August 15, 2017 at 6:34 am #

REPLY ↩

Sure. Sorry, I don't have any examples.

Generally, computers are so fast it might be easier to test all combinations in an exhaustive search.

sunny1304 August 15, 2017 at 3:44 pm #

Hi Json,

Thank you for your awesome tutorial.

I have a question for you.

Is there any guideline on how to decide on neuron number for our network.

for example you used 12 for thr 1st layer and 8 for the second layer.

how do you decide on that ?

Thanks

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee August 15, 2017 at 4:58 pm #

REPLY ↩

No, there is no way to analytically determine the configuration of the network.

I use trial and error. You can grid search, random search, or copy configurations from tutorials or pay

Get Your Start in Machine Learning

yihadad August 16, 2017 at 6:53 pm #

REPLY ↩

Hi Json,

Thanks for a wonderful tutorial.

Run a model generated by a CNN it takes how much ram, cpu ?

Thanks



Jason Brownlee August 17, 2017 at 6:39 am #

It depends on the data you are using to fit the model and the size of the model.

Very large models could be 500MB of RAM or more.

Ankur September 1, 2017 at 3:15 am #

Hi ,

Please let me know , how can i visualise the complete neural network in Keras.....

I am looking for the complete architecture – like number of neurons in the Input Layer, hidden layer , out

Please have a look at the link present below, here someone has created a beautiful visualisation/architecture using neuraienet package in R.

Please let me know, can we create such type of model in KERAS

<https://www.r-bloggers.com/fitting-a-neural-network-in-r-neuralnet-package/>

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee September 1, 2017 at 6:50 am #

REPLY ↩

Get Your Start in Machine Learning

Use the Keras visualization API:

<https://keras.io/visualization/>

ASAD October 17, 2017 at 3:23 am #

REPLY ↩

Hello ANKUR,,,, how are you?

you have try visualization in keras which is suggested by Jason Brownlee?
if you have tried then please send me code i am also trying but didnt work..

please guide me

Adam September 3, 2017 at 1:45 am #

Thank you Dr. Brownlee for the great tutorial,

I have a question about your code:
is the argument metrics=['accuracy'] necessary in the code and does it change the results of the neural network during compiling?

thank you!!

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**
Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee September 3, 2017 at 5:48 am #

No, it just prints out the accuracy of the model at the end of each epoch. Learn more about Keras metrics here:

<https://machinelearningmastery.com/custom-metrics-deep-learning-keras-python/>

PottOfGold September 5, 2017 at 12:14 am #

REPLY ↩

Get Your Start in Machine Learning

Hi Jason,

your work here is really great. It helped me a lot.
I recently stumbled upon one thing I cannot understand:

For the pimas dataset you state:

<>

When I look at the table of the pimas dataset, the examples are in rows and the features in columns, so your input dimension is the number of columns. As far as I can see, you don't change the table.

For neural networks, isn't the input normally: examples = columns, features=rows?

Is this different for Keras? Or can I use both shapes? An if yes, what's the difference in the construction of the net?

Thank you!!



Jason Brownlee September 7, 2017 at 12:36 pm #

No, features are columns, rows are instances or examples.

PottOfGold September 7, 2017 at 3:35 pm #

Thanks! 😊

I had a lot of discussions because of that.

In Andrew Ng new Coursera course it's explained as examples = columns, features=rows, but h
neural networks from scratch.



Jason Brownlee September 9, 2017 at 11:38 am #

I doubt that, I think you may have mixed it up. Columns are never examples.

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**
Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

REPLY ↩

Get Your Start in Machine Learning

PottOfGold October 6, 2017 at 6:26 pm #

Thats what I thought, but I looked it up in the notation for the new coursera course (deeplearning.ai) and there it says: m is the number of examples in the dataset and n is the input size, where X superscript $n \times m$ is the input matrix ...
But either way, you helped me! Thank you. 😊

Lin Li September 16, 2017 at 1:50 am #

REPLY ↩

Hi Jason, thank you so much for your tutorial, it helps me a lot. I need your help for the question below:
I copy the code and run it. Although I got the classification results, there were some warning messages in the process. As follows:

Warning (from warnings module):

File "C:\Users\llfor\AppData\Local\Programs\Python\Python35\lib\site-packages\keras\callbacks.py", line
% delta_t_median)

UserWarning: Method on_batch_end() is slow compared to the batch update (0.386946). Check your ca

I don't know why, and cannot find any answer to this question. I'm looking forward to your reply. Thanks



Jason Brownlee September 16, 2017 at 8:43 am #

Sorry, I have not seen this message before. It looks like a warning, you might be able to ignore it.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**
Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Lin Li September 16, 2017 at 12:24 pm #

REPLY ↩

Thanks for your reply. I'm a start-learner on deep learning. I'd like to put it aside temporarily.

Sagar September 22, 2017 at 2:51 pm #

Get Your Start in Machine Learning

Hi Jason,

Great article, thumbs up for that. I am getting this error when I try to run the file on the command prompt. Any suggestions. Thanks for you response.

```
#####
C:\Work\ML>python keras_first_network.py
Using TensorFlow backend.
2017-09-22 10:11:11.189829: W C:\tf_jenkins\home\workspace\rel-win\M\windows\PY\
36\tensorflow\core\platform\cpu_feature_guard.cc:45] The TensorFlow library wasn
't compiled to use AVX instructions, but these are available on your machine and
could speed up CPU computations.
2017-09-22 10:11:11.190829: W C:\tf_jenkins\home\workspace\rel-win\M\windows\PY\
36\tensorflow\core\platform\cpu_feature_guard.cc:45] The TensorFlow library wasn
't compiled to use AVX2 instructions, but these are available on your machine an
d could speed up CPU computations.
32/768 [>.....] - ETA: 0s
acc: 78.52%
#####
```



Jason Brownlee September 23, 2017 at 5:35 am #

Looks like warning messages that you can ignore.

Sagar September 24, 2017 at 3:52 am #

Thanks I got to know what the problem was. According to section 6 I had set verbose argument to 0 while calling "model.fit()". Now all the epochs are getting printed.



Jason Brownlee September 24, 2017 at 5:17 am #

Get Your Start in Machine Learning ×

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

DEDI V ↩

Get Your Start in Machine Learning

Glad to hear it.

Valentin September 26, 2017 at 6:35 pm #

REPLY ↩

Hi Jason,

Thanks for the amazing article . Clear and straightforward.
I had some problems installing Keras but was advised to prefix
with `tf.contrib.keras`
so I have code like

```
model=tf.contrib.keras.models.Sequential()  
Dense=tf.contrib.keras.layers.Dense
```

Now I try to train Keras on some small datafile to see how things work out:

```
1,1,0,0,8  
1,2,1,0,4  
1,0,0,1,5  
1,0,1,0,7  
0,1,0,0,8  
1,4,1,0,4  
1,0,2,1,1  
1,0,1,0,7
```

The first 4 columns are inputs and the 5-th column is output.
I use the same code for training (adjust number of inputs) as in your article,
but the network only gets to 12.5% accuracy.
Any advise?

Thanks,
Valentin

Get Your Start in Machine Learning



You can master applied Machine Learning
without the math or fancy degree.
Find out how in this *free* and *practical* email
course.

START MY EMAIL COURSE

Jason Brownlee September 27, 2017 at 5:40 am #

Get Your Start in Machine Learning



Thanks Valentin.

I have a good list of suggestions for improving model performance here:

<http://machinelearningmastery.com/improve-deep-learning-performance/>

Priya October 3, 2017 at 2:28 pm #

REPLY ↩

Hi Jason,

I tried replacing the pima data with random data as follows:

```
X_train = np.random.rand(18,61250)
X_test = np.random.rand(18,61250)
Y_train = np.array([0.0, 1.0, 1.0, 0.0, 1.0, 1.0, 1.0, 0.0, 1.0,
0.0, 1.0, 0.0, 0.0, 1.0, 0.0, 0.0, 0.0, 1.0,])
Y_test = np.array([1.0, 0.0, 0.0, 1.0, 1.0, 0.0, 1.0, 1.0, 1.0,
1.0, 0.0, 0.0, 0.0, 1.0, 0.0, 0.0, 1.0, 0.0,])

_, input_size = X_train.shape #put this in input_dim in the first dense layer
```

I took the round() off of the predictions so I could see the full value and then inserted my random test data

```
predictions = model.predict(X_test)
preds = [x[0] for x in predictions]
print(preds)
```

```
model.fit(X_train, Y_train, epochs=100, batch_size=10, verbose=2, validation_data=(X_test,Y_test))
```

I found something slightly odd; I expected the predicted values to be around 0.50, plus or minus some, but instead, I got this:

```
[0.49525392, 0.49652839, 0.49729034, 0.49670222, 0.49342978, 0.49490061, 0.49570397, 0.4962129, 0.49774086, 0.49475089, 0.4958384, 0.49506786,
0.49696651, 0.49869373, 0.49537542, 0.49613148, 0.49636957, 0.49723724]
```

which is near 0.50 but always less than 0.50. I ran this a few times with different random seeds, so it's not coincidental. Would you have any explanation for why it does this?

Thanks,
Priya

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning



Jason Brownlee October 3, 2017 at 3:46 pm #

REPLY ↩

Perhaps calculate the mean of your training data and compare it to the predicted value. It might be simple sampling error.



Priya October 4, 2017 at 1:02 am #

REPLY ↩

I found out I was doing predictions before fitting the model. (I suppose that would mean the network hadn't adjusted to the data's distribution yet.)



Saurabh October 7, 2017 at 5:59 am #

Hello Jason,

I tried to train this model on my laptop, it is working fine. But I tried to train this model on google-cloud with keras and tensorflow and it is failing.

Can you just let me know, which changes are to required for the model, so that I can train this on cloud.

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee October 7, 2017 at 7:37 am #

Sorry, I don't know about google cloud.

I have instructions here for running on AWS:

<https://machinelearningmastery.com/develop-evaluate-large-deep-learning-models-keras-amazon-web-services/>

tobegit3hub October 12, 2017 at 6:40 pm #

REPLY ↩

Get Your Start in Machine Learning

Great post. Thanks for sharing.



Jason Brownlee October 13, 2017 at 5:45 am #

REPLY ↩

You're welcome.

Manoj October 12, 2017 at 11:43 pm #

REPLY ↩

Hi Jason,

Is there a way to store the model, once it is created so that I can use it for different input data sets as an



Jason Brownlee October 13, 2017 at 5:48 am #

Yes, you can save it to file. See this tutorial:

<https://machinelearningmastery.com/save-load-machine-learning-models-python-scikit-learn/>

Cam October 23, 2017 at 6:11 pm #

I get a syntax error for the

model.fit() line in this example. Is it due to library conflicts with theano and tensorflow if i have both installed?



Jason Brownlee October 24, 2017 at 5:28 am #

REPLY ↩

Perhaps ensure your environment is up to date and that you copied the code exactly.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree**. Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

This tutorial can help with setting up your environment:

<http://machinelearningmastery.com/setup-python-environment-machine-learning-deep-learning-anaconda/>

Cam October 24, 2017 at 2:11 pm #

REPLY ↩

Thanks, fixed!



Jason Brownlee October 24, 2017 at 4:01 pm #

REPLY ↩

Glad to hear it.

Diego Quintana October 25, 2017 at 7:37 am #

Hi Jason, thanks for the example.

How would you predict a single element from X? X[0] raises a ValueError

ValueError: Error when checking : expected dense_1_input to have shape (None, 8) but got array with s

Thanks!

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee October 25, 2017 at 3:56 pm #

REPLY ↩

You can reshape it to have 1 row and 8 columns:

```
1 X = X.reshape((1,8))
```

This post will give you further advice:

<https://machinelearningmastery.com/index-slice-reshape-numpy-arrays-machine-learning-python/>

Get Your Start in Machine Learning

Shahbaz Wasti October 28, 2017 at 1:30 pm #

REPLY ↩

Dear Sir ,

I have installed and configured the environment according to your directions but while running the program i have following error

“from keras.utils import np_utils”



Jason Brownlee October 29, 2017 at 5:50 am #

REPLY ↩

What is the error exactly?

Zhengping October 30, 2017 at 12:12 am #

Hi Jason, thanks for the great tutorials. I just learnt and repeated the program in your “Your First Neural Network in Python With Keras Step-By-Step” without problem. Now trying this one, getting stuck at the line “model = Sequential()” when the InteractiveModel class is not defined. tried to google, can’t find a solution. I did import Sequential from keras.models as in ur example. Can you please provide your help.

Zhengping October 30, 2017 at 12:14 am #

I’m running ur examples in Anaconda 4.4.0 environment in visual studio community version. relevant packages have been installed as in ur earlier tutorials instructed.

Zhengping October 30, 2017 at 12:18 am #

REPLY ↩

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
>> # create model
... model = Sequential()
...
Traceback (most recent call last):
File "", line 2, in
NameError: name 'Sequential' is not defined
>>> model.add(Dense(12, input_dim=8, init='uniform', activation='relu'))
Traceback (most recent call last):
File "", line 1, in
AttributeError: 'SVC' object has no attribute 'add'
```



Jason Brownlee October 30, 2017 at 5:39 am #

This does not look good. Perhaps post the error to stack exchange or other keras s
<https://machinelearningmastery.com/get-help-with-keras/>



Jason Brownlee October 30, 2017 at 5:38 am #

Looks like you need to install Keras. I have a tutorial here on how to do that:
<https://machinelearningmastery.com/setup-python-environment-machine-learning-deep-learning-and-keras/>

Akhil October 30, 2017 at 5:04 pm #

Ho Jason,

Thanks a lot for this wonderful tutorial.

I have a question:

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

REPLY ↩

I want to use your code to predict the classification (1 or 0) of unknown samples. Should I create one common csv file having the train (known) as well as the test (unknown) data. Whereas the 'classification' column for the known data will have a known value, 1 or 0, for the unknown data, should I leave the column empty (and let the code decide the outcome)?

Thanks a lot



Jason Brownlee October 31, 2017 at 5:29 am #

REPLY ↩

Great question.

No, you only need the inputs and the model can predict the outputs, call `model.predict(X)`.

Also, this post will give a general idea on how to fit a final model:

<https://machinelearningmastery.com/train-final-machine-learning-model/>

Guilherme November 3, 2017 at 1:26 am #

Hi Jason,

This is really cool! I am blown away! Thanks so much for making it so simple for a beginner to have some

1) where are the weights, can I save and/or retrieve them?

2) if I want to train images with dogs and cats and later ask the neural network whether a new image has as an array and my output result to be "cat" or "dog"?

Thanks again and great job!

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree**. Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee November 3, 2017 at 5:20 am #

REPLY ↩

The weights are in the model, you can save them:

<https://machinelearningmastery.com/save-load-keras-deep-learning-models/>

Get Your Start in Machine Learning

Yes, you would save your model, then call `model.predict()` on the new data.

Michael November 5, 2017 at 8:33 am #

REPLY ↩

Hi Jason,

Are you familiar with a python tool/package that can build neural network as in the tutorial, but suitable for data stream mining?

Thanks,
Michael



Jason Brownlee November 6, 2017 at 4:46 am #

Not really, sorry.

bea November 8, 2017 at 1:58 am #

Hi, there. Could you please clarify why exactly you've built your network with 12 neurons in the
"The first layer has 12 neurons and expects 8 input variables. The second hidden layer has 8 neurons at
class (onset of diabetes or not)..."

Shouldn't it have 8 neurons at the start?

Thanks



Jason Brownlee November 8, 2017 at 9:28 am #

REPLY ↩

The input layer has 8, the first hidden layer has 12. I chose 12 through a little trial and error

Get Your Start in Machine Learning



You can master applied Machine Learning
without the math or fancy degree.
Find out how in this *free* and *practical* email
course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

Guilherme November 9, 2017 at 12:54 am #

REPLY ↩

Hi Jason,

Do you have or else could you recommend a beginner's level image segmentation approach that uses deep learning? For example, I want to train some neural net to automatically "find" a particular feature out of an image.

Thanks!



Jason Brownlee November 9, 2017 at 10:00 am #

Sorry, I don't have image segmentation examples, perhaps in the future.

Andy November 12, 2017 at 6:56 pm #

Hi Jason,

I just started my DL training a few weeks ago. According to what I learned in course, in order to train the and Backward propagation; however, looking at your Keras example, i don't find any of these propagatic mechanism to find the parameters instead of using Forward and Backward propagation?

Thanks!

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee November 13, 2017 at 10:13 am #

REPLY ↩

It is performing those operations under the covers for you.

Get Your Start in Machine Learning

Badr November 13, 2017 at 11:42 am #

REPLY ↩

Hi Jason,

Can you explain why I got the following output:

ValueError Traceback (most recent call last)

in ()

—> 1 model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])

2 model.fit(X, Y, epochs=150, batch_size=10)

3 scores = model.evaluate(X, Y)

4 print("\n%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))

/Users/badrshomrani/anaconda/lib/python3.5/site-packages/keras/models.py in compile(self, optimizer, l

545 metrics=metrics,

546 sample_weight_mode=sample_weight_mode,

—> 547 **kwargs)

548 self.optimizer = self.model.optimizer

549 self.loss = self.model.loss

/Users/badrshomrani/anaconda/lib/python3.5/site-packages/keras/engine/training.py in compile(self, opt

sample_weight_mode, **kwargs)

620 loss_weight = loss_weights_list[i]

621 output_loss = weighted_loss(y_true, y_pred,

—> 622 sample_weight, mask)

623 if len(self.outputs) > 1:

624 self.metrics_tensors.append(output_loss)

/Users/badrshomrani/anaconda/lib/python3.5/site-packages/keras/engine/training.py in weighted(y_true, y_pred, weights, mask)

322 def weighted(y_true, y_pred, weights, mask=None):

323 # score_array has ndim >= 2

—> 324 score_array = fn(y_true, y_pred)

325 if mask is not None:

326 # Cast the mask to floatX to avoid float64 upcasting in theano

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
/Users/badrshomrani/anaconda/lib/python3.5/site-packages/keras/objectives.py in binary_crossentropy(y_true, y_pred)
46
47 def binary_crossentropy(y_true, y_pred):
--> 48 return K.mean(K.binary_crossentropy(y_pred, y_true), axis=-1)
49
50

/Users/badrshomrani/anaconda/lib/python3.5/site-packages/keras/backend/tensorflow_backend.py in binary_crossentropy(output, target, from_logits)
1418 output = tf.clip_by_value(output, epsilon, 1 - epsilon)
1419 output = tf.log(output / (1 - output))
-> 1420 return tf.nn.sigmoid_cross_entropy_with_logits(output, target)
1421
1422

/Users/badrshomrani/anaconda/lib/python3.5/site-packages/tensorflow/python/ops/nn_impl.py in sigmoid_cross_entropy_with_logits(name)
147 # pylint: disable=protected-access
148 nn_ops._ensure_xent_args("sigmoid_cross_entropy_with_logits", _sentinel,
-> 149 labels, logits)
150 # pylint: enable=protected-access
151

/Users/badrshomrani/anaconda/lib/python3.5/site-packages/tensorflow/python/ops/nn_ops.py in _ensure_xent_args(labels, logits, name)
1696 if sentinel is not None:
1697 raise ValueError("Only call %s with "
-> 1698 "named arguments (labels=..., logits=..., ...) " % name)
1699 if labels is None or logits is None:
1700 raise ValueError("Both labels and logits must be provided.")

ValueError: Only call sigmoid_cross_entropy_with_logits with named arguments (labels=..., logits=..., ...)
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee November 14, 2017 at 10:05 am #

REPLY ↩

Perhaps double check you have the latest versions of the keras and tensorflow libraries installed

Get Your Start in Machine Learning

Badr November 14, 2017 at 10:50 am #

REPLY ↩

keras was outdated



Jason Brownlee November 15, 2017 at 9:44 am #

REPLY ↩

Glad to hear you fixed it.

Mikael November 22, 2017 at 8:20 am #

Hi Jason, thanks for your short tutorial, helps a lot to actually get your hands dirty with a simple
I have tried 5 different parameters and got some interesting results to see what would happen. Unfortunately

Test 1 Test 2 Test 3 Test 4 Test 5 Test 6 Test 7

number of layers 3 3 3 3 3 3 4

Train set 768 768 768 768 768 768 768

Iterations 150 100 1000 1000 1000 150 150

Rate of update 10 10 10 5 1 1 5

Errors 173 182 175 139 161 169 177

Values 768 768 768 768 768 768 768

% Error 23,0000% 23,6979% 22,7865% 18,0990% 20,9635% 22,0052% 23,0469%

I can't seem to see a trend here.. That could put me on the right track to adjust my hyperparameters.

Do you have any advice on that?

Get Your Start in Machine Learning



You can master applied Machine Learning
without the math or fancy degree.
Find out how in this *free* and *practical* email
course.

START MY EMAIL COURSE



Jason Brownlee November 22, 2017 at 11:17 am #

Get Your Start in Machine Learning

Something is wrong. Here is a good list of things to try:

<http://machinelearningmastery.com/improve-deep-learning-performance/>

Nikolaos November 28, 2017 at 10:58 am #

REPLY ↩

Hi, I try to implement the above example with fer2013.csv but I receive an error, it is possible to help me to implement this correctly?

```
1 keras.models import Sequential
2 from keras.layers import Dense
3 import numpy
4 import numpy as np
5
6 # fix Random seed for reproducibility
7 numpy.random.seed(7)
8 Y = []
9 X = []
10 #load dataset
11 for line in open("fer2013.csv"):
12     row = line.split(',')
13     Y.append(int(row[0]))
14     X.append([int(p) for p in row[1].split()])
15 X, Y = np.array(X) / 255.0, np.array(Y)
16 print(Y.shape)
17 print(X.shape)
18
19
20 #create model
21 model = Sequential()
22 model.add(Dense(12, input_dim=(35887, 2304), activation='tanh'))
23 model.add(Dense(8, activation='tanh'))
24 model.add(Dense(1, activation='sigmoid'))
25
26 #Compile Model
27 model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
28
29 #Fit Model
30 model.fit(X, Y, epochs=150, batch_size=1)
31
32 # evaluate the model
33 scores = model.evaluate(X, Y)
34 print("\n%s: %.2f%%" % (model.metrics_names[1], scores[1]*100))
35
36 # calculate predictions
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
37 predictions = model.predict(X)
38 # round predictions
39 rounded = [round(x[0]) for x in predictions]
40 print(rounded)
```



Jason Brownlee November 29, 2017 at 8:10 am #

REPLY ↩

Sorry, I cannot debug your code.

What is the problem exactly?

Tanya December 2, 2017 at 12:06 am #

Hello,

i have a a bit general question.

I have to do a forecasting for restaurant sales (meaning that I have to predict 4 meals based on a historical temperature, rain, etc), official holiday and in-off-season. I have to perform that forecasting using neuron. I am unfortunately not a very skilled in python. On my computer I have Python 2.7 and I have install anaconda. Mr. Brownlee. But somehow I can not run the code at all (in Spyder). Can you tell me what kind of version of python I need on my computer and in which environment (jupyterlab,notebook,qtconsole, spyder, etc) I can run the code, so that I will be very thankful for your response

KG

Tanya

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee December 2, 2017 at 9:02 am #

REPLY ↩

Perhaps this tutorial will help you setup and confirm your environment:

<http://machinelearningmastery.com/setup-python-environment-machine-learning-deep-learning-anaconda/>

I would also recommend running code from the command line as IDEs and notebooks can introduce

Get Your Start in Machine Learning

Eliah December 3, 2017 at 10:53 am #

REPLY ↩

Hi Dr. Brownlee.

I looked over the tutorial and I had a question regarding reading the data from a binary file? For instance I working on solving the sliding tiled n-puzzle using neural networks, but I seem to have trouble getting my data which is in a binary file and it generates the number of move required for the n-puzzle to be solve in. Am not sure if you have dealt with this before, but any help would be appreciated.



Jason Brownlee December 4, 2017 at 7:43 am #

REPLY ↩

Sorry, I don't know about your binary file.

Perhaps after you load your data, you can convert it to a numpy array so that you can provide it to a

Eliah December 4, 2017 at 9:28 am #

Thanks for the tip, I'll try it.

Wafaa December 7, 2017 at 4:59 pm #

Thank you very very much for all your great tutorials.

If I wanted to add batch layer after the input layer, how should I do it?

Cuz I applied this tutorial on a different dataset and features and I think I need normalization or standardization and I want to do it the easiest way.

Thank you,

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning



Jason Brownlee December 8, 2017 at 5:35 am #

REPLY ↩

I recommend preparing the data prior to fitting the model.

zaheer December 9, 2017 at 3:03 am #

REPLY ↩

thanks for sharing such nice tutorials, it helped me alot. i want to print the confusion matrix from the above example. and one more question.

if i have

20-input variable

1- class label (binary)

and 400 instances

how i would know , setting up the dense layer parameter in the first layer and hidden layer and output layer



Jason Brownlee December 9, 2017 at 5:44 am #

I recommend trial and error to configure the number of neurons in the hidden layer to see what works best.

zaheer December 9, 2017 at 3:29 am #

C:\Users\zaheer\AppData\Local\Programs\Python\Python36\python.exe C:/Users/zaheer/PycharmProjects/PythonBegin/Bin-CLNCL-Copy.py

Using TensorFlow backend.

Traceback (most recent call last):

File "C:/Users/zaheer/PycharmProjects/PythonBegin/Bin-CLNCL-Copy.py", line 28, in

model.fit(x_train , y_train , epochs=100, batch_size=100)

File "C:\Users\zaheer\AppData\Local\Programs\Python\Python36\lib\site-packages\keras\models.py", line 960, in fit

validation_steps=validation_steps)

File "C:\Users\zaheer\AppData\Local\Programs\Python\Python36\lib\site-packages\keras\engine\training.py", line 1574, in fit

batch_size=batch_size)

File "C:\Users\zaheer\AppData\Local\Programs\Python\Python36\lib\site-packages\keras\engine\training.py", line 1574, in fit

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
exception_prefix='input')
```

```
File "C:\Users\zaheer\AppData\Local\Programs\Python\Python36\lib\site-packages\keras\engine\training.py", line 153, in _standardize_input_data  
str(array.shape))
```

```
ValueError: Error when checking input: expected dense_1_input to have shape (None, 20) but got array with shape (362, 1)
```



Jason Brownlee December 9, 2017 at 5:45 am #

REPLY ↩

Ensure the input shape matches your data.

Anam Zahra December 10, 2017 at 7:40 pm #

Dear Jason! Great job a very simple guide.
I am trying to run the exact code but there is an error
`str(array.shape))`

ValueError: Error when checking target: expected dense_3 to have shape (None, 1) but got array with shape (362, 1)

How can I resolve.

I have windows 10 and spyder.

Get Your Start in Machine Learning



You can master applied Machine Learning
without the math or fancy degree.
Find out how in this *free* and *practical* email
course.

START MY EMAIL COURSE



Jason Brownlee December 11, 2017 at 5:24 am #

Sorry to hear that, perhaps confirm that you have the latest version of Numpy and Keras installed?

nazek hassouneh December 11, 2017 at 7:33 am #

REPLY ↩

Get Your Start in Machine Learning

after run this code , i will calculate the accuracy , how i did , i
i want to split the data set into test data , training data
and evaluate the model and calculate the accuracy
thank dr.

Suchith December 21, 2017 at 2:35 pm #

REPLY ↩

In the model how many hidden layers are there ?



Jason Brownlee December 21, 2017 at 3:35 pm #

There are 2 hidden layers, 1 input layer and 1 output layer.

Amare Mahtesenu December 22, 2017 at 9:55 am #

hi there. this blog is very awesome like the Adrian's pyimagesearch blog. I have one question a
keras frame work with SSD or Yolo architechtures?



Jason Brownlee December 22, 2017 at 4:16 pm #

Thanks for the suggestion, I hope to cover them in the future.

Get Your Start in Machine Learning



You can master applied Machine Learning
without the math or fancy degree.

Find out how in this *free* and *practical* email
course.

START MY EMAIL COURSE

Leave a Reply

Get Your Start in Machine Learning

Name (required)

Email (will not be published) (required)

Website

Welcome to Machine Learning Mastery



Hi, I'm Dr. Jason Brownlee.
My goal is to make practitioners like YOU awesome at applied machine learning.

[Read More](#)

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**
Find out how in this *free* and *practical* email course.


Finally Get Started With Deep Learning

Sick of the fancy math and need for super computers?
Looking for step-by-step tutorials?

[Get Your Start in Machine Learning](#)

Want end-to-end projects?

Get Started With Deep Learning in Python Today!



POPULAR

**Time Series Prediction with LSTM Recurrent Neural Networks in Python with Keras**

JULY 21, 2016

**Your First Machine Learning Project in Python Step-By-Step**

JUNE 10, 2016

**Develop Your First Neural Network in Python With Keras Step-By-Step**

MAY 24, 2016

**Multivariate Time Series Forecasting with LSTMs in Keras**

AUGUST 14, 2017

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

[START MY EMAIL COURSE](#)[Get Your Start in Machine Learning](#)



How to Setup a Python Environment for Machine Learning and Deep Learning with Anaconda

MARCH 13, 2017



Sequence Classification with LSTM Recurrent Neural Networks in Python with Keras

JULY 26, 2016



Time Series Forecasting with the Long Short-Term Memory Network in Python

APRIL 7, 2017



Regression Tutorial with the Keras Deep Learning Library in Python

JUNE 9, 2016



Multi-Class Classification Tutorial with the Keras Deep Learning Library

JUNE 2, 2016



How to Grid Search Hyperparameters for Deep Learning Models in Python With Keras

AUGUST 9, 2016

Get Your Start in Machine Learning



You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

© 2018 Machine Learning Mastery. All Rights Reserved.

[Privacy](#) | [Contact](#) | [About](#)

Get Your Start in Machine Learning