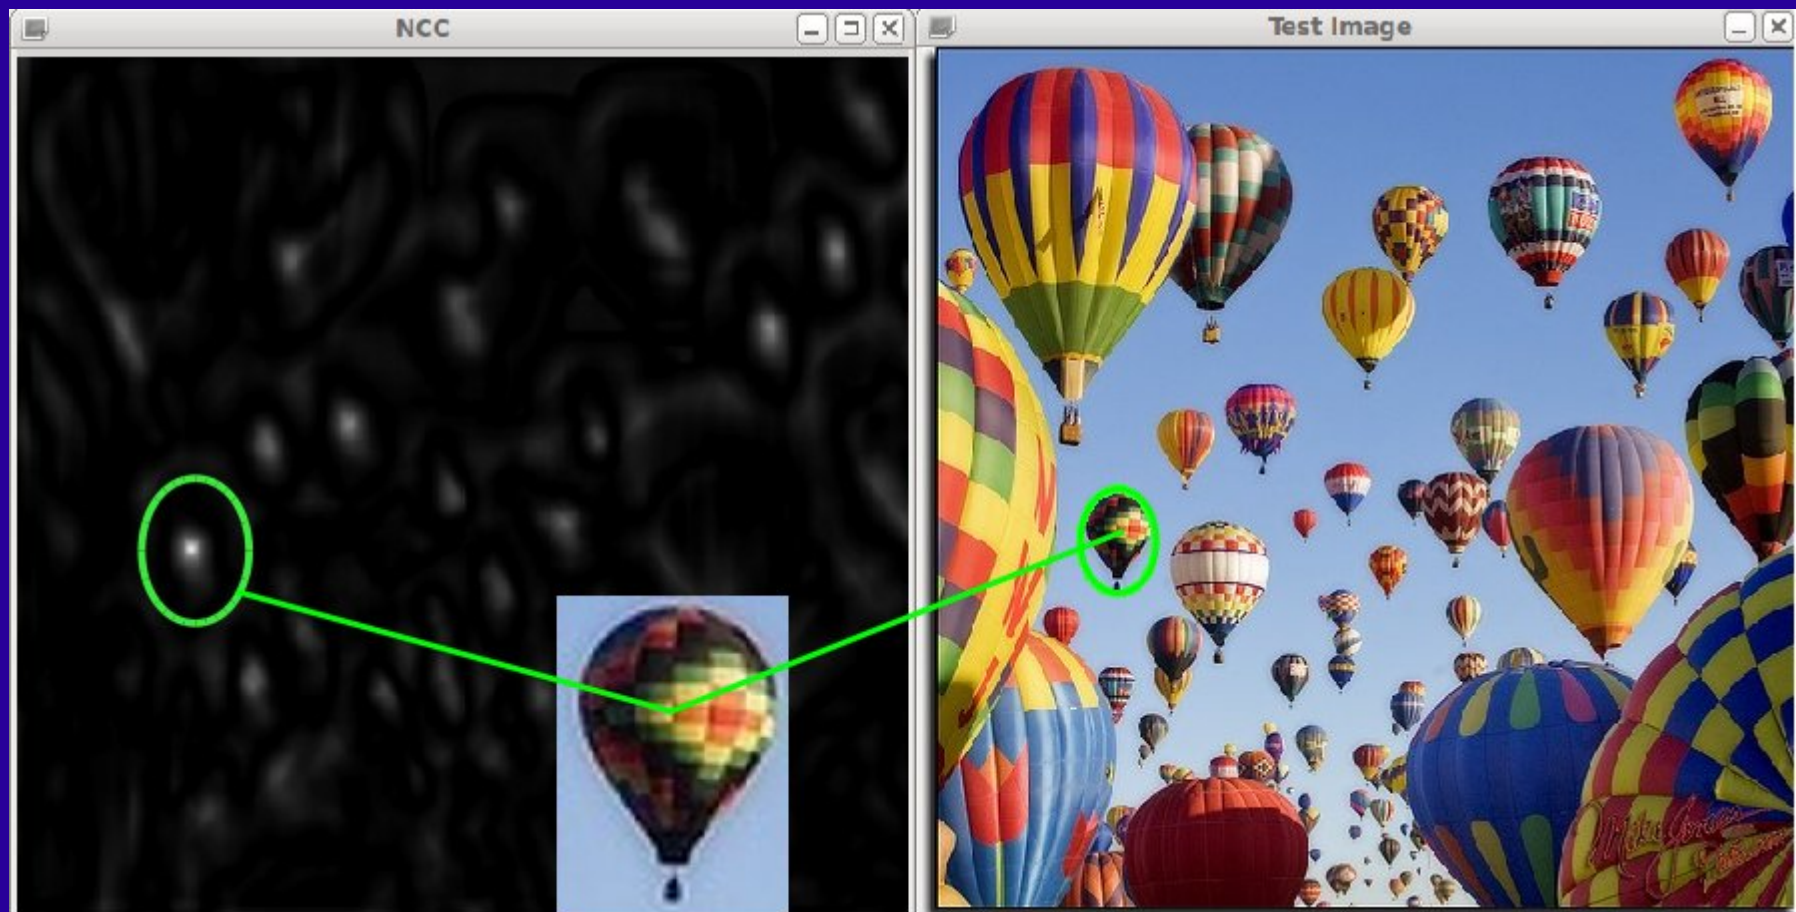


# Computer Vision & Machine Learning for Robots




# Where's Waldo?



# OpenCV Template Match

*30 milliseconds*



# Visual Pattern Matching

- Household objects; e.g. mug on a table, beer in the fridge
- People, faces, pets
- Landmarks for navigation; e.g. AR markers (fiducials)
- Keypoints for Visual-SLAM

# Strategies for Object Detection & Recognition

- Template Matching
- Feature Matching
- Machine Learning (learning from examples)

# Detection vs Recognition

- Detection: is there a face, any face, in this picture? E.g., OpenCV Haar face detector
- Recognition: **who** is that face in the picture? E.g. Eigenfaces

# Vision Software

- OpenCV – Linux, Windows, MacOS X, Android
  - C, C++, Python, no GUI )-:
  - Advanced vision algorithms including SIFT, SURF, face detection, machine learning
- RoboRealm – Windows Only
  - C, C++, Python, VBScript, very nice GUI
  - Many, many vision filters
  - Control for many popular robots and cameras
  - None of the patent-protected algorithms

# Software Continued...


- PCL (Point Cloud Library)

- Linux, Windows, MacOS X, Android
- C++, no Python, no GUI )-:
- 3D planar segmentation, clustering
- [www.pointclouds.org](http://www.pointclouds.org)



- Ecto – Linux only?

- Combine OpenCV & PCL using Python image processing pipelines
- Easily extendible
- [ecto.willowgarage.com](http://ecto.willowgarage.com)





# Appearances are Everything




# Special Challenges for Robotics

- Live video rather than static pictures
- Changes in orientation and position (occlusions, rotations, scale, affine transformations)
- Changes in lighting and background
- Motion of object and/or camera
- Changing appearance (today Joe has on a hat and glasses)
- The world is 3D while pictures are 2D (Kinect/Xtion)
- Pose estimation (e.g. for grasping)
- Speed: at least 10-30 fps


# Facing the Challenge

Courtesy of Honda/UCSD Video Database






# Haar Face Detector on Video





# Template Matching

317x480 pixels, 40 *milliseconds*



# Using pyrDown() Twice




*under 10 milliseconds*



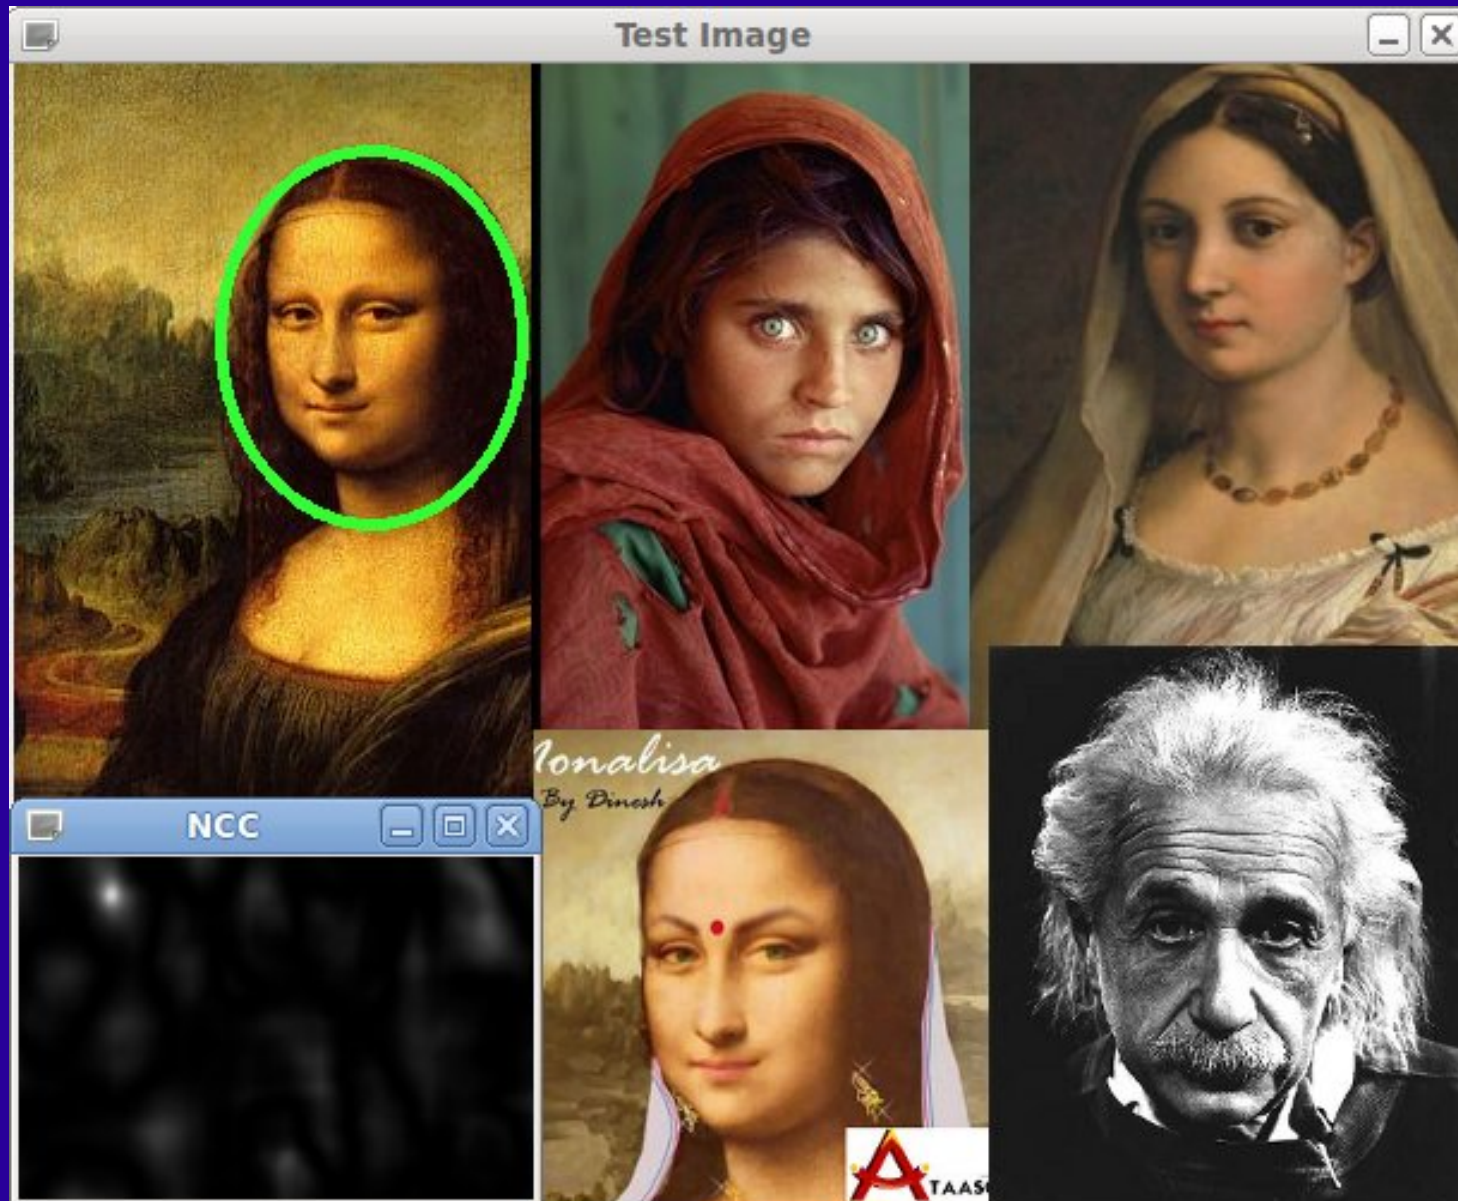


# Template Matching Including Scale & Rotation

*230 milliseconds*




# Template is Object Specific







# Template Matching on Video



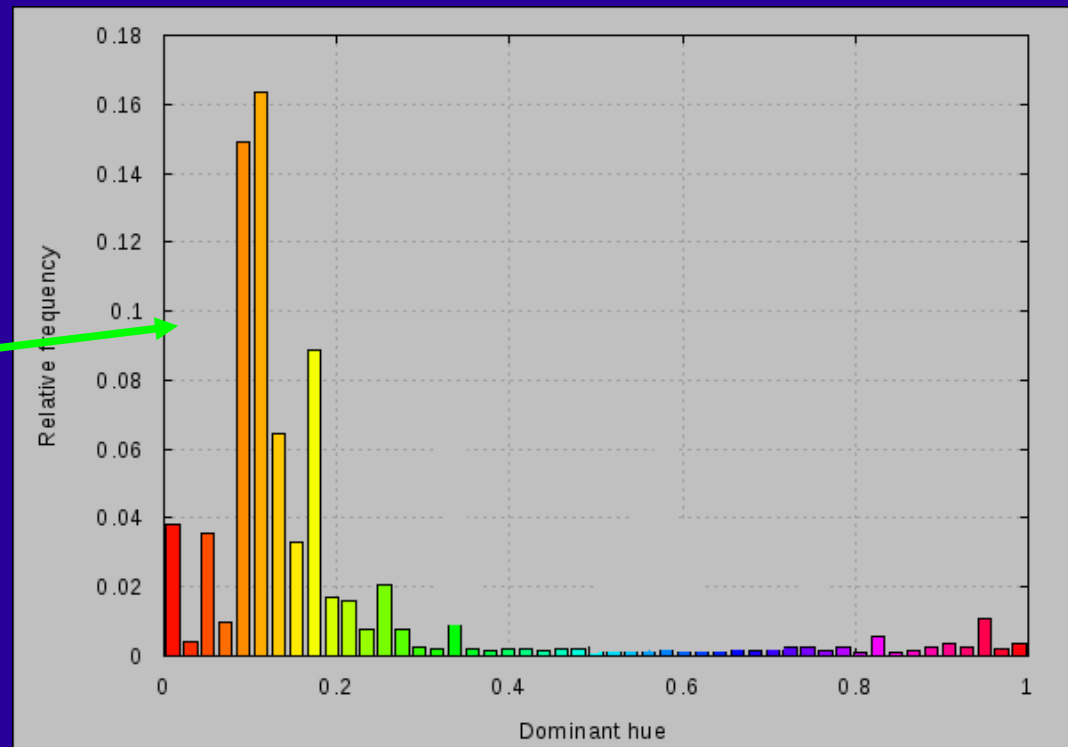
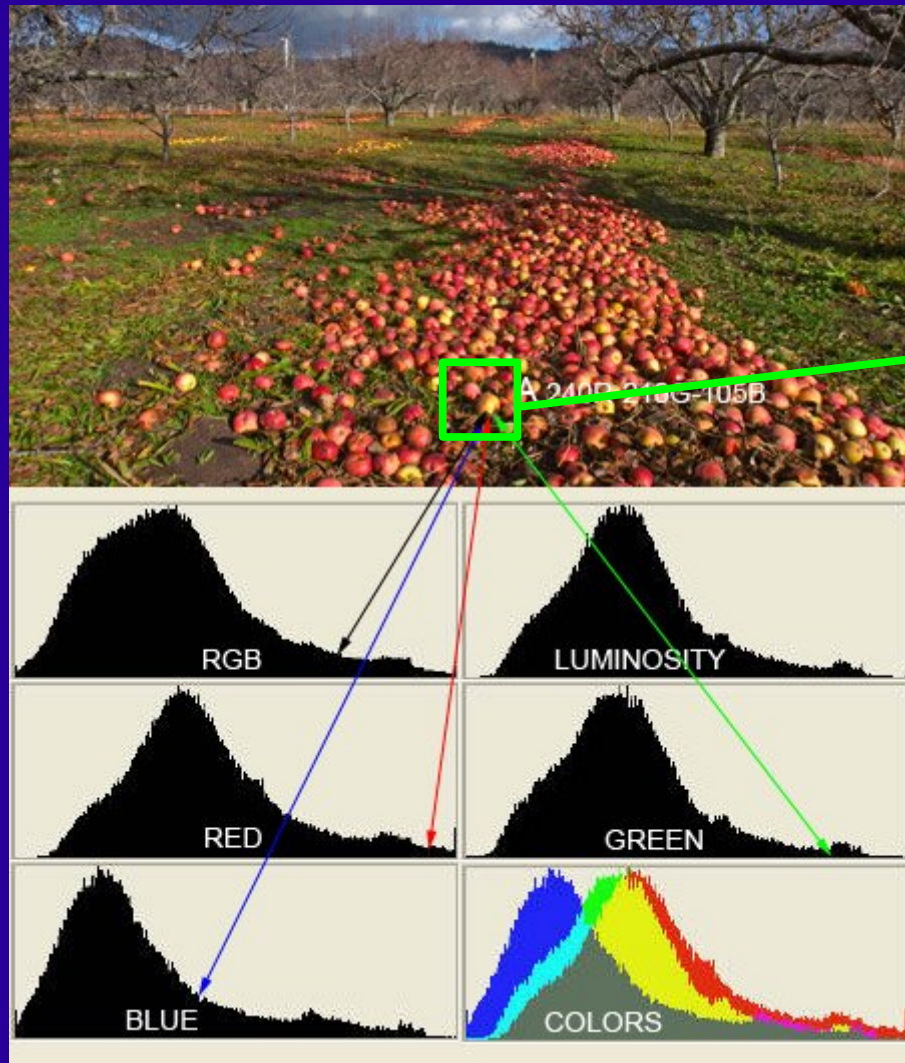
# Template Limitations



# From Pixels to Features

- What is a feature? A summary statistic or descriptor computed from pixel values in an image patch, region of interest (ROI), or even the whole image
- Colors → histograms (e.g. RGB, hue)
- Pixel intensities → gradients, binary comparisons, histograms, interest points (well defined locations)
- Connected regions → contours, segmentation
- Ideally, invariant to scale and rotations



# Color Histograms



0.038 0.01 0.036 0.011 0.15 ... 0.05 0.062 0.041 0.002 0.005




# OpenCV CamShift Tracking



# Interest Point Detection


## Harris Corners, SUSAN, FAST




$$\mathbf{M} = \begin{bmatrix} \left(\frac{\partial f}{\partial x}\right)^2 & \frac{\partial^2 f}{\partial x \partial y} \\ \frac{\partial^2 f}{\partial x \partial y} & \left(\frac{\partial f}{\partial y}\right)^2 \end{bmatrix}$$

$$R = \det \mathbf{M} - k \operatorname{Trace}^2 \mathbf{M}$$

# Interest Points = Keypoints






# Another Keypoint Example







# Keypoint Descriptors



0.038 0.01 0.036 0.011 0.15 ... 0.05 0.062 0.041 0.002 0.005

# Descriptors Live in Feature Space

Distance Measures: Euclidean, Manhattan, Hamming




# Descriptors Available in OpenCV


- SIFT – Scale Invariant Feature Transform
- SURF – Speed Up Robust Features
- BRIEF – Binary Robust Independent Features
- ORB – Oriented FAST + Rotated BRIEF

# OpenCV SURF Keypoints

58 keypoints  
10 ms





474 keypoints  
89 ms








# SURF Keypoint/Descriptor Matching

## FLANN + RANSAC, 85 *ms*



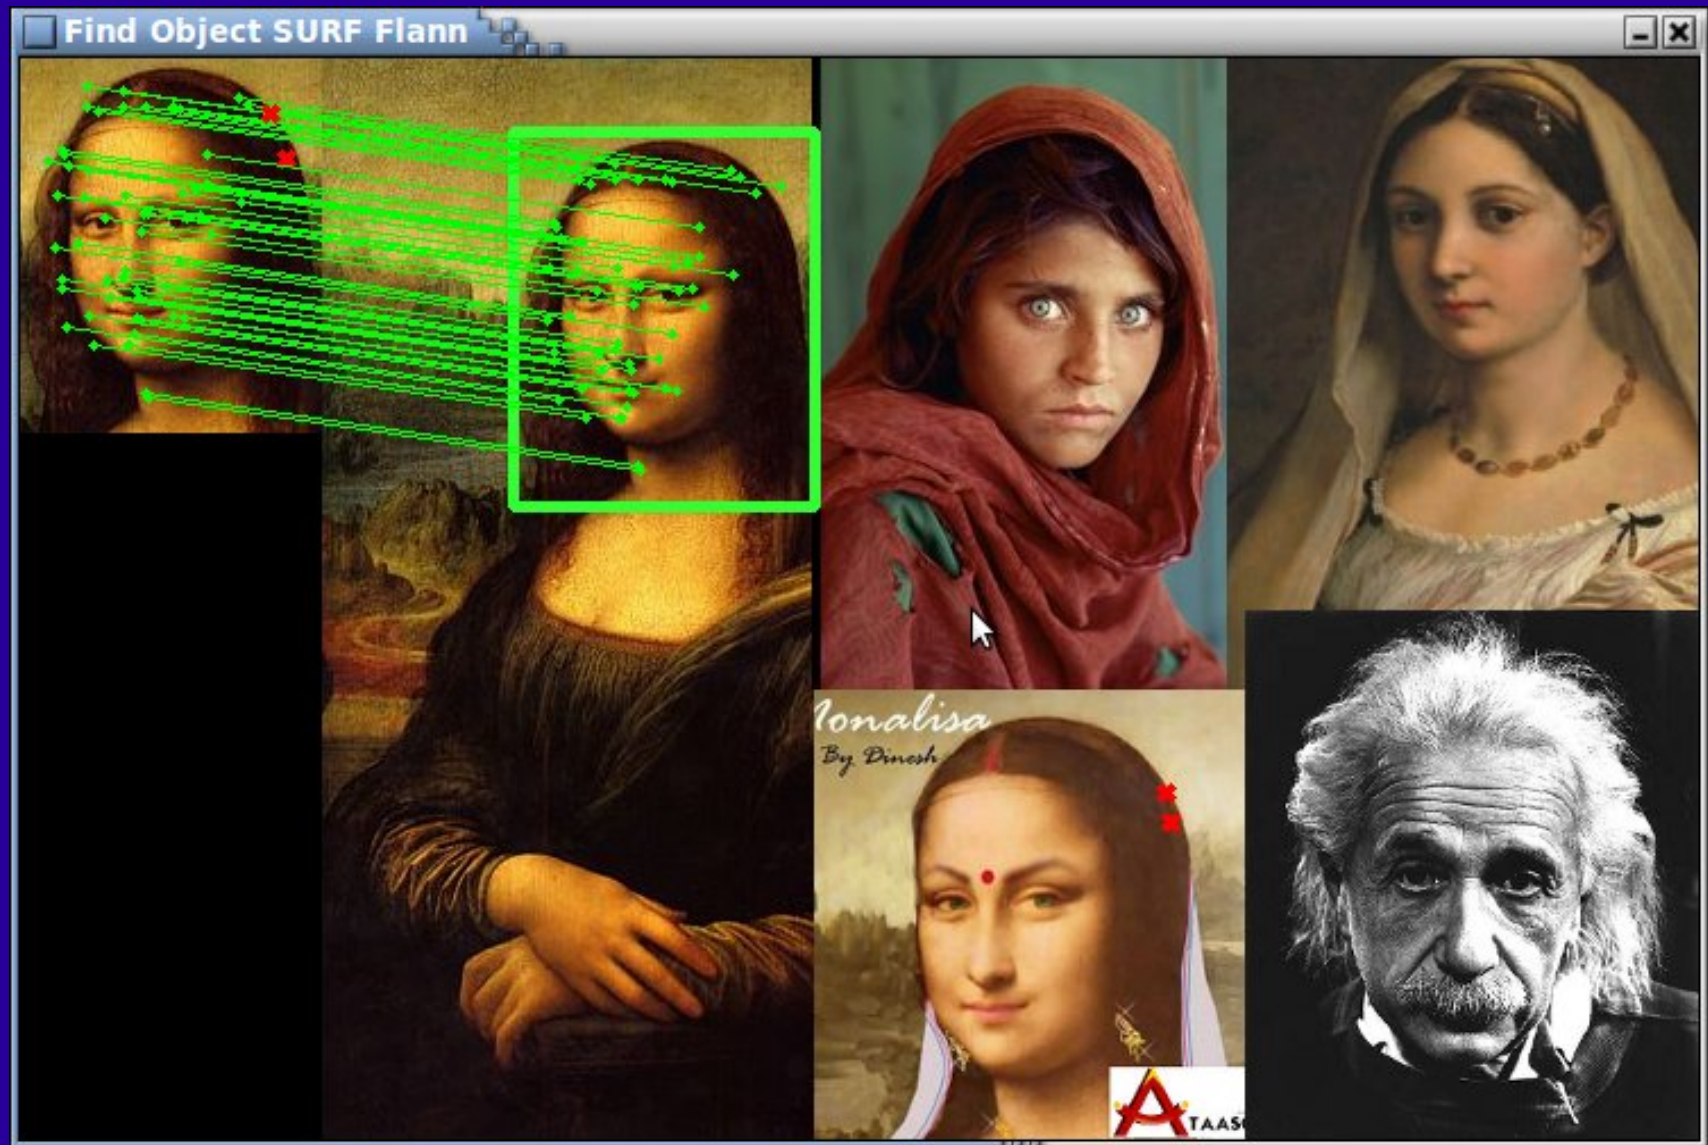
# SURF Matching Examples

Upright = False, 180 ms






# SURF is (Partially) Object Specific



# Partial SURF Match






# Finding Waldo with SURF

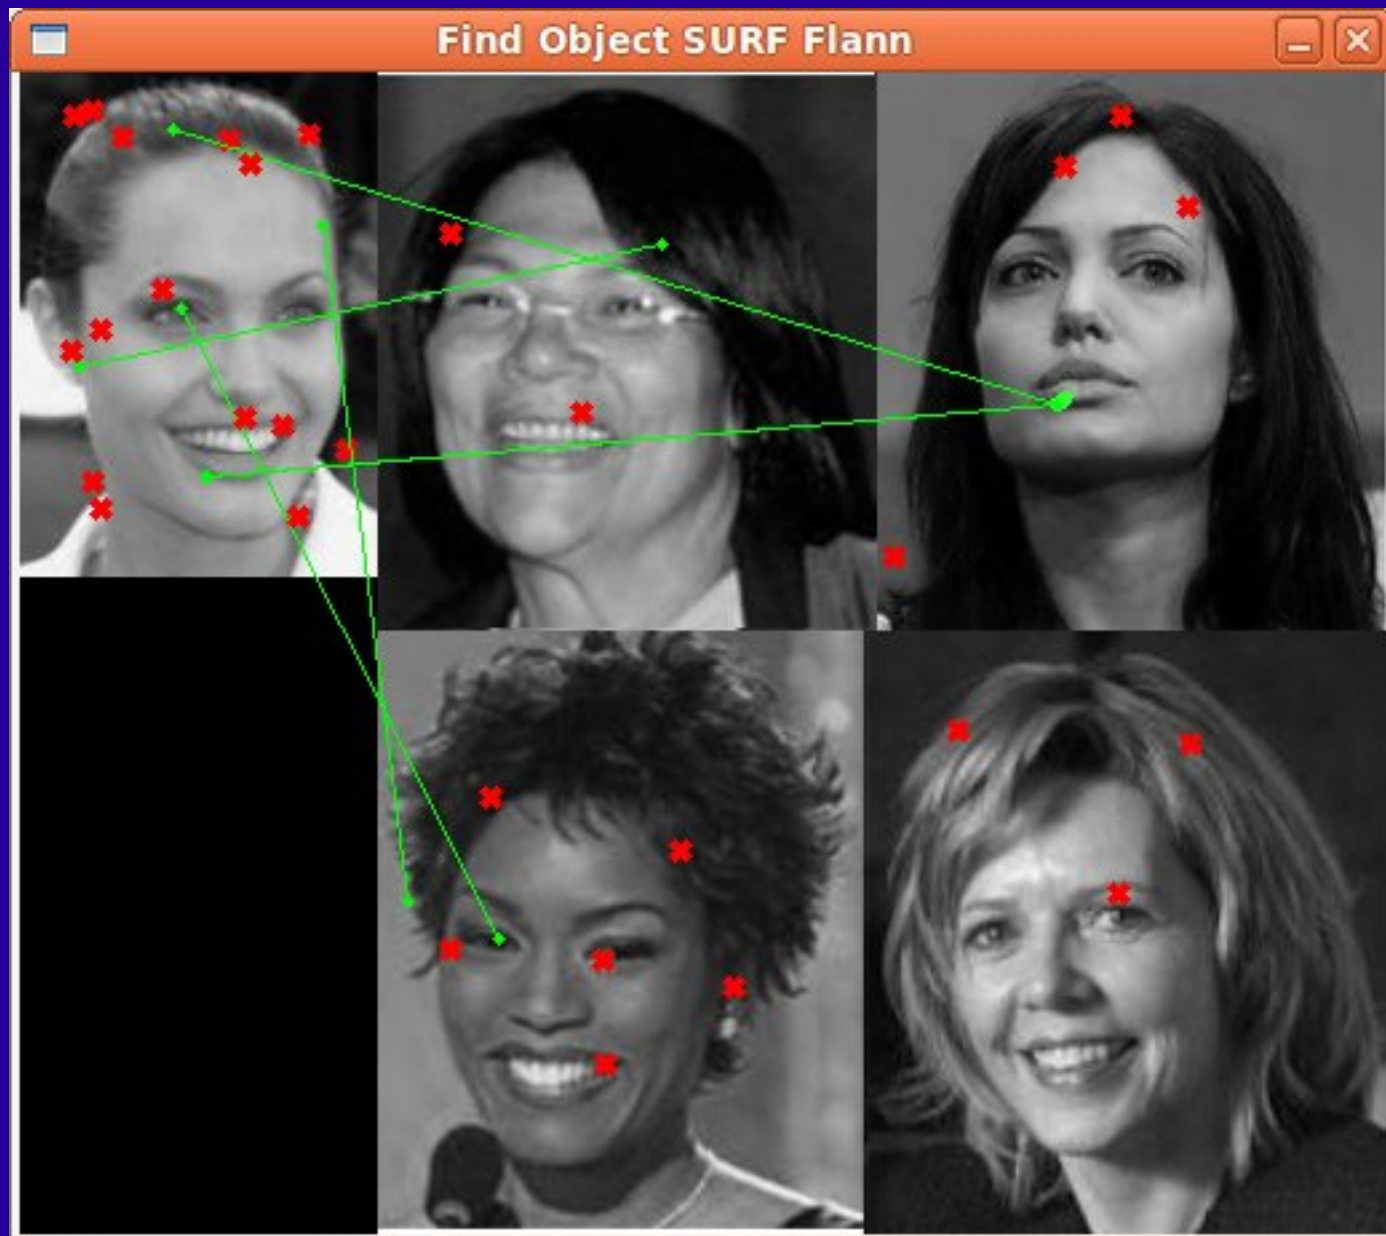


# SURF Matching on Video






# Keypoint Matching Limitations



# Machine Learning & Classification



# Building a Classifier

- Acquire training samples; e.g. Google Images, research databases (AT&T, Yale), recorded or live video
  - Single class: positive and negative samples
  - Multi-class: samples for each class
- Choose a model (e.g. SVM vs ANN)
- Train the classifier on a subset of the samples
- Test the classifier (cross-validate) on the remaining samples
- Apply classifier to new samples

# Machine Learning Software

- OpenCV (<http://opencv.willowgarage.com>)
  - Linux, Windows, MacOS, Android; C, C++, Python
  - kNN, SVM, Random Trees, Neural Networks, PCA, Normal Bayes, Boosting
- Orange (<http://orange.biolab.si>)
  - Linux, Windows, MacOS X; Python
  - Visual Programming, Cross-Validation, Learner Comparison, Feature Selection, Ensembles
  - (No Neural Networks)

# Learning Software Cont'd...


- Scikit-Learn (<http://www.scikit-learn.org>)
  - Integrated with numpy, scipy, matplotlib
  - Linux, Windows, MacOS X; Python
  - Stochastic Gradient Descent, Gaussian Mixture Models, (No Neural Networks)
- PyBrain (<http://www.pybrain.org>)
  - Linux; Python
  - Reinforcement Learning, Sequential Learning, Q-Learning, Recurrent Neural Networks, Belief Networks

# Most Popular Classifiers



- K Nearest Neighbors (kNN)
- Support Vector Machines (SVM)
- Decision Trees (DT)
- Random Forests (Ensemble of DTs)
- Cascades (e.g. Haar face detector)
- Artificial Neural Networks (ANN)
- Adaptive Boosting (AdaBoost)



# k-Nearest Neighbors (kNN) (no training required!)



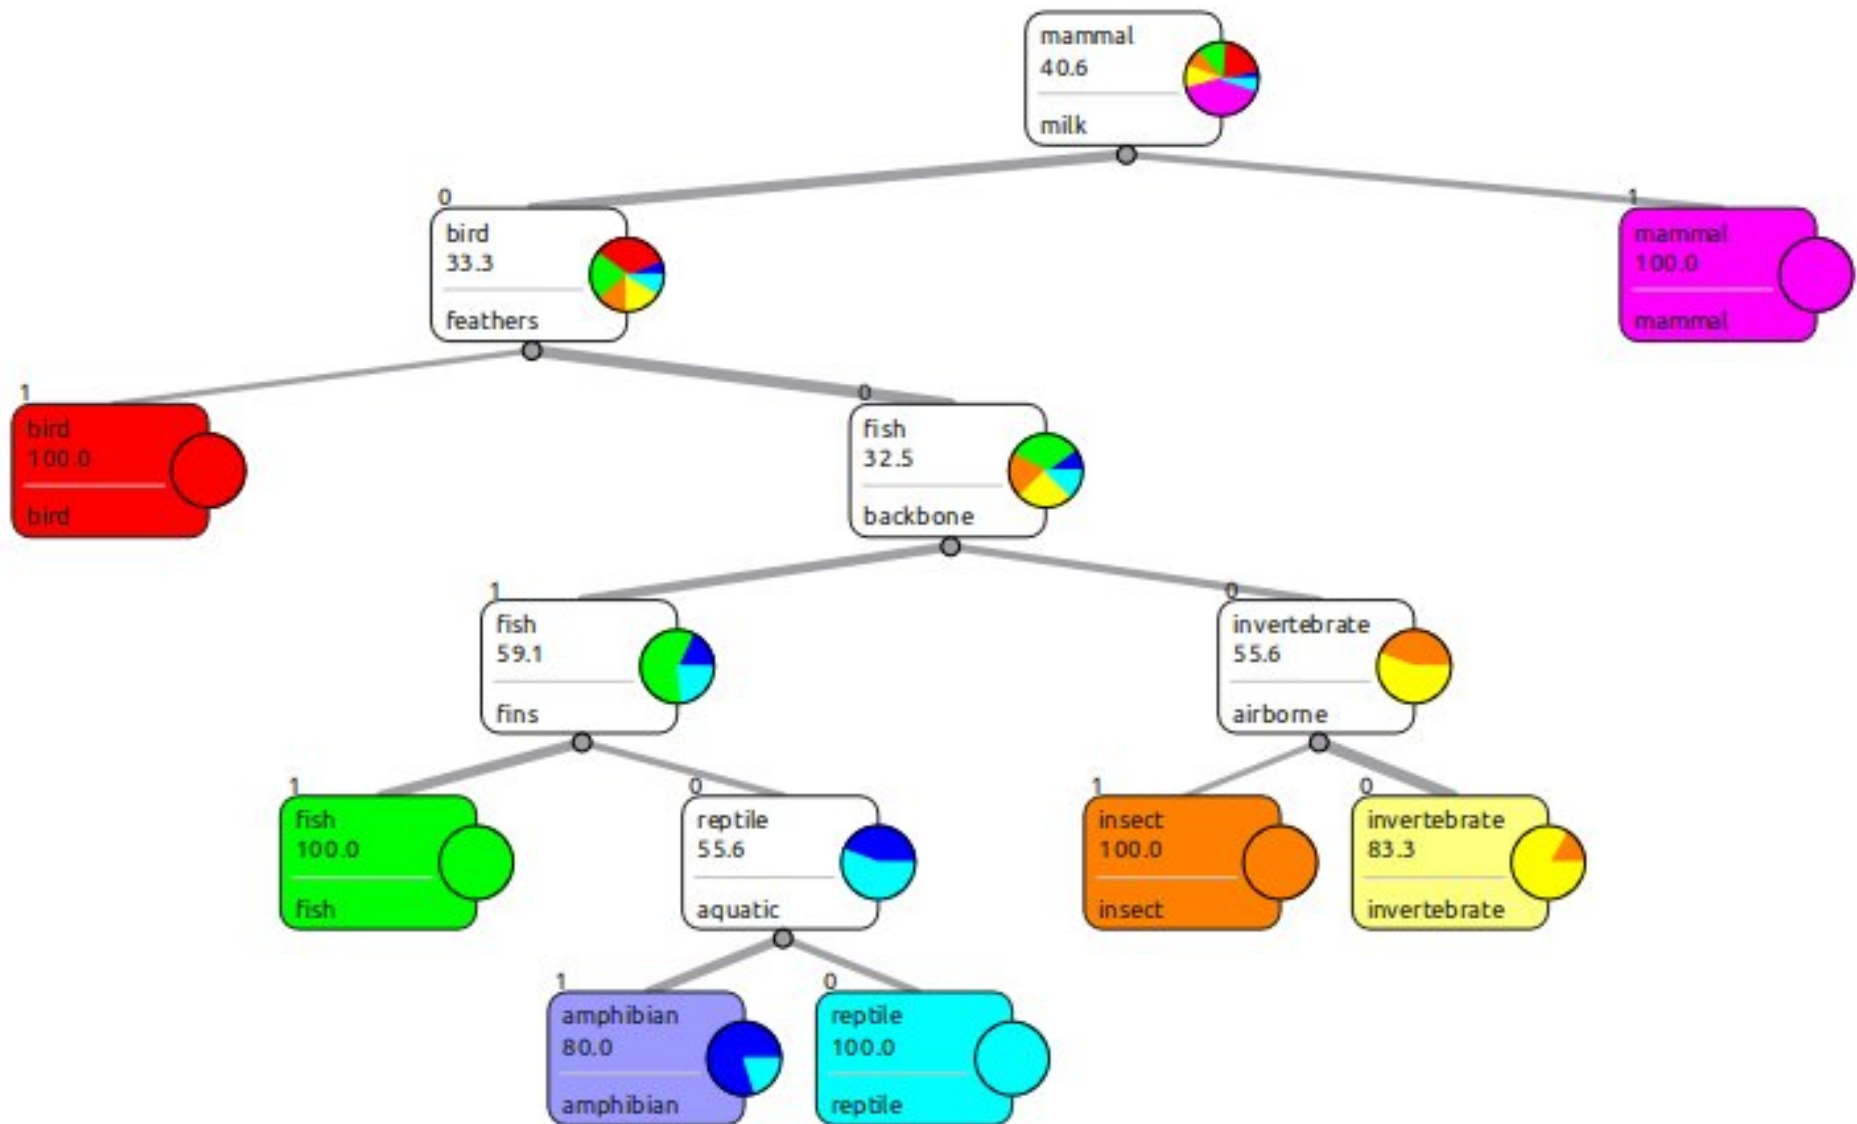
# Support Vector Machines (SVM)



# Decision Tree Data


	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q
1	Name	hair	feathers	eggs	milk	airborne	aquatic	predator	toothed	backbone	breathes	venomous	fins	tail	domestic	catsize	type
2	aardvark	1	0	0	1	0	0	1	1	1	1	0	0	0	0	1	mammal
3	antelope	1	0	0	1	0	0	0	1	1	1	0	0	1	0	1	mammal
4	bass	0	0	1	0	0	1	1	1	1	0	0	1	1	0	0	fish
5	bear	1	0	0	1	0	0	1	1	1	1	0	0	0	0	1	mammal
6	boar	1	0	0	1	0	0	1	1	1	1	0	0	1	0	1	mammal
7	buffalo	1	0	0	1	0	0	0	1	1	1	0	0	1	0	1	mammal
8	calf	1	0	0	1	0	0	0	1	1	1	0	0	1	1	1	mammal
9	carp	0	0	1	0	0	1	0	1	1	0	0	1	1	1	0	fish
10	catfish	0	0	1	0	0	1	1	1	1	0	0	1	1	0	0	fish
11	cavy	1	0	0	1	0	0	0	1	1	1	0	0	0	1	0	mammal
12	cheetah	1	0	0	1	0	0	1	1	1	1	0	0	1	0	1	mammal
13	chicken	0	1	1	0	1	0	0	0	1	1	0	0	1	1	0	bird
14	chub	0	0	1	0	0	1	1	1	1	0	0	1	1	0	0	fish
15	clam	0	0	1	0	0	0	1	0	0	0	0	0	0	0	0	invertebrate
16	crab	0	0	1	0	0	1	1	0	0	0	0	0	0	0	0	invertebrate
17	crayfish	0	0	1	0	0	1	1	0	0	0	0	0	0	0	0	invertebrate
18	crow	0	1	1	0	1	0	1	0	1	1	0	0	1	0	0	bird
19	deer	1	0	0	1	0	0	0	1	1	1	0	0	1	0	1	mammal
20	dogfish	0	0	1	0	0	1	1	1	1	0	0	1	1	0	1	fish
21	dolphin	0	0	0	1	0	1	1	1	1	1	0	1	1	0	1	mammal
22	dove	0	1	1	0	1	0	0	0	1	1	0	0	1	1	0	bird
23	duck	0	1	1	0	1	1	0	0	1	1	0	0	1	0	0	bird
24	elephant	1	0	0	1	0	0	0	1	1	1	0	0	1	0	1	mammal

# Decision Tree Classifier






# Random Forests



# Artificial Neural Networks (ANN)




# Prebuilt Classifiers in OpenCV

## Face & People Detectors




- Haar Face Detector



- HOG Person Detector



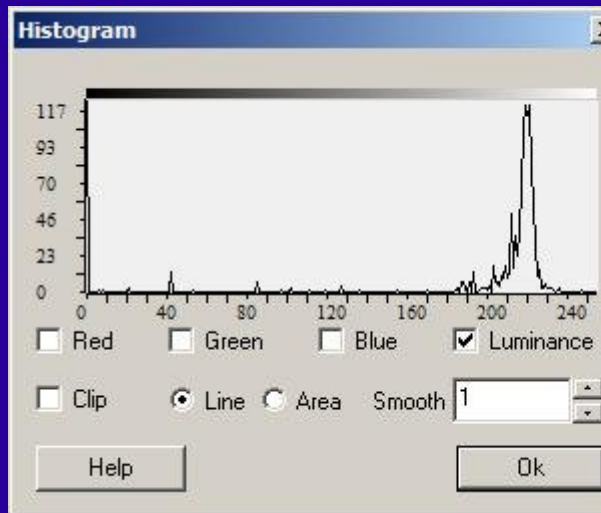
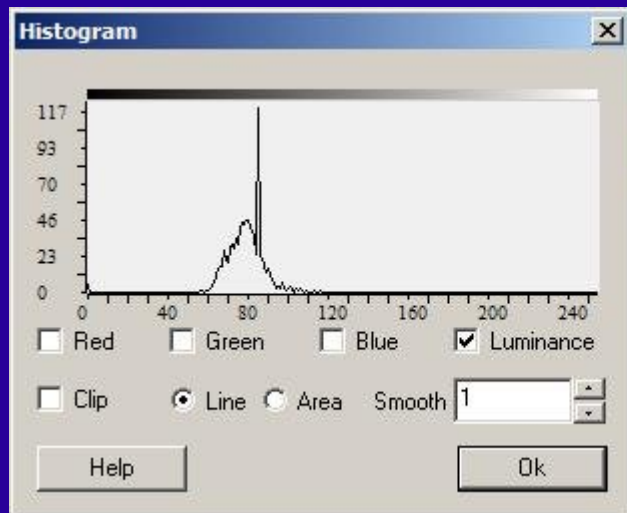
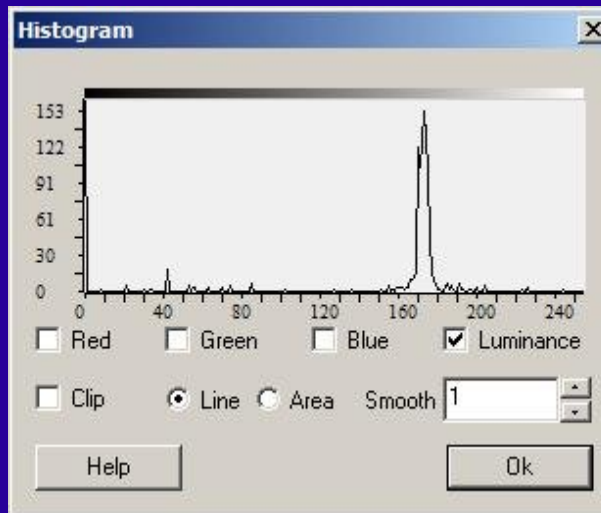
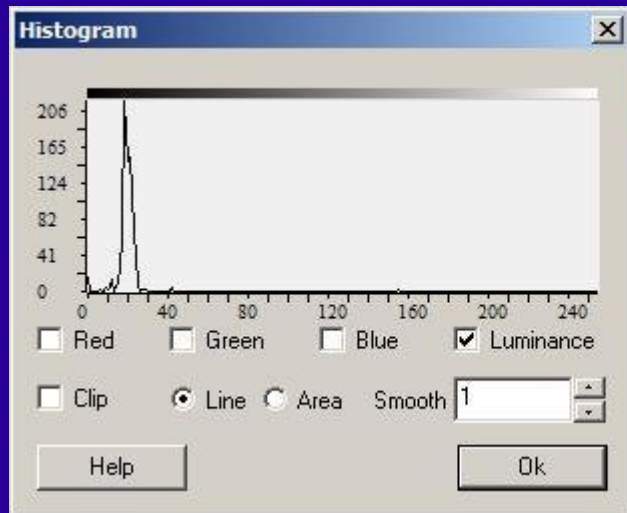
# Automated Color Naming





# Hue Histograms


## RGB $\rightarrow$ HSV $\rightarrow$ Hue



# Preparing the Color Data

- Collect color images, say 20 images per color
- Store images in folders named after each color
- Run training script:
  - For each color folder
    - Read image
    - Convert to HSV
    - Compute Hue histogram
    - Add color label and histogram values to tab-delimited data file

# Test Multiple Learners with Orange



### Test Learners

**Sampling**

- ☒ Cross-validation
- ☐ Leave-one-out
- ☐ Random sampling

Number of folds: 5

Repeat train/test: 10

Relative training set size: 70%

**Evaluation Results**

	Method	CA
1	kNN	0.8732
2	SVM	0.9206
3	Random Forest	0.8175
4	Classification Tree	0.7225
5	Naive Bayes	0.8412






# The AT&T / ORL Face Database




# Using Eigenfaces for Face Recognition

- Convert each  $W \times H$  image to a  $(W \times H)$ -d vector by concatenating the rows; e.g.  $100 \times 100$  image  $\rightarrow$  10,000-d vector
- Run a Principle Components Analysis (PCA) on all training vectors  $\rightarrow$  returns eigenvectors
- Keep  $N$  eigenvectors with largest eigenvalues; e.g.  $N = 64$
- Project sample images onto this feature space: converts each 10,000 element vector to a 64 element vector
- Build a classifier using these feature vectors

# Eigenfaces




# Average and Reconstructed Faces using PCA (64 eigenvectors)





# Eigenface Performance on AT&T faces

40 people, 10 images per person  
(ANN = 94% correct)

 **Test Learners**

### Sampling

☒ Cross-validation

Number of folds:

☐ Leave-one-out

☐ Random sampling

Repeat train/test:

Relative training set size:

70%


### Evaluation Results

	Method	CA
1	kNN	0.9525
2	SVM	0.9700
3	Random Forest	0.6350
4	Classification Tree	0.5925
5	Naive Bayes	0.9275

# The U Sheffield Faces

20 people, 20-40 images per person





## Test Learners

### Sampling

☒ Cross-validation

Number of folds:

☐ Leave-one-out

☐ Random sampling



Repeat train/test:

Relative training set size:


### Evaluation Results

	Method	CA
1	kNN	0.9878
2	SVM	0.9913
3	Random Forest	0.8557
4	Classification Tree	0.7826
5	Naive Bayes	0.9496

# Labeled Faces in the Wild



# Eigenface Failure...

 **Test Learners**

### Sampling

☒ Cross-validation

Number of folds:

☐ Leave-one-out

☐ Random sampling

Repeat train/test:

Relative training set size:

### Evaluation Results

	Method	CA
1	kNN	0.5000
2	SVM	0.6000
3	Random Forest	0.6000
4	Classification Tree	0.5500
5	Naive Bayes	0.3000



# Alternatives to Eigenfaces

Your Contribution Goes Here :-)

- PCL: Viewpoint Feature Histograms (VFH)
- SIFT/SURF + PCA
- Biometrics
- 3D Modeling

# Real Time Classifier Learning

- Select an object any way you can
  - manual selection
  - existing detector (e.g. Haar face detector)
  - motion detection
- Track keypoints using Optical Flow
- Grab positive and negative samples as you go
- Build a custom classifier from current samples
- Re-detect object using classifier
- Store custom classifier for later detection

# Face Tracking using Optical Flow

Video courtesy of Honda/UCSD Video Database

Pi Face Tracker Video

[http://youtu.be/Yw\\_zkLaZNsQ](http://youtu.be/Yw_zkLaZNsQ)

# Predator Algorithm (OpenTLD)

## Zdenek Kalal, Jan 2011

Predator Track-Learn-Detect (TLD)

<http://youtu.be/1GhNXHCQGSM>



# Thanks for Listening!

